

Type-Driven Development with Idris

Edwin Brady



MANNING

The Idris read-eval-print loop (REPL) provides several commands. The most common commands, listed below, are introduced over the course of this book.

Command	Arguments	Description
<expression>	None	Displays the result of evaluating the expression. The variable <code>it</code> contains the result of the most recent evaluation.
:t	<expression>	Displays the type of the expression.
:total	<name>	Displays whether the function with the given name is total.
:doc	<name>	Displays documentation for <code>name</code> .
:let	<definition>	Adds a new definition.
:exec	<expression>	Compiles and executes the expression. If none is given, compiles and executes <code>main</code> .
:c	<output file>	Compiles to an executable with the entry point <code>main</code> .
:r	None	Reloads the current module.
:l	<filename>	Loads a new file.
:module	<module name>	Imports an extra module for use at the REPL.
:printdef	<name>	Displays the definition of <code>name</code> .
:apropos	<word>	Searches function names, types, and documentation for the given word.
:search	<type>	Searches for functions with the given type.
:browse	<namespace>	Displays the names and types defined in the given namespace.
:q	None	Exits the REPL.

Type-Driven Development with Idris

Type-Driven Development with Idris

EDWIN BRADY



MANNING
SHELTER ISLAND

For online information and ordering of this and other Manning books, please visit www.manning.com. The publisher offers discounts on this book when ordered in quantity. For more information, please contact

Special Sales Department
Manning Publications Co.
20 Baldwin Road
PO Box 761
Shelter Island, NY 11964
Email: orders@manning.com

©2017 by Manning Publications Co. All rights reserved.

No part of this publication may be reproduced, stored in a retrieval system, or transmitted, in any form or by means electronic, mechanical, photocopying, or otherwise, without prior written permission of the publisher.

Many of the designations used by manufacturers and sellers to distinguish their products are claimed as trademarks. Where those designations appear in the book, and Manning Publications was aware of a trademark claim, the designations have been printed in initial caps or all caps.

- ⊗ Recognizing the importance of preserving what has been written, it is Manning's policy to have the books we publish printed on acid-free paper, and we exert our best efforts to that end. Recognizing also our responsibility to conserve the resources of our planet, Manning books are printed on paper that is at least 15 percent recycled and processed without the use of elemental chlorine.

 Manning Publications Co.
20 Baldwin Road
PO Box 761
Shelter Island, NY 11964

Development editor: Dan Maharry
Review editor: Aleksandar Dragosavljević
Technical development editor: Andrew Gibson
Project editor: Kevin Sullivan
Copyeditor: Andy Carroll
Proofreader: Katie Tennant
Technical proofreaders: Arnaud Bailly, Nicolas Biri
Typesetter: Dottie Marsico
Cover designer: Marija Tudor

ISBN 9781617293023
Printed in the United States of America
1 2 3 4 5 6 7 8 9 10 – EBM – 22 21 20 19 18 17

brief contents

PART 1	INTRODUCTION	1
1	■ Overview	3
2	■ Getting started with Idris	25
PART 2	CORE IDRIS	53
3	■ Interactive development with types	55
4	■ User-defined data types	87
5	■ Interactive programs: input and output processing	123
6	■ Programming with first-class types	147
7	■ Interfaces: using constrained generic types	182
8	■ Equality: expressing relationships between data	208
9	■ Predicates: expressing assumptions and contracts in types	236
10	■ Views: extending pattern matching	258
PART 3	IDRIS AND THE REAL WORLD	289
11	■ Streams and processes: working with infinite data	291
12	■ Writing programs with state	324
13	■ State machines: verifying protocols in types	352
14	■ Dependent state machines: handling feedback and errors	373
15	■ Type-safe concurrent programming	403

contents

<i>preface</i>	xv
<i>acknowledgments</i>	xvii
<i>about this book</i>	xix
<i>about the author</i>	xxiii
<i>about the cover illustration</i>	xxiv

PART 1 INTRODUCTION.....1

1 Overview 3

1.1 What is a type?	4
1.2 Introducing type-driven development	5
<i>Matrix arithmetic</i>	6
<i>An automated teller machine</i>	7
<i>Concurrent programming</i>	9
<i>Type, define, refine: the process of type-driven development</i>	10
<i>Dependent types</i>	11
1.3 Pure functional programming	13
<i>Purity and referential transparency</i>	13
<i>Side-effecting programs</i>	14
<i>Partial and total functions</i>	16
1.4 A quick tour of Idris	17
<i>The interactive environment</i>	17
<i>Checking types</i>	18
<i>Compiling and running Idris programs</i>	19
<i>Incomplete definitions: working with holes</i>	20
<i>First-class types</i>	22
1.5 Summary	24

2	Getting started with Idris	25
2.1	Basic types	26
	<i>Numeric types and values</i>	27
	<i>Type conversions using cast</i>	28
	<i>Characters and strings</i>	29
	<i>Booleans</i>	30
2.2	Functions: the building blocks of Idris programs	30
	<i>Function types and definitions</i>	31
	<i>Partially applying functions</i>	33
	<i>Writing generic functions: variables in types</i>	33
	<i>Writing generic functions with constrained types</i>	35
	<i>Higher-order function types</i>	36
	<i>Anonymous functions</i>	38
	<i>Local definitions: let and where</i>	39
2.3	Composite types	40
	<i>Tuples</i>	40
	<i>Lists</i>	41
	<i>Functions with lists</i>	43
2.4	A complete Idris program	46
	<i>Whitespace significance: the layout rule</i>	46
	<i>Documentation comments</i>	47
	<i>Interactive programs</i>	48
2.5	Summary	52
PART 2 CORE IDRIS		53
3	Interactive development with types	55
3.1	Interactive editing in Atom	56
	<i>Interactive command summary</i>	57
	<i>Defining functions by pattern matching</i>	57
	<i>Data types and patterns</i>	61
3.2	Adding precision to types: working with vectors	64
	<i>Refining the type of allLengths</i>	65
	<i>Type-directed search: automatic refining</i>	69
	<i>Type, define, refine: sorting a vector</i>	70
3.3	Example: type-driven development of matrix functions	75
	<i>Matrix operations and their types</i>	76
	<i>Transposing a matrix</i>	77
3.4	Implicit arguments: type-level variables	82
	<i>The need for implicit arguments</i>	82
	<i>Bound and unbound implicits</i>	83
	<i>Using implicit arguments in functions</i>	84
3.5	Summary	86
4	User-defined data types	87
4.1	Defining data types	88
	<i>Enumerations</i>	89
	<i>Union types</i>	90
	<i>Recursive types</i>	92
	<i>Generic data types</i>	95

4.2	Defining dependent data types	102
	<i>A first example: classifying vehicles by power source</i>	102
	<i>Defining vectors</i>	104
	<i>Indexing vectors with bounded numbers using Fin</i>	107
4.3	Type-driven implementation of an interactive data store	110
	<i>Representing the store</i>	112
	<i>Interactively maintaining state in main</i>	113
	<i>Commands: parsing user input</i>	115
	<i>Processing commands</i>	118
4.4	Summary	122

5 Interactive programs: input and output processing 123

5.1	Interactive programming with IO	124
	<i>Evaluating and executing interactive programs</i>	125
	<i>Actions and sequencing: the >= operator</i>	127
	<i>Syntactic sugar for sequencing with do notation</i>	129
5.2	Interactive programs and control flow	132
	<i>Producing pure values in interactive definitions</i>	132
	<i>Pattern-matching bindings</i>	134
	<i>Writing interactive definitions with loops</i>	136
5.3	Reading and validating dependent types	138
	<i>Reading a Vect from the console</i>	139
	<i>Reading a Vect of unknown length</i>	140
	<i>Dependent pairs</i>	141
	<i>Validating Vect lengths</i>	143
5.4	Summary	146

6 Programming with first-class types 147

6.1	Type-level functions: calculating types	148
	<i>Type synonyms: giving informative names to complex types</i>	149
	<i>Type-level functions with pattern matching</i>	150
	<i>Using case expressions in types</i>	153
6.2	Defining functions with variable numbers of arguments	155
	<i>An addition function</i>	155
	<i>Formatted output: a type-safe printf function</i>	157
6.3	Enhancing the interactive data store with schemas	161
	<i>Refining the DataStore type</i>	162
	<i>Using a record for the DataStore</i>	164
	<i>Correcting compilation errors using holes</i>	165
	<i>Displaying entries in the store</i>	170
	<i>Parsing entries according to</i>	

the schema 171 • *Updating the schema* 175 • *Sequencing expressions with Maybe using do notation* 177

6.4 Summary 181

7 Interfaces: using constrained generic types 182

7.1 Generic comparisons with Eq and Ord 183

Testing for equality with Eq 183 • *Defining the Eq constraint using interfaces and implementations* 185
Default method definitions 189 • *Constrained implementations* 189 • *Constrained interfaces: defining orderings with Ord* 191

7.2 Interfaces defined in the Prelude 194

Converting to String with Show 194 • *Defining numeric types* 195 • *Converting between types with Cast* 198

7.3 Interfaces parameterized by Type -> Type 199

Applying a function across a structure with Functor 200
Reducing a structure using Foldable 201 • *Generic do notation using Monad and Applicative* 205

7.4 Summary 207

8 Equality: expressing relationships between data 208

8.1 Guaranteeing equivalence of data with equality types 209

Implementing exactLength, first attempt 210 • *Expressing equality of Nats as a type* 211 • *Testing for equality of Nats* 212 • *Functions as proofs: manipulating equalities* 215
Implementing exactLength, second attempt 216 • *Equality in general: the = type* 218

8.2 Equality in practice: types and reasoning 220

Reversing a vector 220 • *Type checking and evaluation* 221
The rewrite construct: rewriting a type using equality 223
Delegating proofs and rewriting to holes 224 • *Appending vectors, revisited* 225

8.3 The empty type and decidability 227

Void: a type with no values 228 • *Decidability: checking properties with precision* 229 • *DecEq: an interface for decidable equality* 233

8.4 Summary 234

9 Predicates: expressing assumptions and contracts in types 236

- 9.1 Membership tests: the `Elem` predicate 237
Removing an element from a Vect 238 • *The `Elem` type: guaranteeing a value is in a vector* 239 • *Removing an element from a Vect: types as contracts* 241 • *auto-implicit arguments: automatically constructing proofs* 244 • *Decidable predicates: deciding membership of a vector* 245
- 9.2 Expressing program state in types: a guessing game 250
Representing the game's state 250 • *A top-level game function* 251 • *A predicate for validating user input: `ValidInput`* 251 • *Processing a guess* 253 • *Deciding input validity: checking `ValidInput`* 255 • *Completing the top-level game implementation* 255
- 9.3 Summary 257

10 Views: extending pattern matching 258

- 10.1 Defining and using views 259
Matching the last item in a list 260 • *Building views: covering functions* 262 • *with blocks: syntax for extended pattern matching* 262 • *Example: reversing a list using a view* 264 • *Example: merge sort* 266
- 10.2 Recursive views: termination and efficiency 271
“Snoc” lists: traversing a list in reverse 271 • *Recursive views and the `with` construct* 274 • *Traversing multiple arguments: nested with blocks* 275 • *More traversals: `Data.List.Views`* 277
- 10.3 Data abstraction: hiding the structure of data using views 280
Digression: modules in Idris 280 • *The data store, revisited* 282 • *Traversing the store’s contents with a view* 284
- 10.4 Summary 288

PART 3 IDRIS AND THE REAL WORLD 289

11 Streams and processes: working with infinite data 291

- 11.1 Streams: generating and processing infinite lists 292
Labeling elements in a List 293 • *Producing an infinite list of numbers* 295 • *Digression: what does it mean for a function to*

<i>be total?</i>	296	▪ <i>Processing infinite lists</i>	297	▪ <i>The Stream data type</i>	299	▪ <i>An arithmetic quiz using streams of random numbers</i>	301			
11.2	Infinite processes: writing interactive total programs	305								
	<i>Describing infinite processes</i>	306	▪ <i>Executing infinite processes</i>	307	▪ <i>Executing infinite processes as total functions</i>	308	▪ <i>Generating infinite structures using Lazy types</i>	309	▪ <i>Extending do notation for InfIO</i>	311
	<i>A total arithmetic quiz</i>	311								
11.3	Interactive programs with termination	314								
	<i>Refining InfIO: introducing termination</i>	314	▪ <i>Domain-specific commands</i>	317	▪ <i>Sequencing Commands with do notation</i>	320				
11.4	Summary	323								

12 Writing programs with state 324

12.1	Working with mutable state	325							
	<i>The tree-traversal example</i>	326	▪ <i>Representing mutable state using a pair</i>	328	▪ <i>State, a type for describing stateful operations</i>	329	▪ <i>Tree traversal with State</i>	331	
12.2	A custom implementation of State	333							
	<i>Defining State and runState</i>	333	▪ <i>Defining Functor, Applicative, and Monad implementations for State</i>	335					
12.3	A complete program with state: working with records	340							
	<i>Interactive programs with state: the arithmetic quiz revisited</i>	340	▪ <i>Complex state: defining nested records</i>	343					
	<i>Updating record field values</i>	344	▪ <i>Updating record fields by applying functions</i>	346	▪ <i>Implementing the quiz</i>	346			
	<i>Running interactive and stateful programs: executing the quiz</i>	348							
12.4	Summary	351							

13 State machines: verifying protocols in types 352

13.1	State machines: tracking state in types	353							
	<i>Finite state machines: modeling a door as a type</i>	354							
	<i>Interactive development of sequences of door operations</i>	356							
	<i>Infinite states: modeling a vending machine</i>	358							
	<i>A verified vending machine description</i>	360							

13.2 Dependent types in state: implementing a stack 363

Representing stack operations in a state machine 364

Implementing the stack using Vect 366 • *Using a stack interactively: a stack-based calculator* 367

13.3 Summary 371

14 *Dependent state machines: handling feedback and errors* 373

14.1 Dealing with errors in state transitions 374

Refining the door model: representing failure 375 • *A verified, error-checking, door-protocol description* 378

14.2 Security properties in types: modeling an ATM 382

Defining states for the ATM 383 • *Defining a type for the ATM* 384 • *Simulating an ATM at the console: executing ATMCmd* 387 • *Refining preconditions using auto-implicits* 388

14.3 A verified guessing game: describing rules in types 390

Defining an abstract game state and operations 391 • *Defining a type for the game state* 392 • *Implementing the game* 395
Defining a concrete game state 397 • *Running the game: executing GameLoop* 399

14.4 Summary 402

15 *Type-safe concurrent programming* 403

15.1 Primitives for concurrent programming in Idris 404

Defining concurrent processes 406 • *The Channels library: primitive message passing* 407 • *Problems with channels: type errors and blocking* 410

15.2 Defining a type for safe message passing 411

Describing message-passing processes in a type 412 • *Making processes total using Inf* 415 • *Guaranteeing responses using a state machine and Inf* 418 • *Generic message-passing processes* 422 • *Defining a module for Process* 426 • *Example 1: List processing* 427 • *Example 2: A word-counting process* 429

15.3 Summary 433

appendix A *Installing Idris and editor modes* 435

appendix B *Interactive editing commands* 438

appendix C *REPL commands* 439

appendix D *Further reading* 441

index 445

preface

Computers are everywhere, and we rely on software daily. As well as running our desktop and laptop computers, software controls our communications, banking, transport infrastructure, and even our domestic appliances. Even so, it's considered a fact of life that software is unreliable. If a laptop or mobile phone fails, it's merely inconvenient and requires a restart (possibly accompanied by cursing over losing the last few minutes of work). If, on the other hand, the software controlling a business-critical application or server fails, significant time and money can be lost. For safety-critical systems, the repercussions could be even worse.

For many years, therefore, computer science researchers have been searching for ways to improve the robustness and safety of software. One approach among many is to use *types* to describe what a program is supposed to do. In particular, by using *dependent types*, you describe precise properties of a program. The idea is that if you can express a program's intention in its type, and the program successfully type-checks, then the program must behave as intended. An important (if ambitious and long-term) goal of the Idris programming language is to make the results of this research accessible to software developers in general, and correspondingly reduce the possibility of critical software failures.

Initially, this book's focus was programming in Idris: showing how to use its type system to guarantee important properties of programs. Under the guidance of development editor Dan Maharry, and thanks to the efforts of technical development editor Andrew Gibson, it has evolved to become as much about the process of programming with dependent types as about how the resulting programs work. You'll learn about the fundamentals of dependent types, how to use types to define programs interactively, and how to refine programs and types as your understanding of a

problem evolves. You'll also learn about some real-world applications of type-driven development, particularly in dealing with state, protocols, and concurrency.

Idris itself arose as a result of my own research into program verification and language design with dependent types. After spending several years immersed in the concept of programming with dependent types, I felt there was a need for a language designed for developers as well as researchers. I hope that you have as much fun learning about type-driven development with Idris as I have had developing it!

acknowledgments

Many people have helped in the writing of this book, and it wouldn't exist without them. In particular, I thank Dan Maharry, who encouraged me to reveal the ideas of type-driven development much more clearly. The mantra of "*type, define, refine*," which appears throughout the book, was Dan's suggestion. I also owe many thanks to Andrew Gibson, who has meticulously worked through all the examples and exercises throughout the book, making sure they work, checking that the exercises are solvable, and suggesting many improvements to the text and explanations. Overall, I'd like to thank the team at Manning Publications for helping to make this book a reality.

The design of Idris owes much to several decades of research into type theory, functional programming, and language design. I'm grateful to James McKinna and Conor McBride in particular for teaching me the fundamentals of type theory when I was a graduate student at Durham University, and for their continued advice and encouragement since. I'd also like to thank the researchers and developers responsible for the languages and systems that have inspired my work, namely, tools such as Haskell, Epigram, Agda, and Coq. Idris couldn't exist without the work that has come before, and I can only hope that it, in turn, inspires others in the future. See appendix D for some references to the work that inspired Idris.

Several colleagues and students at the University of St. Andrews and elsewhere have provided useful feedback on earlier drafts of chapters and have been patient while I worked on the book instead of working on other things. In particular, I would like to thank Ozgur Akgun, Nicola Botta, Sam Elliot, Simon Fowler, Nicolas Gaglani (who contributed the extension to the Atom editor that you'll use throughout the book), Jan de Muijnck-Hughes, Markus Pfeiffer, Chris Schwaab, and Matúš Tejiščák

for their comments and suggestions. My sincere apologies to anyone else I've forgotten to name!

Readers who have purchased early access and reviewers of earlier drafts have contributed many useful comments and suggestions. Those reviewers include Alexander A. Myltsev, Álvaro Falquina, Arnaud Bailly, Carsten Jørgensen, Christine Koppelt, Giovanni Ruggiero, Ian Dees, Juan Gabriel Bono, Mattias Lundell, Phil de Joux, Rintcius Blok, Satadru Roy, Sergey Selyugin, Todd Fine, and Vitaly Bragilevsky.

I couldn't have implemented Idris on my own. Since I began developing the current version in late 2011, there have been many contributors, but most of all I would like to thank David Christiansen, who is responsible for much of the polish in the Idris REPL and the interactive editing tools; he has also worked hard to help newcomers to the project. Thanks are also due to the other contributors: Ozgur Akgun, Ahmad Salim Al-Sibahi, Edward Chadwick Amsden, Michael R. Bernstein, Jan Bessai, Nicola Botta, Vitaly Bragilevsky, Jakob Brünker, Alyssa Carter, Carter Charbonneau, Aaron Craelius, Jason Dagit, Adam Sandberg Eriksson, Guglielmo Fachini, Simon Fowler, Zack Grannan, Sean Hunt, Cezar Ionescu, Heath Johns, Irene Knapp, Paul Koerbitz, Niklas Larsson, Shea Levy, Mathnerd314, Hannes Mehnert, Mekeor Melire, Melissa Mozifian, Jan de Muijnck-Hughes, Dominic Mulligan, Echo Nolan, Tom Prince, raichoo, Philip Rasmussen, Aistis Raulinaitis, Reynir Reynisson, Seo Sanghyeon, Benjamin Saunders, Alexander Shabalin, Jeremy W. Sherman, Timo Petteri Sinnemäki, JP Smith, startling, Chetan T, Matúš Tejiščák, Dirk Ullrich, Leif Warner, Daniel Waterworth, Eric Weinstein, Jonas Westerlund, Björn Aili, and Zheng Jihui.

Finally, thanks go to my parents, whose purchase of a BBC Micro in 1983 set me off on this path; and to Emma, for waiting so patiently for me to finish this, and for bringing me coffee to keep me going.

about this book

Type-Driven Development with Idris is about making types work for you. Types are often seen as a tool for checking for errors, with the programmer writing a complete program first and using the type checker to detect errors. In type-driven development, you use types as a tool for constructing programs, and the type checker as your assistant to guide you to a complete and working program.

This book begins by describing what you can express with types; then, it introduces the core features of the Idris programming language. Finally, it describes some more-practical applications of type-driven development.

Who should read this book

This book is aimed at developers who want to learn about the state of the art in using sophisticated type systems to help develop robust software. It aims to provide an accessible introduction to dependent types, and to show how modern type-based techniques can be applied to real-world problems.

Readers will ideally already be familiar with functional programming concepts such as closures and higher-order functions, although the book introduces these and other concepts as necessary. Knowledge of another functional programming language such as Haskell, OCaml, or Scala will be particularly helpful, though none is assumed.

Roadmap

This book is divided into three parts. Part 1 (chapters 1 and 2) introduces the concepts and gives a tour of the Idris programming language:

- Chapter 1 introduces type-driven development and gets you started with the Idris environment.

- Chapter 2 covers the basics of Idris programming, including primitive types and structuring Idris programs.

Part 2 (chapters 3–10) introduces the core language features of Idris:

- Chapter 3 discusses interactive development using the Atom editor and describes how using more-precise types means that the type checker can help you write programs.
- Chapter 4 explains how to define your own data types and presents a first example of writing a larger interactive program in the type-driven style.
- Chapter 5 describes interactive programs in more depth, including how to use types to help validate user inputs to interactive programs.
- Chapter 6 introduces type-level programming, showing how to write functions that calculate types and how to use them in practice.
- Chapter 7 describes how to use interfaces to write programs with generic types.
- Chapter 8 explains how you can use types to express relationships between data, particularly to describe properties of data and to guarantee that functions behave a certain way.
- Chapter 9 explains further how types can express contracts that functions must satisfy, including an example that shows how to use types to describe the state of a system.
- Chapter 10 introduces views, which are alternative ways of inspecting and traversing data structures.

Part 3 (chapters 11–15) describes some applications of Idris to real-world software development, particularly working with state and interactive programs:

- Chapter 11 describes how to work with potentially infinite data, such as streams, and how to write and reason about interactive programs that could run indefinitely.
- Chapter 12 explains how to write programs with state and how to represent and manipulate complex state using records.
- Chapter 13 shows how to express a state machine in an Idris type, and how to use the type to guarantee that programs follow protocols correctly.
- Chapter 14 describes more-sophisticated state machines, how to handle errors and feedback from an environment, and how to represent security properties of a system in its type.
- Chapter 15 concludes the book with a worked example: type-driven development of a small library for concurrent programming.

In general, each chapter builds on concepts introduced in earlier chapters, so it's intended that you read the chapters in order. Most importantly, the book describes the *process* of type-driven development and constructing programs interactively from a type. Therefore, I strongly recommend working through the examples on a computer

as you read. Furthermore, if you’re reading the eBook, type in the examples—don’t just copy and paste.

There are exercises throughout each chapter, so, as you read, make sure that you complete the exercises to reinforce your understanding. Sample solutions are available online from the book’s website at www.manning.com/books/type-driven-development-with-idris.

There are four appendices: Appendix A describes how to install Idris and the Atom editor mode, which we’ll use throughout the book. Appendix B summarizes the interactive editing commands supported by Atom. Appendix C summarizes the commands you can use in the Idris environment. Finally, appendix D gives references to some of the work that inspired Idris, and where you can learn more about the theoretical background and related tools.

Code conventions and downloads

This book contains many examples of source code both in numbered listings and inline with normal text. In both cases, source code is formatted in a fixed-width font like this to separate it from ordinary text.

In many cases, the original source code has been reformatted; I’ve added line breaks and reworked indentation to accommodate the available page space in the book. Additionally, comments in the source code have often been removed from the listings when the code is described in the text. Code annotations accompany many of the listings, highlighting important concepts.

All of the code in this book is available online from the book’s website (www.manning.com/books/type-driven-development-with-idris) and has been tested with Idris 1.0. The code is also available in a Git repository here: <https://github.com/edwinb/TypeDD-Samples>.

Author Online

Purchase of *Type-Driven Development with Idris* includes free access to a private web forum run by Manning Publications where you can make comments about the book, ask technical questions, and receive help from the author and from other users. To access the forum and subscribe to it, point your web browser to www.manning.com/books/type-driven-development-with-idris. This page provides information on how to get on the forum once you’re registered, what kind of help is available, and the rules of conduct on the forum.

Manning’s commitment to our readers is to provide a venue where a meaningful dialog between individual readers and between readers and the author can take place. It’s not a commitment to any specific amount of participation on the part of the author, whose contribution to the forum remains voluntary (and unpaid). We suggest you try asking him challenging questions lest his interest stray!

The Author Online forum and the archives of previous discussions will be accessible from the publisher’s website as long as the book is in print.

Other online resources

If you'd like to learn more about Idris, you can find more resources on the Idris website: <http://idris-lang.org/>. You can also find help in several other places:

- The idris-lang Google Group is an active group discussing all aspects of Idris. The group welcomes questions from beginners and more-advanced users alike.
- There is an IRC channel, #idris on irc.freenode.net, which is similarly open to questions.
- You can ask and answer questions using the Idris tag on Stack Overflow.

about the author



EDWIN BRADY leads the design and implementation of the Idris programming language. He is a lecturer in Computer Science at the University of St. Andrews in Scotland, and he regularly speaks at conferences. When he's not doing that, you might find him playing a game of Go, watching a game of cricket, or somewhere up a hill in the middle of Scotland.

about the cover illustration

The figure on the cover of *Type-Driven Development with Idris* is captioned “La Gasconne,” or “A woman from Gascony.” The illustration is taken from a collection of works by many artists, edited by Louis Curmer and published in Paris in 1841. The title of the collection is *Les Français peints par eux-mêmes*, which translates as *The French People Painted by Themselves*. Each illustration is finely drawn and colored by hand, and the rich variety of drawings in the collection reminds us vividly of how culturally apart the world’s regions, towns, villages, and neighborhoods were just 200 years ago. Isolated from each other, people spoke different dialects and languages. In the streets or in the countryside, it was easy to identify where they lived and what their trade or station in life was just by their dress.

Dress codes have changed since then, and the diversity by region, so rich at the time, has faded away. It’s now hard to tell apart the inhabitants of different continents, let alone different towns or regions. Perhaps we have traded cultural diversity for a more varied personal life—certainly for a more varied and fast-paced technological life.

At a time when it’s hard to tell one computer book from another, Manning celebrates the inventiveness and initiative of the computer business with book covers based on the rich diversity of regional life of two centuries ago, brought back to life by pictures from collections such as this one.

Part 1

Introduction

I

In this first part, you'll get started with Idris and learn about the ideas behind type-driven development. I'll take you on a brief tour of the Idris environment, and you'll write some simple but complete programs.

In the first chapter, I'll explain more about what I mean by type-driven development. Most importantly, I'll define what I mean by "type" and give several examples of how you can use expressive types to describe the intended purpose of your programs more precisely. I'll also introduce the two most distinctive features of the Idris language: *holes*, which stand for parts of programs that are yet to be written, and the use of types as a *first-class* language construct.

Before you get too deeply into type-driven development in Idris, it's important to have a solid grasp of the fundamentals of the language. Therefore, in chapter 2, I'll discuss some of the primitive language constructs, many of which will be familiar to you from other languages, and show how to construct complete programs in Idris.

Overview



This chapter covers

- Introducing type-driven development
- The essence of pure functional programming
- First steps with Idris

This book teaches a new approach to building robust software, *type-driven development*, using the Idris programming language. Traditionally, types are seen as a tool for checking for errors, with the programmer writing a complete program first and using either the compiler or the runtime system to detect type errors. In type-driven development, we use types as a tool for constructing programs. We put the *type* first, treating it as a plan for a program, and use the compiler and type checker as our assistant, guiding us to a complete and working program that satisfies the type. The more expressive the type is that we give up front, the more confidence we can have that the resulting program will be correct.

TYPES AND TESTS The name “type-driven development” suggests an analogy to test-driven development. There’s a similarity, in that writing tests first helps establish a program’s purpose and whether it satisfies some basic requirements. The difference is that, unlike tests, which can usually only be used to show the *presence* of errors, types (used appropriately) can show the *absence* of errors. But although types *reduce* the need for tests, they rarely eliminate it entirely.

Idris is a relatively young programming language, designed from the beginning to support type-driven development. A prototype implementation first appeared in 2008, with development of the current implementation beginning in 2011. It builds on decades of research into the theoretical and practical foundations of programming languages and type systems.

In Idris, types are a *first-class* language construct. Types can be manipulated, used, passed as arguments to functions, and returned from functions just like any other value, such as numbers, strings, or lists. This is a simple but powerful idea:

- It allows relationships to be expressed between values; for example, that two lists have the same length.
- It allows assumptions to be made explicit and checkable by the compiler. For example, if you assume that a list is non-empty, Idris can ensure this assumption always holds *before the program is run*.
- If desired, it allows program behavior to be formally stated and proven correct.

In this chapter, I'll introduce the Idris programming language and give a brief tour of its features and environment. I'll also provide an overview of type-driven development, discussing why types matter in programming languages and how they can be used to guide software development. But first, it's important to understand exactly what we mean when we talk about "types."

1.1 What is a type?

We're taught from an early age to recognize and distinguish *types* of objects. As a young child, you may have had a shape-sorter toy. This consists of a box with variously shaped holes in the top (see figure 1.1) and some shapes that fit through the holes. Sometimes they're equipped with a small plastic hammer. The idea is to fit each shape (think of this as a "value") into the appropriate hole (think of this as a "type"), possibly with coercion from the hammer.

In programming, types are a means of classifying values. For example, the values 94, "thing", and [1, 2, 3, 4, 5] could respectively be classified as an integer, a string, and a list of integers. Just as you can't put a square shape in a round hole in the shape sorter, you can't use a string like "thing" in a part of a program where you need an integer.

All modern programming languages classify values by type, although they differ enormously in when and how they do so (for example, whether they're checked statically at compile time or dynamically at runtime, whether conversions between types are automatic or not, and so on).




Figure 1.1 The top of a shape-sorter toy.
The shapes correspond to the types of objects that will fit through the holes.

Types serve several important roles:

- For a *machine*, types describe how bit patterns in memory are to be interpreted.
- For a *compiler* or *interpreter*, types help ensure that bit patterns are interpreted consistently when a program runs.
- For a *programmer*, types help name and organize concepts, aiding documentation and supporting interactive editing environments.

From our point of view in this book, the most important purpose of types is the third. Types help programmers in several ways:

- By allowing for the naming and organization of concepts (such as `Square`, `Circle`, `Triangle`, and `Hexagon`)
- By providing explicit documentation of the purposes of variables, functions, and programs
- By driving code completion in an interactive editing environment

As you'll see, type-driven development makes extensive use of code completion in particular. Although all modern, statically typed languages support code completion to some extent, the expressivity of the Idris type system leads to powerful automatic code generation.

1.2 *Introducing type-driven development*

Type-driven development is a style of programming in which we write types first and use those types to guide the definition of functions. The overall process is to write the necessary data types, and then, for each function, do the following:

- 1 Write the input and output types.
- 2 Define the function, using the structure of the input types to guide the implementation.
- 3 Refine and edit the type and function definition as necessary.

In type-driven development, instead of thinking of types in terms of *checking*, with the type checker criticizing you when you make a mistake, you can think of types as being a *plan*, with the type checker acting as your guide, leading you to a working, robust program. Starting with a type and an empty function body, you gradually add details to the definition until it's complete, regularly using the

Types as models

When you write a program, you'll often have a conceptual model in your head (or, if you're disciplined, even on paper) of how it's supposed to work, how the components interact, and how the data is organized. This model is likely to be quite vague at first and will become more precise as the program evolves and your understanding of the concept develops.

Types allow you to make these models explicit in code and ensure that your implementation of a program matches the model in your head. Idris has an expressive type system that allows you to describe a model as precisely as you need, and to refine the model at the same time as developing the implementation.

compiler to check that the program so far satisfies the type. Idris, as you'll soon see, strongly encourages this style of programming by allowing *incomplete* function definitions to be checked, and by providing an expressive language for describing types.

To illustrate further, in this section I'll show some examples of how you can use types to describe in detail what a program is intended to do: matrix arithmetic, modeling an automated teller machine (ATM), and writing concurrent programs. Then, I'll summarize the process of type-driven development and introduce the concept of *dependent types*, which will allow you to express detailed properties of your programs.

1.2.1 Matrix arithmetic

A *matrix* is a rectangular grid of numbers, arranged in rows and columns. They have several scientific applications, and in programming they have applications in cryptography, 3D graphics, machine learning, and data analytics. The following, for example, is a 3×4 matrix:

$$\begin{pmatrix} 1 & 2 & 3 & 4 \\ 5 & 6 & 7 & 8 \\ 9 & 10 & 11 & 12 \end{pmatrix}$$

You can implement various arithmetic operations on matrices, such as addition and multiplication. To *add* two matrices, you add the corresponding elements, as you see here:

$$\begin{pmatrix} 1 & 2 \\ 3 & 4 \\ 5 & 6 \end{pmatrix} + \begin{pmatrix} 7 & 8 \\ 9 & 10 \\ 11 & 12 \end{pmatrix} = \begin{pmatrix} 8 & 10 \\ 12 & 14 \\ 16 & 18 \end{pmatrix}$$

When programming with matrices, if you begin by defining a `Matrix` data type, then addition requires two inputs of type `Matrix` and gives an output of type `Matrix`. But because adding matrices involves adding corresponding elements of the inputs, what happens if the two inputs have different dimensions, as here?

$$\begin{pmatrix} 1 & 2 \\ 3 & 4 \\ 5 & 6 \end{pmatrix} + \begin{pmatrix} 7 & 8 \\ 9 & 10 \end{pmatrix} = ???$$

It's likely that if you're trying to add matrices of different dimensions, then you've made a mistake somewhere. So, instead of using a `Matrix` type, you could *refine* the

type so that it includes the dimensions of the matrix, and require that the two input matrices have the same dimensions:

- The first example of a 3×4 matrix now has type `Matrix 3 4`.
- The first (correct) example of addition takes two inputs of type `Matrix 3 2` and gives an output of type `Matrix 3 2`.

By including the dimensions in the type of a matrix, you can describe the input and output types of addition in such a way that it's a *type error* to try to add matrices of different sizes. If you try to add a `Matrix 3 2` and a `Matrix 2 2`, your program won't compile, let alone run.

If you include the dimensions of a matrix in its type, then you need to think about the relationship between the dimensions of the input and output for *every* matrix operation. For example, transposing a matrix involves switching the rows to columns and vice versa, so if you transpose a 3×2 matrix, you'll end up with a 2×3 matrix:

$$\begin{pmatrix} 1 & 2 \\ 3 & 4 \\ 5 & 6 \end{pmatrix} \dots \text{transposed to} \dots \begin{pmatrix} 1 & 3 & 5 \\ 2 & 4 & 6 \end{pmatrix}$$

The input type of this transposition is `Matrix 3 2`, and the output type is `Matrix 2 3`.

In general, rather than giving *exact* dimensions in the type, we'll use *variables* to describe the relationship between the dimensions of the inputs and the dimensions of the outputs. Table 1.1 shows the relationships between the dimensions of inputs and outputs for three matrix operations: addition, multiplication, and transposition.

Table 1.1 Input and output types for matrix operations. The names `x`, `y`, and `z` describe, in general, how the dimensions of the inputs and outputs are related.

Operation	Input types	Output type
Add	<code>Matrix x y</code> , <code>Matrix x y</code>	<code>Matrix x y</code>
Multiply	<code>Matrix x y</code> , <code>Matrix y z</code>	<code>Matrix x z</code>
Transpose	<code>Matrix x y</code>	<code>Matrix y x</code>

We'll look at matrices in depth in chapter 3, where we'll work through an implementation of matrix transposition in detail.

1.2.2 An automated teller machine

As well as using types to describe the relationships between the inputs and outputs of functions, as with matrix operations, you can describe precisely *when* operations are valid. For example, if you're implementing software to drive an ATM, you'll want to guarantee that the machine will dispense cash only after a user has entered a card and validated their personal identification number (PIN).

To see how this works, we'll need to consider the possible *states* that an ATM can be in:

- Ready—The ATM is ready and waiting for a user to insert a card.
- CardInserted—The ATM is waiting for a user, having inserted a card, to enter their PIN.
- Session—A validated session is in progress, with the ATM, having validated the user's PIN, ready to dispense cash.

An ATM supports several basic operations, each of which is valid only when the machine is in a specific state, and each of which might change the state of the machine, as illustrated in figure 1.2. These are the basic operations:

- InsertCard—Waits for the user to insert a card
- EjectCard—Ejects a card from the machine
- GetPIN—Prompts the user to enter a PIN
- CheckPIN—Checks whether an entered PIN is correct
- Dispense—Dispenses cash




Figure 1.2 The states and valid operations on an ATM. Each operation is valid only in specific states and can change the state of the machine. CheckPIN changes the state only if the entered PIN is correct.

Whether an operation is valid or not depends on the state of the machine. For example, InsertCard is valid only in the Ready state, because that's the only state where there's no card already in the machine. Also, Dispense is valid only in the Session state, because that's the only state where there's a validated card in the machine.

Furthermore, executing one of these operations can *change* the state of the machine. For example, InsertCard changes the state from Ready to CardInserted, and CheckPIN changes the state from CardInserted to Session, provided that the entered PIN is correct.

STATE MACHINES AND TYPES Figure 1.2 illustrates a *state machine*, describing how operations affect the overall state of a system. State machines are often present, implicitly, in real-world systems. For example, when you open, read,

and then `close` a file, you change the state of the file with the `open` and `close` operations. As you'll see in chapter 13, types allow you to make these state changes explicit, guarantee that you'll execute operations only when they're valid, and help you use resources correctly.

By defining precise types for each of the operations on the ATM, you can guarantee, by type checking, that the ATM will execute only valid operations. If, for example, you try to implement a program that dispenses cash without validating a PIN, the program won't compile. By defining valid state transitions explicitly in types, you get strong and machine-checkable guarantees about the correctness of their implementation. We'll look at state machines in chapter 13, and then implement the ATM example in chapter 14.

1.2.3 Concurrent programming

A *concurrent* program consists of multiple processes running at the same time and coordinating with each other. Concurrent programs can be responsive and continue to interact with a user while a large computation is running. For example, a user can continue browsing a web page while a large file is downloading. Moreover, by writing concurrent programs we can take full advantage of the processor power of modern CPUs, dividing work among multiple processes on separate CPU cores.

In Idris, processes coordinate with each other by sending and receiving *messages*. Figure 1.3 shows one way this can work, with two processes, `main` and `adder`. The `adder` process waits for a *request* to add numbers from other processes. After it receives a message from `main` asking it to add two numbers, it sends a *response* back with the result.

Despite its advantages, however, concurrent programming is notoriously error prone. The need for processes to interact with each other can greatly increase a system's complexity. For each process, you need to ensure that the messages it sends and




Figure 1.3 Two interacting concurrent processes, `main` and `adder`. The `main` process sends a request to `adder`, which then sends a response back to `main`.

receives are properly coordinated with other processes. If, for example, `main` and `adder` aren't properly coordinated and each is expecting to receive a message from the other at the same time, they'll *deadlock*.

TYPES VERSUS TESTING FOR CONCURRENT PROGRAMS Testing a concurrent program is difficult because, unlike a purely sequential program, there's no guarantee about the order in which operations from different processes will execute. Even if two processes are correctly coordinated when you run a test once, there's no guarantee they'll be correctly coordinated when you next run the test. On the other hand, if you can express the coordination between processes in *types*, you can be sure that a concurrent program that type-checks has properly coordinated processes.

When you write concurrent programs, you'll ideally have a model of how processes should interact. Using types, you can make this model explicit in code. Then, if a concurrent program type-checks, you'll know that it correctly follows the model. In particular, you can do two things:

- Define an interface for `adder` that describes the form of messages it will handle.
- Define a protocol that defines the order of message passing, ensuring that `main` will *always* send a message to `adder` and then receive a reply, and `adder` will always do the opposite.

Concurrent programming is an extensive topic, and there are several ways you can use types to model coordination between processes. We'll look at one example of how to do this in chapter 15.

1.2.4 **Type, define, refine: the process of type-driven development**

In each of these introductory examples, we've discussed in general terms how we might *model* a system: by describing the valid forms of inputs and outputs for matrix operations, the valid states of an interactive system, or the order of transmission of messages between concurrent processes. In each case, to implement the system, you start by trying to find a type that captures the important details of the model, and then define functions to work with that type, refining the type as necessary.

To put it succinctly, you can characterize type-driven development as an iterative process of *type*, *define*, *refine*: writing a type, implementing a function to satisfy that type, and refining the type or definition as you learn more about the problem.

With matrix addition, for example, you do the following:

- *Type*—Write a `Matrix` data type, and use it as the input and output types for an addition function.
- *Define*—Write an addition function that satisfies its input and output types.
- *Refine*—Notice that the input and output types for your addition function allow you to give invalid inputs with different dimensions, and then make the type more precise by including the dimensions of the matrices.

In general, you'll write a type to represent the system you're modeling, define functions using that type, and then refine the type and definition as necessary to capture any missing properties. You'll see a lot more of this type-define-refine process throughout this book, both on a small scale when implementing individual functions, and on a larger scale when deciding how to write function and data types.

1.2.5 Dependent types

In the matrix arithmetic example, we began with a `Matrix` type and then refined it to include the number of rows and columns. This means, for example, that `Matrix 3 4` is the type of 3×4 matrices. In this type, 3 and 4 are ordinary values. A *dependent type*, such as `Matrix`, is a type that's calculated from some other values. In other words, it *depends on other values*.

By including values in a type like this, you can make types as precise as required. For example, some languages have a simple list type, describing lists of objects. You can make this more precise by parameterizing over the element type: a generic list of strings is more precise than a simple list and differs from a list of integers. You can be more precise still with a dependent type: a list of 4 strings differs from a list of 3 strings.

Table 1.2 illustrates how types in Idris can have differing levels of precision even for fundamental operations such as appending lists. Suppose you have two specific input lists of strings:

```
["a", "b", "c", "d"]
["e", "f", "g"]
```

When you append them, you'll expect the following output list:

```
["a", "b", "c", "d", "e", "f", "g"]
```

Using a *simple* type, both input lists have type `AnyList`, as does the output list. Using a *generic* type, you can specify that the input lists are both lists of strings, as is the output list. The more-precise types mean that, for example, the output is clearly related to the input in that the element type is unchanged. Finally, using a *dependent* type, you can specify the sizes of the input and output lists. It's clear from the type that the length of the output list is the sum of the lengths of the input lists. That is, a list of 3 strings appended to a list of 4 strings results in a list of 7 strings.

Table 1.2 Appending specific typed lists. Unlike simple types, where there's no difference between the input and output list types, dependent types allow the length to be encoded in the type.

	Input ["a", "b", "c", "d"]	Input ["e", "f", "g"]	Output type
Simple	<code>AnyList</code>	<code>AnyList</code>	<code>AnyList</code>
Generic	<code>List String</code>	<code>List String</code>	<code>List String</code>
Dependent	<code>Vect 4 String</code>	<code>Vect 3 String</code>	<code>Vect 7 String</code>

LISTS AND VECTORS The syntax for the types in table 1.2 is valid Idris syntax. Idris provides several ways of building list types, with varying levels of precision. In the table, you can see two of these, List and Vect. AnyList is included in the table purely for illustrative purposes and is not defined in Idris. List encodes generic lists with no explicit length, and Vect (short for “vector”) encodes lists with the length explicitly in the type. You’ll see much more of both these types throughout this book.

Table 1.3 illustrates how the input and output types of an append function can be written with increasing levels of precision in Idris. Using *simple* types, you can write the input and output types as AnyList, suggesting that you have no interest in the types of the elements of the list. Using *generic* types, you can write the input and output types as List elem. Here, elem is a *type variable* standing for the element types. Because the type variable is the same for both inputs and the output, the types specify that both the input lists and the output list have a consistent element type. If you append two lists of integers, the types guarantee that the output will also be a list of integers. Finally, using *dependent* types, you can write the inputs as Vect n elem and Vect m elem, where n and m are variables representing the *length* of each list. The output type specifies that the resulting length will be the sum of the lengths of the inputs.

Table 1.3 Appending typed lists, in general. Type variables describe the relationships between the inputs and outputs, even though the exact inputs and outputs are unknown.

	Input 1 type	Input 2 type	Output type
Simple	AnyList	AnyList	AnyList
Generic	List elem	List elem	List elem
Dependent	Vect n elem	Vect m elem	Vect (n + m) elem

TYPE VARIABLES Types often contain *type variables*, like n, m, and elem in table 1.3. These are very much like parameters to generic types in Java or C#, but they’re so common in Idris that they have a very lightweight syntax. In general, concrete type names begin with an uppercase letter, and type variable names begin with a lowercase letter.

In the dependent type for the append function in table 1.3, the parameters n and m are ordinary numeric values, and the + operator is the normal addition operator. All of these could appear in programs just as they’ve appeared here in the types.

Introductory exercises



Throughout this book, exercises will help reinforce the concepts you’ve learned. As a warm-up, take a look at the following selection of function specifications, given purely in the form of input and output types. For each of them, suggest possible operations

that would satisfy the given input and output types. Note that there could be more than one answer in each case.

- 1** Input type: Vect n elem
Output type: Vect n elem
- 2** Input type: Vect n elem
Output type: Vect (n * 2) elem
- 3** Input type: Vect (1 + n) elem
Output type: Vect n elem
- 4** Assume that Bounded n represents a number between zero and n - 1.
Input types: Bounded n, Vect n elem
Output type: elem

1.3 Pure functional programming

Idris is a *pure functional* programming language, so before we begin exploring Idris in depth, we should look at what it means for a language to be *functional*, and what we mean by the concept of *purity*. Unfortunately, there's no universally agreed-on definition of exactly what it means for a programming language to be *functional*, but for our purposes we'll take it to mean the following:

- Programs are composed of functions.
- Program execution consists of the evaluation of functions.
- Functions are a first-class language construct.

This differs from an *imperative* programming language primarily in that functional programming is concerned with the evaluation of functions, rather than the execution of statements.

In a *pure* functional language, the following are also true:

- Functions don't have side effects such as modifying global variables, throwing exceptions, or performing console input or output.
- As a result, for any specific inputs, a function will always give the same result.

You may wonder, very reasonably, how it's possible to write any useful software under these constraints. In fact, far from making it more difficult to write realistic programs, pure functional programming allows you to treat tricky concepts such as state and exceptions with the respect they deserve. Let's explore further.

1.3.1 Purity and referential transparency

The key property of a pure function is that the same inputs always produce the same result. This property is known as *referential transparency*. An expression (such as a function call) in a function is referentially transparent if it can be replaced with its result without changing the behavior of the function. If functions produce only results, with no side effects, this property is clearly true. Referential transparency is a very useful concept in type-driven development, because if a function has no side effects and is



Figure 1.4 A pure function, taking inputs and producing outputs with no observable side effects

defined entirely by its inputs and outputs, then you can look at its input and output types and have a clear idea of the limits of what the function can do.

Figure 1.4 shows example inputs and outputs for the `append` function. It takes two inputs and produces a result, but there's no interaction with a user, such as reading from the keyboard, and no informative output, such as logging or progress bars.

Figure 1.5 shows pure functions in general. There can be no observable side effects when running these programs, other than perhaps making the computer slightly warmer or taking a different amount of time to run.




Figure 1.5 Pure functions, in general, take only inputs and have no observable side effects.

Pure functions are very common in practice, particularly for constructing and manipulating data structures. It's possible to reason about their behavior because the function always gives the same result for the same inputs; these functions are important components of larger programs. The preceding `append` function is pure, and it's a valuable component for any program that works with lists. It produces a list as a result, and because it's pure, you know that it won't require any input, output any logging, or do anything destructive like delete files.

1.3.2 Side-effecting programs

Realistically, programs must have side effects in order to be useful, and you're always going to have to deal with unexpected or erroneous inputs in practical software. At first, this would seem to be impossible in a pure language. There is a way, however: pure functions may not be able to perform side effects, but they can *describe* them.

Consider a function that reads two lists from a file, appends them, prints the resulting list, and returns it. The following listing outlines this function in imperative-style pseudocode, using simple types.

Listing 1.1 Appending lists read from a file (pseudocode)

```

List appendFromFile(File h) {
    list1 = readListFrom(h)
    list2 = readListFrom(h)

    result = append(list1, list2)
    print(result)
  
```




Figure 1.6 A side-effecting program, reading inputs from a file, printing the result, and returning the result

```

        return result
}
  
```

This program takes a file handle as an input and returns a `List` with some side effects. It reads two lists from the given file and prints the list before returning. Figure 1.6 illustrates this for the situation when the file contains the two lists `["a", "b", "c", "d"]` and `["e", "f", "g"]`.

The `appendFromFile` function doesn't satisfy the referential transparency property. Referential transparency requires that an expression can be replaced by its result without changing the program's behavior. Here, however, replacing a call to `appendFromFile` with its result means that nothing will be read from the file, and nothing will be output to the screen. The function's *input* and *output* types tell us that the input is a file and the output is a list, but nothing in the type describes the side effects the function may execute.

In pure functional programming in general, and Idris in particular, you can solve this problem by writing functions that *describe* side effects, rather than functions that *execute* them, and defer the details of execution to the compiler and runtime system. We'll explore this in greater detail in chapter 5; for now, it's sufficient to recognize that a program with side effects has a type that makes this explicit. For example, there's a distinction between the following:

- `String` is the type of a program that results in a `String` and is guaranteed to perform *no* input or output as side effects.
- `IO String` is the type of a program that describes a sequence of input and output operations that result in a `String`.

Type-driven development takes this idea much further. As you'll see from chapter 12 onward, you can define types that describe the specific side effects a program can have, such as console interaction, reading and writing global state, or spawning concurrent processes and sending messages.

1.3.3 Partial and total functions

Idris supports an even stronger property than purity for functions, making a distinction between *partial* and *total* functions. A *total* function is guaranteed to produce a result, meaning that it will return a value in a finite time for every possible well-typed input, and it's guaranteed not to throw any exceptions. A *partial* function, on the other hand, might not return a result for some inputs. Here are a couple of examples:

- The append function is *total* for finite lists, because it will always return a new list.
- The function that returns the first element of a list is *partial*, because it's not defined if the list is empty, and it will therefore crash.

TOTAL FUNCTIONS AND LONG-RUNNING PROGRAMS A total function is guaranteed to produce a *finite prefix* of a potentially infinite result. As you'll see in chapter 11, you can write command shells or servers as total functions that guarantee a response for every user input, indefinitely.

The distinction is important because knowing that a function is total allows you to make much stronger claims about its behavior based on its type. If you have a function with a return type of `String`, for example, you can make different claims depending on whether the function is partial or total.

- *If it's total*—It will return a value of type `String` in finite time.
- *If it's partial*—If it doesn't crash or enter an infinite loop, the value it returns will be a `String`.

In most modern languages, we must assume that functions are partial and can therefore only make the latter, weaker, claim. Idris checks whether functions are total, so we can therefore often make the former, stronger, claim.

Total functions and the halting problem

The *halting problem* is the problem of determining whether a program terminates for some specific input. Thanks to Alan Turing, we know that it's not possible to write a program that solves the halting problem in general. Given this, it's reasonable to wonder how Idris can determine that a function is total, which is essentially checking that a function terminates for *all* inputs.

Although it can't solve the problem in general, Idris can identify a large class of functions that are definitely total. You'll learn more about how it does so, along with some techniques for writing total functions, in chapters 10 and 11.

A useful pattern in type-driven development is to write a type that precisely describes the valid states of a system (like the ATM in section 1.2.2) and that constrains the oper-

ations the system is allowed to perform. A total function with that type is then *guaranteed* by the type checker to perform those operations as precisely as the type requires.

1.4 A quick tour of Idris

The Idris system consists of an interactive environment and a batch mode compiler. In the interactive environment, you can load and type-check source files, evaluate expressions, search libraries, browse documentation, and compile and run complete programs. We'll use these features extensively throughout this book.

In this section, I'll briefly introduce the most important features of the environment, which are evaluation and type checking, and describe how to compile and run Idris programs. I'll also introduce the two most distinctive features of the Idris language itself:

- *Holes*, which stand for incomplete programs
 - The use of types as *first-class* language constructs

As you'll see, by using holes you can define functions incrementally, asking the type checker for contextual information to help complete definitions. Using first-class types, you can be very precise about what a function is intended to do, and even ask the type checker to fill in some of the details of functions for you.

1.4.1 The interactive environment

Much of your interaction with Idris will be through an interactive environment called the *read-eval-print loop*, typically abbreviated as REPL. As the name suggests, the REPL will *read* input from the user, usually in the form of an expression, *evaluate* the expression, and then *print* the result.

Once Idris is installed, you can start the REPL by typing `idris` at a shell prompt. You should see something like the following:

Idris is free software with ABSOLUTELY NO WARRANTY.
For details type :warranty.
Idris>

INSTALLING IDRIS You can find instructions on how to download and install Idris for Linux, OS X, or Windows in appendix A.

You can enter expressions to be evaluated at the `Idris>` prompt. For example, arithmetic expressions work in the conventional way, with the usual precedence rules (that is, `*` and `/` have higher precedence than `+` and `-`):

```
Idris> 2 + 2  
4 : Integer
```

```
Idris> 2.1 * 20
42.0 : Double

Idris> 6 + 8 * 11
94 : Integer
```

You can also manipulate `Strings`. The `++` operator concatenates `Strings`, and the `reverse` function reverses a `String`:

```
Idris> "Hello" ++ " " ++ "World!"
"Hello World!" : String

Idris> reverse "abcdefg"
"gfedcba" : String
```

Notice that Idris prints not only the result of evaluating the expression, but also its *type*. In general, if you see something of the form `x : T`—some expression `x`, a colon, and some other expression `T`—this can be read as “`x` has type `T`.” In the previous examples, you have the following:

- 4 has type `Integer`.
- 42.0 has type `Double`.
- "Hello World!" has type `String`.

1.4.2 Checking types

The REPL provides a number of *commands*, all prefixed by a colon. One of the most commonly useful is `:t`, which allows you to check the *types* of expressions without evaluating them:

```
Idris> :t 2 + 2
2 + 2 : Integer

Idris> :t "Hello!"
"Hello!" : String
```

Types, such as `Integer` and `String`, can be manipulated just like any other value, so you can check their types too:

```
Idris> :t Integer
Integer : Type

Idris> :t String
String : Type
```

It’s natural to wonder what the type of `Type` itself might be. In practice, you’ll never need to worry about this, but for the sake of completeness, let’s take a look:

```
Idris> :t Type
Type : Type 1
```

That is, Type has type Type 1, Type 1 has type Type 2, and so on forever, as far as we're concerned. The good news is that Idris will take care of the details for you, and you can always write Type alone.

1.4.3 Compiling and running Idris programs

As well as evaluating expressions and inspecting the types of functions, you'll want to be able to compile and run complete programs. The following listing shows a minimal Idris program.

Listing 1.2 Hello, Idris World! (Hello.idr)

```
module Main           ← Module header
main : IO ()          ← Function declaration
main = putStrLn "Hello, Idris World!" ← Function definition
```

At this stage, there's no need to worry too much about the syntax or how the program works. For now, you just need to know that Idris source files consist of a module header and a collection of function and data type definitions. They may also import other source files.

WHITESPACE SIGNIFICANCE Whitespace is significant in Idris, so when you type listing 1.2, make sure there are no spaces at the beginning of each line.

Here, the module is called `Main`, and there's only one function definition, called `main`. The entry point to any Idris program is the `main` function in the `Main` module.

To run the program, follow these steps:

- 1 Create a file called `Hello.idr` in a text editor.¹ Idris source files all have the extension `.idr`.
- 2 Enter the code in listing 1.2.
- 3 In the working directory where you saved `Hello.idr`, start up an Idris REPL with the command `idris Hello.idr`.
- 4 At the Idris prompt, type `:exec`.

If all is well, you should see something like the following:

```
$ idris Hello.idr
/ ____/ ____/ ____(____)
/ / / ____/ ____/ / ____/
/ / / / / / / / (____)      Version 1.0
/ ____\____, / / / / / ____/      http://www.idris-lang.org/
                                         Type :? for help

Idris is free software with ABSOLUTELY NO WARRANTY.
For details type :warranty.
Type checking ./Hello.idr
```

¹ I recommend Atom because it has a mode for interactive editing of Idris programs, which we'll use in this book.

```
*Hello> :exec
Hello, Idris World
```

Here, \$ stands for your shell prompt. Alternatively, you can create a standalone executable by invoking the `idris` command with the `-o` option, as follows:

```
$ idris Hello.idr -o Hello
$ ./Hello
Hello, Idris World
```

THE REPL PROMPT The REPL prompt, by default, tells you the name of the file that's currently loaded. The `Idris>` prompt indicates that no file is loaded, whereas the prompt `*Hello>` indicates that the `Hello.idr` file is loaded.

1.4.4 Incomplete definitions: working with holes

Earlier, I compared working with types and values to inserting shapes into a shape-sorter toy. Much as the square shape will only fit through a square hole, the argument `"Hello, Idris World!"` will only fit into a function in a place where a `String` type is expected.

Idris functions themselves can contain *holes*, and a function with a hole is *incomplete*. Only a value of an appropriate type will fit into the hole, just as a square shape will only fit into a square hole in the shape sorter. Here's an incomplete implementation of the "Hello, Idris World!" program:

```
module Main
main : IO ()
main = putStrLn ?greeting
```



If you edit `Hello.idr` to replace the string `"Hello, Idris World!"` with `?greeting` and load it into the Idris REPL, you should see something like the following:

```
Type checking ./Hello.idr
Holes: Main.greeting
*Hello>
```

The syntax `?greeting` introduces a *hole*, which is a part of the program yet to be written. You can type-check programs with holes and evaluate them at the REPL.

Here, when Idris encounters the `?greeting` hole, it creates a new name, `greeting`, that has a type but no definition. You can inspect the type using `:t` at the REPL:

```
*Hello> :t greeting
-----
greeting : String
```

If you try to evaluate it, on the other hand, Idris will show you that it's a hole:

```
*Hello> greeting
?greeting : String
```

Reloading

Instead of exiting the REPL and restarting, you can also reload Hello.idr with the :r REPL command as follows:

```
*Hello> :r
Type checking ./Hello.idr
Holes: Main.greeting
*Hello>
```

Holes allow you to develop programs *incrementally*, writing the parts you know and asking the machine to help you by identifying the types for the parts you don't. For example, let's say you'd like to print a character (with type Char) instead of a String. The putStrLn function requires a String argument, so you can't simply pass a Char to it.

Listing 1.3 A program with a type error

```
module Main

main : IO ()
main = putStrLn 'x'
```

Type error, giving a
character instead of a string

If you try loading this program into the REPL, Idris will report an error:

```
Hello.idr:4:17:When checking right hand side of main:
When checking an application of function Prelude.putStrLn:
    Type mismatch between
        Char (Type of 'x')
    and
        String (Expected type)
```

You have to convert a Char to a String somehow. Even if you don't know exactly how to do this at first, you can start by adding a hole to stand in for a conversion.

```
module Main

main : IO ()
main = putStrLn (?convert 'x')
```

Then you can check the type of the convert hole:

```
*Hello> :t convert
-----
convert : Char -> String
```

This is a function type, taking a Char
as input and returning a String.

The type of the hole, Char -> String, is the type of a function that takes a Char as an input and returns a String as an output. We'll discuss type conversions in more detail in chapter 2, but an appropriate function to complete this definition is cast:

```
main : IO ()
main = putStrLn (cast 'x')
```

1.4.5 First-class types

A *first-class* language construct is one that's treated as a value, with no syntactic restrictions on where it can be used. In other words, a first-class construct can be passed to functions, returned from functions, stored in variables, and so on.

In most statically typed languages, there are restrictions on where types can be used, and there's a strict syntactic separation between types and values. You can't, for example, say `x = int` in the body of a Java method or C function. In Idris, there are no such restrictions, and types are first-class; not only can types be used in the same way as any other language construct, but any construct can appear as part of a type.


This means that you can write functions that compute types, and the return type of a function can differ depending on the input *value* to a function. This idea comes up regularly when programming in Idris, and there are several real-world situations where it's useful:

- A database schema determines the allowed forms of queries on a database.
- A form on a web page determines the number and type of inputs expected.
- A network protocol description determines the types of values that can be sent or received over a network.

In each of these cases, one piece of data tells you about the expected form of some other data. If you've programmed in C, you'll have seen a similar idea with the `printf` function, where one argument is a format string that describes the number and expected types of the remaining arguments. The C type system can't check that the format string is consistent with the arguments, so this check is often hardcoded into C compilers. In Idris, however, you can write a function similar to `printf` directly, by taking advantage of types as first-class constructs. You'll see this specific example in chapter 6.

The following listing illustrates the concept of first-class types with a small example: computing a type from a Boolean input.

Listing 1.4 Calculating a type, given a Boolean value (FCTypes.idr)



```

    ▶ valToString : (x : Bool) -> StringOrInt x -> String
    valToString x val = case x of
        True => cast val
        False => val

```

The argument type is calculated from the value of the input.

The input x was True, so the argument val must be an Int and needs to be converted to a String.

The input x was False, so the argument val must be a String and can be returned directly.

Function syntax

We'll go into much more detail on Idris syntax in the coming chapters. For now, just keep the following in mind:

- A *function type* takes the form $a \rightarrow b \rightarrow \dots \rightarrow t$, where a , b , and so on, are the *input types*, and t is the *output type*. Inputs may also be annotated with names, taking the form $(x : a) \rightarrow (y : b) \rightarrow \dots \rightarrow t$.
- $\text{name} : \text{type}$ declares a new function, name , of type type .
- Functions are defined by *equations*:

```
square x = x * x
```

This defines a function called `square` that multiplies its input by itself.

Here, `StringOrInt` is a function that computes a type. Listing 1.4 uses it in two ways:

- In `getStringOrInt`, `StringOrInt` calculates the return type. If the input is `True`, `getStringOrInt` returns an `Int`; otherwise it returns a `String`.
- In `valToString`, `StringOrInt` calculates an argument type. If the first input is `True`, the second input must be an `Int`; otherwise it must be a `String`.

You can see in detail what's going on by introducing holes in the definition of `valToString`:

```

valToString : (x : Bool) -> StringOrInt x -> String
valToString x val = case x of
    True => ?xtrueType
    False => ?xfalseType

```

Inspecting the type of a hole with `:t` gives you not only the type of the hole itself, but also the types of any local variables in scope. If you check the type of `xtrueType`, you'll see the type of `val`, which is computed when `x` is known to be `True`:

```

*FCTypes> :t xtrueType
x : Bool
val : Int
-----
xtrueType : String

```

So, if `x` is `True`, then `val` must be an `Int`, as computed by the `StringOrInt` function. Similarly, you can check the type of `xfalseType` to see the type of `val` when `x` is known to be `False`:

```
*FCTypes> :t xfalseType
x : Bool
val : String
-----
xfalseType : String
```

This is a small example, but it illustrates a fundamental concept of type-driven development and programming with dependent types: **the idea that the *type* of a variable can be computed from the *value* of another.** In each case, Idris has used `StringOrInt` to refine the type of `val`, given what it knows about the value of `x`.

1.5 Summary

- Types are a means of classifying values. Programming languages use types to decide how to lay out data in memory, and to ensure that data is interpreted consistently.
- A type can be viewed as a specification, so that a language implementation (specifically, its type checker) can check whether a program conforms to that specification.
- Type-driven development is an iterative process of type, define, refine, creating a type to model a system, then defining functions, and finally refining the types as necessary.
- In type-driven development, a type is viewed more like a plan, helping an interactive environment guide the programmer to a working program.
- Dependent types allow you to give more-precise types to programs, and hence more informative plans to the machine.
- In a functional programming language, program execution consists of evaluating functions.
- In a purely functional programming language, additionally, functions have no *side effects*.
- Instead of writing programs that perform side effects, you can write programs that describe side effects, with the side effects stated explicitly in a program's type.
- A total function is guaranteed to produce a result for any well-typed input in finite time.
- Idris is a programming language that's specifically designed to support type-driven development. It's a purely functional programming language with first-class dependent types.
- Idris allows programs to contain holes that stand for incomplete programs.
- In Idris, types are first-class, meaning that they can be stored in variables, passed to functions, or returned from functions like any other value.

Getting started with Idris



This chapter covers

- Using built-in types and functions
- Defining functions
- Structuring Idris programs

When learning any new language, it's important to have a solid grasp of the fundamentals before moving on to the more distinctive features of the language. With this in mind, before we begin exploring dependent types and type-driven development itself, we'll look at some types and values that will be familiar to you from other languages, and you'll see how they work in Idris. You'll also see how to define functions and put these together to build a complete, if simple, Idris program.

If you're already familiar with a pure functional language, particularly Haskell, much of this chapter will seem very familiar. Listing 2.1 shows a simple, but self-contained, Idris program that repeatedly prompts for input from the console and then displays the average length of the words in the input. If you're already comfortable reading this program with the help of the annotations, you can safely skip this chapter, as it deliberately avoids introducing any language features specific to Idris.¹

¹ Comparing Idris to Haskell, the most important difference is that Idris doesn't use lazy evaluation by default.

Even so, I still suggest you browse through this chapter’s tips and notes and read the summary at the end to make sure there aren’t any small details you’ve missed.

Otherwise, don’t worry. By the end of this chapter we’ll have covered all of the necessary features for you to be able to implement similar programs yourself.

Listing 2.1 A complete Idris program to calculate average word length (Average.idr)

All top-level functions must have a type declaration. Argument types may optionally be given names as part of the type declaration. Here, there’s one argument with type String and name str.

The cast function explicitly converts between types. Here, the division operator requires a Double, but the totalLength and numWords variables are Nats.

```
module Main

average : (str : String) -> Double
average str = let numWords = wordCount str
              totalLength = sum (allLengths (words str)) in
                cast totalLength / cast numWords
where
  wordCount : String -> Nat
  wordCount str = length (words str)

  allLengths : List String -> List Nat
  allLengths strs = map length strs

showAverage : String -> String
showAverage str = "The average word length is: " ++
                  show (average str) ++ "\n"

main : IO ()
main = repl "Enter a string: " showAverage
```

The function main, in a module called Main, is the entry point to an Idris program.

The where keyword introduces local function definitions. Here, they’re only visible inside the scope of the average function.

String is a primitive type, unlike some other languages (notably Haskell, where a string is represented as a list of characters).

repl is a function that repeatedly displays a prompt, reads a String from the console, and then displays the result of running a function on that String.

2.1 Basic types

Idris provides some standard basic types and functions for working with various forms, characters, and strings. In this section, I’ll give you an overview of these, along with some examples. These basic types are defined in the *Prelude*, which is a collection of standard types and functions automatically imported by every Idris program.

I’ll show you several example expressions in this section, and it may seem fairly clear what they should do. Nevertheless, instead of simply reading them and nodding, I strongly recommend you type the examples at the Idris REPL. You’ll learn the syntax much more easily by using it than you will by merely reading it.

Along the way, we’ll also encounter a couple of useful REPL features that will allow us to store the results of calculations at the REPL.

THE PRELUDE The types and functions I'll discuss are defined in the *Prelude*. The Prelude is Idris's standard library, always available at the REPL and automatically imported by every Idris program. With the exception of some primitive types and operations, everything in the Prelude is written in Idris itself.

2.1.1 Numeric types and values

Idris provides several basic numeric types, including the following:

- `Int`—A *fixed-width* signed integer type. It's guaranteed to be at least 31 bits wide, but the exact width is system dependent.
- `Integer`—An *unbounded* signed integer type. Unlike `Int`, there's no limit to the size of the numbers that can be represented, other than your machine's memory, but this type is more expensive in terms of performance and memory usage.
- `Nat`—An *unbounded* unsigned integer type. This is very often used for sizes and indexing of data structures, because they can never be negative. You'll see much more of `Nat` later.
- `Double`—A double-precision floating-point type.

SUBTRACTION WITH NATS Because `Nats` can never be negative, a `Nat` can only be subtracted from a larger `Nat`.

We can use standard numeric literals as values for each of these types. For example, the literal `333` can be of type `Int`, `Integer`, `Nat`, or `Double`. The literal `333.0` can be only of type `Double`, due to the explicit decimal point.

You can try some simple calculations at the REPL:

```
Idris> 6 + 3 * 12
42 : Integer
Idris> 6.0 + 3 * 12
42.0 : Double
```

Note that Idris will treat a number as an `Integer` by default, unless there's some context, and both operands must be the same type. Therefore, in the second of the two preceding expressions, the literal `6.0` can be only a `Double`, so the whole expression is a `Double`, and `3` and `12` are also treated as `Doubles`.

When an expression, such as `6 + 3 * 12`, can be one of several types, you can make the type explicit with the notation `<type><expression>`, to say that `type` is the required type of `expression`:

REPL results

The most recent result at the REPL can always be retrieved and used in further calculations by using the special value `it`:

```
Idris> 6 + 3 * 12
42 : Integer
Idris> it * 2
84 : Integer
```

It's also possible to bind expressions to names at the REPL using the `:let` command:

```
Idris> :let x = 100
Idris> x
100 : Integer
Idris> :let y = 200.0
Idris> y
200.0 : Double
```

```
Idris> 6 + 3 * 12
42 : Integer
Idris> the Int (6 + 3 * 12)
42 : Int
Idris> the Double (6 + 3 * 12)
42.0 : Double
```

“THE” EXPRESSIONS the is not built-in syntax but an ordinary Idris function, defined in the Prelude, which takes advantage of first-class types.

2.1.2 Type conversions using cast

Arithmetic operators work on any numeric types, but both inputs and the output must have the *same* type. Sometimes, therefore, you’ll need to convert between types.

Let’s say you’ve defined an Integer and a Double at the REPL:

```
Idris> :let integerval = 6 * 6
Idris> :let doubleval = 0.1
Idris> integerval
36 : Integer
Idris> doubleval
0.1 : Double
```

If you try to add integerval and doubleval, Idris will complain that they aren’t the same type:

```
Idris> integerval + doubleval
(input):1:8-9:When checking an application of function Prelude.Classes.+:
  Type mismatch between
    Double (Type of doubleval)
  and
    Integer (Expected type)
```

To fix this, you can use the cast function, which converts its input to the required type, as long as that conversion is valid. Here, you can cast the Integer to a Double:

```
Idris> cast integerval + doubleval
36.1 : Double
```

Idris supports casting between all the primitive types, and it’s possible to add user-defined casts, as you’ll see in chapter 7. Note that some casts may lose information, such as casting a Double to an Integer.

Specifying the target of a cast

You can also use the to specify which type you want to cast to, as in these examples:

```
Idris> the Integer (cast 9.9)
9 : Integer

Idris> the Double (cast (4 + 4))
8.0 : Double
```

2.1.3 Characters and strings

Idris also provides Unicode characters and strings as primitive types, along with some useful primitive functions for manipulating them.

- Character literals (of type `Char`) are enclosed in single quotation marks, such as `'a'`.
- String literals (of type `String`) are enclosed in double quotation marks, such as `"Hello world!"`

Like many other languages, Idris supports special characters in character and string literals by using *escape sequences*, beginning with a backslash. For example, a newline character is indicated using `\n`:

```
Idris> :t '\n'
'\n' : Char

Idris> :t "Hello world!\n"
"Hello world!\n" : String
```

These are the most common escape sequences:

- `\'` for a literal single quote
- `\"` for a literal double quote
- `\\"` for a literal backslash
- `\n` for a newline character
- `\t` for a tab character

The Prelude defines several useful functions for manipulating `Strings`. You can see some of these in action at the REPL:

```
Idris> length "Hello!"
6 : Nat

Idris> reverse "drawer"
"reward" : String

Idris> substr 6 5 "Hello world"
"world" : String

Idris> "Hello" ++ " " ++ "World"
"Hello World" : String
```

Here's a brief explanation of these functions:

- `length`—Gives the length of its argument as a `Nat` because a `String` can't have a negative length
- `reverse`—Returns a reversed version of its input
- `substr`—Returns a substring of an input string, given a position to start at and the desired length of the substring
- `++`—An operator that concatenates two `Strings`

Notice the syntax of the function calls. In Idris, functions are separated from their arguments by *whitespace*. If the argument is a complex expression, it must be bracketed, as follows:

```
Idris> length ("Hello" ++ " " ++ "World")
11 : Nat
```

FUNCTION SYNTAX Calling functions by separating the arguments with spaces may seem strange at first. There's a good reason for it, though, as you'll discover when we look at function types later in this chapter. In short, it makes manipulating functions much more flexible.

2.1.4 *Booleans*

Idris provides a `Bool` type for representing truth values. A `Bool` can take the value `True` or `False`. The operators `&&` and `||` represent logical *and* and *or*, respectively:

```
Idris> True && False
False : Bool
Idris> True || False
True : Bool
```

The usual comparison operators (`<`, `<=`, `==`, `/=`, `>`, `>=`) are available:

```
Idris> 3 > 2
True : Bool
Idris> 100 == 99
False : Bool
Idris> 100 /= 99
True : Bool
```

INEQUALITY The inequality operator in Idris is `/=`, which follows Haskell syntax, rather than `!=`, which would follow the syntax of languages like C and Java.

There is also an `if...then...else` construct. This is an expression, so it must *always* include both a `then` branch and an `else` branch. For example, you can write an expression that evaluates to a different message as a `String`, depending on the length of a word:

```
Idris> :let word = "programming"
Idris> if length word > 10 then "What a long word!" else "Short word"
"What a long word!" : String
```

2.2 *Functions: the building blocks of Idris programs*

Now that you've seen some basic types and a simple control structure, you can begin defining functions. In this section, you'll write some Idris functions using the basic types you've seen so far, load them into the Idris system, and test them at the REPL. You'll also see how the functional programming style allows you to write more *generic* programs in two ways:

- Using *variables* in function types, so that functions can be written to work with several different types
- Using *higher-order* functions to capture common programming patterns

2.2.1 Function types and definitions

Function types are composed of one or more *input* types and an *output* type. For example, a function that takes an Int as input and returns another Int would be written as Int \rightarrow Int. The following listing shows a simple function definition with this type, the double function.

Listing 2.2 A function to double an Int (Double.idr)

```
double : Int -> Int
double x = x + x
```

The function type states that it takes an Int as input and returns an Int as output.

The function definition gives an equation that defines what it means to double an input.

You can try this function by typing it into a file, Double.idr; loading it into the Idris REPL by typing idris Double.idr at the shell prompt; and then trying some examples at the REPL:

```
*Double> double 47
94 : Int

*Double> double (double 15)
60 : Int
```

Figure 2.1 shows the components of this function definition. All functions in Idris, like double, are introduced with a *type declaration* and then defined by *equations* with a left side and a right side.




Figure 2.1 The components of a function definition

An expression is evaluated by rewriting the expression according to these equations until no further rewriting can be done. Functions, therefore, define *rules* by which expressions can be rewritten. For example, consider the definition of double:

```
double x = x + x
```

This means that whenever the Idris evaluator encounters an expression of the form `double x`, with some expression standing for `x`, it should be rewritten as `x + x`.

So, in the example `double (double 15)`,

- First, the inner `double 15` is rewritten as `15 + 15`.
- `15 + 15` is rewritten as `30`.
- `double 30` is rewritten as `30 + 30`.
- `30 + 30` is rewritten as `60`.

Evaluation order

You might have noticed that instead of choosing to evaluate the inner `double 15` first, you could have chosen the outer `double (double 15)`, which would have reduced to `double 15 + double 15`. Either order is possible, and each would lead to the same result. Idris, by default, will evaluate the *innermost* expression first. In other words, it will evaluate function arguments before function definitions.

There are merits and drawbacks to both choices, and as a result this topic has been debated at length! Now is not the time to revisit this debate, but if you're interested, you can investigate *lazy evaluation*. Idris supports lazy evaluation using explicit types, as you'll see in chapter 11.

Optionally, you can give explicit names in the input types of a function. For example, you could write the type of `double` as follows:

```
double : (value : Int) -> Int
```

This has exactly the same meaning as the previous declaration (`double : Int -> Int`). There are two reasons why you might make the names of the arguments explicit:

- Naming the argument in the type can give the reader some information about the purpose of the argument.
- Naming the argument means you can *refer* to it later.

You'll see more of this in chapter 4 when we begin to explore dependent types in depth. For now, remember the example of first-class types from chapter 1, where I gave the following Idris type for `getStringOrInt`:

```
getStringOrInt : (x : Bool) -> StringOrInt x
```

The first argument, of type `Bool`, was given the name `x`, which then appears in the return type.

TYPE DECLARATIONS ARE REQUIRED! Functions in Idris *must* have an explicit type declaration, like `double : Int -> Int` here. Some other functional languages, most notably Haskell and ML, allow programmers to omit type declarations and have the compiler *infer* the type. In a language with first-class types, however, this generally turns out to be impossible. In any case, it's

undesirable to omit type declarations in type-driven development. Our philosophy is to use types to help us write programs, rather than to use programs to help us infer types!

2.2.2 Partially applying functions

When a function has more than one argument, you can create a specialized version of the function by omitting the later arguments. This is called *partial application*.

For example, assume you have an `add` function that adds two integers, defined as follows in the `Partial.idr` file:

```
add : Int -> Int -> Int
add x y = x + y
```

If you apply the function to two arguments, it will evaluate to an `Int`:

```
*Partial> add 2 3
5 : Int
```

If, on the other hand, you only apply the function to one argument, omitting the second, Idris will return a function of type `Int -> Int`:

```
*Partial> add 2
add 2 : Int -> Int
```

By applying `add` to only one argument, you've created a new specialized function, `add 2`, that adds 2 to its argument. You can see this more explicitly by creating a new function with `:let`:

```
*Partial> :let addTwo = add 2
*Partial> :t addTwo
addTwo : Int -> Int
*Partial> addTwo 3
5 : Int
```

The function application syntax, applying functions to arguments simply by separating the function from the argument with whitespace, gives a particularly concise syntax for partial application. Partial application is common in Idris programs, and you'll see some examples of it in action shortly, in section 2.2.5.

2.2.3 Writing generic functions: variables in types

As well as *concrete* types, such as `Int`, `String`, and `Bool`, function types can contain *variables*. Variables in a function type can be instantiated with different values, just like variables in functions themselves.

For example, let's consider the *identity* function, which returns its input, unchanged. The *identity* function on `Ints` is written as follows:

```
identityInt : Int -> Int
identityInt x = x
```

Similarly, the identity function on Strings is written like this:

```
identityString : String -> String
identityString x = x
```

And here's the identity function on Bools:

```
identityBool : Bool -> Bool
identityBool x = x
```

You may have noticed a pattern here. In each case the definition is the same! You don't need to know anything about `x` because you're returning it unchanged in each case. So, instead of writing an identity function for every type separately, you can write one identity function using a variable, `ty`, at the type level, in place of a concrete type:

```
identity : ty -> ty
identity x = x
```

THE ID FUNCTION In fact, there is an identity function in the Prelude called `id`, defined in exactly the same way as `identity` here.

The `ty` in the type of `identity` is a *variable*, standing for any type. Therefore, `identity` can be called with any input type, and will return an output with the same type as the input.

VARIABLE NAMES IN TYPES Any name that appears in a type declaration, begins with a lowercase letter, and is otherwise undefined is assumed to be a variable. Note that I'm careful to call these *variables*, rather than *type variables*. This is because, with dependent types, variables in types don't necessarily stand for only types, as you'll see in chapter 3.

You've already seen a form of the identity function when working with numeric types: `the` is an identity function. It's defined in the Prelude as follows:

```
the : (ty : Type) -> ty -> ty
the ty x = x
```

It takes an explicit type as its first argument, which is explicitly named `ty`. The type of the second argument is given by the *input value* of the first argument. This is a simple example of dependent types in action, in that the value of an earlier argument gives the type of a later argument. You can see this explicitly at the REPL, by partially applying `the` to only one argument:

```
Idris> :t the Int
the Int : Int -> Int

Idris> :t the String
the String : String -> String

Idris> :t the Bool
the Bool : Bool -> Bool
```

2.2.4 Writing generic functions with constrained types

The first function you saw in section 2.2.1, double, doubles the Int given as input:

```
double : Int -> Int
double x = x + x
```

But what about other numeric types? For example, you could also write a function to double a Nat, or an Integer:

```
doubleNat : Nat -> Nat
doubleNat x = x + x

doubleInteger : Integer -> Integer
doubleInteger x = x + x
```

As with identity, you're probably starting to see a pattern here, so let's see what happens if we try to replace the input and output types with a variable. Put the following in a file called Generic.idr and load it into Idris:

```
double : ty -> ty
double x = x + x
```

You'll find that Idris rejects this definition, with the following error message:

```
Generic.idr:2:8:
When checking right hand side of double with expected type
    ty
ty is not a numeric type
```

The problem is that, unlike identity, double needs to know something about its input x, specifically that it's numeric. You can only use arithmetic operators on numeric types, so you need to *constrain* ty so that it only stands for numeric types. The following listing shows how you can do this.

Listing 2.3 A generic type, constrained to numeric types (Generic.idr)

```
double : Num ty => ty -> ty
double x = x + x
```

The Num ty before the main part of the function type indicates that ty can only stand for numeric types.

The type Num ty => ty -> ty can be read as, “A function with input type ty and output type ty under the constraint that ty is a numeric type.”

TYPE CONSTRAINTS Constraints on generic types can be user-defined using *interfaces*, which we'll cover in depth in chapter 7. Here, Num is an interface provided by Idris. Interfaces can be given implementations for specific types, and the Num interface has implementations for numeric types.

Perhaps surprisingly, arithmetic and comparison operators aren't primitive operators in Idris, but rather functions with constrained generic types. Infix operators such as +,

`==`, and `<=` are really functions with two inputs, as you can see by checking their types at the REPL:

```
Idris> :t (+)
(+) : Num ty => ty -> ty -> ty

Idris> :t (==)
(==) : Eq ty => ty -> ty -> Bool

Idris> :t (<=)
(<=) : Ord ty => ty -> ty -> Bool
```

As well as `Num` for numeric types, here you can see two other constraints provided by Idris:

- `Eq` states that the type must support the equality and inequality operators, `==` and `/=`.
- `Ord` states that the type must support the comparison operators `<`, `<=`, `>`, and `>=`.

Infix operators and operator sections

Infix operators in Idris aren't a primitive part of the syntax, but are defined by functions. Putting operators in brackets, as with `(+)`, `(==)`, and `(<=)` in the REPL example, means that they'll be treated as ordinary function syntax. For example, you can apply `(+)` to one argument:

```
Idris> :t (+) 2
(+ 2 : Integer -> Integer
```

Infix operators can also be partially applied using *operator sections*:

- `(< 3)` gives a function that returns whether its input is less than 3.
- `(3 <)` gives a function that returns whether 3 is less than its input.

An operator in brackets with only one argument is therefore considered to be a function that expects the other missing argument.

2.2.5 Higher-order function types

There are no restrictions on what the argument or return types of a function can be. You've already seen how functions of multiple arguments are really functions that return something with a function type. Similarly, functions can take functions as arguments. Such functions are called *higher-order* functions.

Higher-order functions can be used to create abstractions for repeated programming patterns. For example, say you've defined a `quadruple` function that quadruples its input for any number, using `double`:

```
quadruple : Num a => a -> a
quadruple x = double (double x)
```

Or say you have a `Shape` type that represents any geometric shape, and a function `rotate : Shape -> Shape` that rotates a shape through 90 degrees. You could define a `turn_around` function that rotates a shape through 180 degrees as follows:

```
turn_around : Shape -> Shape
turn_around x = rotate (rotate x)
```

Each of these functions has exactly the same pattern, but they work on different input types. You can capture this pattern using a higher-order function to apply a function to an argument twice. The next listing gives a definition of a `twice` function, along with new definitions of `quadruple` and `rotate`.

Listing 2.4 Defining `quadruple` and `rotate` using a higher-order function (HOF.idr)

`twice` takes a function as its first argument, and the argument to apply that function to as its second.

```
⇒ twice : (a -> a) -> a -> a
twice f x = f (f x)
```

```
Shape : Type
rotate : Shape -> Shape
```

```
quadruple : Num a => a -> a
quadruple = twice double
```

```
turn_around : Shape -> Shape
turn_around = twice rotate
```

The definition follows exactly the same pattern as the initial definitions of `quadruple` and `turn_around` but with a generic function `f`.

These are type declarations with no definitions.

This implements `quadruple` by directly instantiating “`twice`” with “`double`”.

This implements `turn_around` by directly instantiating `twice` with “`rotate`”.

In chapter 1, I introduced the concept of “holes,” which are incomplete function definitions. The type declarations with no definitions in listing 2.4, `Shape` and `rotate`, are treated as holes. They allow you to try an idea (such as how to implement `turn_around` in terms of `rotate`) without fully defining the types and functions.

Partial application in definitions

In listing 2.4, `quadruple` and `turn_around` have function types, `Num a => a -> a` and `Shape -> Shape`, respectively, but in their definitions neither has an argument.

The only requirement when checking a definition is that both sides of the definition must have the same type. You can check that this is the case here by looking at the types of the left and right sides of the definition at the REPL. You have the following definition:

```
turn_around = twice rotate
```

(continued)

By checking the types at the REPL, you can see that both `turn_around` and `twice_rotate` have the same type:

```
Idris> :t turn_around
turn_around : Shape -> Shape

Idris> :t twice_rotate
twice_rotate : Shape -> Shape
```

The definitions of `quadruple` and `turn_around` both use partial application, as described in section 2.2.2.

Another common use of partial application is in constructing arguments for higher-order functions. Consider this example, using `HOF.idr` and adding the definition of `add` from section 2.2.2:

```
*HOF> twice (add 5) 10
20 : Int
```

This uses a partial application of the `add` function to add 5 to an `Int`, twice. Because `twice` requires a function of *one* argument, and `add` takes *two* arguments, you can apply `add` to one argument so that it's usable in an application of `twice`.

You could also use an operator section, as described at the end of section 2.2.4:

```
*HOF> twice (5 +) 10
20 : Integer
```

Note that, in the absence of any other type information, Idris has defaulted to `Integer`, as described in section 2.1.1.

2.2.6 Anonymous functions

When using higher-order functions, it's often useful to pass an *anonymous* function as an argument. An anonymous function is generally a small function that you only expect to use once, so there's no need to create a top-level definition for it.

For example, you could pass an anonymous function that squares its input to `twice`:

```
*HOF> twice (\x => x * x) 2
16 : Integer
```

Anonymous functions are introduced with a backslash \ followed by a list of arguments. If you check the type of the preceding anonymous function, you'll see that it has a function type:

```
*HOF> :t \x => x * x
\x => x * x : Integer -> Integer
```

Anonymous functions can take more than one argument, and arguments can optionally be given explicit types:

```
*HOF> :t \x : Int, y : Int => x + y
\x, y => x + y : Int -> Int -> Int
```

Note that the output doesn't show the types explicitly.

2.2.7 Local definitions: let and where

As functions get larger, it's usually a good idea to break them down into smaller definitions. Idris provides two constructs for locally defining variables and functions: let and where.

Figure 2.2 illustrates the syntax for let bindings, which define local variables.

If you evaluate this expression at the REPL, you'll see the following:

```
Idris> let x = 50 in x + x
100 : Integer
```

The next listing shows a larger example of let in action. It defines a function, longer, that takes two Strings and returns the length of the longer one. It uses let to record the length of each input.

Listing 2.5 Local variables with let (Let_Where.idr)

```
longer : String -> String -> Nat
longer word1 word2
  = let len1 = length word1
    len2 = length word2 in
      if len1 > len2 then len1 else len2
```

Records the length of the first word
Records the length of the second word
Returns whichever of the lengths is longest.

MULTIPLE LETS There can be several definitions in a let block. In listing 2.5, for example, there are two local variables defined in the let block in longer.

Whereas let blocks contain local variable definitions, where blocks contain local function definitions. Listing 2.6 shows where in action. It defines a function to calculate the length of the hypotenuse of a triangle, using the Pythagorean Theorem and a local square function.

Listing 2.6 Local function definitions with where (Let_Where.idr)

```
pythagoras : Double -> Double -> Double
pythagoras x y = sqrt (square x + square y) ← sqrt is defined in the Prelude.
where
  square : Double -> Double
  square x = x * x
```

This definition is only visible within the scope of pythagoras.




Figure 2.2 A local variable definition: in the expression after the `in` keyword, `x` has the value 50.

Generally, `let` is useful for breaking a complex expression into smaller subexpressions, and `where` is useful for defining more-complex functions that are only relevant in the local context.

2.3 Composite types

Composite types are composed of other types. In this section, we'll look at two of the most common composite types provided by Idris: tuples and lists.

2.3.1 Tuples

A *tuple* is a *fixed-size* collection, where each value in the collection can have a different type. For example, a *pair* of an `Integer` and a `String` can be written as follows:

```
Idris> (94, "Pages")
(94, "Pages") : (Integer, String)
```

Tuples are written as a bracketed, comma-separated list of values. Notice that the type of the pair `(94, "Pages")` is `(Integer, String)`. Tuple types are written using the same syntax as tuple values.

The `fst` and `snd` functions extract the first and second items, respectively, from a pair:

```
Idris> :let mypair = (94, "Pages")
Idris> fst mypair
94 : Integer
Idris> snd mypair
"Pages" : String
```

Both `fst` and `snd` have *generic* types, because pairs can contain *any* types. You can check the type of each at the REPL:

```
Idris> :t fst
fst : (a, b) -> a
Idris> :t snd
snd : (a, b) -> b
```

You can read the type of `fst`, for example, as “Given a pair of an `a` and a `b`, return the value that has type `a`.” In these types, you know that both `a` and `b` are variables because they begin with lowercase letters.

Tuples can have any number of components, including zero:

```
Idris> ('x', 8, String)
('x', 8, String) : (Char, Integer, Type)
Idris> ()
() : ()
```

The empty tuple, `()`, is often referred to as “unit” and its type as “the unit type.” Notice that the syntax is overloaded, and Idris will decide whether `()` means unit or the unit type from the context.

Colors in the REPL

You may have noticed that some values and types in the REPL are colored differently, particularly when evaluating the empty tuple `()`. This is *semantic highlighting*, and it indicates whether a subexpression is a type, value, function, or variable. By default, the REPL displays each as follows:

- Types are blue.
- Values (more precisely, data constructors, as I'll explain in chapter 3) are red.
- Functions are green.
- Variables are magenta.

If these colors aren't to your liking or are hard to distinguish (for example, if you're color blind), you can change the settings with the `:colour` command.

Tuples can also be arbitrarily deeply nested:

```
Idris> (('x', 8), (String, 'y', 100), "Hello")
((('x', 8), (String, 'y', 100), "Hello")
 : ((Char, Integer), (Type, Char, Integer), String))
```

Tuples and pairs

Internally, all tuples other than the empty tuple are stored as nested pairs. That is, if you write `(1, 2, 3, 4)`, Idris will treat this in the same way as `(1, (2, (3, 4)))`. The REPL will always display a tuple in the non-nested form:

```
Idris> (1, (2, (3, 4)))
(1, 2, 3, 4) : (Integer, Integer, Integer, Integer)
```

2.3.2 Lists

Lists, like tuples, are collections of values. Unlike tuples, lists can be any size, but every element must have the same type. Lists are written as comma-separated lists of values in square brackets, as follows.

```
Idris> [1, 2, 3, 4]
[1, 2, 3, 4] : List Integer

Idris> ["One", "Two", "Three", "Four"]
["One", "Two", "Three", "Four"] : List String
```

The type of each of these expressions, `List Integer` and `List String`, indicates the element type that Idris has *inferred* for the list. In type-driven development, we typically give types first, and then write a corresponding value or function that satisfies this type. At the REPL, it would be inconvenient to do this for every value, so Idris will try

to infer a type for the given value. Unfortunately, it's not always possible. For example, if you give it an empty list, Idris doesn't know what the element type should be:

```
Idris> []
(input):Can't infer argument elem to []
```

This error message means that Idris can't work out the element type (which happens to be named `elem`) for the empty list `[]`. The problem can be resolved in this case by giving an explicit type, using the:

```
Idris> the (List Int) []
[] : List Int
```

Like strings, lists can be concatenated with the `++` operator, provided that both operands have the same element type:

```
Idris> [1, 2, 3] ++ [4, 5, 6, 7]
[1, 2, 3, 4, 5, 6, 7] : List Integer
```

You can add an element to the front of a list using the `::` (pronounced “cons”) operator:

```
Idris> 1 :: [2, 3, 4]
[1, 2, 3, 4] : List Integer
Idris> 1 :: 2 :: 3 :: 4 :: []
[1, 2, 3, 4] : List Integer
```

Syntactic sugar for lists

The `::` operator is the primitive operator for constructing lists from a head element and a tail, and `Nil` is a primitive name for the empty list. A list can therefore take one of the following two canonical forms:

- `Nil`, the empty list
- `x :: xs`, where `x` is an element, and `xs` is another list

Because this can get quite verbose, Idris provides *syntactic sugar* for lists. List literals consisting of comma-separated elements inside square brackets are *desugared* to these primitive forms. For example, `[]` is desugared directly to `Nil`, and `[1, 2, 3]` is desugared to `1 :: (2 :: (3 :: Nil))`.

There's also a more concise notation for ranges of numbers. Here are a few examples:

- `[1..5]` expands to the list `[1, 2, 3, 4, 5]`.
- `[1,3..9]` expands to the list `[1, 3, 5, 7, 9]`.
- `[5,4..1]` expands to the list `[5, 4, 3, 2, 1]`.

More generally, `[n..m]` gives an increasing list of numbers between `n` and `m`, and `[n,m..p]` gives a list of numbers between `n` and `p` with the step given by the difference between `n` and `m`.

2.3.3 Functions with lists

We'll discuss lists in more depth in the next chapter, including how to define functions over lists, but there are several useful functions defined in the Prelude. Let's take a look at some of these.

The `words` function, of type `String -> List String`, converts a string into a list of the whitespace-separated components of the string:²

```
Idris> words "'Twas brillig, and the slithy toves"
[''Twas", "brillig,", "and", "the", "slithy", "toves"] : List String
```

The `unwords` function, of type `List String -> String`, does the opposite, converting a list of words into a string where the words are separated by a space:

```
Idris> unwords ["One", "two", "three", "four!"]
"One two three four!" : String
```

You've already seen a `length` function for calculating the lengths of strings. There's also an *overloaded* `length` function, of type `List a -> Nat`, that gives the length of a list:

```
Idris> length ["One", "two", "three", "four!"]
4 : Nat
```

You can use `length` and `words` to write a word-count function for strings:

```
wordCount : String -> Nat
wordCount str = length (words str)
```

The `map` function is a higher-order function that applies a function to every element in a list. It has the type `(a -> b) -> List a -> List b`. This example finds the lengths of every word in a list:

```
Idris> map length (words "How long are these words?")
[3, 4, 3, 5, 6] : List Nat
```

You can use `map` and `length` to write a function that gets the length of every element in a list of Strings:

```
allLengths : List String -> List Nat
allLengths strs = map length strs
```

Type of map

If you check the type of `map` at the REPL, you'll see something slightly different:

```
Idris> :t map
map : Functor f => (a -> b) -> f a -> f b
```

² The `words` name comes from a similar function in the Haskell libraries. Many Idris function names follow Haskell terminology.

(continued)

The reason for this is that `map` can work on a variety of structures, not just lists, so it has a constrained generic type. You'll learn about `Functor` in chapter 7, but for the moment it's fine to read `f` as `List` in this type.

The `filter` function is another higher-order function that filters a list according to a Boolean function. It has type `(a -> Bool) -> List a -> List a` and returns a new list of everything in the input list for which the function returns True. For example, here's how you find all the numbers larger than 10 in a list:

```
Idris> filter (> 10) [1,11,2,12,3,13,4,14]
[11, 12, 13, 14] : List Integer
```

Overloading functions

You've seen the `length` function on both strings and lists. This works because Idris allows function names to be *overloaded* to work on multiple types. You can see what's happening by checking the type of `length` at the REPL:

```
*lists> :t length
Prelude.List.length : List a -> Nat
Prelude.Strings.length : String -> Nat
```

In fact, there are two functions called `length`. The prefixes `Prelude.List` and `Prelude.Strings` are the namespaces these functions are defined in. Idris decides which `length` function is required from the context in which it's used.

The `sum` function, of type `Num a => List a -> a`, calculates the sum of a list of numbers:

```
Idris> sum [1..100]
5050 : Integer
```

The type of `sum` states that every element of the input `List` must have the same type, `a`, and that type is constrained by `Num`.

You now know enough to be able to define a function named `average` that computes the average length of the words in a string. This function is defined in the next listing, which shows the complete Idris file `Average.idr`.

Listing 2.7 Calculating the average word length in a string (Average.idr)

This is a module declaration.

```
module Average
export
average : String -> Double
average str = let numWords = wordCount str
             totalLength = sum (allLengths (words str)) in
```

The `export` keyword means that the definition of `average` is exported from the module.

```

cast totalLength / cast numWords
where
  wordCount : String -> Nat
  wordCount str = length (words str)

  allLengths : List String -> List Nat
  allLengths strs = map length strs

```

A function to calculate the number of words in a string.

The casts are needed to convert Nat to Double.

You only need these definitions in the scope of average, so you put them in a where block.

A function to get the lengths of each word in a list.

Listing 2.7 also introduces the module and export keywords. A module declaration can optionally be put at the top of a file. This declares a namespace in which every function is defined. Conventionally, module names are the same as the filename (without the .idr extension). The export keyword allows the average function to be used by other modules that import Average.idr.

As usual, you can try this at the REPL by loading Average.idr into Idris, and evaluating as follows:

```
*Average> average "How long are these words?"
4.2 : Double
```

Modules and namespaces

By adding a module declaration to the top of Average.idr, you declare a *namespace* for the definitions in the module. Here, it means that the *fully qualified* name of the average function is Average.average. A module declaration must be the first thing in the file. If there's no declaration, Idris calls the module Main.

Modules allow you to divide larger Idris programs logically into several source files, each with their own purpose. They can be imported with an `import` statement. For example:

- `import Average` will import the definitions from Average.idr, provided that Average.idr is either in the current directory or in some other path that Idris can find.
- `import Utils.Average` will import the definitions from Average.idr in a subdirectory called Utils, provided that the file and subdirectory exist.

Modules themselves can be combined into *packages* and distributed separately. Technically, the Prelude is defined in a module called Prelude, which itself imports several other modules, and which is part of a package called prelude. You can learn more about packages and how to create your own from the Idris package documentation at <http://idris-lang.org/documentation/packages>.

2.4 A complete Idris program

So far, you've seen how to write functions with built-in types and some basic operations on those types. Functions are the basic building blocks of Idris programs, so now that you've written some simple functions, it's time to see how to put these together to build a complete program.

2.4.1 Whitespace significance: the layout rule

Whitespace, specifically indentation, is significant in Idris programs. Unlike some other languages, there are no braces or semicolons to indicate where expressions, type declarations, and definitions begin and end. Instead, in any list of definitions and declarations, all must begin in precisely the same column. Listing 2.8 illustrates where definitions and declarations begin and end according to this rule in a file containing the previous definition of `average`.

SPACES AND TABS One complication with whitespace significance is that tab sizes can be set differently in different editors, and Idris expects tabs and spaces to be used consistently. To avoid any confusion with tab sizes, I strongly recommend you set your editor to replace tabs with spaces!

Listing 2.8 The layout rule applied to average (Average.idr)

Type declaration
begins in column 0

```
↳ average : String -> Double
    average str = let numWords = wordCount str
                  totalLength = sum (allLengths (w
                      cast totalLength / cast numWords
    where
        wordCount : String -> Nat
        wordCount str = length (words str)
    allLengths : List String -> List Nat
    allLengths strs = map length strs
```

The type declaration ends, and the definition of wordCount begins in column 4.

The type declaration of wordCount begins in column 4.

Type declaration ends, and a new definition begins in column 0. The definition of numWords begins in column 18.

The definition of numWords ends, and a new definition of totalLength begins in column 18. The definition of totalLength ends at the keyword "in".

The definition ends, and the type declaration of allLengths begins in column 4

If, for example, `allLengths` was indented one extra space, as in listing 2.9, it would be considered a continuation of the previous definition of `wordCount`, and would therefore be invalid.

Listing 2.9 The layout rule, applied incorrectly

```
wordCount : String -> Nat
wordCount str = length (words str)

allLengths : List String -> List Nat
allLengths strs = map length strs
```

allLengths is indented one space too far,
so this line is considered a continuation
of the definition of wordCount.

2.4.2 Documentation comments

As with any other language, it's good form to comment definitions to give the reader of the code an idea of the purposes of functions and document how they work. Idris provides three kinds of comments:

- Single-line comments, introduced with -- (two minus signs). These comments continue to the end of the line.
- Multiline nested comments, introduced with { - and ending with - }.
- Documentation comments, which are used to provide documentation for functions and types at the REPL.

The first two types of comments are conventional, and merely cause the commented section to be ignored (the syntax is identical to the syntax of comments in Haskell).

Documentation comments, on the other hand, make documentation available at the REPL, accessible with the :doc command. You can look at the documentation for some of the types and functions we've encountered so far. For example, :doc fst produces the following output:

```
Idris> :doc fst
Prelude.Basics.fst : (a, b) -> a
    Return the first element of a pair.

The function is Total
```

This output includes the fully qualified name of fst, showing that it is defined in the module Prelude.Basics, and states that the function is total, meaning that it's guaranteed to produce a result for all inputs.

You can also get documentation for types. For example, :doc List gives the following output:

```
Idris> :doc List
Data type Prelude.List.List : Type -> Type
    Generic lists

Constructors:
    Nil : List elem
        The empty list

    (::) : elem -> List elem -> List elem
        A non-empty list, consisting of a head element and the rest of
        the list.
    infixr 7
```

Again, this gives the fully qualified name of the type `Prelude.List.List`. It also gives the *constructors*, which are the primitive ways of constructing lists. Finally, for the `::` operator, it gives the *fixity*, which states the operator is *right associative* (`infixr`) and has precedence level 7. I'll describe precedence and the associativity of operators in more detail in chapter 3.

Documentation comments, which produce this documentation, are introduced with three vertical bars, `|||`. For example, you could document `average` as follows:

```
||| Calculate the average length of words in a string.
average : String -> Double
```

Then, `:doc average` would produce the following output:

```
*Average> :doc average
Average.average : (str : String) -> Double
Calculate the average length of words in a string.

The function is Total
```

You can refer to the argument of `average` by giving it a name, `str`, and referring to that name in the comment with `@str`:

```
||| Calculate the average length of words in a string.
||| @str a string containing words separated by whitespace.
average : (str : String) -> Double
average str = let numWords = wordCount str
              totalLength = sum (allLengths (words str)) in
              cast totalLength / cast numWords
```

This makes `:doc average` produce some more informative output:

```
*Average> :doc average
Main.average : (str : String) -> Double
Calculate the average length of words in a string.
Arguments:
  str : String -- a string containing words separated by whitespace.

The function is Total
```

TOTALITY CHECKING Notice that `:doc average` reports that `average` is total. Idris checks every definition for totality. The result of totality checking has several interesting implications in type-driven development, which we'll discuss throughout the book, and particularly in chapters 10 and 11.

2.4.3 *Interactive programs*

The entry point to a compiled Idris program is the `main` function, defined in a `Main` module. That is, it's the function with the fully qualified name `Main.main`. It must have the type `IO ()`, meaning that it returns an `IO` action that produces an empty tuple.

You've already seen the “Hello, Idris World!” program:

```
main : IO ()
main = putStrLn "Hello Idris World!"
```

Here, `putStrLn` is a function of type `String -> IO ()` that takes a string as an argument and returns an IO action that outputs that string. We'll discuss IO actions in some depth in chapter 5, but even before then you'll be able to write complete interactive Idris programs using the `repl` function (and some variants on it, as you'll see in chapter 4) provided by the Prelude:

```
Idris> :doc repl
Prelude.Interactive.repl : (prompt : String) ->
    A basic read-eval-print loop
    Arguments:
        prompt : String -- the prompt to show
        onInput : String -> String -- the function to run on reading
            input, returning a String to output
```

This allows you to write programs that repeatedly display a prompt, read some input, and produce some output by running a function of type `String -> String` on it. For example, the next listing is a program that repeatedly reads a string and then prints the string in reverse.

Listing 2.10 Reversing strings interactively (Reverse.idr)

```
module Main

main : IO ()
main = repl "> " reverse
```

You can compile and run this program at the REPL with the `:exec` command. Note that the program will loop indefinitely, but you can quit by interrupting the program with Ctrl-C:

```
*reverse> :exec
> hello!
!olleh> goodbye
eybdoog>
```

To conclude this chapter, you'll write a program that imports the `Average` module, reads a string from the console, and displays the average number of letters in each word in the string. One difficulty is that the `average` function returns a `Double`, but `repl` requires a function of type `String -> String`, so you can't use `average` directly. In general, though, values can be converted to `String` using the `show` function. Let's take a look at it using `:doc`:

```
Idris> :doc show
Prelude.Show.show : Show ty => (x : ty) -> String
    Convert a value to its String representation.
```

Note that this is a *constrained* generic type, meaning that the type `ty` must support the `Show` interface, which is true of all the types in the Prelude.

Using this, you can write a `showAverage` function that uses `average` to get the average word length and displays it in a nicely formatted string. The complete program is given in the following listing.

Listing 2.11 Displaying average word lengths interactively (AveMain.idr)

```
module Main
import Average

showAverage : String -> String
showAverage str = "The average word length is: " ++
    show (average str) ++ "\n"

main : IO ()
main = repl "Enter a string: "
    showAverage
```

The code is annotated with several callouts:

- A bracket labeled "Imports the definitions from Average.idr" points to the `import Average` line.
- A bracket labeled "The function that calculates the string to output from some input. It uses show to convert the result of average to a String." points to the `showAverage` definition.
- A bracket labeled "The function to call to calculate the output" points to the `showAverage` call in the `main` function.
- A bracket labeled "The prompt to display" points to the `repl` call in the `main` function.

Again, you can use `:exec` to compile and run this at the REPL, and then try some inputs:

```
*AveMain> :exec
Enter a string: The quick brown fox jumped over the lazy dog
The average word length is: 4
Enter a string: The quick brown fox jumped over the lazy frog
The average word length is: 4.11111
Enter a string:
```

Exercises



- What are the types of the following values?

- `("A", "B", "C")`
- `["A", "B", "C"]`
- `(('A', "B"), 'C')`

You can check your answers with `:t` at the REPL, but try to work them out yourself first.

- Write a `palindrome` function, of type `String -> Bool`, that returns whether the input reads the same backwards as forwards.

Hint: You may find the function `reverse : String -> String` useful.

You can test your answer at the REPL as follows:

```
*ex_2> palindrome "racecar"
True : Bool

*ex_2> palindrome "race car"
False : Bool
```

- 3 Modify the palindrome function so that it's not case sensitive.

Hint: You may find `toLower : String -> String` useful.

You can test your answer at the REPL as follows:

```
*ex_2> palindrome "Racecar"
True : Bool
```

- 4 Modify the palindrome function so that it only returns True for strings longer than 10 characters.

You can test your answer at the REPL as follows:

```
*ex_2> palindrome "racecar"
False : Bool

*ex_2> palindrome "able was i ere i saw elba"
True : Bool
```

- 5 Modify the palindrome function so that it only returns True for strings longer than some length given as an argument.

Hint: Your new function should have type `Nat -> String -> Bool`.

You can test your answer at the REPL as follows:

```
*ex_2> palindrome 10 "racecar"
False : Bool

*ex_2> palindrome 5 "racecar"
True : Bool
```

- 6 Write a `counts` function of type `String -> (Nat, Nat)` that returns a pair of the number of words in the input and the number of characters in the input.

You can test your answer at the REPL as follows:

```
*ex_2> counts "Hello, Idris world!"
(3, 19) : (Nat, Nat)
```

- 7 Write a `top_ten` function of type `Ord a => List a -> List a` that returns the ten largest values in a list. You may find the following Prelude functions useful:

- `take : Nat -> List a -> List a`
- `sort : Ord a => List a -> List a`

Use `:doc` for further information about these functions if you need it.

You can test your answer at the REPL as follows:

```
*ex_2> top_ten [1..100]
[100, 99, 98, 97, 96, 95, 94, 93, 92, 91] : List Integer
```

- 8 Write an `over_length` function of type `Nat -> List String -> Nat` that returns the number of strings in the list longer than the given number of characters.

You can test your answer at the REPL as follows:

```
*ex_2> over_length 3 ["One", "Two", "Three", "Four"]
2 : Nat
```

- 9 For each of `palindrome` and `counts`, write a complete program that prompts for an input, calls the function, and prints its output.

You can test your answer using `:exec` at the REPL:

```
*ex_2_palindrome> :exec
Enter a string: Able was I ere I saw Elba
True
Enter a string: Madam, I'm Adam
False
Enter a string:
```

2.5 Summary

- The Prelude defines a number of basic types and functions and is imported automatically by all Idris programs.
- Idris provides basic numeric types, `Int`, `Integer`, `Nat`, and `Double`, as well as a Boolean type, `Bool`, a character type, `Char`, and a string type, `String`.
- Values can be converted between compatible types using the `cast` function, and can be given explicit types with the `the` function.
- Tuples are fixed-size collections where each element can be a different type.
- Lists are variable size collections where each element has the same type.
- Function types have one or more input types and one output type.
- Function types can be generic, meaning that they can contain variables. These variables can be constrained to allow a smaller set of types.
- Higher-order functions are functions in which one of the arguments is itself a function.
- Functions consist of a required type declaration and a definition. Function definitions are equations defining rewrite rules to be used during evaluation.
- Whitespace is significant in Idris programs. Each definition in a block must begin in exactly the same column.
- Function documentation can be accessed at the REPL with the `:doc` command.
- Idris programs can be divided into separate source files called modules.
- The entry point to an Idris program is the `main` function, which must have type `IO ()`, and be defined in the module `Main`. Simple interactive programs can be written by applying the `repl` function from `main`.

Part 2

Core Idris

N

ow that you have some experience writing programs in Idris, it's time to start your exploration of type-driven development in depth. In this part, you'll learn about the core features of Idris and gain some experience in the process of type-driven development. Rather than showing you complete programs right from the start, I'll show you how to build programs interactively, via a process of type, define, refine:

- *Type*—Write a type for a function.
- *Define*—Create an initial definition for the function, possibly containing holes.
- *Refine*—Complete the definition by filling in holes, possibly modifying the type as your understanding of the problem develops.

In chapter 3, you'll learn the basics of interactive development; and then, in chapter 4, you'll learn to define your own data types and build larger programs around them. Chapter 5 shows how you can write programs that interact with the outside world, using types to separate evaluation from execution. Later chapters introduce more-advanced concepts in type-driven development, including type-level computation in chapter 6, working with constrained generic types in chapter 7, describing and proving properties of programs in chapters 8 and 9, and defining alternative traversals of data structures using views in chapter 10.

By the end of part 2, you'll have learned about all of the core features of Idris.

Interactive development with types

This chapter covers

- Defining functions by pattern matching
- Type-driven interactive editing in Atom
- Adding precision to function types
- Practical programming with vectors

You've now seen how to define simple functions, and how to structure these into complete programs. In this chapter, we'll start a deeper exploration into type-driven development. First, we'll look at how to write more-complex functions with existing types from the Prelude, such as lists. Then, we'll look at using the Idris type system to give functions more-precise types.

In type-driven development, we follow the process of "type, define, refine." You'll see this process in action throughout this chapter as you first write the types and, as far as possible, always have a type-correct, if perhaps incomplete, definition of a function and refine it step by step until it's complete. Each step will be broadly characterized as one of these three:

- *Type*—Either write a type to begin the process, or inspect the type of a hole to decide how to continue the process.

- *Define*—Create the structure of a function definition either by creating an outline of a definition or breaking it down into smaller components.
- *Refine*—Improve an existing definition either by filling in a hole or making its type more precise.

In this chapter, I'll introduce interactive development in the Atom text editor, which assists with this process. Atom provides an interactive editing mode that communicates with a running Idris system and uses types to help direct function development. Atom also provides some structural editing features and contextual information about functions with holes, and, when the types are precise enough, even completes large parts of functions for you. I'll begin this chapter, therefore, by introducing the interactive editing mode in Atom.

EDITOR MODES Although we'll use Atom to edit Idris programs, the interactive features we'll use are provided by Idris itself. The Atom integration works by communicating with an Idris process running in the background. This process is run as a child of the editor, so it's independent of any REPLs you may have running. As such, it's reasonably straightforward to add Idris support to other text editors, and similar editing modes currently exist for Emacs and Vim. In this book, we'll stick to Atom for consistency, but each of the commands directly maps to corresponding commands in other editors.

3.1 *Interactive editing in Atom*

You saw in chapter 1 that Idris programs can contain holes, which stand for parts of a function definition that haven't yet been written. This is one way in which programs can be developed interactively: you write an incomplete definition containing holes, check the types of the holes to see what Idris expects in each, and then continue by filling in the holes with more code. There are several additional ways in which Idris can help you by integrating interactive development features with a text editor:

- *Add definitions*—Given a type declaration, Idris can add a skeleton definition of a function that satisfies that type.
- *Case analysis*—Given a skeleton function definition with arguments, Idris can use the types of those arguments to help define the function by pattern matching.
- *Expression search*—Given a hole with a precise enough type, Idris can try to find an expression that satisfies the hole's type, refining the definition.

In this section, we'll begin writing some more-complex Idris functions, using its interactive editing features to develop those functions, step by step, in a type-directed way. We'll use the Atom text editor, because there's an extension available for editing Idris programs, which can be installed directly from the default Atom distribution. The rest of this chapter assumes that you have the interactive Idris mode up and running. If not, follow the instructions in appendix A for installing Atom and the Idris mode.

3.1.1 Interactive command summary

Interactive editing in Atom involves a number of keyboard commands in the editor, which are summarized in table 3.1.

COMMAND MNEMONICS For each command, the shortcut in Atom is to press Ctrl, Alt, and the first letter of the command.

Table 3.1 Interactive editing commands in Atom

Shortcut	Command	Description
Ctrl-Alt-A	Add definition	Adds a skeleton definition for the name under the cursor
Ctrl-Alt-C	Case split	Splits a definition into pattern-matching clauses for the name under the cursor
Ctrl-Alt-D	Documentation	Displays documentation for the name under the cursor
Ctrl-Alt-L	Lift hole	Lifts a hole to the top level as a new function declaration
Ctrl-Alt-M	Match	Replaces a hole with a case expression that matches on an intermediate result
Ctrl-Alt-R	Reload	Reloads and type-checks the current buffer
Ctrl-Alt-S	Search	Searches for an expression that satisfies the type of the hole name under the cursor
Ctrl-Alt-T	Type-check name	Displays the type of the name under the cursor

3.1.2 Defining functions by pattern matching

All of the function definitions we've looked at so far have involved a single equation to define that function's behavior. For example, in the previous chapter you wrote a function to calculate the lengths of every word in a list:

```
allLengths : List String -> List Nat
allLengths strs = map length strs
```

Here, you used functions defined in the Prelude (`map` and `length`) to inspect the list. At some stage, though, you're going to need a more direct way to inspect values. After all, functions like `map` and `length` need to be defined themselves somehow!

In general, you define functions by pattern matching on the possible values of inputs to a function. For example, you can define a function to invert a `Bool` as follows:

```
invert : Bool -> Bool
invert False = True
invert True = False
```

The possible inputs of type `Bool` are `True` and `False`, so here you implement `invert` by listing the possible inputs and giving the corresponding outputs. Patterns can also contain variables, as illustrated by the following function, which returns "Empty" if

given an empty list or "Non-empty" followed by the value of the tail if given a non-empty list:

```
describeList : List Int -> String
describeList [] = "Empty"
describeList (x :: xs) = "Non-empty, tail = " ++ show xs
```

Figure 3.1 illustrates how the patterns are matched in `describeList` for the inputs `[1]` (which is syntactic sugar for `1 :: []`) and `[2,3,4,5]` (which is syntactic sugar for `2 :: 3 :: 4 :: 5 :: []`).




Figure 3.1 Matching the pattern `(x :: xs)` for inputs `[1]` and `[2,3,4,5]`

NAMING CONVENTIONS FOR LISTS Conventionally, when working with lists and list-like structures, Idris programmers use a name ending in “s” (to suggest a plural), and then use its singular to refer to individual elements. So, if you have a list called `things`, you might refer to an element of the list as `thing`.

A function definition consists of one or more equations matching the possible inputs to that function. You can see how this works for lists if you implement `allLengths` by inspecting the list directly, instead of using `map`.

Listing 3.1 Calculating word lengths by pattern matching on the list (WordLength.idr)

```
allLengths : List String -> List Nat
allLengths [] = []
allLengths (x :: xs) = length x :: allLengths xs
```

If the input list is empty, the output list will also be empty.

If the input list has a head element `x` and a tail `xs`, the output list will have the length of `x` as the head and then, recursively, a list of the lengths of `xs` as the tail.

To see in detail how this definition is constructed, you’ll build it interactively in Atom. Each step can be broadly characterized as *Type* (creating or inspecting a type), *Define* (making a definition or breaking it down into separate clauses), or *Refine* (improving a definition by filling in a hole or making its type more precise).

- 1 **Type**—First, start up Atom and create a `WordLength.idr` file containing the type declaration for `allLengths`. That’s *only* the following:

```
allLengths : List String -> List Nat
```

You should also start up an Idris REPL in a separate terminal so that you can type-check and test your definition:

```
$ idris WordLength.idr
```

```
/ \ / \ / \ / \ ( ) _____  
/ / / / / \ / \ / \ ( ) _____  
_ / / / / \ / \ / \ / \ ( ) _____  
/ \ / \ / \ / \ / \ / \ / \ _____  
  
Version 1.0  
http://www.idris-lang.org/  
Type :? for help  
  
Idris is free software with ABSOLUTELY NO WARRANTY.  
For details type :warranty.  
Holes: Main.allLengths  
*WordLength>
```

- 2 Define**—In Atom, move the cursor over the name `allLengths` and press `Ctrl-Alt-A`. This will add a skeleton definition, and your editor buffer should now contain the following:

```
allLengths : List String -> List Nat  
allLengths xs = ?allLengths rhs
```

The skeleton definition is always a clause with the appropriate number of arguments listed on the left side of the `=`, and with a hole on the right side. Idris uses various heuristics to choose initial names for the arguments. By convention, Idris chooses default names of `xs`, `ys`, or `zs` for Lists.

- 3 **Type**—You can check the types of holes in Atom by pressing Ctrl-Alt-T with the cursor over the hole you want to check. If you check the type of the `allLengths` `rhs` hole, you should see this:

```
xs : List String  
-----  
allLengths rhs : List Nat
```

- 4 **Define**—You'll write this definition by inspecting the list argument, named `xs`, directly. This means that for every form the list can take, you need to explain how to measure word lengths when it's in that form. To tell the editor that you want to inspect the first argument, press Ctrl-Alt-C in Atom with the cursor over the variable `xs` in the first argument position. This expands the definition to give the two forms that the argument `xs` could take:

```
allLengths : List String -> List Nat
allLengths [] = ?allLengths_rhs_1
allLengths (x :: xs) = ?allLengths_rhs_2
```

These are the two *canonical forms* of a list. That is, *every* list must be in one of these two forms: it can either be empty (in the form `[]`), or it can be non-empty, containing a head element and the rest of the list (in the form `(x :: xs)`). It's a good idea at this point to rename `x` and `xs` to something more meaningful than these default names:

```
allLengths : List String -> List Nat
allLengths [] = ?allLengths_rhs_1
allLengths (word :: words) = ?allLengths_rhs_2
```

In each case, there's a new hole on the right side to fill in. You can check the types of these holes; type checking gives the expected return type and the types of any local variables. For example, if you check the type of `allLengths_rhs_2`, you'll see the types of the local variables `word` and `words`, as well as the expected return type:

```
word : String
words : List String
-----
allLengths_rhs_2 : List Nat
```

- 5 *Refine*—Idris has now told you which patterns are needed. Your job is to complete the definition by filling in the holes on the right side. In the case where the input is the empty list, the output is also the empty list, because there are no words for which to measure the length:

```
allLengths [] = []
```

- 6 *Refine*—In the case where the input is non-empty, there's a word as the first element (`word`) followed by the remainder of the list (`words`). You need to return a list with the length of `word` as its first element. For the moment, you can add a new hole (`?rest`) for the remainder of the list:

```
allLengths : List String -> List Nat
allLengths [] = []
allLengths (word :: words) = length word :: ?rest
```

You can even test this incomplete definition at the REPL. Note that the REPL doesn't reload files automatically, because it runs independently of the interactive editing in Atom, so you'll need to reload explicitly using the `:r` command:

```
*WordLength> :r
Type Checking ./WordLength.idr
Holes: Main.rest
*WordLength> allLengths ["Hello", "Interactive", "Editors"]
5 :: ?rest : List Nat
```

For the hole `rest`, you need to calculate the lengths of the words in `words`. You can do this with a recursive call to `allLengths`, to complete the definition:

```
allLengths : List String -> List Nat
allLengths [] = []
allLengths (word :: words) = length word :: allLengths words
```

You now have a complete definition, which you can test at the REPL after reloading:

```
*WordLength> :r
Type Checking ./WordLength.idr
*WordLength> allLengths ["Hello", "Interactive", "Editors"]
[5, 11, 7] : List Nat
```

It's also a good idea to check whether Idris believes the definition is total.

```
*WordLength> :total allLengths
Main.allLengths is Total
```

Totality checking

When Idris has successfully type-checked a function, it also checks whether it believes the function is total. If a function is total, it's guaranteed to produce a result for any well-typed input, in finite time. Thanks to the halting problem, which we discussed in chapter 1, Idris can't decide in general whether a function is total, but by analyzing a function's syntax, it can decide that a function is total in many specific cases.

We'll discuss totality checking in much more detail in chapters 10 and 11. For the moment, it's sufficient to know that Idris will consider a function total if

- It has clauses that cover all possible well-typed inputs
- All recursive calls converge on a *base case*

As you'll see in chapter 11 in particular, the definition of totality also allows interactive programs that run indefinitely, such as servers and interactive loops, provided they continue to produce intermediate results in finite time.

Idris believes that `allLengths` is total because there are clauses for all possible well-typed inputs, and the argument to the recursive call to `allLengths` is smaller (that is, closer to the base case) than the input.

3.1.3 Data types and patterns

When you press Ctrl-Alt-C in Atom with the cursor over a variable on the left side of a definition, it performs a *case split* on that variable, giving the possible patterns the variable can match. But where do these patterns come from?

Each data type has one or more *constructors*, which are the primitive ways of building values in that data type and give the patterns that can be matched for that data type. For `List`, there are two:

- `Nil`, which constructs an empty list
- `::`, an infix operator that constructs a list from a head element and a tail

Additionally, as you saw in chapter 2, there's syntactic sugar for lists that allows a list to be written as a comma-separated list of values in square brackets. Hence, `Nil` can also be written as `[]`.

For any data type, you can find the constructors and thus the patterns to match using `:doc` at the REPL prompt:

```
Idris> :doc List
Data type Prelude.List.List : (elem : Type) -> Type
  Generic lists

Constructors:
  Nil : List elem
    Empty list
  (::) : (x : elem) -> (xs : List elem) -> List elem
```

```
A non-empty list, consisting of a head element and the rest of
the list.
infixr 7
```

DOCUMENTATION IN ATOM You can get documentation directly in Atom by pressing Ctrl-Alt-D, with the cursor over the name for which you want documentation.

For `Bool`, for example, `:doc` reveals that the constructors are `False` and `True`:

```
Idris> :doc Bool
Data type Prelude.Bool.Bool : Type
    Boolean Data Type

Constructors:
    False : Bool
    True : Bool
```

Therefore, if you write a function that takes a `Bool` as an input, you can provide explicit cases for the inputs `False` and `True`.

For example, to write the exclusive OR operator, you could follow these steps:

- 1 *Type*—Start by giving a type:

```
xor : Bool -> Bool -> Bool
```

- 2 *Define*—Press Ctrl-Alt-A with the cursor over `xor` to add a skeleton definition:

```
xor : Bool -> Bool -> Bool
xor x y = ?xor_rhs
```

- 3 *Define*—Press Ctrl-Alt-C over the `x` to give the two possible cases for `x`:

```
xor : Bool -> Bool -> Bool
xor False y = ?xor_rhs_1
xor True y = ?xor_rhs_2
```

- 4 *Refine*—Complete the definition by filling in the right sides:

```
xor : Bool -> Bool -> Bool
xor False y = y
xor True y = not y
```

TYPE CHECKING IN ATOM While developing a function, and particularly when writing clauses by hand rather than using interactive editing features, it can be a good idea to type-check what you have so far. The Ctrl-Alt-R command rechecks the current buffer using the running Idris process. If it loads successfully, the Atom status bar will report “File loaded successfully.”

The `Nat` type, which represents unbounded unsigned integers, is also defined by primitive constructors. In Idris, a natural number is defined as being either zero, or one more than (that is, the *successor* of) another natural number.

```
Idris> :doc Nat
Data type Prelude.Nat.Nat : Type
    Natural numbers: unbounded, unsigned integers which can be
    pattern matched.
```

Constructors:

```
Z : Nat
Zero

S : Nat -> Nat
Successor
```

DATA TYPES AND CONSTRUCTORS Data types are defined in terms of their constructors, as you'll see in detail in chapter 4. The constructors of a data type are the primitive ways of building that data type, so in the case of Nat, every value of type Nat must be either zero or the successor of another Nat. The number 3, for example, is written in primitive form as S (S (S Z)). That is, it's the successor (S) of 2 (written as S (S Z)).

Therefore, if you write a function that takes a Nat as an input, you can provide explicit cases for the number zero (Z) or a number greater than zero (S k, where k is any non-negative number). For example, to write an isEven function that returns True if a natural number input is divisible by 2, and returns False otherwise, you could define it recursively (if inefficiently) as follows:

- 1 *Type*—Start by giving a type:

```
isEven : Nat -> Bool
```

- 2 *Define*—Press Ctrl-Alt-A to add a skeleton definition:

```
isEven : Nat -> Bool
isEven k = ?isEven_rhs
```

NAMING CONVENTIONS As a naming convention, Idris chooses k by default for variables of type Nat. Naming conventions can be set by the programmer when defining data types, and you'll see how to do this in chapter 4. In any case, it's usually a good idea to rename these variables to something more informative.

- 3 *Define*—Press Ctrl-Alt-C over the k to give the two possible cases for k:

```
isEven : Nat -> Bool
isEven Z = ?isEven_rhs_1
isEven (S k) = ?isEven_rhs_2
```

To complete the definition, you have to explain what to return when the input is zero (Z) or when the input is non-zero (if the input takes the form S k, then k is a variable standing for a number that's one smaller than the input).

- 4 *Refine*—Complete the definition by filling in the right sides:

```
isEven : Nat -> Bool
isEven Z = True
isEven (S k) = not (isEven k)
```

You've defined this recursively. Zero is an even number, so you return True for the input Z. If a number is even, its successor is odd, and vice versa, so you return not (isEven k) for the input S k.

Mutually defined functions

Idris processes input files from top to bottom and requires types and functions to be defined before use. This is necessary due to complications that arise with dependent types, where the definition of a function can affect a type.

It's nevertheless sometimes useful to define two or more functions in terms of each other. This can be achieved in a `mutual` block. For example, you could define `isEven` in terms of an `isOdd` function, and vice versa:

```
mutual
    isEven : Nat -> Bool
    isEven Z = True
    isEven (S k) = isOdd k

    isOdd : Nat -> Bool
    isOdd Z = False
    isOdd (S k) = isEven k
```

3.2 Adding precision to types: working with vectors

In chapter 1, we discussed how having types as a first-class language construct allows us to define more-precise types. As one example, you saw how lists could be given more-precise types by including the number of elements in a list in its type, as well as the type of the elements. In Idris, a list that includes both the number and the type of elements in its type is called a *vector*, defined as the data type `Vect`. The following listing shows some example vectors.

Listing 3.2 Vectors: lists with lengths encoded in the type (Vectors.idr)

```
import Data.Vect

fourInts : Vect 4 Int
fourInts = [0, 1, 2, 3]

sixInts : Vect 6 Int
sixInts = [4, 5, 6, 7, 8, 9]

tenInts : Vect 10 Int
tenInts = fourInts ++ sixInts
```

Appending vectors using `++`
also adds their lengths in the
type of the result.

`Vect` isn't defined in the Prelude, but it can be made available by importing the `Data.Vect` library module. Modules are imported with an `import` statement at the top of a source file:

```
import Data.Vect
```

PACKAGES: PRELUDE AND BASE Idris modules can be combined into *packages*, from which individual modules can be imported. The Prelude is defined in a package called `prelude`, from which all modules are imported automatically. Idris programs also have access to a package called `base`, which defines

several commonly useful data structures and algorithms including `Vect`, but from which modules must be imported explicitly. Up-to-date documentation for the packages distributed with Idris is available at www.idris-lang.org/documentation.

3.2.1 Refining the type of `allLengths`

To see how `Vect` works, you can refine the type of the `allLengths` function from section 3.1.2 to use a `Vect` instead of a `List`, and redefine the function.

To do so, create a new source file called `WordLength_vec.idr` containing only the line `import Data.Vect`, and load it into the REPL. You can then check the documentation for `Vect`:

```
*WordLength_vec> :doc Vect
Data type Data.Vect.Vect : Nat -> Type -> Type
  Vectors: Generic lists with explicit length in the type

Constructors:
  Nil : Vect 0 a
    Empty vector

  (::) : (x : a) -> (xs : Vect k a) -> Vect (S k) a
    A non-empty vector of length S k, consisting of a head element
    and the rest of the list, of length k.
  infixr 7
```

Notice that it has the same constructors as `List`, but they have different types that give explicit lengths. Lengths are given as `Nat`, because they can't be negative:

- The type of `Nil` explicitly states that the length is `z` (displayed as the numeric literal `0` here).
- The type of `::` explicitly states that the length is `S k`, given an element and a tail of length `k`.

The same syntactic sugar applies as for `List`, translating a bracketed list of values such as `[1, 2, 3]` to a sequence of `::` and `Nil`: `1 :: 2 :: 3 :: Nil`, in this case. In fact, this syntactic sugar applies to any data type with constructors called `Nil` and `::`.

Overloading names

The constructor names for both `List` and `Vect` are the same, `Nil` and `::`. Names can be overloaded, provided that different things with the same name are defined in separate namespaces, which in practice usually means separate modules. Idris will infer the appropriate namespace from the context in which the name is used.

You can explicitly indicate which of `List` or `Vect` is required using the:

```
Idris> the (List _) ["Hello", "There"]
["Hello", "There"] : List String

Idris> the (Vect _ _) ["Hello", "There"]
["Hello", "There"] : Vect 2 String
```

(continued)

The underscore (`_`) in the preceding expressions indicates to Idris that you would like it to infer a value for that argument. You can use `_` in an expression whenever there's only one valid value that would stand for that expression. You'll see more about this in section 3.4.

To define `allLengths` using a `Vect`, you can follow a similar procedure as when defining it using `List`. The difference is that you must consider how the lengths of the input and output are related.

Figure 3.2 shows how, in the output list, there's always an entry corresponding to an entry in the input list. Therefore, you can make it explicit in the type that the output vector has the same length as the input vector:

```
allLengths : Vect len String -> Vect len Nat
```




Figure 3.2 Calculating the lengths of words in a list. Notice that each input has a corresponding output, so the length of the output vector is always the same as the length of the input vector.

The `len` that appears in the input is a variable at the type level, standing for the length of the input. Because the output uses the same variable at the type level, it's explicit in the type that the output has the same length as the input. Here's how you can write the function:

- 1 *Type*—Create an Atom buffer with the following contents to import `Data.Vect` and give the type of `allLengths`:

```
import Data.Vect
allLengths : Vect len String -> Vect len Nat
```

- 2 *Define*—As before, create a skeleton definition with `Ctrl-Alt-A`:

```
allLengths : Vect len String -> Vect len Nat
allLengths xs = ?allLengths_rhs
```

- 3 *Define*—As before, press `Ctrl-Alt-C` over the `xs` to tell Idris that you'd like to define the function by pattern matching on `xs`:

```
allLengths : Vect len String -> Vect len Nat
allLengths [] = ?allLengths_rhs_1
allLengths (x :: xs) = ?allLengths_rhs_2
```

As before, it's a good idea at this point to rename the variables `x` and `xs` to something more meaningful:

```
allLengths : Vect len String -> Vect len Nat
```

```
allLengths [] = ?allLengths_rhs_1
allLengths (word :: words) = ?allLengths_rhs_2
```

- 4 Type**—If you now inspect the types of the holes `allLengths_rhs_1` and `allLengths_rhs_2`, you'll see more information than in the List version, because the types are more precise. For example, in `allLengths_rhs_1` you can see that the only valid result is a vector with zero elements:

```
-----
allLengths_rhs_1 : Vect 0 Nat
```

In `allLengths_rhs_2` you can see how the lengths of the pattern variables and output relate to each other, given some natural number `n`:

```
word : String
k : Nat
words : Vect k String
-----
allLengths_rhs_2 : Vect (S k) Nat
```

That is, in the pattern `(word :: words)`, `word` is a `String`, `words` is a vector of `k` `Strings`, and for the output you need to provide a vector of `Nat` of length $1 + k$, represented as `S k`.

- 5 Refine**—For `allLengths_rhs_1`, there's only one vector that has length zero, which is the empty vector, so there's only one value you can use to refine the definition by filling in the hole:

```
allLengths : Vect len String -> Vect len Nat
allLengths [] = []
allLengths (word :: words) = ?allLengths_rhs_2
```

- 6 Refine**—For `allLengths_rhs_2`, the required type is `Vect (S k) Nat`, so the result must be a non-empty vector (using `::`). Also, you can make a value of type `Vect k Nat` by calling `allLengths` recursively. You can leave a hole for the first element in the resulting list, refining the definition by hand as follows:

```
allLengths : Vect len String -> Vect len Nat
allLengths [] = []
allLengths (word :: words) = ?wordlen :: allLengths words
```

You still have one hole in this result, `?wordlen`, which will be the length of the first word:

```
word : String
k : Nat
words : Vect k String
-----
wordlen : Nat
```

- 7 Refine**—To complete the definition, fill in the remaining hole by calculating the length of the word `x`:

```
allLengths : Vect len String -> Vect len Nat
allLengths [] = []
allLengths (word :: words) = length word :: allLengths words
```

The more precise type, describing how the lengths of the input and output relate, means that the interactive editing mode can tell you more about the expressions you're looking for. You can also be more confident that the program behaves as intended by ruling out any program that doesn't preserve length by type-checking.

NAT AND DATA STRUCTURES You might notice a direct correspondence between the constructors of Vect and the constructors of Nat. When you add an element to a Vect with `::`, you add an S constructor to its length. In practice, capturing the size of data structures like this is a very common use of Nat.

To illustrate how the more precise type rules out some incorrect programs, consider the following implementation of `allLengths`, using List instead of Vect:

```
allLengths : List String -> List Nat
allLengths xs = []
```

This is well typed and it will be accepted by Idris, but it won't work as intended because there's no guarantee that the output list has an entry corresponding to each entry in the input. On the other hand, the following program with the more precise type is *not* well typed and will not be accepted by Idris:

```
allLengths : Vect n String -> Vect n Nat
allLengths xs = []
```

This results in the following type error, which states that an empty vector was given when a vector of length `n` was needed:

```
WordLength_vec.idr:4:14:When checking right hand side of allLengths:
Type mismatch between
    Vect 0 Nat (Type of [])
and
    Vect n Nat (Expected type)
```

As with the previous List-based version of `allLengths`, you can check that your new definition is total at the REPL:

```
*WordLength_vec> :total allLengths
Main.allLengths is Total
```

If, for example, you remove the case for the empty list, you have a definition that is well typed but partial:

```
allLengths : Vect len String -> Vect len Nat
allLengths (word :: words) = length word :: allLengths words
```

When you check this for totality, you'll see this:

```
*WordLength_vec> :total allLengths
Main.allLengths is not total as there are missing cases
```

Totality annotations

For added confidence in a function's correctness, you can annotate in the source code that a function must be total. For example, you can write this:

```
total allLengths : Vect len String -> Vect len Nat
allLengths [] = []
allLengths (word :: words) = length word :: allLengths words
```

The `total` keyword before the type declaration means that Idris will report an error if the definition isn't total. For example, if you remove the `allLengths []` case, this is what Idris will report:

```
WordLength_vec.idr:5:1:
Main.allLengths is not total as there are missing cases
```

3.2.2 Type-directed search: automatic refining

After step 3 in the previous section, you had the patterns for `allLengths` and holes for the right sides, which you provided by refining directly:

```
allLengths : Vect n String -> Vect n Nat
allLengths [] = ?allLengths_rhs_1
allLengths (word :: words) = ?allLengths_rhs_2
```

Take another look at the types and the local variables for the `allLengths_rhs_1` and `allLengths_rhs_2` holes:

```
-----
allLengths_rhs_1 : Vect 0 Nat
  word : String
  k : Nat
  words : Vect k String
-----
allLengths_rhs_2 : Vect (S k) Nat
```

By looking carefully at the types, you could see how to construct values to fill these holes. But not only do you have more information here, so does Idris!

Given enough information in the type, Idris can search for a valid expression that satisfies the type. In Atom, press Ctrl-Alt-S over the `allLengths_rhs_1` hole, and you should see that the definition has changed:

```
allLengths : Vect n String -> Vect n Nat
allLengths [] = []
allLengths (word :: words) = ?allLengths_rhs_2
```

Because there's only one possible value for a vector of length zero, Idris has refined this automatically.

You can also try an expression search on the `allLengths_rhs_2` hole. Press Ctrl-Alt-S with the cursor over `allLengths_rhs_2`, and you should see this:

```
allLengths : Vect n String -> Vect n Nat
allLengths [] = []
allLengths (word :: words) = 0 :: allLengths words
```

The required type was `Vect (S k) Nat` so, as before, Idris realized that the only possible result would be a non-empty vector. It also realized that it could find a value of type `Vect k Nat` by calling `allLengths` recursively on `words`.

For the `Nat` at the head of the vector, Idris has found the first value that satisfies the type, `0`, but this isn't exactly what you want, so you can replace it with a hole—`?vecthead`:

```
allLengths : Vect n String -> Vect n Nat
allLengths [] = []
allLengths (word :: words) = ?vecthead :: allLengths words
```

Checking the type of `?vecthead` confirms that you're looking for a `Nat`:

```
word : String
k : Nat
words : Vect k String
-----
vecthead : Nat
```

As before, you can complete the definition by filling this hole with `length word`. So, not only does the more precise type give you more confidence in the correctness of the program and give you more information when writing the program, it also gives Idris some information, allowing it to write a good bit of the program for you.

3.2.3 Type, define, refine: sorting a vector

For all the functions you've written so far in this chapter, you've followed this process:

- 1 Write a type.
- 2 Create a skeleton definition.
- 3 Pattern match on an argument.
- 4 Fill in the holes on the right side with a combination of type-driven expression search and refining holes by hand.

Usually, there's a little more work to do, however. For example, you may find you need to create additional helper functions, inspect intermediate results, or refine the type you initially gave for a function.

You can see this in practice by creating a function that returns a sorted version of an input vector. You can use *insertion sort*, which is a simple sorting algorithm that's easily implemented in a functional style, informally described as follows:

- Given an empty vector, return an empty vector.
- Given the head and tail of a vector, sort the tail of the vector and then insert the head into the sorted tail such that the result remains sorted.

You can write this interactively, beginning with the skeleton definition shown in listing 3.3. Open an Atom buffer and put this code into a file called VecSort.idr. Remember that you can create the skeleton definition of `insSort` from the type with Ctrl-Alt-A.

**Listing 3.3 A skeleton definition of `insSort` on vectors with an initial type
(VecSort.idr)**

```
import Data.Vect
insSort : Vect n elem -> Vect n elem
insSort xs = ?insSort_rhs
```

This type explicitly states that the output must have the same length as the input. The element type, `elem`, is given by a type-level variable, so it stands for any type.

As you work through this process, I recommend that you inspect the types of each of the holes that arise using Ctrl-Alt-T, and ensure that you understand the types of the variables and the holes.

DEVELOPMENT WORKFLOW It's typically useful to have an Atom window open for interactively editing a file, as well as a terminal window with a REPL open for testing, evaluation, checking documentation, and so on.

Having written the type for this function, implement it by doing the following:

- 1 *Define*—Case-split on `xs` with Ctrl-Alt-C, leading to this:

```
insSort : Vect n elem -> Vect n elem
insSort [] = ?insSort_rhs_1
insSort (x :: xs) = ?insSort_rhs_2
```

- 2 *Refine*—Try an expression search on `?insSort_rhs_1`, leading to this:

```
insSort : Vect n elem -> Vect n elem
insSort [] = []
insSort (x :: xs) = ?insSort_rhs_2
```

A sorted empty vector is itself an empty vector, as expected.

- 3 *Refine*—Trying an expression search on `?insSort_rhs_2` is, unfortunately, less effective:

```
insSort : Vect n elem -> Vect n elem
insSort [] = []
insSort (x :: xs) = x :: xs
```

Although Idris knows how long the vector should be, and it has local variables of the correct types, the overall type of `insSort` isn't precise enough for Idris to fill in the hole with the program you intend.

EXPRESSION SEARCH As this example demonstrates, although expression search can often lead you to a valid function, it's not a substitute for understanding how the program works! You need to understand the algorithm, but you can use expression search to help fill in the details.

- 4 Define**—When you’re writing a function over a recursive data type, it’s often effective to make recursive calls to the recursive parts of the structure. Here, you can sort the tail with a recursive call to `insSort xs` and bind the result to a locally defined variable:

```
insSort : Vect n elem -> Vect n elem
insSort [] = []
insSort (x :: xs) = let xsSorted = insSort xs in
    ?insSort_rhs_2
```

- 5 Type**—In `?insSort_rhs_2` you’re going to insert the head `x` into the sorted tail, `xsSorted`. Because this is going to be a little more complex, you can lift the hole to a top-level definition by pressing Ctrl-Alt-L with the cursor over `?insSort_rhs_2`, which leads to the following:

```
insSort_rhs_2 : (x : elem) -> (xs : Vect k elem) ->
    (xsSorted : Vect k elem) ->
    Vect (S k) elem

insSort : Vect n elem -> Vect n elem
insSort [] = []
insSort (x :: xs) = let xsSorted = insSort xs in
    insSort_rhs_2 x xs xsSorted
```

This has created a new top-level function with a new type but no implementation, and it has replaced the hole with a call to the new function. The arguments to the new function are the local variables that were in scope in the hole `?insSort_rhs_2`.

- 6 Refine**—Once you realize that the job of the new function is to insert `x` into the `xsSorted` vector, you can edit the name and type of `insSort_rhs_2` to reflect this. You can remove the arguments that you aren’t going to need:

```
insert : (x : elem) -> (xsSorted : Vect k elem) -> Vect (S k) elem

insSort : Vect n elem -> Vect n elem
insSort [] = []
insSort (x :: xs) = let xsSorted = insSort xs in
    insert x xsSorted
```

LIFTING DEFINITIONS When lifting a definition with Ctrl-Alt-L, Idris will generate a new definition with the same name as the hole, using *all* of the local variables in the generated type. In this case, you know you aren’t going to need all of them, so you can edit out the unnecessary `xs` argument.

- 7 Define**—You now have to define `insert`. Again, create a skeleton definition:

```
insert : (x : elem) -> (xsSorted : Vect k elem) -> Vect (S k) elem
insert x xsSorted = ?insert_rhs
```

Then case-split on `xsSorted`, leading to this:

```
insert : (x : elem) -> (xsSorted : Vect k elem) -> Vect (S k) elem
insert x [] = ?insert_rhs_1
insert x (y :: xs) = ?insert_rhs_2
```

- 8 Refine**—An expression search on `insert_rhs_1` leads to the following:

```
insert : (x : elem) -> (xsSorted : Vect k elem) -> Vect (S k) elem
insert x [] = [x]
insert x (y :: xs) = ?insert_rhs_2
```

This works because Idris knows it's looking for a vector with one element of type `elem`, and the only thing available with type `elem` is `x`.

- 9 Refine**—For the `(y :: xs)` case, with the hole `?insert_rhs_2`, you have a problem. There are two cases to consider:

- If $x < y$, the result should be $x :: y :: xs$, because the result won't be ordered if x is inserted after y .
- Otherwise, the result should begin with y , and then have x inserted into the tail xs .

Your problem is that you know nothing about the element type `elem`. You need to constrain it so that you know you can compare elements of type `elem`. You can refine the type of `insert` (and correspondingly `insSort`) so that you can do the necessary comparison:

```
insert : Ord elem =>
          (x : elem) -> (xsSorted : Vect k elem) -> Vect (S k) elem
insert x [] = [x]
insert x (y :: xs) = ?insert_rhs_2

insSort : Ord elem => Vect n elem -> Vect n elem
insSort [] = []
insSort (x :: xs) = let xsSorted = insSort xs in
                      insert x xsSorted
```

Remember from chapter 2 that you constrain generic types by placing constraints such as `Ord elem` before `=>` in the type. You'll see more about this in chapter 7.

- 10 Define**—You now need to check $x < y$ and act on the result. One way to do this is with an `if...then...else` construct:

```
insert : Ord elem =>
          (x : elem) -> (xsSorted : Vect k elem) -> Vect (S k) elem
insert x [] = [x]
insert x (y :: xs) = if x < y then x :: y :: xs
                      else y :: insert x xs
```

Alternatively, you can use interactive editing to give more structure to the definition, and insert a `case` construct to match on an intermediate result. Press Ctrl-Alt-M with the cursor over `?insert_rhs_2`. This introduces a new `case` expression with a placeholder for the value to be inspected (so the function won't type-check yet):

```
insert : Ord elem =>
          (x : elem) -> (xsSorted : Vect k elem) -> Vect (S k) elem
insert x [] = [x]
```

```
insert x (y :: xs) = case _ of
    case_val => ?insert_rhs_2
```

The `_` stands for an expression you need to provide in order for the function to type-check successfully. You'll need to fill in the `_` with the expression you want to match:

```
insert : Ord elem =>
    (x : elem) -> (xsSorted : Vect k elem) -> Vect (S k) elem
insert x [] = [x]
insert x (y :: xs) = case x < y of
    case_val => ?insert_rhs_2
```

- 11 Define**—You can now case-split on `case_val` in the usual way, with Ctrl-Alt-C, leading to this:

```
insert : Ord elem =>
    (x : elem) -> (xsSorted : Vect k elem) -> Vect (S k) elem
insert x [] = [x]
insert x (y :: xs) = case x < y of
    False => ?insert_rhs_1
    True => ?insert_rhs_3
```

- 12 Refine**—Finally, you complete the implementation by filling in the new holes `insert_rhs_1` and `insert_rhs_3`. Unfortunately, expression search won't help you much here because there's not enough information in the type, so you need to fill these in by hand:

```
insert : Ord elem =>
    (x : elem) -> (xsSorted : Vect k elem) -> Vect (S k) elem
insert x [] = [x]
insert x (y :: xs) = case x < y of
    False => y :: insert x xs
    True => x :: y :: xs
```

Once the definition is complete, you can test it at the REPL, like this:

```
*VecSort> insSort [1,3,2,9,7,6,4,5,8]
[1, 2, 3, 4, 5, 6, 7, 8, 9] : Vect 9 Integer
```

Remember totality checking!

Don't forget to check that `insSort` is total:

```
*VecSort> :total insSort
Main.insSort is Total
```

It's a good habit to check that the functions you define are total. If a function is type-correct, but not total, it might appear to work when you test it, but there might still be a subtle error on some unusual inputs, such as a missing pattern or possible non-termination.

To summarize, following the type-define-refine process, you've done the following:

- 1 Written a type for `insSort`.
- 2 Tried to define `insSort` until you encountered the need to `insert`, which you lifted to a new top-level definition with its own type.
- 3 Tried to define `insert` until you encountered the need to `compare`, at which point you refined the types to support the ordering constraint on `elem`.
- 4 Continued to define `insert` using the refined type, and hence completed the implementation of `insSort`.

Exercises



To conclude this section, here are some exercises to test your understanding of the interactive editing mode and pattern matching on `List` and `Vect`.

The following functions, or some variant on them, are defined in the Prelude or in `Data.Vect`:

- 1 `length` : `List a` \rightarrow `Nat`
- 2 `reverse` : `List a` \rightarrow `List a`
- 3 `map` : $(a \rightarrow b)$ \rightarrow `List a` \rightarrow `List b`
- 4 `map` : $(a \rightarrow b)$ \rightarrow `Vect n a` \rightarrow `Vect n b`

For each of them, define your own version using interactive editing in Atom. Note that you'll need to use different names (such as `my_length`, `my_reverse`, `my_map`) to avoid clashing with the names in the Prelude. You can test your answers at the REPL as follows:

```
*ex_3_2> my_length [1..10]
10 : Nat

*ex_3_2> my_reverse [1..10]
[10, 9, 8, 7, 6, 5, 4, 3, 2, 1] : List Integer

*ex_3_2> my_map (* 2) [1..10]
[2, 4, 6, 8, 10, 12, 14, 16, 18, 20] : List Integer

*ex_3_2> my_vect_map length ["Hot", "Dog", "Jumping", "Frog"]
[3, 3, 7, 4] : Vect 4 Nat
```

Don't forget to check that your definitions are total!

3.3 Example: type-driven development of matrix functions

The main reason you might use vectors, with the length explicit in the type, as opposed to lists, is to have the lengths of the vectors help guide you to a working function more quickly. This can be particularly helpful when you work with two-dimensional vectors. These, in turn, can be helpful for implementing operations on matrices, which have several applications in programming, such as 3D graphics.

A *matrix*, in mathematics, is a rectangular array of numbers arranged in rows and columns. Figure 3.3 shows an example 3×4 matrix in both mathematical notation,

and in Idris notation as a vector of vectors. Notice that when representing a matrix as nested vectors, the dimensions of the matrix become explicit in the type.

$$\begin{pmatrix} 1 & 2 & 3 & 4 \\ 5 & 6 & 7 & 8 \\ 9 & 10 & 11 & 12 \end{pmatrix}$$

3 × 4 matrix

[[1, 2, 3, 4],
 [5, 6, 7, 8],
 [9, 10, 11, 12]]

Vect 3 (Vect 4 Int)

Figure 3.3 Representation of a matrix as two-dimensional vectors. On the left, the matrix is in mathematical notation. The same matrix is represented in Idris on the right.

3.3.1 Matrix operations and their types

When implementing operations on matrices, such as addition and multiplication, it's important to check that the dimensions of the vectors you're working with are appropriate for the operations. For example:

- When adding matrices, each matrix must have exactly the same dimensions. Addition works by adding corresponding elements in each matrix. For example, you can add two 3×2 matrices as follows:

$$\begin{pmatrix} 1 & 2 \\ 3 & 4 \\ 5 & 6 \end{pmatrix} + \begin{pmatrix} 7 & 8 \\ 9 & 10 \\ 11 & 12 \end{pmatrix} = \begin{pmatrix} 8 & 10 \\ 12 & 14 \\ 16 & 18 \end{pmatrix}$$

The following addition of a 2×2 matrix to a 3×2 matrix is invalid because there are no corresponding elements in the third row:

$$\begin{pmatrix} 1 & 2 \\ 3 & 4 \\ 5 & 6 \end{pmatrix} + \begin{pmatrix} 7 & 8 \\ 9 & 10 \end{pmatrix} = ???$$

So, the type of matrix addition could be as follows:

```
addMatrix : Num numType =>
            Vect rows (Vect cols numType) ->
            Vect rows (Vect cols numType) ->
            Vect rows (Vect cols numType)
```

In other words, for some numeric type `numType`, adding a `rows × cols` matrix to a `rows × cols` matrix results in a `rows × cols` matrix.

- When multiplying matrices, the number of columns in the left matrix must be the same as the number of rows in the right matrix. Then, multiplication works as follows:

$$\begin{pmatrix} 1 & 2 \\ 3 & 4 \\ 5 & 6 \end{pmatrix} \times \begin{pmatrix} 7 & 8 & 9 & 10 \\ 11 & 12 & 13 & 14 \end{pmatrix} = \begin{pmatrix} 29 & 32 & 35 & 38 \\ 65 & 72 & 79 & 86 \\ 101 & 112 & 123 & 134 \end{pmatrix}$$

Row 2 column 3 of result = $3 \times 9 + 4 \times 13 = 79$

Here, multiplying a 3×2 matrix by a 2×4 matrix results in a 3×4 matrix. The value in row x , column y in the result is the sum of the product of corresponding elements in row x of the left input, and column y of the right input. So the type of matrix multiplication could be as follows:

```
multMatrix : Num numType =>
    Vect n (Vect m numType) -> Vect m (Vect p numType) ->
    Vect n (Vect p numType)
```

In other words, for some numeric type `numType`, multiplying an $n \times m$ matrix by an $m \times p$ matrix results in an $n \times p$ matrix.

3.3.2 Transposing a matrix

A useful operation when manipulating matrices is transposition, which converts rows to columns and vice versa. For example, a 3×2 matrix becomes a 2×3 matrix:

$$\begin{pmatrix} 1 & 2 \\ 3 & 4 \\ 5 & 6 \end{pmatrix} \dots \text{transposed to} \dots \begin{pmatrix} 1 & 3 & 5 \\ 2 & 4 & 6 \end{pmatrix}$$

You can write a `transposeMat` function that, in general, converts an $m \times n$ matrix into an $n \times m$ matrix, representing matrices as nested vectors. As usual, you can write the function interactively, where each step is broadly characterized as one of type, define, or refine. From this point, I'll generally assume you're comfortable with the interactive commands in Atom, and I'll describe the overall type-driven process rather than the specifics of building the function.

- 1 **Type**—Begin by giving a type for `transposeMat`:

```
transposeMat : Vect m (Vect n elem) -> Vect n (Vect m elem)
```

For matrix arithmetic, in the types of `addMatrix` and `multMatrix`, you need to constrain the element type to be numeric. Here, though, the element type of the matrix, `elem`, could be anything. You're not going to inspect it or use it at any point in the implementation of `transposeMat`; you merely change the rows to columns and columns to rows.

- 2 Define**—Create a skeleton definition, and then case-split on the input vector:

```
transposeMat : Vect m (Vect n elem) -> Vect n (Vect m elem)
transposeMat [] = ?transposeMat_rhs_1
transposeMat (x :: xs) = ?transposeMat_rhs_2
```

- 3 Type**—When Idris creates new holes, either from a case split or an incomplete result of an expression search, it's always a good idea to inspect the types of those holes to get some insight into how to proceed. First, take a look at the type of `?transposeMat_rhs_1`:

```
elem : Type
n : Nat
-----
transposeMat_rhs_1 : Vect n (Vect 0 elem)
```

Here, you're trying to convert a $0 \times n$ vector into an $n \times 0$ vector, so you need to create n copies of an empty vector. We'll return to this case later; for now, you can rename the hole to `createEmpties` and lift it to a top-level function with Ctrl-Alt-L:

```
createEmpties : Vect n (Vect 0 elem)

transposeMat : Vect m (Vect n elem) -> Vect n (Vect m elem)
transposeMat [] = createEmpties
transposeMat (x :: xs) = ?transposeMat_rhs_2
```

- 4 Type**—Look at the type of `?transposeMat_rhs_2`:

```
elem : Type
n : Nat
x : Vect n elem
k : Nat
xs : Vect k (Vect n elem)
-----
transposeMat_rhs_2 : Vect n (Vect (S k) elem)
```

You have `xs`, which is a $k \times n$ matrix, and you need to make an $n \times (S k)$ matrix.

- 5 Define**—One insight you need here is that it's likely easier to build an $n \times (S k)$ matrix from an $n \times k$ matrix, because at least one of the dimensions is correct. You can create such a matrix by transposing `xs`:

```
transposeMat (x :: xs) = let xsTrans = transposeMat xs in
                        ?transposeMat_rhs_2
```

- 6 Type**—You can also lift `?transposeMat_rhs_2` to a top-level function, renaming it `transposeHelper`. This results in the following:

```
transposeHelper : (x : Vect n elem) -> (xs : Vect k (Vect n elem)) ->
                  (xsTrans : Vect n (Vect k elem)) -> Vect n (Vect (S k) elem)
```

The type for `transposeHelper` is generated from the types of the local variables you have access to: `x`, `xs`, and `xsTrans`. It will take these variables as inputs, and produce a `Vect n (Vect (S k) elem)` as output.

- 7 **Type**—At this stage, it's good to look more closely at the types of the variables `x`, `xs`, and `xsTrans` and try to use these types to visualize what you need to do to complete `transposeMat`.

For the sake of visualization, let's take $n = 4$ and $k = 2$. Figure 3.4 shows an original two-dimensional vector, of the form $(x :: xs)$, where `x` is the first row and `xs` is the rest of the rows, built with these dimensions. It identifies the components `x` and `xs`, and it shows the result of transposing `xs` and the expected result of the whole operation.




Figure 3.4 The components of the vector you're transposing (`x` and `xs`), along with the result of transposing `xs`, and the expected overall result

What you need to do, therefore, is add each element of `x` to the vector in the corresponding element of `xsTrans`; you won't need `xs` in `transposeHelper`. If you delete this by hand in the type of `transposeHelper` and the application in `transposeMat`, you get this:

```
transposeHelper : (x : Vect n elem) -> (xsTrans : Vect n (Vect k elem)) ->
Vect n (Vect (S k) elem)

transposeMat : Vect m (Vect n elem) -> Vect n (Vect m elem)
transposeMat [] = createEmpties
transposeMat (x :: xs) = let xsTrans = transposeMat xs in
                        transposeHelper x xsTrans
```

- 8 **Define**—To implement `transposeHelper`, you can add a skeleton definition, pattern match on `x` and `xsTrans`, and then use expression search to complete the definition. There's enough information for Idris to fill in the details itself:

```
transposeHelper : (x : Vect n elem) -> (xsTrans : Vect n (Vect k elem)) ->
Vect n (Vect (S k) elem)
transposeHelper [] [] = []
transposeHelper (x :: xs) (y :: ys) = (x :: y) :: transposeHelper xs ys
```

Rather than typing this in directly, try to build it using the interactive commands. It's possible to write this function from the type using only Ctrl-Alt-A, Ctrl-Alt-C, Ctrl-Alt-S, and cursor movements.

CASE SPLITTING ON VECTORS If you case-split on `x` and then case-split on `xsTrans`, notice that Idris only gives one possible pattern for `xsTrans`. This is because the types of `x` and `xsTrans` state that both must have the same length.

- 9 *Define*—All that remains is to implement `createEmpties`. You can implement this using a library function, `replicate`:

```
*transpose> :doc Vect.replicate
Data.Vect.replicate : (n : Nat) -> (x : a) -> Vect n a
  Repeat some value n times
  Arguments:
    n : Nat -- the number of times to repeat it
    x : a -- the value to repeat
```

If you create a skeleton definition of `createEmpties` from its type, you'll see the following:

```
createEmpties : Vect n (Vect 0 elem)
createEmpties = ?createEmpties_rhs
```

You need to call `replicate` to build a vector of `n` empty lists. Unfortunately, because there are no local variables available from the patterns on the left side, the natural definition results in an error message:

```
createEmpties : Vect n (Vect 0 elem)
createEmpties = replicate n [] -- "No such variable n"
```

The problem is that `n` is a type-level variable, and not accessible to the definition of `createEmpties`. Shortly, in section 3.4, you'll see how to handle type-level variables in general, and how you might write `createEmpties` directly. For the moment, because the type dictates that there's only one valid value for the length argument to `replicate`, you can use an underscore instead:

```
createEmpties : Vect n (Vect 0 elem)
createEmpties = replicate _ []
```

The implementation of `transposeMat` is now complete. For reference, the complete definition is given in listing 3.4. You can test it at the REPL:

```
*transpose> transposeMat [[1,2], [3,4], [5,6]]
[[1, 3, 5], [2, 4, 6]] : Vect 2 (Vect 3 Integer)
```

Listing 3.4 Complete definition of matrix transposition (Matrix.idr)

```
createEmpties : Vect n (Vect 0 elem)
createEmpties = replicate _ []

transposeHelper : (x : Vect n elem) ->
  (xsTrans : Vect n (Vect k elem)) ->
  Vect n (Vect (S k) elem)
transposeHelper [] [] = []
transposeHelper (x :: xs) (y :: ys) = (x :: y) :: transposeHelper xs ys

transposeMat : Vect m (Vect n elem) -> Vect n (Vect m elem)
transposeMat [] = createEmpties
transposeMat (x :: xs) = let xsTrans = transposeMat xs in
  transposeHelper x xsTrans
```

Code reuse

When building a definition interactively using the type-define-refine process, it's good to look out for parts of the definition that could be made more generic, or that could instead be implemented using existing library functions.

For example, `transposeHelper` has a structure very similar to the library function `zipWith`, which applies a function to corresponding elements in two vectors and is defined as follows:

```
zipWith : (a -> b -> c) -> Vect n a -> Vect n b -> Vect n c
zipWith f []           []           = []
zipWith f (x :: xs) (y :: ys) = f x y :: zipWith f xs ys
```

Exercises



- Reimplement `transposeMat` using `zipWith` instead of `transposeHelper`.

You can test your answer at the REPL as follows:

```
*ex_3_3_3> transposeMat [[1,2], [3,4], [5,6]]
[[1, 3, 5], [2, 4, 6]] : Vect 2 (Vect 3 Integer)
```

- Implement `addMatrix` : `Num a => Vect n (Vect m a) -> Vect n (Vect m a) -> Vect n (Vect m a)`.

You can test your answer at the REPL as follows:

```
*ex_3_3_3> addMatrix [[1,2], [3,4]] [[5,6], [7,8]]
[[6, 8], [10, 12]] : Vect 2 (Vect 2 Integer)
```

- Implement a function for multiplying matrices, following the description given in section 3.3.1.

Hint: This definition is quite tricky and involves multiple steps. Consider the following:

- You have a left matrix of dimensions $n \times m$, and a right matrix of dimensions $m \times p$. A good start is to use `transposeMat` on the *right* matrix.
- Remember that you can use Ctrl-Alt-L to lift holes to top-level functions.
- Remember to pay close attention to the types of the local variables and the types of the holes.
- Remember to use Ctrl-Alt-S to search for expressions, and pay close attention to the types of any resulting holes.

You can test your answer at the REPL as follows:

```
*ex_3_3_3> multMatrix [[1,2], [3,4], [5,6]] [[7,8,9,10], [11,12,13,14]]
[[29, 32, 35, 38],
 [65, 72, 79, 86],
 [101, 112, 123, 134]] : Vect 3 (Vect 4 Integer)
```

3.4 Implicit arguments: type-level variables

You've now seen several definitions with variables at the type level that can stand either for types or values. For example:

```
reverse : List elem -> List elem
```

Here, `elem` is a type-level variable standing for the element type of the list. It appears twice, in the input type and the return type, so the element type of each must be the same.

```
append : Vect n elem -> Vect m elem -> Vect (n + m) elem
```

Here, `elem` is again a type-level variable standing for the element type of the vectors. `n` and `m` are type-level variables standing for the lengths of the input vectors, and they're used again in the output to describe how the length of the output relates to the length of the inputs.

These type-level variables aren't declared anywhere else. Because types are first class, type-level variables can also be brought into scope and used in definitions. These type-level variables are referred to as *implicit* arguments to the functions `reverse` and `append`. In this section, you'll see how implicit arguments work, and how to use them in definitions.

3.4.1 The need for implicit arguments

To illustrate the need for implicit arguments, let's take a look at how you might define `append` without them. You could make the `elem`, `n`, and `m` arguments to `append` explicit, resulting in the following definition:

```
append : (elem : Type) -> (n : Nat) -> (m : Nat) ->
          Vect n elem -> Vect m elem -> Vect (n + m) elem
append elem Z m [] ys = ys
append elem (S k) m (x :: xs) ys = x :: append elem k m xs ys
```

But if you did so, you'd also have to be explicit about the element type and lengths when calling `append`:

```
*Append_expl> append Char 2 2 ['a','b'] ['c','d']
['a', 'b', 'c', 'd'] : Vect 4 Char
```

Given the types of the arguments `['a', 'b']` and `['c', 'd']`, there's only one possible value for each of the arguments `elem` (which must be a `Char`), `n` (which must be 2, from the length of `['a', 'b']`), and `m` (which must also be 2, from the length of `['c', 'd']`). Any other value for any of these would not be well typed.

Because there is enough information in the types of the vector arguments, Idris can infer the values for the `a`, `n`, and `m` arguments. You can therefore write this:

```
*Append> append _ _ _ ['a','b'] ['c','d']
['a', 'b', 'c', 'd'] : Vect 4 Char
```

Implicit values

An underscore (`_`) in a function call means that you want Idris to work out an *implicit value* for the argument, given the information in the rest of the expression:

```
*Append> append _ _ _ ['a','b'] ['c','d']
['a', 'b', 'c', 'd'] : Vect 4 Char
```

Idris will report an error if it can't:

```
*Append> append _ _ _ _ ['c','d']
(input) : Can't infer argument n to append,
          Can't infer explicit argument to append
```

Here, Idris reports that it can't work out the length of the first vector, or the first vector itself. Unlike *holes*, which stand for parts of expressions that aren't yet written, underscores stand for parts of expressions for which there's only one valid value. It's an error if Idris can't infer a unique value for an underscore.

Using implicit arguments avoids the need for writing details explicitly that can be inferred by Idris. Making `elem`, `n`, and `m` implicit in the type of `append` means that you can refer to them directly in the type where necessary, without needing to give explicit values when calling the function.

3.4.2 Bound and unbound implicits

Let's take another look at the types of `reverse` and `append`, with implicit arguments:

```
reverse : List elem -> List elem
append : Vect n elem -> Vect m elem -> Vect (n + m) elem
```

The names `elem`, `n`, and `m` are called *unbound implicits*. This is because their names are used directly, without being declared (or *bound*) anywhere else. You could also have written these types as follows:

```
reverse : {elem : Type} -> List elem -> List elem
append : {elem : Type} -> {n : Nat} -> {m : Nat} ->
         Vect n elem -> Vect m elem -> Vect (n + m) elem
```

Here, the implicit arguments have been explicitly *bound* in the type. The notation $\{x : S\} \rightarrow T$ denotes a function type where the argument is intended to be *inferred* by Idris, rather than written directly by the programmer.

When you write a type with *unbound* implicits, Idris will look for undefined names in the type, and turn them internally into *bound* implicits. Consider this example:

```
append : Vect n elem -> Vect m elem -> Vect (n + m) elem
```

First, Idris identifies that `elem`, `n`, and `m` are undefined, so it internally rewrites the type as follows:

```
append : {elem : _} -> {n : _} -> {m : _} ->
         Vect n elem -> Vect m elem -> Vect (n + m) elem
```

Note that it has not attempted to fill in types for the new arguments, but instead has given them as underscores in the hope that it will be able to infer them from some other information in the rest of the type. Here, this results in the following:

```
append : {elem : Type} -> {n : Nat} -> {m : Nat} ->
          Vect n elem -> Vect m elem -> Vect (n + m) elem
```

Unbound implicit names

In practice, Idris will not treat every undefined name as an unbound implicit—only names that begin with a lowercase letter and that appear either on their own or in a function argument position. Given the following,

```
test : f m a -> b -> a
```

`m` appears in an argument position, as does `a`. `b` appears on its own, and `f` only appears in a function position. As a result, `m`, `a`, and `b` are treated as unbound implicits. `f` isn't treated as an unbound implicit, meaning that it *must* be defined elsewhere for this type to be valid.

Normally, for conciseness, you leave implicits unbound, but in some situations, it's useful to use bound implicits instead:

- For clarity and readability, it can be useful to give explicit types to implicit arguments.
- If there's a dependency between an implicit argument and some other argument, you may need to use a bound implicit to make this clear to Idris.

3.4.3 Using implicit arguments in functions

Internally, Idris treats implicit arguments like any other argument, but with the notational convenience that the programmer doesn't need to provide them explicitly. As a result, you can refer to implicit arguments inside function definitions, and even case-split on them.

For example, how can you find the length of a vector? You could do this by case splitting on the vector itself:

```
length : Vect n elem -> Nat
length [] = Z
length (x :: xs) = 1 + length xs
```

Because the length is part of the type, you could also refer to it directly:

```
length : Vect n elem -> Nat
length {n} xs = n
```

The notation `{n}` in a pattern brings the implicit argument `n` into scope, allowing you to use it directly.

More generally, you can give explicit values for implicit arguments by using the notation `{n = value}`, where `n` is the name of an implicit argument:

```
*Append> append {elem = Char} {n = 2} {m = 3}
append : Vect 2 Char -> Vect 3 Char -> Vect 5 Char
```

Here, you've *partially applied* `append` to its implicit arguments only, giving a specialized function for appending a vector of two Chars to a vector of three Chars.

This notation can also be used on the left side of a definition to case-split on an implicit argument. For example, to implement `createEmpties` in section 3.3.2, you could have written it directly by case splitting on the length `n` after bringing it into scope:

```
createEmpties : Vect n (Vect 0 a)
createEmpties {n} = ?createEmpties_rhs
```

If you case-split on `n` in Atom, you'll see this:

```
createEmpties : Vect n (Vect 0 a)
createEmpties {n = Z} = ?createEmpties_rhs_1
createEmpties {n = (S k)} = ?createEmpties_rhs_2
```

Finally, you can complete the definition with an expression search for both of the remaining holes:

```
createEmpties : Vect n (Vect 0 a)
createEmpties {n = Z} = []
createEmpties {n = (S k)} = [] :: createEmpties
```

Note that in the recursive call, `createEmpties` is sufficient. There's no need to provide an explicit value for the length because only one value (`k`) would type-check. On the other hand, if you try at the REPL without a value for the length, Idris will report an error:

```
*transpose> createEmpties
(input):Can't infer argument n to createEmpties,
          Can't infer argument a to createEmpties
```

You can solve this either by giving explicit values for `n` and `elem`, or giving a target type for the expression:

```
*transpose> createEmpties {a=Int} {n=4}
[] : Vect 4 (Vect 0 Int)

*transpose> the (Vect 4 (Vect 0 Int)) createEmpties
[] : Vect 4 (Vect 0 Int)
```

Type and argument erasure

Because implicit arguments are, internally, treated the same as any other argument, you might wonder what happens at runtime. Generally, you use implicit arguments to give precise types to programs, so does this mean that type information has to be compiled and present at runtime?

Fortunately, the Idris compiler is aware of this problem. It will analyze a program before compiling it so that any arguments that are only used for type checking won't be present at runtime.

3.5 Summary

- Functions in Idris are defined by collections of pattern-matching equations.
- Patterns arise from the constructors of a data type.
- The Atom text editor provides an interactive editing mode that uses the type to help direct function implementation. Similar modes are available for Emacs and Vim.
- Interactive editing commands provide natural tools for following the process of type, define, refine.
- Interactive editing provides a command for searching for a valid expression that satisfies the type of a hole.
- More-precise types, such as `Vect`, give more information to the compiler both to help check that a function is correct, and to help constrain expression searches.
- Matrices are two-dimensional vectors, with the dimensions encoded in the type. Operations on matrices such as addition and multiplication can be given types that precisely describe how the operations affect the dimensions.
- The type-define-refine process helps you to implement matrix operations with precise types by using the type to direct the implementation of each subexpression and create appropriate helper functions.
- Type-level variables are implicit arguments to functions, which can be brought into scope and used like any other arguments by enclosing them in braces `{ }.`.

User-defined data types



This chapter covers

- Defining your own data types
- Understanding different forms of data types
- Writing larger interactive programs in the type-driven style

Type-driven development involves not only giving precise types to functions, as you've seen so far, but also thinking about exactly how data is structured. In a sense, programming (pure functional programming, in particular) is about transforming data from one form to another. Types allow us to describe the form of that data, and the more precise we make these descriptions, the more guidance the language can give in implementing transformations on that data.

Several useful data types are distributed as part of the Idris libraries, many of which we've used so far, such as `List`, `Bool`, and `Vect`. Other than being defined directly in Idris, there's nothing special about these data types. In any realistic program, you'll need to define your own data types to capture the specific requirements of the problem you're solving and the specific forms of data you're working with. Not only that, but there's a significant payoff to thinking carefully about the design of data types: the more precisely types capture the requirements of a problem, the more benefit you'll get from interactive type-directed editing.

In this chapter, therefore, we'll look at how to define new data types. You'll see the various forms of data types in a number of small example functions. We'll also start to work on a larger example, an interactive data store, that we'll extend in the coming chapters. At first, we'll store only strings, accessing them by an integer index, but even in this small example you'll see how user-defined types and dependent types can help build an interactive interface, dealing safely with possible runtime errors.

4.1 Defining data types

Data types are defined by a *type constructor* and one or more *data constructors*. In fact, you've already seen these when you used `:doc` to see the details of a data type. For example, the following listing shows the output of `:doc List`, annotated to highlight the type and data constructors.

Listing 4.1 Type and data constructors from `:doc`

```
Data type Prelude.List.List : (elem : Type) -> Type
  Generic lists

  Constructors:
    Nil : List elem
      Empty list

    (::) : elem -> List elem -> List elem
      A non-empty list, consisting of a head element and the rest of
      the list.
```

The data constructor `Nil` takes no arguments and returns an empty list.

The type constructor for `List` has a function type, which takes a `Type` as an input and returns a `Type`.

The data constructor `(::)` takes two arguments and returns a list.

Using data constructors is the canonical way of building the types given by the type constructor. In the case of `List`, this means that every value with a type of the form `List elem` is either `Nil` or takes the form `x :: xs` for some element `x` and the remainder of the list, `xs`.

We'll classify types into five basic groups, although they're all defined with the same syntax:

- *Enumerated types*—Types defined by giving the possible values directly
- *Union types*—Enumerated types that carry additional data with each value
- *Recursive types*—Union types that are defined in terms of themselves
- *Generic types*—Types that are parameterized over some other types
- *Dependent types*—Types that are computed from some other value

In this section, you'll see how to define enumerated types, union types, recursive types, and generic types; we'll discuss dependent types in the next section. If you've programmed in a functional language before, or in any language that allows you to define generic types, the types we discuss in this section should be familiar, even if the notation is new.

NAMING CONVENTIONS By convention, I'll use an initial capital letter for both type constructors and data constructors, and an initial lowercase letter for functions. There's no requirement to do so, but it gives a useful visual indication to the reader of the code.

4.1.1 Enumerations

An *enumerated* type is directly defined by giving the valid values for that type. The simplest example is `Bool`, which is defined as follows in the Prelude:

```
data Bool = False | True
```

To define an enumerated data type, you give the `data` keyword to introduce the declaration, and then give the name of the type constructor (in this case, `Bool`) and list the names of the data constructors (in this case `True` and `False`).

Figure 4.1 shows another example, defining an enumerated type for representing the four cardinal points on a compass.




Figure 4.1 Defining a `Direction` data type (`Direction.idr`)

Data constructor names are separated by a vertical bar, `|`, and there's no restriction on how the declaration is laid out (other than the usual layout rule that all declarations begin in precisely the same column). For example, they could each be on a different line:

```
data Direction = North
  | East
  | South
  | West
```

Once you've defined a data type, you can use it to define functions interactively. For example, you could define a `turnClockwise` function as follows, with the usual process of type, define, refine:

- 1 *Type*—Write a function type with `Direction` as the input and output, and then create a skeleton definition:

```
turnClockwise : Direction -> Direction
turnClockwise x = ?turnClockwise_rhs
```

- 2 *Define*—Define the function by case splitting on `x`:

```
turnClockwise : Direction -> Direction
turnClockwise North = ?turnClockwise_rhs_1
```

```
turnClockwise East = ?turnClockwise_rhs_2
turnClockwise South = ?turnClockwise_rhs_3
turnClockwise West = ?turnClockwise_rhs_4
```

- 3 Refine—Fill in the holes on the right sides:

```
turnClockwise : Direction -> Direction
turnClockwise North = East
turnClockwise East = South
turnClockwise South = West
turnClockwise West = North
```

4.1.2 Union types

A *union type* is an extension of an enumerated type in which the constructors of the type can themselves carry data. For example, you could create an enumeration type for shapes:

```
data Shape = Triangle | Rectangle | Circle
```

You may want to store more information with a shape, so that you can draw it, for example, or calculate its area. This information will differ depending on the shape:

- For a triangle, you might want to know the length of its base and its height.
- For a rectangle, you might want to know its length and height.
- For a circle, you might want to know its radius.

To represent this information, each of the data constructors, `Triangle`, `Rectangle`, and `Circle`, can be given argument types that carry this data using `Doubles` to represent dimensions. Figure 4.2 shows some example shapes and their representations.




Figure 4.2 Example shapes represented as a union type. `Triangle` takes two Doubles, for base and height; `Rectangle` takes two Doubles, for width and height; `Circle` takes a Double for radius.

You can represent shapes in this form as an Idris data type:

```
data Shape = Triangle Double Double
           | Rectangle Double Double
           | Circle Double
```

Listing 4.2 shows how you might define an area function that calculates the area of a shape for each of these possibilities. As an exercise, rather than typing in this function directly, try to construct it using the interactive editing tools in Atom.

Listing 4.2 Defining a Shape type and calculating its area (Shape.idr)

```
data Shape = Triangle Double Double
           | Rectangle Double Double
           | Circle Double

area : Shape -> Double
area (Triangle base height) = 0.5 * base * height
area (Rectangle length height) = length * height
area (Circle radius) = pi * radius * radius
```



If you check the documentation for Shape using :doc, you can see how this data declaration translates into type and data constructors:

```
*Shape> :doc Shape
Data type Main.Shape : Type

Constructors:
  Triangle : Double -> Double -> Shape
  Rectangle : Double -> Double -> Shape
  Circle : Double -> Shape
```

When you define new data types, it's also a good idea to provide documentation that will be displayed with :doc, using documentation comments. In this case, it helps indicate what each Double is for. Documentation comments are laid out in data declarations by giving the comment before each constructor:

```
||| Represents shapes
data Shape = ||| A triangle, with its base length and height
            Triangle Double Double
            ||| A rectangle, with its length and height
            Rectangle Double Double
            ||| A circle, with its radius
            Circle Double
```

This is rendered as follows with :doc:

```
*Shape> :doc Shape
Data type Main.Shape : Type
  Represents shapes

Constructors:
  Triangle : Double -> Double -> Shape
    A triangle, with its base length and height

  Rectangle : Double -> Double -> Shape
    A rectangle, with its length and height

  Circle : Double -> Shape
    A circle, with its radius
```

Data type syntax

There are two forms of data declaration. In one, as you've already seen, you list the data constructors and the types of their arguments:

```
data Shape = Triangle Double Double
           | Rectangle Double Double
           | Circle Double
```

It's also possible to define data types by giving their type and data constructors directly, in the form shown by :doc. You could define Shape as follows:

```
data Shape : Type where
  Triangle : Double -> Double -> Shape
  Rectangle : Double -> Double -> Shape
  Circle : Double -> Shape
```

This is identical to the previous declaration. In this case, it's a little more verbose, but this syntax is more general and flexible. You'll see more of this shortly when we define dependent types.

I'll use both syntaxes throughout the book. Generally, I'll use the concise form, unless I need the additional flexibility.

4.1.3 Recursive types

Types can also be *recursive*, that is, defined in terms of themselves. For example, Nat is defined recursively in the Prelude as follows:

```
data Nat = Z | S Nat
```

The Prelude also defines functions and notation to allow Nat to be used like any other numeric type, so rather than writing `S (S (S (S Z)))`, you can simply write 4. Nevertheless, in its primitive form it's defined using data constructors.

Nat and efficiency

It would be reasonable to be concerned about the efficiency of Nat given that it's defined in terms of constructors. There's no need to worry, though, for three reasons:

- In practice, Nat is primarily used *structurally*, to describe the size of a data structure such as Vect in its type, so the size of a Nat corresponds to the size of the data structure itself.
- When a program is compiled, types are erased, so a Nat that appears at the type level, such as in the length of a Vect, is erased.
- Internally, the compiler optimizes the representation of Nat so that it's really stored as a machine integer.

You can use a recursive type to extend the Shape example from the previous section to represent larger pictures. We'll define a picture as being one of the following:

- A primitive shape
- A combination of two other pictures
- A picture rotated through an angle
- A picture translated to a different location

Note that three of these are defined in terms of pictures themselves. You could define a picture type following the preceding informal description, as shown next.

Listing 4.3 Defining a Picture type recursively, consisting of Shapes and smaller Pictures (Picture.idr)

Uses the Shape type defined earlier as the primitive

```
data Picture = Primitive Shape
              | Combine Picture Picture
              | Rotate Double Picture
              | Translate Double Double Picture
```

Builds a picture by rotating another picture through an angle

Builds a picture by combining two smaller pictures

Builds a picture by moving a picture to another location

Figure 4.3 shows an example of the sort of picture you could represent with this data type. For each of the primitive shapes, we'll consider its location to be the top left of an imaginary box that bounds the shape.




Figure 4.3 An example picture combining three shapes translated to different positions

We know that there are three subpictures, so to represent this in code, you can begin (define) by using `Combine` to put three subpictures together:

```
testPicture : Picture
testPicture = Combine ?pic1 (Combine ?pic2 ?pic3)
```

To continue, you know that each subpicture is translated to a specific position, so you can fill in those details (refine), leaving holes for the primitive shapes themselves:

```
testPicture : Picture
testPicture = Combine (Translate 5 5 ?rectangle)
    (Combine (Translate 35 5 ?circle)
        (Translate 15 25 ?triangle))
```

Finally, you can fill in (refine) the details of the individual primitive shapes. One way to do this is to use Ctrl-Alt-L in Atom to lift the holes to the top level (type), and then fill in the definition, resulting in the following final definition.

Listing 4.4 The picture from figure 4.3 represented in code (Picture.idr)

```
rectangle : Picture
rectangle = Primitive (Rectangle 20 10)

circle : Picture
circle = Primitive (Circle 5)

triangle : Picture
triangle = Primitive (Triangle 10 10)

testPicture : Picture
testPicture = Combine (Translate 5 5 rectangle)
    (Combine (Translate 35 5 circle)
        (Translate 15 25 triangle))
```

Because Rectangle is a data constructor of Shape rather than Picture, you need to build the picture with a Primitive.

As usual, you write functions over the Picture data type by case splitting. To write a function that calculates the area of every primitive shape in a picture, you can begin by writing a type:

```
pictureArea : Picture -> Double
```

Then, create a skeleton definition and case-split on its argument. You should reach the following:

```
pictureArea : Picture -> Double
pictureArea (Primitive x) = ?pictureArea_rhs_1
pictureArea (Combine x y) = ?pictureArea_rhs_2
pictureArea (Rotate x y) = ?pictureArea_rhs_3
pictureArea (Translate x y z) = ?pictureArea_rhs_4
```

This has given you an outline definition, showing you the forms the input can take and giving holes on the right side.

The names of the variables x, y, and z, chosen by Idris when creating the patterns for pictureArea, are not particularly informative. You can tell Idris how to choose better default names using the %name directive:

```
%name Shape shape, shape1, shape2
%name Picture pic, pic1, pic2
```

Now, when Idris needs to choose a variable name for a variable of type Shape, it will choose shape by default, followed by shape1 if shape is already in scope, followed by shape2. Similarly, it will choose pic, pic1, or pic2 for variables of type Picture.

After adding %name directives, case splitting on the argument will lead to patterns with more-informative variable names:

```
pictureArea : Picture -> Double
pictureArea (Primitive shape) = ?pictureArea_rhs_1
pictureArea (Combine pic pic1) = ?pictureArea_rhs_2
pictureArea (Rotate x pic) = ?pictureArea_rhs_3
pictureArea (Translate x y pic) = ?pictureArea_rhs_4
```

The completed definition is given in listing 4.5. For every Picture you encounter in the structure, you recursively call pictureArea, and when you encounter a Shape, you call the area function defined previously.

Listing 4.5 Calculating the total area of all shapes in a Picture (Picture.idr)

Uses the area function defined earlier to calculate the area of a primitive shape

```
pictureArea : Picture -> Double
pictureArea (Primitive shape) = area shape
pictureArea (Combine pic pic1) = pictureArea pic + pictureArea pic1
pictureArea (Rotate x pic) = pictureArea pic
pictureArea (Translate x y pic) = pictureArea pic
```

When two pictures are combined, the area is the sum of the areas of those pictures.

When a picture is rotated, the area is the area of the rotated picture.

When a picture is translated, the area is the area of the translated picture.

It's always a good idea to test the resulting definition at the REPL:

```
*Picture> pictureArea testPicture
328.5398163397473 : Double
```

Ininitely recursive types

Recursive data types, just like recursive functions, need at least one non-recursive case in order to be useful, so at least one of the constructors needs to have a non-recursive argument. If you don't do this, you'll never be able to construct an element of that type. For example:

```
data Infinite = Forever Infinite
```

It is, however, possible to work with infinite streams of data using a generic Inf type, which we'll explore in chapter 11.

4.1.4 Generic data types

A *generic* data type is a type that's *parameterized* over some other type. Just like generic function types, which you saw in chapter 2, generic data types allow you to capture common patterns of data.

To illustrate the need for generic data types, consider a function that returns the area of the largest triangle in a `Picture`, as defined in the previous section. At first, you might write the following type:

```
biggestTriangle : Picture -> Double
```

But what should it return if there are no triangles in the picture? You could return some kind of sentinel value, like a negative size, but this would be against the spirit of type-driven development because you'd be using `Double` to represent something that isn't really a number. Idris also doesn't have a null value. Instead, you could refine the type of `biggestTriangle`, introducing a new union type for capturing the possibility that there are no triangles:

```
data Biggest = NoTriangle | Size Double
biggestTriangle : Picture -> Biggest
```

I'll leave the definition of `biggestTriangle` as an exercise.

You might also want to write a type that represents the possibility of failure. For example, you could write a safe division function for `Double` that returns an error if dividing by zero:

```
data DivResult = DivByZero | Result Double
safeDivide : Double -> Double -> DivResult
safeDivide x y = if y == 0 then DivByZero
                 else Result (x / y)
```

Both `Biggest` and `DivResult` have the same structure! Instead of defining multiple types of this form, you can define one single generic type. In fact, such a generic type exists in the Prelude, called `Maybe`. The following listing shows the definition of `Maybe`.

Listing 4.6 A generic type, `Maybe`, that captures the possibility of failure

```
data Maybe valtype =
    Nothing
    | Just valtype
```

The name `valtype` here is a type-level variable standing for any type you wish to use `Maybe` with.

Nothing indicates that no value is stored.

Just is a constructor that takes one argument and indicates that a single value is stored.

In a generic type, we use type variables such as `valtype` here to stand for concrete types. You can now define `safeDivide` using `Maybe Double` instead of `DivResult`, instantiating `valtype` with `Double`:

```
safeDivide : Double -> Double -> Maybe Double
safeDivide x y = if y == 0 then Nothing
                 else Just (x / y)
```

Definition of List

We've already written several functions with a generic type, `List`, which is defined in the Prelude as follows:

```
data List elem = Nil | (:) elem (List elem)
```

Generic data types can have more than one parameter, as illustrated in figure 4.4 for `Either`, which is defined in the Prelude and represents a choice between two alternative types.




Figure 4.4 Defining the `Either` data type

Generic types and terminology

You might hear people informally referring to “the `List` type” or “the `Maybe` type.” It's not strictly accurate to do this, however. `List` on its own is not a type, as you can confirm at the REPL:

```
Idris> :t List
List : Type -> Type
```

Technically, `List` has a function type that takes a `Type` as a parameter and returns a `Type`. Although `List Int` is a type because it has been applied to a concrete argument, `List` itself is not. Instead, we'll informally refer to `List` as a generic type.

One useful example of a generic type is a binary tree structure. The following listing shows the definition of binary trees, using a `%name` directive to give naming hints for building definitions interactively.

Listing 4.7 Defining binary trees (Tree.idr)

```
data Tree elem = Empty ← A tree with no data
               | Node (Tree elem) elem (Tree elem) ← A node with a left
                                                         subtree, a value,
                                                         and a right subtree
%name Tree tree, tree1
```

Binary trees are commonly used to store ordered information, where everything in a node's left subtree is *smaller* than the value at the node, and everything in a node's right subtree is *larger* than the value at the node.

Such trees are called *binary search trees*, and you can write a function to insert a value into such a tree, provided that you can order those values:

```
insert : Ord elem => elem -> Tree elem -> Tree elem
insert x tree = ?insert_rhs
```

To write this function, create a Tree.idr file containing the definition in listing 4.7, and do the following:

- 1 *Define*—Case-split on the tree, giving the cases where the tree is `Empty` and where the tree is a `Node`:

```
insert : Ord elem => elem -> Tree elem -> Tree elem
insert x Empty = ?insert_rhs_1
insert x (Node tree y tree1) = ?insert_rhs_2
```

Even with the `%name` directive, the names `tree`, `y`, and `tree1` aren't especially informative, so let's rename them to indicate that they're the left subtree, the value at the node, and the right subtree, respectively:

```
insert : Ord elem => elem -> Tree elem -> Tree elem
insert x Empty = ?insert_rhs_1
insert x (Node left val right) = ?insert_rhs_2
```

- 2 *Refine*—For `?insert_rhs_1`, you create a new tree node with empty subtrees:

```
insert : Ord elem => elem -> Tree elem -> Tree elem
insert x Empty = Node Empty x Empty
insert x (Node left val right) = ?insert_rhs_2
```

- 3 *Define*—For `?insert_rhs_2`, you need to compare the value you're inserting, `x`, with the value at the node, `val`. If `x` is smaller than `val`, you insert it into the left subtree. If it's equal, it's already in the tree, so you return the tree unchanged. If it's greater than `val`, you insert it into the right subtree. To do this, you can use the `compare` function from the Prelude, which returns an element of an `Ordering` enumeration:

```
data Ordering = LT | EQ | GT
compare : Ord a => a -> a -> Ordering
```

You'll perform a match on the intermediate result of `compare x val`. Press Ctrl-Alt-M over `insert_rhs_2`:

```
insert : Ord elem => elem -> Tree elem -> Tree elem
insert x Empty = Node Empty x Empty
insert x (Node left val right) = case _ of
    case_val => ?insert_rhs_2
```

- 4 *Define*—The `case` expression won't type-check until the `_` is replaced with an expression to test, so replace it with `compare x val`:

```

insert : Ord elem => elem -> Tree elem -> Tree elem
insert x Empty = Node Empty x Empty
insert x (Node left val right) = case compare x val of
    case_val => ?insert_rhs_2

```

- 5 *Define*—If you now case-split on `case_val`, you'll get the patterns for LT, EQ, and GT:

```

insert : Ord elem => elem -> Tree elem -> Tree elem
insert x Empty = Node Empty x Empty
insert x (Node left val right) = case compare x val of
    LT => ?insert_rhs_1
    EQ => ?insert_rhs_3
    GT => ?insert_rhs_4

```

The following listing shows the complete definition of this function, after refining the remaining holes.

Listing 4.8 Inserting a value into a binary search tree (Tree.idr)

```

insert : Ord elem => elem -> Tree elem -> Tree elem
insert x Empty = Node Empty x Empty
insert x (Node left val right)
    = case compare x val of
        LT => Node (insert x left) val right
        EQ => Node left val right
        GT => Node left val (insert x right)

x is less than the val at the node, so return a new tree with x inserted into the left subtree.

x is already in the tree, because it's equal to val, so return the original tree.

x is greater than the val at the node, so return a new tree with x inserted into the right subtree.

```

@ patterns

In `insert`, you might have noticed that in the EQ branch, the value you returned was exactly the same as the pattern on the left side. As a notational convenience, you can also *name* patterns:

```

insert x orig@(Node left val right)
    = case compare x val of
        LT => Node (insert x left) val right
        EQ => orig
        GT => Node left val (insert x right)

```

The notation `orig@(Node left val right)` gives the name `orig` to the pattern `Node left val right`. It doesn't change the meaning of the pattern match, but it does mean that you can use the name `orig` on the right side rather than repeating the pattern.

There needs to be an `Ord` constraint on the generic variable `elem` in the type of `insert`, because otherwise you wouldn't be able to use `compare`. An alternative

approach would be to capture the `Ord` constraint in the tree type itself, refining the type to include this extra precision. The following listing shows how to do this by giving the type and data constructors directly.

Listing 4.9 A binary search tree with the ordering constraint in the type (`BSTree.idr`)

Name this type `BSTree` rather than `Tree` because the type is explicitly for representing binary search trees.

Put an `Ord` constraint on the data constructors so you can't have a search tree containing values that aren't ordered.

```
→ data BSTree : Type -> Type where
    Empty : Ord elem => BSTree elem
    Node : Ord elem => (left : BSTree elem) -> (val : elem) ->
           (right : BSTree elem) -> BSTree elem

    insert : elem -> BSTree elem -> BSTree elem
    insert x Empty = Node Empty x Empty
    insert x orig@(Node left val right)
      = case compare x val of
          LT => Node (insert x left) val right
          EQ => orig
          GT => Node left val (insert x right)
```

There's no need for an `Ord` constraint on `insert` because there are already `Ord` constraints on `elem` in `BSTree` itself.

PRECISION AND REUSE Putting a constraint in the tree structure itself makes the type more precise, in that it can now *only* store values that can be compared at the nodes, but at the cost of making it less reusable. This is a trade-off you'll often have to consider when defining new data types. There are various ways of managing this trade-off, such as pairing data with *predicates* that describe the form of that data, as you'll see in chapter 9.

Exercises



- 1 Write a function, `listToTree : Ord a => List a -> Tree a`, that inserts every element of a list into a binary search tree.

You can test this at the REPL as follows:

```
*ex_4_1> listToTree [1,4,3,5,2]
Node (Node Empty 1 Empty)
      2
        (Node (Node Empty 3 (Node Empty 4 Empty)))
          5
            Empty) : Tree Integer
```

- 2 Write a corresponding function, `treeToList : Tree a -> List a`, that *flattens* a tree into a list using *in-order* traversal (that is, all the values in the left subtree of a node should be added to the list before the value at the node, which should be added before the values in the right subtree).

If you have the correct answers to exercises 1 and 2, you should be able to run this:

```
*ex_4_1> treeToList (listToTree [4,1,8,7,2,3,9,5,6])
```

```
[1, 2, 3, 4, 5, 6, 7, 8, 9] : List Integer
```

- 3 An integer arithmetic expression can take one of the following forms:

- A single integer
- Addition of an expression to an expression
- Subtraction of an expression from an expression
- Multiplication of an expression with an expression

Define a recursive data type, Expr, that can be used to represent such expressions.

Hint: Look at the Picture data type and see how the informal description mapped to the data declaration.

- 4 Write a function, evaluate : Expr -> Int, that evaluates an integer arithmetic expression.

If you have correct answers to 3 and 4, you should be able to try something like the following at the REPL:

```
*ex_4_1> evaluate (Mult (Val 10) (Add (Val 6) (Val 3)))
90 : Int
```

- 5 Write a function, maxMaybe : Ord a => Maybe a -> Maybe a -> Maybe a, that returns the larger of the two inputs, or Nothing if both inputs are Nothing. For example:

```
*ex_4_1> maxMaybe (Just 4) (Just 5)
Just 5 : Maybe Integer
```

```
*ex_4_1> maxMaybe (Just 4) Nothing
Just 4 : Maybe Integer
```

- 6 Write a function, biggestTriangle : Picture -> Maybe Double, that returns the area of the biggest triangle in a picture, or Nothing if there are no triangles.

For example, you can define the following pictures:

```
testPic1 : Picture
testPic1 = Combine (Primitive (Triangle 2 3))
                  (Primitive (Triangle 2 4))

testPic2 : Picture
testPic2 = Combine (Primitive (Rectangle 1 3))
                  (Primitive (Circle 4))
```

Then, test biggestTriangle at the REPL as follows:

```
*ex_4_1> biggestTriangle testPic1
Just 4.0 : Maybe Double
```

```
*ex_4_1> biggestTriangle testPic2
Nothing : Maybe Double
```

4.2 Defining dependent data types

A *dependent* data type is a type that's computed from some other value. You've already seen a dependent type, `Vect`, where the exact type is calculated from the vector's length:

```
Vect : Nat -> Type -> Type
```

In other words, the type of a `Vect` depends on its length. This gives us additional precision in the type, which we used to help direct our programming via the process of type, define, refine. In this section, you'll see how to define dependent types such as `Vect`. The core of the idea is that, because there's no syntactic distinction between types and expressions, types can be computed from *any* expression.

We'll begin with a simple example to illustrate how this works, defining a type for representing vehicles and their properties, depending on their power source, and you'll see how you can use this to constrain the valid inputs to a function to those that make sense. You'll then see how `Vect` itself is defined, along with some useful operations on it.

4.2.1 A first example: classifying vehicles by power source

Dependent types allow you to give more precise information about the *data* constructors of a type, by adding more arguments to the *type* constructor. For example, you might have a data type to represent vehicles (for example, bicycles, cars, and buses), but some operations don't make sense on all values in the type (for example, refueling a bicycle wouldn't work because there's no fuel tank). We'll therefore classify vehicles into those powered by pedal and those powered by petrol, and express this in the type.

The following listing shows how you could express this in Idris.

Listing 4.10 Defining a dependent type for vehicles, with their power source in the type (vehicle.idr)

An enumeration type describing possible power sources for a vehicle

```
→ data PowerSource = Petrol | Pedal
      data Vehicle : PowerSource -> Type where
        Bicycle : Vehicle Pedal
        Car : (fuel : Nat) -> Vehicle Petrol
        Bus : (fuel : Nat) -> Vehicle Petrol
```

A vehicle powered by pedal

A Vehicle's type is annotated with its power source.

A vehicle powered by petrol, with a field for current fuel stocks

You can write functions that will work on all vehicles by using a type variable to stand for the power source. For example, all vehicles have a number of wheels. On the other hand, not all vehicles carry fuel, so it only makes sense to refuel a vehicle whose type indicates it's powered by `Petrol`. Both of these concepts are illustrated in the following listing.

Listing 4.11 Reading and updating properties of Vehicle

```
wheels : Vehicle power -> Nat
wheels Bicycle = 2
wheels (Car fuel) = 4
wheels (Bus fuel) = 4
```

← Use a type variable, power,
because this function works
for all possible vehicle types.


```
refuel : Vehicle Petrol -> Vehicle Petrol
refuel (Car fuel) = Car 100
refuel (Bus fuel) = Bus 200
```

← Refueling only makes sense for vehicles
that carry fuel, so restrict the input and
output type to Vehicle Petrol.

Asserting that inputs are impossible

If you try adding a case for refueling a Bicycle, Idris will report a type error, because the input type is restricted to vehicles powered by petrol. If you use the interactive tools, Idris won't even give a case for Bicycle after a case split with Ctrl-Alt-C. Nevertheless, it can sometimes aid readability to make it explicit that you know the Bicycle case is impossible. You can write this:

```
refuel : Vehicle Petrol -> Vehicle Petrol
refuel (Car fuel) = Car 100
refuel (Bus fuel) = Bus 200
refuel Bicycle impossible
```

If you do this, Idris will check that the case you have marked as impossible would produce a type error.

Similarly, if you assert a case is impossible but Idris believes it's valid, it will report an error:

```
refuel : Vehicle Petrol -> Vehicle Petrol
refuel (Car fuel) = Car 100
refuel (Bus fuel) impossible
```

Here, Idris will report the following:

```
vehicle.idr:15:8:refuel (Bus fuel) is a valid case
```

In general, you should define dependent types by giving the type constructor and the data constructors directly. This gives you a lot of flexibility in the form that the constructors can take. Here, it has allowed you to define types where the data constructors can each take different arguments. You can either write functions that work on all vehicles (like wheels) or functions that only work on some subset of vehicles (like refuel).

DEFINING FAMILIES OF TYPES For Vehicle, you've actually defined two types in one declaration (specifically, Vehicle Pedal and Vehicle Petrol). Dependent data types like Vehicle are therefore sometimes referred to as *families* of types, because you're defining multiple related types at the same time. The power source is an index of the Vehicle family. The index tells you exactly which Vehicle type you mean.

4.2.2 Defining vectors

In chapter 3, we looked at the ways we could use the length information in the type to help drive development of functions on vectors. In this section, we'll look at how `Vect` is defined, along with some of the operations on it.

It's defined in the `Data.Vect` module, as shown in listing 4.12. The type constructor, `Vect`, takes a length and an element type as arguments, so when you define the data constructors, you state explicitly in their types what their lengths are.

Listing 4.12 Defining vectors (`Vect.idr`)

The type constructor `Vect` states that a vector constructor takes two arguments: a `Nat`, which is its length, and a `Type`, which is its element type.

```
▷ data Vect : Nat -> Type -> Type where
    Nil   : Vect Z a
    (::)  : (x : a) -> (xs : Vect k a) -> Vect (S k) a
```

The data constructor `Nil` explicitly states that an empty vector has length `Z`.

`%name Vect xs, ys, zs`

Gives some default names to use in case splits

The data constructor `(::)` explicitly states that adding an element `x` to a vector of length `k` results in a vector of length `S k` (that is, $l + k$).

The `Data.Vect` library includes several utility functions on `Vect`, including concatenation, looking up values by their position in the vector, and various higher-order functions, such as `map`. Instead of importing this, though, we'll use our own definition of `Vect` and try writing some functions by hand. To begin, create a `Vect.idr` file containing only the definition of `Vect` in listing 4.12.

Because `Vect` includes the length explicitly in its type, any function that uses some instance of `Vect` will describe its length properties explicitly in its type. For example, if you define an `append` function on `Vect`, its type will express how the lengths of the inputs and output are related:

```
append : Vect n elem -> Vect m elem -> Vect (n + m) elem
```

TYPE-LEVEL EXPRESSIONS The expression `n + m` in the return type here is an ordinary expression with type `Nat`, using the ordinary `+` operator. Because `Vect`'s first argument is of type `Nat`, you should expect to be able to use any expression of type `Nat`. Remember, types are first-class, so types and expressions are all part of the same language.

Having written the type first, as always, you can define `append` by case splitting on the first argument. You do this interactively as follows:

- 1 *Define*—Begin by creating a skeleton definition and then case splitting on the first argument `xs`:

```
append : Vect n elem -> Vect m elem -> Vect (n + m) elem
append [] ys = ?append_rhs_1
append (x :: xs) ys = ?append_rhs_2
```

- 2 Refine**—Idris has enough information in the type to complete this definition by an expression search on each of the holes, resulting in this:

```
append : Vect n elem -> Vect m elem -> Vect (n + m) elem
append [] ys = ys
append (x :: xs) ys = x :: append xs ys
```

Terminology: parameters and indices

`Vect` defines a family of types, and we say that a `Vect` is indexed by its length and parameterized by an element type. The distinction between parameters and indices is as follows:

- A *parameter* is unchanged across the entire structure. In this case, every element of the vector has the same type.
- An *index* may change across a structure. In this case, every subvector has a different length.

The distinction is most useful when looking at a function's type: you can be certain that the specific value of a parameter can play no part in a function's definition. The index, however, might, as you've already seen in chapter 3 when defining `length` for vectors by looking at the length index, and when defining `createEmpties` for building a vector of empty vectors.

Another common operation on vectors is `zip`, which pairs corresponding elements in two vectors, as illustrated in figure 4.5.




Figure 4.5 Pairing corresponding elements of `[1, 2, 3, 4]` and `["one", "two", "three", "four"]` using `zip`

The name `zip` is intended to suggest the workings of a zip fastener, bringing two sides of a bag or jacket together. Because the length of each input `Vect` is in the type, you need to think about how the lengths of the inputs and output will correspond. A reasonable choice for this would be to require the lengths of both inputs to be the same:

```
zip : Vect n a -> Vect n b -> Vect n (a, b)
```

Having a more precise type for `Vect`, capturing the length in the type, means that you need to decide in advance how the lengths of the inputs to `zip` relate and express this decision in the type. Also, it means you can safely assume that `zip` will only ever be called with equal length lists, because if this assumption is violated, Idris will report a type error.

You can define `zip`, as usual, step by step:

- 1 Define**—Again, you can begin to define this by a case split on the first argument:

```
zip : Vect n a -> Vect n b -> Vect n (a, b)
zip [] ys = ?zip_rhs_1
zip (x :: xs) ys = ?zip_rhs_2
```

- 2 Refine**—You can fill in `?zip_rhs_1` with an expression search, because the only well-typed result is an empty vector:

```
zip : Vect n a -> Vect n b -> Vect n (a, b)
zip [] ys = []
zip (x :: xs) ys = ?zip_rhs_2
```

- 3 Refine**—For the second case, `?zip_rhs_2`, take a look at the type of the hole and see if that gives you further information about what to do:

```
b : Type
a : Type
x : a
k : Nat
xs : Vect k a
ys : Vect (S k) b
-----
zip_rhs_2 : Vect (S k) (a, b)
```

Notice that `ys` has length `S k`, meaning that there must be at least one element. If you case-split on `ys`, Idris won't give you a pattern for the empty list, because it wouldn't be a well-typed value:

```
zip : Vect n a -> Vect n b -> Vect n (a, b)
zip [] ys = []
zip (x :: xs) (y :: ys) = ?zip_rhs_1
```

After the case split, Idris has created a new hole, so let's take a look at the types of the local variables:

```
b : Type
a : Type
x : a
k : Nat
```

```

xs : Vect k a
y : b
ys : Vect k b
-----
zip_rhs_1 : Vect (S k) (a, b)

```

Again, there's enough information to complete the definition with an expression search:

```

zip : Vect n a -> Vect n b -> Vect n (a, b)
zip [] ys = []
zip (x :: xs) (y :: ys) = (x, y) :: zip xs ys

```

Idris has noticed that it needs to build a vector of length $S\ k$, that it can create the appropriate vector of length k with a recursive call, and that it can create the appropriate first element by pairing x and y .

Deconstructing expression searches

To understand what expression search has done, it can be instructive to remove some part of the result and replace it with a hole, to see what type expression search was working with at that point. For example, you can remove the (x, y) :

```

zip : Vect n a -> Vect n b -> Vect n (a, b)
zip [] ys = []
zip (x :: xs) (y :: ys) = ?element :: zip xs ys

```

Then, checking the type of the $?element$ hole, you'll see that at this point, Idris was looking for a pair of a and b :

```

b : Type
a : Type
x : a
k : Nat
xs : Vect k a
y : b
ys : Vect k b
-----
element : (a, b)

```

The only way to make a pair of a and b at this point is to use x and y , so this is what Idris used to construct the pair.

4.2.3 Indexing vectors with bounded numbers using Fin

Because Vects carry their length as part of their type, the type checker has additional knowledge that it can use to check that operations are implemented and used correctly. One example is that if you wish to look up an element in a Vect by its location in the vector, you can know at compile time that the location can't be out of bounds when the program is run.

The `index` function, defined in `Data.Vect`, is a bounds-safe lookup function whose type guarantees that it will never access a location that's outside the bounds of a vector:

```
index : Fin n -> Vect n a -> a
```

The first argument, of type `Fin n`, is an unsigned number that has a non-inclusive *upper bound* of `n`. The name `Fin` suggests that the number is *finitely bounded*. So, for example, when you look up an element by location, you can use a number within the bounds of the vector:

```
Idris> :module Data.Vect
*Data/Vect> Vect.index 3 [1,2,3,4,5]
```

IMPORTING MODULES AT THE REPL In the element lookup example, you import `Data.Vect` at the REPL using the `:module` command, to get access to the `index` function. There are several functions in the Prelude called `index` for indexing different list-like structures, so you have to disambiguate explicitly with `Vect.index` here.

But if you try to use a number outside the bounds, you'll get a type error:

```
*Data/Vect> Vect.index 7 [1,2,3,4,5]
(input):1:14:When checking argument prf to function Data.Fin.fromInteger:
    When using 7 as a literal for a Fin 5
        7 is not strictly less than 5
```

INTEGER LITERAL NOTATION As with `Nat`, you can use integer literals for `Fin`, provided that the compiler can be sure that the literal is within the bounds stated in the type.

If you're reading a number as user input that will be used to `index` a `Vect`, the number won't always be within the bounds of the `Vect`. In practice, you'll often need to convert from an arbitrarily sized `Integer` to a bounded `Fin`.

Importing `Data.Vect` gives you access to the `integerToFin` function, which converts an `Integer` to a `Fin` with some bounds, provided the `Integer` is within bounds. It has the following type:

```
integerToFin : Integer -> (n : Nat) -> Maybe (Fin n)
```

The first argument is the integer to be converted, and the second is the upper bound on the `Fin`. Remember that a type `Fin upper`, for some value of `upper`, represents numbers up to but not including `upper`, so `5` is not a valid `Fin 5`, but `4` is. Here are a couple of examples:

```
*Data/Vect> integerToFin 2 5
Just (FS (FS FZ)) : Maybe (Fin 5)

*Data/Vect> integerToFin 6 5
Nothing : Maybe (Fin 5)
```

FIN CONSTRUCTORS `FZ` and `FS` are constructors of `Fin`, corresponding to `Z` and `S` as constructors of `Nat`. Typically, you can use numeric literals, as with `Nat`.

Using `integerToFin`, you can write a `tryIndex` function that looks up a value in a `Vect` by `Integer` index, using `Maybe` in the type to capture the possibility that the result may be out of range. Begin by creating a `TryIndex.idr` file that imports `Data.Vect`. Then, follow these steps:

- 1 Type**—As ever, start by giving a type:

```
tryIndex : Integer -> Vect n a -> Maybe a
```

Note that this type gives no relationship between the input `Integer` and the length of the `Vect`.

- 2 Define**—You can write the definition by using `integerToFin` to check whether the input is within range:

```
tryIndex : Integer -> Vect n a -> Maybe a
tryIndex {n} i xs = case integerToFin i n of
    case_val => ?tryIndex_rhs
```

Note that you need to bring `n` into scope so that you can pass it to `integerToFin` as the desired bound of the `Fin`.

- 3 Define**—Now, define the function by case splitting on `case_val`. If `integerToFin` returns `Nothing`, the input was out of bounds, so you return `Nothing`:

```
tryIndex : Integer -> Vect n a -> Maybe a
tryIndex {n} i xs = case integerToFin i n of
    Nothing => Nothing
    Just idx => ?tryIndex_rhs_2
```

- 4 Type**—If you inspect the type of `tryIndex_rhs_2`, you'll see that you now have a `Fin n` and a `Vect n a`, so you can safely use `index`:

```
n : Nat
idx : Fin n
a : Type
i : Integer
xs : Vect n a
-----
tryIndex_rhs_2 : Maybe a
```

The end result is as follows:

```
tryIndex : Integer -> Vect n a -> Maybe a
tryIndex {n} i xs = case integerToFin i n of
    Nothing => Nothing
    Just idx => Just (index idx xs)
```

This is a common pattern in dependently typed programming, which you'll see more often in the coming chapters. The type of `index` tells you when it's safe to call it, so if you have an input that's potentially unsafe, you need to check. Once you've converted the `Integer` to a `Fin n`, you know that number must be within bounds, so you don't need to check again.

Exercises



- 1 Extend the Vehicle data type so that it supports unicycles and motorcycles, and update wheels and refuel accordingly.
- 2 Extend the PowerSource and Vehicle data types to support electric vehicles (such as trams or electric cars).
- 3 The take function, on List, has type Nat → List a → List a. What's an appropriate type for the corresponding vectTake function on Vect?

Hint: How do the lengths of the input and output relate? It shouldn't be valid to take more elements than there are in the Vect. Also, remember that you can have *any* expression in a type.

- 4 Implement vectTake. If you've implemented it correctly, with the correct type, you can test your answer at the REPL as follows:

```
*ex_4_2> vectTake 3 [1,2,3,4,5,6,7]
[1, 2, 3] : Vect 3 Integer
```

You should also get a type error if you try to take too many elements:

```
*ex_4_2> vectTake 8 [1,2,3,4,5,6,7]
(input):1:14:When checking argument xs to constructor Main.:::
  Type mismatch between
    Vect 0 a1 (Type of [])
  and
    Vect (S m) a (Expected type)
```

- 5 Write a sumEntries function with the following type:

```
sumEntries : Num a => (pos : Integer) -> Vect n a -> Vect n a -> Maybe a
```

It should return the sum of the entries at position pos in each of the inputs if pos is within bounds, or Nothing otherwise. For example:

```
*ex_4_2> sumEntries 2 [1,2,3,4] [5,6,7,8]
Just 10 : Maybe Integer
```

```
*ex_4_2> sumEntries 4 [1,2,3,4] [5,6,7,8]
Nothing : Maybe Integer
```

Hint: You'll need to call integerToFin, but you'll only need to do it once.

4.3 Type-driven implementation of an interactive data store

To into practice put the ideas you've learned so far, let's now take a look at a larger example program, an interactive data store. In this section, we'll set up the basic infrastructure. We'll revise this program in chapter 6, as you learn more about Idris, to support key-value pairs, with schemas for describing the form of the data.

In our initial implementation, we'll only support storing data as Strings, in memory, accessed by a numeric identifier. It will have a command prompt and support the following commands:

- add [String] adds a string to the document store and responds by printing an identifier by which you can refer to the string.
- get [Identifier] retrieves and prints the string with the given identifier, assuming the identifier exists, or an error message otherwise.
- quit exits the program.

A brief session might go as follows:

```
$ ./datastore
Command: add Even Old New York
ID 0
Command: add Was Once New Amsterdam
ID 1
Command: get 1
Was Once New Amsterdam
Command: get 2
Out of range
Command: add Why They Changed It I Can't Say
ID 2
Command: get 2
Why They Changed It I Can't Say
Command: quit
```

We'll use the type system to guarantee that all accesses to the data store use valid identifiers, and all of our functions will be total so we're sure that the program won't abort due to unexpected input.

The overall approach we'll take, at a high level, again follows the process of type, define, refine:

- *Type*—Design a new data type for the representation of the data store. In type-driven development, even at the highest level, types come first. Before we can implement any part of a data store program, we need to know how we're representing and working with the data.
- *Define*—Implement as much of a main program as we can, leaving holes for parts we can't write immediately, lifting these holes to top-level functions.
- *Refine*—As we go deeper into the implementation and improve our understanding of the problem, we'll refine the implementation and types as necessary.

To begin, create an outline of a DataStore.idr file that contains a module header, an import statement for Data.Vect, and an empty main function, as follows.

Listing 4.13 Outline implementation of the data store (DataStore.idr)

```
module Main
import Data.Vect
main : IO ()
main = ?main_rhs
```

Annotations for Listing 4.13:

- An annotation pointing to the import statement: "Use Vect to store the data so you can keep track of the size of the store in types".
- An annotation pointing to the main function: "Initial implementation of main is empty".

We'll start by defining a type to represent the store, and then we'll write a `main` function that reads user input and updates the store according to user commands. As we progress through the implementation, we'll add new types and functions as necessary, always guided by the Idris type checker.

4.3.1 Representing the store

The data store itself is, initially, a collection of strings. We'll begin by defining a type for the store, including the size of the store (that is, the number of stored items), explicitly and using a `Vect` for the items, as shown in the following listing. You can add this to `DataStore.idr`, above the initial empty definition of `main`.

Listing 4.14 A data type for representing the data store (DataStore.idr)

```
data DataStore : Type where      ← DataStore is the type constructor.
  MkData : (size : Nat) ->
    (items : Vect size String) -> DataStore
```

MkData is the data constructor, which gives the canonical way of constructing a data store. You can think of its arguments, size, and items as the fields of a record.

The type explicitly states that the length of this Vect is the size of the store.

You can access the size and content of a data store by writing functions that pattern match on the data store and extract the appropriate fields. These are shown in the following listing.

Listing 4.15 Projecting out the size and content of a data store (DataStore.idr)

```
size : DataStore -> Nat
size (MkData size' items') = size'

items : (store : DataStore) -> Vect (size store) String ←
items (MkData size' items') = items'
```

The length in the type of the Vect projected out of the store is given by the size projected out of the store.

In this listing, the length of the `Vect` in the type of `items` is calculated by a function, `size`.

RECORDS A data type with one constructor, like `DataStore`, is essentially a record with fields for its data. In chapter 6, you'll see a more concise syntax for records that avoids the need to write explicit projection functions such as `size` and `items`.

You'll also need to add new data items to the store, as in listing 4.16. This adds items to the *end* of the store, rather than at the beginning using `::` directly. The reason for this is that we plan to access items by their integer index; if you add items at the beginning, new items will always have an index of zero, and everything else will be shifted along one place.

Listing 4.16 Adding a new entry to the data store (DataStore.idr)

You can use `_` here, rather than giving the size explicitly, because Idris can work out the new size at compile time from the type of `addToData`.

```
addToStore : DataStore -> String -> DataStore
addToStore (MkData size items) newItem = MkData _ (addToData items)
  where
    addToData : Vect old String -> Vect (S old) String
    addToData [] = [newItem]                                ←
    addToData (item :: items) = item :: addToData items
```

This type states that `addToData` always increases the length of the store by 1.

In a `where` block, functions have access to the pattern variables of the outer function, so you can use `newItem` here.

Interactive editing in where blocks

The interactive editing tools work just as effectively in `where` blocks as they do at the top level. For example, try implementing `addToStore` beginning at this point:

```
addToStore : DataStore -> String -> DataStore
addToStore (MkData size items) newItem
  = MkData _ (addToData items)
  where
    addToData : Vect old String -> Vect (S old) String
```

You can use Ctrl-Alt-A to add a definition for `addToData`, and expression search with Ctrl-Alt-S is aware that `newItem` is in scope.

4.3.2 Interactively maintaining state in main

When you implement the `main` function for your data store, you'll need to read input from the user, maintain the state of the store itself, and allow the user to exit. In the last chapter, we used the Prelude function `repl` to write a simple interactive program that repeatedly read input, ran a function on it, and displayed the output:

```
repl : String -> (String -> String) -> IO ()
```

Unfortunately, this only allows simple interactions that repeat forever.

For more-complex programs that maintain state, the Prelude provides another function, `replWith`, which implements a read-eval-print loop that carries some state. `:doc` describes it as follows:

```
Idris> :doc replWith
Prelude.Interactive.replWith : (state : a) ->
  (prompt : String) ->
  (onInput : a -> String -> Maybe (String, a)) -> IO ()
```

```

A basic read-eval-print loop, maintaining a state
Arguments:
    state : a -- the input state

    prompt : String -- the prompt to show

    onInput : a -> String -> Maybe (String, a) -- the function to
    run on reading input, returning a String to output and a new
    state. Returns Nothing if the repl should exit

```

On each iteration through the loop, it calls the `onInput` argument, which itself takes two arguments:

- The current state, of some generic type `a`
- The `String` entered at the prompt

The value the `onInput` function should return is of type `Maybe (String, a)`, meaning that it can be in one of the following forms:

- `Nothing`, if it wants the loop to exit
- `Just (output, newState)`, if it wants to print some `output` and update the state to `newState` for the next iteration

The next listing shows a simple example of this in action: an interactive program that reads an integer from the console and displays a running total of the sum of the inputs. If it reads a negative value, it will exit.

Listing 4.17 Interactive program to sum input values until a negative value is read (`SumInputs.idr`)

Casts the input `String` to an `Integer`. Default value is `0`
if the input is not a valid number. Idris knows that `val`
must be an `Integer` because it's used later in a context
where only `Integer` is well typed.

Negative input
received, so returns
`Nothing`, which
exits the loop

```

sumInputs : Integer -> String -> Maybe (String, Integer)
sumInputs tot inp
  = let val = cast inp in
    if val < 0
      then Nothing
      else let newVal = tot + val in
            Just ("Subtotal: " ++ show newVal ++ "\n", newVal)

```

Initializes
the loop
with `0` as
the initial
state

```

main : IO ()
main = replWith 0 "Value: " sumInputs

```

Calculates a new state
(`newVal`) and continues the
loop, giving the subtotal as an
output, and the new state

You can use `replWith` to refine the `main` function in the data store. At this stage you have

- A type for the data store (`DataStore`)
- A way of accessing the items in the store (`items`)
- A way of updating the store with new items (`addToStore`)

When you call `replWith`, you need to pass the initial data, a prompt, and a function to process inputs as arguments. You can pass an initialized empty data store and a prompt string, but you don't yet have a function to process user input. Nevertheless, you can refine your definition of `main` to the following, leaving a hole for the input-processing function:

```
main : IO ()
main = replWith (MkData _ []) "Command: " ?processInput
```

Lifting `processInput` shows the type you have to work with:

```
processInput : DataStore -> String -> Maybe (String, DataStore)
main : IO ()
main = replWith (MkData _ []) "Command: " processInput
```

Following the type-driven approach, when you refined the definition of `main` with an application of `replWith`, Idris was able to work out the necessary specialized type for `processInput`.

4.3.3 Commands: parsing user input

To process the `String` input, you somehow have to work out which one of the commands `add`, `get`, or `quit` has been entered. Rather than processing the input string *directly*, it's usually much cleaner to define a new data type that represents the possible commands. This way, you can cleanly separate the parsing of commands from the processing.

You'll therefore define a `Command` data type, which is a union type representing the possible commands and their arguments. Put the following definition in `DataStore.idr` above `processInput`:

```
data Command = Add String
             | Get Integer
             | Quit
```

AVOID USING STRINGS FOR DATA REPRESENTATION The user enters a `String`, but only a certain number of `Strings` are valid input commands. Introducing the `Command` type makes the representation of commands more precise in that only valid commands can be represented. If the user enters a `String` that can't be converted to a `Command`, the types force you to think about how to handle the error. At first, you can leave a hole in the program for error handling, but, ultimately, making a precise type leads to a more robust implementation.

You need to convert the string input by the user into a `Command`. The input may be invalid, however, so the type of the function to parse the command captures this possibility in its type:

```
parse : (input : String) -> Maybe Command
```

You can write a simple parser for input commands by searching for the first space in the input using the `span` function to establish which part of the input is the command and which is the argument. The `span` function works as follows:

```
Idris> :t Strings.span
: (Char -> Bool) -> String -> (String, String)

Idris> span (/= ' ') "Hello world, here is a string"
("Hello", " world, here is a string") : (String, String)
```

The first argument, `(/= ' ')`, is a test that returns a `Bool`. This test returns `True` for any character that's not equal to a space. The second argument is an input string, and `span` will split the string into two parts:

- The first part is the prefix of the input string, where all characters satisfy the test.
- The second part is the remainder of the string. If all characters in the string satisfy the test, this will be empty.

You can define `parse` with Ctrl-Alt-A and then refine its definition to the following:

```
parse : (input : String) -> Maybe Command
parse input = case span (/= ' ') input of
    (cmd, args) => ?parseCommand
```

You can then lift `parseCommand` to a top-level function with the appropriate type, using Ctrl-Alt-L:

```
parseCommand : (cmd : String) -> (args : String) -> (input : String) ->
    Maybe Command

parse : (input : String) -> Maybe Command
parse input = case span (/= ' ') input of
    (cmd, args) => parseCommand cmd args input
```

You won't need the `input` argument because you'll be parsing the input from `cmd` and `args` alone, although Idris has added it because `input` is in scope. You can therefore edit the type:

```
parseCommand : (cmd : String) -> (args : String) -> Maybe Command

parse : (input : String) -> Maybe Command
parse input = case span (/= ' ') input of
    (cmd, args) => parseCommand cmd args
```

Also, you can see that `args`, if it's not empty, will have a leading space, because the first character that `span` encounters that doesn't satisfy the test will be a space. You can remove leading spaces with the `ltrim : String -> String` function, which returns its input with leading whitespace characters removed:

```
parseCommand : (cmd : String) -> (args : String) -> Maybe Command

parse : (input : String) -> Maybe Command
parse input = case span (/= ' ') input of
    (cmd, args) => parseCommand cmd (ltrim args)
```

You can now write `parseCommand` by examining the `cmd` and `args` arguments. You'll need some Prelude functions to complete the definition of `parseCommand`:

- `unpack` : `String -> List Char` converts a `String` into a list of characters.
- `isDigit` : `Char -> Bool` returns whether a `Char` is one of the digits 0–9.
- `all` : `(a -> Bool) -> List a -> Bool` returns whether every entry in a list satisfies a test.

Thus, the expression `all isDigit (unpack val)` returns whether the string `val` consists entirely of digits.

DOCUMENTATION AND PRELUDE FUNCTIONS In general, remember that you can use `:doc` at the REPL, or Ctrl-Alt-D in Atom, to check the documentation for any names, no matter whether they are type constructors, data constructors, or functions.

Listing 4.18 shows how parsing the input works. In particular, notice that pattern matching is very general; as long as patterns are composed of primitive ways of constructing a type (data constructors and primitive values), they are valid. An underscore is a match-anything pattern.

Listing 4.18 Parsing a command and argument string into a Command (DataStore.idr)

```
parseCommand : String -> String -> Maybe Command
parseCommand "add" str = Just (Add str)
parseCommand "get" val = case all isDigit (unpack val) of
    False => Nothing
    True => Just (Get (cast val))
parseCommand "quit" "" = Just Quit
parseCommand _ _ = Nothing
```

This pattern matches an application where the first argument is "add". The string you add to the store will be composed of the entire second argument, str.

This matches any input. Pattern matching works top to bottom; if none of the previous patterns have matched, then the input is invalid, so there's no Command.

This pattern matches an application where the first argument is "get". The parse is valid if the second argument consists entirely of digits.

Now that you can parse a string into a `Command`, you can make more progress with `processInput`, calling `parse`. If it fails, you display an error message and leave the store as it is. Otherwise, you add a hole for processing the `Command`:

```
processInput : DataStore -> String -> Maybe (String, DataStore)
processInput store inp
  = case parse inp of
      Nothing => Just ("Invalid command\n", store)
      Just cmd => ?processCommand
```

4.3.4 Processing commands

One way to proceed with your implementation of `processInput` would be to case-split on `cmd` and process the commands directly:

```
processInput : DataStore -> String -> Maybe (String, DataStore)
processInput store inp
  = case parse inp of
      Nothing => Just ("Invalid command\n", store)
      Just (Add item) => ?processCommand_1
      Just (Get pos) => ?processCommand_2
      Just Quit => ?processCommand_3
```

You can refine the implementation by filling in the easier holes, in the cases for `Add item` and `Quit`. The next listing shows how these holes are refined, leaving `processCommand_2` for the moment.

Listing 4.19 Processing the inputs Add item and Quit (DataStore.idr)

```
processInput : DataStore -> String -> Maybe (String, DataStore)
processInput store inp
  = case parse inp of
      Nothing => Just ("Invalid command\n", store)
      Just (Add item) =>
        Just ("ID " ++ show (size store) ++ "\n", addToStore store item)
      Just (Get pos) => ?processCommand_2
      Just Quit => Nothing
```

Returns a String that gives the position at which this item was added and an updated store using `addToStore`

Quit command exits,
so return Nothing

It's always a good idea to inspect the types of holes Idris has generated to see what variables you have available, what their types are, and what type you need to build. For `?processCommand_2`, you have this:

```
pos : Integer
store : DataStore
input : String
-----
processCommand_2 : Maybe (String, DataStore)
```

This will be slightly more involved than the other cases. You need to do the following:

- Get the items from the store.
- Ensure that the position, `pos`, is in range.
- If it is, extract the item from the specified `pos` and return it along with the store itself.

The type-define-refine process encourages you to write parts of definitions, step by step, constantly type-checking as you go, and constantly inspecting the types of the holes.

Because there are a few details involved in filling the `?processCommand_2` hole, we'll rename it `?getEntry` and lift it to a top-level function before implementing it step by step:

- 1 Type**—Lift `getEntry` to a new top-level function:

```
getEntry : (pos : Integer) -> (store : DataStore) -> (input : String) ->
    Maybe (String, DataStore)
```

- 2 Define**—Create a new skeleton definition:

```
getEntry pos store input = ?getEntry_rhs
```

- 3 Define**—You'll need the `items` in the store, so define a new local variable for these:

```
getEntry pos store input = let store_items = items store in
    ?getEntry_rhs
```

Inspecting the type of `getEntry_rhs` tells you the type of `store_items`:

```
pos : Integer
store : DataStore
input : String
store_items : Vect (size store) String
-----
getEntry_rhs : Maybe (String, DataStore)
```

- 4 Refine**—To retrieve an entry from a `Vect`, you can use `index` as you've previously seen:

```
index : Fin n -> Vect n a -> a
```

To extract an entry from `store_items`, which has type `Vect (size store)`, you'll need a `Fin (size store)`. Unfortunately, all you have available at the moment is an `Integer`. Using `integerToFin`, as described in section 4.2.34.2.3, you can refine the definition. If `integerToFin` returns `Nothing`, the input was out of bounds.

```
getEntry : (pos : Integer) -> (store : DataStore) -> (input : String) ->
    Maybe (String, DataStore)
getEntry pos store input
= let store_items = items store in
    case integerToFin pos (size store) of
        Nothing => Just ("Out of range\n", store)
        Just id => ?getEntry_rhs_2
```

- 5 Refine**—If you inspect the type of `getEntry_rhs_2` now, you'll see that you have the `Fin (size store)` you need:

```
store : DataStore
id : Fin (size store)
pos : Integer
input : String
store_items : Vect (size store) String
-----
```

```
getEntry_rhs_2 : Maybe (String, DataStore)
```

You can now refine to a complete definition:

```
getEntry : (pos : Integer) -> (store : DataStore) -> (input : String)
          Maybe (String, DataStore)
getEntry pos store input
  = let store_items = items store in
    case integerToFin pos (size store) of
      Nothing => Just ("Out of range\n", store)
      Just id => Just (index id store_items ++ "\n", store)
```

- 6 Refine**—As a final refinement, observe that `input` is never used, so you can remove this argument. Don’t forget to remove it from the application of `getEntry` in `processInput` too.

```
getEntry : (pos : Integer) -> (store : DataStore) ->
          Maybe (String, DataStore)
getEntry pos store
  = let store_items = items store in
    case integerToFin pos (size store) of
      Nothing => Just ("Out of range\n", store)
      Just id => Just (index id store_items ++ "\n", store)
```

By using a `Vect` for the store, with its size as part of the type, the type system can ensure that any access of the store by index will be within bounds, because you have to show that the index has the same upper bound as the length of the `Vect`.

For reference, the complete implementation of the data store is given in the following listing, with all the functions we just worked through.

Listing 4.20 Complete implementation of a simple data store (DataStore.idr)

```
module Main

import Data.Vect

data DataStore : Type where
  MkData : (size : Nat) -> (items : Vect size String) -> DataStore

  size : DataStore -> Nat
  size (MkData size' items') = size'

  items : (store : DataStore) -> Vect (size store) String
  items (MkData size' items') = items'

  addToStore : DataStore -> String -> DataStore
  addToStore (MkData size store) newItem = MkData _ (addToData store)
    where
      addToData : Vect oldsize String -> Vect (S oldsize) String
      addToData [] = [newItem]
      addToData (x :: xs) = x :: addToData xs

  data Command = Add String
                | Get Integer
                | Quit
```

```

parseCommand : String -> String -> Maybe Command
parseCommand "add" str = Just (Add str)
parseCommand "get" val = case all isDigit (unpack val) of
                           False => Nothing
                           True => Just (Get (cast val))
parseCommand "quit" "" = Just Quit
parseCommand _ _ = Nothing

parse : (input : String) -> Maybe Command
parse input = case span (/= ' ') input of
                  (cmd, args) => parseCommand cmd (ltrim args)

getEntry : (pos : Integer) -> (store : DataStore) ->
           Maybe (String, DataStore)
getEntry pos store
  = let store_items = items store in
    case integerToFin pos (size store) of
      Nothing => Just ("Out of range\n", store)
      Just id => Just (index id (items store) ++ "\n", store)

processInput : DataStore -> String -> Maybe (String, DataStore)
processInput store input
  = case parse input of
      Nothing => Just ("Invalid command\n", store)
      Just (Add item) =>
        Just ("ID " ++ show (size store) ++ "\n", addToStore store item)
      Just (Get pos) => getEntry pos store
      Just Quit => Nothing

main : IO ()
main = replWith (MkData _ []) "Command: " processInput

```

Exercises



- 1 Add a size command that displays the number of entries in the store.
- 2 Add a search command that displays all the entries in the store containing a given substring.

Hint: Use `Strings.isInfixOf`.

- 3 Extend search to print the location of each result, as well as the string.

You can test your solution at the REPL as follows:

```

*ex_4_3> :exec
Command: add Shearer
ID 0
Command: add Milburn
ID 1
Command: add White
ID 2
Command: size
3
Command: search Mil
1: Milburn

```

4.4 Summary

- Data types are defined in terms of a type constructor and data constructors.
- Enumeration types are defined by listing the data constructors of the type.
- Union types are defined by listing the data constructors of the type, each of which may carry additional information.
- Generic types are parameterized over some other type. In a generic type definition, variables stand in place of concrete types.
- Dependent types can be indexed over any other value.
- Using dependent types, you can classify a larger family of types (such as vehicles) into smaller subgroups (such as vehicles powered by petrol and those powered by pedal) in the same declaration.
- Dependent types allow safety checks to be guaranteed at compile time, such as guaranteeing that all vector accesses are within the bounds of the vector.
- You can write larger programs in the type-driven style, creating new data types where appropriate to help describe components of the system.
- Interactive programs that involve state can be written using the `replWith` function.

5

Interactive programs: input and output processing

This chapter covers

- Writing interactive console programs
- Distinguishing evaluation and execution
- Validating user inputs to dependently typed functions

Idris is a pure language, meaning that functions don't have side effects, such as updating global variables, throwing exceptions, or performing console input or output. Realistically, though, when we put functions together to make complete programs, we'll need to interact with users somehow.

In the preceding chapters, we used the `repl` and `replWith` functions to write simple, looping interactive programs that we can compile and execute without worrying too much about how they work. For all but the simplest programs, however, this approach is very limiting. In this chapter, you'll see how interactive programs work in Idris more generally. You'll see how to process and validate user input, and how you can write interactive programs that work with dependent types.

The key idea that allows us to write interactive programs in Idris, despite it being a pure language, is that we distinguish between *evaluation* and *execution*. We write interactive programs using a generic type `IO`, which describes sequences of actions that are then executed by the Idris runtime system.

5.1 Interactive programming with `IO`

Although you can't write functions that interact with a user directly, you can write functions that *describe* sequences of interactions. Once you have a description of a sequence of interactions, you can pass it to the Idris runtime environment, which will *execute* those actions.

The Prelude provides a generic type, `IO`, that allows you to describe interactive programs that return a value:

```
Idris> :doc IO
IO : (res : Type) -> Type
    Interactive programs, describing I/O actions and returning a
    value.
Arguments:
    res : Type -- The result type of the program
```

Thus, you distinguish *in the type* between functions that describe interactions with a user, and functions that return a value directly.

For example, consider the difference between the function types `String -> Int` and `String -> IO Int`:

- `String -> Int` is the type of a function that takes a `String` as an input and returns an `Int` without any user interaction or side effects.
- `String -> IO Int` is the type of a function that takes a `String` as an input and returns a *description* of an interactive program that produces an `Int`.

These are a couple of examples of functions with these types:

- `length : String -> Int`, which is defined in the Prelude, returns the length of the input `String`.
- `readAndGetLength : String -> IO Int` returns a description of interactive actions that display the input `String` as a prompt, reads another `String` from the console, and then returns the length of that `String`.

As you've seen in earlier chapters, the entry point to an Idris application is `main : IO ()`, and we've used this to write simple interactive loops using the `repl` and `replWith` functions without worrying too much about the details of how these functions work. It's now time to look at `IO` in more detail, and to learn how to write more-complex interactive programs.

Listing 5.1 shows an example of an interactive program that returns a description of the actions to display a prompt, read a user's name, and then display a greeting. We'll cover the details of the syntax throughout this section, but for now, notice the type of `main : IO ()`. This type states that `main` takes no input and returns a description of interactive actions that produce an empty tuple.

Listing 5.1 A simple interactive program, reading a user's name and displaying a greeting (Hello.idr)

```
module Main
main : IO ()
main = do
    putStr "Enter your name: "
    x <- getLine
    putStrLn ("Hello " ++ x ++ "!")
→ getLine reads a String
from the console.
```

The keyword `do` introduces a sequence of interactive actions.

`putStr` displays a String on the console.

`putStrLn` displays a String on the console with a trailing newline.

IO IN HASKELL AND IDRIS If you’re already familiar with Haskell, you’ll find that programming with `IO` in Idris is very similar to writing interactive programs in Haskell. If you understand listing 5.1, you can safely move on to section 5.3, where you’ll see how to validate user input and deal with errors in interactive Idris programs. You may also want to look at section 5.2.2, which discusses pattern-matching bindings.

In this section, you’ll learn how to write interactive programs like this using `IO` in Idris and explore the distinction between evaluating expressions and executing programs, which allows us to write interactive programs without compromising on purity.

5.1.1 Evaluating and executing interactive programs

Functions that return an `IO` type are still considered pure, because they merely *describe* interactive actions. For example, the `putStrLn` function is defined in the Prelude and returns the actions that output a given `String`, plus a newline character, to the console:

```
Idris> :t putStrLn
putStrLn : String -> IO ()
```

When you enter an expression at the REPL, Idris evaluates that expression. If that expression is a description of interactive actions, then to perform those actions requires an additional step, *execution*.

Figure 5.1 illustrates what happens when the expression `putStrLn (show (47 * 2))` is executed. First, Idris calculates that the interactive action is to display "94\n" on the console (that is, the evaluator calculates the exact string that is to be displayed), and then it passes that action to the runtime environment.

You can see what happens by evaluating the expression `putStrLn (show (47 * 2))` at the REPL. This merely shows a description of the actions that a runtime environment can execute. There’s no need to look closely at the exact form of the result here, but you can at least see that the evaluation produces an expression with type `IO ()`:

```
Idris> putStrLn (show (47 * 2))
io_bind (prim_write "94\n") (\__bindx => io_return ()) : IO ()
```




Figure 5.1 Evaluating the expression `putStrLn (show (47 * 2))` and then executing the resulting actions

If you want to *execute* the resulting actions, you'll need to pass this description to the Idris runtime environment. You achieve this with the `:exec` command at the REPL:

```
Idris> :exec putStrLn (show (47 * 2))
94
```

In general, the `:exec` command, given an expression of type `IO ()`, can be understood to do the following:

- 1 Evaluate the expression, producing a description of the interactive actions to execute. In figure 5.1, *evaluating* the expression produces a description of an action: *Write the string "94\n" to the console*.
- 2 Pass the resulting actions to the runtime environment, which executes them. In figure 5.1, *executing* the actions writes the string `"94\n"` to the console.

It's also possible to compile to a standalone executable using the `:c` command at the REPL, which takes the executable name as an argument and generates a program that executes the actions described in `main`. For example, let's go back to listing 5.1, repeated here:

```
module Main

main : IO ()
main = do
    putStrLn "Enter your name: "
    x <- getLine
    putStrLn ("Hello " ++ x ++ "!")
```

If you save this in a file, `Hello.idr`, and load it at the REPL, you can produce an executable as follows:

```
*Hello> :c hello
```

This results in an executable file called `hello` (or, on Windows systems, `hello.exe`) that can be run directly from the shell. For example:

```
$ ./hello
Enter your name: Edwin
Hello Edwin!
```

Here, `main` describes not just one action but a sequence of actions. In general, in interactive programs you need to be able to sequence actions and have commands react to user input. Therefore, Idris provides facilities for combining smaller interactive programs into larger ones. We'll begin by seeing how to do this using a Prelude function, `(>>=)`, and then we'll look at the higher-level syntax we used earlier in listing 5.1, the `do` notation.

5.1.2 Actions and sequencing: the `>>=` operator

In the brief example in listing 5.1, we used three `IO` actions:

- `putStr : String -> IO ()`, which is an action to display a `String` on the console
- `putStrLn : String -> IO ()`, which is an action to display a `String` followed by a newline on the console
- `getLine : IO String`, which is an action to read a `String` from the console

In order to write realistic programs, you'll need to do more than just execute actions. You'll need to be able to sequence actions, and you'll need to be able to process the results of those actions.

Let's say, for example, you want to read a string from the console and then output its length. There's a `length` function in the Prelude, of type `String -> Int`, that you can use to calculate the length of a string, but `getLine` returns a value of type `IO String`, not `String`. The type `IO String` is a description of an action that produces a `String`, not the `String` itself.

In other words, you can get access to the string read from the console when the actions resulting from the function are *executed*, but you don't yet have access to it when the function is *evaluated*. There's no function of type `IO String -> String`. If such a function existed, it would mean that it was possible to know which string was read from the console without actually reading any string from the console!

The Prelude provides a function called `>>=` (intended to be used as an infix operator) that allows the sequencing of `IO` actions, feeding the result of one action as input into the next. It has the following type:

```
(>>=) : IO a -> (a -> IO b) -> IO b
```

Figure 5.2 shows an example application of `>>=` to sequence two actions, feeding the output of the first, `getLine`, as the input into the second, `putStrLn`. Because `getLine` returns a value of type `IO String`, Idris expects that `putStrLn` takes a `String` as its argument.

You can execute any sequence of actions with type `IO ()` at the REPL using the `:exec` command, so you can execute the operations in figure 5.2 as follows:




Figure 5.2 An interactive operation to read a `String` and then echo its contents using the `>>=` operator. The types of each subexpression are indicated.

- 7 *Refine*—Finally, complete the definition using `show` to convert `len` to a `String`, and `putStrLn` to give an action that displays the result:

```
printLength : IO ()
printLength = getLine >>= \input => let len = length input in
                                         putStrLn (show len)
```

You can try this definition at the REPL with `:exec printLength`, telling Idris that you'd like to execute the resulting actions rather than merely evaluate the expression:

```
*PrintLength> :exec printLength
test input           ← Entered by the user
10                  ← Output by the Idris runtime environment
```

Listing 5.2 shows the complete definition, further refined by displaying a prompt. The layout has also been adjusted in this definition to highlight the sequence of actions.

Listing 5.2 A function to display a prompt, read a string, and then display its length, using `>>=` to sequence IO actions (`PrintLength.idr`)

```
printLength : IO ()
printLength = putStr "Input string: " >>= \_ =>
    getLine >>= \input =>
        let len = length input in
            putStrLn (show len)

    getLine returns an IO String, so in
    the rest of the definition, input
    will have type String.
```

putStr returns `IO ()`, so the second argument of `>>=` expects a value of type `()` as its input. The underscore indicates that this value is ignored.

The type of `>>=`

If you check the type of `>>=` at the REPL, you'll see a constrained generic type:

```
Idris> :t (>>=)
(>>=) : Monad m => m a -> (a -> m b) -> m b
```

In practice, this means that the pattern applies more generally than for `IO`, and you'll see more of this later. For now, just read the type variable `m` here as `IO`.

In principle, you can always use `>>=` to sequence `IO` actions, feeding the output of one action as the input to the next. However, the resulting definitions are somewhat ugly, and sequencing actions is particularly common. That's why Idris provides an alternative syntax for sequencing `IO` actions, the `do` notation.

5.1.3 Syntactic sugar for sequencing with do notation

Interactive programs are, by their nature, typically imperative in style, with a sequence of commands, each of which produces a value that can be used later. The `>>=` function captures this idea, but the resulting definitions can be difficult to read.

Instead, you can sequence IO actions inside do blocks, which allow you list the actions that will be run when a block is executed. For example, to print two things in succession you'd do something like this:

```
printTwoThings : IO ()
printTwoThings = do putStrLn "Hello"
                  putStrLn "World"
```

Idris translates the do notation into applications of `>=`. Figure 5.3 shows how this works in the simplest case, where an action is to be executed, followed by more actions.




Figure 5.3 Transforming do notation to an expression using the `>=` operator when sequencing actions. The value produced by the action, of type `ty`, is ignored, as indicated by the underscore.

The result of an action can be assigned to a variable. For example, to assign the result of reading from the console with `getLine` to a variable `x` and then print the result, you can write this:

```
printInput : IO ()
printInput = do x <- getLine
               putStrLn x
```

The notation `x <- getLine` states that the result of the `getLine` action (that is, the String produced by the action of type `IO String`) will be stored in the variable `x`, which you can use in the rest of the do block. Sequencing actions and binding a variable like this with do notation directly translates to applications of `>=`, as illustrated in figure 5.4.




Figure 5.4 Transforming do notation to an expression using the `>=` operator, when binding a variable and sequencing actions

You can also use `let` inside `do` blocks to assign a pure expression to a variable. The following listing shows how `printLength` could be written using `do` notation instead of using `>>=` directly.

Listing 5.3 A function to display a prompt, read a string, and then display its length using do notation to sequence IO actions (PrintLength.idr)

```
printLength : IO ()
printLength = do putStrLn "Input string: "
                input <- getLine
                let len = length input
                putStrLn (show len)
```

Annotations for Listing 5.3:

- A callout points to the line `input <- getLine` with the text: "getLine has type IO String, so input has type String."
- A callout points to the line `let len = length input` with the text: "'length input' has type Nat, so len has type Nat; it was assigned using let. Note that in do blocks, there's no 'in' keyword after the assignment."

let and <- in do blocks

The `printLength` function in listing 5.3 uses two different forms for assigning to variables: using `let` and using `<-`. There's an important difference between these two, which again relies on the distinction between evaluation and execution:

- Use `let x = expression` to assign the result of the *evaluation* of an expression to a variable.
- Use `x <- action` to assign the result of the *execution* of an action to a variable.

In listing 5.3, `getLine` describes an action, so it needs to be executed and the result assigned using `<-`, but `length input` doesn't, so it's assigned using `let`.

Most interactive definitions in Idris are written using `do` notation to control sequencing. In larger programs, you'll also need to respond to user input and direct program flow. In the next section, we'll look at various methods of reading and responding to user input, and at how to implement loops and other forms of control flow in interactive programs.

Exercises



- 1 Using `do` notation, write a `printLonger` program that reads two strings and then displays the length of the longer string.
- 2 Write the same program using `>>=` instead of `do` notation.

You can test your answers at the REPL as follows:

```
*ex_5_1> :exec printLonger
First string: short
Second string: longer
6
```

5.2 Interactive programs and control flow

You've seen how to write basic interactive programs by sequencing existing actions, such as `getLine` to read input from the console and `putStrLn` to display text at the console. But as programs get larger, you'll need more control: you'll need to be able to validate and respond to user input, and you'll need to express loops and other forms of control.

In general, control flow in functions describing interactive actions works exactly the same way as control flow in pure functions, by pattern matching and recursion. Functions that describe interactive actions are pure functions themselves, after all, merely describing the actions that are to be executed later.

In this section, we'll look at some common patterns that are encountered in interactive programs: producing pure values by combining the results of interactive actions, pattern matching on the results of interactive actions, and, finally, putting everything together in an interactive program with loops.

5.2.1 Producing pure values in interactive definitions

As well as describing actions for execution by the runtime system, you'll often want to produce results from interactive programs. You've already seen `getLine`, for example:

```
getLine : IO String
```

The type `IO String` says that this is a function describing actions that produce a `String` as a result.

Often, you'll want to write a function that uses an `IO` action such as `getLine` and then manipulates its result further before returning it. For example, you might want to write a `readNumber` function that reads a `String` from the console and converts it to a `Nat` if the input consists entirely of digits. It produces a value of type `Maybe Nat`:

- If the input consists entirely of digits representing the number `i`, it produces `Just i`. For example, on reading the string "1234", it produces `Just 1234`.
- Otherwise, it produces `Nothing`.

Listing 5.4 shows how to define `readNumber`. After reading an input using `getLine`, it checks whether every character in the input is a digit. If so, it converts the input to a `Nat`, producing a result using `Just`; otherwise, it produces `Nothing`.

Listing 5.4 Reading and validating a number (ReadNum.idr)

```
readNumber : IO (Maybe Nat)
readNumber = do
    input <- getLine
    if all isDigit (unpack input)
        then pure (Just (cast input))
        else pure Nothing
```

getLine has type `IO String`,
so input has type `String`.

Uses `unpack` to convert the input to
`List Char`, to check that each input
character is a digit

The pure function constructs an `IO` action that produces a value, with no other input or output effects.

The pure function is used to produce a value in an interactive program without having any other input or output effects when it's executed. Its purpose is to allow pure values to be constructed by interactive programs, as shown by its type:

```
pure : a -> IO a
```

To understand the need for pure in `readNumber`, you can replace the `then` and `else` branches of the `if` with holes, and inspect their types:

```
readNumber : IO (Maybe Nat)
readNumber = do
    input <- getLine
    if all isDigit (unpack input)
        then ?numberOK
        else ?numberBad
```

If you look at the type of `?numberOK`, you'll see that you need a value of type `IO (Maybe Nat)`:

```
input : String
-----
numberOK : IO (Maybe Nat)
```

You can then refine the `?numberOK` and `?numberBad` holes using `pure`:

```
readNumber : IO (Maybe Nat)
readNumber = do
    input <- getLine
    if all isDigit (unpack input)
        then pure ?numberOK
        else pure ?numberBad
```

Now, if you look at the type of `?numberOK`, you'll see that you need a value of type `Maybe Nat` instead, without the `IO` wrapper:

```
input : String
-----
numberOK : Maybe Nat
```

Type of `pure`

As with `>>=`, if you check the type of `pure` at the REPL, you'll see a constrained generic type, rather than a type that uses `IO` specifically:

```
Idris> :t pure
pure : Applicative f => a -> f a
```

You can read the `f` here as `IO`. We'll get to the exact meaning of `Applicative` and `Monad` in chapter 7; you don't need to understand the details to write interactive programs.

You can try `readNumber` at the REPL by executing it and passing the result on to `println` (described in the sidebar, “Displaying values with `println`”):

```
*ReadNum> :exec readNumber >>= println
100           ←———— Entered by the user
Just 100      ←———— Output by Idris

*ReadNum> :exec readNumber >>= println
bad           ←———— Entered by the user
Nothing       ←———— Output by Idris
```

In the first case, the input is valid, so Idris produces the result `Just 100`. In the second case, the input has nondigit characters, so Idris produces `Nothing`.

Displaying values with `println`

`println` is a combination of `putStrLn` and `show`, and it’s convenient for displaying values at the console directly. It’s defined in the Prelude:

```
println : Show a => a -> IO ()
println x = putStrLn (show x)
```

5.2.2 Pattern-matching bindings

When an interactive action produces a value of some complex data type, such as `readNumber`, which produces a value of type `Maybe Nat`, you’ll often want to pattern-match on the intermediate result. You can do this using a `case` expression, but this can lead to deeply nested definitions. The following listing, for example, shows a function that reads two numbers from the console using `readNumber`, and produces a pair of those numbers if both are valid inputs, or `Nothing` otherwise.

Listing 5.5 Reading and validating a pair of numbers from the console (ReadNum.idr)

```
readNumbers : IO (Maybe (Nat, Nat))
readNumbers =
  do num1 <- readNumber
    case num1 of
      Nothing => pure Nothing
      Just num1_ok =>
        do num2 <- readNumber
          case num2 of
            Nothing => pure Nothing
            Just num2_ok => pure (Just (num1_ok, num2_ok))
```

The diagram illustrates the execution flow of the `readNumbers` function. It starts with the function definition and its body. A callout on the left indicates that the first input, `num1`, has type `Maybe Nat`. The `case num1 of` block branches into two paths: one for `Nothing` (producing `pure Nothing`) and one for `Just num1_ok`. The `Just num1_ok` path leads to another `do` block that reads `num2`. This second `do` block also branches on `case num2 of`: if `Nothing`, it produces `pure Nothing`; if `Just num2_ok`, it produces a tuple `pure (Just (num1_ok, num2_ok))`. Callouts on the right explain that `num1` is invalid (producing `Nothing`) and that `num2` is invalid (producing `Nothing`). A large callout at the bottom left states that both inputs are valid, producing a pair of the inputs read.

This will only get worse as functions get longer and there are more error conditions to deal with. To help with this, Idris provides some concise syntax for matching on intermediate values in `do` notation. First, we’ll take a look at a simple example, reading a pair of `Strings` from the console, and then we’ll revisit `readNumbers` and see how it can be made more concise.

You can write a function that reads two Strings and produces a pair as follows, using `pure` to combine the two inputs:

```
readPair : IO (String, String)
readPair = do str1 <- getLine
             str2 <- getLine
             pure (str1, str2)
```

When you use the result produced by `readPair`, you'll need to pattern-match on it to extract the first and second inputs. For example, to read a pair of Strings using `readPair` and then display both, you'd do this:

```
usePair : IO ()
usePair = do pair <- readPair
            case pair of
                (str1, str2) => putStrLn ("You entered " ++
                                         str1 ++ " and " ++ str2)
```

To make these programs more concise, Idris allows the pattern match and the assignment to be combined on one line, in a pattern-matching binding. The following program has exactly the same behavior as the preceding one:

```
usePair : IO ()
usePair = do (str1, str2) <- readPair
            putStrLn ("You entered " ++ str1 ++ " and " ++ str2)
```

A similar idea works for `readNumbers`. You can pattern-match directly on the result of `readNumber` to check the validity of its result:

```
readNumbers : IO (Maybe (Nat, Nat))
readNumbers =
    do Just num1_ok <- readNumber
       Just num2_ok <- readNumber
       pure (Just (num1_ok, num2_ok))
```

This works as you'd like when the user enters valid numbers:

```
*ReadNum> :exec readNumbers >>= println
10          ← Entered by the user
20          ← Entered by the user
Just (10, 20) ← Output by Idris
```

But, unfortunately, it doesn't deal with the case where `readNumber` produces `Nothing`, and therefore crashes on execution:

```
*ReadNum> :exec readNumbers >>= println
bad          ← Entered by the user
*** ReadNum.idr:26:22:unmatched case in Main.case block in readNumbers
at ReadNum.idr:26:22 ***
```

Idris has noticed this when checking `readNumbers` for totality:

```
*ReadNum> :total readNumbers_v2
Main.readNumbers is possibly not total due to:
  Main.case block in readNumbers at ReadNum.idr:26:22, which is
  not total as there are missing cases
```

REMEMBER TO CHECK TOTALITY! This incomplete definition of `readNumbers` illustrates why it's important to check that functions are total. Even though `readNumbers` type-checks successfully, it could still fail at runtime due to the missing cases.

Listing 5.6 shows how you can deal with the other possibilities in a pattern-matching binding. As well as the binding itself, you provide alternative matches after a vertical bar, showing how the rest of the function should proceed if the default match on `Just num1_ok` or `Just num2_ok` fails. This function has exactly the same behavior as the earlier function in listing 5.5.

Listing 5.6 Reading and validating a pair of numbers from the console, concisely (ReadNum.idr)

```
readNumbers : IO (Maybe (Nat, Nat))
readNumbers =
  do Just num1_ok <- readNumber | Nothing => pure Nothing
    Just num2_ok <- readNumber | Nothing => pure Nothing
    pure (Just (num1_ok, num2_ok))
```

Reads first input. If it's invalid, produces Nothing as a result of the computation.

Reads second input. If it's invalid, produces Nothing as a result of the computation.

Both inputs are valid, so produces a pair of the inputs read

Pattern-matching bindings of this form allow you to express the expected valid behavior of a function in the default matches (`Just num1_ok` and `Just num2_ok` in listing 5.6), dealing with the error cases in the alternative matches.

5.2.3 Writing interactive definitions with loops

Now that you can read, validate, and respond to user input, let's put everything together and write interactive definitions with loops.

You can write loops by writing recursive definitions. The next listing, for example, shows a `countdown` function that calculates a sequence of actions to display a countdown, with a one-second pause between displaying each number.

Listing 5.7 Display a countdown, pausing one second between each iteration (Loops.idr)

```
module Main
import System
```

The System module contains several definitions for interacting with the operating system and environment. You import it here for `usleep`.

```
countdown : (secs : Nat) -> IO ()
countdown Z = putStrLn "Lift off!"
countdown (S secs) = do putStrLn (show (S secs))
                       usleep 1000000
                       countdown secs
```

countdown is implemented by pattern matching on `Nat`, so there are cases for each constructor of `Nat`.

`usleep` describes an action that sleeps for the given number of microseconds.

Continues the countdown

If you try executing this at the REPL using :exec, you'll see a countdown displayed:

```
*Loops> :exec countdown 5
5
4
3
2
1
Lift off!
```

You can also check whether this function is total, that is, guaranteed to produce a result in finite time for all possible inputs:

```
*Loops> :total countdown
Main.countdown is Total
```

In general, you can write an interactive function that describes a loop by making a recursive call to the function with its last action. In `countdown`, as long as the input argument isn't `z`, the program when executed will print the argument and wait a second before making a recursive call on the next smaller number. The number of iterations of the loop is therefore determined by the initial input. Because this is finite, `countdown` must terminate eventually, so Idris reports that it is total.

TOTALITY AND INTERACTIVE PROGRAMS Totality checking is based on evaluation, not execution. The result of totality checking an IO program, therefore, tells you whether Idris will produce a finite sequence of actions, but nothing about the runtime behavior of those actions.

Sometimes, however, the number of iterations is determined by user input. For example, you can write a function to keep executing `countdown` until the user wants to stop, as shown next.

Listing 5.8 Keep running `countdown` until the user doesn't want to run it any more (`Loops.idr`)

```
countdowns : IO ()
countdowns = do putStrLn "Enter starting number: "
              Just startNum <- readNumber
              | Nothing => do putStrLn "Invalid input"
                           countdowns
              countdown startNum
              putStrLn "Another (y/n)? "
              yn <- getLine
              if yn == "y" then countdowns
              else pure ()
```

The diagram shows annotations for the code:

- A callout points to the first line with the text "Calls `countdown` with the user's input".
- A callout points to the line "Just startNum <- readNumber" with the text "Uses `readNumber` to get a valid input, or restart if the input is not valid".
- A callout points to the line "if yn == "y" then countdowns" with the text "Makes a recursive call if the user enters "y"".

This function is *not* total, because there's no guarantee that a user will ever enter anything other than `y`, or even provide any valid input.

```
*Loops> :total countdowns
Main.countdowns is possibly not total due to recursive path:
Main.countdowns
```

Interactive programs that might loop forever, such as countdowns (or, more realistically, servers or operating systems) are not total, at least if we limit the definition to *terminating* programs. More precisely, a total function must either terminate or be guaranteed to produce a finite prefix of some infinite input, within finite time. We'll discuss this further in chapter 11.

Exercises



- 1 Write a function that implements a simple “guess the number” game. It should have the following type:

```
guess : (target : Nat) -> IO ()
```

Here, `target` is the number to be guessed, and `guess` should behave as follows:

- Repeatedly ask the user to guess a number, and display whether the guess is too high, too low, or correct.
- When the guess is correct, exit.

Ideally, `guess` will also report an error message if the input is invalid (for example, if it contains characters that are not digits or are out of range).

- 2 Implement a `main` function that chooses a random number between 1 and 100 and then calls `guess`.

Hint: As a source of random numbers, you could use `time : IO Integer`, defined in the `System` module.

- 3 Extend `guess` so that it counts the number of guesses the user has taken and displays that number before the input is read.

Hint: Refine the type of `guess` to the following:

```
guess : (target : Nat) -> (guesses : Nat) -> IO ()
```

- 4 Implement your own versions of `repl` and `replWith`. Remember that you'll need to use different names to avoid clashing with the names defined in the Prelude.

5.3 *Reading and validating dependent types*

In previous chapters, you've seen several functions with dependent types, particularly using `Vect` to express lengths of vectors in their types. This allows you to state assumptions about the form of inputs to a function in its type and guarantees about the form of its output. For example, you've seen `zip`, which pairs corresponding elements of vectors:

```
zip : Vect n a -> Vect n b -> Vect n (a, b)
```

The type expresses the following:

- *Assumption*—Both input vectors have the same length, `n`.
- *Guarantee*—The output vector will have the same length as the input vectors.

- **Guarantee**—The output vector will consist of pairs of the element type of the first input vector and the element type of the second input vector.

Idris checks that whenever the function is called, the arguments satisfy the assumption and the definition of the function satisfies the guarantee.

So far, we've been testing functions such as `zip` at the REPL. Realistically, though, inputs to functions don't come from such a carefully controlled environment where the Idris type checker is available. In practice, when a complete program is compiled and executed, inputs to functions will originate from some external source: perhaps a field on a web page, or user input at the console.

When a program reads data from some external source, it can't make any assumptions about the form of that data. Rather, the program has to check that the data is of the necessary form. The type of a function tells you exactly what you need to check in order to evaluate it safely.

In this section, we'll write a program that reads two vectors from the console, uses `zip` to pair corresponding elements if the vectors are the same length, and then displays the result. Although this is a simple goal, it demonstrates several important aspects of working with dependent types in interactive programs, and you'll see many more examples of this form later in this book. You'll see, briefly, how the types of the pure parts of programs tell you what you need to check in the interactive parts, and how the type system guides you toward the parts where error checking is necessary.

5.3.1 Reading a Vect from the console


As a first step, you'll need to be able to read a vector from the console. Because vectors express their length in the type, you'll need some way of describing the length of the vector you intend to read. One way is to take the length as an input, such as in the following type:

```
readVectLen : (len : Nat) -> IO (Vect len String)
```

This type states that `readVectLen` takes an intended length as input, and returns the sequence of actions that reads a vector of `String`s of that length. The following listing shows one way you could implement this function.

Listing 5.9 Reading a Vect of known length from the console (ReadVect.idr)

```
readVectLen : (len : Nat) -> IO (Vect len String)
readVectLen Z = pure []
readVectLen (S k) = do x <- getLine
                      xs <- readVectLen k
                      pure (x :: xs)
```



You can try `readVectLen` at the REPL by executing it with a specific length, and then printing the result with `printLn`:

```
*ReadVect> :exec readVectLen 4 >>= printLn
John
Paul
George
Ringo
["John", "Paul", "George", "Ringo"]
```

← **Entered by the user**

← **Output by Idris**

A problem with this approach is that you need to know in advance how long the vector should be, because the length is given as an input. What if, instead, you want to read strings until the user enters a blank line? You can't know in advance how many strings the user will enter, so instead, you'll need to return the length along with a vector of that length.

5.3.2 **Reading a Vect of unknown length**

If you read a vector from the console, terminated by a blank line, you can't know how many elements will be in the resulting vector. In this situation, you can define a new data type that wraps not only the vector but also its length:

```
data VectUnknown : Type -> Type where
  MkVect : (len : Nat) -> Vect len a -> VectUnknown a
```

In type-driven development, we aim to express what we know about data in its type; if we can't know something about data, we need to express this somehow too. This is the purpose of `VectUnknown`; it contains both the vector and its length, meaning that the length doesn't have to be known in the type.

You can construct an example at the REPL:

```
*ReadVect> MkVect 4 ["John", "Paul", "George", "Ringo"]
MkVect 4 ["John", "Paul", "George", "Ringo"] : VectUnknown String
```

In fact, you could leave an underscore in the expression instead of giving the length, 4, explicitly, because Idris can infer this from the length of the given vector:

```
*ReadVect> MkVect _ ["John", "Paul", "George", "Ringo"]
MkVect 4 ["John", "Paul", "George", "Ringo"] : VectUnknown String
```

Having defined `VectUnknown`, instead of writing a function that returns `IO (Vect len String)`, you can write a function that returns `IO (VectUnknown String)`. That is, it returns not only the vector, but also its length:

```
readVect : IO (VectUnknown String)
```

This type states that `readVect` is a sequence of interactive actions that produce a vector of some unknown length, which will be determined at runtime. The following listing shows one possible implementation.

Listing 5.10 Reading a Vect of unknown length from the console (ReadVect.idr)

```

readVect : IO (VectUnknown String)
readVect = do x <- getLine           ← Reads the first line
              if (x == "") 
                then pure (MkVect _ [])
                else do MkVect _ xs <- readVect
                          pure (MkVect _ (x :: xs)) ← Combines xs with
                                                the first line (x) and
                                                wraps it in MkVect

```

If the line read is blank,
returns an empty vector
wrapped in MkVect. Idris
will infer its length.

Reads the rest of the vector
and pattern-matches on the
result to extract xs

To try this, you can define a convenience function, printVect, that displays the contents and length of a VectUnknown at the console:

```

printVect : Show a => VectUnknown a -> IO ()
printVect (MkVect len xs)
  = putStrLn (show xs ++ " (length " ++ show len ++ ")")

```

Then, you can try reading some input at the REPL:

```

*ReadVect> :exec readVect >>= printVect
John
Paul
George
Ringo

```

Entered by the user

```

["John", "Paul", "George", "Ringo"] (length 4)

```

Blank line entered by user

Output by Idris

When dealing with user input, there'll often be some properties of the data that you can't know until runtime. The length of a vector is one example: once you've read the vector, you know its length, and from there you can check it and reason about how it relates to other data. But you could have a similar problem with any dependent data type that's read from user input, and it would be better not to define a new type (like VectUnknown) every time this happens. Instead, Idris provides a more generic solution, *dependent pairs*.

5.3.3 Dependent pairs

You've already seen tuples, introduced in chapter 2, which allow you to combine values of different types, as in this example:

```

mypair : (Int, String)
mypair = (94, "Pages")

```

A *dependent pair* is a more expressive form of this construct, where the type of the second element in a pair can be computed from the value of the first element. For example:

```

anyVect : (n : Nat ** Vect n String)
anyVect = (3 ** ["Rod", "Jane", "Freddy"])

```

Dependent pairs are written with the elements separated by **. Their types are written using the same syntax as their values, except that the first element is given an explicit name (n in the preceding example). Figure 5.5 illustrates the syntax for dependent pair types.




Figure 5.5 Dependent pair syntax. Notice that the first element is given a name that can be used in the type of the second element.

You can also often omit the type of the first element, if Idris can infer it from the type of the second element. For example:

```
anyVect : (n ** Vect n String)
anyVect = (3 ** ["Rod", "Jane", "Freddy"])
```

If you replace the value `["Rod", "Jane", "Freddy"]` with a hole, you can see how the first value, 3, affects its type:

```
anyVect : (n ** Vect n String)
anyVect = (3 ** ?anyVect_rhs)
```

Inspecting the type of `?anyVect_rhs` reveals that the second element must specifically be a vector of length 3, as specified by the first element:

```
-----
anyVect_rhs : Vect 3 String
```

TYPES OF SUBEXPRESSIONS Remember that when trying to understand larger program listings, you can replace subexpressions with a hole, like `?anyVect_rhs` in the type-inference example, to find out their expected types and the types of any local variables that are in scope.

Instead of defining `VectUnknown` as in the section 5.3.2, you can define a function that reads vectors of unknown length by returning a dependent pair of the length, and a vector of that length. The following listing shows how `readVect` could be defined using dependent pairs.

Listing 5.11 Reading a Vect of unknown length from the console, returning a dependent pair (DepPairs.idr)

```
readVect : IO (len ** Vect len String)
readVect = do x <- getLine
             if (x == "")
```

Returns actions that build a dependent pair of some length, and a vector of that length

```

  Returns a dependent pair of zero (inferred by Idris) and a Vect of length zero
  ┌─────────────────────────────────────────────────────────────────────────────────┐
  then pure (_ ** [])
  else do (_ ** xs) <- readVect
         pure (_ ** x :: xs)
  └─────────────────────────────────────────────────────────────────────────┘
  Returns a dependent pair of a length inferred by Idris, and a Vect combining x and xs

```

Again, you can try this at the REPL. You can use `printLn` to display the contents of the dependent pair:

```

*DepsPairs> :exec readVect >>= printLn
Rod
Jane
Freddy
          ← Entered by the user
          ← Blank line entered by user
(3 ** ["Rod", "Jane", "Freddy"])
          ← Output by Idris

```

Now that you have the ability to read vectors of arbitrary and user-defined lengths from the console, we can complete our original goal of writing a program that reads two vectors and zips them together if their lengths match.

5.3.4 Validating Vect lengths

Our goal, as stated at the beginning of this section, is to write a program that does the following:

- 1 Read two input vectors, using `readVect`.
- 2 If they have different lengths, display an error.
- 3 If they have the same lengths, display the result of zipping the vectors together.

The program will take its inputs from the user at the console and display its result on the console. We'll implement it as a `zipInputs` function, as follows:

- 1 *Type*—Because the input and output are entirely at the console, the type states that `zipInputs` takes no arguments and returns interactive actions:

```
zipInputs : IO ()
```

- 2 *Define*—The first step in defining the function is to read two inputs using `readVect`. Leave a hole, `?zipInputs_rhs`, for the rest of the definition:

```

zipInputs : IO ()
zipInputs = do putStrLn "Enter first vector (blank line to end):"
              (len1 ** vec1) <- readVect
              putStrLn "Enter second vector (blank line to end):"
              (len2 ** vec2) <- readVect
              ?zipInputs_rhs

```

- 3 *Type*—Looking at the type of `?zipInputs_rhs`, you can see the types (and hence the lengths) of the vectors that have been read:

```

len2 : Nat
vec2 : Vect len2 String
len1 : Nat
vec1 : Vect len1 String
-----
zipInputs_rhs : IO ()

```

`vec1` has length `len1`, and `vec2` has length `len2`; there's no explicit relationship between these lengths. Indeed, there shouldn't be, because they were read independently. But if you look at the type of `zip`, you'll see that the lengths must be the same before you can use it:

```
zip : Vect n a -> Vect n b -> Vect n (a, b)
```

- 4 *Refine*—Before you can make progress, you'll need to check that the length of the second vector is the same as the length of the first. As a first attempt, you might try the following:

```
if len1 == len2
  then ?zipInputs_rhs1
  else ?zipInputs_rhs2
```

Unfortunately, this doesn't help. If you look at the type of `?zipInputs_rhs1`, you'll see that nothing has changed:

```
len1 : Nat
len2 : Nat
vec2 : Vect len2 String
vec1 : Vect len1 String
-----
zipInputs_rhs1 : IO ()
```

The problem is that the *type* of `len1 == len2, Bool`, tells you nothing about the *meaning* of the operation itself. As far as Idris is concerned, the `==` operation could be implemented in any way (you'll see in chapter 7 how `==` can be defined) and doesn't necessarily guarantee that `len1` and `len2` really are equal.

Instead, you can use the following function defined in `Data.Vect`:

```
exactLength : (len : Nat) -> (input : Vect m a) -> Maybe (Vect len a)
```

This function takes a length, `len`, and a vector, `input`, of any length. If the length of the input vector turns out to be `len`, it returns `Just input`, with its type updated to `Vect len a`. Otherwise, it returns `Nothing`.

IMPLEMENTING EXACTLENGTH To implement `exactLength`, you need a more expressive type than `Bool` for representing the result of an equality test. You'll see how to do this in chapter 8, and we'll discuss the limitations of `Bool` in general.

Using `exactLength`, you can refine the definition as follows:

```
zipInputs : IO ()
zipInputs = do putStrLn "Enter first vector (blank line to end):"
              (len1 ** vec1) <- readVect
              putStrLn "Enter second vector (blank line to end):"
              (len2 ** vec2) <- readVect
              case exactLength len1 vec2 of
                Nothing => ?zipInputs_rhs_1
                Just vec2' => ?zipInputs_rhs_2
```

- 5 Refine**—For `zipInputs_rhs_1`, the inputs are different lengths, so you display an error:

```
case exactLength len1 vec2 of
    Nothing => putStrLn "Vectors are different lengths"
    Just vec2' => ?zipInputs_rhs_2
```

- 6 Refine**—For `zipInputs_rhs_2`, you have a new vector, `vec2'`, which is the same as `vec2` but with its length now guaranteed to be the same as the length of `vec1`, as you can confirm by looking at the type:

```
len1 : Nat
vec2' : Vect len1 String
len2 : Nat
vec2 : Vect len2 String
vec1 : Vect len1 String
-----
zipInputs_rhs_2 : IO ()
```

You can therefore complete the definition by calling `zip` with `vec1` and `vec2'`, and then printing the result:

```
case exactLength len1 vec2 of
    Nothing => putStrLn "Vectors are different lengths"
    Just vec2' => println (zip vec1 vec2')
```

For reference, the complete definition is given in the following listing.

Listing 5.12 Complete definition of `zipInputs` (`DepPairs.idr`)

```
zipInputs : IO ()
zipInputs = do putStrLn "Enter first vector (blank line to end) :"
    (len1 ** vec1) <- readVect
    putStrLn "Enter second vector (blank line to end) :"
    (len2 ** vec2) <- readVect
    case exactLength len1 vec2 of
        Nothing => putStrLn "Vectors are different lengths"
        Just vec2' => println (zip vec1 vec2')
```

Exercises



In these exercises, you'll find the following Prelude functions useful, in addition to the functions discussed earlier in the chapter: `openFile`, `closeFile`, `fEOF`, `fGetLine`, and `writeFile`. Use `:doc` to find out what each of these do. Also, see the sidebar, “Handling I/O errors.”

- 1** Write a function, `readToBlank : IO (List String)`, that reads input from the console until the user enters a blank line.
- 2** Write a function, `readAndSave : IO ()`, that reads input from the console until the user enters a blank line, and then reads a filename from the console and writes the input to that file.

- 3 Write a function, `readVectFile : (filename : String) -> IO (n ** Vect n String)`, that reads the contents of a file into a dependent pair containing a length and a Vect of that length. If there are any errors, it should return an empty vector.

Handling I/O errors

Many of the functions in the Exercises may return errors using `Either`. At first, you can assume the result is successful using a pattern-matching binding, as described in section 5.2.2:

```
do Right h <- openFile filename Read
    Right line <- fGetLine h
    {- rest of code -}
```

Then, to make the function total, handle errors using the notation described in the same section:

```
do Right h <- openFile filename Read
    | Left err => putStrLn (show err)
    Right line <- fGetLine h
    | Left err => putStrLn (show err)
    {- rest of code -}
```

5.4 Summary

- Idris provides a generic `IO` type for describing interactive actions.
- Idris distinguishes between evaluation of pure functions and execution of interactive actions. The `:exec` command at the REPL executes interactive actions.
- You can sequence interactive actions using `do` notation, which translates to applications of the `>>=` operator.
- You can produce pure values from interactive definitions by using the `pure` function.
- Idris provides a concise notation for pattern-matching on the result of an interactive action.
- Dependent types express assumptions about inputs to functions, so you need to validate user inputs to check that they satisfy those assumptions.
- Dependent pairs allow you to pair two values, where the type of the second value is computed from the first value.
- You can use dependent pairs to express that a type's argument, such as the length of a vector, will not be known until the user enters some input.

Programming with first-class types

This chapter covers

- Programming with type-level functions
- Writing functions with varying numbers of arguments
- Using type-level functions to calculate the structure of data
- Refining larger interactive programs

In Idris, as you've seen several times now, types can be manipulated just like any other language construct. For example, they can be stored in variables, passed to functions, or constructed by functions. Furthermore, because they're truly first-class, expressions can compute types, and types can also take any expression as an argument. You've seen several uses of this concept in practice, in particular the ability to store additional information about data in its type, such as the length of a vector.

In this chapter, we'll explore more ways of taking advantage of the first-class nature of types. You'll see how *type-level functions* can be used to give alternative names to types and also to calculate the type of a function from some other data. In

particular, you'll see how you can write a type-safe formatted output function, `printf`. For `printf`, the type (and even number) of arguments to the function are calculated from a *format string* provided as its first argument. This technique, calculating the type of some data (in this case, the arguments to `printf`) based on some other data (in this case, the format string) is often useful. Here are a couple of examples:

- Given an HTML form on a web page, you can calculate the type of a function to process inputs in the form.
- Given a database schema, you can calculate types for queries on that database.

As an example of this concept, you'll see how to use type-level functions to refine the data store we implemented at the end of chapter 4. Previously, you could only store data as `Strings`, but for more flexibility you might want the form of the data to be described by a user rather than hardcoded into the program. Using type-level functions, you'll extend the data store with a *schema* that describes the form of the data, and you'll use that schema to calculate appropriate types for functions to parse and display user data. In doing so, you'll learn more about using holes to help correct errors when refining larger programs.

To begin, we'll look at how to use functions at the type level to calculate types.

6.1 Type-level functions: calculating types

One of the most fundamental features of Idris is that types and expressions are part of the same language—you use the same syntax for both. You've already seen in chapter 4 that expressions can appear in types. For example, in the type of `append` on `Vect`, we have `n` and `m` (both `Nats`), and the resulting length is `n + m`:

```
append : Vect n elem -> Vect m elem -> Vect (n + m) elem
```

Here, `n`, `m`, and `n + m` all have type `Nat`, which can also be used in ordinary expressions. Similarly, you can use types in expressions and therefore write functions that *calculate* types. There are two common situations where you might want to do this:

- *To give more meaningful names to composite types*—For example, if you have a type `(Double, Double)` representing a position as a 2D coordinate, you might prefer to call the type `Position` to make the code more readable and self-documenting.
- *To allow a function's type to vary according to some contextual information*—For example, the type of a value returned by a database query will vary depending on the database schema and the query itself.

In the first case, you can define *type synonyms*, which give alternative names to types. In the second case, you can define *type-level functions* (of which type synonyms are a special case) that calculate types from some input. In this section, we'll take a look at both.

6.1.1 Type synonyms: giving informative names to complex types

Let's say you're writing an application that deals with complex polygons, such as a drawing application. You might represent a polygon as a vector of the coordinates of each corner. A triangle, for example, might be initialized as follows:

```
tri : Vect 3 (Double, Double)
tri = [(0.0, 0.0), (3.0, 0.0), (0.0, 4.0)]
```

The type, `Vect 3 (Double, Double)` says exactly what the form of the data will be, which is useful to the machine, but it doesn't give any indication to the reader what the purpose of the data is. Instead, you can refine the type using a type synonym for a position represented as a 2D coordinate:

```
Position : Type
Position = (Double, Double)
```

Here, you have a function called `Position` that takes no arguments and returns a `Type`. This is an ordinary function; there's nothing special about the way it's declared or implemented. Now, anywhere you can use a `Type`, you can use `Position` to calculate that type. For example, the triangle can be defined with a refined type as follows:

```
tri : Vect 3 Position
tri = [(0.0, 0.0), (3.0, 0.0), (0.0, 4.0)]
```

This kind of function is a type synonym because it provides an alternative name for some other type.

NAMING CONVENTION By convention, we usually use an initial capital letter for functions that are intended to compute types.

Because they're ordinary functions, they can also take arguments. For example, the following listing shows how you can use type synonyms to express more clearly that the `Vect` is intended to represent a polygon.

Listing 6.1 Defining a polygon using type synonyms (TypeSynonym.idr)

```
import Data.Vect

Position : Type
Position = (Double, Double)           ← A type synonym for describing
                                         positions as (x, y) coordinates

Polygon : Nat -> Type
Polygon n = Vect n Position          ← A type synonym for describing
                                         polygons with n corners

tri : Polygon 3
tri = [(0.0, 0.0), (3.0, 0.0), (0.0, 4.0)]
```

Because `Polygon` is an ordinary function, you can evaluate it at the REPL:

```
*TypeSynonyms> Polygon 3
Vect 3 (Double, Double) : Type
```

Also, notice that if you evaluate `tri` at the REPL, Idris will display `tri`'s type in the evaluated form. In other words, evaluation at the REPL evaluates both the expression and the type:

```
*TypeSynonyms> tri
[(0.0, 0.0), (3.0, 0.0), (0.0, 4.0)] : Vect 3 (Double, Double)
```

Using `:t`, on the other hand, displays `tri`'s type:

```
*TypeSynonyms> :t tri
Polygon 3
```

Finally, you can see what happens when you try to define `tri` interactively using expression search in Atom. Enter the following into an Atom buffer, along with the previous definition of `Polygon`:

```
tri : Polygon 3
tri = ?tri_rhs
```

The interactive editing tools in general, and expression search in particular, are aware of how type synonyms are defined, so if you try an expression search on `tri_rhs`, you'll get the same result as if the type were written directly as `Vect 3 (Double, Double)`:

```
tri : Polygon 3
tri = [(?tri_rhs1, ?tri_rhs2), (?tri_rhs3, ?tri_rhs4),
       (?tri_rhs5, ?tri_rhs6)]
```

The type synonyms defined in this section, `Position` and `Polygon`, are really just ordinary functions that happen to be used to compute types. This gives you a lot of flexibility in how you describe types, as you'll see in the rest of this chapter.

6.1.2 Type-level functions with pattern matching

Type synonyms are a special case of *type-level functions*, which are functions that can be used anywhere Idris is expecting a `Type`. There isn't anything special about type-level functions as far as Idris is concerned; they're ordinary functions that happen to return a `Type`, and they can use all of the language constructs available elsewhere. Nevertheless, it's useful to consider them separately, to see how they work in practice.

Because type-level functions are ordinary functions that return a type, you can write them by case splitting. For example, the following listing shows a function that calculates a type from a Boolean input. You saw this function in chapter 1, although in a slightly different form.

Listing 6.2 A function that calculates a type from a `Bool` (`TypeFuns.idr`)

```
StringOrInt : Bool -> Type
StringOrInt False = String
StringOrInt True = Int
```

Using this, you can write a function where the return type is *calculated from* or *depends on* an input. Using `StringOrInt`, you can write functions that return either type, depending on a Boolean flag.

As a small example, you can write a function that takes a Boolean input and returns the string "Ninety four" if it's False, or the integer 94 if it's True. Begin with the type:

- 1 Type**—Begin by giving a type declaration, using `StringOrInt`:

```
getStringOrInt : (isInt : Bool) -> StringOrInt isInt
getStringOrInt isInt = ?getStringOrInt_rhs
```

If you look at the type of `getStringOrInt_rhs` now, you'll see this:

```
isInt : Bool
-----
getStringOrInt_rhs : StringOrInt isInt
```

- 2 Define**—Because `isInt` appears in the required type for `getStringOrInt_rhs`, case splitting on `isInt` will cause the expected return type to change according to the specific value of `isInt` for each case. Case splitting on `isInt` leads to this:

```
getStringOrInt : (isInt : Bool) -> StringOrInt isInt
getStringOrInt False = ?getStringOrInt_rhs_1
getStringOrInt True = ?getStringOrInt_rhs_2
```

- 3 Type**—Looking at the types of the newly created holes, you can see how the expected type is changed in each case:

```
-----
getStringOrInt_rhs_1 : String
-----
getStringOrInt_rhs_2 : Int
```

In `getStringOrInt_rhs_1`, the type is refined to `StringOrInt False` because the pattern for `isInt` is `False`, which evaluates to `String`. Then, in `getStringOrInt_rhs_2`, the type is refined to `StringOrInt True`, which evaluates to `Int`.

- 4 Refine**—To complete the definition, you need to provide values of different types in each case:

```
getStringOrInt : (isInt : Bool) -> StringOrInt isInt
getStringOrInt False = "Ninety four"
getStringOrInt True = 94
```

DEPENDENT PATTERN MATCHING The `getStringOrInt` example illustrates a technique that's often useful when programming with dependent types: *dependent pattern matching*. This refers to a situation where the type of one argument to a function can be determined by inspecting the value of (that is, by case splitting) another. You've already seen an example of this when defining `zip` in chapter 4, where the form of one vector restricted the valid forms of the other, and you'll see a lot more.

Type-level functions can be used anywhere a Type is expected, meaning that they can be used in place of argument types too. For example, you can write a function that converts either a String or an Int to a canonical String representation, according to a Boolean flag. This function will behave as follows:

- If the input is a String, it returns the String with leading and trailing space removed using `trim`.
- If the input is an Int, it converts the input to a String using `cast`.

DOCUMENTATION Remember that you can use `:t` and `:doc` to check the types of functions that are unfamiliar (such as `trim`) at the REPL.

You can define this function with the following steps:

- 1 *Type*—Begin by writing a type for `valToString`, again using `StringOrInt`, but this time to calculate the input type:

```
valToString : (isInt : Bool) -> StringOrInt
valToString isInt y = ?valToString_rhs
```

Inspecting the type of `valToString_rhs`, you'll see the following:

```
isInt : Bool
y : StringOrInt
-----
valToString_rhs : String
```

- 2 *Define*—You can define this by case splitting on `isInt`. The type of `y` is calculated from `isInt`, so if you case-split on `isInt`, you should see refined types for `y` in each resulting case:

```
valToString : (isInt : Bool) -> StringOrInt
valToString False y = ?valToString_rhs_1
valToString True y = ?valToString_rhs_2
```

- 3 *Type*—Inspecting the types of `valToString_rhs_1` and `valToString_rhs_2` shows how the type of `y` is refined in each case:

```
y : String
-----
valToString_rhs_1 : String

y : Int
-----
valToString_rhs_2 : String
```

- 4 *Refine*—To complete the definition, fill in the right side to convert `y` to a trimmed String if it was a String, or a string representation of the number if it was an Int:

```
valToString : (isInt : Bool) -> StringOrInt
valToString False y = trim y
valToString True y = cast y
```

The simple examples in this section, `getStringOrInt` and `valToString`, illustrate a technique that can be used more widely in practice. You'll see some more practical examples later in this chapter: using a format string to calculate a type for formatted output, and a larger example allowing users to calculate a schema, extending the data store you saw in chapter 4.

There's more that you can achieve with type-level expressions, however. The fact that types are first-class means not only that types can be computed like any other value, but also that any expression form can appear in types.

6.1.3 Using case expressions in types

Any expression that can be used in a function can also be used at the type level, and vice versa. For example, you can leave holes in types while your understanding of a function's requirements develops, or you can use more-complex expression forms such as `case`. Let's briefly take a look at how this can work by using a `case` expression in the type of `valToString` instead of a separate `StringOrInt` function:

- Type—Start by giving a type for `valToString`, but because you don't immediately know the input type, you can leave a hole:

```
valToString : (isInt : Bool) -> ?argType -> String
```

- 2 Define**—Even though you have an incomplete type, you can still proceed with the definition by case splitting on `isInt`, because you at least know that `isInt` is a `Bool`:

```
valToString : (isInt : Bool) -> ?argType -> String
valToString False y = ?valToString_rhs_1
valToString True y = ?valToString rhs 2
```

- 3 Type—Inspecting the types of `valToString_rhs_1` and `valToString_rhs_2` will, however, reveal that you don't know the type of `y` yet:

```
y : ?argType  
-----  
valToString rhs 1 : String
```

- 4 Refine—At this stage, you need to refine the type with more details of ?argType. You can fill in ?argType with a case expression:

```
valToString : (isInt : Bool) -> (case _ of
                                         case_val => ?argType) -> String
valToString False y = ?valToString_rhs_1
valToString True v = ?valToString_rhs_2
```

- 5 *Refine*—Remember that when Idris generates a case expression, it leaves an underscore in place of the value the case expression will branch on. You'll need to fill this in before the program will compile. You can compute the argument type from the input `isInt`:

```
valToString False y = ?valToString_rhs_1
valToString True y = ?valToString_rhs_2
```

- 6 Refine**—Case splitting on `case_val` gives you the two possible values `isInt` can take:

```
valToString : (isInt : Bool) -> (case isInt of
                                      False => ?argType_1
                                      True => ?argType_2) -> String
valToString False y = ?valToString_rhs_1
valToString True y = ?valToString_rhs_2
```

You can then complete the type in the same way as your implementation of `StringOrInt`, refining `?argType_1` with `String` and `?argType_2` with `Int`:

```
valToString : (isInt : Bool) -> (case isInt of
                                      False => String
                                      True => Int) -> String
valToString False y = ?valToString_rhs_1
valToString True y = ?valToString_rhs_2
```

- 7 Type**—Inspecting `?valToString_rhs_1` and `?valToString_rhs2` now will show you the new types for the input `y`, calculated from the first argument:

```
y : String
-----
valToString_rhs_1 : String

y : Int
-----
valToString_rhs_2 : String
```

- 8 Refine**—Finally, now that you know the type of the input `y` in each case, you can complete the definition as before:

```
valToString : (isInt : Bool) -> (case isInt of
                                      False => String
                                      True => Int) -> String
valToString False y = trim y
valToString True y = cast y
```

Totality and type-level functions

In general, it's best to consider type-level functions in exactly the same way as ordinary functions, as we've done so far. This isn't always the case, though. There are a couple of technical differences that are useful to know about:

- Type-level functions exist at *compile time only*. There's no runtime representation of `Type`, and no way to inspect a `Type` directly, such as pattern matching.
- Only functions that are total will be evaluated at the type level. A function that isn't total may not terminate, or may not cover all possible inputs. Therefore, to ensure that type-checking itself terminates, functions that are not total are treated as *constants* at the type level, and don't evaluate further.

6.2 Defining functions with variable numbers of arguments

You can use type-level functions to calculate types based on some other known input. Given that function types are themselves types, this means that you can write functions with a different number of arguments depending on some other input. This is similar to some other languages that support variable-length argument lists, but with additional precision in the type because you use the *value* of one argument to compute the *types* of the others.

In this section, we'll see a couple of examples of how this can work:

- As an introductory example, we'll write a function that adds a sequence of numbers, where the first argument is a number used to calculate the type of a function that takes that many inputs.
- We'll use the same technique to write a variant of the `printf` function that produces a formatted output `String` using a format specifier in its first argument to describe the form of later arguments.

6.2.1 An addition function

First, we'll define an addition function that adds together a sequence of numbers given directly as function arguments. Its behavior is characterized by the three examples shown in figure 6.1. An expression of the form `adder numargs val` calculates a function that takes `numargs` additional arguments, which are added to an initial value `val`.

As usual in type-driven development, you'll begin writing `adder` by writing its type, but in this case the type is not something you can construct directly; the type of `adder` is different depending on the value of the first argument. For the examples shown in figure 6.1, you're looking for the following types:

```
adder 0 : Int -> Int
adder 1 : Int -> Int -> Int
adder 2 : Int -> Int -> Int -> Int
...
```

Because types are first-class, and this type differs depending on some value, you'll be able to compute it using a type-level function. You can write an `AdderType` function with the required behavior. The name `AdderType` follows the convention that type-level functions are given names with an initial capital letter, and it indicates that it's used to compute the type of `adder`.




Figure 6.1 Behavior of an addition function with a variable number of arguments

You can use a Nat to give the length of the argument list, both because you can case-split on it conveniently, and because it would be meaningless to have a negative length argument list. The following listing gives the definition of AdderType.

Listing 6.3 A function to calculate a type for adder n (Adder.idr)

```
AdderType : (numargs : Nat) -> Type
AdderType Z = Int
AdderType (S k) = (next : Int) -> AdderType k
```

The type of adder can now be computed by passing its first argument to AdderType. The second argument is the initial value, which we'll call acc as an abbreviation for “accumulator”:

```
adder : (numargs : Nat) -> (acc : Int) -> AdderType numargs
```

Because you calculate the type by case splitting on numargs in AdderType, you can write the definition of adder with a corresponding structure by case splitting on numargs, so that AdderType will be refined for each case. You can implement it with the following steps:

- 1 *Define*—Add a skeleton definition, and then case-split on the first argument:

```
adder : (numargs : Nat) -> (acc : Int) -> AdderType numargs
adder Z acc = ?adder_rhs_1
adder (S k) acc = ?adder_rhs_2
```

- 2 *Refine*—Where the number of additional arguments, numargs, is Z, return the accumulator directly:

```
adder : (numargs : Nat) -> (acc : Int) -> AdderType numargs
adder Z acc = acc
adder (S k) acc = ?adder_rhs_2
```

- 3 *Type*—For adder_rhs_2, inspecting the type of the hole shows that you need to provide a function:

```
k : Nat
acc : Int
-----
adder_rhs_2 : Int -> AdderType k
```

This is a function type because in AdderType when numargs matches a nonzero Nat, the expected type is a function type.

- 4 *Refine*—The only way you have of producing something of type AdderType k in general is by calling adder with k as the first argument, so this type hints that you need to make a recursive call to adder. This is the complete definition:

```

adder : (numargs : Nat) -> (acc : Int) -> AdderType numargs
adder Z acc = acc
adder (S k) acc = \next => adder k (next + acc)

```

Now that you have a complete and working definition, it's a good idea to think about how you might refine either the type or the definition itself. For example, adder could be made generic in the type of numbers it adds.

Listing 6.4 shows a slightly refined version of the adder function that works with any numeric type, not just `Int`. It does this by passing an additional `Type` argument to `AdderType`, and then constraining that to numeric types in the type of adder.

Listing 6.4 A generic adder that works for any numeric type (Adder.idr)

```

→ AdderType : (numargs : Nat) -> Type -> Type
AdderType Z numType = numType
AdderType (S k) numType = (next : numType) -> AdderType k numType

adder : Num numType =>
      (numargs : Nat) -> numType -> AdderType numargs numType ←
adder Z acc = acc
adder (S k) acc = \next => adder k (next + acc)

```

The Type here is the type of
arguments you'll be adding together.


Constrains the types you can add to
numeric types, using `Num numType` and
then passing `numType` to `AdderType`

The adder function illustrates the basic pattern for defining functions with variable numbers of arguments: you've written an `AdderType` function to calculate the desired type of adder, given one of adder's inputs. The pattern can also be applied for larger definitions, as you'll now see when we define a function for formatting output.

6.2.2 Formatted output: a type-safe `printf` function

A larger example of a function with a variable number of arguments is `printf`, found in C and some other languages. It produces formatted output given a format string and a variable number of arguments, as determined by the format string. The format string essentially gives a template string to be output, populated by the remaining arguments. In essence, `printf` has the same overall structure as `adder`, using the format string to calculate the types of the later arguments.

Figure 6.2 shows some examples that characterize the `printf` behavior. Note that rather than producing output to the console, our version of `printf` returns a `String`.



In these format strings, the directive `%d` stands for an `Int`; `%s` stands for a `String`; and anything else is printed literally.

Following our usual process of type, define, refine, we'll begin by thinking about a type for `printf`. As with `adder`, we'll start with the types of the characteristic examples and then work out how to write a function that computes these types. These are the types of the examples in figure 6.2:

```
printf "Hello!" : String
printf "Answer: %d" : Int -> String
printf "%s number %d" : String -> Int -> String
```

FORMAT STRINGS In a full implementation of `printf` as provided by C, there are far more directives available than `%d` and `%s`. Furthermore, the directives can be modified in various ways to indicate further how the output should be formatted (such as leading zeroes in a number). Such details don't add anything to this discussion of type-level functions, however, so we'll omit them here.

To get the type of `printf`, you'll need to use the format string to build the expected types of the arguments. Instead of processing the string directly, you can write a data type describing the possible formats, as follows.

Listing 6.5 Representing format strings as a data type (`Printf.idr`)

```
This represents %d, followed by the rest of the format specifier.
→ data Format = Number Format
    | Str Format
    | Lit String Format
    | End
A literal string, followed by the rest of the format specifier
This represents %s, followed by the rest of the format specifier.
An empty format specifier
```

This gives a clean separation between the *parsing* of the format string and the *processing*, much as we did when parsing the commands to the data store in chapter 4. For example, `Str (Lit " = " (Number End))` would represent the format string `"%s = %d"`, as illustrated in figure 6.3.




Figure 6.3 Translating a format string to a Format description

Intermediate types

In the type-driven development process, we often think about functions in terms of transformations between data types. As such, we often define intermediate types, such as `Format` in this section, to describe intermediate stages of a computation. There are two main reasons for doing this:

- The meaning of `String` is not obvious from the type alone. The function type `Format -> Type` has a more precise meaning than the function type `String -> Type` because it's clear that the input must be a format specification rather than any `String`.
- Defining an intermediate data type gives us access to more interactive editing features, particularly case splitting.

For the moment, we'll work directly with `Format`; we'll define a conversion from `String` to `Format` later. The following listing shows how you can compute the type of `printf` from a `Format` specifier.

Listing 6.6 Calculating the type of `printf` from a `Format` specifier (extending `Printf.idr`)

The `Number` directive means that your `printf` function will need another `Int` argument.

```
PrintfType : Format -> Type
→ PrintfType (Number fmt) = (i : Int) -> PrintfType fmt
PrintfType (Str fmt) = (str : String) -> PrintfType fmt
→ PrintfType (Lit str fmt) = PrintfType fmt
PrintfType End = String
```

The `Str` directive means that your `printf` function will need another `String` argument.

← This gives the return type of `printf`.

No additional argument is needed here because you have a literal string. You can calculate the type from the rest of the `Format`.

Listing 6.7 defines a helper function for building a `String` from a `Format`, along with any necessary additional arguments as calculated by `PrintfType`. This works like `adder`, using an accumulator to build the result.

Listing 6.7 Helper function for `printf`, building a `String` from a `Format` specifier (`Printf.idr`)

The `String` is an accumulator, in which you build the `String` to be returned.

```
→ printfFmt : (fmt : Format) -> (acc : String) -> PrintfType fmt
→ printfFmt (Number fmt) acc = \i => printfFmt fmt (acc ++ show i)
printfFmt (Str fmt) acc = \str => printfFmt fmt (acc ++ str)
printfFmt (Lit lit fmt) acc = printfFmt fmt (acc ++ lit)
printfFmt End acc = acc
```

`PrintfType` calculates a function type from `Number` `fmt`, so you need to build a function here.

At the end, there are no further arguments to read and no further literal inputs, so return the accumulator.

REMEMBER INTERACTIVE EDITING! For examples such as those in listings 6.6 and 6.7, don't just type them in directly. Instead, use the interactive editing tools in Atom to try to reconstruct them yourself. While doing so, make sure you take a look at the types of any holes, and see how far expression search can take you.

Finally, to implement a `printf` that takes a `String` as a format specifier rather than a `Format` structure, you'll need to be able to convert a `String` into a `Format`. The following listing defines the top-level `printf` function that does this conversion.

Listing 6.8 Top-level definition of `printf`, with a conversion from `String` to `Format` (`Printf.idr`)

Use List Char rather than String here so that you can easily match on individual characters.

`strCons` builds a `String` from an initial character and the rest of the string.

```
→ toFormat : (xs : List Char) -> Format
toFormat [] = End
toFormat ('%' :: 'd' :: chars) = Number (toFormat chars)
toFormat ('%' :: 's' :: chars) = Str (toFormat chars)
toFormat ('%' :: chars) = Lit "%" (toFormat chars)
toFormat (c :: chars) = case toFormat chars of
    Lit lit chars' => Lit (strCons c lit) chars' ←
    fmt => Lit (strCons c "") fmt

printf : (fmt : String) -> PrintfType (toFormat (unpack fmt))
printf fmt = printfFmt _ "" ←
            Use an underscore (_) for the format,
            because Idris can infer from the type
            that it must be toFormat (unpack fmt).
```

Exercises

- 1 An $n \times m$ matrix can be represented by nested vectors of `Double`. Define a type synonym: `Matrix : Nat -> Nat -> Type`.

You should be able to use it to define the following matrix:

```
testMatrix : Matrix 2 3
testMatrix = [[0, 0, 0], [0, 0, 0]]
```

- 2 Extend `printf` to support formatting directives for `Char` and `Double`.

You can test your answer at the REPL as follows:

```
*ex_6_2> :t printf "%c %f"
printf "%c %f" : Char -> Double -> String

*ex_6_2> printf "%c %f" 'X' 24
"'X' 24.0" : String
```

- 3 You could implement a vector as nested pairs, with the nesting calculated from the length, as in this example:

```
TupleVect 0 ty = ()
TupleVect 1 ty = (ty, ())
TupleVect 2 ty = (ty, (ty, ()))
...
...
```

Define a type-level function, `TupleVect`, that implements this behavior. Remember to start with the type of `TupleVect`.

When you have the correct answer, the following definition will be valid:

```
test : TupleVect 4 Nat
test = (1,2,3,4,())
```

6.3 Enhancing the interactive data store with schemas

In the interactive data store we developed in chapter 4, you were able to add `Strings` to an in-memory store and retrieve them by index. But what if you want to store more-complex data? And what if you'd like the form of the data to be determined by the *user* before entering any data, rather than by hardcoding it in the program itself?

In the rest of this chapter, we'll extend the data store by adding schemas to describe the form of the data. We'll determine the schema by user input, and we'll use type-level functions to compute the correct type for the data. Figure 6.4 shows two different data stores, with different schemas, that we'll be able to represent with our extended system. Schema 1, at the top, shows a store that requires the data to be of type `(Int, String)`, and schema 2, below, shows a store that requires the data to be of type `(String, String, Int)`. In contrast, in the data store developed in chapter 4, the schema was effectively always `String`.

Schema 1: (Int, String)	
0	(11, "Armstrong, Aldrin and Collins")
1	(17, "Cernan, Evans and Schmitt")
2	(8, "Borman, Lovell and Anders")

Schema 2: (String, String, Int)	
0	("Rain Dogs", "Tom Waits", 1985)
1	("Fog on the Tyne", "Lindisfarne", 1971)
2	("Flood", "They Might Be Giants", 1990)

Figure 6.4 Two different data stores, with different schemas. Schema 1 requires the data to be of type `(Int, String)`, and schema 2 requires the data to be of type `(String, String, Int)`

A typical interaction with the extended system might proceed as follows. Note that we're describing the schema *before* entering any data.

```
Command: schema String String Int
OK
Command: add "Rain Dogs" "Tom Waits" 1985
ID 0
Command: add "Fog on the Tyne" "Lindisfarne" 1971
ID 1
Command: get 1
"Fog on the Tyne", "Lindisfarne", 1971
Command: quit
```

Rather than starting from scratch, we'll begin with the existing data store we implemented in chapter 4 and refine the overall system as necessary. We'll take this approach:

- 1 Refine the representation of `DataStore` to include schema descriptions. This refinement will inevitably mean that our program no longer type-checks.
- 2 Correct any errors, introducing holes for any more-complex errors.
- 3 Fill in the holes to complete the implementation, and add any further features that are now supported as a result of refining the type.

6.3.1 Refining the `DataStore` type

At the moment, the `DataStore` only supports storing `Strings`. We implemented it in chapter 4 using the following type:

```
data DataStore : Type where
    MkData : (size : Nat) -> (items : Vect size String) -> DataStore
```

Instead of using a `Vect size String` for the `items` in the store, we'd like to be flexible in the types of these items. One natural way to do this might be to refine this to a generic version of `DataStore`, parameterizing it over a schema that gives the type of data in the store:

```
data DataStore : Type -> Type where
    MkData : (size : Nat) -> (items : Vect size schema) -> DataStore schema
```

COMMENTING OUT CODE SECTIONS While you're working on the refined `DataStore`, you'll inevitably break the rest of the program, which will no longer type-check. Therefore, I'd suggest commenting out the rest of the code (placing it between `{ - } and - }) until you've finished the new DataStore type and you're ready to move on.`

We want the user to be able to describe and possibly even update the schema, but if we parameterize over a schema type, the schema will be fixed in the type. Instead, we'll create a data type for describing schemas, and a type-level function for translating schema descriptions (possibly given by a user at runtime) into concrete types. The following listing shows an outline of a refined `DataStore` type.

Listing 6.9 Outline of the refined `DataStore` type, with the Schema description and the translation of Schema to concrete types left undefined (`DataStore.idr`)

A data declaration without a body stands for a type that hasn't been defined yet, much like a hole stands for a function that hasn't been defined yet.

⇒ `data Schema`

`SchemaType : Schema -> Type`

```
data DataStore : Type where
    MkData : (schema : Schema) ->
```

Once you've defined `Schema`, you'll be able to fill in this hole to convert a `Schema` to a concrete type.

Store the schema description itself in the data store.

```
(size : Nat) ->
(items : Vect size (SchemaType schema)) -> DataStore
```

← Calculate the required type of the items in the store from the schema.

We'll define Schemas as being some combination of Strings and Ints, according to a definition given by the user. The user will provide a Schema, and we'll translate the Schema to some concrete type using a type-level function, SchemaType.

The next listing shows how you can describe Schemas and convert them into Idris Types using a type-level function, SchemaType.

Listing 6.10 Defining Schema, and converting a Schema to a concrete type

```
infixr 5 .+.

data Schema = SString
| SInt
| (..) Schema Schema

SchemaType : Schema -> Type
SchemaType SString = String
SchemaType SInt = Int
SchemaType (x .+ y) = (SchemaType x, SchemaType y)
```

A schema containing a single String

Idris allows us to introduce new operators (see the sidebar, “Declaring operators”).

A schema combining two smaller schemas

A schema containing a single Int

A type-level function for converting a Schema to a concrete type

You can try this for defining schemas for the two example stores shown in figure 6.4 earlier:

```
*DataStore> SchemaType (SInt .+ SString)
(Int, String) : Type

*DataStore> SchemaType (SString .+ SString .+ SInt)
(String, String, Int) : Type
```

Declaring operators

You can define new operators by giving their fixity and precedence. In listing 6.10 you had this:

```
infixr 5 .+.
```

This introduces a new right-associative infix operator (infixr) with precedence level 5. In general, operators are introduced with the keyword infixl (for left-associative operators), infixr (for right-associative operators) or infix (for non-associative operators), followed by a precedence level and a list of operators.

Even the arithmetic and comparison operators are defined this way, rather than being built-in syntax. They're introduced as follows in the Prelude:

```
infixl 5 ==, !=
infixl 6 <, <=, >, >=
infixl 7 <<, >>
```

(continued)

```
infixl 8 +, -
infixl 9 *, /
```

The `:` and `++` operators on lists are also defined in the Prelude and are declared as follows:

```
infixr 7 :::, ++
```

The new `DataStore` type allows you to store not only the size and the contents of the store, but also a description of the structure of the contents of the store, as a schema. Previously, each entry was always a `String`, but now the form is determined by the user.

Now, because you've changed the definition of `DataStore`, you'll also need to change the functions that access it.

6.3.2 Using a record for the DataStore

In order to be able to use the updated `DataStore` type, you'll need to redefine the functions `size` and `items` to project the relevant fields out of the structure.

The definition of `size` is similar to the previous definition:

```
size : DataStore -> Nat
size (MkData schema' size' items') = size'
```

To define `items`, however, you'll also need to write a function to project a schema out of the store, because you need to know the schema description in order to know the type of the items in the store.

```
schema : DataStore -> Schema
schema (MkData schema' size' items') = schema'

items : (store : DataStore) -> Vect (size store) (SchemaType (schema store))
items (MkData schema' size' items') = items'
```

Writing projection functions like these, which essentially extract fields from records, can get tedious very quickly. Instead, Idris provides a notation for defining *records*, which leads to automatically generated functions for projecting fields from a record. You can define `DataStore` as follows.

Listing 6.11 Implementing DataStore as a record, with automatically generated projection functions

```
record DataStore where
    constructor MkData      ←———— Names the data constructor for DataStore
    schema : Schema
    size : Nat
    items : Vect size (SchemaType schema)
```

Declares fields, which automatically generate projection functions with the same name

A record declaration introduces a new data type, much like a data declaration, but with two differences:

- There can be *only one* constructor.
- The fields give rise to projection functions, automatically generated from the types of the fields.

In the case of DataStore, you can see the types of the MkData data constructor and the projection functions generated from the fields by using :doc:

```
*DataStore> :doc DataStore
Record DataStore

Constructor:
  MkData : (schema : Schema) ->
    (size : Nat) ->
      (items : Vect size (SchemaType schema)) -> DataStore

Projections:
  schema : (rec : DataStore) -> Schema

  size : (rec : DataStore) -> Nat

  items : (rec : DataStore) ->
    Vect (size rec) (SchemaType (schema rec))
```

You can try this by creating a simple test record at the REPL:

```
*DataStore> :let teststore = (MkData (SString .+ SInt) 1 [("Answer", 42)])
*DataStore> :t teststore
teststore : DataStore
```

Next, you can project the schema, the size, and the list of items from this test record:

```
*DataStore> schema teststore
SString .+ SInt : Schema

*DataStore> size teststore
1 : Nat

*DataStore> items teststore
[("Answer", 42)] : Vect 1 (String, Int)
```

Records are actually much more flexible than can be seen in this small example. As well as projecting out the values of fields, Idris provides a syntax for setting fields and updating records. You'll learn more about records when we discuss working with state in chapter 12.

6.3.3 Correcting compilation errors using holes

Now that you have a new definition of DataStore, your old program that uses it will no longer type-check because it relies on the old definition. The next step in refining your data store program, then, is to update the definitions so that the whole program

type-checks again. This doesn't necessarily mean completing the program; at this stage, it's fine to resolve type errors by inserting holes that you'll fill in later.

Earlier, I suggested temporarily commenting out the code after the definition of `DataStore` so that you could work on the refined definition without worrying about compilation errors. Now, we'll work through the remainder of the program, uncommenting definitions bit by bit and repairing them, guided by the type errors Idris gives us.

First, let's uncomment `addToStore`, defined previously as follows:

```
addToStore : DataStore -> String -> DataStore
addToStore (MkData size store) newitem = MkData _ (addToData store)
  where
    addToData : Vect oldsize String -> Vect (S oldsize) String
    addToData [] = [newitem]
    addToData (item :: items) = item :: addToData items
```

On reloading, either by using Ctrl-Alt-R in Atom or the `:r` command at the REPL, Idris reports as follows:

```
DataStore.idr:21:1-11:
When checking left hand side of addToStore:
When checking an application of Main.addToStore:
  Type mismatch between
    Vect size (SchemaType schema) ->
    DataStore (Type of MkData schema size)
  and
    DataStore (Expected type)
```

The first problem here is that you've added an argument to `MkData`; it now requires a schema as well as a size and a vector of items. You can correct this by adding a `schema` argument to `MkData`:

```
addToStore : DataStore -> String -> DataStore
addToStore (MkData schema size store) newitem
  = MkData schema _ (addToData store)
  where
    addToData : Vect oldsize String -> Vect (S oldsize) String
    addToData [] = [newitem]
    addToData (item :: items) = item :: addToData items
```

Idris now reports

```
Type mismatch between
  Vect size (SchemaType schema) (Type of store)
and
  Vect size String (Expected type)
```

The problem is that the data store no longer stores merely `String`s, but stores a type described by the schema. You can correct this by changing the types of `addToStore` and `addToData` so that they work with the correct type. A type-correct definition of `addToStore` is shown in the following listing.

Listing 6.12 A corrected definition of addToStore using the refined DataStore type

```

addToStore : (store : DataStore) -> SchemaType (schema store) -> DataStore
addToStore (MkData schema size store) newitem
    = MkData schema _ (addToData store)
where
    addToData : Vect oldsize (SchemaType schema) ->
        Vect (S oldsize) (SchemaType schema)
    addToData [] = [newitem]
    addToData (item :: items) = item :: addToData items

```

Calculates the type of the item you're adding using SchemaType, from the schema defined in the store

The name schema here refers to the name bound in the preceding clause

If you continue uncommenting definitions one at a time, the next error you'll encounter is in getEntry. It's currently defined as follows, with the erroneous line marked.

Listing 6.13 Old version of getEntry, with an error in the application of index

```

getEntry : (pos : Integer) -> (store : DataStore) ->
    Maybe (String, DataStore)
getEntry pos store
    = let store_items = items store in
        case integerToFin pos (size store) of
            Nothing => Just ("Out of range\n", store)
            Just id => Just (index id (items store)
                ++ "\n", store)

```

There's an error in the application of index because the store no longer represents items as a Vect containing Strings.

The problem is in the last line, where you extract an item from the store, because you're treating the store as a Vect containing Strings. Here's what Idris reports:

When checking an application of function Data.Vect.index:
 Type mismatch between
 Vect (size store)
 (SchemaType (schema store)) (Type of items store)
 and
 Vect (size store) String (Expected type)

The problem is in the application of index, which no longer returns a String. You can correct this error, temporarily, by inserting a hole to convert the result of index into a String, as follows.

Listing 6.14 Correcting getEntry by inserting a hole to convert the contents of the store to a displayable String

```

getEntry : (pos : Integer) -> (store : DataStore) ->
    Maybe (String, DataStore)
getEntry pos store
    = let store_items = items store in
        case integerToFin pos (size store) of
            Nothing => Just ("Out of range\n", store)
            Just id => Just (?display (index id (items store)) ++ "\n",
                store)

```

The ?display hole, when filled in, will be a function that converts the result of index, which is a SchemaType (schema store), into a String that can be displayed.

If you check the type of `display`, you'll see the type of the function you need to fill in the hole, converting a `SchemaType (schema store)` into a `String`:

```
store : DataStore
id : Fin (size store)
pos : Integer
store_items : Vect (size store) (SchemaType (schema store))
-----
display : SchemaType (schema store) -> String
```

We'll return to `?display` shortly. For the moment, `getEntry` type-checks again. The next error is in `processInput`. Here's the current definition.

Listing 6.15 Old version of `processInput`, with an error in the application of `addToString`

```
processInput : DataStore -> String -> Maybe (String, DataStore)
processInput store input
  = case parse input of
      Nothing => Just ("Invalid command\n", store)
      Just (Add item) =>
        Just ("ID " ++ show (size store) ++ "\n", addToString store item)
        Just (Get pos) => getEntry pos store
        Just (Quit) => Nothing
```

There's an error in the application of `addToString` here, because you're passing it a `String` and it now expects an entry as described by the schema type.

This definition has an error similar to `getEntry`, showing that you have a `String` but Idris expected a `SchemaType (schema store)`:

```
When checking an application of function Main.addToString:
  Type mismatch between
    String (Type of item)
  and
    SchemaType (schema store) (Expected type)
```

One possible fix is, as with `getEntry`, to add a hole for converting the `String` to an appropriate `SchemaType (schema store)` in `processInput`:

```
Just ("ID " ++ show (size store) ++ "\n", addToString store (?convert item))
```

Alternatively, you could refine the definition of `Command` so that it only represents valid commands, meaning that any user input that's invalid would lead to a parse error. We'll take this approach, because it involves defining a more precise intermediate type, so you'll check the validity of the input as early as possible.

To achieve this, parameterize `Command` by the schema description, and change the `Add` command so that it takes a `SchemaType` rather than the `String` input directly. Here's the refined definition of `Command`.

Listing 6.16 Command, refined to be parameterized by the schema in the data store

```
→ data Command : Schema -> Type where
    Add : SchemaType schema -> Command schema
    Get : Integer -> Command schema
    Quit : Command schema
```

Command is parameterized by the Schema description, which gives the form of entries that can be added to the store.

The type of Add now makes it explicit that only inputs that conform to the schema can be added.

You can now change parse to take an explicit schema description, and add a hole where necessary to convert String input to SchemaType schema. A new definition of parse that type-checks is shown here.

Listing 6.17 Updating parseCommand so that it parses inputs that conform to the schema

```
parseCommand : (schema : Schema) -> String -> String -> Maybe (Command schema)
parseCommand schema "add" rest = Just (Add (?parseBySchema rest))
parseCommand schema "get" val = case all isDigit (unpack val) of
    False => Nothing
    True => Just (Get (cast val))
parseCommand schema "quit" "" = Just Quit
parseCommand _ _ _ = Nothing
→ parse : (schema : Schema) ->
    (input : String) -> Maybe (Command schema)
parse schema input = case span (/= ' ') input of
    (cmd, args) => parseCommand schema cmd (ltrim args)
```

Adds a hole for converting the String input to the required SchemaType schema

Adds an explicit schema argument that can be passed to parseCommand to tell it the form of valid inputs

The resulting hole has a type that explains that it converts a String to an appropriate instance of the SchemaType schema:

```
schema : Schema
rest : String
-----
parseBySchema : String -> SchemaType schema
```

There's still a problem here, however! This function can't be total because not every String is going to be parsable as a valid instance of the schema. Nevertheless, your goal at the moment is merely to make the overall program type-check again. We'll return to this problem shortly.

You've refined the type of Command, added the schema argument to parse, and inserted holes for displaying entries (?display) and converting user input into entries (?parseBySchema). All that remains is to update processInput and main to use the new definitions. These are minor changes, shown in the next listing. In processInput, you pass the current schema to parse, and in main you set the initial schema to simply accept Strings.

Listing 6.18 Updated processInput and main, with a default schema

```

Adds an extra argument to parse so that it      Because item has type SchemaType (schema store)
knows which schema to use to parse the data   here, it's fine to call addToStore as before.

processInput : DataStore -> String -> Maybe (String, DataStore)
processInput store input
  = case parse (schema store) input of
    Nothing => Just ("Invalid command\n", store)
    Just (Add item) =>
      Just ("ID " ++ show (size store) ++ "\n", addToStore store item) ←
    Just (Get pos) => getEntry pos store
    Just Quit => Nothing

main : IO ()
main = replWith (MkData SString _ [])
  "Command: " processInput

```

Sets the initial schema as SString, representing only Strings. We'll add a way for users to set the schema shortly.

To recap, you've updated the `DataStore` type to allow user-defined schemas, defining it using a record to get field access functions for free, and you've updated the remainder of the program so that it now type-checks, inserting holes temporarily for the parts that are more difficult to correct immediately.

6.3.4 Displaying entries in the store

You now have two holes to fill in before you can execute this program. The first is `?display`, which converts an entry in the store into a `String`, where the schema in the store gives the form of the data:

```

store : DataStore
id : Fin (size store)
pos : Integer
store_items : Vect (size store) (SchemaType (schema store))
-----
display : SchemaType (schema store) -> String

```

In this case, using Ctrl-Alt-L to lift the hole to a top-level function gives you a lot of information that you don't need to implement `display`. All you really need is a schema description and the data.

Instead, you can implement `display` by hand as follows:

- 1 Type**—First, you can write a more generic type than Idris suggested. Rather than specifically using a schema extracted from a `store`, you can display data according to *any* schema:

```
display : SchemaType schema -> String
```

- 2 Define**—In order to define this function, you need to know about the schema itself. Otherwise, you won't know what form the data is in. Add a skeleton definition, and then bring the implicit argument `schema` into scope:

```
display : SchemaType schema -> String
display {schema} item = ?display_rhs
```

- 3 Define**—You can define the function by case splitting on schema. Because the type of item is SchemaType schema, case splitting on schema will give you more information about the expected type of item:

```
display : SchemaType schema -> String
display {schema = SString} item = ?display_rhs_1
display {schema = SInt} item = ?display_rhs_2
display {schema = (x .+. y)} item = ?display_rhs_3
```

- 4 Type**—Inspecting each of the resulting holes (?display_rhs_1, ?display_rhs_2 and ?display_rhs_3) tells you what item must be in each case:

```
item : String
-----
display_rhs_1 : String
item : Int
-----
display_rhs_2 : String
x : Schema
y : Schema
item : (SchemaType x, SchemaType y)
-----
display_rhs_3 : String
```

- 5 Refine**—For ?display_rhs_1 and ?display_rhs_2 you can complete the definition by directly converting item to a String. For ?display_rhs_3, you can case-split on item and recursively display each entry:

```
display : SchemaType schema -> String
display {schema = SString} item = show item
display {schema = SInt} item = show item
display {schema = (x .+. y)} (iteml, itemr)
    = display iteml ++ ", " ++ display itemr
```

Once this definition is complete and the file is reloaded into the Idris REPL, there should be one remaining hole, ?parseBySchema.

6.3.5 Parsing entries according to the schema

The remaining hole, ?parseBySchema, is intended to convert the String the user entered into an appropriate type for the schema. You can see what's expected by looking at its type:

```
rest : String
schema : Schema
-----
parseBySchema : String -> SchemaType schema
```

As you saw earlier, not every String will lead to a valid SchemaType schema, so you can refine this type slightly and create a top-level function that returns a Maybe (SchemaType schema) to reflect the fact that parsing the input might fail, additionally making schema explicit:

```
parseBySchema : (schema : Schema) -> String -> Maybe (SchemaType schema)
```

Then, you can edit `parseCommand` to use this new function. If `parseBySchema` fails (that is, returns `Nothing`), `parseCommand` should also return `Nothing`:

```
parseCommand schema "add" rest = case parseBySchema schema rest of
    Nothing => Nothing
    Just restok => Just (Add restok)
```

To parse complete inputs as described by schemas, you'll need to be able to parse portions of the input according to a subset of the schema. For example, given a schema (`SInt .+ SString`) and an input `100 "Antelopes"`, you'll need to be able to parse the prefix `100` as `SInt`, followed by the remainder, `"Antelopes"`, as `SString`.

You can therefore define your parser with the following two components:

- `parsePrefix` reads a prefix of the input according to the schema and returns the parsed input, if successful, paired with the remainder of the text.
- `parseBySchema` calls `parsePrefix` with some input and ensures that once it has parsed the input according to the schema, there's no input remaining.

The next listing shows the top-level implementation of `parseBySchema` and the type of the `parsePrefix` helper function.

Listing 6.19 Outline implementation of `parseBySchema`, using an undefined `parsePrefix` function to parse a prefix of an input according to a schema

```
parsePrefix : (schema : Schema) -> String -> Maybe (SchemaType schema, String)

parseBySchema : (schema : Schema) -> String -> Maybe (SchemaType schema)
parseBySchema schema input = case parsePrefix schema input of
    Just (res, "") => Just res
    Just _ => Nothing
    Nothing => Nothing
```

Parsing succeeds when
parsePrefix succeeds and the
remainder of the input is empty.

If parsePrefix succeeds but there's input remaining, parsing overall should fail because there's more input than required by the schema.

parsePrefix failed,
so parsing overall
should fail.

Let's take a look at the outline of `parsePrefix`, following the type-define-refine approach and see how the structure of the schema gives us hints about how to proceed with each part of the implementation:

- 1 *Define*—You already have the type, so begin by providing a skeleton definition:

```
parsePrefix : (schema : Schema) -> String ->
    Maybe (SchemaType schema, String)
parsePrefix schema item = ?parsePrefix_rhs
```

- 2 *Define*—If you case-split on `schema`, Idris will generate cases for each possible form of the schema:

```
parsePrefix : (schema : Schema) -> String ->
    Maybe (SchemaType schema, String)
parsePrefix SString input = ?parsePrefix_rhs_1
```

```
parsePrefix SInt input = ?parsePrefix_rhs_2
parsePrefix (x .+ y) input = ?parsePrefix_rhs_3
```

- 3 Type**—Looking at the types of the holes tells you what the return types must be for each form of the schema. For example, in `?parsePrefix_rhs_2` you need to convert the `input` into a `Int`, if possible, paired with the rest of the `input`:

```
input : String
-----
parsePrefix_rhs_2 : Maybe (Int, String)
```

- 4 Refine**—For `?parsePrefix_rhs_2`, if the prefix of the `input` contains digits, you can convert them to an `Int` and return the resulting `Int` and the remainder of the `String`. You can refine it to the following:

```
parsePrefix SInt input = case span isDigit input of
    ("", rest) => Nothing
    (num, rest) => Just (cast num, ltrim rest)
```

If the prefix of the `input` that contains digits is empty, then it's not a valid `Int`, so return `Nothing`. Otherwise, you can convert the prefix to an `Int` and return the rest of the `input`, with leading spaces trimmed using `ltrim`.

- 5 Refine**—You can refine `?parsePrefix_rhs_1` similarly, looking for an opening quotation mark and then reading the string until you reach the closing quotation mark. You'll see the complete definition shortly in listing 6.20.
- 6 Refine**—The type of `?parsePrefix_rhs_3` shows that you need to parse two subschemas and combine the results:

```
x : Schema
y : Schema
input : String
-----
parsePrefix_rhs_3 : Maybe ((SchemaType x, SchemaType y), String)
```

It's a good idea to give `x` and `y` more-meaningful names before proceeding:

```
parsePrefix (schemal .+ schemar) input = ?parsePrefix_rhs_3
```

To refine `parsePrefix_rhs_3`, you can recursively parse the first portion of the `input` according to `schemal`, and if that succeeds, parse the rest of the `input` according to `schemar`. Parse the first portion:

```
parsePrefix (schemal .+ schemar) input
= case parsePrefix schemal input of
    Nothing => Nothing
    Just (l_val, input') => ?parsePrefix_rhs_2
```

If `parsePrefix` on the first part of the schema fails, the whole thing will fail. Otherwise, you have a new hole:

```
schemal : Schema
l_val : SchemaType schemal
input' : String
```

```

schemar : Schema
input : String
-----
parsePrefix_rhs_2 : Maybe ((SchemaType schemal, SchemaType schemar), String)

```

KEEP RELOADING! While following this type-driven approach, you always have a file that type-checks as far as possible. Here, rather than filling the hole completely, you've written a small part of it with a new hole, and checked that what you have type-checks before proceeding.

- 7 *Refine*—Finally, you can complete this case by parsing the remaining input according to schemar:

```

parsePrefix (schemal .+ schemar) input
= case parsePrefix schemal input of
    Nothing => Nothing
    Just (l_val, input') =>
        case parsePrefix schemar input' of
            Nothing => Nothing
            Just (r_val, input'') =>
                Just ((l_val, r_val), input'')

```

NESTED CASE BLOCKS These nested case blocks we've used may seem a little verbose. In section 6.3.7 you'll see one way of writing these more concisely.

The following listing shows the complete implementation of `parsePrefix`, filling in the remaining details, including parsing quoted strings. You now have a complete implementation that can compile and run.

Listing 6.20 Parsing a prefix of an input according to a specific schema

Parses a prefix of the input as a quoted string. If the input doesn't begin with a quote character, parsing should fail.

`getQuoted` returns a quoted prefix of a String, with the input broken down into characters as a List Char.

```

parsePrefix : (schema : Schema) -> String -> Maybe (SchemaType schema, String)
→ parsePrefix SString input = getQuoted (unpack input)
  where
    getQuoted : List Char -> Maybe (String, String) ←
    getQuoted ('"' :: xs)
      = case span (/= '"') xs of
          (quoted, '"' :: rest) => Just (pack quoted, ltrim (pack rest))
          _ => Nothing
    getQuoted _ = Nothing

  parsePrefix SInt input = case span isDigit input of
    ("", rest) => Nothing
    (num, rest) => Just (cast num, ltrim rest)

Uses ltrim to remove any leading whitespace from the remainder of the input. ←

```

Parses a prefix of the input as an integer: it takes the prefix of the string that consists entirely of digits. If there are no digits, parsing should fail.

```

parsePrefix (schemal .+ schemar) input
  = case parsePrefix schemal input of
    Nothing => Nothing
    Just (l_val, input') =>
      case parsePrefix schemar input' of
        Nothing => Nothing
        Just (r_val, input'') => Just ((l_val, r_val), input'')

```

↳ Parses a prefix of the string according to schemal, and then the rest of the string according to schemar. If either part fails, parsing overall should fail.

6.3.6 Updating the schema

Although you now have a complete implementation, it still has no more functionality than the previous version, because the schema is initialized as `SString` in `main`, and you haven't yet implemented any way to update this:

```

main : IO ()
main = replWith (MkData SString _ []) "Command: " processInput

```

You can, at least, try out different schemas by updating `main` and recompiling. For example, you could try a schema that accepts two `Strings` and an `Int`:

```

main : IO ()
main = replWith (MkData (SString .+ SString .+ SInt) _ [])
  "Command: " processInput

```

You can compile and run this using `:exec` at the REPL, and try a couple of example entries:

```

*DataStore> :exec
Command: add "Bob Dylan" "Blonde on Blonde" 1965
ID 0
Command: add "Prefab Sprout" "From Langley Park to Memphis" 1988
ID 1
Command: get 0
"Bob Dylan", "Blonde on Blonde", 1965

```

It would be preferable, however, to allow users to define their own schemas, rather than hardcoding them into `main`. To achieve this, you can add a new command for setting a new schema, updating the `Command` data type.

Listing 6.21 The Command data structure with a new command for updating the schema

```

data Command : Schema -> Type where
  SetSchema : (newschema : Schema) -> Command schema
  Add : SchemaType schema -> Command schema
  Get : Integer -> Command schema
  Quit : Command schema

```

↳ New command to set a new schema. Note that the type expresses no relationship between `newschema` and the existing schema.

You'll also need functions to do the following:

- Update the `DataStore` type to hold the new schema. This should only work when the store is empty, because when you change the schema type, it invalidates the current contents of the store.
- Parse a schema description from user input.

You'll also need to update `parseCommand` and `processInput` to deal with parsing and processing the new command. These new functions are implemented using the same process you followed so far in implementing the extended data store. Listing 6.22 shows how parsing the new command works. It adds a user command, schema, followed by a list of `String` and `Int`, and translates this into the `SetSchema` command.

Listing 6.22 Parsing a Schema description, and extending the parser for Command to support setting a new schema

```

parseSchema : List String -> Maybe Schema
parseSchema ("String" :: xs)
    = case xs of
        [] => Just SString
        _ => case parseSchema xs of
                Nothing => Nothing
                Just xs_sch => Just (SString .+. xs_sch)

parseSchema ("Int" :: xs)
    = case xs of
        [] => Just SInt
        _ => case parseSchema xs of
                Nothing => Nothing
                Just xs_sch => Just (SInt .+. xs_sch)

parseSchema _ = Nothing

parseCommand : (schema : Schema) -> String -> String -> Maybe (Command schema)
{- ... rest of definition as before ... -}
parseCommand schema "schema" rest
    = case parseSchema (words rest) of
        Nothing => Nothing
        Just schemaok => Just (SetSchema schemaok)
parseCommand _ _ _ = Nothing

```

Parses an input where the first word is “String”.

Parsing a schema description may fail if the arguments are invalid, so return something of type `Maybe Schema`.

Parses an input where the first word is “Int”.

Parses the schema description by separating the rest of the input into a list of words.

Listing 6.23 shows how updating the schema works once the new command has been parsed. As a design choice, it will only allow the schema to be updated when the store is empty, because there is no general way of updating the contents of the data store with an arbitrarily updated schema (an alternative would be to empty the store when the user changes the schema).

Listing 6.23 Processing the `SetSchema` command, updating the schema description in the `DataStore`

```

setSchema : (store : DataStore) -> Schema -> Maybe DataStore
setSchema store schema = case size store of
    Z => Just (MkData schema _ [])
    S k => Nothing

processInput : DataStore -> String -> Maybe (String, DataStore)
processInput store input
    = case parse (schema store) input of
        Nothing => Just ("Invalid command\n", store)
        Just (Add item) =>

```

Setting a new schema may fail if there are entries in the store, so return `Maybe DataStore` to capture the possibility of failure.

```

Just ("ID " ++ show (size store) ++ "\n", addToStore store item)
Just (SetSchema schema') =>
  case setSchema store schema' of
    Nothing => Just ("Can't update schema\n", store)
    Just store' => Just ("OK\n", store') ←
  Just (Get pos) => getEntry pos store
  Just Quit => Nothing

```

Updating the schema failed, so report an error and keep the old store.

Updating the schema succeeded, so report success and update to the new store.

Finally, you can compile and run this program and try setting a new schema from user input:

```

*DatalStore> :exec
Command: schema Int String
OK
Command: add 99 "Red balloons"
ID 0
Command: add 76 "Trombones"
ID 1
Command: schema String String Int
Can't update schema when entries in store
Command: get 1
76, "Trombones"

```

In the end, using a data type to describe the schema and using that schema to calculate the types of the operations on the data store has a few consequences:

- On reading user input, you can't add the input to the store if it's not valid according to the schema.
- If you change the schema type, you can't invalidate the contents of the store because the type of the store's contents prevents it. Changing the schema type requires you to have an empty store.
- When writing the parser for user input, you can use the description of the schema to guide the implementation of the parser and to build the correct type for the input.

6.3.7 Sequencing expressions with Maybe using do notation

In the data store program, there are several places where we've used a case block to check the result of a function that returns something with a Maybe type, and passed that result on. For example:

```

parseCommand schema "schema" rest
  = case parseSchema (words rest) of
      Nothing => Nothing
      Just schemaok => Just (SetSchema schemaok)

```

Here, `parseCommand`, which returns something of type `Maybe (Command schema)`, has called `parseSchema`, which returns something of type `Maybe Schema`, and has used a case expression to check the result of the call to `parseSchema`.

If `parseSchema` fails, `parseCommand` also fails. Similarly, if `parseSchema` succeeds, then `parseCommand` also succeeds.

You can see a similar pattern in the `maybeAdd` function in the following listing, which adds two values of type `Maybe Int`, returning `Nothing` if either input is `Nothing`.

Listing 6.24 Adding two Maybe Ints (Maybe.idr)

```
maybeAdd : Maybe Int -> Maybe Int -> Maybe Int
maybeAdd x y = case x of
    Nothing => Nothing
    Just x_val => case y of
        Nothing => Nothing
        Just y_val => Just (x_val + y_val)
```

First input was Nothing, so entire computation returns Nothing

Second input was Nothing, so entire computation returns Nothing

Both inputs of form Just val, so add x_val and y_val and return their sum, wrapped in a Just

You can try `maybeAdd` on a few examples, and you'll see that it adds its inputs if both are of the form `Just val` for some concrete value `val`, and that it returns `Nothing` if either of the inputs is `Nothing`:

```
*Maybe> maybeAdd (Just 3) (Just 4)
Just 7 : Maybe Int

*Maybe> maybeAdd (Just 3) Nothing
Nothing : Maybe Int

*Maybe> maybeAdd Nothing (Just 4)
Nothing : Maybe Int
```

A common pattern is found here, in `parseCommand`, and in several other places throughout the data store implementation:

- 1 Evaluate an expression of type `Maybe ty`. The result is either `Nothing` or `Just x`, where `x` has type `ty`.
- 2 If the result is `Nothing`, the result of the entire computation is `Nothing`.
- 3 If the result is `Just x`, pass `x` to the rest of the computation.

When you see a common pattern, it's a good idea to try to capture that pattern in a higher-order function. In fact, you've already seen an operator that implements a similar pattern in chapter 5, when sequencing `IO` operations:

```
(>>=) : IO a -> (a -> IO b) -> IO b
```

The same operator works for sequencing `Maybe` computations, when defined as follows:

```
(>>=) : Maybe a -> (a -> Maybe b) -> Maybe b
(>>=) Nothing next = Nothing
(>>=) (Just x) next = next x
```

Effectively, it takes the output of the first computation, if successful (that is, returning a `Just`), and passes it on as input to the second. It captures a common pattern of evaluating an expression with a `Maybe` type.

Type of (`>=>`)

Remember from chapter 5 that if you check the type of `(>=>)` at the REPL, you'll see a constrained generic type:

```
Idris> :t (>=>)
(>=>) : Monad m => m a -> (a -> m b) -> m b
```

In general, the `>=>` operator can be used to sequence computations. You'll see how this works in chapter 7 when we discuss interfaces.

Using `(>=>)` as an infix operator, you can rewrite `maybeAdd` more concisely, if somewhat cryptically, as in the following listing.

Listing 6.25 Adding two Maybe Ints using (`>=>`) rather than direct case analysis

```
maybeAdd : Maybe Int -> Maybe Int -> Maybe Int
maybeAdd x y = x >=> \x_val => ← If Idris evaluates the second operand of
                                         >=> here, the definition of >=> means
                                         that x must have the value Just x_val.
If Idris evaluates the second operand of >=> here, y must have the value Just y_val.    ↗ y >=> \y_val =>
                                         Just (x_val + y_val)
```

In chapter 5, you saw that Idris provides a special notation for sequencing computations using `(>=>)`, introduced by the keyword `do`. You can use the same notation here, and write `maybeAdd` as follows.

Listing 6.26 Adding two Maybe Ints using do notation rather than using (`>=>`) directly

```
maybeAdd : Maybe Int -> Maybe Int -> Maybe Int
→ maybeAdd x y = do x_val <- x
                     y_val <- y    ← If y has the form Just y_val,
                                         computation continues;
                                         otherwise, it returns Nothing.
If x has the form Just x_val, computation continues; otherwise, it returns Nothing.    ↗ Just (x_val + y_val)
```

Figure 6.5 shows how an expression using `do` notation is translated into an expression using `(>=>)`. This works exactly the same way as the translation of IO programs written using `do` notation. Just like the translation of list syntax into `Nil` and `(::)`, this translation is purely syntactic. As a result, if you define the `(>=>)` operator in some other context, Idris will allow you to use `do` notation in that context.




Figure 6.5 Transforming do notation to an expression using the `(>=)` operator

With do notation, you could have written the "schema" case for `parseCommand` as follows, letting the do notation, via `(>=)`, take care of the handling of `Nothing` so that you could focus on the successful case of `Just schema`:

```
parseCommand schema "schema" rest
  = do schemaok <- parseSchema (words rest)
      Just (SetSchema schemaok)
```

Exercises



- 1** Update the data store program to support `Char`s in the schema.
- You can test your solution at the REPL as follows:

```
*ex_6_3> :exec
Command: schema Char Int
OK
Command: add x 24
ID 0
Command: add y 17
ID 1
Command: get 0
'x', 24
```

- 2** Modify the `get` command so that, if given no arguments, it prints the entire contents of the data store.

For example:

```
*ex_6_3> :exec
Command: schema Char Int
OK
Command: add x 24
ID 0
Command: add y 17
ID 1
Command: get
0: 'x', 24
1: 'y', 17
```

- 3** Update the data store program so that it uses do notation rather than nested `case` blocks, where appropriate.

6.4 Summary

- Type synonyms are alternative names for existing types, which allow you to give more-descriptive names to types.
- Type-level functions are functions that can be used anywhere that Idris is expecting a Type. Type synonyms are a special case of type-level functions.
- Type-level functions can be used to compute types, and they allow you to write functions with varying numbers of arguments, such as `printf`.
- You can represent a schema for a data store as a type, and then use type-level functions to calculate the types of operations, such as parsing and displaying entries in the data store from the schema description.
- A record is a data type with only one constructor, and with automatically generated functions for projecting fields from the record.
- After refining a data type, you can use holes to correct compilation errors temporarily.
- Using `do` notation, you can sequence computations that use `Maybe` to capture the possibility of errors.



Interfaces: using constrained generic types

This chapter covers

- Constraining generic types using interfaces
- Implementing interfaces for specific contexts
- Using interfaces defined in the Prelude

In chapter 2, you saw that generic function types with type variables could be constrained so that the variables stood for a more limited set of types. For example, you saw the following function for doubling a number in any numeric type:

```
double : Num a => a -> a
double x = x + x
```

The type of `double` includes a constraint, `Num a`, which states that `a` can only stand for numeric types. Therefore, you can use `double` with `Int`, `Nat`, `Integer`, or any other numeric type, but if you try to use it with a non-numeric type, such as `Bool`, Idris will report an error.

You've seen a few such constraints, such as `Eq` for types that support equality tests and `Ord` for types that support comparisons. You've also seen functions that

rely on other constraints that we haven't discussed in detail, such as `map` and `>>=`, which rely on `Functor` and `Monad`, respectively. We haven't yet discussed how these constraints are defined or introduced.

In this chapter, we'll discuss how to define and use constrained generic types using *interfaces*. In the type declaration for `double`, for example, the constraint `Num a` is implemented by an interface, `Num`, which describes arithmetic operations that will be implemented in different ways for different numeric types.

TYPE CLASSES IN HASKELL If you know Haskell, you'll be familiar with Haskell's concept of *type classes*. Interfaces in Idris are similar to type classes in Haskell and are often used in the same way, though there are some differences. The most important are, first, that interfaces in Idris can be parameterized by values of *any* type, and are not limited to types or type constructors, and, second, interfaces in Idris can have multiple implementations, though we won't go into the details in this chapter.

From the perspective of type-driven development, interfaces allow us to give the necessary level of precision to generic types. In general, an interface in Idris describes a collection of generic operations that can be implemented in different ways for different concrete types. For example:

- You could define an interface for operations that serialize and load generic data to and from JSON.
- A graphics library could provide an interface for operations that draw representations of data.

The Prelude includes a wide range of interfaces, and we'll concentrate on these in this chapter. We'll begin by looking in detail at two of the most important.

7.1 Generic comparisons with *Eq* and *Ord*

To begin, we'll look at two interfaces that define generic comparisons, both of which are defined in the Prelude:

- `Eq`, which supports comparing values for equality or inequality.
- `Ord`, which supports comparing values to determine which is larger.

In doing so, you'll learn how to declare interfaces, how to implement those interfaces in specific contexts, and how different interfaces can relate to each other.

7.1.1 Testing for equality with *Eq*

As you saw in chapter 2, Idris provides an operator for testing values for equality, `==`, with a constrained generic type:

```
Idris> :t (==)
(==) : Eq ty => ty -> ty -> Bool
```

In other words, you can compare two values of some generic type `ty` for equality, returning a result as a `Bool`, provided that `ty` satisfies the `Eq` constraint. Similarly, there's an operator for testing for inequality:

```
Idris> :t (/=)
(/=) : Eq ty => ty -> ty -> Bool
```

Specifying multiple constraints

In the example in this section, I've specified exactly one constraint, `Eq ty`. You can also list multiple constraints as a comma-separated list. For example, the following function type specifies that `ty` needs to satisfy the `Num` and `Show` constraints:

```
addAndShow : (Num ty, Show ty) => ty -> ty -> String
```

You'll see more examples later, in section 7.2.2.

Any time you use either of these operators, you must know that its operands are of a type that can be compared for equality. Let's say, for example, you want to write a function that counts the number of occurrences of a specific value, of some generic type `ty`, in a list. We'll create a function called `occurrences` in a file, `Eq.idr`:

- Type**—As always, begin by giving a type for the function and creating a skeleton definition:

```
occurrences : (item : ty) -> (values : List ty) -> Nat
occurrences item xs = ?occurrences_rhs
```

- Define, refine**—You can define the function by case splitting on the input list, `xs`, and refining the hole generated in each case. If it's an empty list, there can be no occurrences of `item`, so return 0:

```
occurrences : (item : ty) -> (values : List ty) -> Nat
occurrences item [] = 0
occurrences item (value :: values) = ?occurrences_rhs_2
```

- Refine**—If the input is a non-empty list, try testing the value at the start of the list and `item` for equality:

```
occurrences : (item : ty) -> (values : List ty) -> Nat
occurrences item [] = 0
occurrences item (value :: values) = case value == item of
                                         case_val => ?occurrences_rhs_2
```

Unfortunately, Idris reports an error:

```
Eq.idr:3:13:
When checking right hand side of occurrences with expected type
Nat

Can't find implementation for Eq ty
```

This problem is similar to the one we encountered in chapter 3 when defining `ins_sort`. In this case, `value` and `item` have type `ty`, but there's no constraint that values of type `ty` are comparable for equality.

- 4 Refine**—The solution, as with `ins_sort`, is to refine the type by adding a constraint:

```
occurrences : Eq ty => (item : ty) -> (values : List ty) -> Nat
```

The constraint `Eq ty` means that you can now use the `==` operator and complete the definition as follows:

```
occurrences : Eq ty => (item : ty) -> (values : List ty) -> Nat
occurrences item [] = 0
occurrences item (value :: values) = case value == item of
                                         False => occurrences item values
                                         True => 1 + occurrences item values
```

The final type of `occurrences` says that it takes a `ty` and a `List ty` as inputs, provided that the type variable `ty` stands for a type that can be compared for equality.

This works fine for built-in types such as `Char` and `Integer`:

```
*Eq> occurrences 'b' ['a','a','b','b','b','c']
3 : Nat

*Eq> occurrences 100 [50,100,100,150]
2 : Nat
```

But what if you have user-defined types? Let's say you have a user-defined type named `Matter`, also defined in `Eq.idr`:

```
data Matter = Solid | Liquid | Gas
```

You'd hope to be able to count the number of occurrences of `Liquid` in a list. Unfortunately, if you try this, Idris reports an error:

```
*Eq> occurrences Liquid [Solid, Liquid, Liquid, Gas]
Can't find implementation for Eq Matter
```

This error message means that Idris doesn't know how to compare values in the user-defined `Matter` type for equality. To correct this, we'll need to take a look at how the `Eq` constraint is defined, and how to explain to Idris how user-defined types such as `Matter` can satisfy it.

7.1.2 Defining the Eq constraint using interfaces and implementations

Constraints such as `Eq` are defined in Idris using interfaces. An interface definition contains a collection of related functions (called *methods* of the interface) that can be given different *implementations* for specific contexts. The `Eq` interface is defined in the Prelude, containing the methods `(==)` and `(/=)`.

Listing 7.1 The Eq interface (defined in the Prelude)

An interface declaration introduces a new interface that can be used to constrain generic functions.

```
interface Eq ty where
  (==) : ty -> ty -> Bool
  (/=) : ty -> ty -> Bool
```

Method declarations introduce functions that must be defined by implementations of the interface.

Defining an interface introduces new top-level functions for each method of the interface. The interface declaration in listing 7.1 introduces the top-level functions `(==)` and `(/=)`. If you check the types of these functions, you'll see the `Eq ty` constraint explicit in their types:

```
(==) : Eq ty => ty -> ty -> Bool
(/=) : Eq ty => ty -> ty -> Bool
```

Figure 7.1 shows the components of the interface declaration. The parameter, `ty` in this case, is assumed to have type `Type` by default, and stands for the generic argument that's to be constrained. The parameter must appear in each of the method declarations.




Figure 7.1 The `Eq` interface declaration

PARAMETER NAMING CONVENTIONS The parameters of an interface (`ty` here) are typically generic type variables, so I typically give them short, generic, names. The name `ty` indicates that the parameter could be *any* type. When there are more parameters, or where the interface has a more specific purpose, I give more-specific names.

An interface declaration, then, provides the types for related functions with a parameter standing for a generic argument. To define these functions for specific cases, you need to provide *implementations*.

To write an implementation, you give the interface and a parameter, along with definitions for each of the methods. The following listing shows an implementation for the `Eq` interface, which explains how the `Matter` type satisfies the constraint.

Listing 7.2 Implementation of Eq for Matter (Eq.idr)

```
Eq Matter where
```

```
  (==) Solid Solid = True
  (==) Liquid Liquid = True
  (==) Gas Gas = True
  (==) _ _ = False
```

An implementation declaration, here explaining how to satisfy the Eq constraint for the Matter type

Implementation of the (==) method for Matter

```
(/=) x y = not (x == y)
```

Implementation of the (/=) method for Matter

INTERACTIVE DEVELOPMENT OF IMPLEMENTATIONS As you'll see shortly, you can use Ctrl-Alt-A to provide a skeleton implementation of `Eq Matter`, as with function definitions.

If you add this implementation to the file in which you defined `occurrences` and `Matter`, you'll now be able to use `occurrences` with a `List Matter`:

```
*Eq> occurrences Liquid [Solid, Liquid, Liquid, Gas]
2 : Nat
```

This provides implementations of the two methods, `(==)` and `(/=)`, which are required to be implemented in order to be able to compare elements of type `Matter` for equality and inequality. Note that the implementation doesn't contain any type declarations, because these are given in the interface declaration.

Interface documentation

If you use `:doc` on an interface at the REPL, or Ctrl-Alt-D over an interface name in Atom, Idris will show documentation for the interface including its parameters, methods, and a list of the known implementations. For example, here's what it looks like for `Eq`:

```
Idris> :doc Eq
Interface Eq
    The Eq interface defines inequality and equality.

Parameters:
    ty

Methods:
    (==) : Eq ty => ty -> ty -> Bool
        infixl 5
        The function is Total
    (/=) : Eq ty => ty -> ty -> Bool
        infixl 5
        The function is Total

Implementations:
    Eq ()
    Eq Int
    Eq Integer
    Eq Double
    [...]
```

You can build the implementation interactively in Atom as follows:

- 1 *Type*—Start by giving the implementation header on its own:

```
Eq Matter where
```

This a Type step because it instantiates the types of `(==)` and `(/=)`, with `Matter` standing for the parameter `a`.

- 2 *Define*—You can also use the interactive editing tools in Atom to build implementations of interfaces. With the cursor over `Eq`, press Ctrl-Alt-A, and Idris will add skeleton definitions for `(==)` and `(/=)` with generic names for the arguments:

```
Eq Matter where
  (==) x y = ?Eq_rhs_1
  (/=) x y = ?Eq_rhs_2
```

- 3 *Type*—If you check the type of `?Eq_rhs_1`, you'll see confirmation that `x` and `y` are of type `Matter`:

```
x : Matter
y : Matter
-----
Eq_rhs_1 : Bool
```

- 4 *Define*—You can begin by defining `(==)` by case splitting on `x`:

```
Eq Matter where
  (==) Solid y = ?Eq_rhs_3
  (==) Liquid y = ?Eq_rhs_4
  (==) Gas y = ?Eq_rhs_5

  (/=) x y = ?Eq_rhs_2
```

- 5 *Refine*—Each value of type `Matter` is only equal to itself, so you can refine the definition as follows, with a catchall case to handle when the inputs aren't equal:

```
Eq Matter where
  (==) Solid Solid = True
  (==) Liquid Liquid = True
  (==) Gas Gas = True
  (==) _ _ = False

  (/=) x y = ?Eq_rhs_2
```

Remember that the order of cases is important, and Idris will try to match the clauses in order, so the catchall case must be at the end.

- 6 *Refine*—To complete the implementation, you need to define `(/=)`. The simplest definition is to use `(==)` and then invert the result:

```
Eq Matter where
  (==) Solid Solid = True
  (==) Liquid Liquid = True
  (==) Gas Gas = True
  (==) _ _ = False

  (/=) x y = not (x == y)
```

When you declare an interface, you introduce a collection of related new generic functions, known as methods, that can be overloaded for specific situations. When you define an implementation of the interface, you must provide definitions for *all* its methods. So, when you defined the Eq instance for Matter, you had to provide definitions for both (==) and (/=).

7.1.3 Default method definitions

Interfaces define a collection of related methods, such as (==) and (/=). In some cases, methods are so closely related that you can define them in terms of other methods in the interface. For example, you'd always expect the result of $x \neq y$ to be the opposite of the result of $x == y$, no matter what the values or even types of x and y .

For this situation, Idris allows you to provide a *default* definition for a method. If an implementation doesn't provide a definition for a method that has a default definition, Idris uses that default. For example, the Eq interface provides defaults for both (==) and (/=), each defined in terms of the other, as follows.

Listing 7.3 Eq interface with default method definitions

```
interface Eq a where
  (==) : a -> a -> Bool
  (/=) : a -> a -> Bool
  (==) x y = not (x /= y)
  (/=) x y = not (x == y)
```

 Default method instance, used if the method isn't present in a specific implementation of an interface

You can therefore provide an implementation of Eq for Matter by providing only a definition of (==), and using the default method implementation for (/=):

```
Eq Matter where
  (==) Solid Solid = True
  (==) Liquid Liquid = True
  (==) Gas Gas = True
  (==) _ _ = False
```

The default method definitions mean that when you're defining an implementation of Eq, you can provide definitions for either or both of (==) and (/=).

7.1.4 Constrained implementations

When writing implementations for generic types, you may discover the need for additional constraints on the parameters of the generic types. For example, to check whether two lists are equal, you'll need to know how to compare the element types for equality.

In chapter 4, you saw the generic type of binary trees, defined as follows in tree.idr:

```
data Tree elem = Empty
              | Node (Tree elem) elem (Tree elem)
```

To check whether two trees are equal, you'll also need to be able to compare the element types. Let's see what happens if you try to define an implementation of `Eq` for `Trees` with a generic element type:

- 1 *Type*—Begin by writing the implementation header, stating the interface you wish to implement and the type for which you're implementing it:

```
Eq (Tree elem) where
```

- 2 *Define*—Add a skeleton definition that gives template definitions for all the methods of the interface:

```
Eq (Tree elem) where
(==) x y = ?Eq_rhs_1
(/=) x y = ?Eq_rhs_2
```

You can use the default definition for `(/=)`, so you can delete the second method:

```
Eq (Tree elem) where
(==) x y = ?Eq_rhs_1
```

- 3 *Define*—As with the `Eq` implementation for `Matter`, you can define `(==)` by case splitting on each argument and using a catchall case for cases where the inputs are unequal. You know that `Empty` equals itself, Nodes are equal if all their arguments are equal, and everything else is not equal:

```
Eq (Tree elem) where
(==) Empty Empty = True
(==) (Node left e right) (Node left' e' right') = ?Eq_rhs_3
(==) _ _ = False
```

For the moment, you've left a hole, `?Eq_rhs_3`, for the details of comparing Nodes.

- 4 *Refine*—You'd expect to be able to complete the definition by refining the `?Eq_rhs_3` hole to say that you can compare each corresponding argument:

```
Eq (Tree elem) where
(==) Empty Empty = True
(==) (Node left e right) (Node left' e' right')
    = left == left' && e == e' && right == right'
(==) _ _ = False
```

But, unfortunately, Idris reports a problem:

```
Can't find implementation for Eq elem
```

You can compare `left` and `left'` for equality, and correspondingly `right` and `right'`, because they're of type `Tree elem`, and the implementation can be recursive; you're currently defining equality for `Tree elem`. But `e` and `e'` are of type `elem`, a generic type, and you don't necessarily know how to compare `elem` for equality.

- 5 *Refine*—The solution is to refine the implementation declaration by *constraining* it to require that `elem` also has an implementation of `Eq` available:

```
Eq elem => Eq (Tree elem) where
  (==) Empty Empty = True
  (==) (Node left e right) (Node left' e' right')
    = left == left' && e == e' && right == right'
  (==) _ _ = False
```

The constraint appears to the left of the arrow, `=>`, the same way as constraints appear in type declarations. Figure 7.2 shows the components of this interface header.

You can read this header as stating that generic trees can be compared for equality, provided that their element type can also be compared for equality. You can introduce interface constraints like this anywhere you introduce a type variable, if you need to constrain that type variable further.




Figure 7.2 Constrained implementation header

LIMITATIONS ON IMPLEMENTATION PARAMETERS You've now seen implementations of `Eq` parameterized by `Matter` and by `Tree elem`, both of which are of type `Type`. But you can't parameterize implementations by *everything* of type `Type`. You're limited to names introduced by either a data or record declaration, or primitive types. In particular, this means you can't parameterize implementations by type synonyms or functions that compute types.

You've already seen constraints on type declarations, and here you've seen one in an implementation definition. In the same way, you can place constraints on interface definitions themselves.

7.1.5 Constrained interfaces: defining orderings with `Ord`

If you place a constraint on an interface definition, you're effectively *extending* an existing interface. In the Prelude, for example, there's an `Ord` interface, shown in the next listing, that extends `Eq` to support ordering of values.

Listing 7.4 The `Ord` Interface, which extends `Eq` (defined in the Prelude)

```
interface Eq ty => Ord ty where
  compare : ty -> ty -> Ordering
  (<) : ty -> ty -> Bool
  (>) : ty -> ty -> Bool
  (≤) : ty -> ty -> Bool
  (≥) : ty -> ty -> Bool
  max : ty -> ty -> ty
  min : ty -> ty -> ty
```

Returns the larger of the two inputs

Declares the `Ord` interface and states that an implementation of `Ord ty` requires an implementation of `Eq ty`

Ordering is defined in the Prelude as either LT, EQ, or GT.

Returns the smaller of the two inputs

All methods, except `compare`, have default definitions, some of which are written in terms of `compare`, and some of which use the `(==)` method provided by the `Eq` interface.

If you have an implementation of `Ord` for some data type, you can sort lists containing that type:

```
sort : Ord ty => List ty -> List ty
```

For example, you might have a data type representing a music collection, with records containing a title, artist, and year of release, and you may wish to sort them first by artist name, then by year of release, and then by title. The following listing gives a definition of this data type with some examples.

Listing 7.5 A record data type and a collection to be sorted (Ord.idr)

```
record Album where
    constructor MkAlbum
    artist : String
    title : String
    year : Integer
    |————— A record declaration (see chapter 6, section 6.3.2)
    |————— Record fields, which give rise to the
    |————— artist, title and year projection functions

help : Album
help = MkAlbum "The Beatles" "Help" 1965

rubbersoul : Album
rubbersoul = MkAlbum "The Beatles" "Rubber Soul" 1965

clouds : Album
clouds = MkAlbum "Joni Mitchell" "Clouds" 1969

hunkydory : Album
hunkydory = MkAlbum "David Bowie" "Hunky Dory" 1971

heroes : Album
heroes = MkAlbum "David Bowie" "Heroes" 1977

collection : List Album
collection = [help, rubbersoul, clouds, hunkydory, heroes]
```

If you try sorting the collection as it stands, Idris will report that it doesn't know how to order values of type `Album`:

```
*Ord> sort collection
Can't find implementation for Ord Album
```

Listing 7.6 shows how you can explain to Idris how to order `Albums`. First, you need to give an `Eq` implementation, because the constraint on the `Ord` interface requires that implementations of `Ord` are also implementations of `Eq`. The `Eq` implementation checks that each field has the same value; the `Ord` implementation compares by artist name, and then by year, and then by title if the first two are equal.

Listing 7.6 Implementations of Eq and Ord for Album

```

Eq Album where
  (==) (MkAlbum artist title year) (MkAlbum artist' title' year')
    = artist == artist' && title == title' && year == year'

Ord Album where
  compare (MkAlbum artist title year) (MkAlbum artist' title' year')
    = case compare artist artist' of
      EQ => case compare year year' of
        EQ => compare title title'
        diff_year => diff_year
      diff_artist => diff_artist

```

The diagram illustrates the implementation of Eq and Ord for the Album type. It shows the Eq implementation constraint and the detailed logic of the Ord implementation.

- Eq Implementation:** The constraint "The Eq implementation is necessary because of the Eq constraint on the Ord interface" is shown above the Eq implementation code.
- Ord Implementation Logic:**
 - If the artists are the same, it compares by year (highlighted in a box).
 - If the artists are different, it returns the result of comparing them directly (highlighted in a box).
 - If the years are the same, it compares by title (highlighted in a box).
 - If the years are different, it returns the result of comparing them directly (highlighted in a box).

Implementing Ord for Album means that you can use the usual comparison operators on values of type Album:

```
*Ord> heroes > clouds
False : Bool

*Ord> help <= rubbersoul
True : Bool
```

It also means that you can use any function with an Ord constraint, such as sort. For example, you can sort the collection and list the titles in sorted order:

```
*Ord> map title (sort collection)
["Hunky Dory", "Heroes", "Clouds", "Help", "Rubber Soul"] : List String
```

Giving constraints on interfaces, like the Eq constraint on Ord, allows you to define *hierarchies* of interfaces. So, for example, if there's an implementation of Ord for some type, you know it's safe to assume that there's also an implementation of Eq for that type. The Prelude defines several interfaces, some arranged into hierarchies, and we'll look at some of the most important in the next section.

Exercises



For these exercises, you'll use the Shape type defined in chapter 4:

```
data Shape = Triangle Double Double
           | Rectangle Double Double
           | Circle Double
```

1 Implement Eq for Shape.

You can test your answer at the REPL as follows:

```
*ex_7_1> Circle 4 == Circle 4
True : Bool

*ex_7_1> Circle 4 == Circle 5
False : Bool

*ex_7_1> Circle 4 == Triangle 3 2
False : Bool
```

2 Implement Ord for Shape. Shapes should be ordered by area, so that shapes with a larger area are considered greater than shapes with a smaller area.

You can test this by trying to sort the following list of Shapes:

```
testShapes : List Shape
testShapes = [Circle 3, Triangle 3 9, Rectangle 2 6, Circle 4,
             Rectangle 2 7]
```

You should see the following when sorting the list at the REPL:

```
*ex_7_1> sort testShapes
[Rectangle 2.0 6.0,
 Triangle 3.0 9.0,
 Rectangle 2.0 7.0,
 Circle 3.0,
 Circle 4.0] : List Shape
```

7.2 Interfaces defined in the Prelude

The Idris Prelude provides several commonly used interfaces, in addition to Eq and Ord, as you've just seen. In this section and the next, I'll briefly describe some of the most important. We've encountered several in passing already: Show, Num, Cast, Functor, and Monad. Here, you'll see how these interfaces are defined, where they might be used, and some examples of how you can write implementations of them for your own types.

The parameters of an interface (in other words, the variables given in the interface header) can have *any* type. If there's no explicit type given in the interface header for a parameter, it's assumed to be of type Type. In this section, we'll look at some interfaces that are parameterized by Types.

7.2.1 Converting to String with Show

In chapter 2, you saw the show function, which converts a value to a String. This is a method of the Show interface, defined in the Prelude and shown in the following listing. Both the show and showPrec methods have default implementations, each defined in terms of the other. Here, we'll only consider show.

Listing 7.7 The Show interface (defined in the Prelude)

```
interface Show ty where
    show : (x : ty) -> String           ← Direct conversion of a value to a String
    showPrec : (d : Prec) -> (x : ty) -> String   ←
    ty is the parameter of the interface
    and is a variable of type Type.          Conversion of a
                                            value to a String in a
                                            precedence context
```

PRECEDENCE CONTEXTS The purpose of `showPrec` is to be able to display complex expressions in parentheses if necessary. This might be useful if, say, you have a representation of arithmetic formulae, where precedence rules say that some subexpressions need to be parenthesized. By default, `showPrec` calls `show` directly, and for most display purposes this is entirely adequate. If you wish to investigate further, take a look at the documentation using `:doc`.

The Prelude provides two functions that use `show` to convert a value to a `String` and then output it to the console, with or without a trailing newline character:

```
printLn : Show ty => ty -> IO ()
print : Show ty => ty -> IO ()
```

You can define a simple `Show` implementation for the `Album` type:

```
Show Album where
    show (MkAlbum artist title year)
        = title ++ " by " ++ artist ++ " (released " ++ show year ++ ")"
```

Then, you can print an `Album` to the console using `printLn`. For example:

```
*Ord> :exec printLn hunkydory
Hunky Dory by David Bowie (released 1971)
```

The `Show` interface is primarily used for debugging output, to display values of complex data types in a human-readable form.

7.2.2 Defining numeric types

The Prelude provides a hierarchy of interfaces with methods for numeric operations. These interfaces divide operators into several groups:

- *The Num interface*—Contains operations that work on all numeric types, including addition, multiplication, and conversion from integer literals
- *The Neg interface*—Contains operations that work on numeric types that can be negative, including negation, subtraction, and absolute value
- *The Integral interface*—Contains operations that work on integer types
- *The Fractional interface*—Contains operations that work on numeric types that can be divided into fractions

The following listing shows how these interfaces are defined.

Listing 7.8 The numeric interface hierarchy (defined in the Prelude)

```

interface Num ty where
  (+) : ty -> ty -> ty
  (*) : ty -> ty -> ty
  fromInteger : Integer -> ty

interface Num ty => Neg ty where
  negate : ty -> ty
  (-) : ty -> ty -> ty
  abs : ty -> ty

interface Num ty => Integral ty where
  div : ty -> ty -> ty
  mod : ty -> ty -> ty

interface Num ty => Fractional ty where
  (/) : ty -> ty -> ty
  recip : ty -> ty

  recip x = 1 / x
  
```

The diagram illustrates the numeric interface hierarchy defined in the Prelude. It shows five interfaces: Num, Neg, Integral, and Fractional, each with its associated operations and a note about its implementation.

- Num**: Values of all numeric types can be added and multiplied.
- Neg**: Converts from an integer literal to the numeric type. Numeric types that can have negative values additionally support subtraction, negation, and absolute value.
- Integral**: Numeric types that are integers additionally support integer division and modulus (remainder).
- Fractional**: Numeric types that can be divided into fractions additionally support division and taking a reciprocal.

It's useful to know which operations are available on which types. Table 7.1 summarizes the numeric interfaces and the implementations that exist for each in the Prelude.

Table 7.1 Summary of numeric interfaces and their implementations

Interface	Description	Implementations
Num	All numeric types	Integer, Int, Nat, Double
Neg	Numeric types that can be negative	Integer, Int, Double
Integral	Integer types	Integer, Int, Nat
Fractional	Numeric types that can be divided into fractions	Double

One method that's particularly worth noting is `fromInteger`. All integer literals in Idris are implicitly converted to the appropriate numeric type using `fromInteger`. As a result, as long as there's an implementation of `Num` for a numeric type, you can use integer literals for that type.

By making new implementations of `Num` and related interfaces, you can use standard arithmetic notation and integer literals for your own types. For example, listing 7.9 shows an `Expr` data type that represents arithmetic expressions, including an “absolute value” operation and an `eval` function that calculates the result of evaluating an expression.

Listing 7.9 An arithmetic expression data type and an evaluator (Expr.idr)

```

data Expr num = Val num
  | Add (Expr num) (Expr num)
  | Sub (Expr num) (Expr num)
  | Mul (Expr num) (Expr num)
  
```

The diagram shows the `Expr` data type definition. It consists of a base case `Val num` and three recursive cases: `Add`, `Sub`, and `Mul`, each taking two `Expr num` arguments. A note indicates that expressions are parameterized by the type of numeric literals.

Expressions are parameterized by the type of numeric literals.

```

| Div (Expr num) (Expr num)
| Abs (Expr num)

eval : (Neg num, Integral num) => Expr num -> num
eval (Val x) = x
eval (Add x y) = eval x + eval y
eval (Sub x y) = eval x - eval y
eval (Mul x y) = eval x * eval y
eval (Div x y) = eval x `div` eval y
eval (Abs x) = abs (eval x)

```

You need num to be negatable, because the evaluator subtracts, and also to support integer division.

To construct an expression, you need to apply the constructors of `Expr` directly. For example, to represent the expression $6 + 3 * 12$, and to evaluate it, you'll need to write it as follows:

```

*Expr> Add (Val 6) (Mul (Val 3) (Val 12))
Add (Val 6) (Mul (Val 3) (Val 12)) : Expr Integer

*Expr> eval (Add (Val 6) (Mul (Val 3) (Val 12)))
42 : Integer

```

If, on the other hand, you make a `Num` implementation for `Expr`, you'll be able to use standard arithmetic notation (with `+`, `*`, and integer literals) to build values of type `Expr`. If, furthermore, you make a `Neg` implementation, you'll be able to use negative numbers and subtraction. The following listing shows how you can do this.

Listing 7.10 Implementations of Num and Neg for Expr (Expr.idr)

```

Num ty => Num (Expr ty) where
  (+) = Add
  (*) = Mul
  fromInteger = Val . fromInteger

Neg ty => Neg (Expr ty) where
  negate x = 0 - x
  (-) = Sub
  abs = Abs

```

You need the constraint `Num ty` because you're using the `ty` implementation of `fromInteger`.

The `(.)` function allows you to compose functions.

Function composition

The `(.)` function gives you a concise notation for composing two functions. It has the following type and definition:

```

(.) : (b -> c) -> (a -> b) -> a -> c
(.) func_bc func_ab x = func_bc (func_ab x)

```

In the example in listing 7.10, you could have written this:

```
fromInteger x = Val (fromInteger x)
```

The `(.)` function allows you instead to write this:

```
fromInteger x = (Val . fromInteger) x
```

(continued)

Finally, rather than have the `x` as an argument on both sides, you can use partial application:

```
fromInteger = Val . fromInteger
```

To construct an expression and evaluate it, you can use standard notation:

```
*Expr> the (Expr _) (6 + 3 * 12)
Add (Val 6) (Mul (Val 3) (Val 12)) : Expr Integer

*Expr> eval (6 + 3 * 12)
42 : Integer
```

In the first case, you need to use `the` to make it clear that the numeric expression should be interpreted as an `Expr`, rather than the default, which would be an `Integer`. In the second case, Idris infers from the type of `eval` that the argument must be an `Expr`.

7.2.3 Converting between types with Cast

In chapter 2, you saw the `cast` function, which is used to convert values between different, compatible types. If you look at the type of `cast`, you'll see that it has a constrained generic type that uses a `Cast` interface:

```
Idris> :t cast
cast : Cast from to => from -> to
```

Unlike the interfaces we've seen so far, `Cast` has *two* parameters rather than one. The next listing gives its definition. Interfaces in Idris can have any number of parameters (even zero!).

Listing 7.11 The Cast interface (defined in the Prelude)

```
interface Cast from to where
  cast : (orig : from) -> to
```

← Cast has two parameters, from and to, both of which are type Type.

As we observed in chapter 2, conversions using `cast` may be lossy. Idris defines casts between `Double` and `Integer`, for example, which may lose precision. The purpose of `cast` is to provide a convenient generic function, with an easy-to-remember name, for conversions.

To define an implementation of an interface with more than one parameter, you need to give both concrete parameters. For example, you could define a cast from `Maybe elem` to `List elem`, because you can think of `Maybe elem` as being a list of either zero or one `elems`:

```
Cast (Maybe elem) (List elem) where
  cast Nothing = []
  cast (Just x) = [x]
```

You could also define a cast in the other direction, from `List elem` to `Maybe elem`, but this would potentially lose information, because you'd need to decide which element to take if the list had more than one element.

Exercises



- 1** Implement `Show` for the `Expr` type defined in section 7.2.2.

Hint: To keep this simple and avoid concerns about precedence, assume all sub-expressions can be written in parentheses.

You can test your answer at the REPL as follows:

```
*ex_7_2> show (the (Expr _) (6 + 3 * 12))
"(6 + (3 * 12))" : String

*ex_7_2> show (the (Expr _) (6 * 3 + 12))
"((6 * 3) + 12)" : String
```

- 2** Implement `Eq` for the `Expr` type. Expressions should be considered equal if their evaluation is equal. So, for example, you should see the following:

```
*Expr> the (Expr _) (2 + 4) == 3 + 3
True : Bool

*Expr> the (Expr _) (2 + 4) == 3 + 4
False : Bool
```

Hint: Start with the implementation header `Eq (Expr ty)` where, and add constraints as you discover you need them.

- 3** Implement `Cast` to allow conversions from `Expr num` to any appropriately constrained type `num`.

Hint: You'll need to evaluate the expression, which should tell you how `num` needs to be constrained.

You can test your answer at the REPL as follows:

```
*ex_7_2> let x : Expr Integer = 6 * 3 + 12 in the Integer (cast x)
30 : Integer
```

7.3 Interfaces parameterized by Type -> Type

In all the interfaces we've seen so far, the parameters have been `Types`. There is, however, no restriction on what the types of the parameters can be. In particular, it's common for interfaces to have parameters of `Type -> Type`. For example, you've already seen the following functions with constrained types:

```
map : Functor f => (a -> b) -> f a -> f b
pure : Applicative f => a -> f a
(>=>) : Monad m => m a -> (a -> m b) -> m b
```

In each case, the parameter `f` or `m` stands for a parameterized type, such as `List` or `IO`. You've seen `map` in the context of `List`, and `pure` and `(>=>)` in the context of `IO`.

In this section, we'll look at how these and other operations are defined generically in the Prelude using interfaces, along with some examples of how to apply them to your own types.

7.3.1 Applying a function across a structure with Functor

In chapter 2 you saw the `map` function, which applies a function to every element in a list:

```
Idris> map (*2) [1,2,3,4]
[2, 4, 6, 8] : List Integer
```

I noted, however, that `map` isn't limited to lists, but rather has a constrained generic type using the `Functor` interface. `Functors` allow you to apply a function uniformly across a generic type. The preceding example applied the "multiply by two" function uniformly across a list of `Integers`.

The following listing shows the definition of the `Functor` interface, which contains `map` as its only method, and its implementation for `List`, as defined in the Prelude.

Listing 7.12 The Functor interface and an implementation for List (defined in the Prelude)

```
interface Functor (f : Type -> Type) where
  map : (func : a -> b) -> f a -> f b
  ↗ Gives the type of f explicitly,
  ↗ because it's not Type

Functor List where
  map func []      = []
  map func (x::xs) = func x :: map func xs
```

So far, the interfaces we've seen have been parameterized by a variable with type `Type`. But the parameter of `Functor` is itself a parameterized type (such as `List`). When the parameter has any type other than `Type`, you need to give the parameter's type explicitly.

It's often useful to provide an implementation of `Functor` for collection data structures. For example, the following listing shows how to define a `Functor` implementation for binary trees, uniformly applying a function across every element that appears at a `Node`.

Listing 7.13 Implementation of Functor for Tree (Tree.idr)

```
data Tree elem = Empty
  | Node (Tree elem) elem (Tree elem)

Functor Tree where
  map func Empty = Empty
  map func (Node left e right)
    = Node (map func left)
      (func e)
      (map func right)
      ↗ Applies func uniformly
      ↗ across the left subtree
      ↗ Applies func to the element at the Node
      ↗ Applies func uniformly across
      ↗ the right subtree
```

The Prelude provides Functor implementations for all types with a single type parameter, where possible, including List, Maybe, and IO. If you import Data.Vect, there's also a Functor implementation for Vect n, as follows.

Listing 7.14 Implementation of Functor for vectors (defined in Data.Vect)

```
Functor (Vect n) where
  map func []      = []
  map func (x :: xs) = func x :: map func xs
```

The n here is an implicit argument and means that this implementation works for vectors of any length.

The parameter in the implementation's header, Vect n, means that you're defining an implementation of Functor for vectors of *any* length. The usual rules for implicit arguments apply, as described in chapter 3: any name beginning with a lowercase letter in a function argument position is treated as an implicit argument. Therefore, n is treated as an implicit argument here.

7.3.2 Reducing a structure using Foldable

If you write down a list of numbers in the form $1 :: 2 :: 3 :: 4 :: []$, you can compute the sum of the numbers by applying the following rules:

- 1 Replace each $::$ with a $+$ (giving $1 + 2 + 3 + 4 + []$).
- 2 Replace the $[]$ with a 0 (giving $1 + 2 + 3 + 4 + 0$).
- 3 Calculate the value of the resulting expression (giving 10).

Or you can compute the product of the numbers by applying the following rules:

- 1 Replace each $::$ with a $*$ (giving $1 * 2 * 3 * 4 * []$).
- 2 Replace the $[]$ with a 1 (giving $1 * 2 * 3 * 4 * 1$).
- 3 Calculate the value of the resulting expression (giving 24).

In general, we're reducing the contents of the list to a single value by replacing the $[]$ with a default, or initial, value, and replacing $::$ with a function of two arguments, which combines each value with the result of reducing the rest of the list. Idris provides two higher-order functions, called *folds* to suggest folding the structure into a single value, to do exactly this:

```
foldr : (elem -> acc -> acc) -> acc -> List elem -> acc
foldl : (acc -> elem -> acc) -> acc -> List elem -> acc
```

The distinction between foldr and foldl is in how the resulting expression is bracketed. In our first example, foldr would calculate the result as $1 + (2 + (3 + (4 + 0)))$, whereas foldl would calculate the result as $((((0 + 1) + 2) + 3) + 4)$. In other words, foldr processes the elements left to right, and foldl processes the elements right to left.

TYPE VARIABLES IN FOLDL AND FOLDR The names of the variables in the types of foldl and foldr suggest their purpose: elem is the element type of the list and acc is the type of the result. The name acc suggests the type of an *accumulating* parameter, in which the eventual result is computed.

In each case, the first argument is the function (or operator) to apply, and the second argument is the initial value. So, you can calculate the two previous examples as follows:

```
Idris> foldr (+) 0 [1,2,3,4]
10 : Integer

Idris> foldr (*) 1 [1,2,3,4]
24 : Integer
```

Or you could use a fold to calculate the total length of the Strings in a List String. For example, the total length of `["One", "Two", "Three"]` should be 11.

Let's write this interactively, using `foldr`, in a file named `Fold.idr`:

- 1 Type, define**—You can begin with a type and a candidate definition, and apply `foldr` to the input list, `xs`, but leave holes for the function and initial value so their types can give you hints as to how to proceed:

```
totalLen : List String -> Nat
totalLen xs = foldr ?sumLength ?initial xs
```

- 2 Type, refine**—The type of `?initial` tells you that the initial value must be a `Nat`:

```
xs : List String
t : Type -> Type
elem : Type
-----
initial : Nat
```

You can initialize it with 0, because the total length of an empty list of strings is 0:

```
totalLen : List String -> Nat
totalLen xs = foldr ?sumLength 0 xs
```

- 3 Type, refine**—The type and context of `?sumLength` tells you that you need to provide a function:

```
xs : List String
t : Type -> Type
elem : Type
-----
sumLength : String -> Nat -> Nat
```

The function you need to provide takes a `String` (standing for the string at a given position in the list) and a `Nat` (standing for the result of folding the rest of the list), and returns a `Nat` (standing for the total length). You can complete the definition as follows:

```
totalLen : List String -> Nat
totalLen xs = foldr (\str, len => length str + len) 0 xs
```

You can test the resulting function at the REPL:

```
*Fold> totalLen ["One", "Two", "Three"]
11 : Nat
```

In the preceding example, I gave a type for `foldr` that was specific to `List`. But if you look at the type at the REPL, you'll see a constrained generic type using a `Foldable` interface:

```
foldr : Foldable t => (elem -> acc -> acc) -> acc -> t elem -> acc
```

An implementation of `Foldable` for a structure explains how to reduce that structure to a single value using an initial value and a function to combine each element with an overall folded structure. The following listing gives the interface definition. There's a default definition for `foldl`, written in terms of `foldr`, so only `foldr` need be implemented.

Listing 7.15 The `Foldable` interface (defined in the Prelude)

```
interface Foldable (t : Type -> Type) where
  foldr : (elem -> acc -> acc) -> acc -> t elem -> acc
  foldl : (acc -> elem -> acc) -> acc -> t elem -> acc
```

Folds a structure, working from left to right

Folds a structure, working from right to left. Implementing `foldl` is optional.

The next listing shows how `Foldable` is implemented in the Prelude for `List`. Note in particular the distinction between the implementations of `foldr` and `foldl`.

Listing 7.16 Implementation of `Foldable` for `List` (defined in the Prelude)

```
Foldable List where
  foldr func acc [] = acc
  foldr func acc (x :: xs) = func x (foldr func acc xs)
```

Applies the function to the first element and the result of folding the tail of the list

```
foldl func acc [] = acc
foldl func acc (x :: xs) = foldl func (func acc x) xs
```

Recursively folds the tail, updating the initial value by applying the function to it and the first element

Because our `Tree` data type is a generic type containing a collection of values, you should be able to provide a `Foldable` implementation:

- 1 Type**—Begin by providing an implementation header and a skeleton definition of `foldr`. You can use the default definition for `foldl`:

```
Foldable Tree where
  foldr func acc tree = ?Foldable_rhs_1
```

- 2 Define, refine**—Case-split on the tree. If the tree is `Empty`, you return the initial value:

```
Foldable Tree where
  foldr func acc Empty = acc
  foldr func acc (Node left e right) = ?Foldable_rhs_3
```

- 3 Refine**—In the `Node` case, because `foldr` works left to right, you begin by folding the left subtree recursively. Figure 7.3 illustrates this with an example, assuming an initial accumulator of 0.




Figure 7.3 Fold the left subtree with the initial accumulator. This gives you $0 + 1 + 2 + 3 = 6$.

In code, it looks like this:

```
Foldable Tree where
  foldr func acc Empty = acc
  foldr func acc (Node left e right)
    = let leftfold = foldr func acc left in
      ?Foldable_rhs_3
```

- 4 *Refine*—Next, take the result of folding the left subtree (`leftfold`) and use it as the initial value when folding the right subtree. Figure 7.4 illustrates this on the same example.




Figure 7.4 Fold the right subtree, initialized with the result of folding the left subtree. This gives you $6 + 5 + 6 + 7 = 24$.

In code, it looks like this:

```
Foldable Tree where
  foldr func acc Empty = acc
  foldr func acc (Node left e right)
    = let leftfold = foldr func acc left
      rightfold = foldr func leftfold right in
        ?Foldable_rhs_3
```

- 5 *Refine*—Finally, apply `func` to the value at the node `e` and the result of folding the right subtree:

```
Foldable Tree where
  foldr func acc Empty = acc
  foldr func acc (Node left e right)
    = let leftfold = foldr func acc left
      rightfold = foldr func leftfold right in
        func e rightfold
```

In the example, this gives you the result $4 + 24 = 28$.

7.3.3 Generic do notation using Monad and Applicative

There are two other interfaces that are useful to know about, and which you've already seen in practice: `Monad` and `Applicative`. In most cases, you're unlikely to need to provide your own implementations, but they do appear throughout the Prelude and base libraries, so it's useful to know their capabilities, and particularly which functions they provide.

In chapter 5, you saw how the `(>>=)` function was used to sequence IO operations:

```
(>>=) : IO a -> (a -> IO b) -> IO b
```

Then, in chapter 6, you saw how the `(>>=)` function was used to sequence `Maybe` computations, abandoning the sequence if any computation returned `Nothing`:

```
(>>=) : Maybe a -> (a -> Maybe b) -> Maybe b
```

Given that these two functions have similar types (directly replacing `IO` with `Maybe`) and similar purposes (sequencing either interactive actions or computations that might fail), you might expect them to be defined in a common interface. The following listing shows the `Monad` interface that provides `(>>=)`, along with a `join` method that combines nested monadic structures.¹

Listing 7.17 The Monad interface (defined in the Prelude)

```
interface Applicative m => Monad (m : Type -> Type) where
  (>>=) : m a -> (a -> m b) -> m b
  join : m (m a) -> m a
```

The Applicative interface supports function application inside a generic type.

Both `(>>=)` and `join` have default definitions, so you can define a `Monad` implementation in terms of either. Here, we'll concentrate on `(>>=)`.

In chapter 6, you saw a definition of `(>>=)` for `Maybe`. In practice, it's defined in the Prelude as follows:

```
Monad Maybe where
  (>>=) Nothing next = Nothing
  (>>=) (Just x) next = next x
```

You've also seen the `pure` function, particularly in the context of `IO` programs to produce a value in an `IO` computation without describing any actions:

```
pure : a -> IO a
```

As with `(>>=)`, `pure` also works in the context of `Maybe`, applying `Just` to its argument:

```
Idris> the (Maybe _) (pure "driven snow")
Just "driven snow" : Maybe String
```

¹ We won't go into detail on `join` here, but it allows you to define a `Monad` implementation for `List` by concatenating lists of lists, among other things.

Again, given that `pure` works in several contexts, you might expect it to be defined in an interface. The following listing shows the definition of the `Applicative` interface, which provides `pure` and a `(<*>)` function that applies a function inside a structure. You'll see an example of `Applicative` in chapter 12.

Listing 7.18 The `Applicative` interface (defined in the Prelude)

```
interface Functor f => Applicative (f : Type -> Type) where
  pure   : a -> f a
  (<*>)  : f (a -> b) -> f a -> f b
```

There are several practical uses of `Monad` and `Applicative`, although as a user of a library, you'll usually just need to know whether there are `Monad` and `Applicative` instances for the provided types, and particularly what the effect of the `(>>=)` operator is.

One interesting implementation of `Monad` in the Prelude is for `List`. The `(>>=)` function for `List` passes *every* value in the input list to the next function in sequence, and combines the results into a new list. I won't go into this in detail in this book, but you can use it to write *nondeterministic* programs.

There's much more to be said on the subject of interfaces, particularly the hierarchy we've looked at briefly in this section covering `Functor`, `Foldable`, `Applicative`, and `Monad`. A deep discussion is beyond the scope of this book, but the interfaces we've discussed in this chapter are the ones you'll encounter most often at first.

Exercises



- 1** Implement `Functor` for `Expr` (defined in section 7.2.2).

You can test your answer at the REPL as follows:

```
*Expr> map (*2) (the (Expr _) (1 + 2 * 3))
Add (Val 2) (Mul (Val 4) (Val 6)) : Expr Integer

*Expr> map show (the (Expr _) (1 + 2 * 3))
Add (Val "1") (Mul (Val "2") (Val "3")) : Expr String
```

- 2** Implement `Eq` and `Foldable` for `Vect`.

Hint: If you import `Data.Vect`, these implementations already exist, so you'll need to define `Vect` by hand to try this exercise.

You can test your answers at the REPL as follows:

```
*ex_7_3> foldr (+) 0 (the (Vect _ _) [1,2,3,4,5])
15 : Integer

*ex_7_3> the (Vect _ _) [1,2,3,4] == [1,2,3,4]
True : Bool

*ex_7_3> the (Vect _ _) [1,2,3,4] == [5,6,7,8]
False : Bool
```

7.4 Summary

- Generic types can be constrained using interfaces, which describe groups of functions that can be applied in a specific context.
- The `Eq` interface provides functions for comparing values for equality and inequality.
- The `Ord` interface provides functions for comparing two values to see which is smaller or larger.
- Implementations of interfaces describe how interfaces can be evaluated in specific contexts.
- Interfaces and implementations can themselves have constraints. Constraints say which interfaces must be implemented for a definition to be valid.
- The Prelude provides several standard interfaces, including `Show` for converting values to a `String`, `Num` for arithmetic operations and numeric literals, and `Cast` for converting between types.
- Interfaces can be parameterized by values of any type. Several in the Prelude are parameterized by values of type `Type -> Type`.
- An implementation of `Functor` for a structure allows a uniform action to be applied to each element in the structure.
- An implementation of `Foldable` allows a structure to be reduced to a single value.
- An implementation of `Monad` allows you to use `do` notation to sequence computations on a structure.



Equality: expressing relationships between data

This chapter covers

- Expressing properties of functions and data
- Checking and guaranteeing equalities between data
- Showing when cases are impossible with the empty type `Void`

You've now seen several ways in which first-class types increase the expressivity of the type system, and the precision of the types we give to functions. You've seen how to use increased precision in types (along with holes) to help write functions, and how to write functions to calculate types. Another way you can use first-class types to increase the precision of your types, and to increase confidence in functions behaving correctly, is to write types specifically to express properties of data and the relationships between data.

In this chapter, we'll look at a simple property, using types to express guarantees that values are equal. You'll also see how to express guarantees that values are *not* equal. Properties such as equality and inequality are sometimes required when

you're defining more-complex functions with dependent types, where the relationship between values might not be immediately obvious to Idris. For example, as you'll see when we define `reverse` on vectors, the input and output vector lengths must be the same, so we'll need to explain to the compiler why the length is preserved.

We'll start by looking at a function we've already used, `exactLength`, and see in some detail how to build it from first principles.

8.1 Guaranteeing equivalence of data with equality types

When you want to compare values for equality, you can use the `==` operator, which returns a value of type `Bool`, given two values in a type `ty` for which there's an implementation of the `Eq` interface:

```
(==) : Eq ty => ty -> ty -> Bool
```

But if you look closely at this type, what does it tell you about the relationships between the inputs (of type `ty`) and the output (of type `Bool`)?

In fact, it tells you nothing at all! Without looking at the specific implementation, you don't know exactly how `==` will behave. Any of the following would be reasonable behavior for an implementation of `==` in the `Eq` interface, at least as far as the type is concerned:

- Always returning `True`
- Always returning `False`
- Returning whether the inputs are *not* equal

It would be surprising to programmers if `==` behaved in any of these ways, but as far as the Idris type checker is concerned, it can make no assumptions other than those given explicitly in the type. If you want to compare values at the type level, therefore, you'll need something more expressive.

In fact, you've already seen an example of a situation where you might want to do this: at the end of chapter 5 you used the `exactLength` function to check whether a `Vect` had a specific length:

```
exactLength : (len : Nat) -> (input : Vect m a) -> Maybe (Vect len a)
```

You used this function to check whether two vectors entered by a user had exactly the same length. Given a specific length, `len`, and a vector of length, `m`, it returns the following:

- `Nothing`, if `len` is not the same as `m`
- `Just input`, if `len` is the same as `m`

In this section, we'll look at how to implement `exactLength` by representing equality as a type, and you'll see in more detail why `==` isn't sufficient. We'll start by trying to implement it using `==` and see where we run into limitations.

8.1.1 Implementing exactLength, first attempt

Rather than importing Data.Vect, which defines exactLength, we'll begin by defining Vect by hand and giving a type and a skeleton definition of exactLength. The following listing shows our starting point, in a file named ExactLength.idr.

Listing 8.1 A definition of Vect and the type and skeleton definition of exactLength (ExactLength.idr)

```
data Vect : Nat -> Type -> Type where
  Nil : Vect Z a
  (::) : a -> Vect k a -> Vect (S k) a

exactLength : (len : Nat) -> (input : Vect m a) -> Maybe (Vect len a)
exactLength len input = ?exactLength_rhs
```

You should expect to be able to implement exactLength by comparing the length of input with the desired length, len. If they're equal, then input has length len, so its type should be considered equivalent to Vect len a, and you can return it directly. Otherwise, you return Nothing.

As a first attempt, you can try the following steps:

- 1 *Define*—Because m doesn't appear on the left side of the skeleton definition, you'll need to bring it into scope explicitly in order to compare len and m. You can do this by updating the definition to the following:

```
exactLength {m} len input = ?exactLength_rhs
```

Recall from chapter 3 that an implicit argument such as m can be used in the definition if it's written inside braces on the left side.

- 2 *Define*—You can now compare m and len for equality using == and inspect the result in a case statement:

```
exactLength : (len : Nat) -> (input : Vect m a) -> Maybe (Vect len a)
exactLength {m} len input = case m == len of
  False => ?exactLength_rhs_1
  True => ?exactLength_rhs_2
```

- 3 *Refine*—For ?exactLength_rhs_1, the lengths are different, so you return Nothing:

```
exactLength {m} len input = case m == len of
  False => Nothing
  True => ?exactLength_rhs_2
```

- 4 *Refine*—For ?exactLength_rhs_2, you'd like to return Just input because the lengths are equal. You can start with Just:

```
exactLength : (len : Nat) -> (input : Vect m a) -> Maybe (Vect len a)
exactLength {m} len input = case m == len of
  False => Nothing
  True => Just ?exactLength_rhs_2
```

- 5 *Type*—Then, you can take a look at the types of the variables in scope, and the type required to fill in the `?exactLength_rhs_2` hole:

```
a : Type
m : Nat
len : Nat
input : Vect m a
-----
exactLength_rhs_2 : Vect len a
```

Even though you've checked that `m` and `len` are equal using `==`, you can't fill in the hole with `input` because it has type `Vect m a`, and the required type is `Vect len a`. The problem, as when defining `zipInputs` at the end of chapter 5, is that the type of `==` isn't informative enough to guarantee that `m` and `len` are equal, even if it returns `True`.

You will, therefore, need to consider alternative approaches to implementing `exactLength`, using a more informative type when comparing `m` and `len`.

A variable's type tells Idris what possible values a variable can have, but it says nothing about where the value has come from. If a variable has type `Bool`, Idris knows that it can have either of the values `True` or `False`, but nothing about the computation that has produced the value. There are lots of possible computations that could produce a result of type `Bool` other than testing for equality. Furthermore, equality is defined by the `Eq` interface, and there are no guarantees in the type about how that interface might be implemented.

Instead, you'll need to create a more precise type for the equality test, where the type guarantees that a comparison between two inputs can only be successful if the inputs really are identical. In the rest of this section, you'll see how to do this from first principles, and how to use the new equality type to implement `exactLength`.

8.1.2 Expressing equality of Nats as a type

Listing 8.2 shows a dependent type, `EqNat`. It has two numbers as arguments, `num1` and `num2`. If you have a value of type `EqNat num1 num2`, you know that `num1` and `num2` must be the same number because the only constructor, `Same`, can only build something with a type of the form `EqNat num num`, where the two arguments are the same.

Listing 8.2 Representing equal Nats as a type (`EqNat.idr`)

```
data EqNat : (num1 : Nat) -> (num2 : Nat) -> Type where
  Same : (num : Nat) -> EqNat num num
```

The only constructor, `Same`, constructs evidence that `num` is equal to itself.

With dependent types, you can use types such as `EqNat` to express additional information about other data, in this case expressing that two `Nat`s are guaranteed to be equal. This is a powerful concept, as you'll soon see, and can take some time to fully appreciate. We'll therefore look at representing and checking equalities in some depth.

First, to see how EqNat works, let's try a few examples at the REPL:

```
*EqNat> Same 4
Same 4 : EqNat 4 4

*EqNat> Same 5
Same 5 : EqNat 5 5

*EqNat> the (EqNat 3 3) (Same _)
Same 3 : EqNat 3 3

*EqNat> Same (2 + 2)
Same 4 : EqNat 4 4
```

Whatever you try, the argument in the type is repeated. It is, nevertheless, perfectly valid to write a *type* with unequal arguments:

```
*EqNat> EqNat 3 4
EqNat 3 4 : Type
```

But if you try to construct a *value* with this type, you can't succeed and will always get a type error:

```
*EqNat> the (EqNat 3 4) (Same _)
(input):1:5:When checking argument value to function Prelude.Basics.the:
    Type mismatch between
        EqNat num num (Type of Same num)
    and
        EqNat 3 4 (Expected type)

    Specifically:
        Type mismatch between
            0
        and
            1
```

This error message indicates that 3 and 4 need to be the same, because they both need to be instantiated for `num` in the type of `Same`. The type `EqNat 3 4` is an *empty type*, meaning that there are *no* values of that type.

SPECIFICITY IN ERROR MESSAGES You'll notice that in error messages, Idris often reports the error in two ways. The first part gives an overall type mismatch. This can become quite large, however, so Idris also reports the specific part of the expression that didn't match. Here, it reports a mismatch between 0 and 1 because of the way `Nat` is defined in terms of `Z` and `S`. The overall mismatch is between `S (S (S Z))` and `S (S (S (S Z)))`, for which the *specific* difference is between `Z` and `S Z`.

8.1.3 Testing for equality of Nats

We'll use `EqNat` to help implement `exactLength`. Because `EqNat num1 num2` is essentially a proof that `num1` must be equal to `num2`, we'll write a function that checks whether the input lengths are equal, and, if they are, express that equality as an instance of `EqNat`.

We'll begin by writing a `checkEqNat` function that either returns a proof that its inputs are the same, in the form of `EqNat`, or `Nothing` if the inputs are different. It has the following type:

```
checkEqNat : (num1 : Nat) -> (num2 : Nat) -> Maybe (EqNat num1 num2)
```

Once implemented, it will behave as in the following examples:

```
*EqNat> checkEqNat 5 5
Just (Same 5) : Maybe (EqNat 5 5)

*EqNat> checkEqNat 1 2
Nothing : Maybe (EqNat 1 2)
```

Because we can only have a value of type `EqNat num1 num2` for a specific `num1` and `num2` if `num1` and `num2` are identical, the type of `checkEqNat` guarantees that if it's successful (that is, returns a value of the form `Just p`), then its inputs really must be equal.

You can implement the function as follows:

- 1 Type**—Begin with a type and a skeleton definition:

```
checkEqNat : (num1 : Nat) -> (num2 : Nat) -> Maybe (EqNat num1 num2)
checkEqNat num1 num2 = ?checkEqNat_rhs
```

- 2 Define**—You can define the function by case splitting on both `Nat` inputs. First, case-split on `num1`:

```
checkEqNat : (num1 : Nat) -> (num2 : Nat) -> Maybe (EqNat num1 num2)
checkEqNat Z num2 = ?checkEqNat_rhs_1
checkEqNat (S k) num2 = ?checkEqNat_rhs_2
```

- 3 Define**—Then, case-split on `num2` in both cases:

```
checkEqNat : (num1 : Nat) -> (num2 : Nat) -> Maybe (EqNat num1 num2)
checkEqNat Z Z = ?checkEqNat_rhs_3
checkEqNat Z (S k) = ?checkEqNat_rhs_4
checkEqNat (S k) Z = ?checkEqNat_rhs_1
checkEqNat (S k) (S j) = ?checkEqNat_rhs_5
```

- 4 Refine**—If you take a look at the type of the `?checkEqNat_rhs_3` hole, you'll see that you need to provide evidence that `0` is the same as itself, so you can fill this in with `Just (Same 0)`. For `?checkEqNat_rhs_4` and `?checkEqNat_rhs_1`, you can't provide any evidence that `0` is the same as a non-zero number, so return `Nothing`. You now have this:

```
checkEqNat : (num1 : Nat) -> (num2 : Nat) -> Maybe (EqNat num1 num2)
checkEqNat Z Z = Just (Same 0)
checkEqNat Z (S k) = Nothing
checkEqNat (S k) Z = Nothing
checkEqNat (S k) (S j) = ?checkEqNat_rhs_5
```

- 5 Refine**—The result for `?checkEqNat_rhs_5` depends on whether `k` is equal to `j`, which you can determine by making a recursive call and then case splitting on the result:

```
checkEqNat (S k) (S j) = case checkEqNat k j of
                           case_val => ?checkEqNat_rhs_5
```

- 6 Define**—Case splitting on `case_val` leads to cases for when `k` and `j` aren't equal (`Nothing`) and when they are equal (`Just x`). You can rename the `Just x` produced by the case split to something more informative, and fill in the case where the recursive call produces `Nothing`:

```
checkEqNat (S k) (S j) = case checkEqNat k j of
                           Nothing => Nothing
                           Just eq => ?checkEqNat_rhs_2
```

- 7 Type**—Looking at the type of `checkEqNat_rhs_2` gives you some information about how to proceed:

```
k : Nat
j : Nat
eq : EqNat k j
-----
checkEqNat_rhs_2 : Maybe (EqNat (S k) (S j))
```

The type of `eq` is `EqNat k j`, and you're looking for something of type `Maybe (EqNat (S k) (S j))`.

- 8 Refine**—If `k` and `j` are equal, you know that `S k` and `S j` must be equal, so you can return a value constructed with `Just`:

```
checkEqNat (S k) (S j) = case checkEqNat k j of
                           Nothing => Nothing
                           Just eq => Just ?checkEqNat_rhs_2
```

- 9 Refine**—To complete the definition, rename the remaining hole, and lift it to a top-level definition, which you'll implement in the next section.

```
sameS : (eq : EqNat k j) -> EqNat (S k) (S j)

checkEqNat : (num1 : Nat) -> (num2 : Nat) -> Maybe (EqNat num1 num2)
checkEqNat Z Z = Just (Same Z)
checkEqNat Z (S k) = Nothing
checkEqNat (S k) Z = Nothing
checkEqNat (S k) (S j) = case checkEqNat k j of
                           Nothing => Nothing
                           Just eq => Just (sameS eq)
```

You can see from the type of the lifted definition, `sameS`, that it's a function that takes evidence that `k` and `j` are equal, and returns evidence that `S k` and `S j` are equal. By expressing equality between numbers as a dependent data type, `EqNat`, you're able to write functions like `sameS` that take an instance of `EqNat` as an input and manipulate them, essentially deducing additional information about equalities.

8.1.4 Functions as proofs: manipulating equalities

It's impossible to create an instance of `EqNat k j` when `k` and `j` are different, which means you can think of `sameS` as a proof that if `k` and `j` are equal, `S k` and `S j` are also equal.

Let's now try implementing `sameS`. For clarity, we'll make the `Nat` arguments explicit:

```
sameS : (k : Nat) -> (j : Nat) -> (eq : EqNat k j) -> EqNat (S k) (S j)
sameS k j eq = ?sameS_rhs
```

You can implement `sameS` with the following steps:

- 1 Define**—Begin by creating a skeleton definition:

```
sameS : (k : Nat) -> (j : Nat) -> (eq : EqNat k j) -> EqNat (S k) (S j)
sameS k j eq = ?sameS_rhs
```

- 2 Define**—Next, in Atom, ask to case-split on `eq`:

```
sameS : (k : Nat) -> (j : Nat) -> (eq : EqNat k j) -> EqNat (S k) (S j)
sameS k k (Same k) = ?sameS_rhs_1
```

Notice that `k` appears three times on the left side of this definition! Because you've expressed a relationship between `k` and `j` using `eq` in the type of `sameS`, and you've case-split on `eq`, Idris has noticed that both `Nat` inputs must be the same. Not only that, if you try to give a different value, it will report an error. If, instead, you write this,

```
sameS k j (Same k) = ?sameS_rhs_1
```

then Idris will report the following:

```
EqNat.idr:15:7:When checking left hand side of sameS:
Type mismatch between
      j (Inferred value)
and
      k (Given value)
```

In other words, because the type states that both `Nat` inputs must be the same, Idris isn't happy that they're different. So, revert to the left side that Idris generated after the case split on `eq`:

```
sameS : (k : Nat) -> (j : Nat) -> (eq : EqNat k j) -> EqNat (S k) (S j)
sameS k k (Same k) = ?sameS_rhs_1
```

- 3 Type**—If you check the type of `?sameS_rhs_1`, you'll see that you need to provide evidence that `S k` is the same as itself:

```
k : Nat
-----
sameS_rhs_1 : EqNat (S k) (S k)
```

- 4 Refine**—You can therefore complete the definition as follows:

```
sameS : (k : Nat) -> (j : Nat) -> (eq : EqNat k j) -> EqNat (S k) (S j)
sameS k k (Same k) = Same (S k)
```

PROOFS IN IDRIS In principle, you can state and try to prove complex properties of any function in Idris. For example, you could write a function whose type states that reversing a list twice yields the original list. In practice, however, you’ll rarely need to manipulate equalities much more complex than the implementation of `sames`. Nevertheless, you’ll see a bit more about manipulating equalities, and where they arise when defining functions, in section 8.2.

Listing 8.3 gives the complete definition of `checkEqNat`, using the version of `sames` with explicit arguments. You could also write this function without using `sames`, instead using a case split on `eq`. You could also use `do` notation, as described at the end of chapter 6, to make the definition more concise. As exercises, try reimplementing it in each of these ways.

Listing 8.3 Testing for equality of Nats with EqNat (EqNat.idr)

```
sames : (k : Nat) -> (j : Nat) -> (eq : EqNat k j) -> EqNat (S k) (S j)
sames k k (Same k) = Same (S k)
```

```
checkEqNat : (num1 : Nat) -> (num2 : Nat) -> Maybe (EqNat num1 num2)
checkEqNat Z Z = Just (Same 0)
checkEqNat Z (S k) = Nothing
checkEqNat (S k) Z = Nothing
checkEqNat (S k) (S j) = case checkEqNat k j of
```

Z and S are the canonical
constructors for Nat and
can never be equal.

```
    Nothing => Nothing
    Just eq => Just (sames _ _ eq)
```

S k and S j are equal if k and j are equal. You
use `sames` to provide evidence for this.

Unlike `==`, `checkEqNat` expresses the relationship between its inputs and its output precisely in its type. Using this, we can make another attempt to implement `exactLength`.

8.1.5 Implementing `exactLength`, second attempt

Earlier, we reached the following point in implementing `exactLength`, before establishing that a Boolean comparison wasn’t sufficient:

```
exactLength : (len : Nat) -> (input : Vect m a) -> Maybe (Vect len a)
exactLength {m} len input = case m == len of
    False => ?exactLength_rhs_1
    True => ?exactLength_rhs_2
```

Instead of using the Boolean comparison operator `==` to compare `m` and `len`, you can try using `checkEqNat m len`. This will return a value of type `Maybe (EqNat m len)`, so if `m` and `len` are equal, you’ll have some additional information in the type that states the meaning of the result of the comparison precisely. In this second attempt, you can implement the function as follows:

- 1 *Define*—As before, begin with a type and a skeleton definition, bringing `m` into scope on the left side:

```
exactLength : (len : Nat) -> (input : Vect m a) -> Maybe (Vect len a)
exactLength {m} len input = ?exactLength_rhs
```

- 2 Define**—Instead of defining the function by case splitting on the result of `m == len`, case-split on the result of `checkEqNat`:

```
exactLength : (len : Nat) -> (input : Vect m a) -> Maybe (Vect len a)
exactLength {m} len input = case checkEqNat m len of
    Nothing => ?exactLength_rhs_1
    Just eq_nat => ?exactLength_rhs_2
```

- 3 Refine**—If `checkEqNat` returns `Nothing`, the inputs were different, so the input vector is the wrong length:

```
exactLength : (len : Nat) -> (input : Vect m a) -> Maybe (Vect len a)
exactLength {m} len input = case checkEqNat m len of
    Nothing => Nothing
    Just eq_nat => Just ?exactLength_rhs_2
```

- 4 Type**—If `checkEqNat` returns `Just eq_nat`, the lengths are equal, and `eq_nat` provides evidence that they're equal. For `?exactLength_rhs_2`, you have this:

```
m : Nat
len : Nat
eq_nat : EqNat m len
a : Type
input : Vect m a
-----
exactLength_rhs_2 : Maybe (Vect len a)
```

- 5 Define**—As before, you still need to provide a result of type `Maybe (Vect len a)`, and all you have available is `input` of type `Vect m a`. But you also have `eq_nat`, which provides evidence that `m` and `len` are equal. If you case-split on `eq_nat`, you'll get the following:

```
exactLength : (len : Nat) -> (input : Vect m a) -> Maybe (Vect len a)
exactLength {m} len input = case checkEqNat m len of
    Nothing => Nothing
    Just (Same len) => ?exactLength_rhs_1
```

Then, when inspecting the type of the new hole, `?exactLength_rhs_1`, you'll see this:

```
m : Nat
a : Type
input : Vect len a
-----
exactLength_rhs_1 : Maybe (Vect len a)
```

Because `eq_nat` can only take the form `Same len`, and the type of `Same len` forces `m` to be identical to `len`, Idris has refined the required type to be `Maybe (Vect len a)`.

- 6 Refine**—From here, it's easy to complete the definition:

```
exactLength : (len : Nat) -> (input : Vect m a) -> Maybe (Vect len a)
exactLength {m} len input = case checkEqNat m len of
    Nothing => Nothing
    Just (Same len) => Just input
```

This isn't exactly the definition used in the Prelude, however. Instead, the Prelude uses a generic equality type, built into Idris.

8.1.6 Equality in general: the `=` type

Rather than defining a specific equality type and function for every possible type you'll need to compare, such as `Nat` with `EqNat` and `checkEqNat` here, Idris provides a generic equality type. This is built into Idris's syntax, so you can't define this yourself (because `=` is a reserved symbol), but conceptually it would be defined as follows.

Listing 8.4 Conceptual definition of a generic equality type

```
data (=) : a -> b -> Type where
    Refl : x = x
```

Like Same, but x is an implicit argument

= takes two arguments: one of some generic type a, and the other of some generic type b.

The name `Refl` is short for *reflexive*, a mathematical term that (informally) means that a value is equal to itself. As with `EqNat`, you can try some examples at the REPL:

```
Idris> the ("Hello" = "Hello") Refl
Refl : "Hello" = "Hello"

Idris> the (True = True) Refl
Refl : True = True
```

If you give more-complex expressions as part of the type, Idris will evaluate them. For example, if you give the expression $2 + 2$ as part of a type, $2 + 2 = 4$, Idris will evaluate it:

```
Idris> the (2 + 2 = 4) Refl
Refl : 4 = 4
```

As before, if you try to use `Refl` to build an instance of an equality type using two unequal values, you'll get an error:

```
Idris> the (True = False) Refl
(input):1:5:When checking argument value to function Prelude.Basics.the:
    Type mismatch between
        x = x (Type of Refl)
    and
        True = False (Expected type)

    Specifically:
        Type mismatch between
            True
        and
            False
```

Using the built-in equality type, rather than `EqNat`, you can define `checkEqNat` as follows.

**Listing 8.5 Checking equality between Nats using the generic equality type
(CheckEqMaybe.idr)**

```
checkEqNat : (num1 : Nat) -> (num2 : Nat) -> Maybe (num1 = num2)
checkEqNat Z Z = Just Refl
checkEqNat Z (S k) = Nothing
checkEqNat (S k) Z = Nothing
checkEqNat (S k) (S j) = case checkEqNat k j of
    Nothing => Nothing
    Just prf => Just (cong prf)      ← cong is a generic
                                         version of sameS.
```

Congruence

In listing 8.5, you used `cong` to convert a value of type $k = j$ to a value of type $S k = S j$. It has the following type:

```
cong : {func : a -> b} -> x = y -> func x = func y
```

In other words, given some (implicit) function, `func`, if you have two values that are equal, then applying `func` to those values gives an equal result. This does the same job as `sameS`, using the generic equality type.

Exercises



- Implement the following function, which states that if you add the same value onto the front of equal lists, the resulting lists are also equal:

```
same_cons : {xs : List a} -> {ys : List a} ->
            xs = ys -> x :: xs = x :: ys
```

Because this function represents an equality proof, it's sufficient to know that your definition type-checks and is total:

```
*ex_8_1> :total same_cons
Main.same_cons is Total
```

- Implement the following function, which states that if two values, x and y , are equal, and two lists, xs and ys , are equal, then the two lists $x :: xs$ and $y :: ys$ must also be equal:

```
same_lists : {xs : List a} -> {ys : List a} ->
             x = y -> xs = ys -> x :: xs = y :: ys
```

Again, it's sufficient to know that your definition type-checks.

- Define a type, `ThreeEq`, that expresses that *three* values must be equal.

Hint: `ThreeEq` should have the type $a \rightarrow b \rightarrow c \rightarrow \text{Type}$.

- Implement the following function, which uses `ThreeEq`:

```
allSameS : (x, y, z : Nat) -> ThreeEq x y z -> ThreeEq (S x) (S y) (S z)
```

What does this type mean?

8.2 Equality in practice: types and reasoning

Equality proofs, and functions that manipulate them, can be useful when defining functions with dependent types. You saw a small example of this when implementing checkEqNat, where you wrote a sameS function to show that adding one to equal numbers results in equal numbers.

Reasoning about equality often becomes necessary when writing functions on types that are indexed by numbers. For example, when manipulating vectors, you might need to prove the equivalence between two expressions with natural numbers that appear in the types of the vectors. To see how this can happen and how to deal with it, we'll look at how to implement a function that reverses a Vect.

8.2.1 Reversing a vector

If you import Data.Vect, you'll get access to a function that reverses a vector, with the following type:

```
reverse : Vect n elem -> Vect n elem
```

The type states that reversing a vector preserves the length of the input vector. You'd expect it to be fairly straightforward to implement this function using the following rules:

- If the input vector is empty, return an empty vector.
- If the input vector isn't empty and has a head x and a tail xs , reverse xs and append x .

Let's see what happens if we try to implement our own version, myReverse, in the file ReverseVec.idr:

- 1 *Type, define*—Begin with a type and a skeleton definition, and case-split on the input:

```
myReverse : Vect n elem -> Vect n elem
myReverse [] = ?myReverse_rhs_1
myReverse (x :: xs) = ?myReverse_rhs_2
```

- 2 *Refine*—For $?myReverse_rhs_1$, return an empty vector:

```
myReverse : Vect n elem -> Vect n elem
myReverse [] = []
myReverse (x :: xs) = ?myReverse_rhs_2
```

- 3 *Refine failure*—For $?myReverse_rhs_2$, you'd like to be able to reverse xs and append x , as follows, but unfortunately this fails:

```
myReverse : Vect n elem -> Vect n elem
myReverse [] = []
myReverse (x :: xs) = myReverse xs ++ [x]
```

This line has
a type error

The error message tells you that Idris has found a mismatch between the type of the value you've given, Vect (k + 1), and the expected type, Vect (S k):

```
ReverseVec.idr:6:12:
When checking right hand side of myReverse with expected type
  Vect (S k) elem

Type mismatch between
  Vect (k + 1) elem (Type of myReverse xs ++ [x])
and
  Vect (S k) elem (Expected type)
```

This seems surprising, because our knowledge of the behavior of addition suggests that $S k$ and $k + 1$ will always evaluate to the same result, whatever the value of k .

We'll postpone the completion of this definition and take a closer look at how type checking works in Idris in order to understand what has gone wrong and how we might correct it.

8.2.2 Type checking and evaluation

When Idris type-checks an expression, it will look at the *expected* type of the expression, and check that the type of the *given* expression matches it, after evaluating both. The following code fragment will type-check:

```
test1 : Vect 4 Int
test1 = [1, 2, 3, 4]

test2: Vect (2 + 2) Int
test2 = test1
```

Even though `test1` and `test2` have different expressions in their types, these expressions evaluate to the same result, so you can define `test2` to return `test1`.

You can see at the REPL how the Idris type checker evaluates types containing type variables by using an anonymous function (explained in chapter 2) to introduce variables with unknown values. Consider this example:

```
*ReverseVec> \k, elem => Vect (1 + k) elem
\k => \elem => Vect (S k) elem : Nat -> Type -> Type
```

Here, Idris has evaluated $1 + k$ in the type to $S k$. But if you try swapping the arguments to $+$, you'll get a different result:

```
*ReverseVec> \k, elem => Vect (k + 1) elem
\k => \elem => Vect (plus k 1) elem : Nat -> Type -> Type
```

Here, Idris has evaluated $k + 1$ to $\text{plus } k \ 1$, where plus is the Prelude function that implements addition on Nat . To see the reason for the difference, you need to look at the definition of plus . You can do this using the REPL command `:printdef`, which prints the pattern-matching definition of the function given as its argument:

```
*ReverseVec> :printdef plus
plus : Nat -> Nat -> Nat
plus 0 right = right
plus (S left) right = S (plus left right)
```

Because plus is defined by pattern-matching on its first argument, Idris can't evaluate plus k 1 any further. To do so, it would need to know what form k takes, but at the moment, neither of the clauses for plus matches.

Returning to our problem with defining myReverse, let's take a look at the current state:

```
myReverse : Vect n elem -> Vect n elem
myReverse [] = []
myReverse (x :: xs) = myReverse xs ++ [x]
```

If you rewrite the definition using let to build each component of the result, and leave a hole for the return value, you can see in more detail what problem you have to solve:

```
myReverse : Vect n elem -> Vect n elem
myReverse [] = []
myReverse (x :: xs) = let rev_xs = myReverse xs
                      result = rev_xs ++ [x] in
                      ?myReverse_rhs_2
```

Checking the type of ?myReverse_rhs_2 shows the type of each component and the required type of the result:

```
elem : Type
x : elem
k : Nat
xs : Vect k elem
rev_xs : Vect k elem
result : Vect (plus k 1) elem
-----
myReverse_rhs_2 : Vect (S k) elem
```

For the intended result, you have an expression with type Vect (k + 1) elem, but you need an expression with type Vect (S k) elem.

You know, from trying the evaluation at the REPL previously, that Vect (1 + k) elem would evaluate to the type you need, so you need to be able to explain to Idris that 1 + k is equal to k + 1; or, when evaluated, that S k is equal to plus k 1. The Idris library provides a function that can help you:

```
plusCommutative : (left : Nat) -> (right : Nat) -> left + right = right + left
```

If you check the type of plusCommutative at the REPL with values 1 and k for left and right, respectively, you'll see exactly the equality you need:

```
*ReverseVec> :t \k => plusCommutative 1 k
\k => plusCommutative 1 k : (k : Nat) -> S k = plus k 1
```

You can think of this expression's type as a “rewrite rule” that allows you to replace one value with another. If you can find a way to apply this rule to rewrite the expected type to Vect (plus k 1) a, you'll be able to complete the definition.

8.2.3 The rewrite construct: rewriting a type using equality

We'll resume our definition of `myReverse` at the following point, using `let` to name the result we'd like to return:

```
myReverse : Vect n elem -> Vect n elem
myReverse [] = []
myReverse (x :: xs) = let result = myReverse xs ++ [x] in
                      ?myReverse_rhs_2
```

At this stage, these are the types:

```
elem : Type
x : elem
k : Nat
xs : Vect k elem
result : Vect (plus k 1) elem
-----
myReverse_rhs_2 : Vect (S k) elem
```

In order to complete the definition, you can use the information given by `plusCommutative 1 k` to rewrite the type of `?myReverse_rhs_2`. You'll want to replace any `S k` in the type of `?myReverse_rhs_2` with `plus k 1`, so that you can return `result`. Idris provides a syntax for using equality proofs to rewrite types, illustrated in figure 8.1.




Figure 8.1 Rewriting a type using an equality proof

You can therefore implement `myReverse` using the `rewrite` construct to update the type of `?myReverse_rhs_2`, taking the following steps:

- 1 *Define*—Before rewriting the type using `plusCommutative 1 k`, you'll need to bring `k` into scope, and `k` arises from pattern matching on the length of the input vector:

```
myReverse : Vect n elem -> Vect n elem
myReverse [] = []
myReverse {n = S k} (x :: xs)
  = let result = myReverse xs ++ [x] in
    ?myReverse_rhs_2
```

- 2 *Refine, type*—Now, you can apply the rewriting rule:

```
myReverse : Vect n elem -> Vect n elem
myReverse [] = []
myReverse {n = S k} (x :: xs)
```

```
= let result = myReverse xs ++ [x] in
  rewrite plusCommutative 1 k in ?myReverse_rhs_2
```

If you look at the resulting type for `?myReverse_rhs_2`, you can see the effect the rewrite has had:

```
elem : Type
k : Nat
x : elem
xs : Vect k elem
result : Vect (plus k 1) elem
_rewrite_rule : plus k 1 = S k
-----
myReverse_rhs_2 : Vect (plus k 1) elem
```

- 3 Refine**—Finally, you can complete the definition using expression search:

```
myReverse : Vect n elem -> Vect n elem
myReverse [] = []
myReverse {n = S k} (x :: xs)
= let result = myReverse xs ++ [x] in
  rewrite plusCommutative 1 k in result
```

Using `rewrite`, you've replaced an expression in the type (`S k`) with an equivalent expression (`plus k 1`), which allows you to use `result`. But although this has allowed you to write the function, this definition still leaves something to be desired: the part of the definition that computes the result (`myReverse xs ++ [x]`) has become rather lost in the details of the proof. You can improve this by delegating the details of the proof, using a hole.

8.2.4 Delegating proofs and rewriting to holes

Instead of applying the `rewrite` inside the definition of `myReverse`, you can use a hole along with interactive editing to generate a helper function that contains the details of the proof. Let's start again, with our initial (failing) definition of `myReverse`:

```
myReverse : Vect n elem -> Vect n elem
myReverse [] = []
myReverse (x :: xs) = myReverse xs ++ [x]
```

You can correct this definition using the following steps:

- 1 Refine, type**—Add a hole to the right side, which takes the initial attempt as an argument:

```
myReverse : Vect n elem -> Vect n elem
myReverse [] = []
myReverse (x :: xs) = ?reverseProof (myReverse xs ++ [x])
```

If you check the type of `?reverseProof`, you'll see exactly how you need to rewrite the type for this definition to be accepted:

```
elem : Type
x : elem
k : Nat
```

```

xs : Vect k elem
-----
reverseProof : Vect (plus k 1) elem -> Vect (S k) elem

2 Type—Using Ctrl-Alt-L in Atom, lift ?reverseProof to a top-level function:

reverseProof : (x : elem) -> (xs : Vect k elem) ->
               Vect (k + 1) elem -> Vect (S k) elem

myReverse : Vect n elem -> Vect n elem
myReverse [] = []
myReverse (x :: xs) = reverseProof x xs (myReverse xs ++ [x])

3 Define—Define reverseProof as follows, using the same application of rewrite
as in your previous definition of myReverse:

reverseProof : (x : elem) -> (xs : Vect k elem) ->
               Vect (k + 1) elem -> Vect (S k) elem
reverseProof {k} x xs result = rewrite plusCommutative 1 k in result

4 Refine—Finally, because you don't use x or xs in reverseProof, and because
reverseProof will only be used by myReverse, you can tidy up the definition as
follows:

myReverse : Vect n elem -> Vect n elem
myReverse [] = []
myReverse (x :: xs) = reverseProof (myReverse xs ++ [x])
where
  reverseProof : Vect (k + 1) elem -> Vect (S k) elem
  reverseProof {k} result = rewrite plusCommutative 1 k in result

```

By introducing the hole ?reverseProof, you've been able to keep the relevant computation part of myReverse separate from the details of the proof.

8.2.5 Appending vectors, revisited

You can often avoid the need for rewriting in types by taking care in how you write function types. For example, in chapter 4, you saw how to define a function that appends vectors with the following type:

```
append : Vect n elem -> Vect m elem -> Vect (n + m) elem
```

The order of arguments to the + operator in the return type turns out to be important because of the definition of +. If you begin to implement this (in the file AppendVec.idr) by creating a skeleton definition and then case splitting on the first argument, you'll reach the following state:

```

append : Vect n elem -> Vect m elem -> Vect (n + m) elem
append [] ys = ?append_rhs_1
append (x :: xs) ys = ?append_rhs_2

```

Then, if you check the type of ?append_rhs_1, you'll see the following:

```

elem : Type
m : Nat
ys : Vect m elem

```

```
-----  
append_rhs_1 : Vect m elem
```

In this case, the first argument, `[]`, has type `Vect 0 elem`, and the second argument, `ys`, has type `Vect m elem`. According to the return type of `append`, `?append_rhs_1` should have type `Vect (0 + m) elem`, which reduces to `Vect m elem` by the definition of `+`.

Take a look at what happens if you swap the arguments `n` and `m`, as follows:

```
append : Vect n elem -> Vect m elem -> Vect (m + n) elem  
append [] ys = ?append_rhs_1  
append (x :: xs) ys = ?append_rhs_2
```

You now see a different type for `?append_rhs_1`:

```
elem : Type  
m : Nat  
ys : Vect m elem  
-----  
append_rhs_1 : Vect (plus m 0) elem
```

As in the definition of `myReverse`, Idris can't reduce `plus m 0` any further, because `plus` is defined by pattern-matching on its first argument, and the form of `m` is unknown. Because of this, you can't simply return `ys` for this case, and you'll need to rewrite the type of `append_rhs_1`. You'll have a similar problem for `append_rhs_2`; the following listing shows the definition of `append`, with holes in place of the necessary rewrites.

Listing 8.6 Implementing `append` on vectors with the arguments to `+` swapped in the return type (AppendVec.idr)

`?append_nil` stands for a proof that it's valid to return `ys` here.

```
append : Vect n elem -> Vect m elem -> Vect (m + n) elem  
append [] ys = ?append_nil ys  
append (x :: xs) ys = ?append_xs (x :: append xs ys) ←
```

`?append_xs` stands for a proof that it's valid to return `x :: append xs ys` here.

The next listing shows a completed definition of `append`, with definitions of `append_nil` and `append_xs` that rewrite the types for each case.

Listing 8.7 Completing `append` on vectors by adding rewriting proofs for `append_nil` and `append_xs` (AppendVec.idr)

```
append_nil : Vect m elem -> Vect (plus m 0) elem  
append_nil {m} xs = rewrite plusZeroRightNeutral m in xs ←
```

`sym` is a function that reverses the direction of a rewrite.

```
append_xs : Vect (S (m + k)) elem -> Vect (plus m (S k)) elem  
append_xs {m} {k} xs = rewrite sym  
                      (plusSuccRightSucc m k) in xs ←
```

```
append : Vect n elem -> Vect m elem -> Vect (m + n) elem  
append [] ys = append_nil ys  
append (x :: xs) ys = append_xs (x :: append xs ys)
```

These rewrites use several definitions from the Prelude. Two of them assist with rewriting expressions using Nat:

```
plusZeroRightNeutral : (left : Nat) -> left + 0 = left
plusSuccRightSucc : (left : Nat) -> (right : Nat) ->
                     S (left + right) = left + S right
```

The third, sym, allows you to apply a rewrite rule in reverse:

```
sym : left = right -> right = left
```

Essentially, plusZeroRightNeutral and plusSuccRightSucc together explain that the behavior of plus is identical if the arguments are given in the opposite order.

Exercises



- Using plusZeroRightNeutral and plusSuccRightSucc, write your own version of plusCommutes:

```
myPlusCommutes : (n : Nat) -> (m : Nat) -> n + m = m + n
```

Hint: Write this by case splitting on n. In the case of S k, you can rewrite with a recursive call to myPlusCommutes k m, and rewrites can be nested.

- The implementation of myReverse you wrote earlier is inefficient because it needs to traverse the entire vector to append a single element on every iteration. You can write a better definition as follows, using a helper function, reverse', that takes an accumulating argument to build the reversed list:

```
myReverse : Vect n a -> Vect n a
myReverse xs = reverse' [] xs
  where reverse' : Vect n a -> Vect m a -> Vect (n+m) a
        reverse' acc [] = ?reverseProof_nil acc
        reverse' acc (x :: xs)
          = ?reverseProof_xs (reverse' (x :: acc) xs)
```

Complete this definition by implementing the holes ?reverseProof_nil and ?reverseProof_xs.

You can test your answer at the REPL as follows:

```
*ex_8_2> myReverse [1,2,3,4]
[4, 3, 2, 1] : Vect 4 Integer
```

8.3 The empty type and decidability

You can use the equality type, =, to write functions with types that state that two values are guaranteed to be equal, and then use this guarantee elsewhere in the program. This works because the only way to construct a value with an equality type is to use Refl, and Refl will only construct a value with a type of the form x = x:

```
Idris> :t Refl
Refl : x = x
```

```
Idris> Refl {x = 94}
Refl : 94 = 94
```

But what if you want to say the opposite, that two values are guaranteed *not* to be equal? When you construct a value in a type, you’re effectively giving evidence that an element of that type exists. To show that two values x and y are *not* equal, you need to be able to give evidence that an element of the type $x = y$ *can’t* exist.

In this section, you’ll see how to use the *empty type*, `Void`, to express that something is impossible. If a function returns a value of type `Void`, that can only mean that it’s impossible to construct values of its inputs (or, logically, that the types of its inputs express a *contradiction*.) We’ll use `Void` to express guarantees in types that values *can’t* be equal, and then use it to write a more precise type for `checkEqNat`, which guarantees that

- If its inputs are equal, it will produce a proof that they are equal
- If its inputs are not equal, it will produce a proof that they are not equal

First, let’s take a look at how `Void` is defined and used in the simplest case.

8.3.1 Void: a type with no values

In order to express that something *can’t* happen, the Prelude provides a type with no values, `Void`. The complete definition of `Void` is as follows:

```
data Void : Type where
```

You can’t write values of type `Void` directly, because there aren’t any! As a result, if you have a function that returns something of type `Void`, it must be because one of its arguments is also impossible to construct.

Just as you can use `=` to write functions that express facts about how functions *do* behave, you can use `Void` to express facts about how functions *don’t* behave. For example, you can show that $2 + 2$ doesn’t equal 5.

Let’s write a function in a file named `Void.idr`:

- 1 *Type*—Begin by writing the appropriate type:

```
twoPlusTwoNotFive : 2 + 2 = 5 -> Void
```

You can read this “if $2 + 2 = 5$, then return an element of the empty type.”

- 2 *Define*—Adding a skeleton definition gives you this:

```
twoPlusTwoNotFive : 2 + 2 = 5 -> Void
twoPlusTwoNotFive prf = ?twoPlusTwoNotFive_rhs
```

If you look at the type of `prf`, you’ll see that it’s a proof that $4 = 5$:

```
prf : 4 = 5
-----
twoPlusTwoNotFive_rhs : Void
```

- 3 *Define*—You can try to define the function by a case split on `prf`. Idris produces this:

```
twoPlusTwoNotFive : 2 + 2 = 5 -> Void
twoPlusTwoNotFive Refl impossible
```

This definition is now complete. Idris has produced one case, and noticed that the only possible input, `Refl`, can never be valid. Recall from chapter 4 that the `impossible` keyword means that the pattern clause must not type-check.

Void and total functions

It's important to check that a function that returns `Void` is a total function, if you really want to believe that it takes an impossible input:

```
*Void> :total twoPlusTwoNotFive
Main.twoPlusTwoNotFive is Total
```

Otherwise, you could write a function that claims to return `Void` by looping forever:

```
loop : Void
loop = loop
```

This function isn't total, so you can't believe that it really does produce an element of `Void`. Idris reports the following:

```
*Void> :total loop
Main.loop is possibly not total due to recursive path:
    Main.loop
```

Similarly, you can write a function that shows that a number can never be equal to its successor:

```
valueNotSuc : (x : Nat) -> x = S x -> Void
valueNotSuc _ Refl impossible
```

If you were able to provide a value of the empty type, you'd be able to produce a value of *any* type. In other words, if you have a proof that an impossible value has happened, you can do anything. The Prelude provides a function, `void`, that expresses this:

```
void : Void -> a
```

It may seem strange, and of little practical use, to be writing functions merely to show that something can't happen. But if you know that something *can't* happen, you can use this knowledge to express limitations about what *can* happen. In other words, you can express more precisely what a function is intended to do.

8.3.2 Decidability: checking properties with precision

Previously, when you wrote `checkEqNat`, you used `Maybe` for the result:

```
checkEqNat : (num1 : Nat) -> (num2 : Nat) -> Maybe (num1 = num2)
```

So, if `checkEqNat` returns a value of the form `Just p`, you can be certain that `num1` and `num2` are equal, and that `p` represents a proof that they're equal. But you can't say the opposite: that if `checkEqNat` returns `Nothing`, then `num1` and `num2` are guaranteed *not* to be equal.

The following definition would, for example, be perfectly valid, though not very useful:

```
checkEqNat : (num1 : Nat) -> (num2 : Nat) -> Maybe (num1 = num2)
checkEqNat num1 num2 = Nothing
```

Instead, to make this type more precise, you'd need a way of stating that for any pair of numbers, `num1` and `num2`, you'll always be able to produce either a proof that they're equal (of type `num1 = num2`) or a proof that they're not equal (of type `num1 = num2 -> Void`). That is, you'd like to state that checking whether `num1 = num2` is a *decidable* property.

DECIDABILITY A property of some values is decidable if you can always say whether the property holds or not for specific values. For example, checking equality on `Nat` is decidable, because for any two natural numbers you can *always* decide whether they are equal or not.

Listing 8.8 shows the `Dec` generic type, which is defined in the Prelude. Like `Maybe`, `Dec` has a constructor (`Yes`) that carries a value. Unlike `Maybe`, it also has a constructor (`No`) that carries a proof that no value of its argument type can exist.

Listing 8.8 Dec: precisely stating that a property is decidable

```
→ data Dec : (prop : Type) -> Type where
  Yes : (prf : prop) -> Dec prop
  No : (contra : prop -> Void) -> Dec prop
```

The type `Dec prop` states that you can decide whether `prop` is either guaranteed to hold, or guaranteed to be impossible.

`contra` is a proof that the property, `prop`, doesn't hold, because it's a function that returns a value of the empty type `Void`, given a `prop`.

<div style="position: absolute; left: 0; top: 50%;

- If the inputs num1 and num2 are equal, checkEqNat is guaranteed to produce a result of the form Yes prf, where prf has type num1 = num2.
- If the inputs num1 and num2 are *not* equal, checkEqNat is guaranteed to produce a result of the form No contra, where contra has type num1 = num2 \rightarrow Void.

These are guarantees because the type of checkEqNat gives a direct link between the inputs and the output type:

```
checkEqNat : (num1 : Nat) -> (num2 : Nat) -> Dec (num1 = num2)
```

You can write this function interactively, in a new file called CheckEqDec.idr:

- 1 *Define, refine*—You can mostly follow the same steps as earlier, when you defined checkEqNat using Maybe. But instead of Just, use Yes, and instead of Nothing, use No. No requires proofs that the inputs are unequal, so you can leave holes for these for the moment:

```
checkEqNat : (num1 : Nat) -> (num2 : Nat) -> Dec (num1 = num2)
checkEqNat Z Z = Yes Refl
checkEqNat Z (S k) = No ?zeroNotSuc
checkEqNat (S k) Z = No ?sucNotZero
checkEqNat (S k) (S j) = case checkEqNat k j of
    Yes prf => Yes (cong prf)
    No contra => No ?noRec
```

You can test this at the REPL, even without the proofs, though you may see holes in the results:

```
*CheckEqDec> checkEqNat 3 3
Yes Refl : Dec (3 = 3)

*CheckEqDec> checkEqNat 3 4
No ?noRec : Dec (3 = 4)

*CheckEqDec> checkEqNat 3 0
No ?sucNotZero : Dec (3 = 0)
```

- 2 *Type*—If you look at the types of the holes, you’ll see what you need to prove to complete the definition, as in this example:

```
k : Nat
-----
zeroNotSuc : (0 = S k) -> Void
```

- 3 *Define, refine*—You can lift ?zeroNotSuc and ?sucNotZero to top-level definitions with Ctrl-Alt-L and implement them using case splitting and impossible, as with twoPlusTwoNotFive in the previous section, because it’s never possible for zero to be equal to a nonzero number:

```
zeroNotSuc : (0 = S k) -> Void
zeroNotSuc Refl impossible

sucNotZero : (S k = 0) -> Void
sucNotZero Refl impossible
```

- 4 *Type, define*—For `?noRec`, you can take a look at the type to see what you need to show:

```
k : Nat
j : Nat
contra : (k = j) -> Void
```

```
noRec : (S k = S j) -> Void
```

The type of contra tells you that k is guaranteed not to be equal to j . Given that, you have to show that $S\ k$ can't equal $S\ j$. Again, you can lift this to a top-level definition with Ctrl-Alt-L:

```

noRec : (contra : (k = j) -> Void) -> (S k = S j) -> Void
noRec contra prf = ?noRec rhs

```

- 5 Define—The type of `noRec` says that if you have a proof that $k = j$ is impossible, and a proof that $S k = S j$, you can produce an element of the empty type. Logically, the two inputs contradict each other. You can write this function by case splitting on `prf`:

```

noRec : (contra : (k = j) -> Void) -> (S k = S j) -> Void
noRec contra Refl = ?noRec rhs 1

```

- 6 **Type**—The only possible value `prf` can take is `Ref1`, and the only way it can take the value `Ref1` is if `s k` and `s j` are equal, and therefore `k` and `j` are equal. Recognizing this, Idris refines the type of `contra`, as you can see by inspecting the type of `noRec rhs 1`:

```
j : Nat
contra : (j = j) -> Void
-----
noRec rhs 1 : Void
```

- 7 **Refine**—To complete the definition, you can use `contra` to produce an element of the empty type; an expression search will find this:

```

noRec : (contra : (k = j) -> Void) -> (S k = S j) -> Void
noRec contra Refl = contra Refl

```

The following listing shows the complete definition of `checkEqNat`, including the helper functions `zeroNotSuc`, `sucNotZero`, and `noRec`.

Listing 8.9 Checking whether Nats are equal, with a precise type (`CheckEqDec.idr`)

<pre>zeroNotSuc : (0 = S k) -> Void zeroNotSuc Refl impossible</pre> <pre>sucNotZero : (S k = 0) -> Void sucNotZero Refl impossible</pre> <pre>noRec : (contra : (k = j) -> Void) -> (S k = S j) -> Void noRec contra Refl = contra Refl</pre>	<p>Given a proof that zero equals a nonzero number, produce a value of the empty type.</p> <p>Given a proof that a nonzero number equals zero, produce a value of the empty type.</p> <p>Given a proof that two numbers aren't equal, and a proof that their successors are equal, produce a value of the empty type.</p>
---	---

```

checkEqNat : (num1 : Nat) -> (num2 : Nat) -> Dec (num1 = num2)
checkEqNat Z Z = Yes Refl
checkEqNat Z (S k) = No zeroNotSuc
checkEqNat (S k) Z = No sucNotZero
checkEqNat (S k) (S j) = case checkEqNat k j of
    Yes prf => Yes (cong prf)
    No contra => No (noRec contra)

```

PROVING INPUTS ARE IMPOSSIBLE WITH VOID When you run `checkEqNat`, you aren't really going to produce a value of the empty type using `zeroNotSuc`, `sucNotZero`, or `noRec`. Essentially, a function that produces a value of type `Void` can be seen as a proof that its arguments can't all be provided at the same time. In the case of `noRec`, the type of the functions says that if you can provide *both* a proof that `k` doesn't equal `j` *and* a proof that `S k = S j`, then there's a contradiction, and you can therefore have a value of type `Void`.

By using `Dec`, you've been able to write explicitly in the type what `checkEqNat` is supposed to do: return either a proof that the inputs are equal or a proof that they're not. The real benefit of this comes not in `checkEqNat` itself, however, but in the functions that use it, because they not only have the result of the equality test, but also a proof that the equality test has worked correctly.

Being able to guarantee that two values are equal (or different) is commonly useful in type-driven development, because showing relationships between larger structures depends on showing relationships between the individual components. Because of this, the Idris Prelude provides an interface, `DecEq`, with a generic function, `decEq`, for deciding equality.

8.3.3 DecEq: an interface for decidable equality

Rather than providing specific functions like `checkEqNat` for each type, the Idris Prelude provides an interface, `DecEq`. The next listing shows how `DecEq` is defined. There are implementations for all of the types defined in the Prelude.

Listing 8.10 The DecEq interface (defined in the Prelude)

```

interface DecEq ty where
  decEq : (val1 : ty) -> (val2 : ty) -> Dec (val1 = val2)

  Given two values, val1 and val2, return either a proof
  that they are equal or a proof that they are different.

```

Instead of defining and using a special-purpose function, `checkEqNat`, to define `exactLength`, you could use `decEq`. The following listing shows how to do this, and it's the definition of `exactLength` used in `Data.Vect`.

Listing 8.11 Implementing exactLength using decEq (ExactLengthDec.idr)

```

exactLength : (len : Nat) -> (input : Vect m a) -> Maybe (Vect len a)
exactLength {m} len input = case decEq m len of
    Yes Refl => Just input
    No contra => Nothing

```

Using `decEq` rather than the Boolean equality operator, `==`, gives you a strong guarantee about how the equality test works. You can be sure that if `decEq` returns a value of the form `Yes prf`, then the inputs really are structurally equal.

In the next chapter, you'll see how to describe relationships between larger data structures (such as showing that a value is an element of a list) and how to use `decEq` to build proofs of these relationships.

Exercises



- 1** Implement the following functions:

```
headUnequal : DecEq a => {xs : Vect n a} -> {ys : Vect n a} ->
  (contra : (x = y) -> Void) -> ((x :: xs) = (y :: ys)) -> Void
tailUnequal : DecEq a => {xs : Vect n a} -> {ys : Vect n a} ->
  (contra : (xs = ys) -> Void) -> ((x :: xs) = (y :: ys)) -> Void
```

The first states that if the first elements of two vectors are unequal, then the vectors must be unequal. The second states that if there are differences in the tails of two vectors, then the vectors must be unequal.

If you have a correct solution, both `headUnequal` and `tailUnequal` should type-check and be total:

```
*ex_8_3> :total headUnequal
Main.headUnequal is Total

*ex_8_3> :total tailUnequal
Main.tailUnequal is Total
```

- 2** Implement `DecEq` for `Vect`. Begin with the following implementation header:

```
DecEq a => DecEq (Vect n a) where
```

Hint: You'll find `headUnequal` and `tailUnequal` useful here. Remember to check the types of the holes as you write the definition. You should also use your own definition of `Vect` rather than importing `Data.Vect`, because the library provides a `DecEq` implementation.

You can test your answer at the REPL as follows:

```
*ex_8_3> decEq (the (Vect _ _) [1,2,3]) [1,2,3]
Yes Refl : Dec ([1, 2, 3] = [1, 2, 3])
```

8.4

Summary

- You can write a type to express a proof that two values must be equal.
- You can write functions that take equality types as inputs to prove additional properties of data.
- Using `Maybe`, you can test for equality of values at runtime.
- The generic `=` type allows you to describe equality between values of any type.
- Proof requirements arise naturally when programming with dependent types, such as to show that lengths are preserved when reversing a vector.

- The `rewrite` construct allows you to update a type using an equality proof.
- Using holes and interactive editing, you can delegate the details of rewriting types to a separate function.
- `Void` is a type with no values, used to show that inputs to a function can't all occur at once.
- A property is decidable if you can always say whether the property holds for some specific values.
- Using `Dec`, you can compute at runtime whether a property is guaranteed to hold or guaranteed not to hold.

Predicates: expressing assumptions and contracts in types

This chapter covers

- Describing and checking membership of a vector using a predicate
- Using predicates to describe contracts for function inputs and outputs
- Reasoning about system state in types

Dependent types like `EqNat` and `=`, which you saw in the previous chapter, are used entirely for describing relationships between data. These types are often referred to as *predicates*, which are data types that exist entirely to describe a property of some data. If you can construct a value for a predicate, then you know the property described by that predicate must be true.

In this chapter, you'll see how to express more-complex relationships between data using predicates. By expressing relationships between data in types, you can be explicit about the *assumptions* you're making about the inputs to a function, and have those assumptions checked by the type checker when those functions are

called. You can even think of these assumptions as expressing compile time *contracts* that other arguments must satisfy before anything can call the function.

In practice, you'll often write functions that make assumptions about the form of some other data, having (hopefully!) checked those assumptions beforehand. Here are a few examples:

- A function that reads from a file assumes that the file handle describes an open file.
- A function that processes data it has received from a network assumes that the data follows the appropriate protocol.
- A function that retrieves user data (say, for a customer on a website) assumes that the user has successfully authenticated.

You need to ensure that you check any necessary assumptions before you call these functions, particularly when security or user privacy is at stake. In Idris, you can make this kind of assumption explicit in a type, and have the type checker ensure that you've properly checked the assumption in advance. An advantage of expressing assumptions in types is that they are *guaranteed* to hold, even as a system evolves over time.

In this chapter, we'll take an in-depth look at a specific small example: removing a value from a vector if that value is contained within the vector. We'll look at how to express in a type that a vector contains a specific element, how to test this property at runtime, and how to use such properties in practice to write a larger program with some aspects of its behavior expressed in its types.

9.1 Membership tests: the *Elem* predicate

Using `=`, `Void`, and `Dec`, you can describe *properties* that your programs satisfy in their types, making the types precise as a result. Typically, you'll use `=`, `Void`, and `Dec` as building blocks for describing and checking relationships between pieces of data. Describing relationships in types allows you to state assumptions that functions make and, more importantly, it allows Idris to check whether those assumptions are satisfied or violated.

For example, you might want to

- Remove an element from a vector only under the assumption that the element is present in the vector
- Search for a value in a list only under the assumption that the list is ordered
- Run a database query only under the assumption that the inputs have been validated
- Send a message to a machine on the network only under the assumption that you have an open connection

In this section, we'll examine the first of these in more detail. We'll begin by trying to write a function that removes an element from a vector, and then we'll see how we can refine the type to state and check any necessary assumptions that make the type more precise. Finally, we'll use the function in a simple interactive program.

9.1.1 Removing an element from a Vect

If you have a Vect containing a number of elements, you'd expect the following of a removeElem function that removes a specific element from that vector:

- The input vector should have at least one element.
- The output vector's length should be one less than the input vector's length.

Because Vect is parameterized by the element type contained in the vector and is indexed by the length of the vector, you can (indeed, you must) express these length properties in the type of removeElem. Your first attempt might look like this:

```
removeElem : (value : a) -> (xs : Vect (S n) a) -> Vect n a
```

Let's see what happens if you try to write this function:

- 1 *Define*—If you add a skeleton definition and case-split on xs, Idris will give you only one pattern, because xs can't be an empty vector:

```
removeElem : (value : a) -> (xs : Vect (S n) a) -> Vect n a
removeElem value (x :: xs) = ?removeElem_rhs_1
```

- 2 *Refine*—If the value you're looking for is equal to x, you can remove it from the list, returning xs. You'll need to refine the type before you can compare value and x; you can use decidable equality so that you can be certain that the equality test is accurate. This is the refined type:

```
removeElem : DecEq a => (value : a) -> (xs : Vect (S n) a) -> Vect n a
```

- 3 *Refine*—You can now use decEq to compare value and x, and discard x if they're equal:

```
removeElem : DecEq a => (value : a) -> (xs : Vect (S n) a) -> Vect n a
removeElem value (x :: xs) = case decEq value x of
    Yes prf => xs
    No contra => ?removeElem_rhs_3
```

- 4 *Refine failure*—For ?removeElem_rhs_3, you might hope to be able to remove value from xs recursively, but if you try this, you'll get an error:

```
removeElem : DecEq a => (value : a) -> (xs : Vect (S n) a) -> Vect n a
removeElem value (x :: xs) = case decEq value x of
    Yes prf => xs
    No contra => x :: removeElem value xs
```

Idris reports the following:

```
When checking right hand side of Main.case block in removeElem
at removeelem.idr:4:35 with expected type
    Vect n a
```

```
When checking argument xs to Main.removeElem:
```

```
Type mismatch between
```

```
    Vect n a (Type of xs)
```

```
and
```

```
    Vect (S k) a (Expected type)
```

The problem is that `removeElem` requires a vector that's guaranteed to be non-empty, but `xs` may be empty! You can see this from its type, `Vect n a`: the `n` could stand for any natural number, including zero.

This problem arises because there's no guarantee that `value` will appear in the vector, so it's possible to reach the end of the vector without encountering it. If this happens, there's no value to remove, so you can't satisfy the type.

You'll need to refine the type further in order to be able to write this function. You can try one of the following:

- Rewrite the type so that `removeElem` returns `Maybe (Vect n a)`, returning `Nothing` if the value doesn't appear in the input.
- Rewrite the type so that `removeElem` returns a dependent pair, `(newLength ** Vect newLength a)`.
- Express a *precondition* on `removeElem` that the value is guaranteed to be in the vector.

You've already seen how to express possible failure with `Maybe` (in chapter 4), and how to express an unknown length using dependent pairs (in chapter 5). The third option, however, would express the purpose of `removeElem` precisely. To achieve this, you'll need to write a type that describes the relationship between a value and a vector that contains that value.

If you can describe in a type that a vector contains a specific element, you'll be able to use that type to express a *contract* on `removeElem` expressing that you can only use it if you know the element is in the vector. In the rest of this section, you'll implement this type, `Elem`, use it to make the type of `removeElem` more precise, and learn more about how to use `Elem` in practice.

9.1.2 The `Elem` type: guaranteeing a value is in a vector

In the previous chapter, you saw how to express that two values are guaranteed to be equal by using either a specific `EqNat` type or the generic `=` type. The existence of a value in one of these types is, essentially, a proof that two values are equal. You can do something similar to guarantee that a value is in a vector.

Our goal is to define a type, `Elem`, with the following type constructor:

```
Elem : (value : a) -> (xs : Vect k a) -> Type
```

If we have a `value`, and a vector, `xs`, that contains that value, we should be able to construct an element of the type `Elem value xs`. For example, we should be able to construct the following:

```
oneInVector : Elem 1 [1,2,3]
maryInVector : Elem "Mary" ["Peter", "Paul", "Mary"]
```

We should also be able to construct functions of the following types, which show that a specific value is *not* contained in a vector:

```
fourNotInVector : Elem 4 [1,2,3] -> Void
peteNotInVector : Elem "Pete" ["John", "Paul", "George", "Ringo"] -> Void
```

The following listing shows how the `Elem` dependent type is defined in `Data.Vect`.

Listing 9.1 The `Elem` dependent type, expressing that a value is guaranteed to be contained in a vector (defined in `Data.Vect`)

Here states that `x` is the first value in a vector of the form `x :: xs`.

```
data Elem : a -> Vect k a -> Type where
  Here : Elem x (x :: xs)
  There : (later : Elem x xs) -> Elem x (y :: xs)
```

There states that if you know that `x` occurs in a vector `xs`, then `x` must also occur in a vector `y :: xs`.

The value `Here` can be read as a proof that a value is the first value in a vector, as in this example:

```
oneInVector : Elem 1 [1,2,3]
oneInVector = Here
```

The constructor `There`, given an argument that shows a value `x` is in the vector `xs`, can be read as a proof that `x` must also be in the vector `y :: xs` for any value `y`.

To illustrate this, let's try to write `maryInVector`:

```
maryInVector : Elem "Mary" ["Peter", "Paul", "Mary"]
```

1 Define—Create a skeleton definition with a hole for the right side:

```
maryInVector : Elem "Mary" ["Peter", "Paul", "Mary"]
maryInVector = ?maryInVector_rhs
```

2 Refine, type—You can't use `Here`, because "Mary" isn't the first element of the vector, so try using `There` and leaving a hole for its argument:

```
maryInVector : Elem "Mary" ["Peter", "Paul", "Mary"]
maryInVector = There ?maryInVector_rhs
```

If you check the type of `?maryInVector_rhs`, you'll see that you now have a smaller problem:

```
-----  
maryInVector_rhs : Elem "Mary" ["Paul", "Mary"]
```

3 Refine, type—If you do the same again, applying `There` and leaving a hole for its argument, you get the following program:

```
maryInVector : Elem "Mary" ["Peter", "Paul", "Mary"]
maryInVector = There (There ?maryInVector_rhs)
```

You also now have a simpler type for `?maryInVector_rhs`:

```
-----  
maryInVector_rhs : Elem "Mary" ["Mary"]
```

- 4 *Refine*—"Mary" is now the first element in the vector, so you can refine `?mary-InVector_rhs` using `Here`:

```
maryInVector : Elem "Mary" ["Peter", "Paul", "Mary"]
maryInVector = There (There Here)
```

EXPRESSION SEARCH An expression search with Ctrl-Alt-S will successfully find the definition of `maryInVector`, and it's often useful for constructing values of dependent types like `Elem` and `=`.

You can use `Elem` and types like it that describe relationships between data to express *contracts* on the form of data expected as input to a function. These contracts, being expressed as types, can be verified by Idris using type checking; if a function call violates a contract, the program will not compile.

More constructively, if you express the types of your functions precisely enough, with contracts expressed using types like `Elem`, you know that once your program compiles, every contract must be satisfied. You can use `Elem` to express a contract on `removeElem` that specifies when the inputs are valid.

9.1.3 Removing an element from a Vect: types as contracts

The difficulty we had when writing `removeElem` was that we couldn't make a recursive call on the tail of the vector because we couldn't guarantee that the element we were trying to remove was in the vector. We'll therefore refine the type of `removeElem`, adding an argument that expresses the contract that the element to be removed must be in the vector. The following listing shows our starting point.

Listing 9.2 Removing an element from a vector, with a contract specified in the type using `Elem` (`RemoveElem.idr`)

```
removeElem : (value : a) ->           ← The value to be removed
              (xs : Vect (S n) a) ->      ← The vector to remove the value from
              (prf : Elem value xs) ->    ← A proof that the value is present
                                              in the vector, expressing a
                                              contract for the caller to satisfy
removeElem value xs prf = ?removeElem_rhs
```

In type-driven development, we aim to use more-precise types to help direct the implementation of functions. Here, the input `prf` gives more precision to the input type of `removeElem`, and you can even case-split on it to see what you learn about the inputs `value` and `xs` from the relationship specified between them by `prf`.

You can write the function as follows:

- 1 *Define*—Begin with the case split on `prf`:

```
removeElem : (value : a) -> (xs : Vect (S n) a) ->
              (prf : Elem value xs) ->
              Vect n a
removeElem value (value :: ys) Here = ?removeElem_rhs_1
removeElem value (y :: ys) (There later) = ?removeElem_rhs_2
```

- 2 Refine**—For the first case, `?removeElem_rhs_1`, you can see from the patterns that Idris has generated that value *must* be the first element in the vector if the proof has the form `Here`. You can therefore refine this case to remove value:

```
removeElem : (value : a) -> (xs : Vect (S n) a) ->
            (prf : Elem value xs) ->
            Vect n a
removeElem value (value :: ys) Here = ys
removeElem value (y :: ys) (There later) = ?removeElem_rhs_2
```

- 3 Type**—If you look at the type of `?removeElem_rhs_2`, you’ll see that it appears you have the same problem as before, in that `ys` has length `n`, and `removeElem` requires a vector of length `S n` for the recursive call:

```
a : Type
value : a
y : a
n : Nat
ys : Vect n a
later : Elem value ys
-----
removeElem_rhs_2 : Vect n a
```

But you have some further information that you didn’t have available earlier: you know from the variable `later` that `value` must occur in `ys`, and this means that `ys` must have a nonzero length. But how can you use this knowledge?

- 4 Define**—Given that the length `n` must be nonzero, the trick is to case-split on `n` (by bringing it into scope using braces on the left side of the definition) and show that it’s impossible for it to be zero:

```
removeElem : (value : a) -> (xs : Vect (S n) a) ->
            (prf : Elem value xs) ->
            Vect n a
removeElem value (value :: ys) Here = ys
removeElem {n = Z} value (y :: ys) (There later) = ?removeElem_rhs_1
removeElem {n = (S k)} value (y :: ys) (There later) = ?removeElem_rhs_3
```

- 5 Refine**—Now, you can complete `?removeElem_rhs_3`. You now know the length of `ys` is nonzero, so you can make a recursive call:

```
removeElem : (value : a) -> (xs : Vect (S n) a) -> (prf : Elem value xs) ->
            Vect n a
removeElem value (value :: ys) Here = ys
removeElem {n = Z} value (y :: ys) (There later) = ?removeElem_rhs_1
removeElem {n = (S k)} value (y :: ys) (There later)
                                         = y :: removeElem value ys later
```

Note that you need to pass `later` to the recursive call, as evidence that `value` is contained within `ys`.

- 6 Type, define**—For the remaining hole, `?removeElem_rhs_1`, take a look at its type and what variables are available:

```
a : Type
value : a
y : a
ys : Vect 0 a
later : Ele� value ys
```

You're looking for an empty vector. That empty vector, if you look at the variables on the left side, should be a vector resulting from removing value from ys. This doesn't make sense, because ys is an empty vector!

You can make this more clear by case splitting on `ys`. Idris produces only one case:

Looking at the type of the new hole, `?removeElem rhs 1`, shows the following:

```
a : Type  
value : a  
y : a  
later : Ele� value []  
-----  
removeElem rhs 1 : Vect 0 a
```

Previously, you've used the `impossible` keyword to rule out a case that doesn't type-check. This case *does* type-check, so you can't use `impossible`. But there's no way you'll ever have a value of type `Elem` value `[]` as an input, because there's no way to construct an element of this type.

- 7 **Refine**—You can complete the definition using a function, `absurd`, defined in the Prelude. It has the following type:

absurd : Uninhabited \Rightarrow $t \Rightarrow$ a

The `Uninhabited` interface, described in the sidebar, can be implemented for any type that has no values (as you saw with `twoplustwo_not_five` in chapter 8). So, you can refine `?remove` `Elem` `rhs_1` as follows:

The Uninhabited interface

If a type has no values, like `2 + 2 = 5` or `Elem x []`, you can provide an implementation of the `Uninhabited` interface for it. `Uninhabited` is defined in the Prelude as follows:

```
interface Uninhabited t where
    uninhabited : t -> Void
```

There's one method, which returns an element of the empty type. For example, you can provide an implementation of `Uninhabited` for `2 + 2 = 5`:

```
Uninhabited (2 + 2 = 5) where
    uninhabited Refl impossible
```

Using `uninhabited`, the Prelude defines `absurd` as follows, using `void`:

```
absurd : Uninhabited t => (h : t) -> a
absurd h = void (uninhabited h)
```

The additional argument to `removeElem`, of type `Elem value xs`, means that `removeElem` can work under the assumption that `value` is in the vector `xs`. It also means that you must provide a proof that `value` is in `xs` when calling the function.

9.1.4 auto-implicit arguments: automatically constructing proofs

If you try running `removeElem` with some specific values for the element and the vector, you'll find you need to provide an additional argument for the proof. Sometimes it might be a proof that's possible to construct:

```
*RemoveElem> removeElem 2 [1,2,3,4,5]
removeElem 2 [1, 2, 3, 4, 5] : Elem 2 [1, 2, 3, 4, 5] -> Vect 4 Integer
```

Sometimes it might not:

```
*RemoveElem> removeElem 7 [1,2,3,4,5]
removeElem 7 [1, 2, 3, 4, 5] : Elem 7 [1, 2, 3, 4, 5] -> Vect 4 Integer
```

In the first case, you can call `removeElem` by explicitly providing a proof of `Elem 2 [1,2,3,4,5]`:

```
*RemoveElem> removeElem 2 [1,2,3,4,5] (There Here)
[1, 3, 4, 5] : Vect 4 Integer
```

The need to provide proofs explicitly like this can add a lot of noise to programs and can harm readability. Idris provides a special kind of implicit argument, marked with the keyword `auto`, to reduce this noise.

You can define a `removeElem_auto` function:

```
removeElem_auto : (value : a) -> (xs : Vect (S n) a) ->
                  {auto prf : Elem value xs} -> Vect n a
removeElem_auto value xs {prf} = removeElem value xs prf
```

The third argument, named `prf`, is an auto-implicit argument. Like the implicit arguments you've already seen, an auto-implicit argument can be brought into scope by writing it in braces, and Idris will attempt to find a value automatically. Unlike ordinary implicits, Idris will search for a value for an auto implicit using the same machinery it uses for expression search with Ctrl-Alt-S in Atom.

When you run `removeElem_auto` with arguments for `value` and `xs`, Idris will try to construct a proof for the argument marked `auto`:

```
*RemoveElem> removeElem_auto 2 [1,2,3,4,5]
[1, 3, 4, 5] : Vect 4 Integer
```

If it can't find a proof, it will report an error:

```
*RemoveElem> removeElem_auto 7 [1,2,3,4,5]
(input):1:17:When checking argument prf to function Main.removeElem_auto:
    Can't find a value of type
        Elem 7 [1, 2, 3, 4, 5]
```

Alternatively, the following listing shows how you could define `removeElem` using an auto implicit directly.

Listing 9.3 Defining `removeElem` using an auto-implicit argument (`RemoveElem.idr`)

```
removeElem : (value : a) -> (xs : Vect (S n) a) ->
    {auto prf : Elem value xs} ->
    Vect n a
removeElem value (value :: ys) {prf = Here} = ys
removeElem {n = Z} value (y :: []) {prf = There later} = absurd later
removeElem {n = (S k)} value (y :: ys) {prf = There later}
    = y :: removeElem value ys
```

`prf` is an auto-implicit argument. Idris will try to find a value automatically for each call, using expression search.

Matches on `prf` by bringing it into scope using braces

Idris will find the value (later) for the proof in this call automatically.

In general, though, you won't know the specific values you're passing to `removeElem`. They could be values read from a user or constructed by another part of the program. We'll therefore need to consider how you can use `removeElem` when the inputs are unknown until runtime.

Just as you wrote a `checkEqNat` function to decide whether two numbers are equal, and later generalized it to `decEq` using an interface, you'll need a function to decide whether a value is contained in a vector.

9.1.5 Decidable predicates: deciding membership of a vector

As you saw in the last chapter, a property is decidable if you can always say whether the property holds for some specific values. Using the following function, you can see that `Elem value xs` is a decidable property for specific values of `value` and `xs`, so `Elem` is a decidable predicate:

```
isElem : DecEq ty => (value : ty) -> (xs : Vect n ty) -> Dec (Elem value xs)
```

The type of `isElem` states that as long as you can decide equality of values in some type `ty`, you can decide whether a value with type `ty` is contained in a vector of types `ty`. Remember that `Dec` has the following constructors:

- Yes, which takes as its argument a proof that the predicate holds. In this case, that's a value of type `Elem` value `xs` for whichever value and `xs` are passed as inputs to `isElem`.
- No, which takes as its argument a proof that the predicate doesn't hold. In this case, that's a value of type `Elem` value `xs -> Void`.

`isElem` is defined in `Data.Vect`, but it's instructive to see how to write it yourself. The following listing shows our starting point, defining `Elem` by hand in a file named `ElemType.idr`.

Listing 9.4 Defining `Elem` and `isElem` by hand (`ElemType.idr`)

```
data Elem : a -> Vect k a -> Type where
  Here : Elem x (x :: xs)
  There : (later : Elem x xs) -> Elem x (y :: xs)

isElem : DecEq a => (value : a) -> (xs : Vect n a) -> Dec (Elem value xs)
```

- 1 *Define, refine*—Define the function by case splitting on the input vector `xs`:

```
isElem : DecEq a => (value : a) -> (xs : Vect n a) -> Dec (Elem value xs)
isElem value [] = ?isElem_rhs_1
isElem value (x :: xs) = ?isElem_rhs_2
```

- 2 *Refine*—For `?isElem_rhs_1`, `value` is clearly not in the empty vector, so you can return `No`. Remember that `No` takes as an argument a proof that you *can't* construct the predicate. You can leave a hole for this argument for the moment:

```
isElem : DecEq a => (value : a) -> (xs : Vect n a) -> Dec (Elem value xs)
isElem value [] = No ?notInNil
isElem value (x :: xs) = ?isElem_rhs_2
```

If you check the type of `?notInNil`, you'll see that to fill in this hole, you need to provide a proof that there *can't* be an element of the empty vector:

```
a : Type
value : a
constraint : DecEq a
-----
notInNil : Elem value [] -> Void
```

We'll return to this hole shortly.

- 3 *Refine*—For `?isElem_rhs_2`, if `value` and `x` are equal, you've found the value you're looking for. You can use `decEq`, case-split on the result, and if the result is of the form `Yes prf`, return `Yes` with a hole for the proof that the value is the first in the vector:

```
isElem : DecEq a => (value : a) -> (xs : Vect n a) -> Dec (Elem value xs)
```

```

isElem value [] = No ?notInNil
isElem value (x :: xs) = case decEq value x of
    Yes prf => Yes ?isElem_rhs_1
    No notHere => ?isElem_rhs_3

```

- 4 *Type, refine*—You might like to fill the `?isElem_rhs_1` hole with `Here`, but if you check its type, you'll see that you can't quite use `Here` yet:

```

a : Type
value : a
x : a
prf : value = x
k : Nat
xs : Vect k a
constraint : DecEq a
-----
isElem_rhs_1 : Elem value (x :: xs)

```

You can't use `Here` because the type you're looking for isn't of the form `Elem value (value :: xs)`. But `prf` tells you that `value` and `x` *must* be the same, so if you case-split on `prf`, you'll get this:

```

isElem value (x :: xs) = case decEq value x of
    Yes Refl => Yes ?isElem_rhs_2
    No notHere => ?isElem_rhs_3

```

The type of the newly created hole, `?isElem_rhs_2`, is now in the form you need:

```

a : Type
value : a
k : Nat
xs : Vect k a
constraint : DecEq a
-----
isElem_rhs_2 : Elem value (value :: xs)

```

You can fill in the `?isElem_rhs_2` using expression search in Atom:

```

isElem : DecEq a => (value : a) -> (xs : Vect n a) -> Dec (Elem value xs)
isElem value [] = No ?notInNil
isElem value (x :: xs) = case decEq value x of
    Yes Refl => Yes Here
    No notHere => ?isElem_rhs_3

```

- 5 *Refine*—For `?isElem_rhs_3`, `value` and `x` aren't equal (you know this because `decEq value x` has returned a proof), so you can search recursively in `xs`:

```

isElem : DecEq a => (value : a) -> (xs : Vect n a) -> Dec (Elem value xs)
isElem value [] = No ?notInNil
isElem value (x :: xs) = case decEq value x of
    Yes Refl => Yes Here
    No notHere => case isElem value xs of
        Yes prf => Yes ?isElem_rhs_1
        No notThere => No ?isElem_rhs_2

```

Expression search will find the necessary proof for `?isElem_rhs_1`. You can leave a hole for the `No` case for now:

```
isElem : DecEq a => (value : a) -> (xs : Vect n a) -> Dec (Elem value xs)
isElem value [] = No ?notInNil
isElem value (x :: xs) = case decEq value x of
    Yes Refl => Yes Here
    No notHere => case isElem value xs of
        Yes prf => Yes (There prf)
        No notThere => No ?notInTail
```

At this stage, you can test the definition at the REPL. If a value is in a vector, you'll see `Yes` and a proof:

```
*ElemType> isElem 3 [1,2,3,4,5]
Yes (There (There Here)) : Dec (Elem 3 [1, 2, 3, 4, 5])
```

If not, you'll see `No`, with a hole for the proof that the element is missing:

```
*ElemType> isElem 7 [1,2,3,4,5]
No ?notInTail : Dec (Elem 7 [1, 2, 3, 4, 5])
```

To complete the definition, you'll need to complete `notInNil` and `notInTail`.

- 6 *Define, refine*—You can complete `?notInNil` by lifting it to a top-level function with Ctrl-Alt-L and then case splitting on its argument. Idris notices that neither input is possible:

```
notInNil : Elem value [] -> Void
notInNil Here impossible
notInNil (There _) impossible
```

- 7 *Define*—You can complete `?notInTail` similarly, but you'll need to work a bit harder to show that each case is impossible. Lifting to a top-level definition and then case splitting on the argument leads to the following:

```
notInTail : (notThere : Elem value xs -> Void) ->
            (notHere : (value = x) -> Void) ->
            Elem value (x :: xs) -> Void
notInTail notThere notHere Here = ?notInTail_rhs_1
notInTail notThere notHere (There later) = ?notInTail_rhs_2
```

For each hole, remember to check its type and the type of its local variables, because these will often give a strong hint as to how to proceed.

- 8 *Refine*—For each case, you can use either `notHere` or `notThere` to produce the value of type `Void` that you need:

```
notInTail : (notThere : Elem value xs -> Void) ->
            (notHere : (value = x) -> Void) ->
            Elem value (x :: xs) -> Void
notInTail notThere notHere Here = notHere Refl
notInTail notThere notHere (There later) = notThere later
```

Here's the completed definition, for reference.

Listing 9.5 Complete definition of `isElem` (`ElemType.idr`)

```
notInNil : Elem value [] -> Void
notInNil Here impossible
notInNil (There _) impossible

notInTail : (notThere : Elem value xs -> Void) ->
            (notHere : (value = x) -> Void) -> Elem value (x :: xs) -> Void
notInTail notThere notHere Here = notHere Refl
notInTail notThere notHere (There later) = notThere later

isElem : DecEq a => (value : a) -> (xs : Vect n a) -> Dec (Elem value xs)
isElem value [] = No notInNil
isElem value (x :: xs)
  = case decEq value x of
    Yes Refl => Yes Here
    No notHere => case isElem value xs of
      Yes prf => Yes (There prf)
      No notThere => No (notInTail notThere notHere)
```

For comparison, listing 9.6 shows how you could define a Boolean test for checking vector membership. Here, I've used the `Eq` interface, so we don't have any guarantees from the type about how the equality test behaves. Nevertheless, `elem` follows a structure similar to `decElem`.

Listing 9.6 A Boolean test for whether a value is in a vector (`ElemBool.idr`)

```
elem : Eq ty => (value : ty) -> (xs : Vect n ty) -> Bool
elem value [] = False
elem value (x :: xs) = case value == x of
  False => elem value xs
  True => True
```

Both definitions have a similar structure, but more work is required in `isElem` to show that the impossible cases are really impossible. As a trade-off, there's no need for *any* tests for `isElem`, because the type is sufficiently precise that the implementation must be correct.

Exercises



- 1 Data.List includes a version of `Elem` for `List` that works similarly to `Elem` for `Vect`. How would you define it?
- 2 The following predicate states that a specific value is the last value in a `List`:

```
data Last : List a -> a -> Type where
  LastOne : Last [value] value
  LastCons : (prf : Last xs value) -> Last (x :: xs) value
```

So, for example, you can construct a proof of `Last [1,2,3] 3`:

```
last123 : Last [1,2,3] 3
```

```
last123 = LastCons (LastCons LastOne)
```

Write an `isLast` function that decides whether a value is the last element in a `List`. It should have the following type:

```
isLast : DecEq a => (xs : List a) -> (value : a) -> Dec (Last xs value)
```

You can test your answer at the REPL as follows:

```
*ex_9_1> isLast [1,2,3] 3
Yes (LastCons (LastCons LastOne)) : Dec (Last [1, 2, 3] 3)
```

9.2 Expressing program state in types: a guessing game

In practice, the need for predicates like `Elem` and functions like `removeElem` arises naturally when we write functions that express characteristics of system state, such as vector length, in their types. To conclude this chapter, we'll look at a small example of where this happens: using the type system to encode simple properties of a word-guessing game, Hangman.

In Hangman, a player tries to guess a word by guessing one letter at a time. If they guess all the letters in the word, they win. Otherwise, they're allowed a limited number of incorrect guesses, after which they lose.

9.2.1 Representing the game's state

Listing 9.7 shows how we can represent the current state of a word game as a data type, `WordState`, in Idris. There are two important parts of the state that capture the features of the game, and they're written as parts of the type:

- The number of guesses the player has remaining
- The number of letters the player still has to guess

Listing 9.7 The game state (Hangman.idr)

```
data WordState : (guesses_remaining : Nat) -> (letters : Nat) -> Type where
  MkWordState : (word : String) ->
    (missing : Vect letters Char) ->
      WordState guesses_remaining letters
```

The word the player
is trying to guess

Lettters in the word that
have not yet been
guessed correctly

guesses_remaining is an implicit
argument of `MkWordState`, which
we'll keep track of in the type.

A game is finished if either the number of guesses remaining is zero (in which case the player has lost) or the number of letters remaining is zero (in which case the player has won). Listing 9.8 shows how we can represent this in a data type. We can be certain that this captures *only* games that are won or lost, because we're including the number of guesses and the number of letters remaining as arguments to `WordState`.

Listing 9.8 Defining precisely when a game is in a finished state (Hangman.idr)

```
data Finished : Type where
    Lost : (game : WordState 0 (S letters)) -> Finished
    Won : (game : WordState (S guesses) 0) -> Finished
```

A game is lost if zero guesses remain.

A game is won if zero letters remain to be guessed.

The WordState dependent type stores the core data required by the game. By including the core components of the rules—the number of guesses and number of letters—as arguments to WordState, we’ll be able to see exactly how any function that uses a WordState implements the rules of the game.

9.2.2 A top-level game function

We’ll take a top-down approach, defining a top-level function that implements a complete game. The following listing gives the starting point.

Listing 9.9 Top-level game function (Hangman.idr)

```
game : WordState (S guesses) (S letters) -> IO Finished
game st = ?game_rhs
```

?game_rhs needs to read a letter and updates the game state according to the guess.

The type of game states that a game can proceed if there’s at least one guess remaining (S guesses) and at least one letter still to guess (S letters). It returns IO Finished, meaning that it performs interactive actions that produce game data in a Finished state.

To implement the game, you’ll need to read a single letter from the user (any other input is invalid), check whether the letter input by the user is in the target word in the game state, update the game state, and loop if the game isn’t complete. You can update the state in one of the following ways:

- The letter is in the word, in which case you continue the game with the same number of guesses and one fewer letter.
- The letter isn’t in the word, in which case you continue the game with one fewer guess and the same number of letters.

The numbers of guesses and letters used and remaining are explicitly recorded in the game state’s type. The next step, therefore, is to write a function that captures these state updates in its type.

9.2.3 A predicate for validating user input: ValidInput

You can read user input using the IO action getLine:

```
getLine : IO String
```

But because the user is guessing a letter, the only inputs that are valid are those that consist of exactly one character. For precision, you can make the notion of a valid input explicit in a predicate. Because `String` is a primitive, it's difficult to reason about its individual components in a type, so you can use a `List Char` to represent the input in the predicate and convert as necessary:

```
data ValidInput : List Char -> Type where
    Letter : (c : Char) -> ValidInput [c]
```

To check whether a `String` is a valid input, you can write the following function, which returns either a proof that the input is valid, or a proof that it can never be valid, using `Dec`:

```
isValidString : (s : String) -> Dec (ValidInput (unpack s))
```

You'll see the definition of `isValidString` shortly, in section 9.2.5. As an exercise, you can try writing it yourself, beginning with the following helper function:

```
isValidInput : (cs : List Char) -> Dec (ValidInput cs)
```

Then, instead of using `getLine`, write a `readGuess` function that returns the user's guess, along with an instance of a predicate that guarantees the user's guess is a valid input:

```
readGuess : IO (x ** ValidInput x)
```

`readGuess` returns a dependent pair, where the first element is the input read from the console, and the second is a predicate that the input must satisfy.

CONTRACTS ON RETURN VALUES By returning a value paired with a predicate on that value, the type of `readGuess` is expressing a contract that the return value must satisfy.

The following listing gives the definition of `readGuess`.

**Listing 9.10 Read a guess from the console, which is guaranteed to be valid
(Hangman.idr)**

```
readGuess : IO (x ** ValidInput x)
readGuess = do putStr "Guess:"
              x <- getLine
              case isValidString (toUpper x) of
                Yes prf => pure (_ ** prf)
                No contra => do putStrLn "Invalid guess"
                                readGuess
```

Annotations for the code:

- Reads a string from the console**: Points to the `getLine` call.
- If it's a valid input, returns the input and the proof**: Points to the `pure (_ ** prf)` line.
- Converts the input to uppercase so that you can treat inputs as non-case-sensitive**: Points to the `toUpper` call.
- If it's not a valid input, tries again**: Points to the recursive call `readGuess`.

By using `readGuess`, you can be certain that any valid string read from the console is guaranteed to contain exactly one character.

9.2.4 Processing a guess

Processing a guess will return a game state with one type if the guess is correct, and a different type if the guess is incorrect. You can use the `Either` generic type to represent this. Remember that `Either` is a generic type defined in the Prelude that carries a value of two possible types:

```
data Either a b = Left a | Right b
```

`Either` is often used to represent the result of a computation that might fail, carrying information about the error if it does fail. By convention, we use `Left` for the error case and `Right` for the success case.¹ You can think of an incorrect guess as being the error case, so the next listing shows the type you'll use for a function that processes a guess.

Listing 9.11 Type of a function that processes a user's guess (`Hangman.idr`)

```
processGuess : (letter : Char) ->
              WordState (S guesses) (S letters) ->
              Either (WordState guesses (S letters))
                      (WordState (S guesses) letters)
```

This type states that, given a letter and an input game state, it will either produce a new game state where the number of guesses remaining has decreased (the guess was incorrect), or where the number of letters remaining has decreased (the guess was correct.)

TYPES AS ABSTRACT STATE A value of type `WordState` `guesses` `letters` holds *concrete* information about system state, including the exact word to be guessed and exactly which letters are still missing. The type itself expresses *abstract* information about the game state (guesses remaining and number of missing letters), which allows you to express the rules in function types like `processGuess`.

You can implement `processGuess` as follows:

- 1 **Define**—You need to inspect the input game state, so you can case-split on the input state:

```
processGuess : (letter : Char) ->
              WordState (S guesses) (S letters) ->
              Either (WordState guesses (S letters))
                      (WordState (S guesses) letters)
processGuess letter (MkWordState word missing) = ?guess_rhs_1
```

¹ “Right” also being a synonym of “correct.”

- 2 Define**—If the guessed letter is correct, it'll be in the missing vector, and you'll need to remove it. You can use `isElem` to find whether it's there and define the function by case splitting on the result of `isElem`:

```
processGuess letter (MkWordState word missing)
  = case isElem letter missing of
      Yes prf => ?guess_rhs_2
      No contra => ?guess_rhs_3
```

- 3 Refine**—If the letter is in the vector of missing letters, you'll return a new state with one fewer letters to guess. Otherwise, you'll return a new state with one fewer guess available:

```
processGuess letter (MkWordState word missing)
  = case isElem letter missing of
      Yes prf => Right (MkWordState word ?nextVect)
      No contra => Left (MkWordState word missing)
```

You have a hole, `?nextVect`, for the updated vector of missing letters.

- 4 Type, refine**—The type of `?nextVect` shows that you need to remove an element from the vector:

```
word : String
letter : Char
letters : Nat
missing : Vect (S letters) Char
prf : ELEM letter missing
guesses : Nat
-----
nextVect : Vect letters Char
```

You can complete the definition by using `removeElem` to remove `letter` from `missing`, and Idris will find the necessary proof, `prf`, as an auto implicit. This completes the definition:

```
processGuess : (letter : Char) ->
              WordState (S guesses) (S letters) ->
              Either (WordState guesses (S letters))
                      (WordState (S guesses) letters)
processGuess letter (MkWordState word missing)
  = case isElem letter missing of
      Yes prf => Right (MkWordState word (removeElem letter missing))
      No contra => Left (MkWordState word missing)
```

EXPLICIT ASSUMPTIONS In the game state, you assume that the number of letters still to guess is the same as the length of the vector of missing letters. By putting this in the type of `WordState` and the type of `processGuess`, you can be sure that if you ever violate this assumption, your program will no longer compile.

You can complete the definition of `game` using `processGuess` to update the game state as necessary.

9.2.5 Deciding input validity: checking `ValidInput`

Before completing the implementation of `game`, you'll need to complete your implementation of `isValidString`, which decides whether a string entered by the user is a valid input or not.

Listing 9.12 Showing that an input string must be either a valid or invalid input (Hangman.idr)

```
data ValidInput : List Char -> Type where
    Letter : (c : Char) -> ValidInput [c]           ← An input with one character is valid.
                                                       Remember that [c] desugars to c :: [].

isValidNil : ValidInput [] -> Void
isValidNil (Letter _) impossible               ← None of the constructors of ValidInput
                                               have the type ValidInput [].

isValidTwo : ValidInput (x :: y :: xs) -> Void
isValidTwo (Letter _) impossible             ← None of the constructors of
                                               ValidInput have a type of the
                                               form ValidInput (x :: y :: xs).

isValidInput : (cs : List Char) -> Dec (ValidInput cs)
isValidInput [] = No isValidNil
isValidInput (x :: []) = Yes (Letter x)          ← The only valid case,
                                                       because a constructor
                                                       of ValidInput has a
                                                       type of the form
                                                       ValidInput (x :: []).
isValidInput (x :: (y :: xs)) = No isValidTwo

isValidString : (s : String) -> Dec (ValidInput (unpack s))
isValidString s = isValidInput (unpack s)
```

You can test this definition at the REPL by using `>>=` to take the output of `readGuess`, matching on the first component of the dependent pair, and passing it on to `println`:

```
*Hangman> :exec readGuess >>= \(x ** _) => println x
Guess: badguess           ← Entered by the user
Invalid guess            ← Putput by readGuess
Guess:                   ← Entered by the user
Invalid guess            ← Output by readGuess
Guess: f                  ← Entered by the user
['F']                    ← Output by readGuess
```

Now that you have the ability to read input that's guaranteed to be in a valid form, and the ability to process guesses to update the game state, you can complete the top-level game implementation.

9.2.6 Completing the top-level game implementation

The next listing shows one possible (if basic) way to complete the implementation of `game`. It uses `processGuess` to check the user's input and reports whether a guess is correct or incorrect until the player has won or lost.

Listing 9.13 Top-level game function (Hangman.idr)

You'll need to check guesses and letters to find out if the player has won or lost, so bring them into scope.

```
game : WordState (S guesses) (S letters) -> IO Finished
→ game {guesses} {letters} st
  = do (_ ** Letter letter) <- readGuess
    case processGuess letter st of
      Left l => do putStrLn "Wrong!"
      case guesses of
        Z => pure (Lost l)
        S k => game l
      Right r => do putStrLn "Right!"
      case letters of
        Z => pure (Won r)
        S k => game r
```

The guess was wrong;
check whether there
are any guesses left,
and continue if so.

Extract the letter by
pattern-matching on the
ValidInput predicate.

The guess was correct;
check whether there
are any letters left, and
continue if so.

By including the number of guesses and number of letters remaining in the type, you've essentially written an important rule of the game in the type of the guess function. As a result, certain kinds of errors in the implementation can't happen. For example, as long as you always use the guess function to update the game state, you avoid the following potential problems:

- It's impossible to test the wrong variable (guesses or letters) before continuing the game.² Doing so would cause a type error.
- It's impossible to continue a completed game, because there must be at least one guess remaining and at least one letter to guess.
- It's impossible to finish a game prematurely, when there are both missing letters and guesses still available.

The following listing shows a possible implementation of main that sets up a game with a target word, "Test", so requiring a player to guess the letters 'T', 'E', and 'S'.

Listing 9.14 A main program to set up a game (Hangman.idr)

```
main : IO ()
main = do result <- game {guesses=2}
          (MkWordState "Test" ['T', 'E', 'S'])
          case result of
            Lost (MkWordState word missing) =>
              putStrLn ("You lose. The word was " ++ word)
            Won game =>
              putStrLn "You win!"
```

Sets up a
new game
with 2 wrong
guesses
allowed and
a target
word, "Test"

² I really did make this mistake when writing this example!

9.3 Summary

- You can write types that express assumptions about how values relate.
- The `Elem` dependent type is a predicate that expresses that a value must be contained in a vector.
- By passing a predicate as an argument to a function, you can express a contract that the inputs to the function must follow.
- You can write a function to show that a predicate is decidable, using `Dec`.
- Idris will attempt to find values for arguments marked `auto` by expression search.
- You can capture properties of a system's state (such as the rules of a game) in a type.
- Predicates can describe the validity of user input and ensure that the user input is validated when necessary.

10

Views: extending pattern matching

This chapter covers

- Defining views, which describe alternative forms of pattern matching
- Introducing the `with` construct for working with views
- Describing efficient traversals of data structures
- Hiding the representation of data behind views

In type-driven development, our approach to implementing functions is to write a type, define the structure of the function by case splits on its arguments, and refine the function by filling in holes. In the define step in particular, we use the structure of the input types to drive the structure of the function as a whole.

As you learned in chapter 3, when you case-split on a variable, the patterns arise from the variable’s type. Specifically, the patterns arise from the data constructors that can be used to build values of that type. For example, if you case-split on an `items` variable of type `List elem`, the following patterns will arise:

- `[]`—Represents the empty list
- `(x :: xs)`—Represents a non-empty list with `x` as the first element and `xs` as the list of the remaining elements

Pattern matching, therefore, *deconstructs* variables into their components. Often, though, you'll want to deconstruct variables in different ways. For example, you might want to deconstruct an input list into one of the following forms:

- A list of all but the last element, and then the last element
- Two sublists—the first half and the second half of the input

In this chapter, you'll see how to extend the forms of patterns you can use by defining informative data types, called *views*. Views are dependent types that are parameterized by the data you'd like to match, and they give you new ways of observing that data. For example, views can do the following:

- Describe new forms of patterns, such as allowing you to match against the last element of a list rather than the first
- Define alternative ways of traversing data structures, such as traversing a list in reverse, or by repeatedly splitting the list into halves
- Hide complex data representations behind an abstract interface while still supporting pattern-matching on that data

We'll start by looking at how, using views, we can define alternative ways of matching on lists, such as processing the last element first. We'll then look at defining *efficient* traversals of lists and guaranteeing that they *terminate*, and we'll look at some traversals provided by the Idris library. Finally, we'll use views to help with data abstraction by hiding the structure of data in a separate module and matching and traversing the data using a view.

10.1 Defining and using views

When you write a `(x :: xs)` pattern for matching a list, `x` will take the value of the first element of the list and `xs` will take the value of the tail of the list. You saw the following function in chapter 3 to describe how a list was constructed, illustrated again in figure 10.1:

```
describeList : List Int -> String
describeList [] = "Empty"
describeList (x :: xs) = "Non-empty, tail = " ++ show xs
```




Figure 10.1 Matching the pattern `(x :: xs)` for inputs `[1]` and `[2, 3, 4, 5]`

As a result, matching lists using the pattern `(x :: xs)` means that you'll always traverse the list *forwards*, processing `x` first, and then processing the tail, `xs`. But sometimes it's convenient to be able to traverse the list *backwards*, processing the last element first.

You can add a single element to the end of a list using the `++` operator:

```
Idris> [2,3,4] ++ [5]
[2, 3, 4, 5] : List Integer
```

It would be nice to be able to match patterns of the form `(xs ++ [x])`, where `x` takes the value of the last element of the list and `xs` takes the value of the initial segment of the list. That is, you might like to be able to write a function of the following form, as illustrated in figure 10.2:

```
describeListEnd : List Int -> String
describeListEnd [] = "Empty"
describeListEnd (xs ++ [x]) = "Non-empty, initial portion = " ++ show xs
```

This line does not type check




Figure 10.2 Matching the pattern `(xs ++ [x])` for inputs `[1]` and `[2,3,4,5]`

Unfortunately, if you try to implement `describeListEnd`, it will fail with the following error:

```
DLFail.idr:3:19:
When checking left hand side of describeListEnd:
Can't match on describeListEnd (xs ++ [x])
```

The problem is that you're trying to pattern-match on the result of a function, `++`, and in general there's no way for Idris to automatically deduce the inputs to an arbitrary function from its output. You can, however, extend the forms of the patterns you can use by defining informative data types called *views*. In this section, you'll see how to define views, and I'll introduce the `with` construct, which provides a concise notation for defining functions that use views.

10.1.1 Matching the last item in a list

You can't write `describeListEnd` directly because you can't match a pattern of the form `(xs ++ [x])` directly. You can, however, take advantage of dependent pattern matching to deduce that a particular input must have the form `(xs ++ [x])`. You saw

dependent pattern matching in chapter 6, where inspecting the value of one argument (that is, case-splitting on that argument) can tell you about the form of another. Using dependent pattern matching, you can describe alternative patterns for lists.

The following listing shows the `ListLast` dependent type, which describes two possible forms a `List` can take. A list is either empty, `[]`, or constructed from the initial portion of a list and its final element.

**Listing 10.1 The `ListLast` dependent type, which gives alternative patterns for a list
(`DescribeList.idr`)**

```
data ListLast : List a -> Type where
  Empty : ListLast []
  NonEmpty : (xs : List a) -> (x : a) -> ListLast (xs ++ [x])
```

Using `ListLast`, along with dependent pattern matching, you can define `describeListEnd`. You'll start by defining a helper function that takes an input list, `input`, and an extra argument, `form`, that says whether the list is empty or non-empty:

- 1 *Type*—Begin with the type. You'll use `ListLast` to describe the possible forms that `input` can take:

```
describeHelper : (input : List Int) -> (form : ListLast input) -> String
describeHelper input form = ?describeHelper_rhs
```

- 2 *Define*—You want to describe `input`, so you'll need to inspect its form somehow. Previously, when defining a function to inspect the form of an input, you've case-split on that input. Here, the intention of the `form` argument is to tell you more about `input`, so case-split on that instead:

```
describeHelper : (input : List Int) -> (form : ListLast input) -> String
describeHelper [] Empty = ?describeHelper_rhs_1
describeHelper (xs ++ [x]) (NonEmpty xs x) = ?describeHelper_rhs_2
```

Notice that the case split on `form` has told you more about the form of `input`. In particular, the type of `NonEmpty` means that if `form` has the value `NonEmpty xs x`, then `input` *must* have the value `(xs ++ [x])`.

- 3 *Refine*—Complete the definition by describing the patterns as in our initial attempt at `describeListEnd`:

```
describeHelper : (input : List Int) -> ListLast input -> String
describeHelper [] Empty = "Empty"
describeHelper (xs ++ [x]) (NonEmpty xs x)
  = "Non-empty, initial portion = " ++ show xs
```

`ListLast` is a *view* of lists because it provides an alternative means of viewing the data. It's an ordinary data type, though, and in order to use `ListLast` in practice, you'll need to be able to convert an input list, `xs`, into a value of type `ListLast xs`.

10.1.2 Building views: covering functions

The next listing shows `listLast`, which converts an input list, `xs`, into an instance of a view, `ListLast xs`, which gives access to the last element of `xs`.

Listing 10.2 Describing a List in the form ListLast (DescribeList.idr)

```
total
listLast : (xs : List a) -> ListLast xs
listLast [] = Empty
listLast (x :: xs) = case listLast xs of
    Empty => NonEmpty [] x
    NonEmpty ys y => NonEmpty (x :: ys) y
```

You need to traverse the entire list to find the last element, so make a recursive call.

The total flag means Idris will report an error if `listLast` is not fully defined, ensuring that `listLast` will work for every list.

`listLast` is the *covering function* of the `ListLast` view. A covering function of a view describes how to convert a value (in this case the input list) into a view of that value (in this case, an `xs` list and a value `x`, where `xs ++ [x] = input`).

NAMING CONVENTION FOR COVERING FUNCTIONS By convention, covering functions are given the same name as the view type, but with an initial lowercase letter.

Now that you can describe any `List` in the form `ListLast`, you can complete the definition of `describeListEnd`:

```
describeHelper : (input : List Int) -> ListLast input -> String
describeHelper [] Empty = "Empty"
describeHelper (xs ++ [x]) (NonEmpty xs x)
    = "Non-empty, initial portion = " ++ show xs

describeListEnd : List Int -> String
describeListEnd xs = describeHelper xs (listLast xs)
```

This works as we intended in our initial attempt at `describeListEnd`, with the original patterns in `describeHelper`. Given that you're using a different notion of pattern matching than the default, you should expect to have to use some additional notation to explain how to match the pattern `(xs ++ [x])`, but the overall definition still feels quite verbose.

Because dependent pattern matching in this way is a common programming idiom in Idris, there's a construct for expressing extended pattern matching more concisely: the `with` construct.

10.1.3 with blocks: syntax for extended pattern matching

Using views to generate informative patterns like `(xs ++ [x])` can help the readability of functions and increase your confidence in their correctness because the types tell you exactly what form the inputs must take. But these functions can be a little more verbose because you need to create an extra helper function (like `describeHelper`) to

do the necessary pattern matching. The `with` construct provides a notation for using views more concisely.

Using `with` blocks, you can add extra arguments to the left side of a definition, giving you more arguments to case-split. The easiest way to see how this works is by example, so let's take a look at how you can use a `with` block to define `describeListEnd`:

- 1 Type**—Begin with a top-level type that takes a `List Int` and returns a `String`:

```
describeListEnd : List Int -> String
describeListEnd input = ?describeListEnd_rhs
```

- 2 Define**—Press `Ctrl-Alt-W` with the cursor on the line with the hole `?describeListEnd_rhs`. This adds a new pattern clause as follows (it won't yet type-check):

```
describeListEnd : List Int -> String
describeListEnd input with (_)
    describeListEnd input | with_pat = ?describeListEnd_rhs
```

- 3 Define**—The `with` syntax adds a new argument to the left side of the definition. The brackets after `with` give the value of that argument, and the (indented) clause beneath contains the extra argument, after a vertical bar. Here, you'll match on the result of `listLast input`, so replace the `_` with `listLast input`:

```
describeListEnd : List Int -> String
describeListEnd input with (listLast input)
    describeListEnd input | with_pat = ?describeListEnd_rhs
```

- 4 Type**—If you check the type of `?describeListEnd_rhs`, you can see more detail about how `with` works, looking specifically at the type of `with_pat`:

```
input : List Int
with_pat : ListLast input
-----
describeListEnd_rhs : String
```

The value of `with_pat` is the result of `listLast input`. Figure 10.3 shows the components of the syntax for the `with` construct. Note that the scope of the `with` block is managed by indentation.




Figure 10.3 Syntax for the `with` construct

- 5 Define**—Notice that the types of `input` and `with_pat` are exactly the same as the types of the inputs to `describeHelper` earlier. And as in the definition of `describeHelper`, if you case-split on `with_pat`, you'll learn more about the form of `input`:

```
describeListEnd : List Int -> String
describeListEnd input with (listLast input)
  describeListEnd [] | Empty = ?describeListEnd_rhs_1
  describeListEnd (xs ++ [x]) | (NonEmpty xs x)
    = ?describeListEnd_rhs_2
```

- 6 Refine**—Complete the definition as before, with the description of the patterns:

```
describeListEnd : List Int -> String
describeListEnd input with (listLast input)
  describeListEnd [] | Empty = "Empty"
  describeListEnd (xs ++ [x]) | (NonEmpty xs x)
    = "Non-empty, initial portion = " ++ show xs
```

Effectively, the `with` construct has allowed you to use an intermediate pattern match, on the result of `listLast input`, without needing to define a separate function like `describeHelper`. In turn, matching on the result of `listLast input` gives you more-informative patterns for `input`.

THE DIFFERENCE BETWEEN WITH AND CASE The purpose of the `with` construct is similar to that of a `case` block in that it allows matching on intermediate results. There's one important difference, however: `with` introduces a new pattern to match on the *left side* of a definition. As a result, you can use dependent pattern matching directly. In `describeListEnd`, for example, matching on the result of `listLast input` told you about the form `input` must take.

You can't use just any expression in a pattern because it isn't possible, in general, to decide what the inputs to a function must be, given only its result. Idris therefore allows patterns only when it *can* deduce those inputs, which is in the following cases:

- The pattern consists of a data constructor applied to some arguments. The arguments must also be valid patterns.
- The value of the pattern is *forced* by some other valid pattern. In the case of `describeListEnd`, the value of the pattern `(xs ++ [x])` is forced by the valid pattern `NonEmpty xs x`.

10.1.4 Example: reversing a list using a view

Once you have the ability to pattern-match in different ways using views, you can traverse data structures in new ways. Rather than always traversing a list from left to right, for example, you can use `listLast` to traverse a list from right to left, inspecting the last element first.

You can reverse a list in this way:

- To reverse an empty list `[]`, return `[]`.
- To reverse a list in the form `xs ++ [x]`, reverse `xs` and then add `x` to the front of the list.

You can implement this algorithm fairly directly using the `ListLast` view:

- 1 Type**—Call the function `myReverse` because there's already a `reverse` function in the Prelude:

```
myReverse : List a -> List a
myReverse input = ?myReverse_rhs
```

- 2 Define**—You'll define the function by inspecting the last element of the input first, so you can use `listLast` to match on a value of type `ListLast` `input`. Press Ctrl-Alt-W to insert a `with` block, and add `listLast input` as the value to inspect:

```
myReverse : List a -> List a
myReverse input with (listLast input)
  myReverse input | with_pat = ?myReverse_rhs
```

- 3 Define**—Next, case-split on `with_pat` to give the relevant patterns for `input`:

```
myReverse : List a -> List a
myReverse input with (listLast input)
  myReverse [] | Empty = ?myReverse_rhs_1
  myReverse (xs ++ [x]) | (NonEmpty xs x) = ?myReverse_rhs_2
```

- 4 Refine**—Finally, you can complete the definition as follows:

```
myReverse : List a -> List a
myReverse input with (listLast input)
  myReverse [] | Empty = []
  myReverse (xs ++ [x]) | (NonEmpty xs x) = x :: myReverse xs
```

This is a fairly direct implementation of the algorithm, traversing the list in reverse and constructing a new list by adding the *last* item of the input as the *first* item of the result. There are, nevertheless, two problems:

- The definition is inefficient because it constructs `ListLast` `input` on every recursive call, and constructing `ListLast` `input` requires traversing `input`.
- Idris can't decide whether the definition is total or not:

```
*Reverse> :total myReverse
Main.myReverse is possibly not total due to:
  possibly not total due to recursive path:
    with block in Main.myReverse, with block in Main.myReverse
```

See the sidebar for a brief discussion on totality checking in Idris.

The first problem is important to address, because it's possible to write `myReverse` in linear time, traversing the input list only once. The second problem is important from the type-driven development perspective: as I discussed in chapter 1, if Idris can determine that a function is total, you have a strong guarantee that the type accurately describes what the function will do. If not, you only have a guarantee that the function will produce a value of the given type if it terminates without crashing. Furthermore, if Idris can't determine that a function is total, it can't determine that any functions that call it are total, either.

We'll look at how to address each of these problems, with only minor alterations to `myReverse` itself, in section 10.2.

Totality checking

Idris tries to decide whether a function is always terminating by checking two things:

- There must be patterns for all possible well-typed inputs.
- When there is a recursive call (or a sequence of mutually recursive calls), there must be a *decreasing argument* that converges toward a base case.

To determine which arguments are decreasing, Idris looks at the patterns for the inputs in a definition. If a pattern is in the form of a data constructor, Idris considers the arguments in that pattern to be *smaller* than the input. In `myReverse`, for example, Idris doesn't consider `xs` to be smaller than `(xs ++ [x])`, because `(xs ++ [x])` is not in the form of a data constructor.

This restriction keeps the concept of *decreasing argument* simple for Idris to check. In general, Idris can't tell whether the inputs to a function will always be smaller than the result.

As you'll see in section 10.2, you can work around this restriction by defining *recursive views*.

10.1.5 Example: merge sort

Views allow you to describe matching on data structures any way you like, with as many patterns as you like, provided you can implement a covering function for the view. As a second example of a view, we'll implement the merge sort algorithm on Lists.

Merge sort, at a high level, works as follows:

- If the input is an empty list, return an empty list.
- If the input has one element, it's already sorted, so return the input.
- In all other cases, split the list into two halves (differing in size by at most one), recursively sort those halves, and then merge the two sorted halves into a sorted list (see figure 10.4).




Figure 10.4 Sorting a list using merge sort: split the list into two halves, sort the halves, and then merge the sorted halves back together

If you have two sorted lists, you can merge them together using the `merge` function defined in the Prelude:

```
Idris> :doc merge
Prelude.List.merge : Ord a => List a -> List a
    Merge two sorted lists using the default ordering for the type
    of their elements.
```

Note that it has a generic type and requires the `Ord` interface to be implemented for the element type of the list.

Assuming that the two input lists are sorted, `merge` will produce a sorted list of the elements in the input lists. For example, the lists in figure 10.4 would be merged as follows:

```
Idris> merge [1,3,5] [2,4,6]
[1, 2, 3, 4, 5, 6] : List Integer
```

Listing 10.3 shows how you might ideally like to write a `mergeSort` function that sorts a list using the merge sort algorithm. Unfortunately, as it stands, this won't work because `(lefts ++ rights)` isn't a valid pattern.

Listing 10.3 Initial attempt at `mergeSort` with an invalid pattern (MergeSort.idr)

```
mergeSort : Ord a => List a -> List a
mergeSort [] = []
mergeSort [x] = [x]
mergeSort (lefts ++ rights)
    = merge (mergeSort lefts) (mergeSort rights)
```

This pattern isn't valid because `++` isn't a data constructor, but you'd like to extract the left and right halves of the input.

An empty list is already sorted.

A singleton list is already sorted.

Recursively sorts the left and right halves of the list, and then merges the results into a complete sorted list

Given an input list matched against a pattern, `(lefts ++ rights)`, Idris can't in general deduce what `lefts` and `rights` must be; indeed, there could be several reasonable possibilities if the input list has one or more elements. Here are a couple of examples:

- `[1]` could match against `(lefts ++ rights)` with `lefts` as `[1]` and `rights` as `[]`, or `lefts` as `[]` and `rights` as `[1]`.
- `[1,2,3]` could match against `(lefts ++ rights)` with `lefts` as `[1]` and `rights` as `[2,3]`, or `lefts` as `[1,2]` and `rights` as `[3]`, among many other possibilities.

To match the patterns we want, as shown in listing 10.3, you'll need to create a view of lists, `SplitList`, that gives the patterns you want. The following listing shows the definition of `SplitList` and gives the type for its covering function, `splitList`.

Listing 10.4 A view of lists that gives patterns for empty lists, singleton lists, and concatenating lists (MergeSort.idr)

```
data SplitList : List a -> Type where
  SplitNil : SplitList []
  SplitOne : SplitList [x]
  SplitPair : (lefts : List a) -> (rights : List a) ->
    SplitList (lefts ++ rights)

splitList : (input : List a) -> SplitList input
```

I'll give a definition of the covering function, `splitList`, shortly. For now, note that as long as the implementation of `splitList` is total, you can be sure from its type that it gives valid patterns for empty lists, singleton lists, or concatenations of two lists. You don't, however, have any guarantees in the type about *how* a list is split into pairs; in this case, you need to rely on the implementation to ensure that `lefts` and `rights` differ in size by at most one.

PRECISION OF SPLITPAIR In principle, you could make the type of `SplitPair` more precise and carry a proof that `lefts` and `rights` differ in size by at most one. In fact, the Idris library module `Data.List.Views` exports such a view, called `SplitBalanced`.

You can implement `mergeSort` using the `SplitList` view as follows:

- 1 *Type*—Begin with the type and a skeleton definition:

```
mergeSort : Ord a => List a -> List a
mergeSort input = ?mergeSort_rhs
```

- 2 *Define*—To get the patterns you want, you'll need to use the `SplitList` view you've just created. Add a `with` block and use the `splitList` covering function:

```
mergeSort : Ord a => List a -> List a
mergeSort input with (splitList input)
  mergeSort input | with_pat = ?mergeSort_rhs
```

- 3 *Define*—If you case-split on `with_pat`, you'll get appropriate patterns for `input` arising from the types of the constructors to `SplitList`:

```
mergeSort : Ord a => List a -> List a
mergeSort input with (splitList input)
  mergeSort [] | SplitNil = ?mergeSort_rhs_1
  mergeSort [x] | SplitOne = ?mergeSort_rhs_2
  mergeSort (lefts ++ rights) | (SplitPair lefts rights) = ?mergeSort_rhs_3
```

- 4 *Refine*—You can complete the definition by filling in the right sides for each pattern, directly following the high-level description of the merge sort algorithm I gave at the start of this section:

```
mergeSort : Ord a => List a -> List a
mergeSort input with (splitList input)
  mergeSort [] | SplitNil = []
  mergeSort [x] | SplitOne = [x]
```

```
mergeSort (lefts ++ rights) | (SplitPair lefts rights)
  = merge (mergeSort lefts) (mergeSort rights)
```

Before you can test `mergeSort`, you'll need to implement the covering function `splitList`. The next listing gives a definition of `splitList` that returns a description of the empty list pattern, a singleton list pattern, or a pattern consisting of the concatenation of two lists, where those lists differ in length by at most one.

Listing 10.5 Defining a covering function for `SplitList` (`MergeSort.idr`)

<p>Adds the “total” flag so Idris checks that the definition is total</p>	<p>Uses a second reference to the list as a counter, which steps through the list two elements at a time</p>	<p>The input is empty, so return <code>SplitNil</code>, which gives a pattern for the empty list.</p>
<pre>→ total splitList : (input : List a) -> SplitList input splitList input = splitListHelp input input ←</pre>		
<pre>where splitListHelp : List a -> (input : List a) -> SplitList input splitListHelp _ [] = SplitNil ← splitListHelp _ [x] = SplitOne splitListHelp (_ :: _ :: counter) (item :: items) = case splitListHelp counter items of SplitNil => SplitOne SplitOne {x} => SplitPair [item] [x] SplitPair lefts rights => SplitPair (item :: lefts) rights splitListHelp _ items = SplitPair [] items ←</pre>		
<p>The input has one element, so return <code>SplitOne</code>, which gives a pattern for the singleton list.</p>	<p>There are fewer than two elements in the counter list, so put the remaining items in the second list.</p>	<p>Recursively splits the list, moving the counter along two places in the list, and then adds the first element, item, to the first half</p>

You build two (approximately) equally sized lists by using a second reference to the input list as a *counter* in a helper function, `splitListHelp`, as follows:

- On each recursive call, the counter steps through *two* elements of the list. The pattern `(_ :: _ :: counter)` matches any list with at least two elements, where `counter` is a list containing all but the first two elements.
- When the counter reaches the end of the list (that is, fewer than two elements remain), you must have traversed *half* of the input.

You can now try `splitList` and `mergeSort` at the REPL:

```
*MergeSort> splitList [1]
SplitOne : SplitList [1]

*MergeSort> splitList [1,2,3,4,5]
SplitPair [1, 2] [3, 4, 5] : SplitList [1, 2, 3, 4, 5]

*MergeSort> mergeSort [3,2,1]
[1, 2, 3] : List Integer

*MergeSort> mergeSort [5,1,4,3,2,6,8,7,9]
[1, 2, 3, 4, 5, 6, 7, 8, 9] : List Integer
```

By defining a view that gives possible cases for the input list in terms of how it can be split in half, you can write a definition of `mergeSort` that directly implements the high-level description of the algorithm.

Like `myReverse`, however, Idris can't tell whether `mergeSort` is total:

```
*MergeSort> :total mergeSort
Main.mergeSort is possibly not total due to:
  possibly not total due to recursive path:
    with block in Main.mergeSort, with block in Main.mergeSort
```

Again, the problem is that Idris can't tell that the recursive calls are guaranteed to be on smaller lists than the original input. In the next section, you'll see how to address this problem by defining recursive views that describe the recursive structure of a function, as well as the patterns that define a function.

Exercises



- The `TakeN` view allows traversal of a list several elements at a time:

```
data TakeN : List a -> Type where
  Fewer : TakeN xs
  Exact : (n_xs : List a) -> TakeN (n_xs ++ rest)

takeN : (n : Nat) -> (xs : List a) -> TakeN xs
```

The `Fewer` constructor covers the case where there are fewer than `n` elements. Define the covering function `takeN`.

To check that your definition works, you should be able to run the following function, which groups lists into sublists with `n` elements each:

```
groupByN : (n : Nat) -> (xs : List a) -> List (List a)
groupByN n xs with (takeN n xs)
  groupByN n xs | Fewer = [xs]
  groupByN n (n_xs ++ rest) | (Exact n_xs) = n_xs :: groupByN n rest
```

Here's an example:

```
*ex_10_1> groupByN 3 [1..10]
[[1, 2, 3], [4, 5, 6], [7, 8, 9], [10]] : List (List Integer)
```

- Use `TakeN` to define a function that splits a list into two halves by calculating its length:

```
halves : List a -> (List a, List a)
```

If you have implemented this correctly, you should see the following:

```
*ex_10_1> halves [1..10]
([1, 2, 3, 4, 5], [6, 7, 8, 9, 10]) : (List Integer, List Integer)

*ex_10_1> halves [1]
([], [1]) : (List Integer, List Integer)
```

Hint: Use `div` for dividing a `Nat`.

10.2 Recursive views: termination and efficiency

The purpose of views is to give us new ways to match on data, using the `with` construct for a more concise syntax. When you write a function with a view, you use the following components:

- The original input data
- A *view* of the input data, where the view data type is parameterized by the input data
- A *covering function* for the view, which constructs an instance of the view for the input data

Then, using the `with` construct, you can pattern-match on the view. Dependent pattern matching gives you informative patterns for the original data. You've seen two examples of this in the previous section: reversing a list, and splitting a list in half for merge sort. In both cases, however, Idris couldn't tell that the resulting function was total. Moreover, when you reverse a list, the resulting function is inefficient because it has to reconstruct the view on every recursive call.

In this section, we'll look at how to solve both these problems by defining *recursive* views for describing traversals of data structures. Furthermore, once we've defined a view, we can reuse it for any function that uses the same recursion pattern. The Idris library provides a number of useful views for traversals of data structures, and we'll look at some example functions with some of these. First, though, we'll improve the definition of `myReverse`.

10.2.1 “Snoc” lists: traversing a list in reverse

A *snoc list* is a list where elements are added to the end of the list, rather than the beginning. We can define them as follows, as a generic data type:

```
data SnocList ty = Empty | Snoc (SnocList ty) ty
```

SNOC LIST TERMINOLOGY The name *snoc list* arises because the traditional name (originating with Lisp) for the operator that adds an element to the beginning of a list is *cons*. Therefore, the name for the operator that adds an element to the end of a list is *snoc*.

Using `SnocList`, you traverse the elements in reverse order, because the element (of type `ty`) appears after the list (of type `SnocList ty`). You can easily produce a `List` from a `SnocList` where the elements are in reverse order:

```
reverseSnoc : SnocList ty -> List ty
reverseSnoc Empty = []
reverseSnoc (Snoc xs x) = x :: reverseSnoc xs
```

You can show the relationship between `SnocList` and `List` more precisely by parameterizing `SnocList` over the equivalent `List`. The following listing shows how to define `SnocList` in this way.

Listing 10.6 The SnocList type, parameterized over the equivalent List (SnocList.idr)

```
data SnocList : List a -> Type where
  Empty : SnocList []
  Snoc : (rec : SnocList xs) -> SnocList (xs ++ [x])
  snocList : (xs : List a) -> SnocList xs
```

Empty is equivalent to the list represented by [].

Given a list, xs, snocList builds a SnocList that is equivalent to xs.

Given a SnocList equivalent to a list, xs, and an implicit value, x, Snoc builds a SnocList equivalent to xs ++ [x].

This is very similar in structure to ListLast, which you defined in the previous section. The difference is that Snoc takes a *recursive* argument of type SnocList xs.

Here, SnocList is a *recursive view*, and snocList is its covering function. We'll come to the definition of snocList shortly; first, let's see how you can use this to implement myReverse:

- 1 *Type*—Begin by defining a helper function that takes a list, input, and its equivalent SnocList:

```
myReverseHelper : (input : List a) -> SnocList input -> List a
myReverseHelper input snoc = ?myReverseHelper_rhs
```

- 2 *Define*—If you case-split on snoc, you'll get corresponding patterns for input:

```
myReverseHelper : (input : List a) -> SnocList input -> List a
myReverseHelper [] Empty = ?myReverseHelper_rhs_1
myReverseHelper (xs ++ [x]) (Snoc rec) = ?myReverseHelper_rhs_2
```

- 3 *Refine*—If the input is empty, you'll return the empty list:

```
myReverseHelper : (input : List a) -> SnocList input -> List a
myReverseHelper [] Empty = []
myReverseHelper (xs ++ [x]) (Snoc rec) = ?myReverseHelper_rhs_2
```

- 4 *Refine, type*—Otherwise, you'll recursively reverse xs and add the item x to the front:

```
myReverseHelper : (input : List a) -> SnocList input -> List a
myReverseHelper [] Empty = []
myReverseHelper (xs ++ [x]) (Snoc rec) = x :: myReverseHelper xs ?snocrec
```

There's still a hole, ?snocrec, for the second argument in the recursive call. If you inspect it, you'll see that you need the SnocList that represents xs:

```
a : Type
xs : List a
rec : SnocList xs
x : a
-----
snocrec : SnocList xs
```

- 5 Refine**—Fortunately, you already have a value, `rec`, of type `SnocList xs`, so you can use it directly to complete the definition:

```
myReverseHelper : (input : List a) -> SnocList input -> List a
myReverseHelper [] Empty = []
myReverseHelper (xs ++ [x]) (Snoc rec) = x :: myReverseHelper xs rec
```

- 6 Define**—Finally, you can define `myReverse` by building a `SnocList` and calling `myReverseHelper`:

```
myReverse : List a -> List a
myReverse input = myReverseHelper input (snocList input)
```

You can't test this yet because you haven't implemented `snocList`, but for now notice how this contrasts with the implementation of `myReverse` in section 10.1.4, using `ListLast`. The similarity is that you find the patterns for the input list by matching on a view of `input`. The difference is that the view is recursive, meaning that you don't have to rebuild the view on each recursive call; you already have access to it.

The definition of `myReverseHelper` is total, because the `SnocList` argument is decreasing on each recursive call:

```
*SnocList> :total myReverseHelper
Main.myReverseHelper is Total
```

It now remains to implement `snocList`. As long as you can implement `snocList` by traversing the list only once, you'll have an implementation of `myReverse` that runs in linear time. The following listing shows an implementation of `snocList` that traverses the list only once, using a helper function with an accumulator to build the `SnocList` by adding one element at a time.

Listing 10.7 Implementing the covering function `snocList`

```
snocListHelp : (snoc : SnocList input) -> (rest : List a) ->
              SnocList (input ++ rest)
→ snocListHelp {input} snoc [] = rewrite appendNilRightNeutral input in snoc
  snocListHelp {input} snoc (x :: xs)
    = rewrite appendAssociative input [x] xs in
      snocListHelp (Snoc snoc {x}) xs
  snocList : (xs : List a) -> SnocList xs
  snocList xs = snocListHelp Empty xs
```

Appends an empty list to a
SnocList representing input

Initializes the SnocList as Empty
and then calls `snocListHelp` to
add `xs` an element at a time

Appends an element, `x`, to a
SnocList representing input,
and then appends the
remaining elements, `xs`

The definition of `snocList` is slightly tricky, involving the `rewrite` construct (which you saw in chapter 8) to get the types in the correct form for building the `SnocList`. You'll rewrite with the following functions from the Prelude:

```
appendNilRightNeutral : (l : List a) -> l ++ [] = l
appendAssociative : (l : List a) -> (c : List a) -> (r : List a) ->
    l ++ (c ++ r) = (l ++ c) ++ r
```

As with any complex definition, it's a good idea to try to understand it by replacing subexpressions of the definition with holes, and seeing what the types of those holes are. In this case, it's also useful to remove the rewrite constructs and replace them with holes, to see how the types need rewriting:

```
snocListHelp : SnocList input -> (xs : List a) -> SnocList (input ++ xs)
snocListHelp {input} snoc [] = ?rewriteNil snoc
snocListHelp {input} snoc (x :: xs)
    = ?rewriteCons (snocListHelp (Snoc snoc {x}) xs)
```

You should see the following types for `rewriteNil` and `rewriteCons`:

```
rewriteNil : SnocList input -> SnocList (input ++ [])
rewriteCons : SnocList ((input ++ [x]) ++ xs) -> SnocList (input ++ x :: xs)
```

The good news is that once you've defined `snocList`, you can reuse it for any function that needs to traverse a list in reverse. Furthermore, as you'll see shortly, a `SnocList` view is also defined in the Idris library, along with several others.

10.2.2 Recursive views and the `with` construct

You now have an implementation of `myReverse` that runs in linear time, because it traverses the list once to build the `SnocList` view and then traverses the `SnocList` view once to build the reversed list. You can also confirm that Idris believes it's total:

```
*SnocList> :total myReverse
Main.myReverse is Total
```

The resulting definition isn't quite as concise as the previous definition of `myReverse`, however, because it doesn't use the `with` construct:

```
myReverseHelper : (input : List a) -> SnocList input -> List a
myReverseHelper [] Empty = []
myReverseHelper (xs ++ [x]) (Snoc rec) = x :: myReverseHelper xs rec

myReverse : List a -> List a
myReverse input = myReverseHelper input (snocList input)
```

Let's see what happens if you try to do so:

- 1 Type, define**—Begin with the following skeleton definition:

```
myReverse : List a -> List a
myReverse input with (snocList input)
    myReverse input | with_pat = ?myReverse_rhs
```

- 2 Define**—You can write the function by case splitting on `with_pat` to get the possible patterns for `input`:

```
myReverse : List a -> List a
myReverse input with (snocList input)
```

```
myReverse [] | Empty = ?myReverse_rhs_1
myReverse (xs ++ [x]) | (Snoc rec) = ?myReverse_rhs_2
```

- 3 *Refine*—Fill in the right side as before:

```
myReverse : List a -> List a
myReverse input with (snocList input)
  myReverse [] | Empty = []
  myReverse (xs ++ [x]) | (Snoc rec) = x :: myReverse xs
```

Unfortunately, this calls the *top-level* reverse function, which rebuilds the view using `snocList` input, so you have the same problem as before:

```
*SnocList> :total myReverse
Main.myReverse is possibly not total due to:
  possibly not total due to recursive path:
    with block in Main.myReverse, with block in Main.myReverse
```

- 4 *Refine*—Instead, when you make the recursive call, you can make the call directly to the `with` block, using the `|` notation on the right side:

```
myReverse : List a -> List a
myReverse input with (snocList input)
  myReverse [] | Empty = []
  myReverse (xs ++ [x]) | (Snoc rec) = x :: myReverse xs | rec
```

The call to `myReverse xs | rec` recursively calls `myReverse`, but bypasses the construction of `snocList` input and uses `rec` directly. The resulting definition is total, building the `SnocList` representation of `input`, and traversing that:

```
*SnocList> :total myReverse
Main.myReverse is Total
```

This also has the effect of making `myReverse` run in linear time.

In practice, when you use the `with` construct, Idris introduces a new function definition for the body of the `with` block, like the definition of `myReverseHelper` that you implemented manually earlier.

When you write `myReverse xs | rec`, this is equivalent to writing `myReverseHelper xs rec` in the earlier definition. But by using the `with` construct instead, Idris generates an appropriate type for the helper function.

By using the `with` construct, you can pattern-match and traverse data structures in different ways, with the structure of the matching and traversal given by the type of a view. Moreover, because the views themselves are data structures, Idris can be sure that functions that traverse views are total.

10.2.3 Traversing multiple arguments: nested with blocks

When you write pattern-matching definitions, you often want to match on several inputs at once. So far, using the `with` construct, you've only been matching one value. But like any language construct, `with` blocks can be nested.

To see how this works, let's define an `isSuffix` function:

```
isSuffix : Eq a => List a -> List a -> Bool
```

The result of `isSuffix` should be True if the list in the first argument is a suffix of the second argument. For example:

```
*IsSuffix> isSuffix [7,8,9,10] [1..10]
True : Bool
```

```
*IsSuffix> isSuffix [7,8,9] [1..10]
False : Bool
```

You can define this function by traversing both lists in reverse, taking the following steps:

- 1 Define**—Start with the skeleton definition:

```
isSuffix : Eq a => List a -> List a -> Bool
isSuffix input1 input2 = ?isSuffix_rhs
```

- 2 Define, refine**—Next, match on the first input, using `snocList` so that you process the *last* element first:

```
isSuffix : Eq a => List a -> List a -> Bool
isSuffix input1 input2 with (snocList input1)
  isSuffix [] input2 | Empty = ?isSuffix_rhs_1
  isSuffix (xs ++ [x]) input2 | (Snoc rec) = ?isSuffix_rhs_2
```

You can rename `rec` to `xsrec`, to indicate that it's a recursive view of `xs` when reversed. Then, if the first list is empty, it's trivially a suffix of the second list:

```
isSuffix : Eq a => List a -> List a -> Bool
isSuffix input1 input2 with (snocList input1)
  isSuffix [] input2 | Empty = True
  isSuffix (xs ++ [x]) input2 | (Snoc xsrec) = ?isSuffix_rhs_2
```

- 3 Define**—Next, match on the second input, again using `snocList` to process the last element first. With the cursor over `?isSuffix_rhs_2`, press Ctrl-Alt-W to add a nested `with` block:

```
isSuffix : Eq a => List a -> List a -> Bool
isSuffix input1 input2 with (snocList input1)
  isSuffix [] input2 | Empty = True
  isSuffix (xs ++ [x]) input2 | (Snoc xsrec) with (snocList input2)
    isSuffix (xs ++ [x]) [] | (Snoc xsrec) | Empty = ?isSuffix_rhs_2
    isSuffix (xs ++ [x]) (ys ++ [y]) | (Snoc xsrec) | (Snoc ysrec)
      = ?isSuffix_rhs_3
```

- 4 Refine**—A non-empty list can't be a suffix of an empty list, and if the last two elements of a list are equal, you'll recursively check the rest of the list:

```
isSuffix : Eq a => List a -> List a -> Bool
isSuffix input1 input2 with (snocList input1)
  isSuffix [] input2 | Empty = True
  isSuffix (xs ++ [x]) input2 | (Snoc rec) with (snocList input2)
```

```

isSuffix (xs ++ [x]) [] | (Snoc rec) | Empty = False
isSuffix (xs ++ [x]) (ys ++ [y]) | (Snoc rec) | (Snoc z)
    = if x == y then isSuffix xs ys | xsrec | ysrec
      else False

```

Note that when you call `isSuffix` recursively, you pass *both* of the recursive view arguments, `xsrec` and `ysrec`, to save recomputing them unnecessarily.

You can confirm that this definition is total by asking Idris at the REPL:

```
*IsSuffix> :total isSuffix
Main.isSuffix is Total
```

10.2.4 More traversals: Data.List.Views

In order to help you write total functions, the Idris library provides a number of views for traversing data structures. The `Data.List.Views` module provides several, including the `SnocList` view you've just seen.

For example, listing 10.8 shows the `SplitRec` view, which allows you to recursively traverse a list, processing one half at a time. This is similar to the `SplitList` view you saw in section 10.1.5, but with recursive traversals on the halves of the list.

Listing 10.8 The SplitRec view from Data.List.Views

```

data SplitRec : List a -> Type where
  SplitRecNil : SplitRec []
  SplitRecOne : SplitRec [x]
  SplitRecPair : (lrec : Lazy (SplitRec lefts)) ->           ← The Lazy annotation
                                                               means that the recursive
                                                               argument will only be
                                                               constructed when needed.
                                                               ←
  (rrec : Lazy (SplitRec rights)) ->
  SplitRec (lefts ++ rights)

total splitRec : (xs : List a) -> SplitRec xs           ← SplitRecPair constructs
                                                               two halves of the input
                                                               list, recursively.

                                                               ← The covering function is total, so you can
                                                               use the view to write total functions.

```

The Lazy generic type

The `Lazy` type allows you to postpone a computation until the result is needed. For example, a variable of type `Lazy Int` is a computation that, when evaluated, will produce a value of type `Int`. Idris has the following two functions built-in:

```

Delay : a -> Lazy a
Force : Lazy a -> a

```

When type-checking, Idris will insert applications of `Delay` and `Force` implicitly, as required. Therefore, in practice, you can treat `Lazy` as an annotation that states that a variable will only be evaluated when its result is required. You'll see the definition of `Lazy` in more detail in chapter 11.

You can use `SplitRec` to reimplement `mergeSort` from section 10.1.5 as a total function. The following listing shows our starting point.

Listing 10.9 Starting point for a total implementation of `mergeSort`, using `SplitRec` (`MergeSortView.idr`)

```
import Data.List.Views
mergeSort : Ord a => List a -> List a
mergeSort input = ?mergeSort_rhs
```

Imports `Data.List.Views`
to get access to `SplitRec`

You can implement `mergeSort` using the `SplitRec` view by taking the following steps:

- 1 *Define*—Begin by adding a `with` block, to say you'd like to write the function by using the `SplitRec` view:

```
mergeSort : Ord a => List a -> List a
mergeSort input with (splitRec input)
  mergeSort [] | SplitRecNil = ?mergeSort_rhs_1
  mergeSort [x] | SplitRecOne = ?mergeSort_rhs_2
  mergeSort (lefts ++ rights) | (SplitRecPair lrec rrec)
    = ?mergeSort_rhs_3
```

- 2 *Refine*—The inputs `[]` and `[x]` are already sorted:

```
mergeSort : Ord a => List a -> List a
mergeSort input with (splitRec input)
  mergeSort [] | SplitRecNil = []
  mergeSort [x] | SplitRecOne = [x]
  mergeSort (lefts ++ rights) | (SplitRecPair lrec rrec)
    = ?mergeSort_rhs_3
```

- 3 *Refine*—For the `(lefts ++ rights)` case, you can sort `lefts` and `rights` recursively, and then merge the results:

```
mergeSort : Ord a => List a -> List a
mergeSort input with (splitRec input)
  mergeSort [] | SplitRecNil = []
  mergeSort [x] | SplitRecOne = [x]
  mergeSort (lefts ++ rights) | (SplitRecPair lrec rrec)
    = merge (mergeSort lefts | lrec)
      (mergeSort rights | rrec)
```

The `|` says that, in the recursive calls, you want to bypass constructing the view, because you already have appropriate views for `lefts` and `rights`.

You can confirm that the new definition of `mergeSort` is total, and test it on some examples:

```
*MergeSortView> :total mergeSort
Main.mergeSort is Total

*MergeSortView> mergeSort [3,2,1]
[1, 2, 3] : List Integer

*MergeSortView> mergeSort [5,1,4,3,2,6,8,7,9]
[1, 2, 3, 4, 5, 6, 7, 8, 9] : List Integer
```

Exercises



These exercises use views defined in the Idris library in the modules `Data.List.Views`, `Data.Vect.Views`, and `Data.Nat.Views`. For each view mentioned in the exercises, use `:doc` to find out about the view and its covering function.

For each of these exercises, make sure Idris considers your solution to be total.

- 1** Implement an `equalSuffix` function using the `SnocList` view defined in `Data.List.Views`. It should have the following type:

```
equalSuffix : Eq a => List a -> List a -> List a
```

Its behavior should be to return the maximum equal suffix of the two input lists. Here's an example:

```
*ex_10_2> equalSuffix [1,2,4,5] [1..5]
[4, 5] : List Integer

*ex_10_2> equalSuffix [1,2,4,5,6] [1..5]
[] : List Integer

*ex_10_2> equalSuffix [1,2,4,5,6] [1..6]
[4, 5, 6] : List Integer
```

- 2** Implement `mergeSort` for vectors, using the `SplitRec` view defined in `Data.Vect.Views`.
- 3** Write a `toBinary` function that converts a `Nat` to a `String` containing a binary representation of the `Nat`. You should use the `HalfRec` view defined in `Data.Nat.Views`.

If you have a correct implementation, you should see this:

```
*ex_10_2> toBinary 42
"101010" : String

*ex_10_2> toBinary 94
"1011110" : String
```

Hint: It's okay to return an empty string if the input is `Z`.

- 4** Write a `palindrome` function that returns whether a list is the same when traversed forwards and backwards, using the `vList` view defined in `Data.List.Views`.

If you have a correct implementation, you should see the following:

```
*ex_10_2> palindrome (unpack "abccba")
True : Bool

*ex_10_2> palindrome (unpack "abcba")
True : Bool

*ex_10_2> palindrome (unpack "abcb")
False : Bool
```

Hint: The `vList` view allows you to traverse a list in linear time, processing the first and last elements simultaneously and recursing on the middle of the list.

10.3 Data abstraction: hiding the structure of data using views

The views you've seen so far in this chapter allow you to inspect and traverse data structures in ways beyond the default pattern matching, focusing in particular on `List`. In a sense, views allow you to describe alternative interfaces for building pattern-matching definitions:

- Using `SnocList`, you can see how a list is constructed using `[]` and by adding a single element with `++`, rather than using `[]` and `::::`.
- Using `SplitRec`, you can see how a list is constructed as an empty list, a singleton list, or a concatenation of a pair of lists.

That is, you can find out how a value was constructed by looking at the *view*, rather than by looking directly at the *constructors* of that value. In fact, you often won't need to know what the data constructors of a value are to be able to use a view of a value.

With this in mind, one use of views in practice is to hide the *representation* of data in a module, while still allowing interactive type-driven development of functions that use that data, by case splitting on a view of that data.

THE ORIGIN OF VIEWS The idea of views was proposed by Philip Wadler for Haskell in 1987, in his paper "Views: a way for pattern matching to cohabit with data abstraction." The example in this section is in the spirit of Wadler's paper, which contains several other examples of using views in practice. Views as a programming idiom, using dependent types and a notation similar to the `with` notation in Idris, was later proposed by Conor McBride and James McKinna in their 2004 paper, "The view from the left."

To conclude this chapter, we'll look at this idea in action. We'll revisit the data store example that we implemented in chapters 4 and 6, hide the representation of the data in an Idris module, and export only the following:

- Schema descriptions, explaining the form of data in the store
- A function for initializing a data store, `empty`
- A function for adding an entry to the store, `addToString`
- A view for traversing the contents of the store, `StoreView`, along with its covering function, `storeView`

None of these require users of the module to know anything about the structure of the store itself or the structure of the data contained within it.

Before you implement a module and export the relevant definitions, though, we'll need to discuss briefly how Idris supports data abstraction in modules.

10.3.1 Digression: modules in Idris

With the exception of a small example in chapter 2, the programs you've written in this book have been self-contained in a single main module. As you write larger applications,

however, you'll need a way to organize code into smaller compilation units, and to control which definitions are exported from those units.

The following listing shows a small Idris module that defines a Shape type and exports it, along with its data constructors and a function to calculate the area of a shape.

Listing 10.10 A module defining shapes and area calculations (Shape.idr)

```
module Shape

public export
data Shape = Triangle Double Double
            | Rectangle Double Double
            | Circle Double

private
rectangle_area : Double -> Double -> Double
rectangle_area width height = width * height

export
area : Shape -> Double
area (Triangle base height) = 0.5 * rectangle_area base height
area (Rectangle length height) = rectangle_area length height
area (Circle radius) = pi * radius * radius
```

Exports modifiers,
saying whether the
names are visible
outside this module

Each name defined in this module has an export modifier that explains whether that name is visible to other modules. An export modifier can be one of the following:

- **private**—The name is not exported at all.
- **export**—The name and type are exported, but not the definition. In the case of a data type, this means the type constructor is exported, but the data constructors are private.
- **public export**—The name, type, and complete definition are exported. For data types, this means the data constructors are exported. For functions, this means the functions' definitions are exported.

If there's no export modifier on a function or data type definition, Idris treats it as **private**. In the preceding example, this means that a module that imports Shape can use the names Shape, Triangle, Rectangle, Circle, and area, but not rectangle_area.

EXPORTING FUNCTION DEFINITIONS Exporting a function's definition as well as its type (via **public export**) is important if you want to use the function's behavior in a type. In particular, this is important for type synonyms and type-level functions, which we first used in chapter 6.

The next listing shows an alternative version of the Shape module that keeps the details of the Shape data type abstract, exporting the type but not its constructors.

Listing 10.11 Exporting Shape as an abstract data type (Shape_abs.idr)

```
module Shape_abs
  export
  data Shape = Triangle Double Double
              | Rectangle Double Double
              | Circle Double

  export
  triangle : Double -> Double -> Shape
  triangle = Triangle

  export
  rectangle : Double -> Double -> Shape
  rectangle = Rectangle

  export
  circle : Double -> Shape
  circle = Circle
```

The code is annotated with two callout boxes:

- A box on the right side of the first export statement contains the text "Exports the Shape data type, but not its constructors".
- A box on the right side of the three export statements below it contains the text "Exports functions for building shapes, rather than their constructors".

Here, we've exported the functions `triangle`, `rectangle`, and `circle` for building `Shape`. Rather than using the data constructors directly, other modules will need to use these functions and won't be able to pattern-match on the `Shape` type, because the constructors aren't exported.

Using `export` modifiers, you can implement a module that implements the features of a data store but only exports functions for creating a store, adding items, and traversing the store, without exporting any details about the structure of the store.

10.3.2 The data store, revisited

To illustrate the role of views in data abstraction, we'll create a module that implements a data store, exporting functions for constructing the store. We'll also implement a view for inspecting and traversing the contents of the store.

The following listing shows the `DataStore.idr` module. This is a slight variation on the `DataStore` record you implemented in chapter 6.

Listing 10.12 Data store, with a schema (DataStore.idr)

```
module DataStore
  import Data.Vect
  infixr 5 .+.

  public export
  data Schema = SString | SInt | (.+) Schema Schema

  public export
  SchemaType : Schema -> Type
  SchemaType SString = String
  SchemaType SInt = Int
  SchemaType (x .+ y) = (SchemaType x, SchemaType y)
```

The code is annotated with two callout boxes:

- A box on the right side of the `Schema` definition contains the text "Exports Schema and all of its data constructors".
- A box on the right side of the `SchemaType` definition contains the text "Exports SchemaType and its definition".
- A vertical callout box on the right side of the page contains the text "The Schema type was defined for the data store implementation in chapter 6, along with SchemaType, which translates a Schema into an Idris type." This box spans both annotations.

```
export
record DataStore (schema : Schema) where
  constructor MkData
  size : Nat
  items : Vect size (SchemaType schema)
```

Exports DataStore, but not its constructor or any of its fields

The DataStore record is parameterized by the data schema.

Rather than storing the schema as a field in the record, here you parameterize the record by the data's schema because you don't intend to allow the schema to be updated:

```
export
record DataStore (schema : Schema) where
  constructor MkData
  size : Nat
  items : Vect size (SchemaType schema)
```

The syntax for a parameterized record declaration is similar to the syntax for an interface declaration, with the parameters and their types listed after the record name. This declaration gives rise to a `DataStore` type constructor with the following type:

```
DataStore : Schema -> Type
```

It also gives rise to functions for projecting the size of the store (`size`) and the entries in the store (`items`) out of the record. The functions have the following types:

```
size : DataStore schema -> Nat
items : (rec : DataStore schema) -> Vect (size rec) (SchemaType schema)
```

Because the record has the export modifier `export`, the `DataStore` data type is visible to other modules, but the `size` and `items` projection functions aren't.

Listing 10.13 shows three functions that other modules can use to create a new, empty store with a specific schema (`empty`), or add a new entry to a store (`addToStore`). Each of these functions has the export modifier `export`, meaning that other modules can see their names and types but have no access to their definitions.

Listing 10.13 Functions for accessing the store (DataStore.idr)

```
export
empty : DataStore schema
empty = MkData 0 []

export
addToStore : (value : SchemaType schema) -> (store : DataStore schema) ->
             DataStore schema
addToStore value (MkData _ items) = MkData _ (value :: items)
```

To be able to use this module effectively, you'll also need to traverse the entries in the store. You can build the contents of a store using `empty` to create a new store and `addToStore` to add a new entry. It would therefore be convenient to be able to use

these as patterns to match the contents of a store. When you match on a store, you'll need to deal with the following two cases:

- An empty case that matches a store with no contents
- An addToStore value store case that matches a store where the first entry is given by value, and the remaining items in the store are given by store

To match these cases, you can write a view of DataStore.

10.3.3 Traversing the store's contents with a view

Listing 10.14 shows a StoreView view and its covering function, storeView. They allow you to traverse the contents of a store by seeing how the store was constructed, either with empty or addToStore.

Listing 10.14 A view for traversing the entries in a store (DataStore.idr)

```
public export
data StoreView : DataStore schema -> Type where
    SNil : StoreView empty
    SAdd : (rec : StoreView store) -> StoreView (addToStore value store)

    storeViewHelp : (items : Vect size (SchemaType schema)) ->
        StoreView (MkData size items)
    storeViewHelp [] = SNil
    storeViewHelp (val :: xs) = SAdd (storeViewHelp xs)

export
storeView : (store : DataStore schema) -> StoreView store
storeView (MkData size items) = storeViewHelp items
```

You want to match on the constructors of StoreView when using the view, so export the data constructors.

There's no export annotation, so Idris considers storeViewHelp to be private.

Exports the covering function so that other modules can build the view

The StoreView view gives you access to the contents of the store by pattern matching but hides its internal representation. To use the store, and to traverse its contents, you don't need to know anything about the internal representation.

To illustrate this, let's set up some test data and write some functions to inspect it. The next listing defines a store and populates it with some test data, mapping planet names to the name of the space probe that first visited the planet and the year of the visit.¹

Listing 10.15 A datastore populated with some test data (TestStore.idr)

```
import DataStore

testStore : DataStore (SString .+ SString .+ SInt)
testStore = addToStore ("Mercury", "Mariner 10", 1974)
          (addToStore ("Venus", "Venera", 1961))
```

¹ We won't, however, get into any debates about whether Pluto is a planet here.

```
(addToStore ("Uranus", "Voyager 2", 1986)
(addToStore ("Pluto", "New Horizons", 2015)
empty))
```

The application operator \$

When expressions get deeply nested, as in the `testStore` definition, it can be difficult to keep track of the bracketing. The `$` operator is an infix operator that applies a function to an argument, and you can use it to reduce the need for bracketing.

Using it, you can write the following:

```
testStore = addToStore ("Mercury", "Mariner 10", 1974) $
  addToStore ("Venus", "Venera", 1961) $
  addToStore ("Uranus", "Voyager 2", 1986) $
  addToStore ("Pluto", "New Horizons", 2015) $
empty
```

Writing `$` is therefore equivalent to putting the rest of the expression in brackets. For example, writing `f x $ y z` is exactly equivalent to writing `f x (y z)`.

The following listing shows a basic traversal of the data store, returning a list of entries in the store.

Listing 10.16 A function to convert the store's contents to a list of entries (TestStore.idr)

```
listItems : DataStore Schema -> List (SchemaType schema)
listItems input with (storeView input)
  listItems empty | SNil = []
  listItems (addToStore value store) | (SAdd rec)
    = value :: listItems store | rec
```

The `| rec` bypasses the building of the view on the recursive call, because `rec` is already a view of the rest of the store.

If you call `showItems` with the test data, you'll see the following result:

```
*TestStore> listItems testStore
[("Mercury", "Mariner 10", 1974),
 ("Venus", "Venera", 1961),
 ("Uranus", "Voyager 2", 1986),
 ("Pluto", "New Horizons", 2015)] : List (String, String, Int)
```

More interestingly, you might want to write functions that traverse the data store and filter out certain entries. For example, suppose you want to get a list of the planets that were first visited by a space probe during the twentieth century. You could do this by writing the following function:

```
filterKeys : (test : SchemaType val_schema -> Bool) ->
  DataStore (SString .+ val_schema) -> List String
```

You can think of a schema of the form `(SString .+ val_schema)` as giving a key-value pair, where the key is a `String` and `val_schema` describes the form of the values.

Then, `filterKeys` will apply a function to the value in the pair, and if it returns `True`, it will add the key to a list of `String`. This can find the planets that a probe visited before the year 2000:

```
*TestStore> filterKeys (\x => snd x < 2000) testStore
["Mercury", "Venus", "Uranus"] : List String
```

You can implement `filterKeys` using `StoreView` by taking the following steps:

- 1** *Type, define*—Begin with a type and a skeleton definition:

```
filterKeys : (test : SchemaType val_schema -> Bool) ->
             DataStore (SString .+. val_schema) -> List String
filterKeys test input = ?filterKeys_rhs
```

- 2** *Define*—You'll define the function by traversing the store using `StoreView`, so you can use the `with` construct to build the view, and case-split on it:

```
filterKeys : (test : SchemaType val_schema -> Bool) ->
             DataStore (SString .+. val_schema) -> List String
filterKeys test input with (storeView input)
  filterKeys test empty | SNil = ?filterKeys_rhs_1
  filterKeys test (addToStore value store) | (SAdd rec)
    = ?filterKeys_rhs_2
```

- 3** *Refine*—If the store is empty, there's no value to apply the test to, so return an empty list:

```
filterKeys : (test : SchemaType val_schema -> Bool) ->
             DataStore (SString .+. val_schema) -> List String
filterKeys test input with (storeView input)
  filterKeys test empty | SNil = []
  filterKeys test (addToStore value store) | (SAdd rec)
    = ?filterKeys_rhs_2
```

- 4** *Refine*—Otherwise, because of the schema of the data store, entry must itself be a key-value pair:

```
filterKeys : (test : SchemaType val_schema -> Bool) ->
             DataStore (SString .+. val_schema) -> List String
filterKeys test input with (storeView input)
  filterKeys test empty | SNil = []
  filterKeys test (addToStore (key, value) store) | (SAdd rec)
    = ?filterKeys_rhs_2
```

You'll apply `test` to the `value`. If the result is `True`, you'll keep the key and recursively build the rest of the list. If the result is `False`, you'll omit the key and build the rest of the list:

```
filterKeys : (test : SchemaType val_schema -> Bool) ->
             DataStore (SString .+. val_schema) -> List String
filterKeys test input with (storeView input)
  filterKeys test empty | SNil = []
  filterKeys test (addToStore (key, value) store) | (SAdd rec)
    = if test value
```

```
then key :: filterKeys test store | rec
else filterKeys test store | rec
```

You can try this function with some test filters:

```
*TestStore> filterKeys (\x => fst x == "Voyager 2") testStore
["Uranus"] : List String

*TestStore> filterKeys (\x => snd x > 2000) testStore
["Pluto"] : List String

*TestStore> filterKeys (\x => snd x < 2000) testStore
["Mercury", "Venus", "Uranus"] : List String
```

For both `showItems` and `filterKeys`, you've written a function that traverses the contents of the data store without knowing anything about the internal representation of the store. In each case, you've used a view to deconstruct the data, rather than deconstructing the data directly. If you were to change the internal representation in the `DataStore` module, and correspondingly the implementation of `storeView`, the implementations of `showItems` and `filterKeys` would remain unchanged.

Exercises



- 1 Write a `getValues` function that returns a list of all values in a `DataStore`. It should have the following type:

```
getValues : DataStore (SString .+ val_schema) ->
List (SchemaType val_schema)
```

You can test your definition by writing a function to set up a data store:

```
testStore : DataStore (SString .+ SInt)
testStore = addToStore ("First", 1) $
           addToStore ("Second", 2) $
           empty
```

If you've implemented `getValues` correctly, you should see the following:

```
*ex_10_3> getValues testStore
[1, 2] : List Int
```

- 2 Define a view that allows other modules to inspect the abstract `Shape` data type in listing 10.11. You should be able to use it to complete the following definition:

```
area : Shape -> Double
area s with (shapeView s)
area (triangle base height) | STriangle = ?area_rhs_1
area (rectangle width height) | SRectangle = ?area_rhs_2
area (circle radius) | SCircle = ?area_rhs_3
```

If you have implemented this correctly, you should see the following:

```
*ex_10_3> area (triangle 3 4)
6.0 : Double
*ex_10_3> area (circle 10)
314.1592653589793 : Double
```

10.4 Summary

- A view is a dependent type that describes the possible forms of another data type. Views take advantage of dependent pattern matching to allow you to extend the forms of patterns you can use.
- A covering function builds a view of a value. By convention, its name is the name of the view with an initial lowercase letter.
- The `with` construct allows you to use views directly, without defining an intermediate function.
- You can use views to define alternative traversals of data structures, such as extracting the last element of a list instead of the first.
- Recursive views help you write functions that are guaranteed to terminate, by writing recursive functions that pattern-match on the view.
- Idris provides several views for alternative traversals of `List` in the `Data.List.Views` library. Similar libraries exist for `Vect`, `Nat`, and `String`.
- You can hide the structure of data in a module, while still supporting interactive type-driven programming with that data, by exporting views for traversing data structures.

Part 3

Idris and the real world

I

In part 2, you gained experience in developing programs interactively, guided by types, and you learned about all of the core features of Idris. Now, it's time to apply what you've learned to some more practical examples.

First, in chapter 11, you'll learn about writing programs that deal with potentially infinite structures such as streams. You've learned about the importance of writing total functions, but in chapter 11 you'll see that totality is about more than termination. A function is also total if it produces some portion of a potentially infinite result, which means you can write interactive systems such as servers and read-eval-print loops that run forever, but which are nevertheless total.

Chapters 12–14 deal with state. Real-world programs usually need to deal with global state somehow, and you'll see both how to represent state and how to describe properties of state in such a way that you can guarantee that programs follow protocols accurately. If you're implementing a system with important security properties, such as an ATM for a bank, you can use type-driven development to ensure that those properties are satisfied.

Finally, chapter 15 provides an extended example of type-driven development, showing how to implement a library for concurrent programming with types. You'll start by writing a simple type to capture a specific concurrent programming problem, and then gradually refine it to capture more of the important properties of concurrent programs.

Streams and processes: working with infinite data

This chapter covers

- Generating and processing streams of data
- Distinguishing terminating from productive total functions
- Defining total interactive processes using infinite streams

The functions we've written in this book so far have worked in batch mode, processing all of their inputs and then returning an output. In the previous chapter, we also spent some time discussing why termination is important, and you learned how to use views to help you write programs that are guaranteed to terminate.

But input data doesn't always arrive in a batch, and you'll often want to write programs that *don't* terminate, running indefinitely. For example, it can be convenient to think of input data to an interactive program (such as keypresses, mouse movements, and so on) as a continuous stream of data, processed one element at a time, leading to a stream of output data. In reality, many programs are, in effect, stream processors:

- A read-eval-print loop, such as the Idris environment, processes a potentially infinite stream of user commands, giving an output stream of responses.
- A web server processes a potentially infinite stream of HTTP requests, giving an output stream of responses to be sent over the network.
- A real-time game processes a potentially infinite stream of commands from a controller, giving an output stream of audiovisual actions.

Furthermore, even when you’re writing pure functions that don’t interact with external data sources or devices, streams allow you to write reusable program components by separating the *production* of data from the *consumption* of data. For example, suppose you’re writing a function to determine the square root of a number. You can do this by producing an infinite list of successively closer approximations to a solution, and then writing a separate function to consume that list, finding the first approximation within the desired bounds.

PRODUCING AND CONSUMING DATA A common theme in this chapter is the distinction between programs that consume (or *process*) data, and programs that produce data. All the functions you’ve seen in this book so far have been consumers of data, and in the last chapter we looked at using views to help us write functions that are guaranteed to terminate when consuming data. When you’re writing terminating functions, however, consuming data is only part of the story: a function that generates an infinite stream is never going to terminate, after all. As you’ll see, Idris checks that functions that generate streams are guaranteed to be *productive*, so that any function that consumes the output of a stream generator will always have data to process.

The kinds of programs we write in practice often have a *terminating* component, processing and responding to user input, and a *nonterminating* component, which is an infinite loop that repeatedly invokes the terminating component. In this chapter, you’ll see how to write programs that manage this distinction, both producing and consuming potentially infinite data. We’ll start with one of the most common infinite structures, streams, and later look at how to define total functions describing interactive functions that execute indefinitely.

11.1 Streams: generating and processing infinite lists

Streams are infinite sequences of values, and you can process one value at a time. In this section, you’ll see how to write functions that *produce* an infinite sequence of data, as required, and how to write functions that *consume* finite portions of data produced as a stream.

As a first example, to illustrate the ideas behind streams, we’ll look at how to generate an infinite sequence of numbers, $0, 1, 2, 3, 4, \dots$, and how to process them as needed to label elements in a list, as illustrated in figure 11.1.




Figure 11.1 Labeling the elements of a List by taking elements from an infinite stream of numbers. The stream contains an infinite number of elements, but you only take as many as you need to label the elements in the finite list.

You might use such a function after sorting some data, for example, to attach an explicit index to the data. As you'll see, you can use streams to cleanly separate the production of an infinite list of labels from the consumption of the labels you need for a specific input list.

As well as showing you how to define potentially infinite data types, I'll also introduce the `Stream` data type, provided by the Prelude, and we'll briefly look at some functions on Streams. Finally, we'll look at a larger example using a stream of random numbers to implement an arithmetic game.

11.1.1 Labeling elements in a List

Suppose you want to write a function that labels every element of a List with an integer indicating its position in the list, as in figure 11.1. That is, something like this:

```
label : List a -> List (Integer, a)
```

Running this function on some examples should give the following results:

```
*Streams> label ['a', 'b', 'c']
[(0, 'a'), (1, 'b'), (2, 'c')] : List (Integer, Char)

*Streams> label ["Anderson", "Botham", "Willis", "Trueman"]
[(0, "Anderson"),
 (1, "Botham"),
 (2, "Willis"),
 (3, "Trueman")] : List (Integer, String)
```

The following listing shows one way to write `label` by writing a helper function, `labelFrom` that takes the label for the first element of the list, and then labels the remainder of the list, incrementing the label.

Listing 11.1 Labeling each element of a list with an integer (Label.idr)

```

labelFrom : Integer -> List a -> List (Integer, a)
labelFrom lbl [] = []
labelFrom lbl (val :: vals) = (lbl, val) :: labelFrom (lbl + 1) vals
                                ←
label : List a -> List (Integer a)
label = labelFrom 0

```

Initializes the label at 0

Labels the first element of the list, and then recursively labels the tail

This works as required, but the definition of `labelFrom` combines two components: labeling each element, and generating the label itself. An alternative way of writing `label` would allow you to reuse these two components separately—you could write two functions:

- `countFrom`—Generates an infinite stream of numbers, counting upwards from a given starting point.
- `labelWith`—Takes an infinite stream of labels, and pairs each label with a corresponding element in a *finite* list. It therefore only consumes as much of the infinite stream as necessary to label the elements in the list.

A natural way to try to write `countFrom` might be to generate a `List` of `Integer` from a given starting point:

```

countFrom : Integer -> List Integer
countFrom n = n :: countFrom (n + 1)

```

When Idris runs a compiled program, however, it fully evaluates the *arguments* to a function before it evaluates the function itself. So, unfortunately, if you try to pass the result of `countFrom` to a function that expects a `List`, that function will never run because the result of `countFrom` will never be fully evaluated. If you ask Idris whether this definition of `countFrom` is total, it will tell you that there's a problem:

```

*StreamFail> :total countFrom
Main.countFrom is possibly not total due to recursive path:
    Main.countFrom, Main.countFrom

```

You can see that `countFrom` will never terminate, because it makes a recursive call for every input, but to write `labelWith` you'll only need a finite portion of the result of `countFrom`. What you really need to know about `countFrom`, therefore, is not that it always *terminates*, but rather that it will always produce as many numbers as you need. That is, you need to know that it's *productive* and is guaranteed to generate an indefinitely long sequence of numbers.

As you'll see in the next section, you can use types to distinguish between those expressions for which evaluation is guaranteed to terminate and those expressions for which evaluation is guaranteed to keep producing new values, marking arguments to a data structure as potentially infinite.

11.1.2 Producing an infinite list of numbers

To generate an infinite list of numbers and consume only the finite portion of the list that you need, you can use a new data type, `Inf`, for marking the potentially infinite parts of the structure.

You'll see more details of how `Inf` works shortly. First, though, let's take a look at a data type of infinite lists that uses `Inf`.

Listing 11.2 A data type of infinite lists (InfList.idr)

There's no `Nil` constructor,
so no end to the list.

The `Inf` generic type marks
the `InfList elem` argument
as potentially infinite.

```
→ data InfList : Type -> Type where
  (::) : (value : elem) -> Inf (InfList elem) -> InfList elem
  %name InfList xs, ys, zs
```

←———— Name hints for interactive editing.

`InfList` is similar to the `List` generic type, with two significant differences:

- There's no `Nil` constructor, only a `(::)` constructor, so there's no way to end the list.
- The recursive argument is wrapped in a new data type, `Inf`, that marks the argument as potentially infinite.

To manipulate potentially infinite computations, you can use the `Delay` and `Force` functions. Listing 11.3 gives the types of `Delay` and `Force`. The idea is that you can use `Delay` and `Force` to control exactly when a subexpression is evaluated, so that you only calculate the finite portion of an infinite list that's required for a specific function.

Listing 11.3 The `Inf` abstract data type, for delaying potentially infinite computations

`Inf` is a generic type of potentially infinite computations.

`Delay` is a function that states that its argument should only be evaluated when its result is forced.

```
→ Inf : Type -> Type
```

```
Delay : (value : ty) -> Inf ty
```

```
Force : (computation : Inf ty) -> ty
```

←———— Force is a function that returns the result from a delayed computation.

The following listing shows how you can define `countFrom`, generating an infinite list of `Integers` from a given starting value.

Listing 11.4 Defining `countFrom` as an infinite list (InfList.idr)

```
countFrom : Integer -> InfList Integer
countFrom x = x :: Delay (countFrom (x + 1))
```

←———— The Delay means that the remainder of the list will only be calculated when explicitly requested using Force.

If you try evaluating `countFrom 0` at the REPL to generate an infinite list counting upwards from 0, you'll see the effect of the `Delay`:

```
*InfList> countFrom 0
0 :: Delay (countFrom 1) : InfList Integer
```

You can see that the Idris evaluator has left the argument to `Delay` unevaluated. The evaluator treats `Force` and `Delay` specially: it will only evaluate an argument to `Delay` when explicitly requested to by a `Force`. As a result, despite there being a recursive call to `countFrom` on every input, evaluation at the REPL still terminates. Idris even agrees that it's total:

```
*InfList> :total countFrom
Main.countFrom is Total
```

TERMINOLOGY: RECURSION AND CORECURSION, DATA AND CODATA You may hear Idris programmers referring to functions such as `countFrom` as *corecursive* rather than *recursive*, and infinite lists as *codata* rather than *data*. The distinction between data and codata is that data is *finite* and is intended to be *consumed*, whereas codata is potentially *infinite* and is intended to be *produced*. Whereas recursion operates by taking data and breaking it down toward a base case, corecursion operates by starting at a base case and building up codata.

It may seem surprising that Idris considers `countFrom` to be total, given that it produces an infinite structure. Before we continue discussing how to work with infinite lists, therefore, it's worth investigating in more detail what it means for a function to be total.

11.1.3 Digression: what does it mean for a function to be total?

If a function is total, it will never crash due to a missing case (that is, all well-typed inputs are covered), and it will always return a well-typed result within a finite time. The functions you've written in previous chapters have all taken finite data as inputs, so they're total as long as they terminate for all inputs. But now that you've seen the `Inf` type, you're able to write functions that produce infinite data, and these functions don't terminate! We'll therefore need to refine our understanding of what it means for a function to be total.

Functions that produce infinite data can be used as components of terminating functions, provided they'll always produce a new piece of data on request. In the case of `countFrom`, it will always produce a new `Integer` before making a delayed recursive call.

Figure 11.2 illustrates the structure of `countFrom`. The delayed recursive call to `countFrom` is an argument to `(::)`, meaning that `countFrom` will always produce at least one element of an infinite list before making a recursive call. Therefore, any function that consumes the result of `countFrom` will always have data to work with.




Figure 11.2 Producing values of an infinite structure. The `Delay` means that Idris will only make the recursive call to `countFrom` when it's required by `Force`.

Total functions defined

A total function is a function that, for all well-typed inputs, does one of the following:

- Terminates with a well-typed result
- Produces a non-empty finite prefix of a well-typed infinite result in finite time

We can describe total functions as either *terminating* or *productive*. The halting problem is the difficulty of determining whether a specific program terminates or not, and, thanks to Alan Turing, we know that it's impossible in general to write a program that solves the halting problem. In other words, Idris can't determine whether one of these conditions holds for *all* total functions. Instead, it makes a conservative approximation by analyzing a function's syntax.

Idris considers a function to be total if there are patterns that cover all well-typed inputs, and it can determine that one of the following conditions holds:

- When there's a recursive call (or a sequence of mutually recursive calls), there's a decreasing argument that converges toward a base case.
- When there's a recursive call as an argument to `Delay`, the delayed call will always be an argument to a data constructor (or sequence of nested data constructors) after evaluation, for all inputs.

We discussed the first of these conditions in the previous chapter. The second condition allows us to use functions like `countFrom` in a terminating function. To illustrate this further, it's helpful to see how the resulting infinite list is used. As an example, let's write a function that consumes a finitely long prefix of an `InfList`.

11.1.4 Processing infinite lists

A function that generates an `InfList` is total provided that it's guaranteed to keep producing data whenever data is required. You can see how this works by writing a program that calculates a finite list from the prefix of an infinite list:

```
getPrefix : (count : Nat) -> InfList ty -> List ty
```

`getPrefix` returns a `List` consisting of the first `count` items from an infinite list. It works by recursively taking the next element from the infinite list as long as it needs more elements. You can define it with the following steps:

- 1 *Define*—First, case-split on the `count` argument:

```
getPrefix : (count : Nat) -> InfList a -> List a
getPrefix Z xs = ?getPrefix_rhs_1
getPrefix (S k) xs = ?getPrefix_rhs_2
```

- 2 *Refine*—If you’re taking zero elements from the infinite list, return an empty list:

```
getPrefix : (count : Nat) -> InfList a -> List a
getPrefix Z xs = []
getPrefix (S k) xs = ?getPrefix_rhs_2
```

- 3 *Define*—If you’re taking more than one element, you’ll case-split on the infinite list and then add the first value in the infinite list as the first element of the result.

```
getPrefix : (count : Nat) -> InfList a -> List a
getPrefix Z xs = []
getPrefix (S k) (value :: xs) = value :: ?getPrefix_rhs_1
```

- 4 *Type, refine*—If you look at the type of the hole, `?getPrefix_rhs_1`, you’ll see the following:

```
a : Type
k : Nat
value : a
xs : Inf (InfList a)
-----
getPrefix_rhs_1 : List a
```

You can see from the type of `xs` that it’s an infinite list that has not yet been computed, because it’s an infinite list wrapped in an `Inf`. To complete the definition, you can Force the computation of `xs` and recursively get its prefix:

```
getPrefix : (count : Nat) -> InfList a -> List a
getPrefix Z xs = []
getPrefix (S k) (value :: xs) = value :: getPrefix k (Force xs)
```

The resulting definition is total, according to Idris:

```
*InfList> :total getPrefix
Main.getPrefix is Total
```

Even though one of the inputs is potentially infinite, `getPrefix` will only evaluate as much as is necessary to retrieve `count` elements from the infinite list. Because `count` is a finite number, `getPrefix` will always terminate as long as the `InfList` is guaranteed to continue producing new elements.

In practice, you can omit calls to `Delay` and `Force` and let Idris insert them where required. If, during type checking, Idris encounters a value of type `Inf ty` when it

requires a value of type `ty`, it will add an implicit call to `Force`. Similarly, if it encounters a `ty` when it requires an `Inf ty`, it will add an implicit call to `Delay`. The following listing shows how you can define `countFrom` and `getPrefix` using implicit `Force` and `Delay`.

Listing 11.5 Taking a finite portion of an infinite list, with implicit Force and Delay (InflList.idr)

```
countFrom : Integer -> InfList Integer
countFrom x = x :: countFrom (x + 1)    ← The Idris type checker implicitly inserts
getPrefix : Nat -> InfList a -> List a   the Delay required for the recursive call.
getPrefix Z x = []
getPrefix (S k) (x :: xs) = x :: getPrefix k xs ← The Idris type checker implicitly
                                                inserts the Force required for xs.
```

You can therefore treat `Inf` as an annotation on a type, mark the parts of a data structure that may be infinite, and let the Idris type checker manage the details of when computations must be delayed or forced.

Now that you've seen how to separate the production of data, using `countFrom` to generate an infinite list of numbers, from the consumption of data, using a function like `getPrefix`, we can revisit the definition of `label`. Rather than using our own `InfList` data type and `countFrom`, we'll use a data type defined in the Prelude for this purpose: `Stream`.

11.1.5 The Stream data type

Listing 11.6 shows the definition of `Stream` in the Prelude. It has the same structure as the definition of `InfList` you saw in the previous section. Additionally, the Prelude provides several useful functions for building and manipulating Streams, some of which this listing also shows.

Listing 11.6 The Stream data type and some functions, defined in the Prelude

```
data Stream : Type -> Type where
  (::) : (value : elem) -> Inf (Stream elem) -> Stream elem   ← Generates an
                                                                     infinite list of a
                                                                     specific element
repeat : elem -> Stream elem
take : (n : Nat) -> (xs : Stream elem) -> List elem
iterate : (f : elem -> elem) -> (x : elem) -> Stream elem
```

Takes a specific number of elements from the start of a stream

Generates a stream by repeatedly applying a function

You can see the functions `repeat`, `take`, and `iterate` in action at the REPL. For example, `repeat` generates an infinite sequence of an element, delayed until specifically requested:

```
Idris> repeat 94
94 :: Delay (repeat 94) : Stream Integer
```

Like `getPrefix` on `InfList`, `take` takes a prefix of a Stream of a specific length:

```
Idris> take 10 (repeat 94)
[94, 94, 94, 94, 94, 94, 94, 94, 94] : List Integer
```

`iterate` applies a function repeatedly, generating a stream of the results. For example, starting at 0 and repeatedly applying `(+1)` leads to a sequence of increasing integers, like `countFrom`:

```
Idris> take 10 (iterate (+1) 0)
[0, 1, 2, 3, 4, 5, 6, 7, 8, 9] : List Integer
```

Syntactic sugar for stream generation

Idris provides a concise syntax for generating streams of numbers, similar to the syntax for lists:

```
Idris> take 10 [1..]
[1, 2, 3, 4, 5, 6, 7, 8, 9, 10] : List Integer
```

The syntax `[1..]` generates a Stream counting upwards from 1. This works for any countable numeric type, as in the following example:

```
Idris> the (List Int) take 10 [1..]
[1, 2, 3, 4, 5, 6, 7, 8, 9, 10] : List Int
```

You can also change the increment:

```
Idris> the (List Int) (take 10 [1,3..])
[1, 3, 5, 7, 9, 11, 13, 15, 17, 19] : List Int
```

Putting all of this together, the following listing shows how to define `label` using `iterate` to generate an infinite sequence of integer labels, and a `labelWith` function that consumes enough of an infinite sequence of labels to attach a label to each element of a List.

Listing 11.7 Labeling each element of a List using a Stream (Streams.idr)

Uses a generic type variable, `labelType`, because `labelWith` never inspects the label itself

```
→ labelWith : Stream labelType -> List a -> List (labelType, a)
→ labelWith [] = []
labelWith (lbl :: lbls) (val :: vals) = (lbl, val) :: labelWith lbls vals ←
  label : List a -> List (Integer, a)
  label = labelWith (iterate (+1) 0)
```

Labels the first element of the list with the first element in the stream of labels

Nothing left to label, so returns an empty list

In this definition, you separate the two components of generating the labels and assigning labels to each element of the list.

You can even give a more generic type to `labelWith`, taking a `Stream` `labelType` rather than a `Stream Integer`, and allow labeling by any type for which you can generate a `Stream`. For example, the `cycle` function generates a `Stream` that repeats a given sequence:

```
*Streams> take 10 $ cycle ["a", "b", "c"]
["a", "b", "c", "a", "b", "c", "a", "b", "c", "a"] : List String
```

Using cycle to generate a Stream, you could label each element of a List with a cycling sequence of labels, repeated as many times as necessary to label the entire list:

```
*Streams> labelWith (cycle ["a", "b", "c"]) [1..5]
[("a", 1), ("b", 2), ("c", 3), ("a", 4), ("b", 5)] : List (String, Integer)
```

11.1.6 An arithmetic quiz using streams of random numbers

You can use a Stream in any situation where you need a source of data but don't know in advance how much data you'll need to generate. For example, you could write an interactive arithmetic quiz that takes a source of numbers for the questions. The following listing shows how you might do this.

Listing 11.8 An arithmetic quiz, taking a Stream of numbers for the questions (Arith.idr)

```
quiz : Stream Int -> (score : Nat) -> IO ()  
quiz (num1 :: num2 :: nums) score  
= do putStrLn ("Score so far: " ++ show score)  
     putStrLn (show num1 ++ " * " ++ show num2 ++ "? ")  
     answer <- getLine  
     if cast answer == num1 * num2  
         then do putStrLn "Correct!"  
                 quiz nums (score + 1)  
         else do putStrLn ("Wrong, the answer is " ++ show (num1 * num2))  
                 quiz nums score
```

Takes two numbers from an infinite source of Ints

Correct answer; continues with an increased score

Wrong answer; displays the correct answer and continues with the same score

The quiz function takes an infinite source of `Ints` and the score so far, and then returns an `IO` action that displays the score and a question, reads an answer from the user, and repeats. The questions arise directly from the input stream, so you could try with a sequence of increasing numbers:

```
*Arith> :exec quiz (iterate (+1) 0) 0
Score so far: 0
0 * 1? 0
Correct!
Score so far: 1
2 * 3? 6
Correct!
```

```
Score so far: 2
4 * 5? 20
Correct!
Score so far: 3
```

ABORTING EXECUTION Execution of a repeatedly looping program will continue until some external event causes the program to exit. You can abort execution at the REPL by pressing Ctrl-C.

So far, this isn't especially interesting because you know in advance what the questions will be. Instead, you could write a function that generates a stream of pseudo-random Ints generated from an initial *seed*. The following listing shows one way to generate a random-looking stream of numbers using a *linear congruential generator*.

Listing 11.9 Generating a stream of pseudo-random numbers from a seed (Arith.idr)

```
randoms : Int -> Stream Int
randoms seed = let seed' = 1664525 * seed + 1013904223 in
    (seed' `shiftR` 2) :: randoms seed'
```

Multiples the current seed by one constant, and adds another

shiftR bitwise-shifts an integer right by a given number of places

PSEUDO-RANDOM NUMBER GENERATION The `randoms` function in listing 11.9 generates a random-looking, but predictable, stream of numbers from an initial seed, using a *linear congruential generator*. This is one of the oldest techniques for pseudo-random number generation. It suits our purposes for this example, but it's not suitable for situations where high-quality randomness is required, such as cryptographic applications, due to the distribution and predictability of the generated numbers.

You could try running `quiz` with a stream of numbers generated by `randoms`:

```
*Arith> :exec quiz (randoms 12345) 0
Score so far: 0
1095649041 * -2129715532? no idea
Wrong, the answer is -765659660
Score so far: 0
```

Earlier, the questions were too predictable, but now they're perhaps a little too hard for most of us! Furthermore, in the preceding example, the result has even overflowed the bounds of an `Int`, so the reported answer is incorrect. Instead, the next listing shows one way to process the results of `randoms` so that they're within reasonable bounds for an arithmetic quiz.

Listing 11.10 Generating suitable inputs for `quiz` (Arith.idr)

```
import Data.Primitives.Views
arithInputs : Int -> Stream Int
arithInputs seed = map bound (randoms seed)
```

You need to import this for the Divides view

```
where
  bound : Int -> Int
  bound num with (divides num 12) ←
    bound ((12 * div) + rem) | (DivBy prf) = rem + 1
                                              ↑
                                              Matches on num by describing it
                                              in terms of a dividend and
                                              remainder when divided by 12
```

Safe division with the Divides view

You can reduce an arbitrary `Int` to a value between 1 and 12 by dividing by 12 and checking the remainder. Dividing isn't necessarily safe, since division by 0 is undefined on `Int`, so instead you can use a view of `Int` that explains that a number is either 0 or composed of a multiplication plus a remainder:

```
data Divides : Int -> (d : Int) -> Type where
  DivByZero : Int.Divides x 0
  DivBy : (prf : rem >= 0 && rem < d = True) ->
    Int.Divides ((d * div) + rem) d

divides : (val : Int) -> (d : Int) -> Divides val d
```

Note that in the `DivBy` case, you also have a proof that the remainder is guaranteed to have a value between 0 and the divisor by using the equality type = introduced in chapter 8.

You can use `arithInputs` as a source of random inputs for `quiz` and be sure that all the questions will use numbers between 1 and 12. Here's an example:

```
*Arith> :exec quiz (arithInputs 12345) 0
Score so far: 0
2 * 1? 2
Correct!
Score so far: 1
6 * 2? 18
Wrong, the answer is 12
Score so far: 1
11 * 10? 110
Correct!
Score so far: 2
```

You've successfully used a source of random `Ints` as inputs to the `quiz`, and because `randoms` (and hence `arithInputs`) produces an infinite sequence of numbers, you'll be able to generate new numbers as long as you need to.

There is, however, one remaining problem, which is that `quiz` itself is not total:

```
*Arith> :total quiz
Main.quiz is possibly not total due to recursive path:
  Main.quiz, Main.quiz
```

This shouldn't really surprise you, because there's nothing in the definition of `quiz` that allows it to terminate. Instead, like `countFrom`, `randoms`, and `arithInputs`, it's reading user input and continually *producing* an infinite sequence of `IO` actions.

In practice, interactive programs often have an outer loop, which you can run indefinitely, invoking specific commands, each of which you'd like to terminate so that you can issue the next command. The way to write interactive programs that run indefinitely, therefore, is to distinguish the types of interactive programs that describe *terminating* sequences of actions (commands in our main loop) from the types of interactive programs that describe possibly *infinite* sequences of actions (the main loop itself). We'll explore this further in the next section.

Exercises



- 1** Write an `every_other` function that produces a new Stream from every second element of an input Stream.

If you've implemented this correctly, you should see the following:

```
*ex_11_1> take 10 (every_other [1..])
[2, 4, 6, 8, 10, 12, 14, 16, 18, 20] : List Integer
```

- 2** Write an implementation of Functor (as described in section 7.3.1) for `InfList`.

If you've implemented this correctly, you should see the following, using `countFrom` and `getPrefix` as defined in section 11.1.4:

```
*ex_11_1> getPrefix 10 (map (*2) (countFrom 1))
[2, 4, 6, 8, 10, 12, 14, 16, 18, 20] : List Integer
```

- 3** Define a Face data type that represents the faces of a coin: heads or tails. Then, define this:

```
coinFlips : (count : Nat) -> Stream Int -> List Face
```

This should return a sequence of count coin flips using the stream as a source of randomness. If you've implemented this correctly, you should see something like the following:

```
*ex_11_1> coinFlips 6 (randoms 12345)
[Tails, Heads, Tails, Tails, Heads, Tails] : List Face
```

Hint: It will help to define a function, `getFace : Int -> Face`.

- 4** You can define a function to calculate the square root of a Double as follows:

- 1** Generate a sequence of closer approximations to the square root.
- 2** Take the first approximation that, when squared, is within a desired bound from the original number.

Write a function that generates a sequence of approximations:

```
square_root_approx : (number : Double) -> (approx : Double) -> Stream Double
```

Here, you're looking for the square root of `number`, starting from an approximation, `approx`. You can generate the next approximation using this formula:

```
next = (approx + (number / approx)) / 2
```

If you've implemented this correctly, you should see the following:

```
*ex_11_1> take 3 (square_root_approx 10 10)
[10.0, 5.5, 3.659090909090909] : List Double

*ex_11_1> take 3 (square_root_approx 100 25)
[25.0, 14.5, 10.698275862068964] : List Double
```

- 3 Write a function that finds the first approximation of a square root that's within a desired bound, or within a maximum number of iterations:

```
square_root_bound : (max : Nat) -> (number : Double) -> (bound : Double) ->
(approxs : Stream Double) -> Double
```

This should return the first element of `approxs` if `max` is zero. Otherwise, it should return the first element (let's call it `val`) for which the difference between `val` \times `val` and `number` is smaller than `bound`.

If you've implemented this correctly, you should be able to define `square_root` as follows, and it should be `total`:

```
square_root : (number : Double) -> Double
square_root number = square_root_bound 100 number 0.00000000001
(square_root_approx number number)
```

You can test it with the following values:

```
*ex_11_1> square_root 6
2.449489742783178 : Double

*ex_11_1> square_root 2500
50.0 : Double

*ex_11_1> square_root 2501
50.009999000199954 : Double
```

11.2 Infinite processes: writing interactive total programs

When you write a function to generate a `Stream`, you give a prefix of the `Stream` and generate the remainder recursively. This is similar to `quiz` in that you give the initial `IO` actions to run and then generate the remainder of the `IO` actions recursively. So, you can think of an interactive program as being a program that *produces* a potentially infinite sequence of interactive actions.

In chapter 5, you wrote interactive programs using the `IO` generic type, where `IO` `ty` is the type of terminating interactive actions giving a result of type `ty`. In this section, you'll see how to write nonterminating, but productive (and therefore *total*), interactive programs by defining an `InfIO` type for representing infinite sequences of `IO` actions.

Because `InfIO` describes infinite sequences of actions, the functions must be *productive*, so you can be certain that the resulting programs continue to produce new `IO` actions for execution while continuing to run forever. But `InfIO` merely *describes* sequences of interactive actions, so you'll also need to write a function to *execute* those actions.

Overall, we'll take the following approach to writing interactive total programs:

- 1 Define a type `InfIO` that describes infinite sequences of interactive actions.
- 2 Write nonterminating but productive functions using `InfIO`.
- 3 Define a run function that converts `InfIO` programs to `IO` actions.

Your first implementation of `run` will not, itself, be total. You'll then see how to refine the definition so that even `run` is total, using a data type to describe how long execution should continue. First, though, you'll need to see how to describe infinite processes by defining `InfIO`.

11.2.1 Describing infinite processes

The following listing shows how to define the `InfIO` type. It's similar to `Stream`, except that, in an interactive program, you may want the value produced by the first action to influence the rest of a computation.

Listing 11.11 Infinite interactive processes (`InfIO.idr`)

```
data InfIO : Type where
  Do : IO a
    -> (a -> Inf InfIO)
    -> InfIO
```

There are two arguments to `Do`:

- A description of an `IO` action to execute
- The remainder of the infinite sequence of actions. Because its type is `a -> InfIO`, it can use the result produced by the first action to compute the remaining actions.

By using `Inf` to mark the remainder of the sequence as infinite, you're telling Idris that you expect functions that return a value of type `InfIO` to be productive. In other words, as with functions on `Stream`, delayed recursive calls that produce an `InfIO` must be arguments to a constructor. The following listing shows how this works in a recursive program that repeatedly displays a message.

Listing 11.12 An infinite process that repeatedly displays a message (`InfIO.idr`)

```
loopPrint : String -> InfIO
loopPrint msg = Do (putStrLn msg)
  (\_ => loopPrint msg)
```

The recursive call to `loopPrint` is an argument to the constructor `Do`, and `loopPrint` is guaranteed to produce a constructor (`Do`) as a finite prefix of its result. This satisfies the definition of a productive total function from section 11.1.3, so Idris is happy that `loopPrint` is total:

```
*InfIO> :total loopPrint
Main.loopPrint is Total
```

Recall from chapter 5 that `IO` is a generic type describing interactive actions and that it will be executed by the runtime system. If you try evaluating `loopPrint` at the REPL, you'll see a description of the first `IO` action that will be executed and the delayed remainder of the infinite sequence of actions:

```
*InfIO> loopPrint "Hello"
Do (io_bind (prim_write "Hello!\n") (\__bindx => io_return ()))
    (\underscore => Delay (loopPrint "Hello")) : InfIO
```

Just as with `IO`, for this to be useful in practice, you'll also need to be able to execute infinite sequences of actions.

11.2.2 Executing infinite processes

In chapter 5, you learned how the Idris runtime system will execute programs of type `IO ty`, where `ty` is the type of value produced by an interactive computation. So, in order to execute a value of type `InfIO`, you'll need to start by converting it to an `IO ()`. Here's one way to do this.

**Listing 11.13 Converting an expression of type `InfIO` to an executable `IO` action
(`InfIO.idr`)**

```
run : InfIO -> IO ()
run (Do action cont) = do res <- action
                         run (cont res)
                         | The first action to execute
                         | is the first argument to Do.
                         |
                         | Continues execution by passing
                         | the result of the action, res, to
                         | the rest of the computation, cont
```

Using `run`, you can convert your function that repeatedly prints a message to an `IO` action, and execute it using `:exec`:

```
*InfIO> :exec run (loopPrint "on and on and on...")
on and on and on...
on and on and on...
on and on and on...      | Continues indefinitely
```

Because this runs indefinitely, at least until you abort it by pressing Ctrl-C, you should perhaps not be surprised to find that Idris doesn't consider `run` to be total:

```
*InfIO> :total run
Main.run is possibly not total due to recursive path:
Main.run, Main.run
```

You know that `loopPrint` will continue to produce new `IO` actions to execute, because it's total. This is valuable because a program that continues to execute `IO` actions is going to continue to make visible progress (at least assuming those actions produce some output, which we'll consider further in section 11.3.2). It would be nice if `run` were also total, so that you'd at least know that all possible `IO` actions are processed and

that there isn't any unexpected nontermination caused by the implementation of `run` itself.

This would seem to be impossible: the only way to have a total, nonterminating function is to use the `Inf` type, and `IO` is a type of terminating action that doesn't use `Inf`. And, indeed, if you want functions to execute indefinitely at runtime, you'll need at least *some* way to escape from total functions. You can, however, try to make the escape as quietly as possible.

To achieve this, we'll begin by making a terminating version of `run` that takes as an argument an upper bound on the number of actions it's willing to execute.

11.2.3 Executing infinite processes as total functions

Earlier, in section 11.1.4, you wrote a `getPrefix` function that retrieved a finite portion of an infinite list:

```
getPrefix : (count : Nat) -> InfList a -> List a
```

You can think of the `count` argument as being the “fuel” that allows you to continue processing the infinite list. Once you run out of fuel, you can't process any more of the list. You can do something similar for `run`, giving it an additional argument standing for the number of iterations it will run.

The following listing defines a `Fuel` data type and gives a new, total, definition of `run` that will execute actions as long as fuel remains.

Listing 11.14 Converting an expression of type `InfIO` to an executable `IO` action running for a finite time (`InfIO.idr`)

```

Fuel defines the length of time a process can run.    data Fuel = Dry | More Fuel
                                                                ↑
                                                                tank : Nat -> Fuel
                                                                tank Z = Dry
                                                                tank (S k) = More (tank k)
                                                                ← Generates an amount of fuel. Defines a new type for Fuel instead of using Nat—you'll see why shortly.

                                                                run : Fuel -> InfIO -> IO ()
                                                                run (More fuel) (Do c f) = do res <- c
                                                                run fuel (f res)           ← Consumes one drop of fuel and continues execution
                                                                run Dry p = putStrLn "Out of fuel"   ← No more fuel; abandons execution

```

Now, `run` is total:

```
*InfIO> :total run
Main.run is Total
```

Unfortunately, you still have a problem because you now need to specify an explicit maximum number of actions a program is allowed to execute, so you don't really have indefinitely running processes anymore! For example:

```
*InfIO> :exec run (tank 5) (loopPrint "vroom")
vroom
vroom
```

```
vroom
vroom
vroom
Out of fuel
```

It's valuable to ensure that `run` is total, because it guarantees that the implementation of `run` itself won't be the cause of any unexpected nontermination. If you still want programs to run indefinitely, though, you'll need to find a way to generate fuel indefinitely. You can achieve this by using the `Lazy` data type.

11.2.4 Generating infinite structures using Lazy types

If you have a means of generating infinite Fuel, you can run interactive programs indefinitely. Listing 11.15 shows how you can do this using a single nontotal function, `forever`. You also need to change the definition of `Fuel` so that it's explicit in the type that you only generate Fuel when it's required.

Listing 11.15 Generating infinite fuel (InflO.idr)

```
data Fuel = Dry | More (Lazy Fuel)
forever : Fuel
forever = More forever
```

Lazy, explained shortly, means that the value of the argument will only be calculated when needed.

Generates fuel as needed

FOREVER AND NONTERMINATION It's necessary for `forever` to be nontotal because it (deliberately) introduces nontermination. Fortunately, this is the *only* nontotal function you need in order to be able to execute programs forever.

The purpose of `Lazy` is to control when Idris evaluates an expression. As the name `Lazy` implies, Idris won't evaluate an expression of type `Lazy ty` until it's explicitly requested by `Force`, which returns a value of type `ty`. The Prelude defines `Lazy` similarly to `Inf`, which you defined in section 11.1.2:

```
Lazy : Type -> Type
Delay : (value : ty) -> Lazy ty
Force : (computation : Lazy ty) -> ty
```

Also, like `Inf`, Idris inserts calls to `Delay` and `Force` implicitly. In fact, `Inf` and `Lazy` are sufficiently similar that they're implemented internally using the same underlying type, as the next listing shows. The only difference in practice between `Inf` and `Lazy` is the way the totality checker treats them, as explained in the sidebar.

Listing 11.16 Internal definition of Inf and Lazy

```
data DelayReason = Infinite | LazyValue
data Delayed : DelayReason -> Type -> Type where
  Delay : (val : ty) -> Delayed reason ty
```

A computation is delayed either because it may be infinite, or because it's to be evaluated later.

```

Inf : Type -> Type
Inf ty = Delayed Infinite ty
    Inf is a type synonym for a
    delayed infinite value.

Lazy : Type -> Type
Lazy ty = Delayed LazyValue ty
    Lazy is a type synonym for
    a delayed lazy value.

Force : Delayed reason ty -> ty
    Force causes a delayed
    value to be evaluated.

```

Totality checking with Inf and Lazy

At runtime, `Inf` and `Lazy` behave the same way. The key difference between them is the way the totality checker treats them. Idris detects termination by looking for arguments that converge toward a base case, so it needs to know whether an argument to a constructor is smaller (that is, closer to a base case) than the overall constructor expression:

- If the argument has type `Lazy ty`, for some type `ty`, it's considered smaller than the constructor expression.
- If the argument has type `Inf ty`, for some type `ty`, it's *not* considered smaller than the constructor expression, because it may continue expanding indefinitely. Instead, Idris will check that the overall expression is productive, as discussed in section 11.1.3.

If you used `Inf` for `Fuel`, rather than `Lazy`, `run` would no longer be total because the argument, `fuel`, wouldn't be considered smaller than the expression `More fuel`.

You've implemented three functions: `loopPrint`, which is the interactive program; `run`, which executes interactive programs given fuel; and `forever`, which provides an indefinite quantity of fuel. To summarize:

- `loopPrint` is a total function because it continues to produce `IO` actions indefinitely.
- `run` is a total function because it will consume `IO` actions, executing them as long as there is `Fuel` available.
- `forever` is a nontotal (or partial) function because it will never terminate and doesn't produce any data inside an `Inf` type.

By writing a version of `run` that will process data as long as fuel is available, Idris can guarantee that `run` is total, consuming fuel as it goes. You still have an “escape hatch” that allows you to run interactive programs indefinitely, in the form of the `forever` function. Nevertheless, `forever` is the only function that's not total.

You can still improve the definition of `loopPrint` itself. When we wrote interactive programs in chapter 5, we used `do` notation to help make interactive programs more readable, but we haven't been able to do so using `InfIO`. You can, however, extend `do` notation to support your own data types, like `InfIO`.

11.2.5 Extending do notation for InfIO

As you saw in chapter 5, do notation translates to applications of the (`>=`) operator, as illustrated again in figure 11.3.

```
do x <- action      -----> action >= \x => result  
    result
```

Figure 11.3 Transforming do notation to an expression using the `>=` operator when sequencing actions

You've seen this transformation for `IO` in chapter 5, for `Maybe` in chapter 6, and in general for implementations of the `Monad` interface in chapter 7. In fact, the transformation is purely syntactic, so you can define your own implementations of `(>=)` to use `do` notation for your own types. The following listing shows how you can define `do` notation for `InfIO`.

Listing 11.17 Defining do notation for infinite sequences of actions

```
(>=) : IO a -> (a -> Inf InfIO) -> InfIO
(>=) = Do
loopPrint : String -> InfIO
loopPrint msg = do putStrLn msg
                  loopPrint msg
```

Defines the ($>=$) operator by using Do from InfIO directly

Idris translates this do block to applications of the ($>=$) operator.

Idris translates the `do` block to applications of `(>=)` and decides which version of `(>=)` to use by looking at the required type. Here, because the required type of the overall expression is `InfIO`, it uses the version of `(>=)` that produces a value of type `InfIO`.

The `InfIO` type allows you to describe infinitely running interactive programs, and by defining a `(>=)` operator, you can write those programs in much the same way as programs with `IO`, provided that the final action is to call a function of type `InfIO`.

Now that you've seen that you can write productive interactive programs using `do` notation, we can revisit the arithmetic quiz from section 11.1.6.

11.2.6 A total arithmetic quiz

To conclude this section, we'll update the arithmetic quiz so that it's a total function, and you'll see how you can incorporate this in a complete Idris program. Listing 11.18 shows our starting point, setting up the `InfIO` type and `run` function, as you've seen earlier in this section. You'll need `Data.Primitives.Views` for generating the stream of random numbers. You'll also import the `System` module for the `time` function, which you'll use to help initialize the stream of random numbers.

Listing 11.18 Setting up InfIO (ArithTotal.idr)

```
import Data.Primitives.Views
import System
```

You'll need this for the time function, which you'll use to seed the random stream.

```
%default total
data InfIO : Type where
  Do : IO a -> (a -> Inf InfIO) -> InfIO
  (>>=) : IO a -> (a -> Inf InfIO) -> InfIO
  (>>=) = Do

  data Fuel = Dry | More (Lazy Fuel)

  run : Fuel -> InfIO -> IO ()
  run (More fuel) (Do c f) = do res <- c
                                run fuel (f res)
  run Dry p = putStrLn "Out of fuel"
```

This compiler directive means that all functions are, unless otherwise stated, expected to be total.

The %default total directive

Idris supports a number of compiler *directives* that change certain details about the language. In listing 11.18, the `%default total` directive means that Idris will report an error if there are any functions that it can't guarantee to be total.

You can override this for individual functions with the `partial` keyword. For example, `forever` is not total:

```
partial
forever : Fuel
forever = More forever
```

It's a good idea to use `%default total` in your programs, because if Idris can't determine that a function is either terminating or productive, this can be a sign that there's a problem with the function's definition. Furthermore, explicitly marking which functions are partial means that if there's a problem with nontermination, or a program that crashes due to missing input, you've minimized the number of functions that could cause the problem.

Listing 11.19 shows the next step, implementing `quiz` using `InfIO`. Because `InfIO` is an infinite sequence of `IO` actions, you can write `quiz` as before, with the final step being a recursive call. In fact, the definition is identical to the earlier definition; only the type has changed.

Listing 11.19 Defining a total `quiz` function (ArithTotal.idr)

```
quiz : Stream Int -> (score : Nat) -> InfIO
quiz (num1 :: num2 :: nums) score
  = do putStrLn ("Score so far: " ++ show score)
       putStr (show num1 ++ " * " ++ show num2 ++ "? ")
       answer <- getLine
       if (cast answer == num1 * num2)
         then do putStrLn "Correct!"
                 quiz nums (score + 1)
         else do putStrLn ("Wrong, the answer is " ++ show (num1 * num2))
                 quiz nums score
```

The last step is a recursive call to `quiz`. These calls are productive, because they each follow a sequence of `IO` actions.

Because you're using the `%default total` annotation, you can be sure that `quiz` is total. There are two recursive calls to `quiz`, and Idris can determine that each one is guaranteed to be prefixed by a sequence of `IO` actions, so `quiz` is guaranteed to keep producing `IO` actions indefinitely.

The final step is to write a `main` function that calls `run` to execute `quiz` with a stream of `Int`s. Here are the remaining parts of the implementation.

Listing 11.20 Completing the implementation with `main` (`ArithTotal.idr`)

```

randoms : Int -> Stream Int
randoms seed = let seed' = 1664525 * seed + 1013904223 in
               (seed' `shiftR` 2) :: randoms seed'

arithInputs : Int -> Stream Int
arithInputs seed = map bound (randoms seed)
  where
    bound : Int -> Int
    bound x with (divides x 12)
      bound ((12 * div) + rem) | (DivBy prf) = abs rem + 1

    partial
    forever : Fuel
    forever = More forever
    partial
    main : IO ()
    main = do seed <- time
              run forever (quiz (arithInputs (fromInteger seed)) 0)

    forever needs to be
    marked partial explicitly
    because it doesn't
    terminate and doesn't
    use Inf, and you used
    %default total.

    main needs to be
    marked partial
    because it uses forever.

    You can use the system time
    to initialize the stream of
    random numbers.

```

You've used `randoms` and `arithInputs`, as defined in section 11.1.6, to generate the stream of `Int`s. By using the system time to initialize the stream, you'll get different questions every time you run the program.

In the whole implementation, the only functions that aren't total are `forever` and `main`, the latter only because it needs to use `forever` to generate an indefinite amount of fuel for `run`. Because you used the `%default total` annotation, you needed to mark these as `partial` explicitly. That means you can be sure that the only possible cause of nontermination in the program is the fact that you've deliberately said that the program should run forever.

Other than `forever`, you know that the individual components will be one of the following:

- *Productive*—Such as `quiz` continuing to produce `IO` actions for interpretation and `randoms` continuing to produce new random numbers when required
- *Terminating*—Such as the individual `IO` commands and the generation of the next random number from a seed

This distinction is useful for writing realistic programs such as servers and REPs, which you want to run indefinitely while being sure that each individual action the program executes terminates. Often, though, you'll want more flexibility. At the moment, for example, you have no way to quit the `quiz` cleanly. We'll return to this in the next section.

Exercise



The `repl` function defined in the Prelude isn't total because it's an `IO` action that loops indefinitely. Implement a new version of `repl` using `InfIO`:

```
totalREPL : (prompt : String) -> (action : String -> String) -> InfIO
```

If you've implemented this correctly, `totalREPL` should be total, and you should be able to test it as follows:

```
*ex_11_2> :total totalREPL
Main.totalREPL is Total

*ex_11_2> :exec run forever (totalREPL "\n: " toUpper)

: Hello [user input]
HELLO
: World [user input]
WORLD
```

11.3 Interactive programs with termination

Using `Inf`, you can explicitly control when you want data to be produced, or when you want it to be consumed. As a result, you have a choice between programs that always terminate and programs that continue running forever. To write complete applications, though, you'll need more control. After all, although you want a server to run indefinitely, you'd like to be able to shut it down cleanly when you want to.

So far, the types you've defined using `Inf` have had only one constructor, so they've required you to produce infinite sequences. Instead, you can mix infinite and finite components in a single data type, which means you can describe processes that can run indefinitely, but which are also allowed to terminate. In this section, you'll see how to refine the `InfIO` type to support processes that terminate cleanly. Moreover, you'll see how to introduce more precision into the type, and define a type of processes specifically for console I/O.

11.3.1 Refining `InfIO`: introducing termination

Using `InfIO`, you can write total interactive programs that are guaranteed to keep producing `IO` actions, running indefinitely. The following listing shows a small example of the form you saw in the previous section. This program repeatedly asks for the user's name and displays a greeting.

Listing 11.21 Repeatedly greeting a user, running forever (`Greet.idr`)

```
greet : InfIO
greet = do putStrLn "Enter your name: "
          name <- getLine
          putStrLn ("Hello " ++ name)
          greet
```

Typically, when you write interactive programs, you'll want to provide a means for the user to quit. Unfortunately, the only way a user can quit `greet` is by pressing Ctrl-C. There's no way to write a function in `InfIO` that quits any other way!

Fortunately, you can solve this with a small variation on `InfIO`. The `Inf` type marks a value as *potentially infinite*, rather than guaranteeing that the value is infinite, and you can introduce extra data constructors for potentially infinite types. You can define a new `RunIO` type, shown in the next listing, that adds a `Quit` constructor to describe programs that exit, producing a value.

Listing 11.22 The `RunIO` type, describing potentially infinite processes with an additional `Quit` command (`RunIO.idr`)

`RunIO` is parameterized by the type of value produced by an interactive process, if it terminates.

```
data RunIO : Type -> Type where
    Quit : a -> RunIO a
    Do : IO a -> (a -> Inf (RunIO b)) -> RunIO b

(>>=) : IO a -> (a -> Inf (RunIO b)) -> RunIO b
(>>=) = Do
```

Annotations for Listing 11.22:

- A callout points to the `Do` constructor with the text "Quits, producing a value".
- A callout points to the `Do` constructor with the text "A process consisting of a single IO action, followed by a potentially infinite process".
- A callout points to the `(>>=)` operator with the text "Implements (>>=) to support do notation for RunIO programs".

Using `RunIO`, you can write a version of `greet`, shown in the following listing, that quits when the user gives an empty input.

Listing 11.23 Repeatedly greets a user, exiting on empty input (`RunIO.idr`)

```
greet : RunIO ()
greet = do putStrLn "Enter your name: "
          name <- getLine
          if name == ""
              then do putStrLn "Bye bye!"
                      Quit ()
              else do putStrLn ("Hello " ++ name)
                      greet
```

Annotations for Listing 11.23:

- A callout points to the `if` condition with the text "The input was empty, so displays a message and exits. Here, `greet` is terminating."
- A callout points to the `else` block with the text "There was input, so greets the user and continues. The recursive call follows IO actions, so `greet` is productive."

Depending on the input, `greet` is either terminating or productive. The totality checker accepts `greet` as total, because it satisfies the definition from section 11.1.3 that a total function either terminates or is productive for all well-typed inputs.

Before you can execute `greet`, you'll need to write a new version of `run` that translates a program in `RunIO` to a sequence of `IO` actions for the runtime system to execute. Previously, `run` would only terminate on running out of `Fuel`, but now there are two possible reasons for termination:

- Running out of `Fuel`, as before
- Exiting cleanly, with a result, if the process being executed invokes `Quit`

Interactive programs and infinite types

Using a potentially infinite data type like `RunIO` is not unique to Idris. A similar idea was described by Peter Hancock and Anton Setzer in their 2004 paper “Interactive programs and weakly final coalgebras in dependent type theory,” following on from their earlier work in describing interactive programs with dependent types.

The generic type `Inf` that you use in the type of potentially infinite structures follows a similar idea used in the Agda programming language, described by Nils Anders Danielsson in his 2010 paper “Total Parser Combinators.”

You can distinguish these results in the type:

```
run : Fuel -> RunIO a -> IO (Maybe a)
```

The possible results of `run` correspond to the two possible reasons for termination:

- If `run` runs out of `Fuel`, it returns `Nothing`.
- If `run` executes an action of the form `Quit value`, it returns `Just value`.

The following listing gives the new definition of `run` for `RunIO`.

Listing 11.24 Converting an expression of type `RunIO ty` to an executable action of type `IO ty` (`RunIO.idr`)

```
run : Fuel -> RunIO a -> IO (Maybe a)
run fuel (Quit value) = pure (Just value)           ← Terminates due to invoking the Quit command, so it returns Just value
run (More fuel) (Do c f) = do res <- c
                           run fuel (f res)           ← Terminates due to running out of fuel, so it returns Nothing
run Dry p = pure Nothing
```

Finally, you can write a `main` program that executes `greet` and discards the result:

```
partial
main : IO ()
main = do run forever greet
         pure ()
```

Because `run` now creates an `IO` action that produces a value of type `Maybe ()` when invoked with `greet`, and `main` is expected to create an action that produces a value of type `()`, you need to finish by calling `pure ()`. When you execute `main` at the REPL, you can now exit cleanly by entering an empty string:

```
*RunIO> :exec main
Enter your name: Edwin
Hello Edwin
Enter your name:      ← Empty string
Bye bye!
```

With `RunIO`, you’ve refined `InfIO` with the ability to terminate a process cleanly when required. This gives you more freedom to write interactive programs, but there’s

another way in which `RunIO` arguably gives you *too much* freedom. Specifically, a process described by `RunIO` is a sequence of arbitrary IO actions, giving you several possibilities, including these:

- Reading from and writing to the console
- Opening and closing files
- Opening network connections
- Deleting files

For the programs you've written in this chapter, you're only interested in the first of these. The others are not only unnecessary, but lead to the possibility of remote security vulnerabilities in the third case, and potentially destructive errors in the fourth.

One of the principles of type-driven development, which we discussed in chapter 1, is that we should aim to write types that describe as *precisely* as possible the values that inhabit that type. To conclude this chapter, therefore, we'll look at how we can refine the `RunIO` type to describe only the actions that are necessary to implement the arithmetic quiz.

11.3.2 Domain-specific commands

You only need two IO actions when implementing the arithmetic quiz: reading from and writing to the console. Instead of allowing interactive programs in `RunIO` to execute arbitrary actions, therefore, you can restrict them to only the actions you need. That is, you can prevent your program from executing any interactive action that's outside the *problem domain* in which you're working.

The next listing shows a refined `ConsoleIO` type that describes interactive programs that support only reading from and writing to the console.

Listing 11.25 Interactive programs supporting only console I/O (ArithCmd.idr)

```
data Command : Type -> Type where
    PutStr : String -> Command ()
    GetLine : Command String

data ConsoleIO : Type -> Type where
    Quit : a -> ConsoleIO a
    Do : Command a -> (a -> Inf (ConsoleIO b)) -> ConsoleIO b
    (>>=) : Command a -> (a -> Inf (ConsoleIO b)) -> ConsoleIO b
    (>>=) = Do
```

Defines the IO commands that are available, parameterized by the commands' return types.

Only IO operations defined in Command are allowed in ConsoleIO programs.

Effectively, `Command` defines an interactive interface that `ConsoleIO` programs can use. You can think of it as defining the *capabilities* or *permissions* of interactive programs, eliminating any unnecessary actions.

You now need to refine the implementation of `run` to be able to execute `ConsoleIO` programs.

Listing 11.26 Executing a ConsoleIO program (ArithCmd.idr)

```

runCommand : Command a -> IO a
runCommand (PutStr x) = putStr x
runCommand GetLine = getLine

run : Fuel -> ConsoleIO a -> IO (Maybe a)
run fuel (Quit val) = do pure (Just val)
run (More fuel) (Do c f) = do res <- runCommand c
                             run fuel (f res)
run Dry p = pure Nothing

```

Runs a console I/O command by using the corresponding IO action

Uses runCommand to execute the IO action given by the command c

DOMAIN-SPECIFIC LANGUAGES A *domain-specific language* (DSL) is a language that's specialized for a particular class of problems. DSLs typically aim to provide only the operations that are needed when working in a specific problem domain in a notation that's accessible to experts in that domain, while eliminating any redundant operations.

In a sense, ConsoleIO defines a DSL for writing interactive console programs, in that it restricts the programmer to only the interactive actions that are needed and eliminates unnecessary actions such as file processing or network communication.

Listing 11.27 shows how you can modify quiz to run as a ConsoleIO program. By looking at the type of quiz and the definitions of ConsoleIO and run, you have a guarantee that quiz will only execute console I/O actions. There's no way it can open or close files, communicate over a network, or perform any other kind of interactive operation.

Listing 11.27 The arithmetic quiz, written as a ConsoleIO program (ArithCmd.idr)

You'll allow the player to quit, so quiz returns the player's score on exit.

PutStr and GetLine are valid Commands in ConsoleIO.

```

quiz : Stream Int -> (score : Nat) -> ConsoleIO Nat
quiz (num1 :: num2 :: nums) score
= do PutStr ("Score so far: " ++ show score ++ "\n")
    PutStr (show num1 ++ " * " ++ show num2 ++ "? ")
    answer <- GetLine
    if toLower answer == "quit" then Quit score else
        if (cast answer == num1 * num2)
            then do PutStr "Correct!\n"
                    quiz nums (score + 1)
        else do PutStr ("Wrong, the answer is " ++
                        show (num1 * num2) ++ "\n")
                    quiz nums score

```

Entering “quit” will exit the quiz.

To complete the implementation, you'll need to implement a main function. The following listing shows a new implementation of main that executes the quiz and then displays the player's final score after they enter quit.

Listing 11.28 A main function that executes the quiz and displays the final score (ArithCmd.idr)

```
partial
main : IO ()
main = do seed <- time
    Just score <- run forever (quiz (arithInputs (fromInteger seed)) 0)
    | Nothing => putStrLn "Ran out of fuel"
    putStrLn ("Final score: " ++ show score) ←

This won't happen because you used
forever, but it's good practice to cover
all possible results of run, nevertheless. ←

→ Displays the score
produced from the
result of quiz
```

You can still refine the definition of `quiz` slightly. As functions get larger, it's good practice to break them down into smaller functions, each with clearly defined roles. Here, for example, you can lift out functions for reporting whether an answer is correct or wrong, as in the following listing. These functions must themselves be either productive (finishing by calling `quiz`, as they do here) or terminating if `quiz` is to remain total.

Listing 11.29 Lifting out components of `quiz` into separate functions (ArithCmd.idr)

```
mutual
  correct : Stream Int -> (score : Nat) -> ConsoleIO Nat
  correct nums score
    = do PutStr "Correct!\n"
        quiz nums (score + 1)

  wrong : Stream Int -> Int -> (score : Nat) -> ConsoleIO Nat
  wrong nums ans score
    = do PutStr ("Wrong, the answer is " ++ show ans ++ "\n")
        quiz nums score

  quiz : Stream Int -> (score : Nat) -> ConsoleIO Nat
  quiz (num1 :: num2 :: nums) score
    = do PutStr ("Score so far: " ++ show score ++ "\n")
        PutStr (show num1 ++ " * " ++ show num2 ++ "? ")
        answer <- GetLine
        if toLower answer == "quit" then Quit score else
          if (cast answer == num1 * num2)
            then correct nums score
            else wrong nums (num1 * num2) score
```

In a mutual block (see chapter 3), definitions can refer to each other. You need mutual because `correct` and `wrong` both call `quiz`, and vice versa.

After `quiz` type-checks successfully, you can make several guarantees about its behavior by looking at the types and checking for totality:

- It will execute no interactive action other than `putStr` and `getLine`.
- It will either exit immediately or execute at least one interactive action, because it is productive.
- Every action executed will return a result in finite time, because it is total.

11.3.3 Sequencing Commands with do notation

In implementing quiz, you used two types to construct the function: Command and ConsoleIO:

- Command describes single commands, which terminate.
- ConsoleIO describes sequences of terminating commands, which might be infinite.

So, you can have *single, finite* commands, or *sequences of infinite* commands. But it would also be useful to be able to construct composite commands; that is, sequences of commands that are guaranteed to terminate. For example, you might like to write a composite command that displays a prompt, and then reads and parses user input. The next listing shows a type that represents possible user inputs, and a skeleton definition of a function to read and parse input.

Listing 11.30 Defining a type for representing user input, and a composite command for reading and parsing input (ArithCmdDo.idr)

```
data Input = Answer Int
           | QuitCmd
```

An input is either a number that answers the question or a quit command.

```
readInput : (prompt : String) -> Command Input
readInput prompt = ?readInput_rhs
```

The return type, Command Input, means that readInput will not loop indefinitely.

To write this function, you need to do the following:

- 1 Display the prompt.
- 2 Read an input from the console.
- 3 Convert the input string to Input and return it.

Because Command currently supports only single commands, you'll need to extend it to support sequences of commands. The next listing shows the extended definition, including two new data constructors, Pure and Bind, and a correspondingly updated definition of runCommand.

Listing 11.31 Extending Command to allow sequences of commands (ArithCmdDo.idr)

A command that does nothing, returning a value

```
data Command : Type -> Type where
  PutStr : String -> Command ()
  GetLine : Command String
  Pure : ty -> Command ty
  Bind : Command a -> (a -> Command b) -> Command b
```

A sequence of two commands, taking the output of the first and passing it to the second

Runs the first command, then runs the second with the output of the first

You might also want to define `(>>=)` to support do notation for sequencing commands, but the following definition doesn't work as you might expect:

```
(>>=) : Command a -> (a -> Command b) -> Command b
(>>=) = Bind
```

If you try this, Idris will complain that `(>>=)` is already defined, because you've defined do notation for `ConsoleIO`:

```
ArithCmdDo.idr:22:7:Main.>>= already defined
```

Idris allows multiple definitions of the same function name, as long as they're in separate *namespaces*. You've already seen this with `List` and `Vect`, for example, where each has constructors called `Nil` and `(::)` that Idris disambiguates according to the context in which you use them.

The namespace is given by the module where you define functions. Namespaces are also *hierarchical*, so you can introduce further namespaces inside a module. You can have multiple definitions of functions called `(>>=)` in one module by introducing new namespaces for each. The following listing shows how you can define new namespaces for each `(>>=)`.

Listing 11.32 Creating two definitions of `(>>=)` in separate namespaces (`ArithCmdDo.idr`)

```
namespace CommandDo
  (>>=) : Command a -> (a -> Command b) -> Command b
  (>>=) = Bind

  Introduces a new,
  nested namespace

namespace ConsoleDo
  (>>=) : Command a -> (a -> Inf (ConsoleIO b)) -> ConsoleIO b
  (>>=) = Do
```

If you check the type of `(>>=)` at the REPL, you'll see all the definitions of `(>>=)` in their respective namespaces, with their types:

```
*ArithCmdDo> :t (>>=)
Main.CommandDo.(>>=) : Command a ->
  (a -> Command b) -> Command b
Main.ConsoleDo.(>>=) : Command a ->
  (a -> Inf (ConsoleIO b)) -> ConsoleIO b
Prelude.Monad.(>>=) : Monad m => m a -> (a -> m b) -> m b
```

Defining a Monad implementation for Command

You could also define an implementation of the `Monad` interface for `Command`, as described in chapter 7. When possible, this is often preferable because the `Prelude` and base libraries define several functions that work generically with `Monad` implementations. To do this, you'd also need to define implementations of `Functor` and `Applicative`. You'll see an example of how to do this for a similar type in the next chapter.

(continued)

You can't, however, define an implementation of `Monad` for `ConsoleIO`, because the type of `ConsoleDo.(>>=)` doesn't fit the type of the `(>>=)` method in the `Monad` interface.

Using `CommandDo.(>>=)` to provide `do` notation, you can complete the definition of `readInput`.

Listing 11.33 Implementing `readInput` by sequence Command (ArithCmdDo.idr)

```
readInput : (prompt : String) -> Command Input
readInput prompt = do PutStr prompt
                      answer <- GetLine
                      if toLower answer == "quit"
                        then Pure QuitCmd
                        else Pure (Answer (cast answer))
```

Finally, you can use `readInput` in the main `quiz` function to encapsulate the details of displaying a prompt and parsing the user's input, shown in the final definition.

Listing 11.34 Defining `quiz` using `readInput` as a composite command to display a prompt and read user input (ArithCmdDo.idr)

```
quiz : Stream Int -> (score : Nat) -> ConsoleIO Nat
quiz (num1 :: num2 :: nums) score
  = do PutStr ("Score so far: " ++ show score ++ "\n")
       input <- readInput (show num1 ++ " * " ++ show num2 ++ "? ")
  → case input of
      Answer answer => if answer == num1 * num2
                           then correct nums score
                           else wrong nums (num1 * num2) score
      QuitCmd => Quit score
```

You can case-split on input because `readInput` has parsed it as a more informative type than simply a `String`.

In this final definition, you distinguish between terminating sequences of commands (using `Command`), and potentially nonterminating console I/O programs (using `ConsoleIO`). Syntactically, you write functions the same way in each, but the type tells you whether the function is allowed to run indefinitely, or whether it must eventually terminate.

Exercises



- 1 Update `quiz` so that it keeps track of the total number of questions and then returns the number of correct answers and the number of questions attempted. A sample run might look something like this:

```
*ex_11_3> :exec
Score so far: 0 / 0
9 * 11? 99
Correct!
Score so far: 1 / 1
6 * 9? 42
Wrong, the answer is 54
Score so far: 1 / 2
10 * 2? 20
Correct!
Score so far: 2 / 3
7 * 2? quit
Final score: 2 / 3
```

- 2 Extend the `Command` type from section 11.3.2 so that it also supports reading and writing files.

Hint: Look at the types of `readFile` and `writeFile` in the Prelude to decide what types your data constructors should have.

- 3 Use your extended `Command` type to implement an interactive “shell” that supports the following commands:

- `cat [filename]`, which reads a file and displays its contents
- `copy [source] [destination]`, which reads a source file and writes its contents to a destination file
- `exit`, which exits the shell

11.4 Summary

- You can generate infinite data using `Inf` to state which parts of a structure are potentially infinite.
- A total function either terminates with a well-typed input or produces a prefix of a well-typed infinite result in finite time.
- You can process infinite structures by using finite data to determine how much of the infinite structure to use.
- The Prelude defines a `Stream` data type for constructing infinite lists.
- You can define processes as infinite sequences of `IO` actions.
- To execute an infinite process, you define a function that takes an argument stating how long the process should run.
- The partial function `forever` allows processes to run indefinitely by generating an infinite amount of `Fuel`, using the `Lazy` type.
- By implementing `(>>=)`, you can extend `do` notation for your own data types to make programs easier to read and write.
- You can mix finite and infinite structures within the same data type to define potentially infinite processes that may also terminate.

Writing programs with state



This chapter covers

- Using the `State` type to describe mutable state
- Implementing custom types for state management
- Defining and using records for global system state

Idris is a pure language, so variables are immutable. Once a variable is defined with a value, nothing can update it. This might suggest that writing programs that manipulate state is difficult, or even impossible, or that Idris programmers in general aren't interested in state. In practice, the opposite is true.

In type-driven development, a function's type tells you exactly what a function can do in terms of its allowed inputs and outputs. So, if you want to write a function that manipulates state, you can do that, but you need to be explicit about it in the function's type. In fact, we've already done this in earlier chapters:

- In chapter 4, we implemented an interactive data store using global state.
- In chapter 9, we implemented a word-guessing game using global state to hold the target word, the letters guessed, and the number of guesses still available.
- In chapter 11, we implemented an arithmetic game using global state to hold the user's current score.

In each case, we implemented state by writing a recursive function that took the current state of the overall program as an argument.

Almost all real-world applications need to manipulate state to some extent. Sometimes, as in the previous examples, state is *global* and is used throughout an application. Sometimes, state is *local* to an algorithm; for example, a graph-traversal algorithm will hold the nodes it has visited in local state to avoid visiting nodes more than once. In this chapter, we'll look at how we can manage mutable state in Idris, both for state that's local to an algorithm, and for representing overall system state.

STATE MANAGEMENT AND DEPENDENT TYPES We won't use many dependent types in this chapter; using dependent types in state introduces some complications, as well as opportunities for precision in describing *state transition systems* and *protocols*. We'll consider these opportunities in the next two chapters, but we'll begin here by looking at how state works in general.

Previously, we've used types to describe interactive programs in terms of sequences of commands, using `IO` in chapter 5 and `ConsoleIO` in chapter 11. Stateful programs can work in the same way, using types to describe the operations a stateful program can perform. We'll begin by looking at the generic `State` type defined in the Idris library and then at how we can define types like `State` ourselves. Finally, we'll look at how we can structure a complete application with state, refining the arithmetic quiz from chapter 11.

12.1 Working with mutable state

Even though Idris is a pure language, it's often useful to work with state. It's particularly useful, when writing functions with a complex data structure, such as a tree or a graph, to be able to read and write local state as you're traversing the structure. In this section, you'll see how to manage mutable state.

COMPARING STATE IN IDRIS AND HASKELL This section describes the `State` type for capturing local mutable state in Idris. If you're familiar with Haskell, you'll find that you can use `State` in Idris in the same way as Haskell, by importing `Control.Monad.State`. If you're familiar with the `State` type in Haskell, you can safely move ahead to section 12.2, where I'll describe a custom implementation of `State`.

In the previous chapter, you wrote a function to attach labels to elements of a list, taking the labels from a `Stream`. It's repeated for reference here.

Listing 12.1 Labeling each element of a List using a Stream (Streams.idr)

```
labelWith : Stream labelType -> List a -> List (labelType, a)
labelWith [] = []
labelWith (lbl :: lbls) (val :: vals) = (lbl, val) :: labelWith lbls vals
```

For every label in the stream, pair it with the corresponding element of the list.

In this section, you'll implement a function similar to `label` a binary tree. First, you'll see how to implement mutable state by hand, with each function returning a pair that stores an updated state along with the result of a computation. You'll then see how to encapsulate the state using a `State` type defined in the Idris base library. First, though, I'll describe the example of tree traversal in more detail.

12.1.1 The tree-traversal example

Figure 12.1 shows the result of a function that labels a binary tree, and we'll use this as a running example throughout this section. This function labels the nodes *depth first*, left to right, so the deepest, leftmost node takes the first label, and the deepest, rightmost node takes the last label.




Figure 12.1 Labeling a tree, depth first. Each node is labeled with an integer.

The following listing gives the definition of binary trees that we'll use for labeling, along with `testTree`, a representation of the specific tree we'll label (from the example in figure 12.1).

Listing 12.2 Definition of binary trees, and an example (TreeLabel.idr)

A tree node with a value, a left subtree, and a right subtree

```

data Tree a = Empty           ← An empty binary tree
            | Node (Tree a) a (Tree a)

testTree : Tree String        ← A representation of the example tree
testTree = Node (Node (Node Empty "Jim" Empty) "Fred"
                (Node Empty "Sheila" Empty)) "Alice"
                (Node Empty "Bob" (Node Empty "Eve" Empty))

→ flatten : Tree a -> List a
flatten Empty = []
flatten (Node left val right) = flatten left ++ val :: flatten right
  
```

Converts the tree to a list by traversing the tree, depth first, left to right

It's convenient to define `flatten` so that you can easily see the ordering in which the labels should be applied:

```
*TreeLabel> flatten testTree
["Jim", "Fred", "Sheila", "Alice", "Bob", "Eve"] : List String
```

Once you've written a `treeLabel` function that labels the nodes in the tree according to the elements in a stream, you should be able to run it as follows:

```
*TreeLabel> flatten (treeLabel testTree)
[(1, "Jim"),
 (2, "Fred"),
 (3, "Sheila"),
 (4, "Alice"),
 (5, "Bob"),
 (6, "Eve")] : List (Integer, String)
```

When you wrote the function to label lists, in listing 12.1, you had a direct correspondence between the structure of the stream of labels and the list you were labeling. That is, for every `(::)` in the `List`, you could take the first element of the stream as the label, and then recursively label the rest of the list. With `Tree`, it's a bit more complicated, because when you label the left subtree, you don't know in advance how many elements you'll need to take from the stream. Figure 12.2 illustrates labeling the subtrees in this example.




Figure 12.2 Labeling subtrees. The left subtree is labeled from 1 to 3, and the right subtree is labeled from 5 to 6.

Before labeling the right subtree, you need to know how many elements you took from the stream when labeling the left subtree. Not only does the labeling function need to return the labeled tree, it also needs to return some information about where to start labeling the rest of the tree.

A natural way to do this might be for the labeling function to take the stream of labels as an input, as before, and to return a *pair*, containing

- The labeled tree
- A new stream of labels that you can use to label the rest of the tree

We'll begin by implementing tree labeling this way, using a pair to represent the result of the operation and the state of the labels. Then, you'll see how the `State` type defined in the Idris library encapsulates the state management details in this kind of algorithm.

12.1.2 Representing mutable state using a pair

Listing 12.3 defines a helper function to label a tree with a stream of labels, and uses a pair to represent the state of the stream after each subtree is labeled. The helper function returns the unused part of the stream of labels, so that when you've labeled one subtree, you know where to start labeling the next subtree.

Listing 12.3 Labeling a Tree with a stream of labels (TreeLabel.idr)

```

treeLabelWith : Stream labelType -> Tree a ->
  (Stream labelType, Tree (labelType, a))
treeLabelWith Empty = (Empty, Empty)
treeLabelWith lbls (Node left val right)
  = let (lblsLeft, left_labelled) = treeLabelWith lbls left
    (lblsRight, right_labelled) = treeLabelWith lblsLeft right
      in
        (lblsRight, Node left_labelled (lblThis, val) right_labelled)

treeLabel : Tree a -> Tree (Integer, a)
treeLabel tree = snd (treeLabelWith [1..] tree)
  
```

Returns the unused portion of the stream, as well as the labeled tree

Labels the left subtree, which gives you a new subtree and a new stream. Uses the first element, `lblThis`, to label the current node.

Labels the right subtree, using the stream returned after labeling the left subtree

Initializes with a stream of integers, and returns only the labeled tree using `snd`

Returns the stream given by labeling the right subtree, and a labeled node

If you try labeling the example tree, you'll see that the labels are applied in the order you expect:

```
*TreeLabel> flatten (treeLabel testTree)
[(1, "Jim"),
 (2, "Fred"),
 (3, "Sheila"),
 (4, "Alice"),
 (5, "Bob"),
 (6, "Eve")] : List (Integer, String)
```

You can check that labeling also preserves the structure of the tree by omitting the call to `flatten`:

```
*TreeLabel> treeLabel testTree
Node (Node Empty (1, "Jim") Empty)
  (2, "Fred")
    (Node Empty (3, "Sheila") Empty)
  (4, "Alice")
    (Node Empty
      (5, "Bob")
        (Node Empty (6, "Eve") Empty)) : Tree (Integer, String)
```

In the current definition of `treeLabelWith`, you need to keep track of the state of the stream of labels. Labeling a subtree not only gives you a tree with labels attached to the nodes, it also gives you a new stream of labels for labeling the next part of the tree.

As you traverse the tree, you keep track of the state of the stream by passing it as an argument, and returning the updated state. Essentially, the function uses *local mutable state*, with the state explicitly threaded through the definition. Although this works, there are two problems with this approach:

- It's *error-prone* because you need to ensure that the correct state is propagated correctly to the recursive calls. Notice that `left_labelled` and `right_labelled` have the same type, for example, so it's easy to use the wrong one!
- It's *hard to read* because the details of the algorithm are hidden behind the details of state management.

It's often useful to have local mutable state, and like any concept you use regularly, it's a good idea to create an abstraction that captures the concept. You can improve the definition of `treeLabel`, making it less error-prone and more readable, by using a type that captures the concept of state explicitly.

12.1.3 State, a type for describing stateful operations

The only reason you're passing the stream of labels around in the definition of `treeLabelWith` is that when you encounter a value at a node, you need to associate it with a label value. In an imperative language, you might pass a mutable variable to `treeLabelWith` and update it as you encounter each node. Because Idris is a purely functional language, you don't have mutable variables, but the base library does provide a type for describing sequences of stateful operations, in the `Control.Monad.State` module. `Control.Monad.State` exports the following relevant definitions.

Listing 12.4 State and associated functions, defined in `Control.Monad.State`

A type describing sequences of stateful operations with a state of type <code>stateType</code> , producing a result of type <code>ty</code>	Runs a sequence of stateful operations, producing a pair of the result and the final state
→ <code>State : (stateType : Type) -> (ty : Type) -> Type</code> <code>runState : State stateType a -> stateType -> (a, stateType)</code>	
→ <code>get : State stateType stateType</code> <code>put : stateType -> State stateType ()</code>	Writes a new state, producing a result of type <code>()</code>
→ <code>(>>=) : State stateType a -> (a -> State stateType b) -> State stateType b</code>	Sequences get and put with do notation. There's a Monad implementation for State that defines <code>(>>=)</code> .

Reads the current state, producing a result of type `stateType`

Just as a value of `IO ty` describes a sequence of *interactive* operations that produce a value of type `ty`, a value of type `State Nat ty` describes a sequence of operations that read and write a mutable state of type `Nat`. Listing 12.5 gives a small example of a function that works with mutable state. It reads the state using `get` and then updates it using `put`, increasing the `Nat` state by the given value.

Listing 12.5 A stateful function that increases a state by a given value (State.idr)

```
import Control.Monad.State           ←———— Needed for the State type

increase : Nat -> State Nat ()
increase inc = do current <- get      ←———— Assigns the value of the state to current
                  put (current + inc)   ←———— Updates the value of the state
```

A value of type `State Nat ()` is a description of stateful operations using a state of type `Nat`. You can execute it using `runState` by passing it the operations and an initial state. For example, you can execute `increase 5` with an initial state of 89:

```
*State> runState (increase 5) 89
(), 94) : ((), Nat)
```

The result is a pair of the value produced by the stateful operations, in this case the unit value `()`, and the final state, in this case 94. There are also two variants of `runState`:

- `evalState`—Returns only the value produced by the sequence of operations:

```
*State> :t evalState
evalState : State stateType a -> stateType -> a

*State> evalState (increase 5) 89
() : ()
```

- `execState`—Returns only the final state after the sequence of operations:

```
*State> :t execState
execState : State stateType a -> stateType -> stateType

*State> execState (increase 5) 89
94 : Nat
```

Generic types of State

If you check the types of `get` and `put`, you'll see that they use constrained generic types:

```
*TreeLabelState> :t get
get : MonadState stateType m => m stateType

*TreeLabelState> :t put
put : MonadState stateType m => stateType -> m ()
```

This gives library authors more flexibility in defining stateful programs. The details of the `MonadState` interface are beyond the scope of this book, but you can read more in the Idris library documentation (<http://idris-lang.org/documentation>). In this example, you can read `m` as `State stateType`.

Although `State` encapsulates sequences of stateful operations, internally it's defined using pure functions. Essentially, it encapsulates the implementation pattern you used to pass the state around in `treeLabelWith`.

Using `State`, you can reimplement `treeLabelWith`, hiding the internal details of state management and only reading and updating the stream of labels when necessary.

12.1.4 Tree traversal with State

The next listing shows how you can define `treeLabelWith` by keeping the stream of labels as state, reading it to get the next label for a node.

**Listing 12.6 Defining `treeLabelWith` as a sequence of stateful operations
(`TreeLabelState.idr`)**

```
treeLabelWith : Tree a -> State (Stream labelType) (Tree (labelType, a))
treeLabelWith Empty = pure Empty
treeLabelWith (Node left val right)
  = do left_labelled <- treeLabelWith left
    -----> (this :: rest) <- get
    put rest
    right_labelled <- treeLabelWith right
    pure (Node left_labelled (this, val) right_labelled)
```

Gets the current stream of labels. You'll use
the first label, this, to label the current node.

Sets the new stream of labels to be
the tail of the current stream, rest

In this definition, you can see the details of the labeling algorithm more clearly than in the previous definition. Here, you leave the internal details of state management to the implementation of `State`.

In order to run this function and actually carry out the tree labeling, you'll need to provide an initial state. The following listing defines a `treeLabel` function that initializes the state with an infinite stream of integers, counting upwards from 1.

**Listing 12.7 Top-level function to label tree nodes, depth first, left to right
(`TreeLabelState.idr`)**

```
treeLabel : Tree a -> Tree (Integer, a)
treeLabel tree = evalState (treeLabelWith tree) [1..]
```

evalState discards the
final state, so it returns
only the labeled tree.

As before, you can test this at the REPL. It behaves the same way as the previous implementation of `treeLabel` on the test input:

```
*TreeLabelState> flatten (treeLabel testTree)
[(1, "Jim"),
 (2, "Fred"),
 (3, "Sheila"),
 (4, "Alice"),
 (5, "Bob"),
 (6, "Eve")] : List (Integer, String)
```

Like `IO`, which you first encountered in chapter 5, `State` gives you a way of writing functions with side effects (here, modifying mutable state) by *describing* sequences of operations and *executing* them separately:

- A value of type `IO ty` is a description of a sequence of interactive actions producing a result of type `ty`. It's executed by the runtime system by compilation, or by `:exec` at the REPL.
- A value of type `State stateType ty` is a description of a sequence of actions that read and write a state of type `stateType`, producing a result of type `ty`. It's executed by using one of `runState`, `execState`, or `evalState`.

Many interesting programs follow this pattern, defining a type for describing sequences of commands and a separate function for executing those commands. Indeed, you've already seen one in chapter 11, when you defined the `ConsoleIO` type for describing indefinitely running interactive programs. You'll see more examples in the remaining chapters, so in the rest of this chapter we'll look at how to implement custom types for representing state and interaction.

Exercises



- 1** Write a function that updates a state by applying a function to the current state:

```
update : (stateType -> stateType) -> State stateType ()
```

You should be able to use `update` to reimplement `increase`:

```
increase : Nat -> State Nat ()
increase x = update (+x)
```

You can test your answer at the REPL as follows:

```
*ex_12_1> runState (increase 5) 89
(((), 94) : (((), Nat))
```

- 2** Write a function that uses `State` to count the number of occurrences of `Empty` in a tree. It should have the following type:

```
countEmpty : Tree a -> State Nat ()
```

You can test your answer at the REPL with `testTree` as follows:

```
*ex_12_1> execState (countEmpty testTree) 0
7 : Nat
```

- 3** Write a function that counts the number of occurrences of both `Empty` and `Node` in a tree, using `State` to store the count of each in a pair. It should have the following type:

```
countEmptyNode : Tree a -> State (Nat, Nat) ()
```

You can test your answer at the REPL with `testTree` as follows:

```
*ex_12_1> execState (countEmptyNode testTree) (0, 0)
(7, 6) : (Nat, Nat)
```

12.2 A custom implementation of State

In the previous section, you saw that the `State` type gives you a generic way of implementing algorithms that use state. You used it to maintain a stream of labels as local mutable state, which you could access by reading (using `get`) and writing (using `put`) as necessary. Like `IO`, which separates the *description* of an interactive program from its execution at runtime, `State` separates the description of a stateful program from its execution with a concrete state.

We'll look at more examples of the same pattern, separating the description of a program from its execution, in the remaining chapters, so before we go any further, let's explore how we can define the `State` type ourselves, along with `runState` for executing stateful operations. In this section, you'll see one way of defining `State`, and how to provide implementations of some interfaces for `State`: `Functor`, `Applicative`, and `Monad`. By implementing these interfaces, you'll be able to use some generic library functions with `State`.

12.2.1 Defining State and runState

The following listing shows one way you could define `State` by hand, with a `Get` data constructor for describing the operation that reads state, and a `Put` data constructor for describing the operation that writes state.

Listing 12.8 A type for describing stateful operations (`TreeLabelType.idr`)

```
data State : (stateType : Type) -> Type -> Type where
    Get : State stateType stateType
    Put : stateType -> State stateType ()
    Pure : ty -> State stateType ty
    Bind : State stateType a -> (a -> State stateType b) -> State stateType b
```

An operation that produces a value

Describes the operation that gets the current state

Describes the operation that puts a new state

Sequences stateful operations, passing the result of the first as the input to the next

Naming conventions reminder

Remember that, by convention, type and data constructor names in Idris begin with a capital letter. I won't deviate from this convention here, so if you want the same names as those exported by `Control.Monad.State`, you'll need to define the following functions:

```
get : State stateType stateType
get = Get

put : stateType -> State stateType ()
put = Put

pure : ty -> State stateType ty
pure = Pure
```

You can support do notation for State by defining `(>=)`. You can do this either by implementing the Monad interface for `(>=)`, or by defining `(>=)` directly:

```
(>>=) : State stateType a -> (a -> State stateType b) -> State stateType b  
(>>=) = Bind
```

Implementing interfaces for State

Defining `(>=)` means that you can use do notation for programs in State. As you saw in chapter 7, `(>=)` is also a method of the Monad interface, which also requires implementations of the Functor and Applicative interfaces. It's defined here as a standalone function to avoid the need to implement Functor and Applicative first for this example.

Where possible, it's a very good idea to implement the Monad interface instead, because that gives you access to several constrained generic functions defined by the Idris library. For example, you'd be able to use the `when` function, which executes an operation when a condition is met, and `traverse`, which applies a computation across a structure. You'll see how to do this for `State` in section 12.2.2.

Using this version of `State`, and defining the functions `get`, `put`, and `pure`, which directly use the data constructors, listing 12.9 shows how you can define `treeLabelWith`. This version is exactly the same as the previous one, as you'd expect, because it uses the same names for the functions that manipulate the state.

Listing 12.9 Defining treeLabelWith as a sequence of stateful operations (TreeLabelTextType.idr)

```
treeLabelWith : Tree a -> State (Stream labelType) (Tree (labelType, a))
treeLabelWith Empty = Pure Empty
treeLabelWith (Node left val right)
    = do left_labelled <- treeLabelWith left
        (this :: rest) <- get
        put rest
        right_labelled <- treeLabelWith right
        pure (Node left labelled (this, val) right labelled)
```

Gets the next label for labeling
the node from the stream

In order to run it, you'll need to define a function that converts the description of the stateful operations into the tree-labeling function. The following listing shows the definition of `runState`, which takes a description of a stateful program and an initial state, and returns the value produced by the stateful program and a final state.

Listing 12.10 Running a labeling operation (TreeLabelType.idr)

```
runState : State stateType a -> (st : stateType) -> (a, stateType)
runState Get st = (st, st)                                ◀
runState (Put newState) st = (((), newState))           ◀-----
```

Returns a pair of the value produced by the stateful operations, and the final state

Put produces the unit value and updates the state to the new value.

Get produces the current state and leaves the state unchanged.

```

runState (Pure x) st = (x, st)
runState (Bind cmd prog) st = let (val, nextState) = runState cmd st in
                                runState (prog val) nextState

```

Runs the first stateful command and then runs the rest of the program with the updated state and the output of the first command

When you run sequences of stateful operations, defined using Bind, you need to take the nextState state returned by running cmd, and pass it to runState when executing the rest of the operations. You calculate the rest of the operations by taking the val returned by running cmd and passing it to prog. This encapsulates the state management that you had to implement by hand (three times!) in your first implementation of treeLabelWith, and it's similar to the way the Control.Monad.State module implements State itself.

The following listing shows a definition of treeLabel that uses the new implementation of treeLabelWith, initializing the stream with [1..].

Listing 12.11 The tree-labeling function, which calls run with an initial stream of labels (TreeLabelType.idr)

```

treeLabel : Tree a -> Tree (Integer, a)
treeLabel tree = fst (runState (treeLabelWith tree) [1..])

```

Uses fst to extract the labeled tree and discard the final state

12.2.2 Defining Functor, Applicative, and Monad implementations for State

Implementing ($\gg=$) for State means that you can use do notation, which gives a clear, readable notation for writing functions that describe sequences of operations. But do notation is all it gives us.

Rather than defining ($\gg=$) as a standalone function, it's a good idea to implement the Functor, Applicative, and Monad interfaces for State. In addition to providing do notation via the Monad interface, this will give you access to generic functions defined in the library. For example, when and traverse are generic functions. In the context of IO, they behave as follows:

- when evaluates a computation if a condition is True. It could be used to run some IO actions only on a specific user input.
- traverse is similar to map and applies a computation across a structure. For example, you could print every element of a List to the console.

You can find out more about these functions, especially their types, with :doc. The following listing shows them in action in the context of an IO computation.

Listing 12.12 Using when and traverse (Traverse.idr)

```

crew : List String
crew = ["Lister", "Rimmer", "Kryten", "Cat"]

```

```

main : IO ()
main = do putStrLn "Display Crew? "
          x <- getLine
          when (x == "yes") $
            do traverse putStrLn crew
               pure ()
putStrLn "Done"

```

Evaluates the computation after \$ only if this condition is true. \$ is the application operator.

For everything in the crew list, evaluates putStrLn

The application operator \$

Remember from chapter 10 that the \$ operator is an infix operator that applies a function to an argument. Its primary purpose is to reduce the need for bracketing. In listing 12.12, you could also have written the application of when with explicit brackets, as follows:

```

when (x == "yes")
  (do traverse putStrLn crew
     pure ())

```

If you implement Functor, Applicative, and Monad for State, you'll be able to use these and other similar functions in functions that use State. The next listing shows an example of what you can do, giving a function that adds Integers from a list to a running total, provided the Integer is positive. At the moment, this will fail to type-check.

Listing 12.13 Adding positive integers from a list to a state (StateMonad.idr)

Returns whether the integer was successfully added

```

addIfPositive : Integer -> State Integer Bool
addIfPositive val = do when (val > 0) $
                      do current <- get
                         put (current + val)
                      pure (val > 0)

addPositives : List Integer -> State Integer Nat
addPositives vals = do added <- traverse addIfPositive vals
                       pure (length (filter id added))

```

For every integer in vals, the value in added corresponds to whether that integer was added to the state.

Increments the state with the given value, if it's positive

filter id added gives the elements of added that are True, so this returns the number of successfully added Integers.

This will fail because you don't have implementations of Functor or Applicative for State:

```

StateMonad.idr:42:15:
When checking right hand side of addIfPositive with expected type
State Integer Bool

```

When checking an application of function Main.>>=:
Can't find implementation for Applicative (State Integer)

Your ultimate goal here is to implement Monad for State, which also requires implementations of Functor and Applicative. The next listing shows the Monad interface, as defined in the Prelude.

Listing 12.14 The Monad interface

```
interface Applicative m => Monad (m : Type -> Type) where
  (=>) : m a -> (a -> m b) -> m b
  join : m (m a) -> m a
```

“Flattens” a nested structure. See the sidebar.

Passes the output of the first operation as the input to the second

The join method

We haven’t looked at `join` in detail, but you can use it to flatten nested structures into a single structure. For example, there are implementations of `Monad` for `List` and `Maybe`, so you can try `join` on examples of each:

```
Idris> join [[1,2,3], [4,5,6]]
[1, 2, 3, 4, 5, 6] : List Integer

Idris> join (Just (Just "One"))
Just "One" : Maybe String

Idris> join (Just (Nothing {a=String}))
Nothing : Maybe String
```

For `List`, `join` will concatenate the nested lists. For `Maybe`, `join` will find the single value nested in the structure, if any.

Both methods, `(=>)` and `join`, have default definitions, so you can implement `Monad` by defining one or both of these. Here, we’ll only use `(=>)`.

If you want to provide an implementation for `Monad`, you also need to implement `Applicative`, because it’s a parent interface of `Monad`. Similarly, `Applicative` has a parent interface, `Functor`. The following listing shows both interfaces as defined in the Prelude.

Listing 12.15 The Functor and Applicative interfaces

```
interface Functor (f : Type -> Type) where
  map : (func : a -> b) -> f a -> f b

interface Functor f => Applicative (f : Type -> Type) where
  pure : a -> f a
  (<*>) : f (a -> b) -> f a -> f b
```

Applies a function to an argument, where the function and argument are inside a structure

The `(<*>)` method allows you to, for example, have a stateful function that returns a function (of type `a -> b`), have another stateful function that returns an argument (of type `a`), and apply the function to the argument.

You can begin by implementing `Functor` as follows:

- Type, define**—Write the implementation header and create a skeleton definition of `map`. Remember that you can create skeleton definitions for interface methods by pressing Ctrl-Alt-A with the cursor over the interface name:

```
Functor (State stateType) where
    map func x = ?Functor_rhs_1
```

- Type, refine**—The type of the `?Functor_rhs_1` hole tells you that `x` is a stateful computation:

```
stateType : Type
b : Type
a : Type
func : a -> b
x : State stateType a
-----
Functor_rhs_1 : State stateType b
```

You can continue the definition by extracting the value from the computation `x` using `Bind`:

```
Functor (State stateType) where
    map func x = Bind x (\val => ?Functor_rhs_1)
```

- Refine**—To complete the definition, pass `val` to `func`, and use `Pure` to convert the result to a stateful computation:

```
Functor (State stateType) where
    map func x = Bind x (\val => Pure (func val))
```

Listing 12.16 shows the definitions of `Applicative` and `Monad` for `State`. You implement `Applicative` in a way similar to `Functor`, using `Bind` to extract the values you need from stateful computations.

Listing 12.16 Implementing Applicative and Monad for State (StateMonad.idr)

```
Applicative (State stateType) where
    pure = Pure
    (<*>) f a = Bind f (\f' =>
        Bind a (\a' =>
            Pure (f' a')))

    Applies Bind directly
```

The diagram illustrates the implementation of `Applicative` and `Monad` for `State`. It shows the code for `Applicative` and `Monad` with annotations explaining the flow of state through `Bind` operations:

- Applicative Implementation:**
 - `pure` is defined as `Pure`.
 - `(<*>)` is defined as `Bind f (\f' => Bind a (\a' => Pure (f' a')))`.
 - A callout labeled "Applies Bind directly" points to the `Bind` operation in the `<*>` definition.
- Annotations for `Bind` operations:**
 - "Gets the function to apply from f" points to the `f'` parameter in the inner `Bind` expression.
 - "Gets the argument to pass to the function from a" points to the `a'` parameter in the inner `Bind` expression.
 - "Applies the function to the argument and creates a stateful computation containing the result" points to the final `Pure` expression.
- Monad Implementation:**
 - `Monad (State stateType) where`
 - `(>>=)` is defined as `Bind`.

You've used `Bind` in the implementations of `Functor` and `Applicative` because you don't have do notation available yet. You need a `Monad` implementation to provide it,

and you need to have `Functor` and `Applicative` implementations to have a `Monad` implementation.

But you could use `do` notation by defining all the implementations together, in a mutual block, as listing 12.17 shows. In a mutual block, definitions can refer to each other, so the implementations of `Functor` and `Applicative` can rely on the implementation of `Monad`.

**Listing 12.17 Defining Functor, Applicative, and Monad implementations together
(StateMonad.idr)**

```
mutual
  Functor (State stateType) where
    map func x = do val <- x
                    pure (func val)
  pure given by definition from Applicative.
  Applicative (State stateType) where
    pure = Pure
    ( <*> ) f a = do f' <- f
                      a' <- a
                      pure (f' a')
  Monad (State stateType) where
    ( >>= ) = Bind
```

Each implementation needs to begin in the same column to be within the scope of the mutual block.

do notation implementation given by (`>>=`) from Monad.

Now that you've implemented these interfaces, you can try the earlier definition of `addPositives` from listing 12.13:

```
*StateMonad> runState (addPositives [-4, 3, -8, 9, 8]) 0
(3, 20) : (Nat, Integer)
```

The Effects library: combining State, IO, and other side effects

Given that you have `Monad` implementations for `State` and `IO`, sequencing stateful computations and interactive computations respectively, it's reasonable to wonder whether you can sequence both at once, in the same function—what about interactive programs that also manipulate state?

You'll see one way to do this in the next section. As a more general solution, though, Idris provides a library called `Effects` that supports combining different kinds of side effects like `State` and `IO` in types, as well as other effects such as exceptions and nondeterminism. You can find more details in the `Effects` library tutorial (<http://idris-lang.org/documentation/effects>).

You've now seen how to encapsulate the details of state manipulation by describing sequences of stateful operations as `State`, and executing them using `runState`. You've also seen how to implement `Functor`, `Applicative`, and `Monad` for `State`.

The states you've used in the examples so far have been fairly small: a single stream of labels, or a single `Integer`. More generally, state can get fairly complex:

- State may consist of several complex components, stored as a *record*. We'll discuss this in the rest of this chapter, where you'll see how to write a complete interactive program with state.
- You might want to use a *dependent* type in your `State`, in which case updating the state will also update its type! For example, if you add an element to a `Vect 8 Int`, it would become a `Vect 9 Int`. We'll discuss this in the next chapter.

An advantage of writing a type that expresses sequences of commands, along with a function for running those commands, is that you can make the command type as precise as you need. As you'll see in the next section, you can describe precisely the set of commands you need for a specific application, including commands for interaction at the console and commands for reading and writing components of the application's state. In the next chapter, you'll see how you can precisely describe *in its type* the effect each command has on a system's state.

12.3 A complete program with state: working with records

In the previous chapter, you implemented an arithmetic quiz that presented multiplication problems to the user and kept track of the numbers of correct answers and questions asked. In this section, we'll write a refined version, with the following improvements:

- We'll add a difficulty setting, which specifies the largest number allowed when generating questions.
- Instead of passing the current score as an argument to the `quiz` function, maintaining the state by hand, we'll extend the `Command` type with commands for giving access to the current score and the difficulty setting.

To write this refined version, we'll need to rethink how to represent the game's state. We'll do this using record types. You've already seen some the use of records to represent the data store in chapter 6, with similar examples in chapters 7 and 10. As an application's state grows, though, it can make sense to divide its state into several hierarchical records.

You'll see how to define and use nested records, how to update records with a concise syntax, and how to use a record to store the state of the interactive quiz program. First, though, we'll revisit the `Command` type from chapter 11 and see how you can extend it to support reading and writing system state, in a way similar to the custom `State` type defined in the previous section.

12.3.1 Interactive programs with state: the arithmetic quiz revisited

In chapter 11, you defined a `Command` data type, representing the commands you could use in console I/O programs, and a `ConsoleIO` type, representing possibly infinite interactive processes. You used this to implement an arithmetic quiz, presenting

multiplication problems for a user to answer. Like State, which describes the operations Get and Put for reading and writing state, Command describes the operations GetLine and PutStr for reading from and writing to the console. The following listing recaps the definitions of Command and ConsoleIO.

Listing 12.18 Interactive programs supporting only console I/O (ArithState.idr)

```

data Command : Type -> Type where
    PutStr : String -> Command ()
    GetLine : Command String

    Pure : ty -> Command ty
    Bind : Command a -> (a -> Command b) -> Command b

data ConsoleIO : Type -> Type where
    Quit : a -> ConsoleIO a
    Do : Command a -> (a -> Inf (ConsoleIO b)) -> ConsoleIO b

namespace CommandDo
    (=>) : Command a -> (a -> Command b) -> Command b
    (=>) = Bind

namespace ConsoleDo
    (=>) : Command a -> (a -> Inf (ConsoleIO b)) -> ConsoleIO b
    (=>) = Do

```

Annotations for Listing 12.18:

- A callout box points to the definition of Command with the text: "Defines the valid commands for an interactive program".
- A callout box points to the definition of ConsoleIO with the text: "Interactive programs that either produce a result with Quit, or loop indefinitely".
- A callout box points to the definition of CommandDo with the text: "Defines (=>) to support do notation for Command and ConsoleIO".

IMPLEMENTING MONAD FOR COMMAND As with State, you could implement Functor, Applicative, and Monad for Command instead of implementing (=>) directly. As an exercise, try providing implementations of each. As I noted in the last chapter, however, you can't provide a Monad implementation for ConsoleIO because the type of ConsoleDo. (=>) doesn't fit.

Like runState, which takes a description of stateful operations and returns the result of executing those operations with an initial state, run takes a description of interactive operations and executes them in IO. Listing 12.19 recaps the run function. Remember that you limit how long interactive programs can run by using the Fuel type, and you add a nontotal function, forever, that allows a *total* interactive program to run indefinitely, while only introducing a single nontotal function.

Listing 12.19 Running interactive programs (ArithState.idr)

```

data Fuel = Dry | More (Lazy Fuel)

forever : Fuel
forever = More forever

runCommand : Command a -> IO a
runCommand (PutStr x) = putStr x
runCommand GetLine = getLine
runCommand (Pure val) = pure val
runCommand (Bind c f) = do res <- runCommand c
                           runCommand (f res)

```

An annotation points to the runCommand definition with the text: "Executes an interactive command".

```

run : Fuel -> ConsoleIO a -> IO (Maybe a)
run fuel (Quit val) = do pure (Just val)
run (More fuel) (Do c f) = do res <- runCommand c
                            run fuel (f res)
run Dry p = pure Nothing

```

←
Executes an interactive program as long as fuel remains

If you want your interactive programs to be able to read and write state, in addition to reading from and writing to the console, you can extend the `Command` type with additional commands for manipulating state, and process those commands just as you did with `State`. For the arithmetic game, we'll need to do the following:

- Get random numbers for the questions, given the game's difficulty setting.
- Read the current game state so that you can display the score.
- Update the game state so that you can update the score according to the user's answers.

The next listing shows how you can extend the `Command` data type to include these commands. There's no need to update `ConsoleIO`, because it merely sequences `Commands`.

Listing 12.20 Extending the `Command` type to support game state (`ArithState.idr`)

```

GameState : Type
data Command : Type -> Type where
    PutStr : String -> Command ()
    GetLine : Command String
    GetRandom : Command Int
    GetGameState : Command GameState
    PutGameState : GameState -> Command ()
    Pure : ty -> Command ty
    Bind : Command a -> (a -> Command b) -> Command b

```

←
This is a placeholder. You'll define `GameState` shortly.

←
Returns a random number based on the game's current difficulty level

←
Returns the current game state

←
Sets the game state

`GameState` is undefined for the moment, so you can't yet complete `run` or `runCommand`. You can, however, add pattern clauses with holes for the extra commands so that the totality checker is satisfied:

```

runCommand : Command a -> IO a
runCommand (PutStr x) = putStr x
runCommand GetLine = getLine
runCommand (Pure val) = pure val
runCommand (Bind c f) = do res <- runCommand c
                           runCommand (f res)
runCommand (PutGameState x) = ?runCommand_rhs_1
runCommand GetGameState = ?runCommand_rhs_2
runCommand GetRandom = ?runCommand_rhs_3

```

ADDING MISSING CASES After you've added constructors to `Command`, you can use Ctrl-Alt-A with the cursor over `runCommand` in the type declaration to add the new pattern clauses for `runCommand`.

You'll use the `GameState` type to store the game's state. Before you implement the refined quiz, therefore, you'll need to consider how to define `GameState`.

12.3.2 Complex state: defining nested records

In the refined quiz implementation, you'll use `GameState` to store the following:

- The current score, consisting of the number of questions answered correctly and the number attempted
- The difficulty setting

When there are several components to a program's state like this, it often makes sense to use a record type. Records are convenient because they give rise to projection functions, which allow you to inspect the fields of the record. You can also nest records; the following listing shows how you can represent the current score as a record, and the overall game state as a record, including the nested score record as a field.

Listing 12.21 Representing a game state as nested records (ArithState.idr)

```
record Score where
    constructor MkScore
    correct : Nat
    attempted : Nat

    ← Current score has fields for the
    number of correct answers and
    attempted questions

record GameState where
    constructor MkGameState
    score : Score
    difficulty : Int

    ← The game state record includes the
    score record as a field in that record.

initState : GameState
initState = MkGameState (MkScore 0 0) 12

    ← The initial game state is a score of 0
    out of 0, with a difficulty of 12.
```

Defining `Score` and `GameState` as records automatically generates projection functions for each field: `correct`, `attempted`, `score`, and `difficulty`. For example, you can get the difficulty level like this:

```
*ArithState> difficulty initState
12 : Int
```

Or you can get the number of correct answers so far:

```
*ArithState> correct (score initState)
0 : Nat
```

Records and namespaces

When you define a record, the projection functions are defined in their own namespace, given by the name of the record. For example, the `score` function is defined in a new `GameState` namespace, as you can see with `:doc` at the REPL:

```
*ArithState> :doc score
Main.GameState.score : (rec : GameState) -> Score
```

(continued)

This allows the same field name to be used multiple times within the same module. For example, you can use a field called `title` in two different records in the same file, `Record.idr`:

```
record Book where
    constructor MkBook
    title : String
    author : String

record Album where
    constructor MkAlbum
    title : String
    tracks : List String
```

Idris will decide which version of `title` to use, according to context:

```
*Record> title (MkBook "Breakfast of Champions" "Kurt Vonnegut")
"Breakfast of Champions" : String
```

The following listing shows how you can use projection functions to define an implementation of `Show` for `GameState`.

Listing 12.22 Show implementation for GameState (ArithState.idr)

```
Show GameState where
    show st = show (correct (score st)) ++ "/" ++
        show (attempted (score st)) ++ "\n" ++
        "Difficulty: " ++ show (difficulty st)
```

You can try this at the REPL, using `println` to display the initial game state:

```
*ArithState> :exec println initState
0/0
Difficulty: 12
```

Records, therefore, give you a convenient notation for *inspecting* field values, but when you write programs that use records to hold state, you'll also need to *update* fields. In the quiz, for example, you'll need to increment the score and the number of questions attempted.

12.3.3 Updating record field values

Idris is a pure functional language, so you won't update record fields in-place. Instead, when we say we're updating a record, we really mean that we're constructing a *new* record containing the contents of the old record with a single field changed. For example, the following listing shows one way to return a new record with the difficulty field updated, using pattern matching.

Listing 12.23 Setting a record field by pattern matching (ArithState.idr)

```
setDifficulty : Nat -> GameState -> GameState
setDifficulty newDiff (MkGameState score _) = MkGameState score newDiff
```

If the record has a lot of fields, this can get unwieldy very quickly, because you'd need to write update functions for every field. Not only that, if you were to add a field to a record, you'd need to modify *all* of the update functions. Idris therefore provides a built-in syntax for updating fields in records. Here's an implementation of `setDifficulty` using record-update syntax.

Listing 12.24 Setting a record field using record-update syntax (ArithState.idr)

```
setDifficulty : Nat -> GameState -> GameState
setDifficulty newDiff state = record { difficulty = newDiff } state
```

Figure 12.3 shows the components of the record-update syntax. Note in particular that the record-update syntax itself is first-class, where `record` is a keyword that begins a record update, so the record update has a type. Here, the update has a function type, `GameState -> GameState`, so you can also implement `setDifficulty` as follows:

```
setDifficulty : Nat -> GameState -> GameState
setDifficulty newDiff = record { difficulty = newDiff }
```




Figure 12.3 Syntax for returning a new record with a field updated

You can update nested record fields in a similar way, by giving the path to the field you'd like to update. The following listing shows how you can write the correct and wrong functions, which update the score.

Listing 12.25 Setting nested record fields using record-update syntax (ArithState.idr)

<pre>addWrong : GameState -> GameState addWrong state = record { score->attempted = attempted (score state) + 1 } state ←</pre>	Updates the number of attempted questions by adding 1 to the current value
---	---

```
addCorrect : GameState -> GameState
addCorrect state
```

```
    = record { score->correct = correct (score state) + 1,
                score->attempted = attempted (score state) + 1 } state
```

Updates the number of correct answers by adding 1 to the current value

You can update multiple fields in one go, separating the updates with a comma.

The `score->attempted` notation gives the path to the field you'd like to update, with the *outermost* field name first. So, in this example, you'd like to update the `attempted` field of the `score` field of the `state` record.

12.3.4 Updating record fields by applying functions

The record-update syntax offers a concise notation for specifying a path to a particular record field. It's still a little inconvenient, though, because you've needed to explicitly write the path to each field twice, in different notations:

- First, to find the field to update, using the `score->correct` path notation
- Second, to find the old value, using a function application, `correct (score state)`

Idris therefore provides a notation for updating record fields by directly applying a function to the current value of the field. The next listing shows a concise way of updating the nested record fields in `GameState`.

Listing 12.26 Updating nested record fields by directly applying functions to the current value of the field (`ArithState.idr`)

```
addWrong : GameState -> GameState
addWrong = record { score->attempted $= (+1) }

addCorrect : GameState -> GameState
addCorrect = record { score->correct $= (+1),
                      score->attempted $= (+1) }
```

The `$=` operator in the field update means that the new field value is calculated by applying the function `(+1)` to the current value.

UPDATING RECORDS WITH `$=` You saw the `$` operator, which applies a function to an argument, in chapter 10. The `$=` syntax arises from a combination of the function application operator `$` and the record-update syntax.

This syntax gives you a concise and convenient way of writing functions that update nested record fields, which makes it easier to write programs that manipulate state. Moreover, because it doesn't use pattern matching, it's independent of any other fields in a record, so even if you add fields to the `GameState` record, `addWrong` and `addCorrect` will work without modifications.

12.3.5 Implementing the quiz

Using your new `Command` type and the `GameState` record, you can implement the arithmetic quiz by reading and updating the state as necessary. The next listing shows an outline of the quiz implementation, leaving holes for `correct` and `wrong`, which, respectively, process a correct answer and a wrong answer.

Listing 12.27 Implementing the arithmetic quiz (ArithState.idr)

```

mutual
  correct : ConsoleIO GameState
  correct = ?correct_rhs

  wrong : Int -> ConsoleIO GameState
  wrong ans = ?wrong_rhs

  readInput : (prompt : String) -> Command Input
  quiz : ConsoleIO GameState
  quiz = do num1 <- GetRandom
            num2 <- GetRandom
            st <- GetGameState
            PutStr (show st ++ "\n")

            input <- readInput (show num1 ++ " * " ++ show num2 ++ "? ")
            case input of
              Answer answer => if answer == num1 * num2
                then correct
                else wrong (num1 * num2)
              QuitCmd => Quit st

```

Gets the game state so you can display the current score and difficulty

readInput was defined at the end of chapter 11. It displays a prompt and reads and parses user input.

Gets the next random number from the state

Processes a correct answer

Processes a wrong answer

This is similar to the implementation of `quiz` at the end of chapter 11, but instead of passing the stream of random numbers and the score as arguments, you treat each of them as state that you read and write as required. This simplifies the definition of `quiz`, at the cost of making the definition of `ConsoleIO` more complex.

The next listing shows how you can implement `correct` and `wrong`, each modifying the state using `addCorrect` and `addWrong`, respectively, as defined in the previous section.

Listing 12.28 Processing correct and wrong answers by updating the game state (ArithState.idr)

```

correct : ConsoleIO GameState
correct = do PutStr "Correct!\n"
            st <- GetGameState
            PutGameState (addCorrect st)
            quiz

            wrong : Int -> ConsoleIO GameState
            wrong ans = do PutStr ("Wrong, the answer is " ++ show ans ++ "\n")
                          st <- GetGameState
                          PutGameState (addWrong st)
            quiz

```

Continues with the quiz

Sets a new game state with an updated record

Sets a new game state with an updated record

At this stage, you have a complete description of an interactive arithmetic quiz that retrieves random numbers and the current score from the state. It's also total:

```
*ArithState> :total quiz
Main.quiz is Total
```

But to *run quiz*, you'll need to extend `runCommand` to support your new commands.

12.3.6 Running interactive and stateful programs: executing the quiz

As with the `runState` function you wrote for processing operations in the custom `State` type in section 12.2, the updated `run` function will need to process the current game state. It will also need to read from a stream of random numbers (for `GetRandom`) and perform console I/O. The following listing shows how these are all captured in a new type for `runCommand`.

Listing 12.29 A new type and skeleton definition for `runCommand` (`ArithState.idr`)

```
The stream of available random
numbers, before running the command
runCommand : Stream Int -> GameState -> Command a ->
IO (a, Stream Int, GameState)
runCommand = ?runCommand_rhs
```

The game state,
before running
the command

The command to run

Returns the result of the
command and an updated
stream of random numbers
and game state

Listing 12.30 gives the complete definition of `runCommand`. Note that, in each case, you need to return the result of the command as well as show how each command affects the random number stream and the game state.

Listing 12.30 The complete definition of `runCommand` (`ArithState.idr`)

```
runCommand : Stream Int -> GameState -> Command a ->
IO (a, Stream Int, GameState)
runCommand rnds state (PutStr x) = do putStr x
                                         pure ((()), rnds, state)
runCommand rnds state GetLine = do str <- getLine
                                         pure (str, rnds, state)
```

No change to
the random
number
stream or
game state

Takes the
first random
number from
the stream and
converts it
to a number
between 1 and
the difficulty
level

Remember to import
Data.Primitives.Views
to be able to use
divides.

The result of the command
is the current game state.

The argument to PutGameState
becomes the new game state.

```
runCommand (val :: rnds) state GetRandom
    = pure (getRandom val (difficulty state), rnds, state)
  where
    getRandom : Int -> Int -> Int
    getRandom val max with (divides val max)
      getRandom val 0 | DivByZero = 1
      getRandom ((max * div) + rem) max | (DivBy prf) = abs rem + 1

runCommand rnds state GetGameState
    = pure (state, rnds, state)
runCommand rnds state (PutGameState newState)
    = pure ((()), rnds, newState)
```

```

runCommand rnds state (Pure val)
  = pure (val, rnds, state)
runCommand rnds state (Bind c f)
  = do (res, newRnds, newState) <- runCommand rnds state c
       runCommand newRnds newState (f res) ← Takes the states from running
                                         the first command and passes
                                         them on to the next command
No change to the random number stream or game state

```

Similarly, you need to update `run` to take a stream of random integers and an initial game state. Like `runCommand`, `run` also returns the result of running the program, along with the updated stream and game state. Here's the new implementation of `run`, which supports the game state.

Listing 12.31 Running a ConsoleIO program consisting of a potentially infinite stream of Command (ArithState.idr)

```

run : Fuel -> Stream Int -> GameState -> ConsoleIO a ->
      IO (Maybe a, Stream Int, GameState)
run fuel rnds state (Quit val) = do pure (Just val, rnds, state)
run (More fuel) rnds state (Do c f)
  = do (res, newRnds, newState) <- runCommand rnds state c
       run fuel newRnds newState (f res)
run Dry rnds state p = pure (Nothing, rnds, state)

```

As in chapter 11, you use `Fuel` to say how long you're willing to allow a potentially infinite interactive program to run, so the portion of the return type that represents the result of running the program, of type `ConsoleIO a`, has type `Maybe a`, to capture the possibility of running out of `Fuel`.

Finally, the next listing shows how you update the main program to initialize the stream of random numbers and the game state.

Listing 12.32 A main program that initializes the arithmetic quiz with a random number stream and an initial state (ArithState.idr)

```

randoms : Int -> Stream Int ← Generates an infinite stream of random numbers
randoms seed = let seed' = 1664525 * seed + 1013904223 in
               (seed' `shiftR` 2) :: randoms seed'

partial                                     ← Needed because you use forever
main : IO ()
main = do seed <- time                  ← Remember to import System to be able to use time.
         (Just score, _, state) <-
             run forever (randoms (fromInteger seed)) initState quiz
             | _ => putStrLn "Ran out of fuel"
             putStrLn ("Final score: " ++ show state)

```

As with the implementation of `quiz` at the end of chapter 11, you separate terminating sequences of commands (using the `Command` type) from possible nonterminating sequences of console I/O operations (using the `ConsoleIO` type). In addition, you extend the `Command` type to allow reading and writing of the game's state much like

the implementation of State. If you define a specific type for the commands that an application can execute, you can make that type as precise as you need, possibly describing the effect of each operation on an abstraction of a system's state. You'll see more of what you can achieve by following this pattern and the guarantees you can make about stateful, interactive programs in the next chapter.

Exercises



- 1 Write an `updateGameState` function with the following type:

```
updateGameState : (GameState -> GameState) -> Command ()
```

You can test it by using it in the definitions of `correct` and `wrong` instead of `GetGameState` and `PutGameState`. For example:

```
correct : ConsoleIO GameState
correct = do PutStr "Correct!\n"
            updateGameState addCorrect
            quiz
```

- 2 Implement the Functor, Applicative, and Monad interfaces for `Command`.
- 3 You could define records for representing an article on a social news website as follows, along with the number of times the article has been upvoted or downvoted:

```
record Votes where
    constructor MkVotes
    upvotes : Integer
    downvotes : Integer

record Article where
    constructor MkArticle
    title : String
    url : String
    score : Votes

initPage : (title : String) -> (url : String) -> Article
initPage title url = MkArticle title url (MkVotes 0 0)
```

Write a function to calculate the overall score of a given article, where the score is calculated from the number of downvotes subtracted from the number of upvotes. It should have the following type:

```
getScore : Article -> Integer
```

You can test it with the following example articles:

```
badSite : Article
badSite = MkArticle "Bad Page" "http://example.com/bad" (MkVotes 5 47)

goodSite : Article
goodSite = MkArticle "Good Page" "http://example.com/good" (MkVotes 101 7)
```

At the REPL, you should see the following:

```
*ex_12_3> getScore goodSite
94 : Integer
```

```
*ex_12_3> getScore badSite  
-42 : Integer
```

- 4 Write `addUpvote` and `addDownvote` functions that modify an article's score up or down. They should have the following types:

```
addUpvote : Article -> Article  
addDownvote : Article -> Article
```

You can test these at the REPL as follows:

```
*ex_12_3> addUpvote goodSite  
MkArticle "Good Page"  
    "http://example.com/good"  
    (MkVotes 102 7) : Article  
  
*ex_12_3> addDownvote badSite  
MkArticle "Bad Page"  
    "http://example.com/bad"  
    (MkVotes 5 48) : Article  
  
*ex_12_3> getScore (addUpvote goodSite)  
95 : Integer
```

12.4 Summary

- Many algorithms read and write state. For example, when labeling the nodes of a tree in depth-first order, you can keep track of the state of the labels while traversing the tree.
- You can manage state by using the generic `State` type to describe operations on the state along with a `runState` function to execute those operations.
- You can define your own `State` type as a sequence of `Get` and `Put` operations.
- Defining `Functor`, `Applicative`, and `Monad` for the `State` type gives you access to several generic functions from the Idris library.
- When writing interactive programs with state, you can define a `Command` type to describe an interface consisting of console I/O and state management operations.
- Record types can represent more complex nested state.
- Idris provides a concise syntax for assigning new values to fields in nested records.
- Idris also provides a syntax (using `$=`) for updating a field in a nested record by applying a function to the value in that field.

13

State machines: verifying protocols in types

This chapter covers

- Specifying protocols in types
- Describing preconditions and postconditions of operations
- Using dependent types in state

In the previous chapter, you saw how to manage mutable state by defining a type for representing sequences of commands in a system, and a function for running those commands. This follows a common pattern: the data type *describes* a sequence of operations, and the function *interprets* that sequence in a particular context. For example, `State` describes sequences of stateful operations, and `runState` interprets those operations with a specific initial state.

In this chapter, we'll look at one of the advantages of using a type for describing sequences of operations and keeping the execution function separate. It allows you to make the descriptions more precise, so that certain operations can only be run when the state has a specific form. For example, some operations require access to a resource, such as a file handle or database connection, before they're executed:

- You need an open file handle to read from a file successfully.
- You need a connection to a database before you can run a query on the database.

When you write programs that work with resources like this, you're really working with a *state machine*. A database client might have two states, such as `Closed` and `Connected`, referring to its connection status to a database. Some operations (such as querying the database) are only valid in the `Connected` state; some (such as connecting to the database) are only valid in the `Closed` state; and some (such as connecting and closing) also change the state of the system. Figure 13.1 illustrates this system.




Figure 13.1 A state transition diagram showing the high-level operation of a database. It has two possible states, `Closed` and `Connected`. Its three operations, `Connect`, `Query`, and `Close`, are only valid in specific states.

State machines like the one illustrated in figure 13.1 exist, implicitly, in lots of real-world systems. When you're implementing communicating systems, for example, whether over a network or using concurrent processes, you need to make sure each party is following the same communication pattern, or the system could deadlock or behave in some other unexpected way. Each party follows a state machine where sending or receiving a message puts the overall system into a new state, so it's important that each party follows a clearly defined protocol. In Idris, we have an expressive type system, so if there's a model for a protocol, it's a good idea to express that in a type, so that you can use the type to help implement the protocol accurately.

In this chapter, you'll see how to make state machines like the one illustrated in figure 13.1 *explicit* in types. In this way, you can be sure that any function that correctly describes a sequence of actions follows the protocol defined by a state machine. Not only that, you can take a type-driven approach to defining sequences of actions using holes and interactive development. We'll begin with some fairly abstract examples to illustrate how you can describe state machines in types, modeling the states and operations on a door and a vending machine.

13.1 State machines: tracking state in types

You've previously implemented programs with state by defining a type that describes commands for reading and writing state. With dependent types, you can make the types of these commands more precise and include any relevant details about the state of the system in the type itself.

For example, let's consider how to represent the state of a door with a doorbell. A door can be in one of two states, open (represented as `DoorOpen`) or closed (represented as `DoorClosed`), and we'll allow three operations:

- Opening the door, which moves the system from the `DoorClosed` state to the `DoorOpen` state
- Closing the door, which moves the system from the `DoorOpen` state to the `DoorClosed` state
- Ringing the doorbell, which we'll only allow when the door is in the `DoorClosed` state

Figure 13.2 is a state transition diagram that shows the states the system can be in and how each operation modifies the overall state.




Figure 13.2 A state transition diagram showing the states and operations on a door

If you can define these state transitions in a type, then a well-typed description of a sequence of operations must correctly follow the rules shown in the state transition diagram. Furthermore, you'll be able to use holes and interactive editing to find out which operations are valid at a particular point in a sequence.

In this section, you'll see how to define state machines like the door in a dependent type. First, we'll implement a model of the door, and then we'll model more-complex states in a model of a simplified vending machine. In each case, we'll focus on the *model* of the state transitions, rather than a concrete implementation of the machine.

13.1.1 Finite state machines: modeling a door as a type

The state machine in figure 13.2 describes a *protocol* for correct use of a door by saying which operations are valid in which state, and how those operations affect the state. Listing 13.1 shows one way to represent the possible operations. This also includes a `(>=)` constructor for sequencing and a `Pure` constructor for producing pure values.

Listing 13.1 Representing operations on a door as a command type (Door.idr)

```

data DoorCmd : Type where
  Open : DoorCmd ()
  Close : DoorCmd ()
  RingBell : DoorCmd ()

  Pure : ty -> DoorCmd ty
  (>=) : DoorCmd a -> (a -> DoorCmd b) -> DoorCmd b
  
```

Changes the state of the door
from `DoorClosed` to `DoorOpen`

Changes the state of the door
from `DoorOpen` to `DoorClosed`

REMINDER: `(>=)` AND DO NOTATION Remember that do notation translates into applications of `(>=)`.

With `DoorCmd`, you can write functions like the following, which describes a sequence of operations for ringing a doorbell and opening and then closing the door, correctly following the door-usage protocol:

```
doorProg : DoorCmd ()  
doorProg = do RingBell  
             Open  
             Close
```

Unfortunately, you can also describe *invalid* sequences of operations that don't follow the protocol, such as the following, where you attempt to open a door twice, and then ring the doorbell when the door is already open:

```
doorProgBad : DoorCmd ()  
doorProgBad = do Open  
                 Open  
                 RingBell
```

You can avoid this, and limit functions with `DoorCmd` to valid sequences of operations that do follow the protocol, by keeping track of the door's state in the type of the `DoorCmd` operations. The following listing shows how to do this, describing exactly the state transitions represented in figure 13.2 in the types of the commands.

Listing 13.2 Modeling the door state machine in a type, describing state transitions in the types of the commands (Door.idr)

```

data DoorState = DoorClosed | DoorOpen

data DoorCmd : Type ->
    DoorState ->           ← The type of the result of the operation
    DoorState ->           ← The state of the door before the operation
    DoorState ->           ← The state of the door after the operation
    Type where
        Open : DoorCmd      () DoorClosed DoorOpen
        Close : DoorCmd     () DoorOpen   DoorClosed
        RingBell : DoorCmd  () DoorClosed DoorClosed

        Pure : ty -> DoorCmd ty state state
        (>>=) : DoorCmd a state1 state2 ->
            (a -> DoorCmd b state2 state3) ->
            DoorCmd b state1 state3

```

Defines the two possible states of a door

Produces a value without affecting the state

Sequences two operations. The output state of the first gives the input state of the second.

Combined operation goes from the input state of the first operation to the output state of the second

Each command's type takes three arguments:

- The type of the value produced by the command
 - The *input* state of the door; that is, the state the door must be in *before* you can execute the operation
 - The *output* state of the door; that is, the state the door will be in *after* you execute the operation

An implementation of the following function would therefore describe a sequence of actions that begins and ends with the door closed:

```
doorProg : DoorCmd () DoorClosed DoorClosed
```

ARGUMENT ORDER IN DOORCMD Notice that the type that a sequence of operations produces is the first argument to `DoorCmd`, and it's followed by the input and output states. This is a common convention when defining types for describing state transitions, and it will become important in chapter 14 when we look at more-complex state machines that deal with errors and feedback from the environment.

In general, if you have a value of type `DoorType ty beforeState afterState`, it describes a sequence of door actions that produces a value of type `ty`; it begins with the door in the state `beforeState`; and it ends with the door in the state `afterState`.

13.1.2 Interactive development of sequences of door operations

To see how the types in `DoorCmd` can help you write sequences of operations correctly, let's reimplement `doorProg`. We'll write this in the same way as before: ring the doorbell, open the door, and close the door.

If you write it incrementally, you'll see how the type shows the changes in the state of the door throughout the sequence of actions:

- 1 *Define*—Begin with the skeleton definition:

```
doorProg : DoorCmd () DoorClosed DoorClosed
doorProg = ?doorProg_rhs
```

- 2 *Refine, type*—Add an action to ring the doorbell:

```
doorProg : DoorCmd () DoorClosed DoorClosed
doorProg = do RingBell
            ?doorProg_rhs
```

If you check the type of `?doorProg_rhs` now, you'll see that it should be a sequence of actions that begins and ends with the door in the `DoorClosed` state:

```
-----
doorProg_rhs : DoorCmd () DoorClosed DoorClosed
```

- 3 *Refine, type*—Next, add an action to open the door:

```
doorProg : DoorCmd () DoorClosed DoorClosed
doorProg = do RingBell
            Open
            ?doorProg_rhs
```

If you check the type of `?doorProg_rhs` now, you'll see that it should begin with the door in the `DoorOpen` state instead:

```
-----
doorProg_rhs : DoorCmd () DoorOpen DoorClosed
```

- 4 Refine** failure—If you add an extra Open now, with the door already in the DoorOpen state, you’ll get a type error:

```
doorProg : DoorCmd () DoorClosed DoorClosed
doorProg = do RingBell
            Open
            Open
            ?doorProg_rhs
```

The error says that the type of Open is an operation that starts in the DoorClosed state, but the expected type starts in the DoorOpen state:

```
Door.idr:20:15:
When checking right hand side of doorProg with expected type
    DoorCmd () DoorClosed DoorClosed

When checking an application of constructor Main.>>=:
    Type mismatch between
        DoorCmd () DoorClosed DoorOpen (Type of Open)
    and
        DoorCmd a DoorOpen state2 (Expected type)

    Specifically:
        Type mismatch between
            DoorClosed
        and
            DoorOpen
```

- 5 Refine**—Instead, complete the definition by closing the door:

```
doorProg : DoorCmd () DoorClosed DoorClosed
doorProg = do RingBell
            Open
            Close
```

Defining preconditions and postconditions in types

The type of doorProg includes input and output states that give preconditions and postconditions for the sequence (the door must be closed both before and after the sequence). If the definition violates either, you’ll get a type error.

For example, you might forget to close the door:

```
doorProg : DoorCmd () DoorClosed DoorClosed
doorProg = do RingBell
            Open
```

In this case, you’ll get a type error:

```
Door.idr:18:15:
When checking right hand side of doorProg with expected type
    DoorCmd () DoorClosed DoorClosed

When checking an application of constructor Main.>>=:
    Type mismatch between
        DoorCmd () DoorClosed DoorOpen (Type of Open)
```

(continued)

and

```
DoorCmd () DoorClosed DoorClosed (Expected type)
```

Specifically:

Type mismatch between

DoorOpen

and

DoorClosed

The error refers to the final step and says that Open moves from DoorClosed to DoorOpen, but the expected type is to move from DoorClosed to DoorClosed.

By defining DoorCmd in this way, with the input and output states explicit in the type, you've defined what it means for a sequence of door operations to be valid. And by writing doorProg incrementally, with a sequence of steps and a hole for the rest of the definition, you can see the state of the door at each stage by looking at the type of the hole.

The door has exactly two states, DoorClosed and DoorOpen, and you can describe exactly when you change states from one to the other in the types of the door operations. But not all systems have an exact number of states that you can determine in advance. Next, we'll look at how you can model systems with an infinite number of possible states.

13.1.3 Infinite states: modeling a vending machine

In this section, we'll model a vending machine using type-driven development, writing types that explicitly describe the input and output states of each operation. As a simplification, the machine accepts only one type of coin (a £1 coin) and dispenses one product (a chocolate bar). Even so, there could be an arbitrarily large number of coins or chocolate bars in the machine, so the number of possible states is not finite.

Table 13.1 describes the basic operations of a vending machine, along with the state of the machine before and after each operation.

Table 13.1 Vending machine operations, with input and output states represented as Nat

Coins (before)	Chocolate (before)	Operation	Coins (after)	Chocolate (after)
pounds	chocs	Insert coin	S pounds	chocs
S pounds	S chocs	Vend chocolate	pounds	chocs
pounds	chocs	Return coins	Z	chocs

As with the door example, each operation has a precondition and a postcondition:

- *Precondition*—The number of coins and amount of chocolate that must be in the machine before the operation

- *Postcondition*—The number of coins and amount of chocolate in the machine after the operation.

You can represent the state of the machine as a pair of two Nats, the first representing the number of coins in the machine and the second representing the number of chocolates:

```
VendState : Type
VendState = (Nat, Nat)
```

The next listing shows a representation of the vending machine state as an Idris type, with the state transitions from table 13.1 explicitly written in the types of the MachineCmd operations.

Listing 13.3 Modeling the vending machine in a type, describing state transitions in the types of commands (Vending.idr)

```
VendState : Type
VendState = (Nat, Nat)           A type synonym for the machine state:  
                                a pair of the number of £1 coins and  
                                the number of chocolates

data MachineCmd : Type ->
    VendState -> VendState          Machine state before the  
                                    operation (precondition)

                                         VendState -> VendState          Machine state after the  
                                         Type where                  operation (postcondition)
                                         InsertCoin : MachineCmd () (pounds, chocs)      (S pounds, chocs)
                                         Vend       : MachineCmd () (S pounds, S chocs) (pounds, chocs)
                                         GetCoins   : MachineCmd () (pounds, chocs)      (Z, chocs)
```

To complete the model, you'll need to be able to sequence commands. You'll also need to be able to read user input: the commands you're defining describe what the *machine* does, but there's also a user interface that consists of the following:

- A coin slot
- A vend button, for dispensing chocolate
- A change button, for returning any unused coins

You can model these operations in a data type for describing possible user inputs. Listing 13.4 shows the complete model of the vending machine, including additional operations for displaying a message (Display), refilling the machine (Refill), and reading user actions (GetInput).

Listing 13.4 The complete model of vending machine state (Vending.idr)

```
data Input = COIN               Defines the possible
                           | VEND
                           | CHANGE
                           | REFILL Nat
```

```

data MachineCmd : Type -> VendState -> VendState -> Type where
  InsertCoin : MachineCmd () (pounds, chocs)      (S pounds, chocs)
  Vend       : MachineCmd () (S pounds, S chocs)   (pounds, chocs)
  GetCoins   : MachineCmd () (pounds, chocs)      (Z, chocs)
  Refill     : (bars : Nat) ->
                MachineCmd () (Z, chocs)           (Z, bars + chocs)

  Display    : String -> MachineCmd () state state
  GetInput   : MachineCmd (Maybe Input) state state

  Pure : ty -> MachineCmd ty state state
  (>>=) : MachineCmd a state1 state2 ->
            (a -> MachineCmd b state2 state3) ->
            MachineCmd b state1 state3

```

Refilling the machine is only valid if there are no coins in the machine.

Displaying a message doesn't affect the state.

Reading user input doesn't affect the state. Returns Maybe Input to account for possible invalid inputs.

```

data MachineIO : VendState -> Type where
  Do : MachineCmd a state1 state2 ->
        (a -> Inf (MachineIO state2)) -> MachineIO state1

namespace MachineDo
  (>>=) : MachineCmd a state1 state2 ->
            (a -> Inf (MachineIO state2)) -> MachineIO state1
  (>>=) = Do

```

An infinite sequence of machine state transitions. The type gives the starting state of the machine.

Supports do notation for infinite sequences of machine state transitions

13.1.4 A verified vending machine description

Listing 13.5 shows the outline of a function that describes verified sequences of operations for a vending machine using the state transitions defined by `MachineCmd`. As long as it type-checks, you know that you've correctly sequenced the operations, and you'll never execute an operation without its precondition being satisfied.

Listing 13.5 A main loop that reads and processes user input to the vending machine (Vending.idr)

```

mutual
  vend : MachineIO (pounds, chocs)
  vend = ?vend_rhs

  refill : (num : Nat) -> MachineIO (pounds, chocs)
  refill = ?refill_rhs

  machineLoop : MachineIO (pounds, chocs)
  machineLoop =
    do Just x <- GetInput
       | Nothing => do Display "Invalid input"
                     machineLoop
    case x of
      COIN => do InsertCoin
                  machineLoop
      VEND => vend

```

User input could be invalid, so check here.

vend and refill need to check their preconditions are satisfied.

A pattern-matching binding alternative (see chapter 5). This branch is executed if GetInput returns Nothing.

```

CHANGE => do GetCoins
            Display "Change returned"
            machineLoop
REFILL num => refill num

```

There are holes for vend and refill. In each case, you need to check that the number of coins and chocolates satisfy their preconditions. If you try to Vend without checking the precondition, Idris will report an error:

```

vend : MachineIO (pounds, chocs)
vend = do Vend
          ←
          Display "Enjoy!"
          machineLoop

```

Doesn't type-check because there may not be coins or chocolate in the machine

Idris will report an error because you haven't checked whether there's a coin in the machine and a chocolate bar available, so the precondition might not be satisfied:

```

Vending.idr:67:13:
When checking right hand side of vend with expected type
    MachineIO (pounds, chocs)

When checking an application of function Main.MachineDo.>>=:
    Type mismatch between
        MachineCmd ()
                    (S pounds1, S chocs2)
                    (pounds1, chocs2) (Type of Vend)
    and
        MachineCmd () (pounds, chocs) (pounds1, chocs2) (Expected type)

    Specifically:
    Type mismatch between
        S chocs1
    and
        chocs

```

The error says that the input state must be of the form `(S pounds1, S chocs2)`, but instead it's of the form `(pounds, chocs)`.

You can solve this problem by pattern matching on the implicit arguments, pounds and chocs, to ensure they're in the right form, or display an error otherwise. The following listing shows definitions of vend and refill that do this.

Listing 13.6 Adding definitions of vend and refill that check that their preconditions are satisfied (Vending.idr)

```

vend : MachineIO (pounds, chocs)
vend {pounds = S p} {chocs = S c} ←
  = do Vend
      Display "Enjoy!"
      machineLoop
vend {pounds = Z} ← No money in the machine; can't vend
  = do Display "Insert a coin"
      machineLoop
vend {chocs = Z} ← No chocolate in the machine; can't vend
  = do Display "Out of stock"

```

A coin and a chocolate are available, so vend and continue.

```

machineLoop

refill : (num : Nat) -> MachineIO (pounds, chocs)
refill {pounds = z} num
  = do Refill num
       machineLoop
refill _ = do Display "Can't refill: Coins in machine"
             machineLoop

```

←

Refill only allows restocking with chocolate when there are no coins in the machine.

With both the door and the vending machine, we've used types to *model* the states of a physical system. In each case, the type gives an abstraction of the state a system is in before and after each operation, and values in the type describe the valid sequences of operations. We haven't implemented a run function to execute the state transitions for either DoorCmd or MachineCmd, but in the code accompanying this book, which is available online, you'll find code that implements a console simulation of the vending machine.

In the next section, you'll see a more concrete example of tracking state in the type, implementing a stack data structure. I'll use this example to illustrate how you can execute commands in practice.

Exercises



- 1 Change the RingBell operation so that it works in any state, rather than only when the door is closed. You can test your answer by seeing that the following function type-checks:

```

doorProg : DoorCmd () DoorClosed DoorClosed
doorProg = do RingBell
             Open
             RingBell
             Close

```

- 2 The following (incomplete) type defines a command for a guessing game, where the input and output states are the number of remaining guesses allowed:

```

data GuessCmd : Type -> Nat -> Nat -> Type where
  Try : Integer -> GuessCmd Ordering ?in_state ?out_state

  Pure : ty -> GuessCmd ty state state
  (=>) : GuessCmd a state1 state2 ->
    (a -> GuessCmd b state2 state3) ->
    GuessCmd b state1 state3

```

The Try command returns an Ordering that says whether the guess was too high, too low, or correct, and that changes the number of available guesses. Complete the type of Try so that you can only make a guess when there's at least one guess allowed, and so that guessing reduces the number of guesses available.

If you have a correct answer, the following definition should type-check:

```

threeGuesses : GuessCmd () 3 0
threeGuesses = do Try 10

```

```
Try 20
Try 15
Pure ()
```

Also, the following definition shouldn't type-check:

```
noGuesses : GuessCmd () 0 0
noGuesses = do Try 10
              Pure ()
```

- 3 The following type defines the possible states of matter:

```
data Matter = Solid | Liquid | Gas
```

Define a `MatterCmd` type in such a way that the following definitions type-check:

```
iceSteam : MatterCmd () Solid Gas
iceSteam = do Melt
              Boil

steamIce : MatterCmd () Gas Solid
steamIce = do Condense
              Freeze
```

Additionally, the following definition should *not* type-check:

```
overMelt : MatterCmd () Solid Gas
overMelt = do Melt
              Melt
```

13.2 Dependent types in state: implementing a stack

You've seen how to model state transitions in a type for two abstract examples: a door (representing whether it was open or closed in its type) and a vending machine (representing its contents in its type). Storing this abstract information in the type of the operations is particularly useful when you also have *concrete* data that relates to that abstract data. For example, if you're describing data of a specific size, and the type of an operation tells you how it changes the size of the data, you can use a `Vect` as a concrete representation. You'll know the required length of the input and output `Vect` from the type of each operation.

In this section, you'll see how this works by implementing operations on a *stack* data structure. A stack is a last-in, first-out data structure where you can add items to and remove them from the top of the stack, and only the top item is ever accessible. A stack supports three operations:

- Push—Adds a new item to the top of the stack
- Pop—Removes the top item from the stack, provided that the stack isn't empty
- Top—Inspects the top item on the stack, provided that the stack isn't empty

Like the operations on the vending machine, each of these operations has a precondition that describes the necessary input state and a postcondition describing the output state. Table 13.2 describes these, giving the required stack size before each operation and the resulting stack size after the operation.

Table 13.2 Stack operations, with input and output stack sizes represented as Nat

Stack size (before)	Operation	Stack size (after)
height	Push element	S height
S height	Pop element	height
S height	Inspect top element	S height

You'll express the preconditions and postconditions in the types of each operation. Once you've defined the operations on a stack, you'll implement a function to run sequences of stack operations using a concrete representation of a stack with its height in its type. Because you're using the stack's height in the state transitions, a good concrete representation of a stack is a `Vect`. You know, for example, that a stack of `Integer` of height 10, contains exactly 10 integers, so you can represent this as a value of type `Vect 10 Integer`.

Finally, you'll see an example of a stack in action, implementing a stack-based interactive calculator.

13.2.1 Representing stack operations in a state machine

As with `DoorCmd` and `MachineCmd` in section 13.1, we'll describe operations on a stack in a dependent type and put the important properties of the input and output states explicitly in the type. Here, the property of the state that interests us is the height of the stack.

Listing 13.7 shows how you can express the operations in table 13.2 in code, describing how each operation affects the height of the stack. For this example, you'll only store `Integer` values on the stack, but you could extend `StackCmd` to allow generic stacks by parameterizing over the element type in the stack.

Listing 13.7 Representing operations on a stack data structure with the input and output heights of the stack in the type (Stack.idr)

You'll use a `Vect` to represent the stack, so import `Data.Vect` here.

→ import Data.Vect

```
data StackCmd : Type -> Nat -> Nat -> Type where
    Push : Integer -> StackCmd () height (S height)
    Pop : StackCmd Integer (S height) height
    Top : StackCmd Integer (S height) (S height)

    Pure : ty -> StackCmd ty height height
    (>>=) : StackCmd a height1 height2 ->
        (a -> StackCmd b height2 height3) ->
        StackCmd b height1 height3
```

Pop requires there to be at least one element on the stack, and it decreases the height of the stack by 1.

Push increases the height of the stack by 1.

Top requires there to be at least one element on the stack, and it preserves the height of the stack.

You're using a `Vect` to represent the stack, so every time you add an element to the vector or remove an element, you'll change the vector's type. You're therefore representing dependently typed mutable state by putting the relevant arguments to the type (the length of the `Vect`) in the `StateCmd` type itself.

Using `StackCmd`, you can write sequences of stack operations where the input and output heights of the stack are explicit in the types. For example, the following function pushes two integers, pops two integers, and then returns their sum:

```
testAdd : StackCmd Integer 0 0
testAdd = do Push 10
            Push 20
            val1 <- Pop
            val2 <- Pop
            Pure (val1 + val2)
```

The types of the constructors in `StackCmd` ensure that there will always be an element on the stack when you try to `Pop`. For example, if you only push one integer in `testAdd`, Idris will report an error:

```
testAdd : StackCmd Integer 0 0
testAdd = do Push 10
            val1 <- Pop
            val2 <- Pop
            Pure (val1 + val2)
```

There's only one element on the stack, so `Pop` doesn't type-check.

When you try to define `testAdd` like this, Idris reports an error:

```
Stack.idr:27:22:
When checking right hand side of testAdd with expected type
  StackCmd Integer 0 0

When checking an application of constructor Main.>>=:
  Type mismatch between
    StackCmd Integer (S height) height (Type of Pop)
  and
    StackCmd a 0 height2 (Expected type)

  Specifically:
    Type mismatch between
      S height
    and
      0
```

This error, and particularly the mismatch between `S height` and `0`, means that you have a stack of height `0`, but `Pop` needs a stack that contains at least one element.

This approach is similar to the stateful functions defined in chapter 12, here using `Push` and `Pop` to describe how you're modifying and querying the state. As with the earlier descriptions of sequences of stateful operations, you'll need to write a separate function to run those sequences.

13.2.2 Implementing the stack using Vect

Listing 13.8 shows how to implement a function that executes stack operations. This is similar to runState, which you saw in chapter 12, but here you take an input Vect of the correct height as the contents of the stack.

Listing 13.8 Executing a sequence of actions on a stack, using a Vect to represent the stack's contents

```
runStack : (stk : Vect inHeight Integer) ->
           StackCmd ty inHeight outHeight ->
           (ty, Vect outHeight Integer)
runStack stk (Push val) = ((), val :: stk)
runStack (val :: stk) Pop = (val, stk)
runStack (val :: stk) Top = (val, val :: stk)

runStack stk (Pure x) = (x, stk)
runStack stk (cmd >>= next)
  = let (cmdRes, newStk) = runStack stk cmd in
    runStack newStk (next cmdRes)
```

If you try runStack with testAdd, passing it an initial empty stack, you'll see that it returns the sum of the two elements that you push, and that the final stack is empty:

```
*Stack> runStack [] testAdd
(30, []) : (Integer, Vect 0 Integer)
```

You can also define functions like the following, which adds the top two elements on the stack, putting the result back onto the stack:

```
doAdd : StackCmd () (S (S height)) (S height)
doAdd = do val1 <- Pop
          val2 <- Pop
          Push (val1 + val2)
```

The input state S (S height) means that the stack must have at least two elements on it, but, otherwise, it could be any height. If you try executing doAdd with an initial stack containing two elements, you'll see that it results in a stack containing a single element that's the sum of the two input elements:

```
*Stack> runStack [2,3] doAdd
((), [5]) : ((), Vect 1 Integer)
```

If the input state contains more than two elements, you'll see that it results in a stack with a height one smaller than the input height. For example, an input stack of [2, 3, 4] results in an output stack with the value [2 + 3, 4]:

```
*Stack> runStack [2,3,4] doAdd
((), [5, 4]) : ((), Vect 2 Integer)
```

You can add the two elements on the resulting stack with another call to doAdd:

```
*Stack> runStack [2,3,4] (do doAdd; doAdd)
(((), [9])) : (((), Vect 1 Integer)
```

But trying one more doAdd would result in a type error, because there's only one element left on the stack:

```
*Stack> runStack [2,3,4] (do doAdd; doAdd; doAdd)
(input):1:34:When checking an application of constructor Main.>>=:
  Type mismatch between
    StackCmd () (S (S height)) (S height) (Type of doAdd)
  and
    StackCmd ty 1 outHeight (Expected type)

  Specifically:
    Type mismatch between
      S height
    and
      0
```

This error means that you needed `S (S height)` elements on the stack (that is, at least two elements) but you only had `S height` (that is, at least one, but not necessarily any more). By putting the height of the stack in the type, therefore, you've explicitly specified the preconditions and postconditions on each operation, so you get a type error if you violate any of these.

13.2.3 Using a stack interactively: a stack-based calculator

If you add commands for reading from and writing to the console, you can write a console application for manipulating the stack and implement a stack-based calculator. A user can either enter a number, which pushes the number onto the stack, or add, which adds the top two stack items, pushes the result onto the stack, and displays the result. A typical session might go as follows:

```
*StackIO> :exec
> 3
> 4
> 5
User pushes three values
onto the stack: 3, 4 and 5

> add
9
> add
12
> add
          ← Error, because there's only one item (12) on the
Fewer than two items on the stack
```

Annotations in the session transcript:

- A bracket labeled "User pushes three values onto the stack: 3, 4 and 5" spans the entries `> 3`, `> 4`, and `> 5`.
- A bracket labeled "Adds the top two stack items, displays and pushes the result" spans the entries `> add` and `9`.
- An arrow labeled "Error, because there's only one item (12) on the Fewer than two items on the stack" points to the entry `> add`.

Figure 13.3 shows how each of the valid inputs in this session affects the contents of the stack. Every time the user enters an integer, the stack size grows by one, and every time the user enters add, the stack size decreases by one, as long as there are two items to add.




Figure 13.3 How each user input affects the contents of the stack

To implement this interactive stack program, you'll need to extend `StackCmd` to support reading from and writing to the console. The following listing shows `StackCmd` in a new file, `StackIO.idr`, extended with two commands: `GetStr` and `PutStr`.

Listing 13.9 Extending StackCmd to support console I/O with the commands GetStr and PutStr (StackIO.idr)

```
data StackCmd : Type -> Nat -> Nat -> Type where
    Push : Integer -> StackCmd () height (S height)
    Pop : StackCmd Integer (S height) height
    Top : StackCmd Integer (S height) (S height)

    GetStr : StackCmd String height height
    PutStr : String -> StackCmd () height height

    Pure : ty -> StackCmd ty height height
    (>=) : StackCmd a height1 height2 ->
        (a -> StackCmd b height2 height3) ->
        StackCmd b height1 height3
```

Neither `GetStr` nor `PutStr` use the stack, so the height remains the same.

DEPENDENT STATES IN THE EFFECTS LIBRARY I mentioned the Effects library in chapter 12, which allows you to combine effects like state and console I/O without having to define a new type, like `StackCmd` here. The Effects library supports descriptions of state transitions and dependent state as in `StackCmd`. I won't describe the Effects library further in this book, but learning about the principles of dependent state here will mean that you'll be able to learn how to use the more flexible Effects library more readily.

You'll also need to update `runStack` to support the two new commands. Because `GetStr` and `PutStr` describe interactive actions, you'll need to update the type of `runStack` to return IO actions. Here's the updated `runStack`.

Listing 13.10 Updating runStack to support the interactive commands GetStr and PutStr (StackIO.idr)

```
runStack : (stk : Vect inHeight Integer) ->
            StackCmd ty inHeight outHeight ->
            IO (ty, Vect outHeight Integer)
runStack stk (Push val) = pure (((), val :: stk))
```

```

runStack (val :: stk) Pop = pure (val, stk)
runStack (val :: stk) Top = pure (val, val :: stk)
runStack stk GetStr = do x <- getLine
                         pure (x, stk)
runStack stk (PutStr x) = do putStrLn x
                             pure ((), stk)
runStack stk (Pure x) = pure (x, stk)
runStack stk (x >>= f) = do (x', newStk) <- runStack stk x
                           runStack newStk (f x')

```

As with the vending machine, you'll describe infinite sequences of `StackCmd` operations in total functions by defining a separate `StackIO` type for describing infinite streams of stack operations. The following listing shows how you can define `StackIO` and how to run `StackIO` sequences, given an initial state for the stack.

Listing 13.11 Defining infinite sequences of interactive stack operations (StackIO.idr)

```

data StackIO : Nat -> Type where
  Do : StackCmd a height1 height2 ->
    (a -> Inf (StackIO height2)) -> StackIO height1

namespace StackDo
  (>>=) : StackCmd a height1 height2 ->
    (a -> Inf (StackIO height2)) -> StackIO height1
  (>>=) = Do

data Fuel = Dry | More (Lazy Fuel)

partial
forever : Fuel
forever = More forever

run : Fuel -> Vect height Integer -> StackIO height -> IO ()
run (More fuel) stk (Do c f)
  = do (res, newStk) <- runStack stk c
      run fuel newStk (f res)
run Dry stk p = pure ()

```

The Nat argument is the initial height of the stack for the infinite sequence.

Supports do notation for StackIO

forever allows you to run a total program indefinitely by giving an infinite supply of Fuel. See chapter 11 for the full details.

The input Vect must have a number of items given by the initial stack height.

The interactive calculator follows a similar pattern to the implementation of the vending machine. The next listing shows an outline of the main loop, which reads an input, parses it into a command type, and processes the command if the input is valid.

Listing 13.12 Outline of an interactive stack-based calculator (StackIO.idr)

```
data StkInput = Number Integer
              | Add

strToInput : String -> Maybe StkInput

mutual
  tryAdd : StackIO height
  stackCalc : StackIO height
  stackCalc = do PutStr "> "
                input <- GetStr
```

← Describes possible user inputs:
entering a number or the add

← Parses the input read from the console.
Returns Maybe because input could be invalid.

← Adds two numbers at the top of the
stack, if present, and then loops

Main loop of the interactive calculator

```

    case strToInput input of
        Nothing => do PutStr "Invalid input\n"
                      stackCalc
        Just (Number x) => do Push x
                               stackCalc
        Just Add => tryAdd

main : IO ()
main = run forever [] stackCalc

```

You still need to define `strToInput`, which parses user input, and `tryAdd`, which adds the two elements on the top of the stack, if possible. The following listing shows the definition of `strToInput`.

Listing 13.13 Reading user input for the stack-based calculator (StackIO.idr)

```

Empty input is considered invalid. → strToInput : String -> Maybe RPNIInput
                                         strToInput "" = Nothing
                                         strToInput "add" = Just Add
                                         strToInput x = if all isDigit (unpack x)
                                            then Just (Number (cast x))
                                            else Nothing
                                         If the input is the string “add”, parse as the Add command.
                                         If the input consists entirely of digits, parse as Number.

```

Finally, the next listing shows the definition of `tryAdd`. Like `vend` and `refill` in the vending machine implementation, you need to match on the initial state to make sure that there are enough items on the stack to add.

Listing 13.14 Adding the top two elements on the stack, if they’re present (StackIO.idr)

```

Adding is only valid if there are at least two elements on the stack. → doAdd, defined earlier, has a precondition in its type that there are two elements on the stack.

tryAdd : StackIO height
tryAdd {height = (S (S h))} = do doAdd
                                result <- Top
                                PutStr (show result ++ "\n")
                                stackCalc
                                Inspects the top item on the stack so that you can display it as the result
tryAdd = do PutStr "Fewer than two items on the stack\n"
           stackCalc
           If the earlier case doesn’t match, there aren’t enough items on the stack to add.

Continues with the main loop →

```

You can check that `stackCalc` is total at the REPL:

```
*StackIO> :total stackCalc
Main.stackCalc is Total
```

By separating the looping component (`StackIO`) from the terminating component (`StackCmd`), and by giving precise types to the operations, you can be sure that `stackCalc` has at least the following properties, as long as it’s total:

- It will continue running indefinitely.
- It will never crash due to user input that isn't handled.
- It will never crash due to a stack overflow.

Exercises



- 1 Add user commands to the stack-based calculator for subtract and multiply. You can test these as follows:

```
*ex_13_2> :exec  
> 5  
> 3  
> subtract  
2  
> 8  
> multiply  
16
```

- 2 Add a negate user command to the stack-based calculator for negating the top item on the stack. You can test this as follows:

```
> 10  
> negate  
-10
```

- 3 Add a discard user command that removes the top item from the stack. You can test this as follows:

```
> 3  
> 4  
> discard  
Discarded 4  
> add  
Fewer than two items on the stack
```

- 4 Add a duplicate user command that duplicates the top item on the stack. You can test this as follows:

```
> 2  
> duplicate  
Duplicated 2  
> add  
4
```

13.3 Summary

- Data types can model state machines by using each data constructor to describe a state transition.
- You can describe how a command changes the state of a system by giving the input and output states of the system as part of the command's type.
- Developing sequences of state transitions interactively, using holes, means you can check the required input and output states of a sequence of commands.

- Types can model infinite state spaces as well as finite states.
- Sequences of commands give verified sequences of state transitions because a sequence of commands will only type-check if it describes a valid sequence of state transitions.
- You can represent mutable dependently typed state by putting the arguments to the dependent type in the state transitions. For example, you can use the length of a vector to represent the height of a stack.

Dependent state machines: handling feedback and errors

This chapter covers

- Handling errors in state transitions
- Developing protocol implementations interactively
- Guaranteeing security properties in types

As you saw in the previous chapter, you can describe the valid state transitions of a state machine in a dependent type, indexed by the required input state of an operation (its precondition) and the output state (its postcondition). By defining valid state transitions in the type, you can be sure that a program that type-checks is guaranteed to describe a valid sequence of state transitions.

You saw two examples, a description of a door and a simulation of a vending machine, and in each case we gave precise types to the operations to describe how they affected the state. But we didn't consider the possibility that any of the operations could fail:

- What if, when you try to open the door, it's jammed? What if, even though you've run the `Open` operation, it's still in the `DoorClosed` state?

- What if, when you insert a coin in the vending machine, the machine rejects the coin?

In almost any realistic setting, when you try to give precise types to describe a state machine, you'll need to consider the possibility of the operation failing, or of an unexpected response:

- Like the `DoorState`, you might represent the state of a file handle (open or closed) in a type, but if you try to open a file, the file might not exist or you might not have permission to open it.
- You might represent a secure communication protocol in a state machine, but whether you can progress in the protocol depends on receiving valid responses from the network to any request you send.
- You might represent the state of a bank's automated teller machine (ATM) in a type (waiting for a card, waiting for PIN entry, and so on), but you can only move to a state where the machine can dispense cash if checking the user's PIN is successful.

In these cases, you're not in complete control of how the system state changes. You can *request* a change in the state (opening a file, sending a message, and so on) but whether and how the state changes in practice depends on the response you receive from the environment. In this chapter, you'll see how to deal with the possibility of an operation failing by allowing a state transition to *depend* on the result of an operation.

You'll also see how to deal with state changes that might depend on user input—we'll look at the state transitions involved in modeling an ATM, where user input determines whether the machine can dispense cash. We'll begin by revisiting the model of the door from chapter 13 and see how to deal with the possibility that the door might jam, so the `Open` operation fails.

14.1 Dealing with errors in state transitions

In the previous chapter, you defined a `DoorCmd` data type for modeling state transitions on a door, as illustrated by the state transition diagram in figure 14.1.

Listing 14.1 recaps the definition of `DoorCmd`, which models the state transition system in figure 14.1.




Figure 14.1 A state transition diagram showing the states and operations on a door

Listing 14.1 Modeling the door state machine in a type

```
data DoorState = DoorClosed | DoorOpen

data DoorCmd : Type -> DoorState -> DoorState -> Type where
  Open : DoorCmd      () DoorClosed DoorOpen
  Close : DoorCmd     () DoorOpen   DoorClosed
  RingBell : DoorCmd  () DoorClosed DoorClosed

  Pure : ty -> DoorCmd ty state state
  (>>=) : DoorCmd a state1 state2 ->
    (a -> DoorCmd b state2 state3) ->
    DoorCmd b state1 state3
```

In this model, you’re in complete control of how each operation moves from one state to another. For example, the type of `Open` states that it *always* starts with a door in the `DoorClosed` state, and it *always* ends with a door in the `DoorOpen` state:

```
Open : DoorCmd () DoorClosed DoorOpen
```

Reality is not always so accommodating, however! If you were to implement this with some real hardware, such as a sliding door operated by pressing a button, you’d need to consider the possibility of hardware problems such as the door jamming. In this section, we’ll refine the door model to capture this possibility of failure, and see how this affects the implementation of programs that follow the protocol.

14.1.1 Refining the door model: representing failure

`Open` could fail due to the door being jammed, so we need a way to represent whether it was successful. We could define an enumeration type to describe the possible results of opening the door:

```
data DoorResult = OK | Jammed
```

Then, instead of producing the unit value, `Open` could return a `DoorResult`. We might try the following type for `Open`:

```
Open : DoorCmd DoorResult DoorClosed DoorOpen
```

Unfortunately, this isn’t quite right because it still says that opening the door causes the door to be in the `DoorOpen` state, whatever happens. Somehow, we need to express that `Open` causes the door to be in either the `DoorClosed` or `DoorOpen` state, depending on the value of the `DoorResult` it produces. Figure 14.2 illustrates the state machine we’d like to implement.




Figure 14.2 A state transition diagram showing the states and operations on a door, where opening the door might fail

You can achieve this by changing the type of `DoorCmd` to allow the output state to be calculated from the return value. Figure 14.3 illustrates how you can refine the type of `DoorCmd` to achieve this.




Figure 14.3 New type for `DoorCmd`, where the output state of an operation is computed from the return value of the operation

Here, you've given a name to the return type of the operation, `ty`, and said that the output state is computed by a function that takes a `ty` as an input. Now, when you define the type of `Open` (and indeed all of the `DoorCmd` operations), you give an *expression* for the output state, explaining how the output state is computed from the return value, of type `DoorResult`:

```

Open : DoorCmd DoorResult DoorClosed
      => (\res => case res of
           OK => DoorOpen
           Jammed => DoorClosed)
  
```

This encodes exactly what the state transition diagram in figure 14.2 illustrates. That is, the output state of `Open` can be one of the following:

- `DoorOpen`—If `Open` returns `OK`
- `DoorClosed`—If `Open` returns `Jammed`

Although you won't know the value of `res` until you run the operation, you can at least use the type to explain the possible states the door will be in given the result of `Open`. Listing 14.2 shows the complete `DoorCmd` type declaration after this refinement. It also adds `Display`, so you can display logging messages if necessary.

Listing 14.2 The refined DoorCmd type, allowing the output state of each operation to be computed from the return value of the operation (DoorJam.idr)

Calculates the output state from the return value of Open

You use `const` to say that the output state is not dependent on the return value.

```
data DoorCmd : (ty : Type) -> DoorState -> (ty -> DoorState) -> Type where
  Open : DoorCmd DoorResult DoorClosed
    (\res => case res of
      OK => DoorOpen
      Jammed => DoorClosed)

  Close : DoorCmd () DoorOpen (const DoorClosed)
  RingBell : DoorCmd () DoorClosed (const DoorClosed)

  Display : String -> DoorCmd () state (const state)

  Pure : (res : ty) -> DoorCmd ty (state_fn res) state_fn
  (>>=) : DoorCmd a state1 state2_fn ->
    ((res : a) -> DoorCmd b (state2_fn res) state3_fn) ->
    DoorCmd b state1 state3_fn
```

This type for `Pure` means that the value `res` can be used to compute the output state of a sequence of operations.

The `(>>=)` operator needs to compute the intermediate state from the output of the first operation.

Calculating output state with `const`

This is the type of `const`, defined in the Prelude:

```
*DoorJam> :t const
const : a -> b -> a
```

It ignores its second argument and returns its first. So, if you say `const DoorClosed` for the output state of an operation, that gives you a function that ignores the result of the function and always returns `DoorClosed`.

In the previous definition of `DoorCmd`, in listing 14.1, you used the `(>>=)` operator to explain that the output state of the first operation should be the input state of the second:

```
(>>=) : DoorCmd a state1 state2 -> (a -> DoorCmd b state2 state3) ->
  DoorCmd b state1 state3
```

It's now slightly more complicated, because the *return value* of the first operation affects the input state of the second:

```
(>>=) : DoorCmd a state1 state2_fn ->
  ((res : a) -> DoorCmd b (state2_fn res) state3_fn) ->
  DoorCmd b state1 state3_fn
```

This works as follows:

- 1 The first operation returns a value of type `a`, and the output state is computed from a `state2_fn` function once you know the result of the operation.
- 2 When you come to the second operation, you'll know the result of the first operation, named `res`, so it has an input state of `state2_fn res`.
- 3 The combined operation has an input state of `state1` (the input state of the first operation) and an output state computed from the result of the *second* operation, using `state3_fn`.

Defining `DoorCmd` this way gives you more precision in defining the state transitions, and it means that when you define functions using `DoorCmd`, the types of the operations require you to execute any necessary checks before continuing. For example, after you try to Open a door, you can't execute any further door operations until you've checked the result. We'll look at how this works by revisiting our earlier example, `doorProg`.

14.1.2 A verified, error-checking, door-protocol description

In chapter 13, you implemented a function using `DoorCmd` as a sequence of actions to ring the bell, and open and then close the door, and you used the types to verify that the sequence of actions was valid. You wrote `doorProg` as follows:

```
doorProg : DoorCmd () DoorClosed DoorClosed
doorProg = do RingBell
             Open
             Close
```

Now that you've refined the type of `DoorCmd` so that the output state is computed from the result of the operation, you'll need to write the type differently:

```
doorProg : DoorCmd () DoorClosed (const DoorClosed)
```

That is, the output state isn't affected by the result, so you use `const`, which ignores its second argument and returns its first unchanged. But if you try implementing `doorProg` as before, without checking the result of `Open`, you'll get an error:

```
doorProg : DoorCmd () DoorClosed (const DoorClosed)
doorProg = do RingBell
             Open
             Close
```

The error happens when you try to use `Close`. Its type requires that the input state is `DoorOpen`, but in fact its input state is computed from the result of `Open`:

```
When checking an application of constructor Main.>>=:
Type mismatch between
    DoorCmd () DoorOpen (const DoorClosed) (Type of Close)
and
    DoorCmd ()
        (\res =>
```

```

    case res of
        OK => DoorOpen
        Jammed => DoorClosed) _)
        (\value => DoorClosed) (Expected type)

```

To see how to avoid this problem, you can develop doorProg interactively, beginning from the following point:

```

doorProg : DoorCmd () DoorClosed (const DoorClosed)
doorProg = do RingBell
            Open
            ?doorProg_rhs

```

DEBUGGING TYPE ERRORS USING HOLES If you get a type error that's hard to understand on its own, it's often a good idea to replace the offending part of the program with a hole (as we did by replacing Close with ?doorProg_rhs), to see what the expected type is, along with any local variables in scope.

- 1 *Type, refine*—If you check the type of ?doorProg_rhs, you'll see this:

```

-----
doorProg_rhs : DoorCmd ()
              (case _ of
                  OK => DoorOpen
                  Jammed => DoorClosed)
                  (\value => DoorClosed)
-----
```

The output state you see here arises from the definition of const in the Prelude, and it's a function that ignores its argument value and returns DoorClosed. The input state you're looking for is calculated from some value, `_`. This value is the result of Open, which you haven't named. Let's call it `jam`:

```

doorProg : DoorCmd () DoorClosed (const DoorClosed)
doorProg = do RingBell
            jam <- Open
            ?doorProg_rhs

```

- 2 *Type*—You'll now see that the input state of the next operation is calculated from the value of `jam`:

```

jam : DoorResult
-----
doorProg_rhs : DoorCmd ()
              (case jam of
                  OK => DoorOpen
                  Jammed => DoorClosed)
                  (\value => DoorClosed)
-----
```

- 3 *Define, type*—Because the input state depends on the value of `jam`, you can make progress by inspecting the value of `jam`:

```

doorProg : DoorCmd () DoorClosed (const DoorClosed)
doorProg = do RingBell
            jam <- Open

```

```
case jam of
    case_val => ?doorProg_rhs
```

You'll now see that the input state depends on the value of `case_val`:

```
case_val : DoorResult
jam : DoorResult
-----
doorProg_rhs : DoorCmd ()
    (case case_val of
        OK => DoorOpen
        Jammed => DoorClosed)
    (\value => DoorClosed)
```

- 4 Define, type**—If you case-split on `case_val`, you should see that each branch of the case has a different type, calculated from the specific value of `case_val`:

```
doorProg : DoorCmd () DoorClosed (const DoorClosed)
doorProg = do RingBell
    jam <- Open
    case jam of
        OK => ?doorProg_rhs_1
        Jammed => ?doorProg_rhs_2
```

In `?doorProg_rhs_1`, for example, `jam` has the value `OK`, so the door must have successfully opened:

```
jam : DoorResult
-----
doorProg_rhs_1 : DoorCmd () DoorOpen (\value => DoorClosed)
```

In `?doorProg_rhs_2`, on the other hand, the door is jammed, so it's still in the `DoorClosed` state:

```
jam : DoorResult
-----
doorProg_rhs_2 : DoorCmd () DoorClosed (\value => DoorClosed)
```

- 5 Refine**—To complete the definition, you can display a log message in each case, and if the door is open, close it again:

```
doorProg : DoorCmd () DoorClosed (const DoorClosed)
doorProg = do RingBell
    jam <- Open
    case jam of
        OK => do Display "Glad To Be Of Service"
                  Close
        Jammed => Display "Door Jammed"
```

The type of `Open` means that you need to check the state of the door before you execute any further operations that need to know the door's state. In particular, you can't `Close` the door unless you've successfully opened it. You don't have to check immediately, though. For example, you can display a message between opening the door and checking the result:

```

doorProg : DoorCmd () DoorClosed (const DoorClosed)
doorProg = do RingBell
    jam <- Open
    Display "Trying to open the door"
    case jam of
        OK => do Display "Glad To Be Of Service"
        Close
    Jammed => Display "Door Jammed"

```

This is valid, because the precondition on `Display` doesn't require the door to be in a specific state; any will do, and `Display` won't change the state.

Using pattern-matching bindings, which you first saw in chapter 5, you can also define `doorProg` more concisely, as follows:

```

doorProg : DoorCmd () DoorClosed (const DoorClosed)
doorProg = do RingBell
    OK <- Open | Jammed => Display "Door Jammed"
    Display "Glad To Be Of Service"
    Close

```

This gives a default path through the sequence of actions, when `Open` returns `OK`, and an alternative action when `Open` returns `Jammed`. Using pattern-matching bindings makes it easier to write longer sequences of actions, where some of the actions might fail. For example, you can open and close the door twice and abandon the sequence if either fails:

```

doorProg : DoorCmd () DoorClosed (const DoorClosed)
doorProg = do RingBell
    OK <- Open | Jammed => Display "Door Jammed"
    Display "Glad To Be Of Service"
    Close
    OK <- Open | Jammed => Display "Door Jammed"
    Display "Glad To Be Of Service"
    Close

```

This example *describes* a protocol in the type, and it explicitly says where an operation might fail. In `doorProg`, the type of `Open` means that you need to check its result before you can proceed with any further operations that change the state.

The type of Pure

The type of `Pure` in `DoorCmd` allows you to define functions like the following, where the call to `Pure` changes the state:

```

logOpen : DoorCmd DoorResult DoorClosed
    (\res => case res of
        OK => DoorOpen
        Jammed => DoorClosed)
logOpen = do Display "Trying to open the door"
    OK <- Open | Jammed => do Display "Jammed"
        Pure Jammed
    Display "Success"
    Pure OK

```

(continued)

If you replace the last line, `Pure OK`, with a hole, `?pure_ok`, you'll see that it has an input state of `DoorOpen`, and the output state (of the entire `logOpen` function) needs to be a function that computes its output state:

```
pure_ok : DoorCmd DoorResult DoorOpen
          (\res => case res of
              OK => DoorOpen
              Jammed => DoorClosed)
```

The type of `Pure` is designed to work in this situation:

```
Pure : (res : ty) -> DoorCmd ty (state_fn res) state_fn
```

Here, `state_fn` is the function containing the case block, and `Pure` must take `OK` as an argument to have the correct input state for `?pure_ok`.

You now have a definition of `DoorCmd` that precisely describes the protocol for opening and closing doors, capturing the possibility of failure. But you haven't yet seen how the corresponding `run` function works, which is where the result of an `Open` operation would be produced in practice. We'll look at this next, in the context of a larger example: an ATM. We'll also look at how state can change according to user input.

14.2 Security properties in types: modeling an ATM

You can use explicit states in the types of operations to guarantee, by type checking, that a system will only execute security-critical operations when it's in a valid state to do so. For example, an ATM should only dispense cash when a user has inserted their card and entered a correct PIN. This is a typical sequence of operations on an ATM:

- 1 A user inserts their bank card.
- 2 The machine prompts the user for their PIN, to check that the user is entitled to use the card.
- 3 If PIN entry is successful, the machine prompts the user for an amount of money, and then dispenses cash.

If the user enters the correct PIN, the machine will be in a state to dispense cash; otherwise, it won't be. In this section, we'll define a model for an ATM and see how to change the state of the machine, in its type, based on user input.

SECURITY PROPERTY OF THE ATM In this model, we'll omit some of the finer details of banking, such as accessing and updating the user's bank account, and checking the PIN securely, which would be done with separate state machines. We'll focus on one important security property we want to maintain: the machine must only dispense cash when there is a validated card in the machine.

As with the door model, we'll begin by defining the possible states of the ATM and the operations that can change the ATM's state. Once we know how the operations

affect the state, we can define a data type for representing the operations on the machine.

14.2.1 Defining states for the ATM

An ATM is either waiting for a user to begin an interaction, waiting for the user to enter their PIN, or ready to dispense cash after validating a PIN. So an ATM, in our model, can be in one of the following states:

- Ready—The ATM is ready and waiting for a card to be inserted.
- CardInserted—There is a card inside the ATM, but the system has not yet checked a PIN entry against the card.
- Session—There is a card inside the ATM and the user has entered a valid PIN for the card, so a validated session is in progress.

We'll validate the card by checking that the user inputs a correct PIN. In the vending machine in chapter 13, you also had to check that an input was valid, but in that case you could check the command *locally*. Here, we'll assume that there's an external service to check the PIN, so we won't know until runtime which inputs will result in which states.

The machine supports the following basic operations, each of which may have preconditions and postconditions on the state of the machine:

- InsertCard—Waits for a user to insert a card
- EjectCard—Ejects a card from the machine, as long as there's a card in the machine
- GetPIN—Reads a user's PIN, as long as there's a card in the machine
- CheckPIN—Checks whether an entered PIN is valid
- Dispense—Dispenses cash as long as there's a validated card in the machine
- GetAmount—Reads from the user an amount to dispense
- Message—Displays a message to the user

Figure 14.4 illustrates how these operations affect the state of the machine.




Figure 14.4 A state machine describing the states and operations on an ATM. This omits operations (such as `GetAmount`) that are valid in all states.

Having defined the states and seen how the high-level operations that the machine performs can affect the states, we're now in a position to define a type for the ATM that describes the state transitions illustrated in figure 14.4.

14.2.2 Defining a type for the ATM

Listing 14.3 defines an `ATMCmd` type that represents the state transitions of operations on an ATM in Idris code. It also includes `GetAmount` and `Message`, which are valid in all states and don't affect the state, and the usual operations `Pure` and `(>>=)`. The type of `EjectCard` is slightly simplified; we'll refine this in section 14.2.4.

Listing 14.3 A type for representing the commands of an ATM, and how they affect the ATM's state (`ATM.idr`)

```
import Data.Vect

PIN : Type
PIN = Vect 4 Char

A PIN is
exactly four
characters. ↗

data ATMState = Ready | CardInserted | Session ←
The possible
states of the ATM

data PINCheck = CorrectPIN | IncorrectPIN ←
The possible results
of checking a PIN:
valid or invalid

data ATMCmd : (ty : Type) -> ATMState -> (ty -> ATMState) -> Type where
  InsertCard : ATMCmd () Ready (const CardInserted)
  EjectCard : ATMCmd () state (const Ready)
  GetPIN : ATMCmd PIN CardInserted (const CardInserted)

  CheckPIN : PIN -> ATMCmd PINCheck CardInserted
    (\check => case check of
      CorrectPIN => Session
      IncorrectPIN => CardInserted)
  GetAmount : ATMCmd Nat state (const state)

  Dispense : (amount : Nat) -> ATMCmd () Session (const Session) ←
  Message : String -> ATMCmd () state (const state)
  Pure : (res : ty) -> ATMCmd ty (state_fn res) state_fn
  (>>=) : ATMCmd a state1 state2_fn ->
    ((res : a) -> ATMCmd b (state2_fn res) state3_fn) ->
    ATMCmd b state1 state3_fn

This isn't quite right, because it allows the
card to be ejected even if there's no card
in the machine. We'll refine it shortly. ↗
The machine will only
dispense money if there's a
validated card in the machine. ↗

Only moves to the Session state
if the PIN check succeeds ↗
```

Using `ATMCmd`, you can write a function that describes a session on an ATM, from the user inserting a card to the machine dispensing cash. Listing 14.4 shows the beginning of an `atm` function that waits for a user to insert a card, prompts for a PIN, and then checks the result. I've left a hole for the rest of the sequence, in which we'll check the PIN and dispense cash if the PIN is valid.

Listing 14.4 An atm function describing a sequence of operations on an ATM (ATM.idr)

```
atm : ATMCmd () Ready (const Ready)
atm = do InsertCard
    pin <- GetPIN
    pinOK <- CheckPIN pin
    ?atm_rhs
```

Checks whether the PIN is valid.
The next state of the machine
depends on the value of pinOK.

You can complete atm as follows:

- Type, define**—If you check the type of the ?atm_rhs hole, you'll see that you begin in a state that depends on the value of pinOK, and you need to end in the Ready state:

```
pin : Vect 4 Char
pinOK : PINCheck
-----
atm_rhs : ATMCmd ()
(case pinOK of
    CorrectPIN => Session
    IncorrectPIN => CardInserted)
(\value => Ready)
```

The type here suggests you can proceed by checking the value of pinOK:

```
atm : ATMCmd () Ready (const Ready)
atm = do InsertCard
    pin <- GetPIN
    pinOK <- CheckPIN pin
    case pinOK of
        CorrectPIN => ?atm_rhs_1
        IncorrectPIN => ?atm_rhs_2
```

- Type**—If you check ?atm_rhs_1 and ?atm_rhs_2, you'll see that the state is different in each case. In ?atm_rhs_1, the PIN was found to be valid, so you have a validated session:

```
pinOK : PINCheck
pin : Vect 4 Char
-----
atm_rhs_1 : ATMCmd () Session (\value => Ready)
```

In ?atm_rhs_2, on the other hand, the PIN was found to be invalid, so you're still in the CardInserted state:

```
pinOK : PINCheck
pin : Vect 4 Char
-----
atm_rhs_2 : ATMCmd () CardInserted (\value => Ready)
```

- Refine**—In ?atm_rhs_1, you can now dispense cash, so you can prompt for an amount and dispense that:

```
case pinOK of
    CorrectPIN => do cash <- GetAmount
                    Dispense cash
```

```
?atm_rhs_1
IncorrectPIN => ?atm_rhs_2
```

- 4 Type, refine**—The input state of `?atm_rhs_1` is still `Session`, so before you finish you need to get back to the `Ready` state:

```
pin : Vect 4 Char
pinOK : PINCheck
cash : Nat
-----
atm_rhs_1 : ATMCmd () Session (\value => Ready)
```

You can achieve this by ejecting the card:

```
case pinOK of
    CorrectPIN => do cash <- GetAmount
                      Dispense cash
                      EjectCard
    IncorrectPIN => ?atm_rhs_2
```

- 5 Refine**—For `?atm_rhs_2`, the PIN was invalid, so the simplest thing to do is eject the card, leading to the following completed definition:

```
atm : ATMCmd () Ready (const Ready)
atm = do InsertCard
         pin <- GetPIN
         pinOK <- CheckPIN pin
         case pinOK of
             CorrectPIN => do cash <- GetAmount
                               Dispense cash
                               EjectCard
             IncorrectPIN => EjectCard
```

There are other ways you could define `atm`. For example, it would be helpful to display messages to the user. Also, in practice, ATMs typically don't check the PIN until just before dispensing cash. The next listing shows this alternative way of implementing `atm`.

Listing 14.5 An alternative implementation of `atm`, including messages to the user and checking the PIN later (ATM.idr)

```
atm : ATMCmd () Ready (const Ready)
atm = do InsertCard
         pin <- GetPIN
         cash <- GetAmount
         pinOK <- CheckPIN pin
         Message "Checking Card"
         case pinOK of
             CorrectPIN => do Dispense cash
                               EjectCard
                               Message "Please remove your card and cash"
             IncorrectPIN => do Message "Incorrect PIN"
                               EjectCard
```

You haven't checked the PIN yet, or the state of `pinOK`, but these commands are valid because they don't require the machine to be in a specific state.

As long as you only execute actions when the machine is in the appropriate state, and as long as you make sure that every path through the actions in `atm` ends in the Ready state, you can implement the details however you like. If `atm` type-checks, you can be certain that you've maintained the important security property: the machine will only dispense cash when there's a card in the machine and the PIN has been entered correctly.

14.2.3 Simulating an ATM at the console: executing ATMCmd

To try the `atm` function, you can write a console simulation of an ATM that produces IO actions for an ATMCmd description:

```
runATM : ATMCmd res inState outState_fn -> IO res
```

Given a sequence of ATMCmd that produces a result of type `res`, begins in state `inState`, and computes the result state with `outState_fn`, `runATM` gives a sequence of IO actions that produces a result, `res`. Let's hardcode a single valid PIN for this simulation:

```
testPIN : Vect 4 Char
testPIN = ['1', '2', '3', '4']
```

The following listing shows a console simulation of an ATM that uses this hardcoded PIN. In this simulation, many of the commands prompt for an input on the console and return the value read.

Listing 14.6 A console simulation of an ATM (ATM.idr)

```
readVect : (n : Nat) -> IO (Vect n Char)
readVect Z = do discard <- getLine
                pure []
readVect (S k) = do ch <- getChar
                    chs <- readVect k
                    pure (ch :: chs)

runATM : ATMCmd res inState outState_fn -> IO res
runATM InsertCard = do putStrLn "Please insert your card (press enter)"
                        x <- getLine
                        pure ()
runATM EjectCard = putStrLn "Card ejected"
runATM GetPIN = do putStrLn "Enter PIN: "
                   readVect 4
runATM (CheckPIN pin) = if pin == testPIN
                           then pure CorrectPIN
                           else pure IncorrectPIN
runATM GetAmount = do putStrLn "How much would you like? "
                      x <- getLine
                      pure (cast x)
runATM (Dispense amount) = putStrLn ("Here is " ++ show amount)

runATM (Message msg) = putStrLn msg
runATM (Pure res) = pure res
runATM (x >= f) = do x' <- runATM x
                      runATM (f x')
```

readVect reads a line, keeping the prefix of a specific length and discarding the rest.

Simulates waiting for a card to be inserted by waiting for any input

Checks the given pin against the hardcoded PIN. The value returned here determines the new state of the machine.

You've now defined the ATM as a type, with commands describing each of the state transitions on the ATM, and a separate `runATM` function that interprets those commands in the `IO` context. By separating the description from the implementation, you can write different interpreters for different contexts, as required. In particular, you wouldn't want to hardcode a PIN on a real device!

14.2.4 Refining preconditions using auto-implicits

One feature of the state machine in figure 14.4 that `ATMCmd` type doesn't quite capture is that ejecting the card should only be allowed when there's a card in the machine. Instead, you have the following type:

```
EjectCard : ATMCmd () state (const Ready)
```

That is, you can try to eject a card in *any* input state, even when there's no card in the machine. But there are only two states when it's okay to eject a card: `CardInserted` and `Session`. You shouldn't be able to write the following function, because the machine is ejecting a card in the `Ready` state:

```
badATM : ATMCmd () Ready (const Ready)
badATM = EjectCard
```

Somehow, you need both of the following types to work for `EjectCard`:

```
EjectCard : ATMCmd () CardInserted (const Ready)
EjectCard : ATMCmd () Session (const Ready)
```

A data constructor like `EjectCard` can't have two different types. You can, however, define a predicate on `ATMState` that will allow you to restrict the possible input states of `EjectCard` to those that are valid. We discussed predicates in chapter 9, and you can define a `HasCard` predicate that describes the states in which a machine contains a card:

```
data HasCard : ATMState -> Type where
    HasCI : HasCard CardInserted
    HasSession : HasCard Session
```

You can only construct a value of type `HasCard` state when state is one of `CardInserted` or `Session`, so you can refine the type of `EjectCard` as follows:

```
EjectCard : HasCard state -> ATMCmd () state (const Ready)
```

If you do this, you'll need to give values of type `HasCard` explicitly when using `EjectCard`. For example:

```
insertEject : ATMCmd () Ready (const Ready)
insertEject = do InsertCard
                EjectCard HasCI
```

Having to write explicit values for the predicate will get tedious very quickly. Instead, you can use an `auto` implicit for the predicate, which you also saw in chapter 9:

```
EjectCard : {auto prf : HasCard state} -> ATMCmd () state (const Ready)
```

Now, you can use `EjectCard` as before, and let Idris find the correct value for the predicate by searching through the possible data constructors for `HasCard` to see if any of them are valid:

```
insertEject : ATMCmd () Ready (const Ready)
insertEject = do InsertCard
                EjectCard
```

For `badATM`, Idris shouldn't be able to find a suitable value:

```
badATM : ATMCmd () Ready (const Ready)
badATM = EjectCard
```

In this case, Idris will report an error, saying that it needs to find a value of type `HasCard Ready` for the predicate to `EjectCard`, but it can't find one:

```
When checking argument prf to constructor Main.EjectCard:
  Can't find a value of type
    HasCard Ready
```

All of your other previous definitions, including the two versions of `atm` and the execution function `runATM`, will work without any alteration using this refined version of `EjectCard`.

Exercises



- 1 The following type outlines a security system in which a user can log in with a password and then read a secret message, although there are some gaps:

```
data Access = LoggedOut | LoggedIn
data PwdCheck = Correct | Incorrect

data ShellCmd : (ty : Type) -> Access -> (ty -> Access) -> Type where
  Password : String -> ShellCmd PwdCheck ?password_in ?password_out
  Logout : ShellCmd () ?logout_in ?logout_out
  GetSecret : ShellCmd String ?getsecret_in ?getsecret_out

  PutStr : String -> ShellCmd () state (const state)
  Pure : (res : ty) -> ShellCmd ty (state_fn res) state_fn
  (=>) : ShellCmd a state1 state2_fn ->
    ((res : a) -> ShellCmd b (state2_fn res) state3_fn) ->
      ShellCmd b state1 state3_fn
```

Fill in the holes in the following types:

- `Password`—Reads a password and changes the state to `LoggedIn` or `LoggedOut`, depending on whether the password was correct
- `Logout`—Changes the state from `LoggedIn` to `LoggedOut`
- `GetSecret`—Reads a secret message as long as the state is `LoggedIn`

The following function should type-check if you have the correct answer:

```
session : ShellCmd () LoggedOut (const LoggedOut)
session = do Correct <- Password "wurzel"
```

```

    | Incorrect => PutStr "Wrong password"
msg <- GetSecret
PutStr ("Secret code: " ++ show msg ++ "\n")
Logout

```

The following functions should *not* type-check:

```

sessionBad : ShellCmd () LoggedOut (const LoggedOut)
sessionBad = do Password "wurzel"
    msg <- GetSecret
    PutStr ("Secret code: " ++ show msg ++ "\n")
    Logout

noLogout : ShellCmd () LoggedOut (const LoggedOut)
noLogout = do Correct <- Password "wurzel"
    | Incorrect => PutStr "Wrong password"
    msg <- GetSecret
    PutStr ("Secret code: " ++ show msg ++ "\n")

```

- 2 When inserting a coin into the vending machine defined in chapter 13, the machine could reject the coin. You can represent this by changing the types of MachineCmd and InsertCoin. An operation in MachineCmd can change the state based on its result:

```
data MachineCmd : (ty : Type) -> VendState -> (ty -> VendState) -> Type
```

Then, InsertCoin can return whether or not the coin insertion was successful and change the state accordingly:

```

InsertCoin : MachineCmd CoinResult (pounds, chocs)
    (\res => case res of
        Inserted => (S pounds, chocs)
        Rejected => (pounds, chocs))

```

Define the CoinResult type, and then make this change to MachineCmd in Vending.idr. Also, refine the types of the other commands and the implementation of machineLoop as necessary.

14.3 A verified guessing game: describing rules in types

As a concluding example in this chapter, we'll look at how you can use a type to represent the rules of a game precisely, and be sure that any implementation of the game follows the rules correctly. We'll revisit an example from chapter 9, the word-guessing game Hangman.

To recap how this worked, you defined a WordState type to represent the state of the game. WordState was defined as follows, including the number of guesses and letters remaining as arguments:

```

data WordState : (guesses : Nat) -> (letters : Nat) -> Type where
  MkWordState : (word : String)
    -> (missing : Vect letters Char)
    -> WordState guesses_remaining letters

```

Target word, to be
guessed by the player

Letters still to
be guessed by
the player

You also defined a `Finished` type to express when a game was complete, either because there were no letters left to guess in the word (so the player won), or there were no guesses remaining (so the player lost):

```
data Finished : Type where
  Lost : (game : WordState 0 (S letters)) -> Finished
  Won : (game : WordState (S guesses) 0) -> Finished
```

Given these, you defined a main loop called `game`, which took a `WordState` with both guesses and letters remaining, and looped until the game was complete:

```
game : WordState (S guesses) (S letters) -> IO Finished
```

In the implementation, you used the type to help direct you to a working implementation. But you could also have written an incorrect implementation of the game using this type. For example, the following implementation of `game` would also be well typed, but wrong, because it returns a losing game state in all cases:

```
game : WordState (S guesses) (S letters) -> IO Finished
game state = pure (Lost (MkWordState "ANYTHING" ['A']))
```

Although the type allows you to express the game state precisely and helps you give types to intermediate operations (such as processing a guess), it doesn't guarantee that the implementation follows the rules of the game correctly. In the preceding implementation, it's impossible for the player to win!

In this section, instead of defining a `WordState` type and then trusting that `game` will follow the rules of the game correctly, we'll define the rules of the game precisely in a state type. Just as `DoorCmd` expresses when we can execute operations on a door, and `ATMCmd` expresses when we can execute operations on an ATM, we can define a dependent `GameCmd` type that expresses when it's valid to execute particular operations in a game, and what the effect on those operations will be. As with the door and ATM examples, we'll begin by defining the states and the operations that can be executed on those states.

14.3.1 Defining an abstract game state and operations

First, we'll think about how we can define the rules of the game in abstract terms, without worrying about the details of the implementation. A game can be in one of the following states:

- `NotRunning`—There is no game currently in progress. Either the game hasn't started and there's no word yet to guess, or the game is complete.
- `Running`—There is a game in progress, and the player has a number of guesses remaining and letters still to guess.

In the case of `Running`, we'll annotate the state with the number of guesses and letters remaining, just as we did with `WordState` earlier, because this means we'll be able to

describe precisely when a game has been won (no letters to guess) or lost (no guesses remaining). We can express the possible states in the following data type:

```
data GameState : Type where
    NotRunning : GameState
    Running : (guesses : Nat) -> (letters : Nat) -> GameState
```

Then, we'll support some basic operations for manipulating the game state:

- NewGame—Initializes a game with a word for the player to guess
- Guess—Allows the player to guess a letter
- Won—Declares that the player has won the game
- Lost—Declares that the player has lost the game

Figure 14.5 illustrates how these basic operations affect the game state. There are additional preconditions on Won and Lost: we can only declare that the player has won the game if there are no letters left to guess, and we can only declare that the player has lost if there are no guesses remaining.




Figure 14.5 State transition diagram for Hangman. The Running state also holds the number of letters and guesses remaining. Won requires the number of letters to be zero, and Lost requires the number of guesses to be zero.

The next step is to represent these state transitions precisely in a dependent type, including the specific rules about the number of guesses and letters required for each operation to be valid.

14.3.2 Defining a type for the game state

We'll define a GameCmd type that describes the possible operations you can run that affect GameState.

The next listing shows the types of NewGame, Won, and Lost from figure 14.5. As usual, Pure and ($\gg=$) are included so that you can introduce pure values and sequence operations.

Listing 14.7 Beginning to define GameCmd (Hangman.idr)

```
import Data.Vect
%default total

data GameCmd : (ty : Type) -> GameState -> (ty -> GameState) -> Type where
```

Reports an
error if any
definitions
are not total

You'll use Vect later to
represent missing letters.

You can only start a new game if you're not currently playing a game.

```

NewGame : (word : String) ->
    GameCmd () NotRunning
        (const (Running 6 (length (letters word))))
    ↗ Won : GameCmd () (Running (S guesses) 0)
        (const NotRunning)
    ↗ Lost : GameCmd () (Running 0 (S guesses))
        (const NotRunning)

Pure : (res : ty) -> GameCmd ty (state_fn res) state_fn
(>>=) : GameCmd a state1 state2_fn ->
    ((res : a) -> GameCmd b (state2_fn res) state3_fn) ->
    GameCmd b state1 state3_fn

```

Allows six guesses and counts the number of letters to guess

You can assert that the player has won, as long as there are no letters to guess, and move into the `NotRunning` state.

You can assert that the player has lost, as long as there are no guesses remaining.

You can get the different letters in a word using `letters`, which converts the word to uppercase, then converts it to a `List Char`, and finally removes any duplicate elements:

```

letters : String -> List Char
letters str = nub (map toUpper (unpack str))

```

nub is defined in the Prelude;
it removes duplicate elements from a list.

Listing 14.8 adds a `Guess` operation to `GameCmd`. The type of `Guess` has a precondition and a postcondition that explain how the guess affects the game:

- As a precondition, there must be at least one guess available (`S guesses`) and at least one letter still to guess (`S letters`), or attempting to guess a letter won't type-check.
- As a postcondition, the number of guesses will reduce if the guess was incorrect, and the number of letters will reduce if the guess was correct.

Listing 14.8 Adding a Guess operation (Hangman.idr)

```

data GuessResult = Correct | Incorrect ← The result of a Guess is either Correct or Incorrect.

data GameCmd : (ty : Type) -> GameState -> (ty -> GameState) -> Type where
    NewGame : (word : String) ->
        GameCmd () NotRunning
            (const (Running 6 (length (letters word))))
    Won : GameCmd () (Running (S guesses) 0)
        (const NotRunning)
    Lost : GameCmd () (Running 0 (S guesses))
        (const NotRunning)

    Guess : (c : Char) ->
        GameCmd GuessResult
            (Running (S guesses) (S letters)) ← You can only guess if there are guesses and letters remaining.

```

```

    (\res => case res of
                Correct => Running (S guesses) letters
                Incorrect => Running guesses (S letters)) ←

Pure : (res : ty) -> GameCmd ty (state_fn res) state_fn
(>>=) : GameCmd a state1 state2_fn ->
          ((res : a) -> GameCmd b (state2_fn res) state3_fn) ->
          GameCmd b state1 state3_fn

```

Correct, so there's one fewer letter to guess

Incorrect, so there's one fewer guess remaining

Finally, in order to be able to implement the game with a user interface, you'll need to add commands for displaying the current game state, displaying any messages, and reading a guess from the user:

```

data GameCmd : (ty : Type) -> GameState -> (ty -> GameState) -> Type where
  {- Continued from Listing 14.8 -}
  ShowState : GameCmd () state (const state)
  Message : String -> GameCmd () state (const state)
  ReadGuess : GameCmd Char state (const state)

```

Displaying the game state should display the known letters in the target word and the number of guesses remaining. For example, if the target word is TESTING and you've already guessed T, with six guesses remaining, ShowState should display the following:

```

T--T---
6 guesses left

```

When you actually implement the game, it's useful to support indefinitely long game loops. For example, once you finish a game, a player might want to start a new game. The next listing defines a GameLoop type, using Inf to note that execution might continue indefinitely.

Listing 14.9 A type for describing potentially infinite game loops (Hangman.idr)

```

→ namespace Loop
  data GameLoop : (ty : Type) -> GameState -> (ty -> GameState) -> Type where
    (>>=) : GameCmd a state1 state2_fn ->
              ((res : a) -> Inf (GameLoop b (state2_fn res) state3_fn)) ->
              GameLoop b state1 state3_fn
    Exit : GameLoop () NotRunning (const NotRunning) ←

```

Introduces a new namespace, because you're overloading (>>=)

You can't Exit a game that's still running.

You can use the operations in GameLoop and GameCmd to define the following function, which implements a game loop:

```
gameLoop : GameLoop () (Running (S guesses) (S letters)) (const NotRunning)
```

Once you have a well-typed, total implementation of gameLoop, you'll know that it's a valid implementation of the rules. Neither the game nor the player will be able to

cheat by breaking the rules as they're defined in `GameCmd`. You can only call `gameLoop` on a properly initialized game, with a word to guess, and any implementation *must* be a complete implementation of the game because the only way to finish a `GameLoop` is by calling `Exit`, which requires a game to be in the `NotRunning` state.

14.3.3 Implementing the game

We'll implement `gameLoop` interactively and see how the state of the game progresses by checking types as we go.

To begin, you can create a skeleton definition, bringing `guesses` and `letters` from the type into scope, because you'll need to inspect them to check the player's progress later:

```
gameLoop : GameLoop () (Running (S guesses) (S letters)) (const NotRunning)
gameLoop {guesses} {letters} = ?gameLoop_rhs
```

To implement `gameLoop`, take the following steps:

- 1 *Refine*—At the start of each iteration of `gameLoop`, you'll display the current state of the game using `ShowState`, read a guess from the user, and then check whether it was correct:

```
gameLoop : GameLoop () (Running (S guesses) (S letters)) (const NotRunning)
gameLoop {guesses} {letters} = do
    ShowState
    g <- ReadGuess
    ok <- Guess g
    ?gameLoop_rhs
```

- 2 *Type, define*—If you check the type of `?gameLoop_rhs` now, you'll see that the current state of the game depends on whether the guess was correct or not:

```
letters : Nat
guesses : Nat
g : Char
ok : GuessResult
-----
gameLoop_rhs : GameLoop ()
(case ok of
    Correct => Running (S guesses) letters
    Incorrect => Running guesses (S letters))
(\value => NotRunning)
```

To make progress, you'll need to inspect `ok` to establish which state the game is in:

```
gameLoop : GameLoop () (Running (S guesses) (S letters)) (const NotRunning)
gameLoop {guesses} {letters} = do
    ShowState
    g <- ReadGuess
    ok <- Guess g
    case ok of
        Correct => ?gameLoop_rhs_1
        Incorrect => ?gameLoop_rhs_2
```

- 3 Type, define**—In `?gameLoop_rhs_1`, the guess was correct, so the number of letters remaining is reduced, as you can see by checking its type:

```

ok : GuessResult
letters : Nat
guesses : Nat
g : Char
-----
gameLoop_rhs_1 : GameLoop ()
    (Running (S guesses) letters)
    (\value => NotRunning)

```

You can only continue with `gameLoop` if there are both letters and guesses remaining, because its input state is `Running (S guesses) (S letters)`. To decide how to continue, you'll need to check the current value of `letters`:

```

case ok of
    Correct => case letters of
        Z => ?gameLoop_rhs_3
        S k => ?gameLoop_rhs_4
    Incorrect => ?gameLoop_rhs_2

```

- 4 Refine**—In `?gameLoop_rhs_3`, there are no letters left to guess, so the player has won. You can declare that the player has won with `Won`, moving to the `NotRunning` state. Then display the final state and exit:

```

case ok of
    Correct => case letters of
        Z => do Won
            ShowState
            Exit
        S k => ?gameLoop_rhs_4
    Incorrect => ?gameLoop_rhs_2

```

You need to `Exit` explicitly, because `Exit` is the only way to break out of a `GameLoop`. You can only `Exit` a game in the `NotRunning` state.

- 5 Refine**—In `?gameLoop_rhs_4`, there are still letters to guess, so you can display a message and continue with `gameLoop`:

```

case ok of
    Correct => case letters of
        Z => do Won
            ShowState
            Exit
        S k => do Message "Correct"
            gameLoop
    Incorrect => ?gameLoop_rhs_2

```

The `Incorrect` case works similarly, checking whether guesses are still available and declaring that the player has lost if not. The following listing gives the complete definition, for reference.

Listing 14.10 A complete implementation of gameLoop (Hangman.idr)

```
gameLoop : GameLoop () (Running (S guesses) (S letters)) (const NotRunning)
gameLoop {guesses} {letters} = do
    ShowState
    g <- ReadGuess
    ok <- Guess g
    case ok of
        Correct => case letters of
            Z => do Won
                  ShowState
                  Exit
            S k => do Message "Correct"
                  gameLoop
        Incorrect => case guesses of
            Z => do Lost
                  ShowState
                  Exit
            (S k) => do Message "Incorrect"
                  gameLoop
```

You'll also need to initialize the game. For example, you could write a function to set up a new game, and then initiate the gameLoop:

```
hangman : GameLoop () NotRunning (const NotRunning)
hangman = do NewGame "testing"
            gameLoop
```

So far, you've only defined a data type that *describes* the actions in the game. The gameLoop function describes sequences of actions in a valid game of Hangman that follows the rules. In order to *run* the game, you'll need to define a concrete representation of the game state and a function that translates a GameLoop to a sequence of IO actions.

14.3.4 Defining a concrete game state

In the stack example in chapter 13, we had an abstract state of the stack (the number of items on the stack), and a concrete state represented by a Vect of an appropriate length. Similarly, GameState is the abstract state of a game, describing only whether a game is running, and if so, how many guesses and letters are remaining.

In order to run a game, you'll need to define a corresponding concrete game state that includes the specific target word and which specific letters are still to be guessed. The following listing defines a Game type with a GameState argument, representing the concrete data associated with an abstract game state.

Listing 14.11 Representing the concrete game state (Hangman.idr)

```
data Game : GameState -> Type where
    GameStart : Game NotRunning
    GameWon   : (word : String) -> Game NotRunning
```

The diagram shows two annotations pointing to the GameStart and GameWon constructors. An annotation for GameStart is enclosed in a box and points to the text: "Game is NotRunning because it hasn't started yet". An annotation for GameWon is enclosed in a box and points to the text: "Game is NotRunning because the player has won".

```

→ GameLost   : (word : String) -> Game NotRunning
InProgress : (word : String) -> (guesses : Nat)
            -> (missing : Vect letters Char)
            -> Game (Running guesses letters)
Game is NotRunning
because the player has lost
Game is in progress,
with guesses and
letters remaining

```

It's convenient to define a `Show` implementation for `Game` so that you can easily display a string representation of a game's progress.

Listing 14.12 A `Show` implementation for `Game` (`Hangman.idr`)

```

Show (Game g) where
    show GameStart = "Starting"
    show (GameWon word) = "Game won: word was " ++ word
    show (GameLost word) = "Game lost: word was " ++ word
    show (InProgress word guesses missing)
        = "\n" ++ pack (map hideMissing (unpack word))
        ++ "\n" ++ show guesses ++ " guesses left"
        where hideMissing : Char -> Char
              hideMissing c = if c `elem` missing then '-' else c

```

Only displays the characters that the player has successfully guessed

You can use `Game` to keep track of the concrete game state. When you execute a `GameLoop`, you'll take a concrete game state as input, and return a result along with the updated game state:

```

data Fuel = Dry | More (Lazy Fuel)

The input game state → run : Fuel ->
                           Game instate ->
                           GameLoop ty instate outstate_fn -> ←
                           IO (GameResult ty outstate_fn)

```

A command to update the game state

GameResult, defined shortly, pairs the result value and the output state.

You use `Fuel`, because `run` potentially loops. In particular, when you read a `Guess` from the player, the only valid input is a single alphabetical character, so you'll need to keep asking for input until it's valid.

If you run out of fuel, the `GameResult` needs to say that execution failed. Otherwise, it needs to store the result of the operation and the new state. Crucially, the type of the new state might depend on the result; for example, the number of guesses available is different depending on whether `Guess` returns `Correct` or `Incorrect`. A `GameResult`, therefore, is one of the following:

- A pair of the result produced by executing the game and the output state, with a type calculated from the result
- An error, if `run` ran out of fuel

You can define `GameResult` as follows:

```

data GameResult : (ty : Type) -> (ty -> GameState) -> Type where
    OK : (res : ty) -> Game (outstate_fn res) -> ←
        GameResult ty outstate_fn
    OutOfFuel : GameResult ty outstate_fn

```

As in the definition of GameCmd, the argument to Game is calculated from the result, `res`.

`outstate_fn` is included in the type of `GameResult` because then you're explicit in the type about how you're calculating the output state of Game.

Now that you have a data type to represent the concrete state of a game—which takes the abstract state as an argument—along with a representation of the result, you're ready to implement `run`.

14.3.5 Running the game: executing GameLoop

The next listing outlines the definition of `run` for `GameLoop`. This uses another function, `runCmd`, to execute `GameCmd`. There's a hole for the definition of `runCmd` for the moment.

Listing 14.13 Running the game loop (Hangman.idr)

Runs a command. We'll leave a hole for this, and define it shortly.

```
runCmd : Fuel ->
    Game instate -> GameCmd ty instate outstate_fn ->
    IO (GameResult ty outstate_fn)
runCmd fuel state cmd = ?runCmd_rhs

run : Fuel -> Game instate -> GameLoop ty instate outstate_fn ->
    IO (GameResult ty outstate_fn)
run Dry _ _ = pure OutOfFuel
run (More fuel) st (cmd >>= next)
    = do OK cmdRes newSt <- runCmd fuel st cmd
        | OutOfFuel => pure OutOfFuel
        run fuel newSt (next cmdRes)
run (More fuel) st Exit = pure (OK () st)
```

When successful, run returns the result of the operation (`cmdRes` here) and an updated state (`newSt` here).

First command ran out of fuel

In `run`, when it's successful, you use `pure` to return a pair of the result and the new state. Because you'll often return a result in this form when executing a command, you can define a helper function, `ok`, to make this more concise:

```
ok : (res : ty) -> Game (outstate_fn res) ->
    IO (GameResult ty outstate_fn)
ok res st = pure (OK res st)
```

Using `ok`, you can refine the last clause of `run` to the following:

```
run (More fuel) st Exit = ok () st
```

Listing 14.14 gives an outline definition of `runCmd`, leaving holes for the `Guess` and `ReadGuess` cases. In the other cases, you use `ok` to update the state as required by the type and execute `IO` actions as necessary.

Listing 14.14 Outline definition of `runCmd` (Hangman.idr)

```
runCmd : Fuel -> Game instate -> GameCmd ty instate outstate_fn ->
    IO (GameResult ty outstate_fn)
runCmd fuel state (NewGame word)
    = ok () (InProgress (toUpper word) _ (fromList (letters word)))
```

Creates a new in-progress game, using letters to extract unique letters from the word

Updates to the NonRunning game state

```
runCmd fuel (InProgress word _ missing) Won
⇒      = ok () (GameWon word)
runCmd fuel (InProgress word _ missing) Lost
⇒      = ok () (GameLost word)
```

```

runCmd fuel state (Guess c) = ?runCmd_rhs_4
runCmd fuel state ShowState = do println state
                                ok () state
runCmd fuel state (Message str) = do putStrLn str
                                ok () str
runCmd fuel state ReadGuess = ?runCmd_rhs_7

runCmd fuel state (Pure res) = ok res state
runCmd fuel st (cmd >= next)
    = do OK cmdRes newSt <- runCmd fuel st cmd
        | OutOfFuel => pure OutOfFuel
        runCmd fuel newSt (next cmdRes)

```

**Prints a message
and continues with
current state**

As with run, you execute the first command and continue with the result and state it produces.

When a game is in progress, the game state uses the `InProgress` constructor of `Game`, which has the following type:

```
*Hangman> :t InProgress
InProgress : String ->
              (guesses : Nat) ->
              Vect letters Char -> Game (Running guesses letters)
```

The third argument is a vector of the letters that are still to be guessed. So, in the Guess case, you check whether the guessed character is in the vector of missing letters:

```

runCmd fuel (InProgress word _ missing) (Guess c)
> = case isElem c missing of
    Yes prf => ok Correct (InProgress word _ )
    No contra => ok Incorrect (InProgress word _ )

```

**Correct, so remove the letter from
the vector of missing letters**

Incorrect. The `guesses` argument to `InProgress` is inferred from the type, and decreases.

Returns a proof of whether c is in the vector of missing letters

You defined `removeElem` interactively in chapter 9 in the earlier implementation of Hangman. For convenience, I'll repeat it here:

Finally, you need to define the `ReadGuess` case, which reads a character from the player. The input is only valid if it's an alphabetical character, so you loop until the player enters a valid input:

```
runCmd (More fuel) st ReadGuess = do
    putStrLn "Guess: "
    input <- getLine
    case unpack input of
        [x] => if isAlpha x
            then ok (toUpper x) st
            else do putStrLn "Invalid input"
                    runCmd fuel st ReadGuess
        _ => do putStrLn "Invalid input"
                runCmd fuel st ReadGuess
runCmd Dry _ _ = pure OutOfFuel
```

The annotations are as follows:

- A callout points to the condition `[x]` with the text "Reads a single alphabetical character, so input is valid".
- A callout points to the case `_ =>` with the text "Invalid input, so prompt again but with reduced fuel".
- A callout points to the final line `runCmd Dry _ _ = pure OutOfFuel` with the text "Needed to catch the case where ReadGuess runs out of fuel".

This case might loop indefinitely if the user continues to enter invalid input, so `runCmd` takes `Fuel` as an argument and consumes fuel whenever there's an invalid input. As a result, `runCmd` itself remains total because it either consumes fuel or processes a command on each recursive call. It's important for `runCmd` to be total, because it means you know that executing a `GameCmd` will continue to make progress as long as there are commands to execute.

You're now in a position to write the main program, using `forever` to ensure that, in practice, `run` never runs out of fuel. Add the following to the end of `Hangman.idr`:

```
%default partial
forever : Fuel
forever = More forever
```

The annotations are as follows:

- A callout points to the `%default partial` line with the text "Analogous to `%default total`, this means that all following definitions are allowed to be partial".
- A callout points to the `main` definition with the text "Initializes with the `GameStart` state, meaning that no game is running".

```
main : IO ()
main = do run forever GameStart hangman
         pure ()
```

You should now be able to execute the game at the REPL. Here's an example:

```
*Hangman> :exec
-----
6 guesses left
Guess: t
Correct

T--T--
6 guesses left
Guess: x
Incorrect

T--T--
5 guesses left
Guess: g
Correct
```

```
T--T--G
5 guesses left
Guess: bad
Invalid input
Guess:
```

In this example, we've separated the *description* of the rules, in `GameCmd` and `GameLoop`, from the *execution* of the rules, in `runCmd` and `run`. Essentially, `GameCmd` and `GameLoop` define an interface for constructing a valid game of Hangman, correctly following the rules. Any well-typed total function using these types *must* be a correct implementation of the rules, or it wouldn't have type-checked!

14.4 Summary

- You can get feedback from the environment by using the result of an operation to compute the output state of a command.
- A system might be in a different state after running a command, depending on whether the command was successful.
- Defining preconditions on operations allows you to express security properties in types, such as when it's valid for an ATM to dispense cash.
- A system might change state according to whether a user's input is valid in the current environment, such as whether a PIN or password is correct.
- Predicates and `auto` implicits help you describe valid input states of operations precisely.
- You can describe the rules of a game precisely in a type, so that a function that type-checks must be a valid implementation of the rules.
- You use an abstract state type to describe what operations do, and a concrete state, depending on the abstract state, to describe their corresponding implementations.

Type-safe concurrent programming

This chapter covers

- Using concurrency primitives
- Defining a type for describing concurrent processes
- Using types to ensure concurrent processes communicate consistently

In Idris, a value of type `IO ()` describes a sequence of actions for interacting with the user and operating system, which the runtime system executes *sequentially*. That is, it only executes one action at a time. You can refer to a sequence of interactive actions as a *process*.

As well as executing actions in sequence in a single process, it's often useful to be able to execute multiple processes at the same time, *concurrently*, and to allow those processes to communicate with each other. In this chapter, I'll introduce concurrent programming in Idris.

MESSAGE PASSING Concurrent programming is a large topic, and there are several approaches to it that would fill books of their own. We'll look at some small examples of *message-passing* concurrency, where processes

interact by sending messages to each other. Message passing is supported as a primitive by the Idris runtime system. In effect, sending a message to a process and receiving a reply corresponds to calling a method that returns a result in an object-oriented language.

Concurrent programming has several advantages:

- You can continue to interact with a user while a large computation runs. For example, a user can continue browsing a web page while a large file downloads.
- You can display feedback on the progress of a large computation running in another process, such as displaying a progress bar for a download.
- You can take full advantage of the processing power of modern CPUs, dividing work between multiple processes running on separate CPU cores.

This chapter presents a larger example of type-driven development. First, I'll introduce the primitives for concurrent programming in Idris, and describe the problems that can arise in concurrent processes in general. Then, I'll present an initial attempt at a type for describing concurrent processes. This initial attempt will have some shortcomings, so we'll refine it and arrive at a type that allows processes to communicate with each other safely and consistently.

15.1 **Primitives for concurrent programming in Idris**

The Idris base library provides a module, `System.Concurrency.Channels`, that contains primitives for starting concurrent processes and for allowing those processes to communicate with each other. This allows you, in theory, to write applications that make efficient use of your CPUs and that remain responsive even when executing complex calculations.

But despite its advantages, concurrent programming is notoriously error-prone. The need for multiple processes to interact with each other greatly increases a program's complexity. For example, if you're displaying a progress bar while a file downloads, the process downloading the file needs to coordinate with the process displaying the progress bar so that it knows how much of the file is downloaded. This complexity leads to new ways in which programs can fail at runtime:

- *Deadlock*—Two or more processes are waiting for each other to perform some action before they can continue.
- *Race conditions*—The behavior of a system depends on the ordering of actions in multiple concurrent processes.

The effect of a deadlock is that the processes concerned freeze, no longer accepting input or giving output. Two concurrent processes—let's call them `client` and `server`—could deadlock if `client` is waiting to receive a message from `server` at the same time that `server` is waiting to receive a message from `client`. If this happens, both `client` and `server` will freeze.

Race conditions can be harder to identify. The pseudocode in figure 15.1 for client and server illustrates a race condition, where the value of the shared variable var depends on the order in which the concurrent operations execute. We'll assume that the Read and Write operations respectively read and write the value of a shared mutable variable, so Read var reads the value of the shared variable var.




Figure 15.1 Pseudocode for client and server, in which each process reads and writes a shared variable, var. This could lead to an unexpected result if lines A and C are executed before lines C and D.

Here, client and server execute concurrently, and the final value of var depends on the order in which the operations A, B, C, and D are executed. A will always run before B, and C before D, but otherwise there are six possible orderings for the operations. Table 15.1 lists these orderings and the resulting value of var in each case, assuming an initial value of 1.

Table 15.1 Value of var for each sequence of operations in figure 15.1, given an initial value of 1 for var

Operation order	Value of var
A, B, C, D	3
A, C, B, D	2
A, C, D, B	2
C, A, B, D	2
C, A, D, B	2
C, D, A, B	3

As the table shows, with an initial value for var of 1, there are two possible results for var, depending on the order in which the Read and Write operations run. Here, there are only two processes with two operations each. As programs grow, the likelihood of this kind of nondeterministic result becomes much larger.

Later in this chapter, using a variety of techniques that we've discussed earlier in this book, you'll see how to write concurrent programs in Idris, avoiding problems such as deadlock and race conditions. But first, you need to understand the primitives

that the Idris base library provides for concurrent programming and see the kinds of problems you'll encounter when writing processes that need to coordinate with each other.

15.1.1 Defining concurrent processes

In a complete Idris program, the `main` function, which has type `IO()`, describes the actions that the runtime system will execute when the program is run. The `main` function, therefore, describes the actions that are executed in a single process.

The actions we've used have described console and file I/O operations, but the runtime system also supports actions for starting new processes and sending *messages* between processes. There are primitive operations for the following actions:

- Creating a new process. Each process is associated with a unique *process identifier* (PID).
- Sending a message to a process identified by its PID.
- Receiving a message from another process.

Figure 15.2 illustrates one way we can use message passing to write concurrent processes that communicate with each other. In this figure, `main` and `adder` are two Idris processes, running concurrently, where each can send a message to the other.

The message that `main` sends to `adder` uses the following type:

```
data Message = Add Nat Nat
```

In this example, after `main` sends the message `Add 2 3` to the `adder` process, it expects to receive a reply with the result of adding 2 and 3. In this way, you can use concurrently running processes to implement *services* that respond to requests sent by other processes. Here, we have a service that performs addition, running in a separate process.




Figure 15.2 Message-passing between processes. A `main` process sends an `Add 2 3` message to an `adder` process, which replies by sending a `5` message back to `main`.

I'll use this as a running example of concurrent processes. Next, you'll see how to use the concurrency primitives that Idris provides to implement an adder service in a separate process, and how to use that service from `main`.

15.1.2 The Channels library: primitive message passing

The `System.Concurrency.Channels` module, in the base library, defines data types and actions that allow Idris processes to create new processes and communicate with each other. It defines the following types:

- `PID`—Represents a *process identifier*. Every process is associated with a `PID` that allows you to set up a *communication channel* with the process.
- `Channel`—A link between two processes. It defines a communication channel along which you can send messages.

The `System.Concurrency.Channels` module also defines the following functions for creating new processes and setting up communication channels:

- `spawn`—Creates a new process for executing a sequence of actions of type `IO ()`, and if successful returns its `PID`.
- `connect`—Initiates a communication channel to a running process.
- `listen`—Waits for another process to initiate a communication. When another process connects, it sets up a communication channel with that process.

Finally, the module defines operations for sending and receiving messages on a `Channel`. We'll come to the type definitions in `System.Concurrency.Channels` shortly, but first let's look at how the processes and channels are set up for the `main` and `adder` processes. The `main` process works as follows:

- 1 Starts the adder process using `spawn`
- 2 Sets up a communication channel to `adder` using `connect`
- 3 Sends an `Add 2 3` message on the channel
- 4 Receives a reply on the channel with the result

Correspondingly, the `adder` process works as follows:

- 1 Waits for another process to initiate a communication using `listen` and sets up a new communication channel when another process connects
- 2 Receives a message on the channel containing the numbers to be added
- 3 Sends a message on the channel with the result
- 4 Returns to step 1, waiting for the next request

The `adder` process provides a long-running service that waits for incoming requests and sends replies to those requests. Once `main` has started the `adder` process, any other process can also send requests to `adder`, as long as it knows `adder`'s `PID`.

The first listing shows the type declarations for channels and `PIDs` in `System.Concurrency.Channels` and the functions you can use to create processes and set up communication channels.

Listing 15.1 Channels and PIDs (defined in System.Concurrency.Channels)

A PID is a process identifier that one process can use to initiate communication with another.

```

    data PID : Type
    data Channel : Type
    spawn : (process : IO ()) -> IO (Maybe PID)
    connect : (pid : PID) -> IO (Maybe Channel)
    listen : (timeout : Int) -> IO (Maybe Channel)
  
```

A Channel is a link between two processes, allowing them to communicate.

Starts a new concurrent process, returning a Maybe PID. Returns Nothing if it can't create the process.

Initiates a session for communicating with a process with identifier pid

Initiates a session by waiting for another process to connect

When you listen for a connection or connect to another process, there's no guarantee that you'll succeed in setting up a communication channel. There may be no incoming connection, or the process you're connecting to may no longer be running. So, to capture the possibility of failure, connect and listen return a value of type Maybe Channel.

You can use spawn, listen, and connect to set up adder and main as separate processes. The following listing outlines a program that sets up the processes, leaving holes for the parts of the processes that send messages to each other.

Listing 15.2 Outline of a program that sets up an adder process (AdderChannel.idr)

```

import System.Concurrency.Channels

adder : IO ()
adder = do Just sender_chan <- listen 1
          | Nothing => adder
          ?adder_rhs

main : IO ()
main = do Just adder_id <- spawn adder
          | Nothing => putStrLn "Spawn failed"
          Just chan <- connect adder_id
          | Nothing => putStrLn "Connection failed"
          ?main_rhs
  
```

If there's no incoming connection after 1 second, loops

Waits 1 second for an incoming connection. sender_chan is the name of the channel to receive messages from the sender.

Creates a new process for executing adder

Creates a channel for sending messages to adder

If connect returns Nothing, the adder process is no longer running.

spawn might fail if there aren't enough system resources, so checks the result

Now that you've set up the processes and channels, main can send a message to adder, and adder can reply. You can use the following primitives:

- unsafeSend—Sends a message of *any* type
- unsafeRecv—Receives a message of an *expected* type

As the names imply, these primitives are *unsafe*, because they don't provide a way of checking that the sender and receiver are expecting messages to be sent in a specific

order, or that they're sending and receiving messages of consistent types. Nevertheless, for the moment we'll use them to complete the implementations of `main` and `adder`. Later, in section 15.2, you'll see how to make safe versions that ensure that processes send and receive messages with a consistent protocol.

WHY SUPPORT AN UNSAFE CHANNEL TYPE? It might seem surprising that Idris, a language that's designed to support type-driven development, supports such unsafe concurrency primitives rather than something more sophisticated. The reason is that there are many possible methods for implementing safe concurrent programs in a type-driven way, and by providing unsafe underlying primitives, Idris is not limited to only one of them. You'll see one such method shortly.

The following listing shows the type declarations for these primitive operations, defined in `System.Concurrency.Channels`.

**Listing 15.3 Primitive message passing
(defined in `System.Concurrency.Channels`)**

```
unsafeSend : Channel -> (val : a) -> IO Bool
unsafeRecv : (expected : Type) -> Channel -> IO (Maybe expected)
```

In the next listing, you can complete the definition of `adder` by receiving a request from the sender and then sending a reply. With `unsafeRecv`, you assert that the request is of type `Message`.

Listing 15.4 Complete definition of adder (AdderChannel.idr)

```
data Message = Add Nat Nat
adder : IO ()
adder = do Just sender_chan <- listen 1
           | Nothing => adder
           Just msg <- unsafeRecv Message sender_chan
           | Nothing => adder
           case msg of
             Add x y => do ok <- unsafeSend sender_chan (x + y)
                           adder
Sends a reply on the channel, with the sum of the inputs of type Nat
```

Message is the type that adder expects to receive.

Waits for a message on the channel you've just created, of type Message

Similarly, the following listing shows the complete definition of `main`, which sends a message using `unsafeSend` and receives a reply of type `Nat` with `unsafeRecv`.

Listing 15.5 Complete definition of main (AdderChannel.idr)

```
main : IO ()
main = do Just adder_id <- spawn adder
           | Nothing => putStrLn "Spawn failed"
           Just chan <- connect adder_id
           | Nothing => putStrLn "Connection failed"
           ok <- unsafeSend chan (Add 2 3)
           Just answer <- unsafeRecv Nat chan
Waits for a reply on the channel, of type Nat
```

Sends a message on the channel you've just created, of type Message

```

    | Nothing => putStrLn "Send failed"
printLn answer      ← Prints the result received from adder

```

If you compile and execute `main` using `:exec` at the REPL, you'll see that it receives the result 5 from adder:

```
*AdderChannel> :exec main
5
```

This only works because you've ensured that `main` and `adder` agree on a communication pattern. When `main` sends a message on a channel, the `adder` process is expecting to receive a message on its corresponding channel, and vice versa.

15.1.3 Problems with channels: type errors and blocking

Channels provide a primitive method of setting up links between processes and sending messages across those links. The types of `unsafeSend` and `unsafeRecv` don't, however, provide any kind of guarantees about how processes *coordinate* with each other. As a result, it's easy to make a mistake.

For example, `adder` sends a `Nat` in reply to `main`, but what if `main` is expecting to receive a `String`?

```

main : IO ()
main = do Just adder_id <- spawn adder
          | Nothing => putStrLn "Spawn failed"
          Just chan <- connect adder_id
          | Nothing => putStrLn "Connection failed"
          ok <- unsafeSend chan (Add 2 3)
          Just answer <- unsafeRecv String chan
          | Nothing => putStrLn "Send failed"
printLn answer

```

main expects a String, but adder sends a Nat, so behavior here is undefined.

In this case, executing `main` will behave unpredictably, and will most likely crash, because at runtime there's an inconsistency between the type of the message received and the expected type. `Nothing` in the types of `unsafeSend` and `unsafeRecv` explains how the sends and receives are coordinated between the two processes, so Idris is happy to accept `main` as valid even though the coordination is, in this case, incorrect.

A different problem occurs if the `unsafeSend` and `unsafeReceive` operations don't correspond in each process. For example, `main` might send a second message on the same channel and expect a reply:

```

main : IO ()
main = do Just adder_id <- spawn adder
          | Nothing => putStrLn "Spawn failed"
          Just chan <- connect adder_id
          | Nothing => putStrLn "Connection failed"
          ok <- unsafeSend chan (Add 2 3)
          Just answer <- unsafeRecv Nat chan
          | Nothing => putStrLn "Send failed"
printLn answer

```

main executes successfully up to here.

```

ok <- unsafeSend chan (Add 3 4)

```

```
Just answer <- unsafeRecv Nat chan
| Nothing => putStrLn "Send failed"
putStrLn answer
```

Execution blocks here, because adder only replies to the first message on the channel.

Even though this type-checks successfully, when you attempt to execute it, it will print the first reply from adder but block while waiting for the second. After adder creates the channel using listen, it only replies to one message on that channel.

Although channels themselves are unsafe, you can use them as a primitive for defining type-safe communication. We'll define a type for *describing* the coordination between communicating processes, and then write a run function that executes the description using the unsafe primitives. Even though, ultimately, you'll need to use the primitives, you can encapsulate all the details in a single, descriptive type that you can then use for type-driven development of communicating systems.

15.2 Defining a type for safe message passing

In the Idris runtime system, concurrent processes run independently of each other. There's no shared memory, and the only way processes can communicate with each other is by sending messages to each other. Because there's no shared memory, there are no race conditions caused by accessing shared state simultaneously, but there are several other problems to consider:

- How can you ensure that the type of the message that main sends is the same type that adder expects to receive?
- What happens if main sends a message to adder, but adder doesn't reply?
- What happens if adder has stopped running when main sends a message?
- How can you prevent the situation where main and adder are both waiting for a message from each other?

In this section, you'll see how to solve these problems by defining a Process type, which allows you to describe well-typed communicating processes.

Types and concurrent programming

Support for types in concurrent programming has generally been quite limited in mainstream programming languages, with some exceptions, such as typed channels in Go. One difficulty is that, as well as the types of messages that a channel can carry, you also need to think about the *protocol* for message passing. In other words, as well as what to send (the type), you also need to think about when to send it (the protocol).

There has, nevertheless, been significant research into types for concurrent programming, most notably the study of *session types* that began with Kohei Honda's 1993 paper "Types for Dyadic Interaction." The type we'll implement in this section is an instance of a session type with a minimal protocol where a client sends one message and then receives one reply. If you're interested in exploring further, a recent (2016) paper, "Certifying Data in Multiparty Session Types" by Bernardo Toninho and Nobuko Yoshida, describes a more sophisticated way of using types in concurrent programs.

We won't get this implementation right on the first try, however. As is often the case in type-driven development, we'll find that we need to refine the type to address the problems that become apparent after our first attempt. We'll start by defining a type that's specific to the adder service, and later refine it to support generic services that are guaranteed to respond to requests indefinitely.

15.2.1 Describing message-passing processes in a type

Earlier, I described two problems with the primitive `Channel` type that make concurrent programming, in this primitive form, unsafe:

- There's no way to check that a response to a request has the correct type.
- There's no way to check the correspondence between sending and receiving messages in communicating processes.

We'll solve both of these problems by defining a type for describing processes and then refining it as necessary to support the message-passing features we need.

To start, you can define a process type that supports `IO` actions, constructing pure values, and sequencing.

Listing 15.6 A type for describing processes (`Process.idr`)

```
data Process : Type -> Type where
  Action : IO a -> Process a
  Pure : a -> Process a
  (>=) : Process a -> (a -> Process b) -> Process b
  run : Process t -> IO t
  run (Action act) = act
  run (Pure val) = pure val
  run (act >= next) = do x <- run act
                         run (next x)
```

USING IO IN ACTION By using `IO` in Action, you can include arbitrary `IO` actions in processes, such as writing to the console or reading user input. This is a bit too general, because `IO` actions include, among other things, the unsafe communication primitives. You could restrict this by defining a more precise command type (see chapter 11 for an example), but we'll stick with `IO` for this example.

At the moment, `Process` is nothing more than a wrapper for sequences of `IO` actions. The next step is to extend it to support spawning new processes. You can define a data type for representing processes that can receive a `Message`, using `PID` from `System.Concurrency.Channels`:

```
data MessagePID = MkMessage PID
```

Next, you can add a constructor to `Process` that describes an action that spawns a new process and returns the `MessagePID` of that process, if it was successful:

```
Spawn : Process () -> Process (Maybe MessagePID)
```

You also need to extend run to be able to execute the new Spawn command. This spawns a new process using the spawn primitive and then returns a MessagePID containing the new PID:

```
run (Spawn proc) = do Just pid <- spawn (run proc)
                      | Nothing => pure Nothing
                      pure (Just (MkMessage pid))
```

ADDING NEW CONSTRUCTORS Remember that after you've added the Spawn constructor, you can add the missing cases to run in Atom by pressing Ctrl-Alt-A with the cursor over the name run.

Next, you can add commands to allow processes to send messages to each other. In the previous examples, the main process sent requests of type Message and waited for corresponding replies of type Nat. You can encapsulate this behavior in a single Request command:

```
Request : MessagePID -> Message -> Process (Maybe Nat)
```

The reason for returning Maybe Nat, rather than Nat, is that you don't have any guarantee that the process to which MessagePID refers is still running. When you run a Request, you'll need to connect to the process that services the request, send it a message, and then wait for the reply:

```
run (Request (MkMessage process) msg)
    = do Just chan <- connect process      ┏━━━ Connects to the server process
          | _ => pure Nothing
          ok <- unsafeSend chan msg        ┏━━━ Connecting succeeded, so sends the request
          if ok then do Just x <- unsafeRecv Nat chan
                      | Nothing => pure Nothing
                      pure (Just x)           ┏━━━ No reply received,
          else pure Nothing                ┏━━━ Reply successfully
                                            received
                                            | Sending the message
                                            failed, so no result
```

ENCAPSULATING PRIMITIVES IN THE PROCESS TYPE You still need to use unsafeSend and unsafeRecv, but by encapsulating them in the Process data type, you know there's only one place in your program where you use the unsafe primitives. You need to be careful to get this definition right, but once you do, you know that any message-passing program implemented in terms of the Process type will follow the message-passing protocol correctly.

The adder process waited for an incoming message, calculated a result, and sent a response back to the requester. You can encapsulate this behavior in a single Respond command:

```
Respond : ((msg : Message) -> Process Nat) -> Process (Maybe Message)
```

This takes a function as an argument, which, when given a message received from a requester, calculates the Nat to send back. It returns a value of type Maybe Message, which is either Nothing, if it didn't process an incoming message, or of the form Just

`msg`, if it processed an incoming message (`msg`). This is useful if you need to do any further processing with the incoming message, even after sending a response.

When you run the `Respond` command, you'll wait for one second for a message and then, if there is a message, calculate the response and send it back:

```
run (Respond calc)
  = do Just sender <- listen 1    ← Waits for 1 second for an incoming connection
      | Nothing => pure Nothing ← No incoming connection, so does nothing
      No message received, so does nothing
      Just msg <- unsafeRecv Message sender
      | Nothing => pure Nothing
      res <- run (calc msg)      ← Calculates the response to the message
      unsafeSend sender res     ← Sends the response, of type Nat,
      pure (Just msg)           ← back to the requesting process
```

ALTERNATIVE IMPLEMENTATIONS FOR RESPOND This implementation of the `Respond` case waits for 1 second if there's no incoming message. An alternative, and more flexible, implementation might allow the user to specify a timeout. For example, if there's no incoming request, it might not make sense to continue waiting if a process has other work to do.

Listing 15.7 shows how you can define `adder` and `main` using `Process`. We'll call them `procAdder` and `procMain` to distinguish them from the earlier versions. In `procAdder`, you use `Respond` to explain how to respond to a `Message`, and in `procMain` you use `Request` to send a message to a spawned process.

Listing 15.7 Implementing a type-safe adder process (Process.idr)

```
procAdder : Process ()
procAdder = do Respond (\msg => case msg of
                                Add x y => Pure (x + y)) ←
Continues waiting for responses
          procAdder
procMain : Process ()
procMain = do Just adder_id <- Spawn procAdder
             | Nothing => Action (putStrLn "Spawn failed")
             Just answer <- Request adder_id (Add 2 3)
             | Nothing => Action (putStrLn "Request failed")
             Action (printLn answer)

Spawns a process that must send and receive messages according to the Process protocol
Sends a request. If successful, the response is given by answer.
```

Responds to a message of the form `Add x y` by sending the response `x + y`

You can try this at the REPL, using `run` to translate `procMain` to a sequence of `IO` actions:

```
*Process> :exec run procMain
5
```

Unlike the previous version, `procMain` can't expect to receive a `String` rather than a `Nat`, because the type of `Request` doesn't allow it. You've also encapsulated the

communication protocol on a channel using Request and Respond, so you know that you won't send or receive too many messages after creating a channel.

As a first attempt, this is an improvement over the primitive implementation with Channel1, but there are a number of ways you can improve it. For example, procAdder is not total:

```
*Process> :total procAdder
Main.procAdder is possibly not total due to recursive path:
  Main.procAdder
```

This is potentially a problem, because a process that isn't total may not successfully respond to requests. As a first refinement, you can modify the Process type, and correspondingly the definition of run, so that indefinitely running processes like procAdder are total.

15.2.2 Making processes total using Inf

As you saw in chapter 11, you can mark parts of data as potentially infinite using Inf:

```
Inf : Type -> Type
```

You can then say a function is total if it produces a *finite prefix* of constructors of a well-typed infinite result in finite time. In practice, this means that any time you use a value with an Inf type, it needs to be an argument to a data constructor or a nested sequence of data constructors.

You've seen various ways to use Inf to define potentially infinite processes in chapters 11 and 12. Here, you can use it to explicitly mark the parts of a process that loop by adding the following constructor to Process:

```
Loop : Inf (Process a) -> Process a
```

For reference, the next listing shows the current definition of Process, including Loop, defined in a new file, ProcessLoop.idr.

Listing 15.8 New Process type, extended with Loop (ProcessLoop.idr)

```
data Message = Add Nat Nat    ← The type of messages that a process can send
data MessagePID = MkMessage PID    ← PIDs for processes that can respond to messages
data Process : Type -> Type where
  Request : MessagePID -> Message -> Process (Maybe Nat)
  Respond : ((msg : Message) -> Process Nat) -> Process (Maybe Message)
  Spawn : Process () -> Process (Maybe MessagePID)
  Loop : Inf (Process a) -> Process a
  Action : IO a -> Process a
  Pure : a -> Process a
  (>>=) : Process a -> (a -> Process b) -> Process b
```

Descriptions of processes that can loop indefinitely

Explicitly loops, executing a potentially infinite process

Using Loop, you can define procAdder as follows, explicitly noting that the recursive call to procAdder is a potentially infinite process:

```
procAdder : Process ()
procAdder = do Respond (\msg => case msg of
                                Add x y => Pure (x + y))
              Loop procAdder
```

This version of procAdder is total:

```
*ProcessLoop> :total procAdder
Main.procAdder is Total
```

By using an explicit Loop constructor, you can mark the infinite parts of a Process so that you can at least be sure that any infinite recursion is intended. Moreover, as you'll see in the next section, it will allow you to refine Process further so that you can control exactly *when* a process is allowed to loop.

You'll also need to extend run to support Loop. The simplest way is to execute the action directly:

```
run (Loop act) = run act
```

Unfortunately, this new definition of run isn't total because the totality checker (correctly!) doesn't believe that act is a smaller sequence than Loop act:

```
*ProcessLoop> :total run
Main.run is possibly not total due to recursive path:
  Main.run, Main.run, Main.run
```

As with the infinite processes in chapter 11, you can define a Fuel data type to give an explicit execution limit to run. Every time you Loop, you reduce the amount of Fuel available. The following listing shows how you can extend run so that it terminates when it runs out of Fuel, following the pattern you've already seen in chapter 11.

Listing 15.9 New run function, with an execution limit (ProcessLoop.idr)

```
data Fuel = Dry | More (Lazy Fuel)

run : Fuel -> Process t -> IO (Maybe t)
run Dry _ = pure Nothing
run fuel (Request (MkMessage process) msg)
    = do Just chan <- connect process
        | _ => pure (Just Nothing)
        ok <- unsafeSend chan msg
        if ok then do Just x <- unsafeRecv Nat chan
                     | Nothing => pure (Just Nothing)
                     pure (Just (Just x))
                     else pure (Just Nothing)
    run fuel (Respond calc)
        = do Just sender <- listen 1
            | Nothing => pure (Just Nothing)
            Just msg <- unsafeRecv Message sender
            | Nothing => pure (Just Nothing)
```

The code is annotated with two callouts:

- A callout pointing to the line "run Dry _ = pure Nothing" contains the text "Returns Nothing when the process is out of fuel".
- A callout pointing to the line "pure (Just (Just x))" contains the text "Uses Just when the process still has fuel".

The process consumes fuel on every Loop.

```

Just res <- run fuel (calc msg)
    | Nothing => pure Nothing
unsafeSend sender res
pure (Just (Just msg))

run (More fuel) (Loop act) = run fuel act
run fuel (Spawn proc) = do Just pid <- spawn (do run fuel proc
                                                pure ())
    | Nothing => pure Nothing
    pure (Just (Just (MkMessage pid)))
run fuel (Action act) = do res <- act
                           pure (Just res)
run fuel (Pure val) = pure (Just val)
run fuel (act >>= next) = do Just x <- run fuel act
                                | Nothing => pure Nothing
                                run fuel (next x)

```

When you run recursively, you need to check that the result was valid before continuing.

Remember that you can generate an infinite amount of Fuel and allow processes to run indefinitely by using a single partial function, forever:

```

partial
forever : Fuel
forever = More forever

```

Using a single forever function to say how long an indefinite process is allowed to run means that you minimize the number of nontotal functions you need. Because run is total, you know that it will continue executing process actions as long as there are actions to execute. For convenience, you can also define a function for initiating a process and discarding its result:

```

partial
runProc : Process () -> IO ()
runProc proc = do run forever proc
                  pure ()

```

Then, you can try executing procMain as follows, which will display the answer 5 as before:

```
*ProcessLoop> :exec runProc procMain
5
```

Using Loop, you can write processes that loop forever and that are total by being explicit about when they loop. Unfortunately, though, there's still no guarantee that a looping process will respond to any messages at all. For example, you could define procAdder as follows:

```

procAdderBad1 : Process ()
procAdderBad1 = do Action (putStrLn "I'm out of the office today")
                  Loop procAdderBad1

```

Or even like this:

```

procAdderBad2 : Process ()
procAdderBad2 = Loop procAdderBad2

```

Both of these programs type-check, and are both checked as total, but neither will respond to any messages because there's no Respond command. In the case of procAdderBad2, it's total because the recursive call to procAdderBad2 is an argument to the Loop constructor, so it will produce a finite prefix of constructors. Being total using Loop is, therefore, not enough to guarantee that a process will respond to a request.

THE MEANING OF TOTALITY, IN PRACTICE Totality means that you're guaranteed that a function behaves in exactly the way described by its type, so if the type isn't precise enough, neither is the guarantee! With Process, the type isn't precise enough to guarantee that a process contains a Respond command before any Loop.

Furthermore, the Process type is specific to the problem of writing a concurrent service to add numbers. What if you want to write different services? You don't want to have to write a different Process type for every kind of service you might want to Spawn.

To solve these problems, you'll need to refine the Process type in two more ways:

- In section 15.2.3, you'll refine it to ensure that a server process responds to requests on every iteration of a Loop.
- In section 15.2.4, you'll see how to use first-class types to allow Process to respond to *any* kind of message.

15.2.3 Guaranteeing responses using a state machine and Inf

In chapter 13, you saw how to guarantee that systems would execute necessary actions in the correct order by representing a state machine in a type. A server process like adder could be in one of several states, depending on whether or not it has received and processed a request:

- NoRequest—It has not yet serviced any requests.
- Sent—It has sent a response to a request.
- Complete—It has completed an iteration of a loop and is ready to service the next request.

Figure 15.3 illustrates how the Respond and Loop commands affect the state of a process.




Figure 15.3 A state transition diagram showing the states and operations in a server process. A process begins in the NoRequest state and *must* end in the Complete state, meaning that it has responded to at least one request.

If you have a process that begins in the `NoRequest` state and ends in the `Complete` state, you can be certain that it has replied to a request, because the only way to reach the `Complete` state is by calling `Respond`. You can also be certain that it's continuing to receive requests, because the only way to reach the `Complete` state is by calling `Loop`. By expressing the state of a process in its type, you can then make stronger guarantees about how that process behaves.

You can refine the type of `Process` to represent the states before and after the process is executed:

```
data ProcState = NoRequest | Sent | Complete
data Process : Type ->
    (in_state : ProcState) -> ← The state of the process before it's executed
    (out_state : ProcState) -> ← The state of the process after it's executed
    Type
```

The type of the result
of the process

The states in the type give the preconditions and postconditions of a process. For example:

- `Process () NoRequest Complete` is the type of a process that responds to a request and then loops.
- `Process () NoRequest Sent` is the type of a process that responds to one or more requests and then terminates.
- `Process () NoRequest NoRequest` is the type of a process that responds to *no* requests and then terminates.

Listing 15.10 shows the refined `Process`, where the type of each command explains how it affects the overall process state. In this definition, where there's no precondition on the state and no change in the state, the input and output states are both `st`.

Listing 15.10 Annotating the `Process` type with its input and output states (`ProcessState.idr`)

Request now returns a `Nat`
rather than a `Maybe Nat`.

```
data Process : Type ->
    (in_state : ProcState) ->
    (out_state : ProcState) ->
    Type where
    Request : MessagePID -> Message -> Process Nat st st
    Respond : ((msg : Message) -> Process Nat NoRequest NoRequest) -> ←
                Process (Maybe Message) st Sent
    Spawn : Process () NoRequest Complete -> ←
                Process (Maybe MessagePID) st st
```

When processing, stays in the
`NoRequest` state to ensure you don't
start processing a new response
before completing this one

You can respond to a request
at any time, and the resulting
state is `Sent`.

You can only spawn a
process if it's a looping
process (that is, it goes
from `NoRequest` to
`Complete`).

```

Loop : Inf (Process a NoRequest Complete) ->
    Process a Sent Complete
Action : IO a -> Process a st st
Pure : a -> Process a st st
(>>=) : Process a st1 st2 -> (a -> Process b st2 st3) ->
    Process b st1 st3

```

You can only loop with a process that goes from NoRequest to Complete.

You can only Loop after sending a message, and a Loop puts the process into the Complete state.

Sequences operations, going from the input state of the first operation to the output state of the second

RETURN TYPE OF REQUEST Earlier, you sent requests of type `Message` and received responses of type `Maybe Nat`. You used `Maybe` because you had no guarantee that the service was still running, so the request could fail. Now you've set up the `Process` state so that services will *always* respond to requests. If you send a request, you're guaranteed a response of type `Nat` in finite time.

When you use this new definition, there's no way for a function to invoke `Loop` unless the function can satisfy the precondition that it has sent a response to a request. Moreover, `Loop` is also the only way for a process to reach the `Complete` state. As a result, you can only invoke `Loop` on a process that's guaranteed to be looping, because the process must begin in the `NoRequest` state and end in the `Complete` state.

You can still define `procAdder` as before, because each command satisfies the precondition, and its type now states that it *must* respond to a request and then loop:

```

procAdder : Process () NoRequest Complete
procAdder = do Respond (\msg => case msg of
                                Add x y => Pure (x + y))
              Loop procAdder

```

The two incorrect versions defined earlier, however, no longer type-check, because the commands don't satisfy the preconditions given by `Process` when you attempt to use them. For example, you can try the following definition:

```

procAdderBad1 : Process () NoRequest Complete
procAdderBad1 = do Action (putStrLn "I'm out of the office today")
                  Loop procAdder

```

Idris reports an error because there's no `Respond` before the `Loop`:

```

ProcessState.idr:63:21:
When checking right hand side of procAdderBad1 with expected type
  Process () NoRequest Complete

When checking an application of constructor Main.>>=:
  Type mismatch between
    Process a Sent Complete (Type of Loop _)
  and
    Process () NoRequest Complete (Expected type)

  Specifically:
    Type mismatch between

```

Sent
and
NoRequest

This error message means that when you call `Loop`, the process is supposed to be in the `Sent` state, but at this point it's in the `NoRequest` state, not having sent any response yet. You'll get a similar error message for the same reason with the following definition:

```
procAdderBad2 : Process () NoRequest Complete
procAdderBad2 = Loop procAdderBad2
```

In order to execute programs using the refined `Process`, you'll need to modify `run` and `runProc`. First, you need to modify their types:

```
run : Fuel -> Process t in_state out_state -> IO (Maybe t)
runProc : Process () in_state out_state -> IO ()
```

The definitions mostly remain the same as the previous versions. The one change is in the definition of the `Request` case in `run`, now that you know a `Request` will always receive a reply in finite time:

```
run fuel (Request (MkMessage process) msg)
= do Just chan <- connect process
     | _ => pure Nothing
     ok <- unsafeSend chan msg
     if ok then do Just x <- unsafeRecv Nat chan
                  | Nothing => pure Nothing
                  pure (Just x)
     else pure Nothing
```

If communication fails, abandons execution

RETURN VALUE OF RUN If `run` fails for any reason, it returns `Nothing`. Up to now, this could only happen if it ran out of `Fuel`. You've now set up `Process` so that senders and receivers are coordinated, so, at least in theory, communication can't fail. If communication *does* fail, there's either an error in the implementation of `run` or a more serious runtime error, so you can also return `Nothing` in this case.

You'll also need to modify the type of `procMain`, to be consistent with the refined `Process` type. This type explicitly states that `procMain` isn't intended to respond to any incoming requests because it ends in the `NoRequest` state:

```
procMain : Process () NoRequest NoRequest
```

It's convenient to define type synonyms for clients, like `procMain`, and for services, like `procAdder`. They both use `Process`, but they differ in how they affect the state of the process:

```
Service : Type -> Type
Service a = Process a NoRequest Complete

Client : Type -> Type
Client a = Process a NoRequest NoRequest
```

The following listing shows the refined definitions of procAdder and procMain using these type synonyms for client and server processes.

Listing 15.11 Refined definitions of procAdder and procMain (ProcessState.idr)

<p>A Service is a Process that begins in the NoRequest state and ends in the Complete state.</p> <pre style="border: 1px solid black; padding: 5px;">⇒ procAdder : Service () procAdder = do Respond (\msg => case msg of Add x y => Pure (x + y)) Loop procAdder</pre>	<p>A Client is a Process that begins and ends in the NoRequest state.</p> <pre style="border: 1px solid black; padding: 5px;">procMain : Client () procMain = do Just adder_id <- Spawn procAdder Nothing => Action (putStrLn "Spawn failed") answer <- Request adder_id (Add 2 3) Action (println answer)</pre>
---	---

If you try this at the REPL, you'll see that it displays 5 as before:

```
*ProcessState> :exec runProc procMain
5
```

You now have a definition of Process with the following guarantees, ensured by the preconditions and postconditions in the definition of Process:

- All requests of type Message are sent responses of type Nat.
- Every process started with Spawn is guaranteed to loop indefinitely and respond to requests on every iteration.
- Therefore, every time a process sends a Request to a service started with Spawn, it will receive a response in finite time as long as the service is defined by a total function.

This means you can write type-safe concurrent programs that can't deadlock, because every request is guaranteed to receive a response eventually. But at this stage, it only allows you to write one kind of service—one that receives a Message and sends back a Nat. It would be far more useful if you could define *generic* message-passing processes with user-defined interactions between the sender and receiver. As you'll see, you can achieve this with a final small refinement to Process.

15.2.4 Generic message-passing processes

When a process receives a request of the form Add x y, it sends back a response of type Nat. You can express this relationship between the request and the response types in a type-level function:

```
AdderType : Message -> Type
AdderType (Add x y) = Nat
```

This function describes the interface that Process supports: if it receives a message of the form Add x y, it will send a response of type Nat. You could define other interfaces

this way; for example, the following listing gives a description of an interface to a process that responds to requests to perform operations on lists.

Listing 15.12 Describing an interface for List operations

```
data ListAction : Type where
    Length : List elem -> ListAction
    Append : List elem -> List elem -> ListAction
```

The type of messages received by a process that manipulates lists concurrently

```
ListType : ListAction -> Type
ListType (Length xs) = Nat           ← Getting the length of a list results in a Nat.
ListType (Append {elem} xs ys) = List elem   ← Appending two lists with element type elem results in a List elem.
```

In general, an interface to a process is a function like AdderType or ListType that calculates a response type from a request. Instead of defining a specific type that processes can send and receive, you can include the interface as part of a process's type by adding an additional argument for the interface, as figure 15.4 shows.




Figure 15.4 A refined Process type, including the interface that the process responds to as part of the type

REQUEST TYPE OF PROCESS The iface argument to Process includes a type variable, reqType. This is an *implicit* argument, and it defines the type of messages the process can receive. Idris will infer reqType from the iface argument. For example, in procAdder, iface is AdderType, so reqType must be Message.

We'll come to the refined definition of Process shortly. Once it's defined, the procAdder service will respond to an interface defined by AdderType:

```
procAdder : Process AdderType () NoRequest Complete
```

Some processes, like procMain, don't respond to any requests. You can make this explicit in the type by defining their interfaces as follows:

```
NoRecv : Void -> Type
NoRecv = const Void
```

Remember from chapter 8 that `Void` is the empty type, with no values. Because you can never construct a value of type `Void`, a process that provides a `NoRecv` interface can never receive a request. You can use it in the following type for `procMain`:

```
procMain : Process NoRecv () NoRequest NoRequest
```

You'll also need to redefine the type synonyms `Service` and `Client` to include the interface description. A `Service` has an interface, but a `Client` receives no requests:

```
Service : (iface : reqType -> Type) -> Type -> Type
Service iface a = Process iface a NoRequest Complete

Client : Type -> Type
Client a = Process NoRecv a NoRequest NoRequest
```

When you create a new process, you get a PID for the new process as a `MessagePID`. You should only send messages to a process when the messages match the interface of that process, so you can refine `MessagePID` to include the interface it supports in its type:

```
data MessagePID : (iface : reqType -> Type) -> Type where
    MkMessage : PID -> MessagePID iface
```

Now, if you have a PID of type `MessagePID AdderType`, you know that you can send it messages of type `Message`, because that's the input type of `AdderType`.

Putting all of this together, you can refine `Process` to describe its own interface and to be explicit about when it's safe to send a request of a particular type to another `Process`. The next listing shows the refined types for `Request`, `Respond`, and `Spawn`.

Listing 15.13 Refining `Process` to include its interface in the type, part 1 (`ProcessInterface.idr`)

```
data ProcState = NoRequest | Sent | Complete

data Process : (iface : reqType -> Type) ->
    Type ->
    (in_state : ProcState) ->
    (out_state : ProcState) ->
    Type where
    Request : MessagePID service_iface ->
        (msg : service_reqType) ->
        Process iface (service_iface msg) st st
    Respond : ((msg : reqType) ->
        Process iface (iface msg NoRequest NoRequest) ->
        Process iface (Maybe reqType) st Sent
    Spawn : Process service_iface () NoRequest Complete ->
        Process iface (Maybe (MessagePID service_iface)) st st
    {-- continued in Listing 15.14 --}
```

The PID for a process with an interface described by `service_iface`

The req has type `service_reqType`, and `service_iface` has type `service_reqType -> Type`.

The reply from the service will be calculated by applying `service_iface` to the request to ensure that the reply type corresponds with the request.

If you spawn a server with an interface described by `service_iface`, the PID has type `MessagePID service_iface`.

If you receive a req message, the required response type is calculated by `iface req`.

The following listing completes the refined definition of Process, adding Loop, Action, Pure, and ($\gg=$). In each case, all you need to do is add an iface argument to Process.

Listing 15.14 Refining Process to include its interface in the type, part 2 (ProcessIFace.idr)

```
data Process : (iface : reqType -> Type) ->
    Type ->
    (in_state : ProcState) ->
    (out_state : ProcState) ->
    Type where
    {-- continued from Listing 15.13 --}
    Loop : Inf (Process iface a NoRequest Complete) ->
        Process iface a Sent Complete
    Action : IO a -> Process iface a st st
    Pure : a -> Process iface a st st
    ( $\gg=$ ) : Process iface a st1 st2 -> (a -> Process iface b st2 st3) ->
        Process iface b st1 st3
```

In Loop and ($\gg=$), the type explicitly states that the interface does not change.

Finally, you need to update run and runProc for the refined Process definition. Listing 15.15 shows the changes you need to make to run. You need only modify the cases for Request and Respond to be explicit about the types of messages a process expects to receive.

Listing 15.15 Updating run for the refined Process (ProcessIFace.idr)

Brings service_iface into scope so that you can calculate the expected response type

Calculates the expected response type from the message you sent

```
run : Fuel -> Process iface t in_state out_state -> IO (Maybe t)
run fuel (Request {service_iface} (MkMessage process) msg)
    = do Just chan <- connect process
        | _ => pure Nothing
        ok <- unsafeSend chan msg
        if ok then do Just x <- unsafeRecv (service_iface msg) chan
            | Nothing => pure Nothing
            pure (Just x)
        else pure Nothing
run fuel (Respond {reqType} f)
    = do Just sender <- listen 1
        | Nothing => pure (Just Nothing)
        Just msg <- unsafeRecv reqType sender
        | Nothing => pure (Just Nothing)
        Just res <- run fuel (f msg)
        | Nothing => pure Nothing
        unsafeSend sender res
        pure (Just (Just msg))
```

You expect to receive a message that satisfies the interface.

Brings reqType into scope so that you can say it's the expected type of received messages

For runProc, you need only change its type to add the iface argument to Process:

```
partial
runProc : Process iface () in_state out_state -> IO ()
```

```
runProc proc = do run forever proc
                  pure ()
```

When designing data types, especially types that express strong guarantees like `Process`, it's often a good idea to begin by trying to solve a specific problem before moving on to a more general solution. Here, we started with a type for `Process` that only supported specific message and response types (`Message` and `Nat`). Only after that worked did we use type-level functions to make a *generic* `Process` type.

15.2.5 Defining a module for Process

Once you've defined a generic type, it's useful to define a new module to make that type and its supporting functions available to other users. We'll define a new module, `ProcessLib.idr`, that defines `Process` and supporting definitions and exports them as necessary.

The next listing shows the overall structure of the module, omitting the definitions but adding *export* modifiers to each declaration.

Listing 15.16 Defining Process in a module, omitting definitions (ProcessLib.idr)

```
module ProcessLib
import System.Concurrency.Channels

%default total
export
data MessagePID : (iface : reqType -> Type) -> Type where

  public export
  data ProcState = NoRequest | Sent | Complete

  public export
  data Process : (iface : reqType -> Type) ->
    Type ->
    (in_state : ProcState) ->
    (out_state : ProcState) ->
    Type where

  public export
  data Fuel
```

Unless stated otherwise, all definitions in the file must be total.

Exports the MessagePID type but not the constructors

Exports the remaining types and their constructors

The complete definitions are the same as those you've already seen for `MessagePID`, `ProcState`, `Process`, and `Fuel`. Remember from chapter 10 that for data declarations, an `export` modifier can be one of the following:

- `export`—The type constructor is exported but not the data constructors.
- `public export`—The type and data constructors are exported.

The following listing shows how you can export the supporting functions.

**Listing 15.17 Supporting function types for Process, omitting definitions
(ProcessLib.idr)**

```

export partial
forever : Fuel

export
run : Fuel -> Process iface t in_state out_state -> IO (Maybe t)

public export
NoRecv : Void -> Type

public export
Service : (iface : reqType -> Type) -> Type -> Type

public export
Client : Type -> Type

export partial
runProc : Process iface () in_state out_state -> IO ()

```

Exports the types of run and runProc but not their definitions

Exports the types and definitions of the remaining functions

For functions, an export modifier can be one of the following:

- `export`—The type is exported but not the definition.
- `public export`—The type and definition are exported.

Unless there's a specific reason to export a definition as well as a type, it's better to use `export`, hiding the details of the definition. Here, you use `public export` for `Client` and `Service` because these are type synonyms, and other modules will need to know that these are defined in terms of `Process`.

Now that you've defined `Process` and a separate `ProcessLib` module that exports the relevant definitions, we can try more examples. To conclude this section, we'll look at two examples of implementing concurrent programs using this generic `Process` type. First, we'll implement a process using `ListType`, which was defined earlier in listing 15.12, and then we'll look at a larger example using concurrency to run a process in the background to count words in a file.

15.2.6 Example 1: List processing

To demonstrate how you can use `Process` to define services other than `procAdder`, we'll start with a service that responds to requests to carry out functions on `List`. The interface of the service is defined by a `ListType` function. It provides two operations: `Length` and `Append`.

```

data ListAction : Type where
    Length : List elem -> ListAction
    Append : List elem -> List elem -> ListAction

ListType : ListAction -> Type
ListType (Length xs) = Nat
ListType (Append {elem} xs ys) = List elem

```

We'll define a procList service that responds to requests on this interface. It has the following type:

```
procList : Service ListType ()
```

You can define procList incrementally, taking the following steps:

- 1 *Define, type*—As with procAdder, you can implement procList as a loop that responds to a request on each iteration:

```
procList : Service ListType ()
procList = do Respond (\msg => ?procList_rhs)
             Loop procList
```

Looking at the ?procList_rhs type here, you can see that the type you need to produce is calculated from the msg you receive:

```
msg : ListAction
-----
procList_rhs : Process ListType (ListType msg) NoRequest NoRequest
```

- 2 *Define*—Because the type you need to produce depends on the value of msg, you can continue the definition with a case statement, inspecting msg:

```
procList : Service ListType ()
procList = do Respond (\msg => case msg of
                           case_val => ?procList_rhs)
             Loop procList
```

Case splitting on case_val produces this:

```
procList : Service ListType ()
procList = do Respond (\msg => case msg of
                           Length xs => ?procList_rhs_1
                           Append xs ys => ?procList_rhs_2)
             Loop procList
```

- 3 *Type, refine*—For ?procList_rhs_1, if you check the type, you'll see that you need to produce a Nat for the result of Length xs:

```
msg : ListAction
a : Type
xs : List elem
-----
procList_rhs_1 : Process ListType Nat NoRequest NoRequest
```

You can refine ?procList_rhs_1 as follows:

```
procList : Service ListType ()
procList = do Respond (\msg => case msg of
                           Length xs => Pure (length xs)
                           Append xs ys => ?procList_rhs_2)
             Loop procList
```

- 4 *Refine*—To refine ?procList_rhs_2, you need to provide a List elem, and you can complete the definition as follows:

```

procList : Service ListType ()
procList = do Respond (\msg => case msg of
                                Length xs => Pure (length xs)
                                Append xs ys => Pure (xs ++ ys))
                  Loop procList

```

Having completed `procList`, you can try it by spawning it in a process and sending it requests. The following listing defines a process that sends two requests to an instance of `procList` and displays their results.

Listing 15.18 A main program that uses the `procList` service (ListProc.idr)

```

procMain : Client ()
procMain = do Just list <- Spawn procList ← Sets up the procList process
             | Nothing => Action (putStrLn "Spawn failed")
             len <- Request list (Length [1,2,3])
             Action (println len)

Invokes the Length command, returning a Nat ←
             app <- Request list (Append [1,2,3] [4,5,6]) ← Invokes the Append command, returning a List Integer
             Action (println app)

```

You can try this at the REPL as follows:

```
*ListProc> :exec runProc procMain
3
[1, 2, 3, 4, 5, 6]
```

Like `procAdder`, `procList` loops, waiting for incoming requests, and processes them as necessary, but it doesn't do any other computation while waiting for a request. Concurrent processes become far more useful if, rather than spending their time idling and waiting for requests from other processes, they also do some computation. In the next example, you'll see how to do this.

15.2.7 Example 2: A word-counting process

When you define services, you can define separate requests for initiating an action and getting the result of that action. For example, if you're defining a word-count service, you could allow a client to take the following steps:

- 1 Send a request to a word-count service to load a file and count the number of words in the file.
- 2 While the word-count service is processing, the client continues its own work, such as reading input and producing output.
- 3 Send a second request to the word-count service to ask how many words were in the file.

In this example, you'll define the word-count service around the `WCData` record and `doCount` function defined in listing 15.19. This function takes the contents of a file, in a `String`, and produces a structure containing the numbers of words and lines in that content.

Listing 15.19 A small function to count the number of words and lines in a String (WordCount.idr)

```

import ProcessLib           ← Imports this so that you can create a
record WCData where          ← concurrent word-counting process
    constructor MkWCData
    wordCount : Nat
    lineCount : Nat
doCount : (content : String) -> WCData
doCount content = let lcount = length (lines content)
                   wcount = length (words content) in
                   MkWCData lcount wcount

```

See chapter 12 for more on records. This record contains fields for the number of words and lines in a file.

Returns a structure containing the number of words and lines in the content

You can see an example of this in action at the REPL:

```
*WordCount> doCount "test test\ntest"
MkWCData 2 3 : WCData
```

The goal is to implement a process that provides word counting as a service. Rather than loading and counting the words in a single request, you can provide two commands:

- CountFile—Given a filename, loads that file and counts the number of words in it
- GetData—Given a filename, returns the WCData structure for that file, as long as the file has already been processed with CountFile

Rather than returning the WCData structure itself, CountFile will return immediately and continue loading the file in a separate process. That is, one request begins the task and another retrieves the result. This will allow the requester to continue its own work while the word-count service is processing the file. The following listing shows the interface and a skeleton definition for the word-count service.

Listing 15.20 Interface for the word-count service (WordCount.idr)

```

data WC = CountFile String
        | GetData String
WCType : WC -> Type
WCType (CountFile x) = ()
WCType (GetData x) = Maybe WCData

```

Returns () because it merely initiates the processing of a file

Returns Maybe WCData because it will fail if the given file has not yet been processed

```

wcService : (loaded : List (String, WCData)) ->
            Service WCType ()
wcService loaded = ?wcService_rhs

```

wcService loops indefinitely and takes a list of loaded files as an argument.

We'll come to the definition of wcService in a moment. The next listing shows how you can invoke it and continue to execute interactive actions in the foreground while wcService is processing a file in the background.

Listing 15.21 Using the word-count service (WordCount.idr)

```

procMain : Client ()
→ procMain = do Just wc <- Spawn (wcService [])
               | Nothing => Action (putStrLn "Spawn failed")
               Action (putStrLn "Counting test.txt")
               Request wc (CountFile "test.txt") ← Initiates the word count for a file
Starts a new
process for the
word-count
service
               Action (putStrLn "Processing")
               Just wadata <- Request wc (GetData "test.txt")
               | Nothing => Action (putStrLn "File error")
               Action (putStrLn ("Words: " ++ show (wordCount wadata)))
               Action (putStrLn ("Lines: " ++ show (lineCount wadata)))
               Does some work while the
               word-count process is running
               Gets the result of the
               word count of the file

```

Listing 15.22 presents an incomplete implementation of `wcService` that shows how it responds to the commands `CountFile` and `GetData`. Two parts of the definition are missing:

- Loading and processing the requested file
- Looping, with the new file information added to the input, loaded

Listing 15.22 Responding to commands in wcService (WordCount.idr)

```

wcService : (loaded : List (String, WCData)) -> Service WCType ()
wcService loaded
= do Respond (\msg => case msg of
    →
        CountFile fname => Pure ()
        GetData fname =>
            Pure (lookup fname loaded))
    ?wcService_rhs
Returns () immediately so that the
requester can continue processing
Looks up the word-count data in
the existing list of processed files

```

To process the input, you can look at the return value from `Respond`. Remember that `Respond` has the following type:

```

Respond : ((msg : reqType) ->
            Process iface (iface msg) NoRequest NoRequest) ->
            Process iface (Maybe reqType) st Sent

```

The return value from `Respond`, of type `Maybe reqType`, tells you which message, if any, was received. If `wcService` received a `CountFile` command, it could load and process the necessary file before processing its next input.

The next listing shows a further refinement of `wcService`, still including a hole for the function that processes the file.

Listing 15.23 Incomplete implementation of wcService (WordCount.idr)

```

wcService : List (String, WCData) -> Service WCType ()
wcService loaded
= do msg <- Respond (\msg => case msg of
                                CountFile fname => Pure ()
                                GetData fname =>
                                    Pure (lookup fname loaded))
    newLoaded <- case msg of
        Just (CountFile fname) =>
            ?countFile loaded fname
        _ => Pure loaded
    Loop (wcService newLoaded)

```

Calculates the new list of loaded file data, given the msg received

Continues with the new list of loaded files

Respond returns the msg received, which you can now process further.

If asked to process a file, process it here. We'll define countFile shortly.

To see what you need to do to complete wcService, you can check the type of ?countFile:

```

loaded : List (String, WCData)
fname : String
msg2 : Maybe WC
st2 : ProcState
a : Type
-----
countFile : List (String, WCData) ->
    String -> Process WCType (List (String, WCData)) Sent Sent

```

countFile needs to be a function that takes the current list of processed file data and a filename, and then returns an updated list of processed file data. The next listing shows how to define it using doCount, defined earlier, to process the file's contents.

Listing 15.24 Loading a file and counting words (WordCount.idr)

```

countFile : List (String, WCData) -> String ->
    Process WCType (List (String, WCData)) Sent Sent
countFile files fname =
    do Right content <- Action (readFile fname)
       | Left err => Pure files
    let count = doCount content
    Action (putStrLn ("Counting complete for " ++ fname))
    Pure ((fname, doCount content) :: files)

```

Prints a message to the console when processing of the file is complete

Reading the file failed, so doesn't update the files list

Adds the data for the newly processed file to the files list

UPDATING THE HOLE FOR COUNTFILE Remember that you need to define countFile before you use it in wcService. Once you've defined countFile, don't forget to replace the ?countFile hole with a call to countFile.

Now that you've defined `countFile`, you can try executing `procMain`, which starts `wcService`, asks to count the words in a file, `test.txt`, and then displays the result. You'll need to create a `test.txt` file with content like the following:

```
test test
test
test test test
test
```

You can execute `procMain` at the REPL as follows:

```
*WordCount> :exec runProc procMain
Counting test.txt
Processing
Counting complete for test.txt
Words: 4
Lines: 7
```

The diagram shows the sequence of messages exchanged between the client and the service. The first two lines ('Counting test.txt' and 'Processing') are labeled 'Displayed by procMain before requesting to process test.txt'. The third line ('Counting complete for test.txt') is labeled 'Displayed by procMain while wcService is processing'. The last two lines ('Words: 4' and 'Lines: 7') are labeled 'Displayed by wcService when it has finished processing'.

With `Process`, you've defined a type that allows you to describe concurrently executing processes and to explain how processes can send each other messages safely, following a protocol:

- A *service* must respond to every message it receives, on a specific interface, and continue responding in a loop.
- A *client* can then send a message to a process, with the correct interface, and be sure of receiving a reply of the correct type.

This doesn't solve all possible concurrent programming problems, but you've defined a type that encapsulates the behavior of one kind of concurrent program. If a function describing a `Process` type-checks and is total, you can be confident that it won't deadlock and that all requests will receive replies. If you later refine `Process` further, such as by allowing more-sophisticated descriptions of interactions between processes, you'll be able to implement more-sophisticated models of concurrent programs.

15.3 Summary

- Concurrent programming involves multiple processes executing simultaneously.
- Processes in Idris cooperate with each other by sending messages.
- The `System.Concurrency.Channels` library provides primitive, but unsafe, operations for message passing.
- Primitive operations are unsafe because they provide no guarantees about when processes send and receive messages, or about the correspondence between the types of sent and received messages.
- You can define a type for describing safe message-passing processes, implemented using `Channel` as a primitive.
- Using `Inf`, you can guarantee that looping processes continue to perform IO actions.

- By defining a state machine in the type, you can be sure that a process will respond to messages on every iteration of a loop.
- The type of a process can be generic and describe the types of messages to which a process will respond.

appendix A

Installing Idris and editor modes

This appendix explains how to install Idris. There are prebuilt binary distributions for Mac and Windows, or you can install from source on any Unix-like operating system. It also describes how to install the Idris mode in the Atom text editor.

The Idris compiler and environment

At the time of writing, Idris is available as a binary distribution for Windows and Mac OS from <http://idris-lang.org/download>. This includes the compiler and REPL, along with the Prelude, base libraries, and a variety of contributed libraries that support various data structures, networking, and programming with side effects.

In the following sections, I'll describe how to install Idris either as binaries or from source. In either case, to check that Idris has installed successfully, you can run `idris --version`, which reports which Idris version is installed:

```
$ idris --version
1.0
```

Mac OS

The latest version of Idris is always available as a binary package, downloadable at <http://idris-lang.org/pkgs/idris-current.pkg>. To be able to compile it and run your programs, you'll also need to install Xcode, which is available from the App Store.

To install Idris, therefore, you can follow these steps:

- 1 Download idris-current.pkg.
- 2 Ensure that Xcode is installed from the App Store.
- 3 Open idris-current.pkg and follow the instructions. This will install the Idris binary to /usr/local/bin/idris.

You can also install the latest version via Homebrew (<http://brew.sh/>):

```
$ brew install idris
```

Windows

Prebuilt binaries of Idris are available from <https://github.com/idris-lang/Idris-dev/wiki/Windows-Binaries>. These binaries include everything necessary for compiling and running Idris programs.

Unix-like platforms, installing from source

Idris is implemented in Haskell, and a source distribution is available from the Haskell package manager, Hackage (<http://hackage.haskell.org/>). This is a good option if you’re using an operating system for which there are no binary releases available, or if you need more control over how Idris is set up.

To install it, follow these steps:

- 1 Install the Haskell Platform, available from www.haskell.org/downloads.
- 2 The Haskell Platform includes a command-line tool, cabal, for installing applications and libraries from the central Hackage repository. Assuming you’re in a Unix-like environment, you can install Idris to an executable at /usr/local/bin/idris with the following commands:

```
$ cabal update
$ cabal install idris --program-prefix=/usr/local
```

Editor modes

In this section, I’ll describe how to install the Idris mode for type-driven interactive editing in the Atom text editor, used throughout this book. I’ll also give you some pointers to editor modes for Emacs and Vim.

Atom

Atom is available for all major platforms from <http://atom.io/>. You can install the extension for editing Idris programs as follows:

- 1 Go to Preferences.
- 2 Go to the Install tab.
- 3 Search for the language-idris package.
- 4 There should be one search result. Click the Install button for this result.

- 5 You'll need to tell the `language-idris` package where the Idris binary is installed. To do this, first go to the Packages tab.
- 6 Click on Settings in the `language-idris` box.
- 7 Under the Settings heading, enter the path where the Idris binary is installed. On this page, you can also change the keyboard shortcuts and various other features.

Other editors

At the time of writing, editor modes are also available for Vim and Emacs:

- Vim extension—<https://github.com/idris-hackers/idris-vim>
- Emacs mode—<https://github.com/idris-hackers/idris-mode>

In each case, installation, configuration, and usage instructions are provided.

appendix B

Interactive editing commands

Throughout the book, I describe the interactive construction of Idris programs via editing commands in Atom. The following table summarizes these commands for easy reference.

Table 1 Interactive editing commands in Atom

Shortcut	Command	Description
Ctrl-Alt-A	Add definition	Adds a skeleton definition for the name under the cursor
Ctrl-Alt-C	Case split	Splits a definition into pattern-matching clauses for the name under the cursor
Ctrl-Alt-D	Documentation	Displays documentation for the name under the cursor
Ctrl-Alt-L	Lift hole	Lifts a hole to the top level as a new function declaration
Ctrl-Alt-M	Match	Replaces a hole with a case expression that matches on an intermediate result
Ctrl-Alt-R	Reload	Reloads and type-checks the current buffer
Ctrl-Alt-S	Search	Searches for an expression that satisfies the type of the hole name under the cursor
Ctrl-Alt-T	Type-check name	Displays the type of the name under the cursor
Ctrl-Alt-W	with block insertion	Adds a with block after the current line, containing a new pattern-matching clause with an extra argument

appendix C

REPL commands

The Idris read-eval-print loop (REPL) provides several commands for evaluating and inspecting expressions and types, compiling programs, and searching documentation, among other things. I introduce several of these throughout the book; the following table lists the most commonly used commands, but there are several others available. For more details, type `:?` at the REPL.

Table 1 Idris REPL commands

Command	Arguments	Description
<code><expression></code>	None	Displays the result of evaluating the expression. The variable <code>it</code> contains the result of the most recent evaluation.
<code>:t</code>	<code><expression></code>	Displays the type of the expression.
<code>:total</code>	<code><name></code>	Displays whether the function with the given name is total.
<code>:doc</code>	<code><name></code>	Displays documentation for <code>name</code> .
<code>:let</code>	<code><definition></code>	Adds a new definition.
<code>:exec</code>	<code><expression></code>	Compiles and executes the expression. If none is given, compiles and executes <code>main</code> .
<code>:c</code>	<code><output file></code>	Compiles to an executable with the entry point <code>main</code> .
<code>:r</code>	None	Reloads the current module.
<code>:l</code>	<code><filename></code>	Loads a new file.
<code>:module</code>	<code><module name></code>	Imports an extra module for use at the REPL.

Table 1 Idris REPL commands (*continued*)

Command	Arguments	Description
:printdef	<name>	Displays the definition of name.
:apropos	<word>	Searches function names, types, and documentation for the given word.
:search	<type>	Searches for functions with the given type.
:browse	<namespace>	Displays the names and types defined in the given namespace.
:q	None	Exits the REPL.

appendix D

Further reading

This appendix lists several other resources through which you can learn more about functional programming, types, and Idris' theoretical foundations. The resources are grouped by topic, with a brief comment on each.

Functional programming in Haskell

Idris is heavily inspired by Haskell, including its syntax, language features, and many of its standard libraries. In particular, Idris interfaces are closely related to Haskell type classes. If you'd like to learn about Haskell in more depth, you can take a look at the following books:

- *Learn Haskell* by Will Kurt (Manning, 2017) covers Haskell and functional programming concepts, including several practical examples, with particular emphasis on Haskell's type system.
- *Programming in Haskell, 2nd ed.* by Graham Hutton (Cambridge University Press, 2016) introduces the core features of Haskell and pure functional programming, and it covers many of the more advanced features of Haskell.
- *Thinking Functionally with Haskell* by Richard Bird (Cambridge University Press, 2014) teaches programming from first principles using Haskell, particularly emphasizing techniques for reasoning mathematically about programs.

Other languages and tools with expressive type systems

Several other languages have arisen as a result of academic research into using types to reason about program correctness. These are some examples:

- *Agda* (<http://wiki.portal.chalmers.se/agda/pmwiki.php>)—Supports type-driven development using dependent types in the same way as Idris, but with a stronger emphasis on theorem proving.
- *F** (www.fstar-lang.org/)—A functional programming language that aims to support program verification using refinement types, which are types augmented with a predicate that describes properties of values in that type.
- *Coq* (<https://coq.inria.fr/>)—A proof management system based on dependent types with support for extracting functional programs from proofs.

Theoretical foundations

Idris is based on decades of research in dependent type theory. You can learn more about the theoretical foundations from the following sources:

- *Type Theory and Functional Programming* by Simon Thompson (Addison Wesley, 1991). This book is now out of print, but it's available online from the author's website (www.cs.kent.ac.uk/people/staff/sjt/TTFP/). It provides an accessible introduction to type theory, including its foundations and applications.
- *Software Foundations* by Benjamin Pierce et al. (2016; www.cis.upenn.edu/~bcpierce/sf/current/index.html). This is a course on the mathematical underpinnings of software, from first principles, using Coq. It's instructive to reproduce the examples in Idris!
- *Types and Programming Languages* by Benjamin Pierce (MIT Press, 2002). This textbook provides a comprehensive introduction to type systems and the basic theory of programming languages.
- “Propositions as Types” by Philip Wadler (Communications of the ACM, December 2015; <http://cacm.acm.org/magazines/2015/12/194626-propositions-as-types>). This paper gives a deep but accessible account of the relationship between logic and programming languages, and in particular the idea that a type in a programming language corresponds to a *logical proposition* and that a program corresponds to a *proof* of that proposition.

Total functional programming

Throughout this book, I talk about the value of writing *total* functions, and in chapter 10 I introduce *views*, which provide one way of writing total functions. There are several other techniques, some of which are described in the following papers:

- “Elementary Strong Functional Programming” by David Turner (*Lecture Notes in Computer Science* 1022 (1995):1–13). In this paper, David Turner argues in favor of writing total functions, describing many of the benefits and giving some basic techniques for writing total functions.
- “Why Dependent Types Matter” by Thorsten Altenkirch, Conor McBride and James McKinna (2004). This paper, available online (www.cs.nott.ac.uk/~psztxa/publ/ydttm.pdf), describes an earlier dependently typed language, Epigram, emphasizing the importance of totality in dependently typed programming.

- “Modelling general recursion in type theory” by Ana Bove and Venanzio Capretta (*Mathematical Structures in Computer Science* 15 (2005):671–708). Here, the authors describe a way to formalize recursive programs by defining a data type to describe the computational structure of the program.

Types for concurrency

Chapter 15 presents a way of implementing type-safe concurrent programs, but it uses a very simplified form of concurrent program. You can learn more about types for concurrent programming by reading about *session types*, first described by Kohei Honda in 1993. There are many sources, but the following papers provide a starting point:

- “Types for Dyadic Interaction” by Kohei Honda (*Lecture Notes in Computer Science* 715 (1993):509–523). This paper introduced the concept of session types, describing a type system for encoding interactions between two processes.
- “Multiparty Asynchronous Session Types” by Kohei Honda, Nobuko Yoshida, and Marco Carbone (*Principles of Programming Languages*, 2008). This paper extends the idea of session types to communicating systems with more than two participants.

index

Symbols

_ (underscore) 66, 83, 160
:: operator 88, 104, 164
:printdef command 221
.r command 166
.function 197
/= method 186, 189
%default total annotation 312–313
+ operator 225
++ operator 164, 260
= type 218–219
== operator 183, 185, 209
>>= operator 127–129
\$ operator 285, 336
\$= operator 346

A

absurd function 243
accumulating parameter 201
Add command 168
addCorrect 346
addDownvote function 351
adder function 157, 406–407
AdderType function 155–156
addition function 155–157
addPositives 339
addToData function 166
addToStore function 166, 280, 283
addUpvote function 351
addWrong function 346
allLengths function, refining type of 65–69
anonymous functions 38–39

Append operation 427
AppendVec.idr file 225
Applicative interface
defining for State 335–340
generic do notation using 205–206
arguments
auto-implicit 244–245
defining functions with variable numbers of 155–161
addition function 155–157
printf function 157
arithInputs 303, 313
arithmetic notation 197
arithmetic quiz example
infinite lists 301–304
infinite processes 311–313
state 340–351
executing quiz 348–350
implementing quiz 346–348
nested records 343–344
updating record field values 344–346
updating record fields by applying functions 346
ATM example 382–390
defining states 383–384
defining type for 384–387
refining preconditions using auto-implicits 388–389
simulating at console 387–388
ATMCmd type 387–388

Atom 56–64, 436–437
commands 57
data types and patterns 61–64
defining functions by pattern matching 57–61
auto keyword 244
auto-implicits
overview 245
refining preconditions using 388–389

B

Bind constructor 320
Bool type 209
Booleans 30
bound implicit arguments 83–84
bracketing 336
built-in types 26, 185

C

cabal tool, Haskell 436
canonical constructors 216
case block 177
case expressions, using in types 153–154
case splitting 188, 274, 428
case statement 210
cast function 28, 198
Cast interface, converting between types with 198–199
catchall case 190

Channels library 407–410
 character literals 29–30
 checkEqNat function 216–217,
 229, 231, 245
 Closed state 353
 Command type 340
 comments 47–48
 communication pattern 410
 composite types 40–45
 lists
 functions with 43–45
 overview 41–42
 tuples 40–41
 concurrent programming 9–10
 connect 407–408
 Connected state 353
 ConsoleIO programs 317–318
 constants 154
 constrained generic types
 comparisons with Eq and
 Ord 183–194
 constrained implementations 189–191
 default method definitions 189
 defining Eq constraint 185–189
 Ord 191–193
 testing for equality with Eq 183–185
 interfaces defined in Prelude 194–199
 Cast 198–199
 defining numeric types 195–198
 Show 194–195
 interfaces parameterized by Type -> Type 199–206
 Foldable 201–204
 Functor 200–201
 Monad and Applicative 205–206
 constrained types 35–36
 control flow 132–138
 pattern-matching bindings 134–136
 producing pure values in interactive definitions 132–134
 writing interactive definitions with loops 136–138
 Control.Monad.State 325, 333
 correct function 345
 CountFile command 430–432
 countFrom function 294, 296

covering function 262, 277
 cycle function 301

D

data abstraction 280–287
 data store 282–284
 Traversing store's contents with views 284–287
 Data.List.Views module 277–278
 Data.Primitives.Views 311, 348
 DataStore type
 refining 162–164
 using records for 164–165
 deadlock 404
 Dec type 229–233
 decEq function 233–234
 decidability 245–249
 Dec 229–233
 DecEq 233–234
 decreasing argument 266
 Delay function 295, 309
 dependent state machines
 describing rules in types 390–402
 defining abstract game state and operations 391–392
 defining concrete game state 397–399
 defining type for game state 392–395
 implementing game 395–397
 running game 399–402
 errors in state transitions 374–382
 security properties in types 382–390
 defining states 383–384
 defining type 384–387
 refining preconditions using auto-implicits 388–389
 simulating at console 387–388
 dependent types 11–12
 defining 102–110
 defining vectors 104–107
 indexing vectors with bounded numbers 107–109
 vehicle classification example 102–103
 reading and validating 138–146
 dependent pairs 141–143
 reading vectors from console 139–140
 reading vectors of unknown length 140–141
 validating vector lengths 143–145
 describeHelper function 262, 264
 describeListEnd 260–261, 263
 DivBy 303
 Do constructor 306
 do keyword 179
 do notation 129–131, 311, 334
 extending for infinite processes 311
 sequencing commands with 320–322
 sequencing expressions with Maybe type 177–180
 using Monad and Applicative 205–206
 doAdd 366
 doCount function 429
 door operations example
 interactive development of sequences of operations 356–358
 modeling door as a type 354–356
 representing failure 375–378
 verified, error-checking, door-protocol description 378–382
 DoorClosed state 353, 358, 375
 DoorCmd type 355, 374, 378
 DoorOpen state 354, 358, 375
 doorProg 356, 379
 DoorResult type 375
 DoorState 374

E

Effects library 339
 Ele� predicate 237–250
 auto-implicit arguments 244–245
 deciding membership of vectors 245–249
 guaranteeing a value is in vectors 239–241
 removing elements from Vect 238–239, 241–244

Emacs mode 437
 empty function 283
 empty type 212, 228–229
 enumerated types, defining 89–90
E
 Eq constraint 182
 Eq interface 211
 constrained implementations 189–191
 default method definitions 189
 defining using interfaces and implementations 185–189
 testing for equality 183–185
 Eq ty constant 186
 EqNat type 211, 239
 equal expressions 199
 equality types
 decidability
 Dec 229–233
 DecEq 233–234
 empty type 228–229
 guaranteeing equivalence of data with 209–219
 = type 218–219
 exactLength 210–211, 216–218
 expressing equality of Nats as a type 211–212
 manipulating equalities 215–216
 testing for equality of Nats 212–214
 reasoning about equality 220–227
 appending vectors 225–227
 delegating proofs and rewriting to holes 224–225
 reversing vectors 220–221
 rewrite construct 223–224
 type checking and evaluation 221–222
 equalSuffix function 279
 evalState 330–331
 exactLength function 209–211, 216–218, 233
 ExactLength.idr file 210
 execState 330
 export modifier 281
 Expr type 196

F

Face type 304
 False value 211
 filterKeys 286–287
 Fin argument, indexing with bounded numbers 107–109
 Finished type 391
 finite prefix 415
 first-class types 22–24
 defining functions with variable numbers of arguments 155–161
 addition function 155–157
 printf function 157
 schemas 161–180
 DataStore type 162–165
 displaying entries in store 170–171
 holes 165–170
 parsing entries according to 171–175
 sequencing expressions with Maybe using do notation 177–180
 updating 175–177
 type-level functions 148–154
 type synonyms 149–150
 using case expressions in types 153–154
 with pattern matching 150–153
 Foldable interface, reducing structure using 201–204
 Force function 295, 309
 forever function 310, 312, 417
 format string 148
 Format type 159
 Fractional interface 195
 fromInteger method 196
 Fuel type 308, 341
 function definitions
 local 39–40
 overview 31–33
 functions 30–40, 75–81
 anonymous 38–39
 defining by pattern matching 57–61
 defining with variable numbers of arguments 155–161
 addition function 155–157
 Formatted output, a type-safe printf function 157

higher-order 36–38
 local definitions 39–40
 partial 16–17
 partially applying 33
 total 16–17, 296–297, 308–309
 type-level 148–154
 type synonyms 149–150
 using case expressions in types 153–154
 with pattern matching 150–153
 types of 31
 constrained 35–36
 overview 31–33
 variables in 33–34
 using implicit arguments in 84–85
 writing generic 33–36
F
 Functor interface 200, 337
 applying functions across structure with 200–201
 defining for State 335–340

G

GameLoop type 251, 391, 399–402
 GameState 342–343
 generic types, defining 95–100
 get command 180
 GetData command 430–431
 getEntry command 167–168
 GetLine command 318, 341
 getPrefix function 298–299, 308
 GetRandom 348
 GetStr command 368
 getStringOrInt 151
 getValues function 287
 global state 325
 greet function 315
 guess function 256
 guessing game example 390–402
 defining abstract game state and operations 391–392
 defining concrete game state 397–399
 defining type for game state 392–395
 implementing 395–397
 running 399–402

H

Hangman guessing game
example 250
completing top-level game
implementation 255
deciding input validity 255
predicate for validating user
input 251–252
processing guesses 253–254
representing game's
state 250–251
top-level game function 251
Haskell 436
headUnequal 234
helper function 328
Here value 240
hierarchical namespaces 321
higher-order functions 36–38
holes 20–21
correcting compilation
errors using 165–170
delegating proofs and rewriting
to 224–225
Homebrew 436

I

Idris programming language
17–24
as pure functional program-
ming language 13–17
partial and total functions
16–17
purity and referential
transparency 13–14
side-effecting programs
14–15
checking types 18–19
comments 47–48
compiling and running
programs 19–20
composite types 40–45
functions with lists 43–45
lists 41–42
tuples 40–41
editor modes 436–437
Atom 436–437
other editors 437
first-class types 22–24
functions 30–40
anonymous 38–39
higher-order 36–38
local definitions 39–40
partially applying 33

types and definitions of
31–33
writing generic with con-
strained types 35–36
writing generic with vari-
ables in types 33–34
holes 20–21
installing compiler and
environment 435–436
Mac OS 435–436
Unix-like platforms 436
Windows 436
interactivity 48–50
REPL (read-eval-print
loop) 17–18
types 26–30
Booleans 30
characters and strings
29–30
converting 28
numeric 27–28
whitespace 46
implicit arguments 82–85
bound and unbound 83–84
need for 82–83
using in functions 84–85
impossible keyword 229
Inf ty type 298
Inf type 295–296, 308
guaranteeing responses
using state machine
and 418–422
making processes total using
415–418
infinite lists 292–305
arithmetic quiz
example 301–304
labeling elements in 293–294
processing 297–299
producing 295–296
Stream data type 299–301
total functions 296–297
infinite processes 305–314
arithmetic quiz example
311–313
describing 306–307
executing 307–308
extending do notation for
311
generating infinite struc-
tures using Lazy
types 309–310
InflIO type 305
infix operator 163, 336
InflList 295
input and output processing
control flow 132–138
pattern-matching bindings
134–136
producing pure values in
interactive definitions
132–134
writing interactive defini-
tions with loops
136–138
IO generic type 124–131
`>=` operator 127–129
do notation 129–131
evaluating and executing
interactive
programs 125–127
reading and validating
dependent types
138–146
dependent pairs 141–143
reading vectors from
console 139–140
reading vectors of
unknown length
140–141
validating vector lengths
143–145
input function 272
Integral interface 195
interactive data store 110–121
interactively maintaining
state in main 113–115
parsing user input 115–117
processing commands
118–120
representing store 112–113
interactive operations 329
interactive programming
control flow 132–138
reading and validating
dependent types
138–146
with IO generic type
124–131
interface declarations 186
interfaces
defined in Prelude 194–199
Cast 198–199
defining numeric types
195–198
Show 194–195
generic comparisons
183–194
constrained implemen-
tations 189–191

interfaces, generic comparisons
(continued)
 default method definitions
 189
 defining Eq using interfaces and implementations
 185–189
 Ord 191–193
 testing for equality with
 Eq 183–185
 parameterized by Type ->
 Type 199–206
 Foldable 201–204
 Functor 200–201
 Monad and Applicative
 205–206
 IO actions 301, 305, 397
 IO Finished 251
 IO generic type 124–131
 >>= operator 127–129
 do notation 129–131
 evaluating and executing
 interactive programs
 125–127
 isInt 153
 isList function 250
 isSuffix function 276–277
 isValidString 252, 255
 it variable 439
 items function 164
 iterate function 299

J

join method 205

L

labelFrom function 293–294
 labelWith function 294, 300
 language-idris package 437
 Lazy annotation 277
 Lazy types, generating infinite
 structures using
 309–310
 Length operation 427
 let construct 39–40
 linear congruential generator
 302
 List Char 160
 List elem 198
 listen 407–408
 ListLast 261, 272
 lists 41–42
 functions with 43–45

matching last item in
 260–261
 reversing 264–266
 snoc 271–274
 ListType function 427
 local mutable state 329
 local state 325
 logOpen function 382
 Loop constructor 416
 loopPrint function 306–307,
 310
 loops, writing interactive definitions with
 136–138

M

MachineCmd 360, 362
 main function
 interactively maintaining
 state in main 113–115
 overview 318, 406
 map function 200
 matrices
 arithmetic involving 6–7
 functions 75–81
 operations and types 76–77
 transposing 77–81
 Matter type 185
 Maybe elem 198
 Maybe type, using do notation
 177–180
 maybeAdd function 178–179
 membership tests 237–250
 auto-implicit arguments
 244–245
 deciding membership of
 vectors 245–249
 guaranteeing a value is in
 vectors 239–241
 removing elements from
 Vect 238–239, 241–244
 merge sorting 266–270
 mergeSort function 267, 269,
 278
 MessagePID 412
 method declarations 186
 MkData 165
 MkWordState argument 250
 modules, in Idris 280–282
 Monad interface
 defining for State 335–340
 generic do notation using
 205–206
 MonadState interface 330

mutable state 325–332
 representing using pairs
 328–329
 State type 329–331
 tree-traversal example
 326–327
 tree-traversal with State
 331–332
 mutual block 339
 myReverse function 222–224,
 265
 myReverseHelper function 273

N

namespaces 321
 Nats
 expressing equality of as a
 type 211–212
 overview 156, 227, 413
 testing for equality of
 212–214
 Neg interface 195
 nextState 335
 Nil constructor 295
 nondeterministic programs
 206
 nonterminating component
 292
 noRec function 233
 NoRecv interface 424
 NoRequest state 420
 notInNil 248
 notInTail 248
 NotRunning state 391
 Num interface 195
 numargs 155–156
 numeric types
 and values 27–28
 defining 195–198

O

occurrences function 184
 Ord interface
 defining orderings with
 191–193
 overview 182, 267

P

pairs, representing mutable
 state 328–329
 palindrome function 279

parameterized types 199
 parseBySchema 172
 parseCommand 169, 172
 parsePrefix 172
 partial functions 16–17
 partially application of
 functions 33
 pattern matching
 extending with views
 data abstraction 280–287
 defining and using views
 259–270
 recursive views 271–279
 syntax for extended 262–264
 type-level functions with
 150–153
 pattern-matching bindings
 134–136
 PID (process identifier)
 406–407
 plus function 221
 plusCommutative type 222
 plusSuccRightSucc 227
 plusZeroRightNeutral 227
 Polygon function 149
 Position function 149–150
 predicates
 Elem predicate 237–250
 auto-implicit arguments
 244–245
 deciding membership of
 vectors 245–249
 guaranteeing a value is in
 vectors 239–241
 removing elements from
 Vect 238–239, 241–244
 Hangman guessing game
 example 250
 completing top-level game
 implementation 255
 deciding input validity 255
 predicate for validating
 user input 251–252
 processing guesses
 253–254
 representing game's state
 250–251
 top-level game function
 251
 primitives, for concurrent
 programming 404–411
 Channels library 407–410
 defining concurrent pro-
 cesses 406–407
 type errors and blocking
 410–411

printf function 148, 155, 157
 PrintType 159
 println 195, 344
 problem domain 317
 procAdder 415–416
 process identifier. *See* PID
 processInput 176
 ProcessLib module 427
 procMain 414, 421
 ProcState 426
 public export 281
 Pure constructor 320, 354
 pure function 205
 pure functional programming
 languages 13–17
 partial and total functions
 16–17
 purity 13–14
 referential transparency
 13–14
 side-effecting programs
 14–15
See also Idris programming
 language
 purity 13–14
 PutStr command 318, 341, 368

Q

Quit command 316
 quiz function 301–302, 313,
 319, 340

R

race conditions 404–405
 randoms function 302, 313
 Read operation 405
 readGuess function 252
 readInput 320, 322
 record declaration 165
 recursion 296
 recursive types, defining 92–95
 recursive views 271–279
 Data.List.Views module
 277–278
 nested with blocks 275–277
 snoc lists 271–274
 with construct 274–275
 referential transparency 13–14
 Refl (reflexive) 218
 removeElem function 238–239
 repeat function 299
 REPL (read-eval-print loop)
 17–18, 439

repl function 314
 Respond command 414
 return values 377
 reverse function 265
 rewrite construct 223–224, 235,
 273
 rewriteCons 274
 rewriteNil 274
 RingBell operation 362
 run function 306–307, 310,
 362, 382, 411
 runCommand 342, 399
 RunIO type 315
 Running state 391
 runProc 421
 runState function 330, 333–334,
 341, 351, 366

S

sameS function 214–215, 220
 schemas 161–180
 correcting compilation errors
 using holes 165–170
 DataStore type
 refining 162–164
 using records for 164–165
 displaying entries in store
 170–171
 overview 148
 parsing entries according
 to 171–175
 sequencing expressions with
 Maybe using do
 notation 177–180
 updating 175–177
 SchemaType function 163, 169
 Sent state 421
 setDifficulty 345
 SetSchema command 176
 Shape type 194
 show function 194
 Show interface, converting to
 String with 194–195
 showPrec function 194
 ShowState 394
 side-effecting programs 14–15
 size function 164
 sort function 193
 sorting vectors 70–75
 spawn 407–408
 SplitList 267
 SplitNil 269
 SplitOne 269

SplitRec 278
 SplitRecPair 277
 stack-based calculator example 367–371
 StackCmd 364, 368
 stacks 363–371
 implementing using Vect 366–367
 representing operations 364–365
 stack-based calculator example 367–371
 state
 arithmetic quiz 340–351
 defining nested records 343–344
 executing 348–350
 implementing 346–348
 updating record field values 344–346
 updating record fields by applying functions 346
 custom implementation of 333–340
 defining Functor, Applicative, and Monad interfaces 335–340
 defining State and runState 333
 mutable state 325–332
 representing using pairs 328–329
 State type 329–331
 tree-traversal example 326–327
 tree-traversal with State 331–332
 state machines
 dependent
 describing rules in types 390–402
 errors in state transitions 374–382
 security properties in types 382–390
 guaranteeing responses using Inf and 418–422
 implementing stacks 363–371
 representing operations 364–365
 stack-based calculator example 367–371
 using Vect 366–367

tracking state in types 353–363
 door operations example 354–358
 vending machine example 358, 360–362
 State Nat type 329
 State type 329–331
 defining 333
 tree-traversal with 331–332
 storeView function 280
 StoreView view 284
 Stream data type 299–301
 Stream type 293
 streams, infinite lists 292–305
 arithmetic quiz example 301–304
 labeling elements in 293–294
 processing 297–299
 producing 295–296
 Stream data type 299–301
 total functions 296–297
 string literals 29–30
 StringOrInt 151, 154
 strToInput 370
 sucNotZero function 232
 System module 311
 System.Concurrency.Channels module 404, 407

T

tailUnequal 234
 take function 299
 takeN function 270
 termination 314–323
 domain-specific commands 317–319
 sequencing commands with do notation 320–322
 testAdd 365
 testStore 285
 testTree 326
 There constructor 240
 ThreeEq type 219
 time function 311
 toBinary function 279
 total functions
 executing infinite processes as 308–309
 overview 16–17, 296–297
 totalREPL 314
 traverse function 335
 Tree data type 203
 Tree elem 191
 treeLabel function 327
 treeLabelWith 328–329
 tree-traversal example
 overview 326–327
 with State 331–332
 trim function 152
 True value 211
 Try command 362
 tuples 40–41
 TupleVect type 161
 Type -> Type parameterized interfaces 199–206
 applying functions across structure with Functor 200–201
 generic do notation using Monad and Applicative 205–206
 reducing structure using Foldable 201–204
 type synonyms 148–150
 type-checking 277
 type-driven development 5–13
 automated teller machine example 7–9
 concurrent programming 9–10
 dependent types 11–12
 matrices, arithmetic involving 6–7
 process of 10–11
 types, defined 4–5
 type-level functions 147
 types 26–30
 Booleans 30
 calculating 148–154
 type synonyms 149–150
 type-level functions with pattern matching 150–153
 using case expressions in types 153–154
 characters and strings 29–30
 checking 18–19
 checking and evaluation of 221–222
 composite 40–45
 lists 41–45
 tuples 40–41
 constrained generic comparisons with Eq and Ord 183–194
 interfaces defined in Prelude 194–199

types, constrained generic
(continued)
 interfaces parameterized
 by Type → Type
 199–206
 converting 28
 defined 4–5
 dependent 11–12
 describing rules in 390–402
 defining abstract game state and operations 391–392
 defining concrete game state 397–399
 defining type for game state 392–395
 implementing game 395–397
 running game 399–402
 empty 228–229
 equality
 decidability 229–234
 empty type 228–229
 guaranteeing equivalence of data with 209–219
 reasoning about equality 220–227
 first-class 22–24
 defining functions with variable numbers of arguments 155–161
 schemas 161–180
 type-level functions 148–154
 function 31–33
 constrained 35–36
 variables in 33–34
 in Atom 61–64
 numeric 27–28, 195–198
 security properties in 382–390
 defining states 383–384
 defining type 384–387
 refining preconditions using auto-implicits 388–389
 simulating at console 387–388
 tracking state in 353–363
 door operations example 354–358
 vending machine example 358, 360–362
 user-defined
 defining 88–101

defining dependent 102–110
 interactive data store 110–121
 type-safe concurrent programming
 primitives for 404–411
 Channels library 407–410
 defining concurrent processes 406–407
 type errors and blocking 410–411
 type for safe message passing 411–433
 defining module 426–427
 describing processes 412–415
 generic processes 422–426
 guaranteeing responses using state machine and Inf 418–422
 list processing example 427–429
 making processes total using Inf 415–418
 word-counting process example 429–433

parsing user input 115–117
 processing commands 118–120
 representing store 112–113

V

ValidInput predicate
 deciding input validity 255
 overview 251–252
 valToString 152–153
 Vect size String 162
 Vect type, implementing stacks using 366–367
 vectors 64–75
 appending 225–227
 automatic refining 69–70
 deciding membership of 245–249
 defining 104–107
 guaranteeing a value is in 239–241
 indexing with bounded numbers using Fin 107–109
 reading
 from console 139–140
 of unknown length 140–141
 refining type of allLengths function 65–69
 removing elements from 238–239, 241–244
 reversing 220–221
 sorting 70–75
 validating lengths 143–145
 vending machine example
 modeling vending machine 358
 verified description 360–362
 views
 data abstraction 280–287
 data store 282–284
 modules in Idris 280–282
 traversing store's contents 284–287
 defining and using 259–270
 building views 262
 matching last item in list 260–261
 merge sorting 266–270
 reversing lists 264–266
 with blocks 262–264
 recursive 271–279

U

unbound implicit arguments 83–84
 underscore (_) 66, 83, 160
 Uninhabited interface 243
 union types, defining 90–92
 unsafeReceive 408–410
 unsafeSend 408, 410
 updateGameState function 350
 user-defined data types
 defining 88–101
 enumerated types 89–90
 generic types 95–100
 recursive types 92–95
 union types 90–92
 defining dependent 102–110
 defining vectors 104–107
 indexing vectors with bounded numbers 107–109
 vehicle classification example 102–103
 interactive data store 110–121
 interactively maintaining state in main 113–115

views, recursive (*continued*)
 Data.List.Views module
 277–278
 nested with blocks
 275–277
 snoc lists 271–274
 with construct 274–275
Vim extension 437
VList view 279
Void type
 empty type 228–229
 overview 228, 424

W

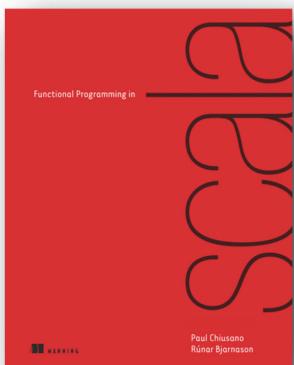
WCData 430
wcService 430–431
when function 335
where construct 39–40
whitespace 46
with blocks
 nested 275–277
overview 263
syntax for extended pattern
 matching 262–264

with construct 260, 274–275
WordState type 250–251, 390
Write operation 405
wrong function 345

Z

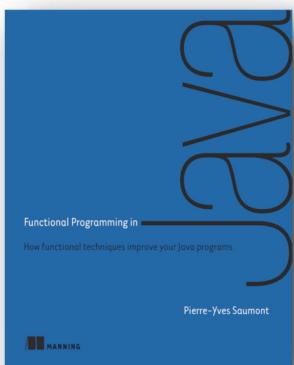
zeroNotSuc function 232
zipInputs 211

MORE TITLES FROM MANNING



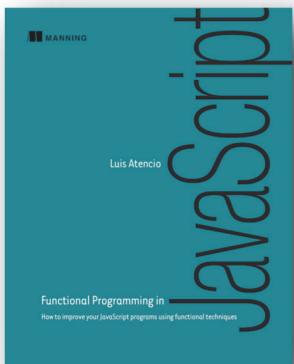
Functional Programming in Scala
by Paul Chiusano and Rúnar Bjarnason

ISBN: 9781617290657
320 pages
\$44.99
September 2014



Functional Programming in Java
How functional techniques improve
your Java programs
by Pierre-Yves Saumont

ISBN: 9781617292736
472 pages
\$49.99
January 2017

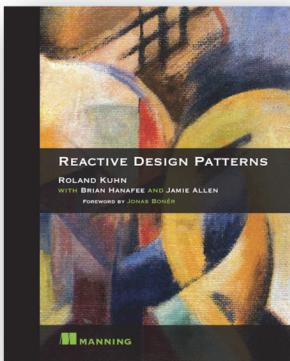


Functional Programming in JavaScript
How to improve your JavaScript programs using
functional techniques
by Luis Atencio

ISBN: 9781617292828
272 pages
\$44.99
June 2016

For ordering information go to www.manning.com

MORE TITLES FROM MANNING



Reactive Design Patterns

by Roland Kuhn

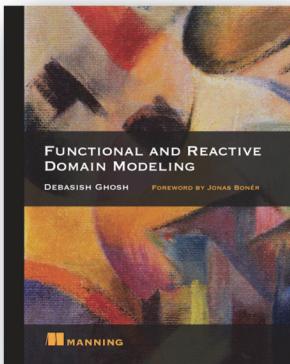
with Brian Hanafee and Jamie Allen

ISBN: 9781617291807

392 pages

\$49.99

February 2017



Functional and Reactive Domain Modeling


by Debasish Ghosh

ISBN: 9781617292248

320 pages

\$59.99

October 2016



The Joy of Clojure, Second Edition

by Michael Fugus and Chris Houser

ISBN: 9781617291418

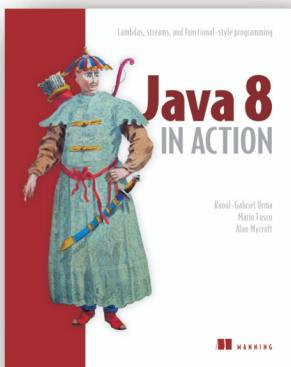
520 pages

\$49.99

May 2014

For ordering information go to www.manning.com

MORE TITLES FROM MANNING



Java 8 in Action

*Lambdas, streams, and functional-style
programming*

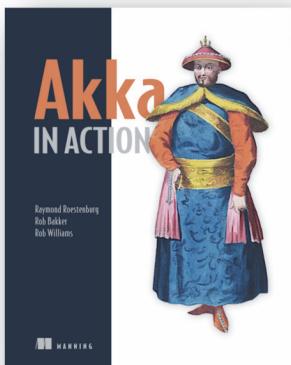
by Raoul-Gabriel Istrate, Mario Fusco,
and Alan Mycroft

ISBN: 9781617291999

424 pages

\$49.99

August 2014



Akka in Action

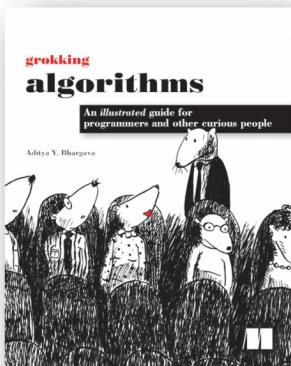
by Raymond Roestenburg, Rob Bakker,
and Rob Williams

ISBN: 9781617291012

448 pages

\$49.99

September 2016



Grokking Algorithms

*An illustrated guide for
programmers and other curious people*

by Aditya Y. Bhargava

ISBN: 9781617292231

256 pages

\$44.99

May 2016

For ordering information go to www.manning.com

Atom commands used for the interactive construction of Idris projects featured in this book are detailed below.

Shortcut	Command	Description
Ctrl-Alt-A	Add definition	Adds a skeleton definition for the name under the cursor
Ctrl-Alt-C	Case split	Splits a definition into pattern-matching clauses for the name under the cursor
Ctrl-Alt-D	Documentation	Displays documentation for the name under the cursor
Ctrl-Alt-L	Lift hole	Lifts a hole to the top level as a new function declaration
Ctrl-Alt-M	Match	Replaces a hole with a case expression that matches on an intermediate result
Ctrl-Alt-R	Reload	Reloads and type-checks the current buffer
Ctrl-Alt-S	Search	Searches for an expression that satisfies the type of the hole name under the cursor
Ctrl-Alt-T	Type-check name	Displays the type of the name under the cursor
Ctrl-Alt-W	with block insertion	Adds a with block after the current line, containing a new pattern-matching clause with an extra argument

Type-Driven Development with Idris

Edwin Brady

Stop fighting type errors! Type-driven development is an approach to coding that embraces types as the foundation of your code—essentially as built-in documentation your compiler can use to check data relationships and other assumptions. With this approach, you can define specifications early in development and write code that's easy to maintain, test, and extend. Idris is a Haskell-like language with first-class, dependent types that's perfect for learning type-driven programming techniques you can apply in any codebase.

Type-Driven Development with Idris teaches you how to improve the performance and accuracy of your code by taking advantage of a state-of-the-art type system. In this book, you'll learn type-driven development of real-world software, as well as how to handle side-effects, interaction, state, and concurrency. By the end, you'll be able to develop robust and verified software in Idris and apply type-driven development methods to other languages.

What's Inside

- Understanding dependent types
- Types as first-class language constructs
- Types as a guide to program construction
- Expressing relationships between data

Written for programmers with knowledge of functional programming concepts.

Edwin Brady leads the design and implementation of the Idris language.

To download their free eBook in PDF, ePUB, and Kindle formats,
owners of this book should visit
manning.com/books/type-driven-development-with-idris



“This book will turn your approach to software upside-down, in the best way.”

—Ian Dees, New Relic

“Highly recommended for anyone developing software with serious safety requirements.”

—Arnaud Bailly, GorillaSpace

“After reading this book, TDD took on a new meaning for me.”

—Giovanni Ruggiero, Eligotech

“A clear and complete view of type-driven development that reveals the power of dependent types.”

—Nicolas Biri
Luxembourg Institute of Science and Technology

ISBN-13: 978-1-61729-302-3
ISBN-10: 1-61729-302-4

5 4 9 9 9



9 7 8 1 6 1 7 2 9 3 0 2 3